

API-231 / GIS-PubPol

Meeting 05 (Lab Exercise + Problem Set 3)

Yuri M. Zhukov
Visiting Associate Professor of Public Policy
Harvard Kennedy School

February 8, 2024

Analysis and Geoprocessing with NYC OpenData



Figure 1: A lot of cool stuff here

Overview of lab exercise and problem set

1. Lab exercise
 - a) Map of **bicycle crashes** (per capita) in each NYC community district
 - b) Map of **bike lanes** per district (normalized for population)
2. Problem set
 - a) Map of **rat activity** (per capita) in each NYC community district

We will first make **this map**

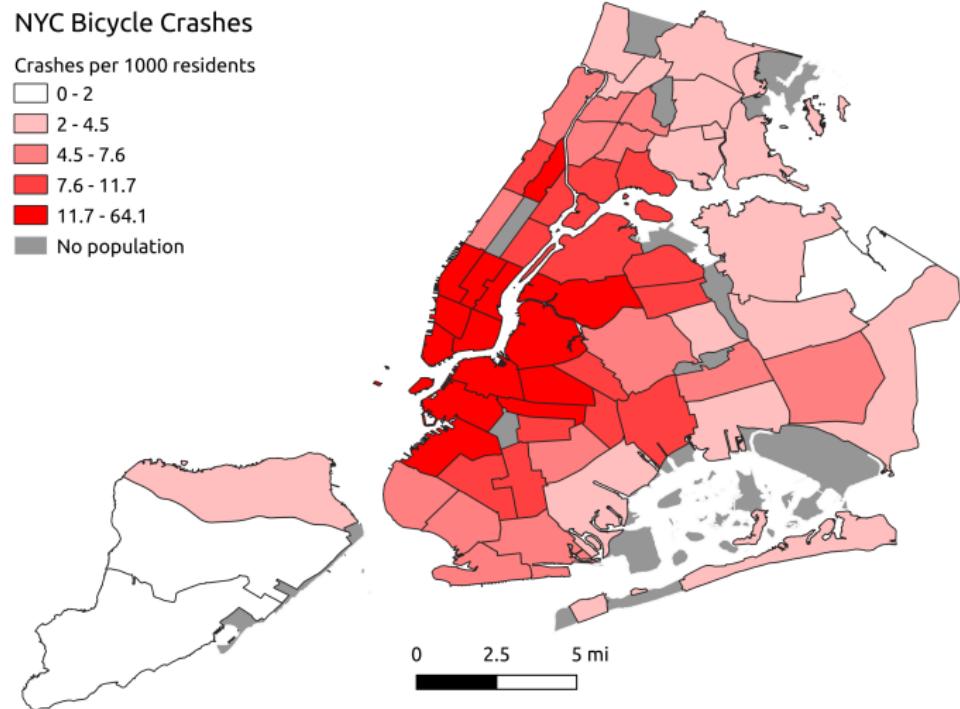


Figure 2: Bike crashes per capita in NYC

This will require joining community districts polygons ...

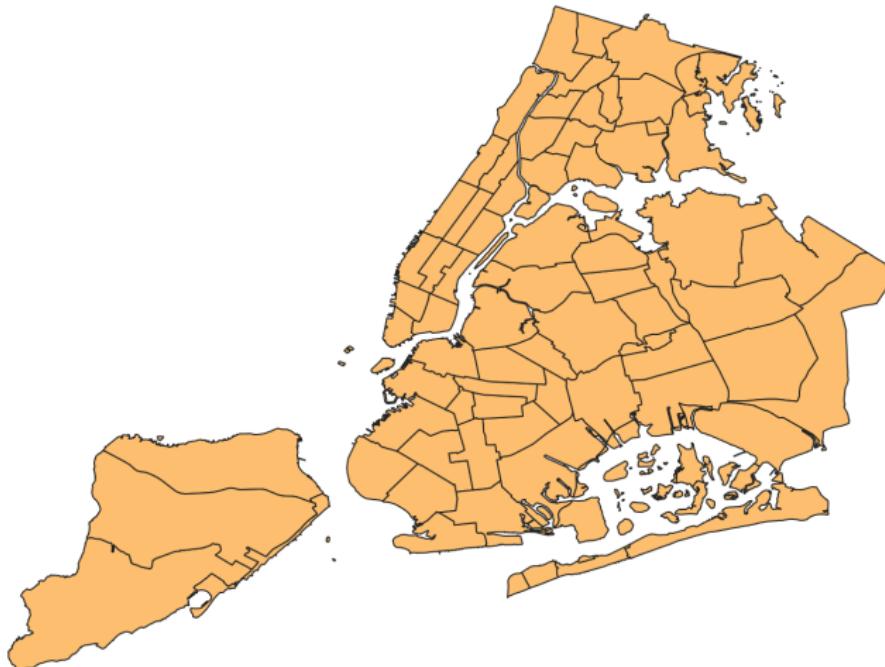


Figure 3: Community districts in NYC

... to a (non-geographic) spreadsheet of population statistics

The screenshot shows a LibreOffice Calc spreadsheet titled "NYC_Pop.csv". The data consists of 27 rows of population statistics for New York City neighborhoods. The columns are labeled A through I, and the rows are numbered 1 through 27. The data includes columns for boro_code, Borough, cd_code, cd_name, and various population counts from 1970 to 2010.

A	B	C	D	E	F	G	H	I
boro_code	Borough	cd_code	cd_name	pop_1970	pop_1980	pop_1990	pop_2000	pop_2010
1	2:Bronx	1:Melrose	Mott Haven, Port Morris	138557	78441	77214	82159	91497
2	2:Bronx	2:Hunts Point, Longwood		99493	34399	39443	46824	52246
3	2:Bronx	3:Morrisania, Crotona Park East		150636	53635	57162	68574	79762
4	2:Bronx	4:Hibridge, Concourse Village		144207	114312	119962	139563	146441
5	2:Bronx	5:University Hts., Fordham, Mt. Hope		121807	107995	118435	128313	128200
6	2:Bronx	6:East Tremont, Belmont		114137	65016	68061	75688	83268
7	2:Bronx	7:Bedford Park, Norwood, Fordham		113764	116827	128588	141411	139286
8	2:Bronx	8:Riverdale, Kingsbridge, Marble Hill		103543	98275	97030	101332	101731
9	2:Bronx	9:Soundview, Parkchester		166442	167627	155970	167859	172298
10	2:Bronx	10:Troy Hills Nc., Co-op City, Pelham Bay		84948	106516	108093	115948	120392
11	2:Bronx	11:Pelham Parkway, Morris Park, Laconia		105980	99088	97842	110706	113232
12	2:Bronx	12:Wakefield, Williamsbridge		135010	128226	129620	149077	152344
13	3:Brooklyn	1:Williamsburg, Greenpoint		178990	142942	155972	160238	173083
14	3:Brooklyn	2:Brooklyn Heights, Fort Greene		110221	92732	94534	98620	99617
15	3:Brooklyn	3:Bedford Stuyvesant		203380	133379	138896	143867	152985
16	3:Brooklyn	4:Bushwick		137902	92497	102572	104358	112634
17	3:Brooklyn	5:East New York, Starrett City		170791	154931	161350	173198	182896
18	3:Brooklyn	6:Park Slope, Carroll Gardens		138933	110228	102724	104054	104709
19	3:Brooklyn	7:Sunset Park, Windsor Terrace		111607	98567	102553	120063	126230
20	3:Brooklyn	8:Crown Heights North		121821	88796	96400	96076	96317
21	3:Brooklyn	9:Crown Heights South, Wrigate		101047	96669	110715	104014	98429
22	3:Brooklyn	10:Bay Ridge, Dyker Heights		129822	118187	110612	122542	124491
23	3:Brooklyn	11:Bensonhurst, Bath Beach		170119	155072	149994	172129	181981
24	3:Brooklyn	12:Borough Park, Ocean Parkway		166301	155899	160018	185046	191382
25	3:Brooklyn	13:Coney Island, Brighton Beach		97750	100030	102596	106120	104278
26	3:Brooklyn	14:Flatbush, Midwood		137041	143859	159825	168806	160664

Figure 4: NYC population data

... followed by a point-in-polygon analysis and some field calculations

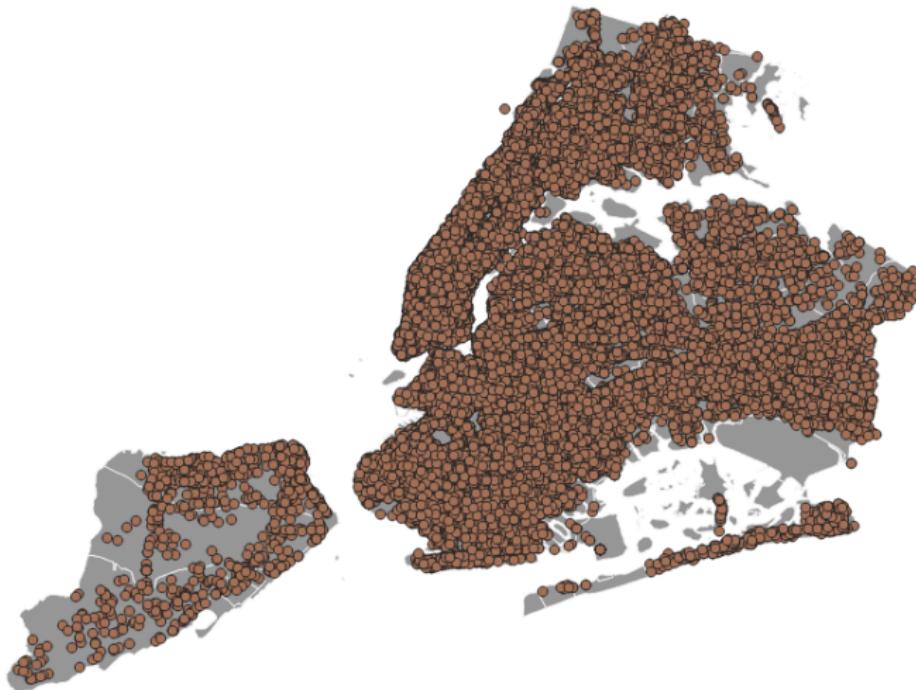


Figure 5: NYC bike crash data

We will then make **this map**

NYC Bicycle Lanes

Miles per 1000 residents

- 0 - 0.066
- 0.066 - 0.19
- 0.19 - 0.283
- 0.283 - 0.432
- 0.432 - 0.899
- No population

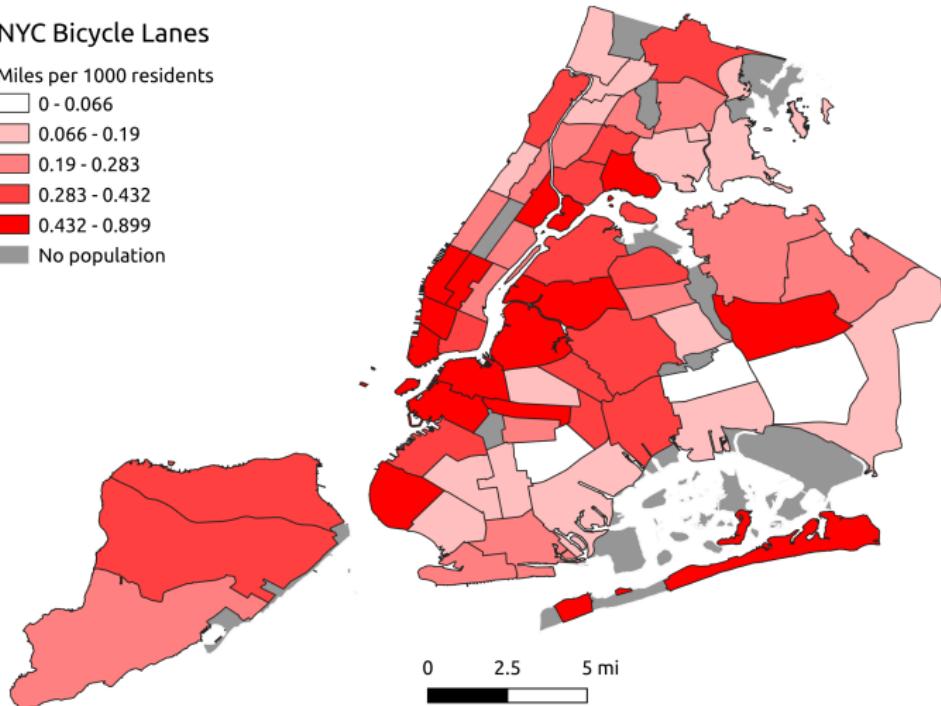


Figure 6: Bike lanes per capita in NYC

This will require some line-in-polygon analysis



Figure 7: Bike lanes in NYC

You will make this map for your **problem set**

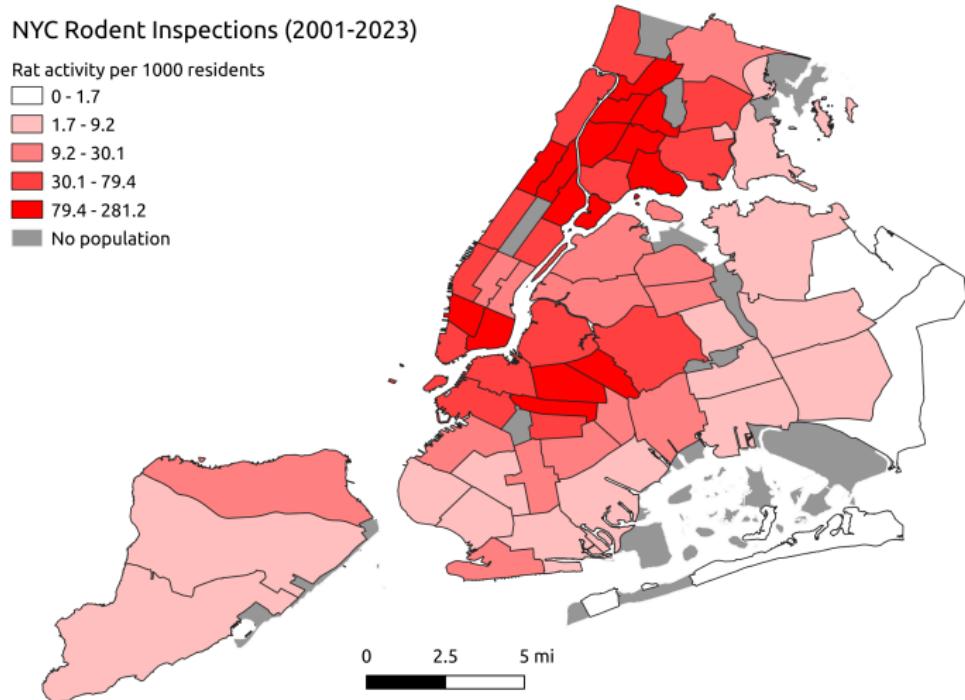


Figure 8: Rat activity per capita in NYC

This will involve (again) point-in-polygon analysis

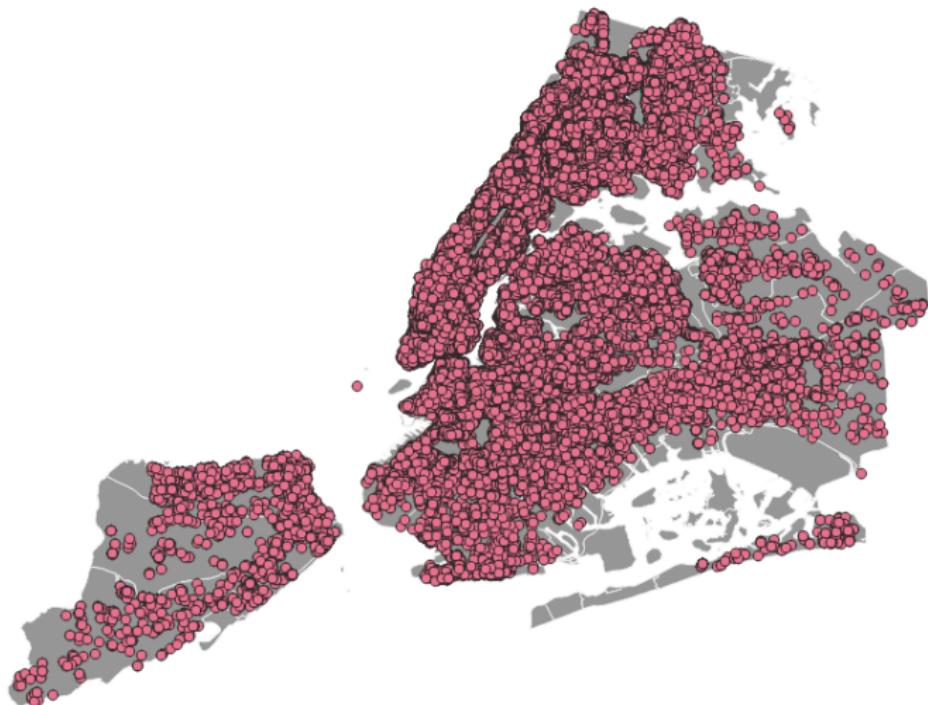


Figure 9: Rodent inspections in NYC

You can make these plots in QGIS or in R. Instructions for both are below.

Bicycle Crashes per 1000 Residents

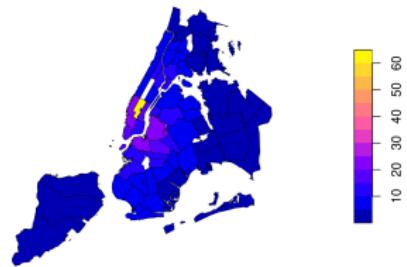


Figure 10: Map 1 in R

Miles of Bicycle Lane per 1000 Residents

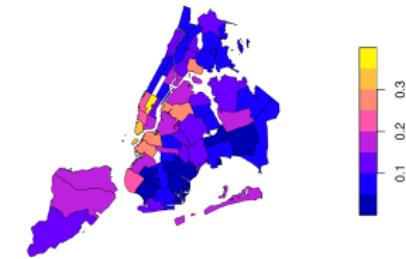


Figure 11: Map 2 in R

Rat Activity per 1000 Residents

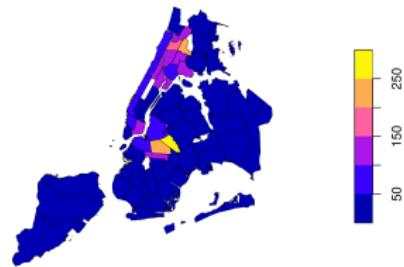


Figure 12: PS03 in R

We will put multiple types of data on same map:

Category	Type	Format	Data source
Community district borders	Vector (polygon)	.geojson	NYC OpenData
Community district population	Table	.csv	NYC OpenData
Vehicle collisions involving bicycles	Vector (point)	.csv	NYC OpenData
Bike routes	Vector (polyline)	.geojson	NYC OpenData
Rodent inspections finding "rat activity"	Vector (point)	.csv	NYC OpenData

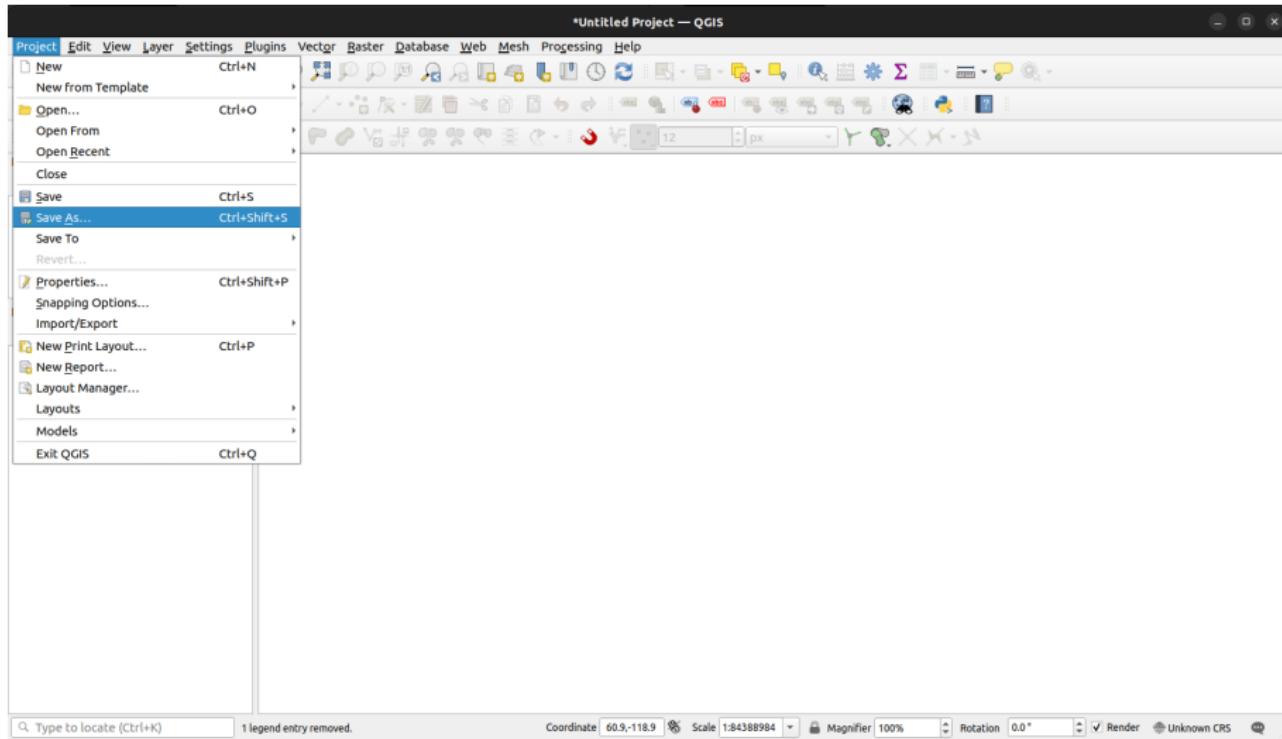
These are all in the PS03.zip file posted on Canvas.

Let's open QGIS...

QGIS

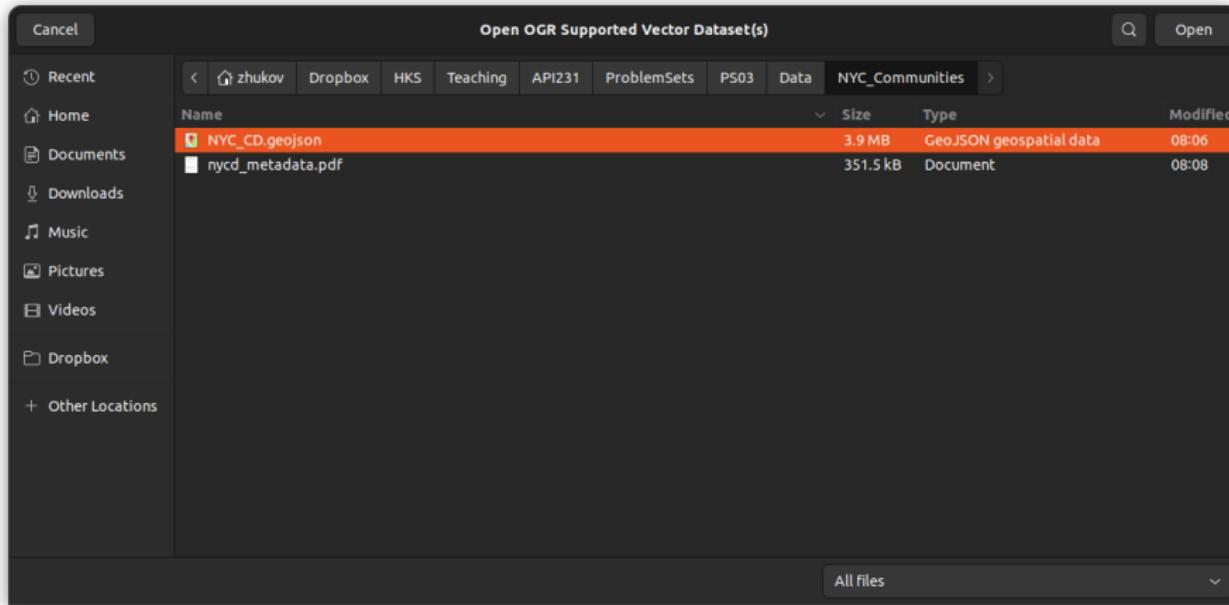
Always save your progress!

Go to Project → Save As...

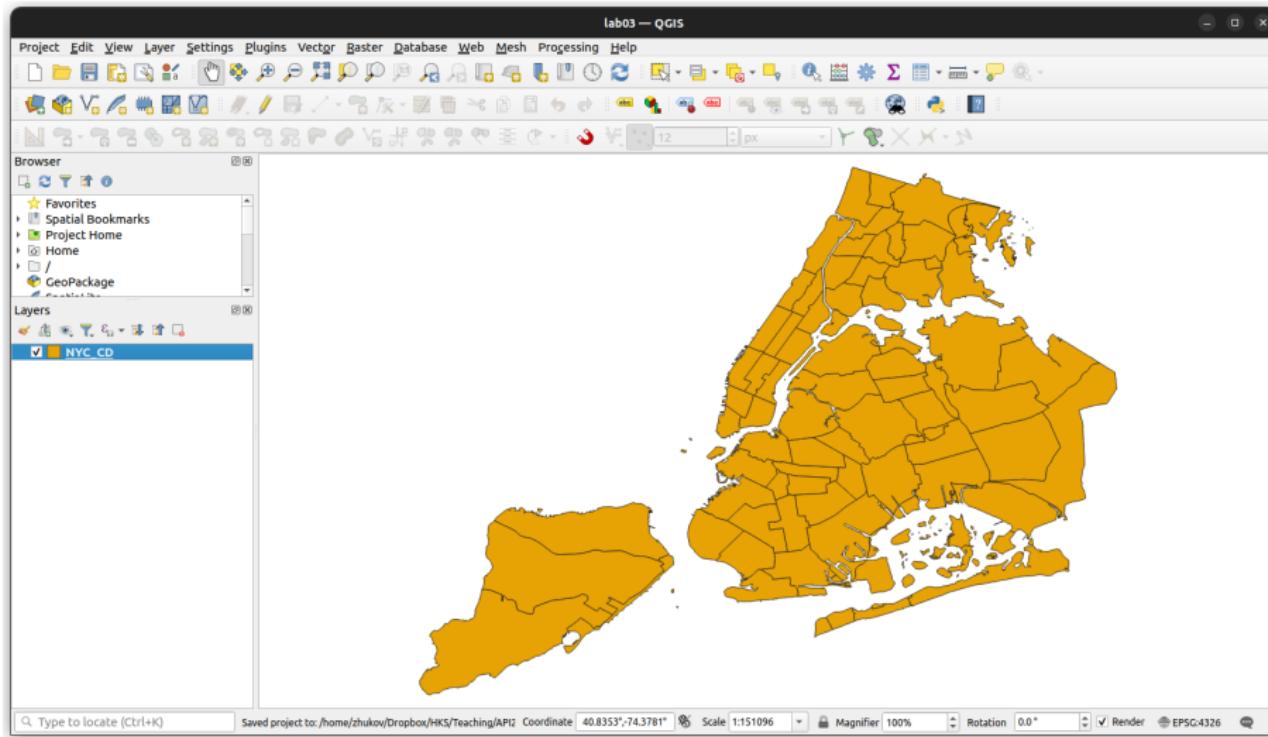


Map 1

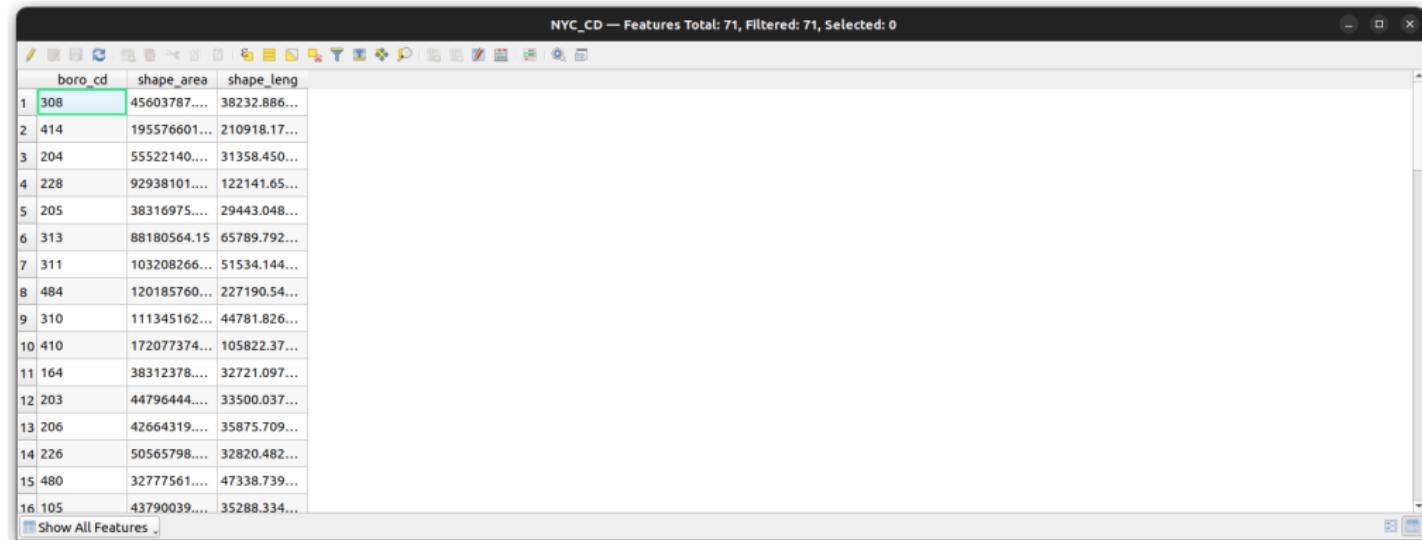
Load community districts: Layer → Add Layer → Add Vector Layer...
Open the file **NYC_CD.geojson** in the **NYC_Communities** folder. Click Add



This dataset includes 71 polygons, including New York's 59 community districts and 12 areas outside of community districts' jurisdiction (parks, waterfront)



Open the Attribute Table for this layer. We see a borough-district code (boro_cd) and geometric statistics (shape_area, shape_leng)



The screenshot shows the QGIS Attribute Table for the 'NYC_CD' layer. The table has three columns: 'boro_cd', 'shape_area', and 'shape_leng'. The first row, which contains the value '308' in the 'boro_cd' column, is highlighted with a green border. The table displays 16 rows of data, each representing a geographic feature with its corresponding code and area/leng values.

	boro_cd	shape_area	shape_leng
1	308	45603787...	38232.886...
2	414	195576601...	210918.17...
3	204	55522140...	31358.450...
4	228	92938101...	122141.65...
5	205	38316975...	29443.048...
6	313	88180564.15	65789.792...
7	311	103208266...	51534.144...
8	484	120185760...	227190.54...
9	310	111345162...	44781.826...
10	410	172077374...	105822.37...
11	164	38312378...	32721.097...
12	203	44796444...	33500.037...
13	206	42664319...	35875.709...
14	226	50565798...	32820.482...
15	480	32777561...	47338.739...
16	105	43790039...	35288.334...

For a definition of boro_cd, we can look at the metadata for NYC_CD (nycd_metadata.pdf). The metadata seem out of date (BoroCD \neq boro_cd), but there's important info here about this code is constructed

FIELD BoroCD

- * **ALIAS** BoroCD
- * **DATA TYPE** SmallInteger
- * **WIDTH** 2
- * **PRECISION** 0
- * **SCALE** 0

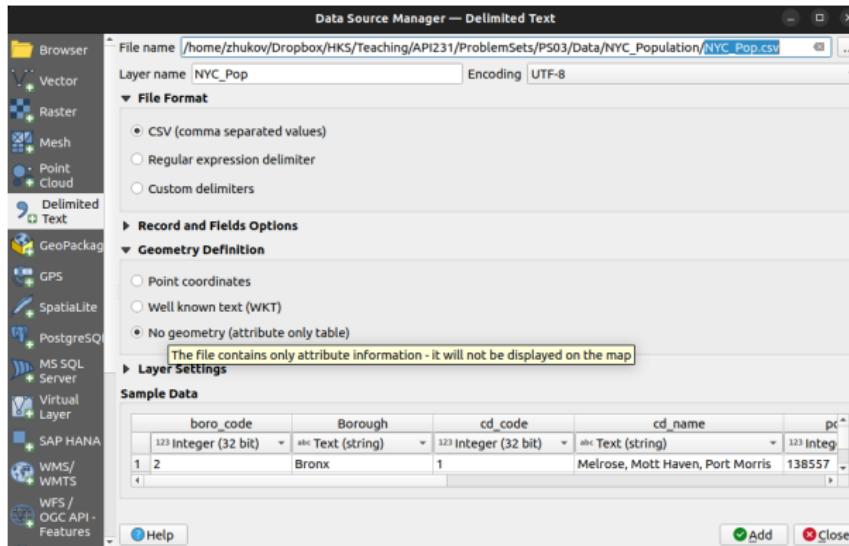
FIELD DESCRIPTION

Community district number preceded by BoroCode.

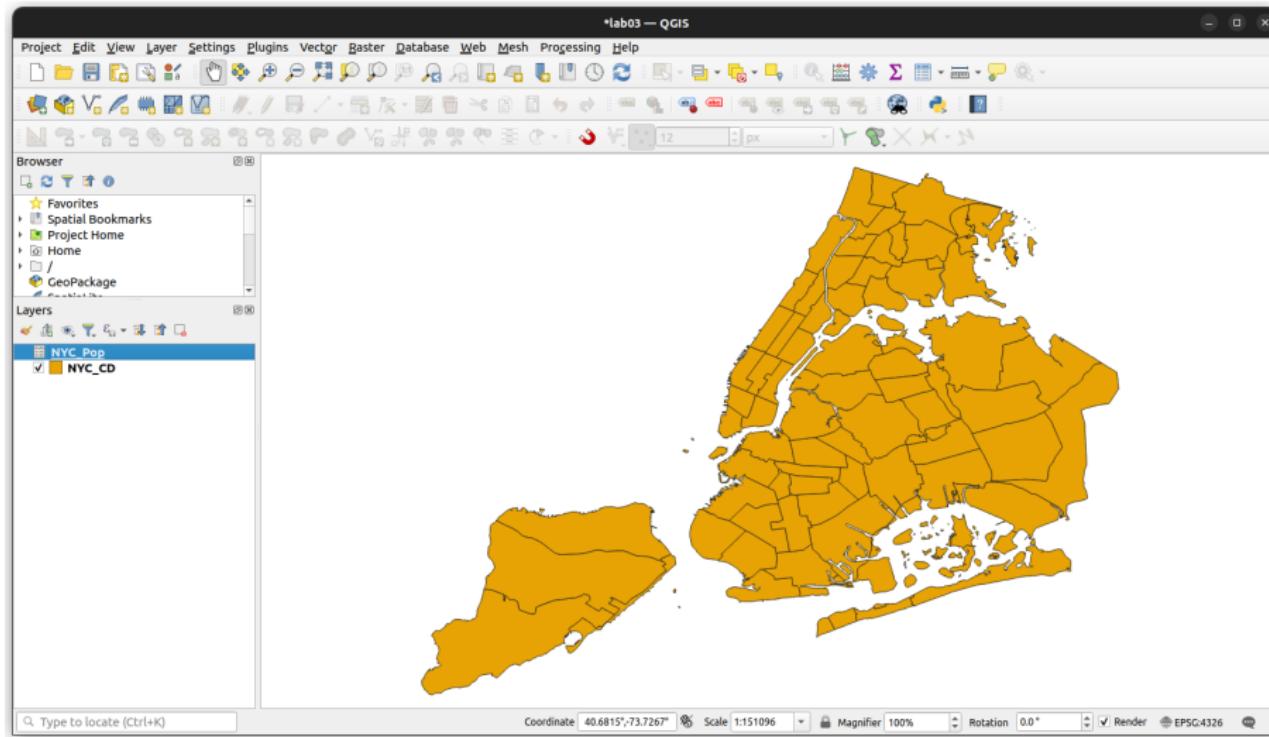
[Hide Field BoroCD](#) ▲

Load population data:

- Layer → Add Layer → Add Delimited Text Layer...
- Navigate to the file NY_Pop.csv in the NYC_Population folder
- Set Geometry Definition = No geometry (attribute table only)
- Click Add



The new layer is visible in the side menu. But it doesn't appear on the map, since it's not a spatial dataset. Right-click NYC_Pop → Open Attribute Table



This is a table of community district names and population statistics. Maybe we can link this table to NYC_CD... but there is no common variable

NYC_Pop — Features Total: 59, Filtered: 59, Selected: 0

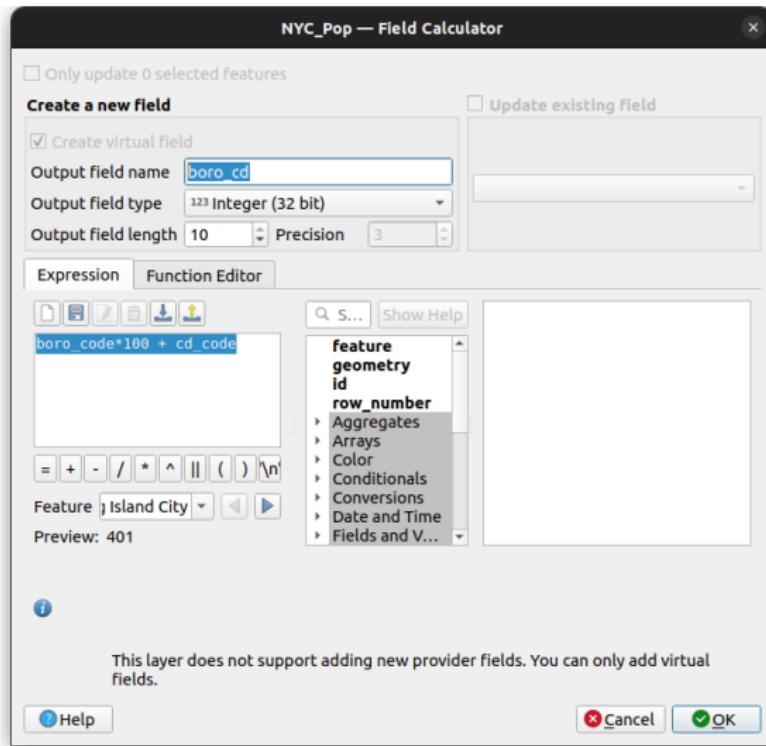
	boro_code	Borough	cd_code	cd_name	pop_1970	pop_1980	pop_1990	pop_2000	pop_2010
1	2	Bronx	1	Melrose, M...	138557	78441	77214	82159	91497
2	2	Bronx	2	Hunts Poin...	99493	34399	39443	46824	52246
3	2	Bronx	3	Morrisania,...	150636	53635	57162	68574	79762
4	2	Bronx	4	Highbridge...	144207	114312	119962	139563	146441
5	2	Bronx	5	University ...	121807	107995	118435	128313	128200
6	2	Bronx	6	East Tremo...	114137	65016	68061	75688	83268
7	2	Bronx	7	Bedford Pa...	113764	116827	128588	141411	139286
8	2	Bronx	8	Riverdale, ...	103543	98275	97030	101332	101731
9	2	Bronx	9	Soundview,...	166442	167627	155970	167859	172298
10	2	Bronx	10	Throgs Nk,...	84948	106516	108093	115948	120392
11	2	Bronx	11	Pelham Pk...	105980	99080	97842	110706	113232
12	2	Bronx	12	Wakefield, ...	135010	128226	129620	149077	152344
13	3	Brooklyn	1	Williamsbu...	179390	142942	155972	160338	173083
14	3	Brooklyn	2	Brooklyn H...	110221	92732	94534	98620	99617
15	3	Brooklyn	3	Bedford St...	203380	133379	138696	143867	152985
16	3	Brooklyn	4	Bushwick	137902	92497	102572	104358	112634

Wait. NYC_CD has a boro_cd code, which is “Community district number preceded by BoroCode”. Maybe we can create a boro_cd code in NYC_Pop, using boro_code and cd_code? Open Field Calculator

The screenshot shows the QGIS Field Calculator dialog for the NYC_Pop layer. The title bar indicates "NYC_Pop — Features Total: 59, Filtered: 59, Selected: 0". The dialog contains a toolbar at the top and a table below with columns: boro_code, Borough, cd_code, cd_name, Open field calculator (Ctrl+I), pop_1990, pop_2000, and pop_2010. The "Open field calculator" column is highlighted with a green border. The table lists 16 rows of data for Bronx and Brooklyn community districts. Row 16 is currently selected, indicated by a green highlight in the first column.

	boro_code	Borough	cd_code	cd_name	Open field calculator (Ctrl+I)	pop_1990	pop_2000	pop_2010
1	2	Bronx	1	Melrose, M...	138557	78441	77214	82159
2	2	Bronx	2	Hunts Poin...	99493	34399	39443	46824
3	2	Bronx	3	Morrisania,...	150636	53635	57162	68574
4	2	Bronx	4	Highbridge...	144207	114312	119962	139563
5	2	Bronx	5	University ...	121807	107995	118435	128313
6	2	Bronx	6	East Tremo...	114137	65016	68061	75688
7	2	Bronx	7	Bedford Pa...	113764	116827	128588	141411
8	2	Bronx	8	Riverdale, ...	103543	98275	97030	101332
9	2	Bronx	9	Soundview,...	166442	167627	155970	167859
10	2	Bronx	10	Throgs Nk,...	84948	106516	108093	115948
11	2	Bronx	11	Pelham Pk...	105980	99080	97842	110706
12	2	Bronx	12	Wakefield, ...	135010	128226	129620	149077
13	3	Brooklyn	1	Williamsbu...	179390	142942	155972	160338
14	3	Brooklyn	2	Brooklyn H...	110221	92732	94534	98620
15	3	Brooklyn	3	Bedford St...	203380	133379	138696	143867
16	3	Brooklyn	4	Bushwick	137902	92497	102572	104358

Create a new variable, called `boro_cd`, of type Integer. For the Expression, write `boro_code * 100 + cd_code`. Click OK

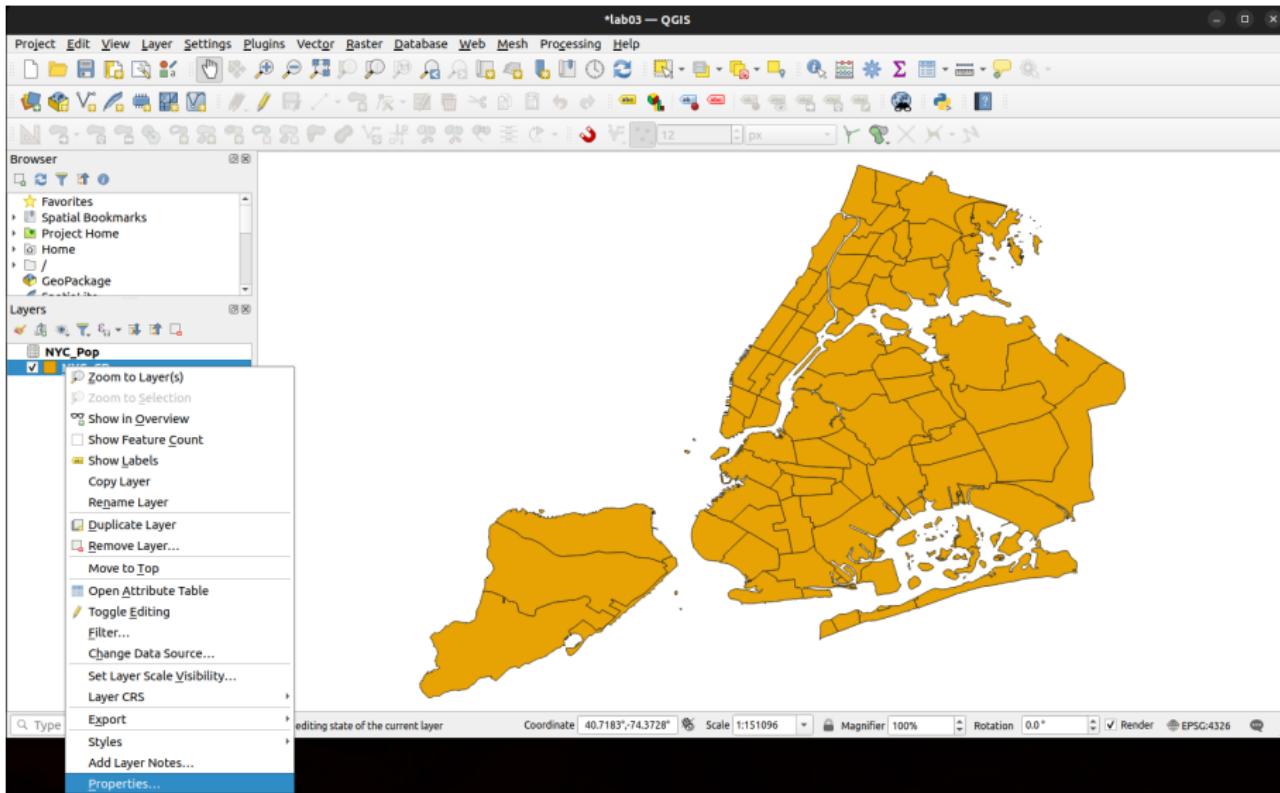


You should now see the variable `boro_cd` in the attribute table.

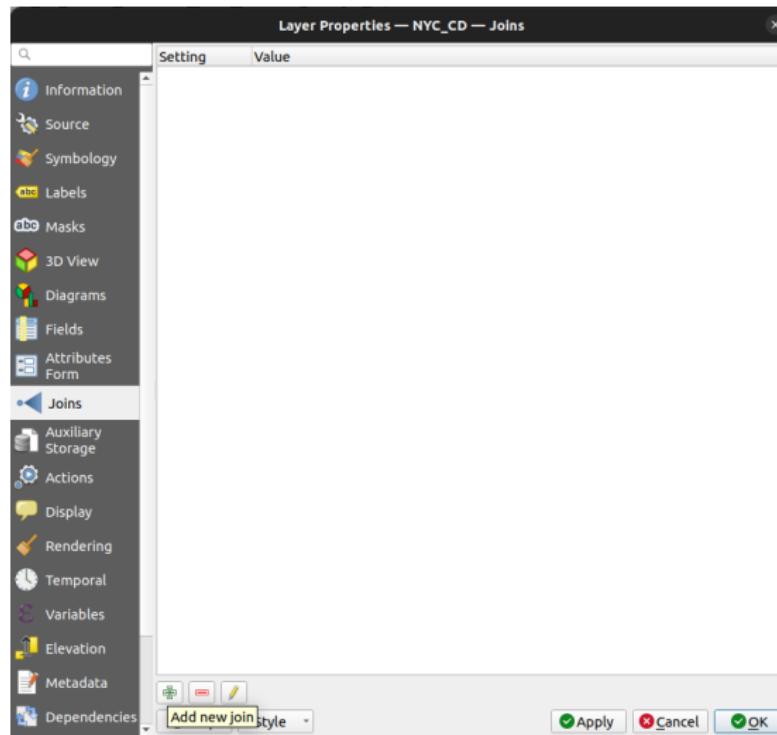
NYC_Pop — Features Total: 59, Filtered: 59, Selected: 0

	boro_code	Borough	cd_code	cd_name	pop_1970	pop_1980	pop_1990	pop_2000	pop_2010	boro_cd
1	2	Bronx	1	Melrose, M...	138557	78441	77214	82159	91497	201
2	2	Bronx	2	Hunts Poin...	99493	34399	39443	46824	52246	202
3	2	Bronx	3	Morrisania,...	150636	53635	57162	68574	79762	203
4	2	Bronx	4	Highbridge...	144207	114312	119962	139563	146441	204
5	2	Bronx	5	University ...	121807	107995	118435	128313	128200	205
6	2	Bronx	6	East Tremo...	114137	65016	68061	75688	83268	206
7	2	Bronx	7	Bedford Pa...	113764	116827	128588	141411	139286	207
8	2	Bronx	8	Riverdale, ...	103543	98275	97030	101332	101731	208
9	2	Bronx	9	Soundview,...	166442	167627	155970	167859	172298	209
10	2	Bronx	10	Throgs Nk,...	84948	106516	108093	115948	120392	210
11	2	Bronx	11	Pelham Pk...	105980	99080	97842	110706	113232	211
12	2	Bronx	12	Wakefield, ...	135010	128226	129620	149077	152344	212
13	3	Brooklyn	1	Williamsbu...	179390	142942	155972	160338	173083	301
14	3	Brooklyn	2	Brooklyn H...	110221	92732	94534	98620	99617	302
15	3	Brooklyn	3	Bedford St...	203380	133379	138696	143867	152985	303
16	3	Brooklyn	4	Bushwick	137902	92497	102572	104358	112634	304

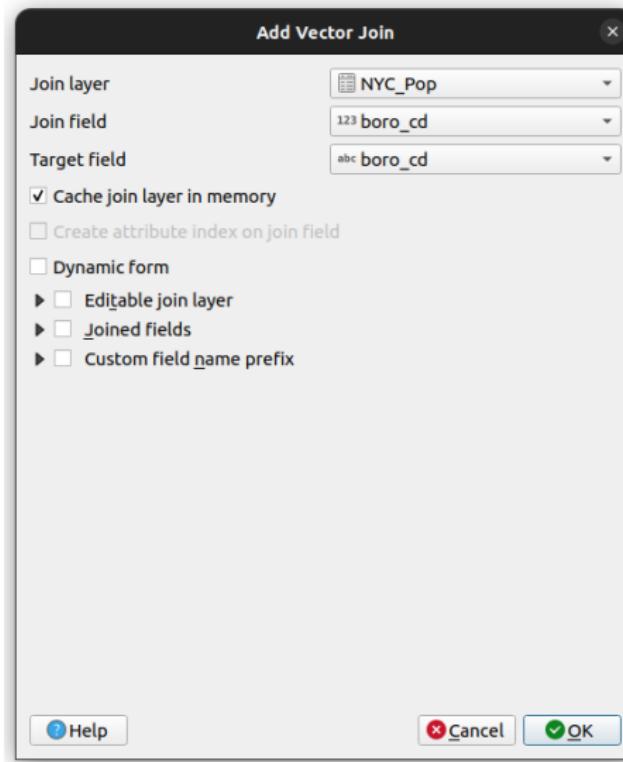
Let's join NYC_CD and NYC_Pop. Go to the Properties for NYC_CD



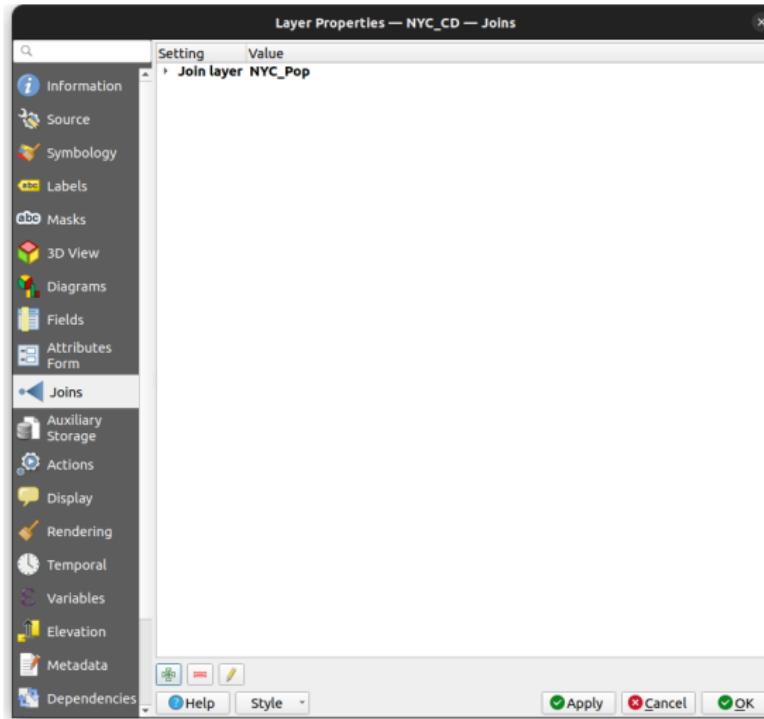
Open the Joins tab in Properties. Click on the “add” button (+)



Select NYC_Pop as the join layer, and boro_cd as the join and target field. Click OK



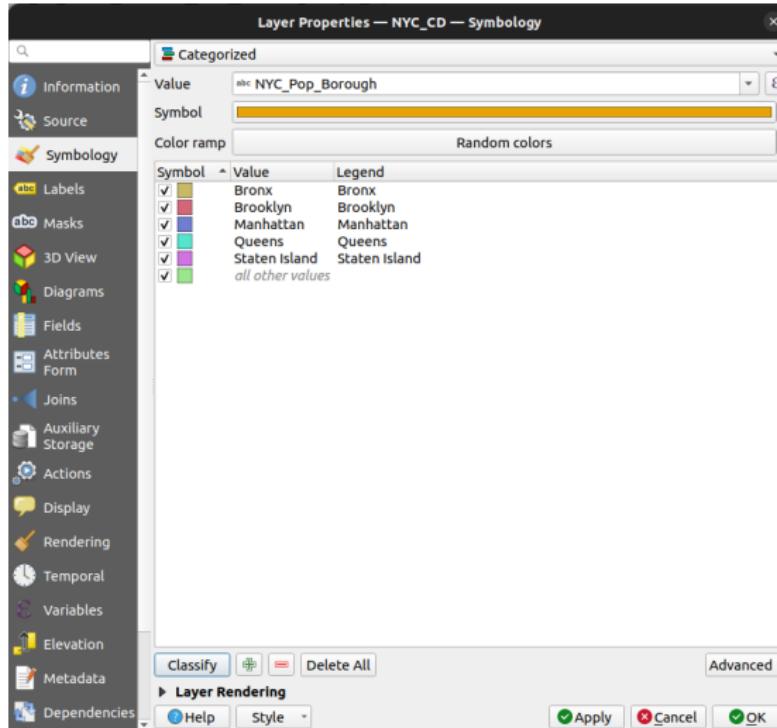
You should now see a new join in the Joins tab



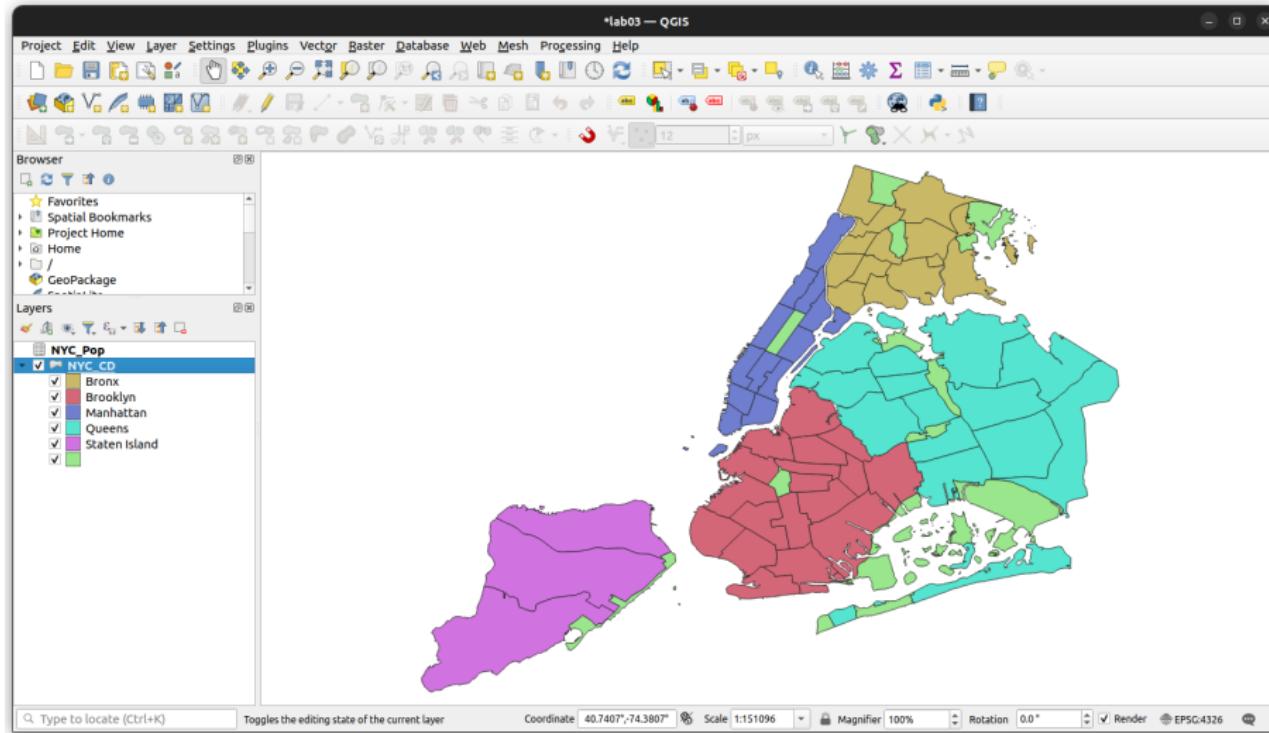
Now go to the Symbology tab in Properties.

Set Symbology type = Categorized, and Value = NYC_Pop_Borough.

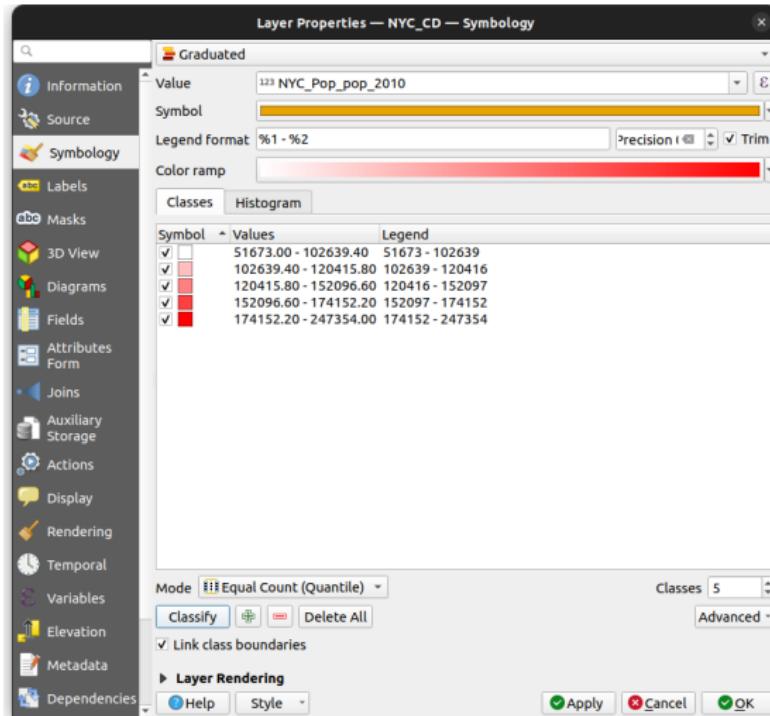
Click Classify, then OK



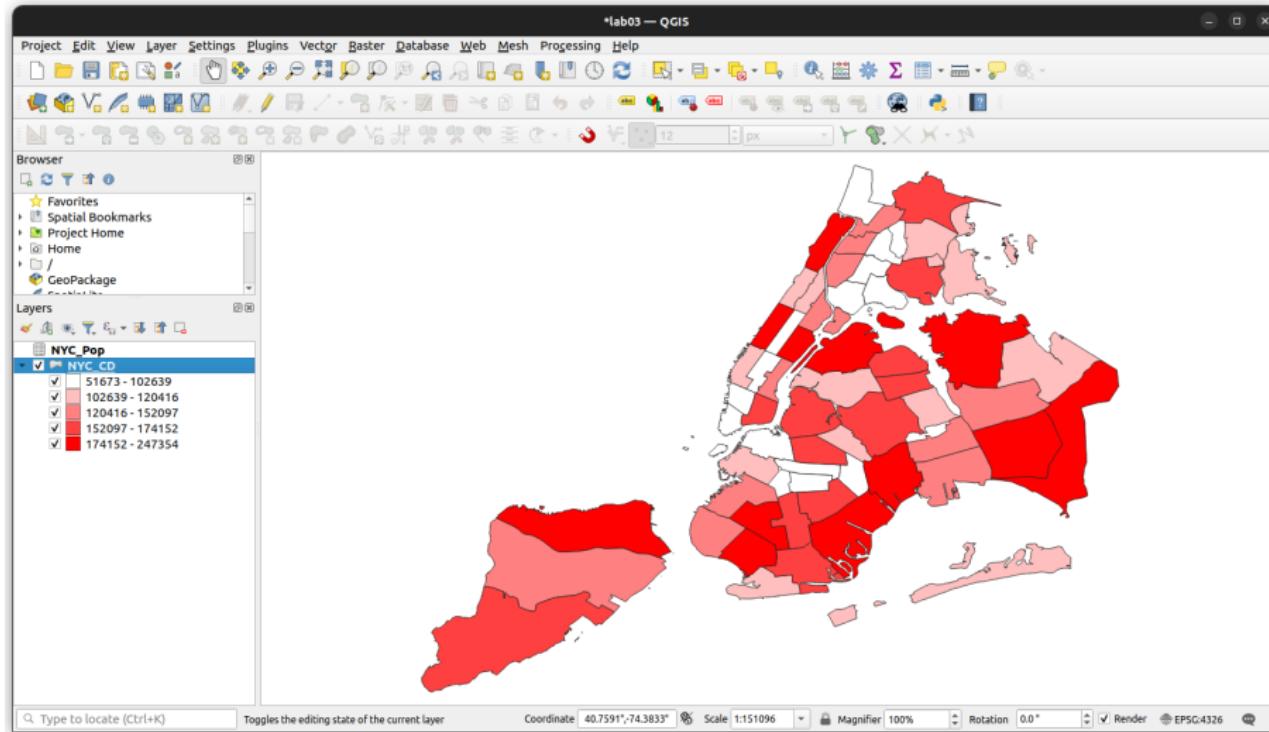
The districts should now be colored by borough



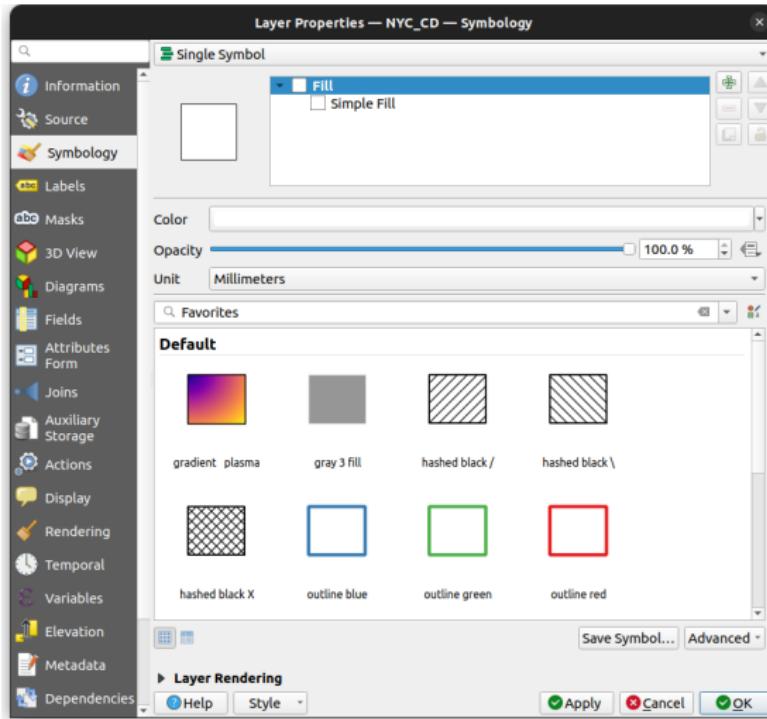
Let's now color them by population size. Go back to Properties → Symbology.
Set Symbology type = Graduated, and Value = NYC_Pop_pop_2010.
Click Classify, then OK



The districts should now be colored by their population in 2010

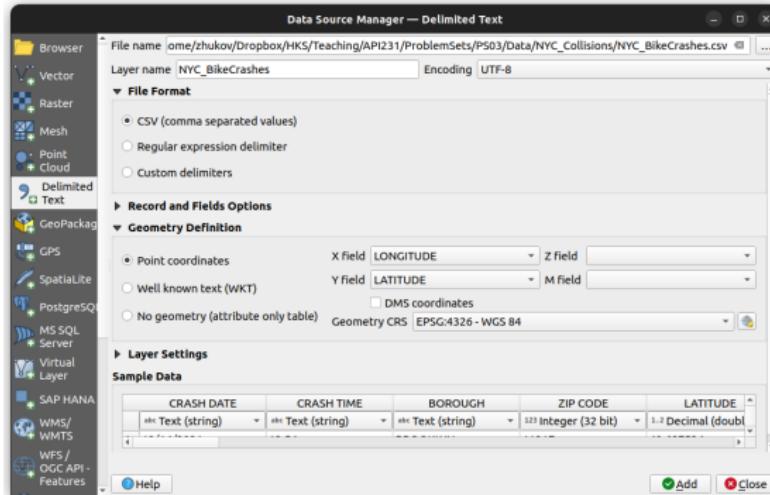


Set the color scheme back to Single symbol in Symbology

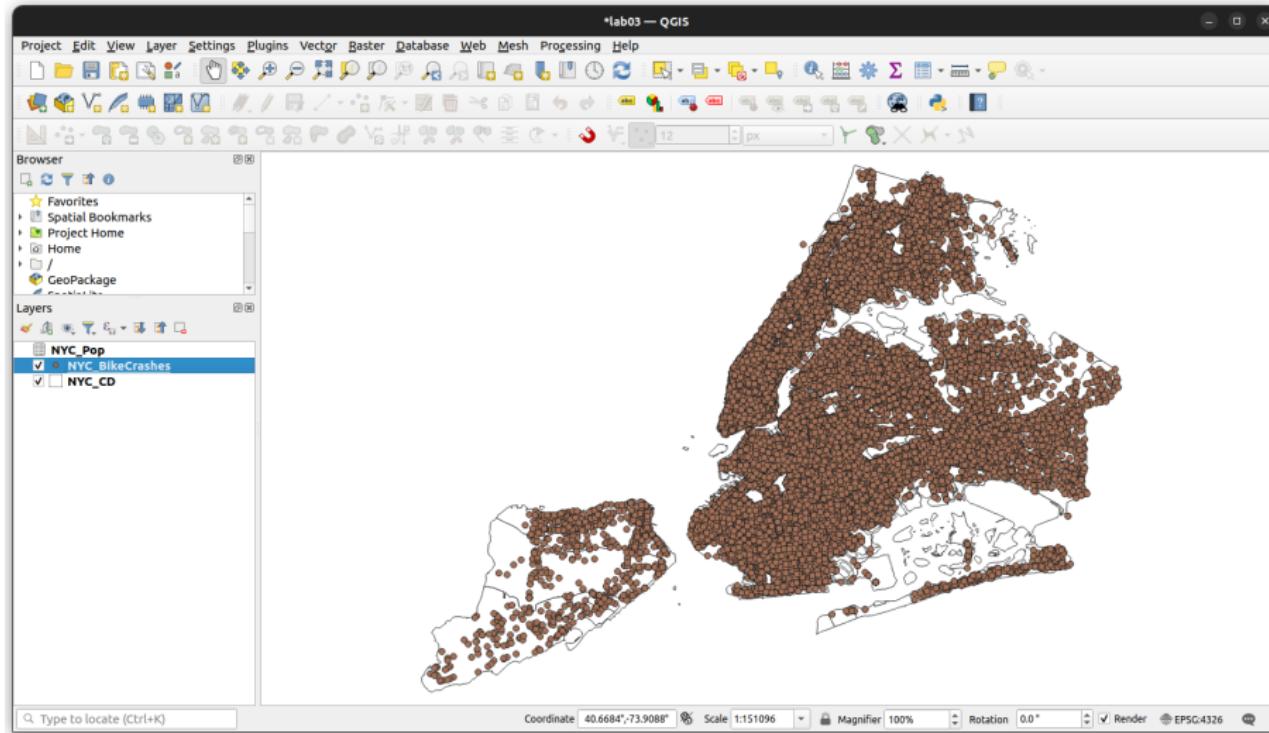


Load data on bicycle crashes:

- Go to Add Delimited Text Layer
- Navigate to NYC_BikeCrashes.csv in the NYC_Collisions folder
- Set Geometry Definition = Point coordinates
- Set X field = LONGITUDE, Y field = LATITUDE.
- Click Add

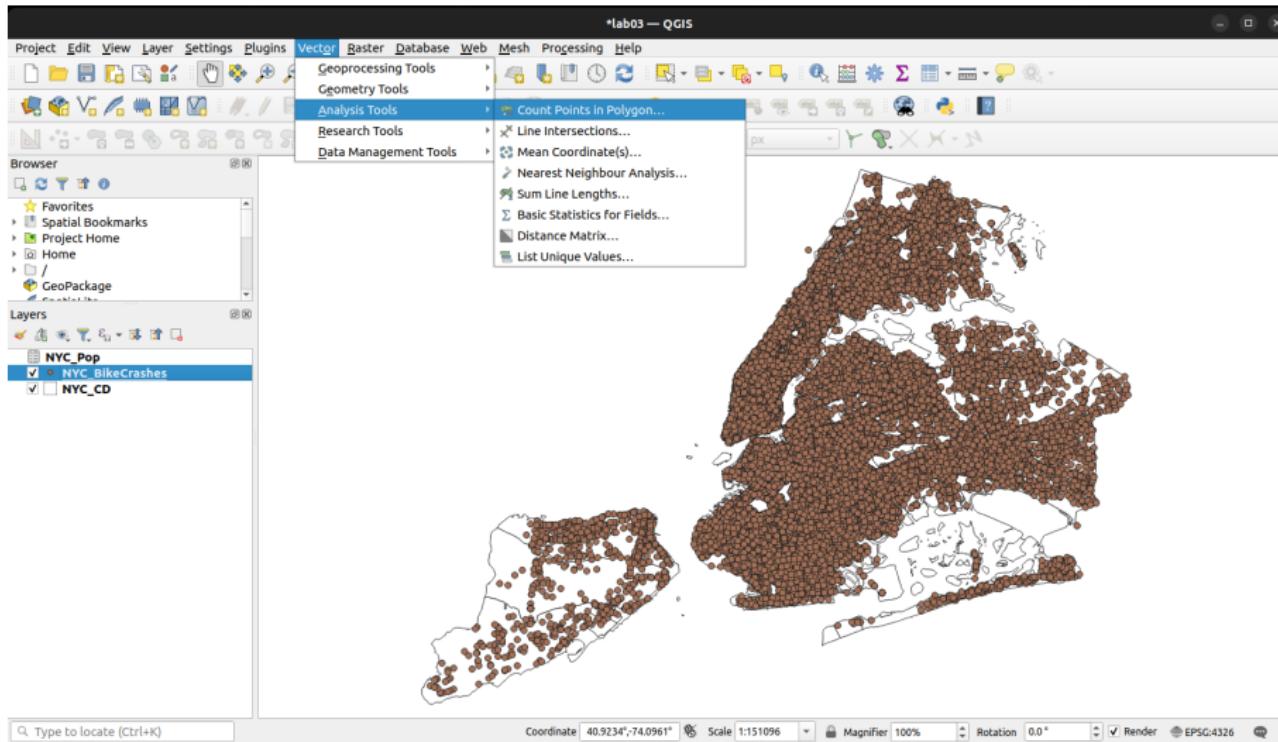


The new layer should be visible in the map, as points.

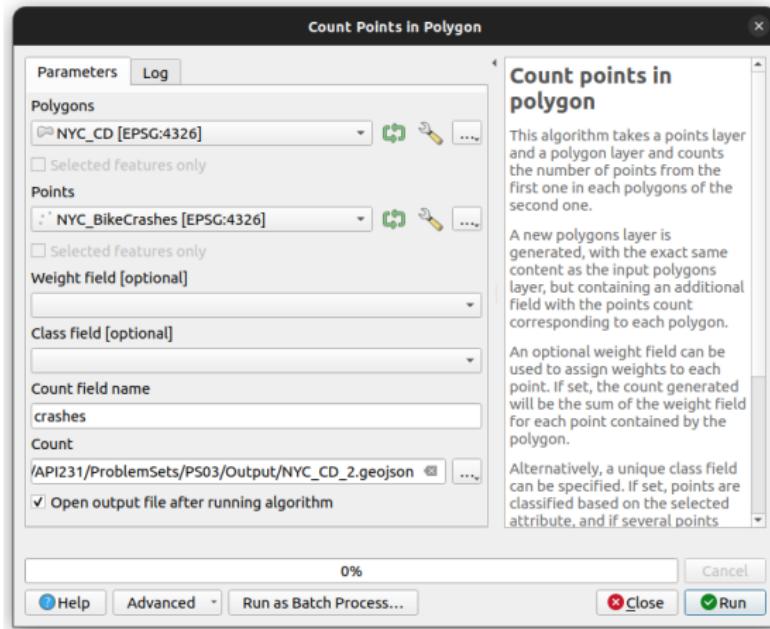


Navigate to the Count Points in Polygon tool.

Vector menu → Analysis Tools → Count Points in Polygon

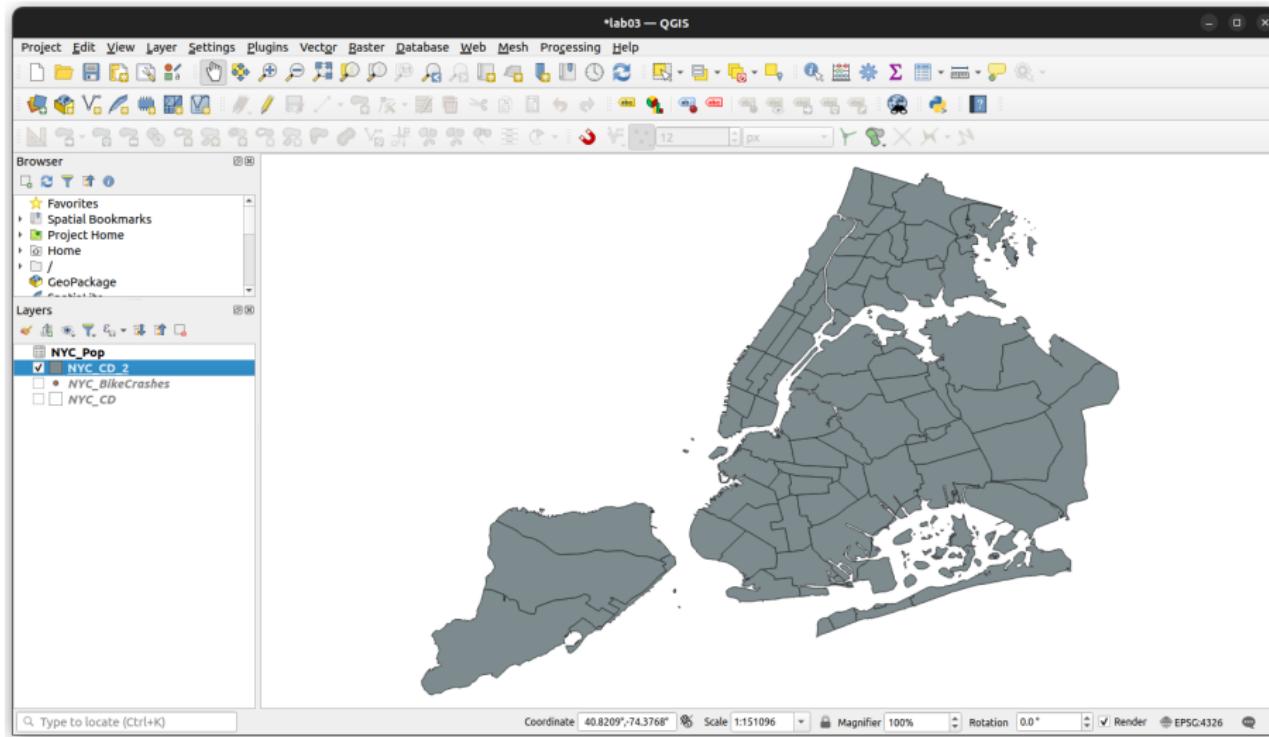


Select Polygons = NYC_CD, Points = NYC_BikeCrashes. Name the count field crashes, and save the output file as NYC_CD_2.geojson. Click Run

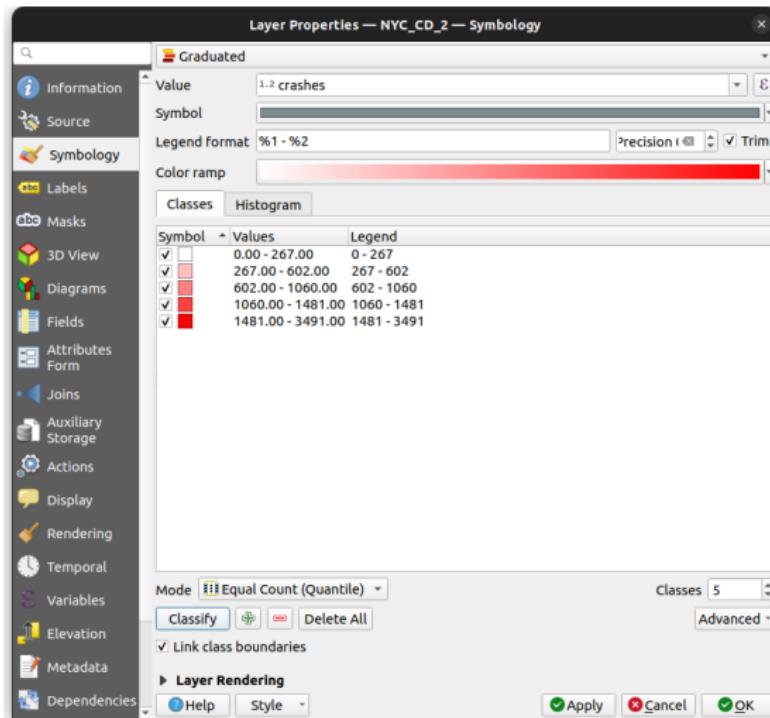


The new layer should appear on the map as `NYC_CD_2`.

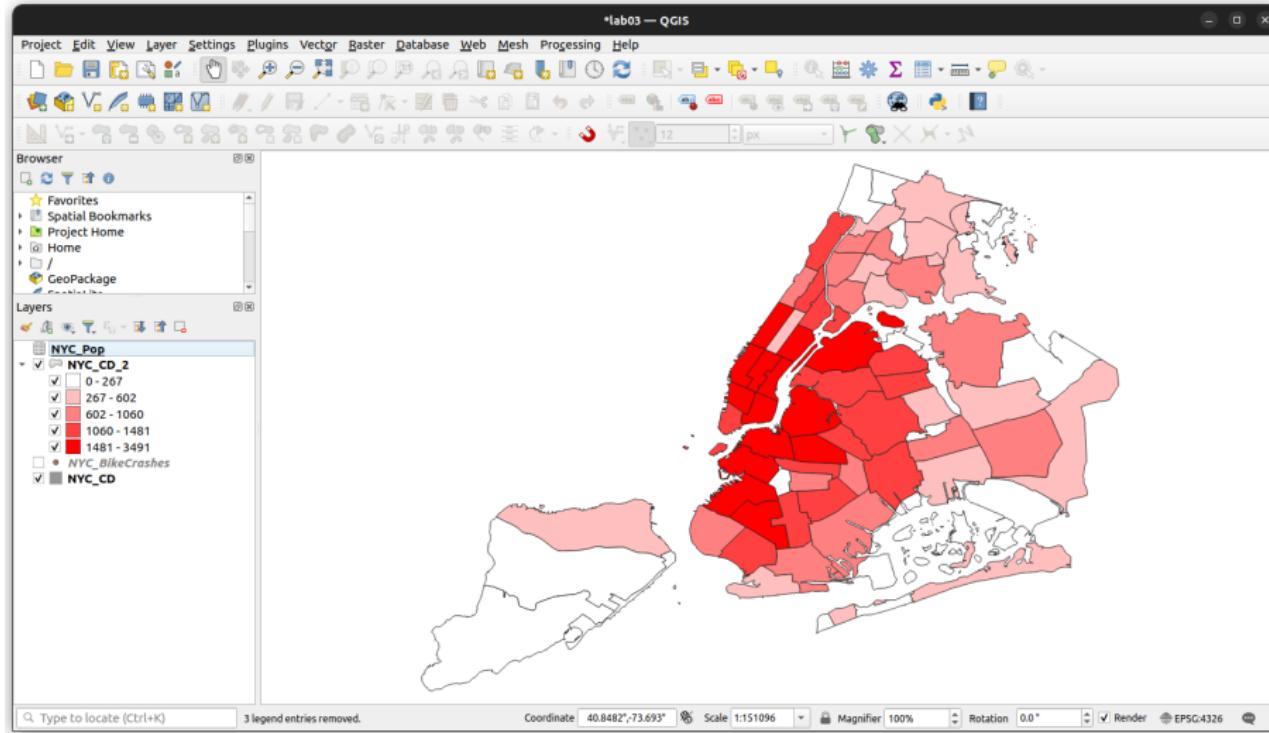
Hide the other two layers by unchecking boxes next to them in the menu.



Let's color the districts by number of bike crashes. Go to the Symbology tab in the new layer's Properties. Set Symbology type = Graduated, and Value = crashes. Click Classify, then OK



In which parts of NYC are cars hitting the most cyclists?



Let's combine this with population data to get a per-capita rate of bike crashes.
Open the Attribute Table for NYC_CD_2. Open the Field Calculator

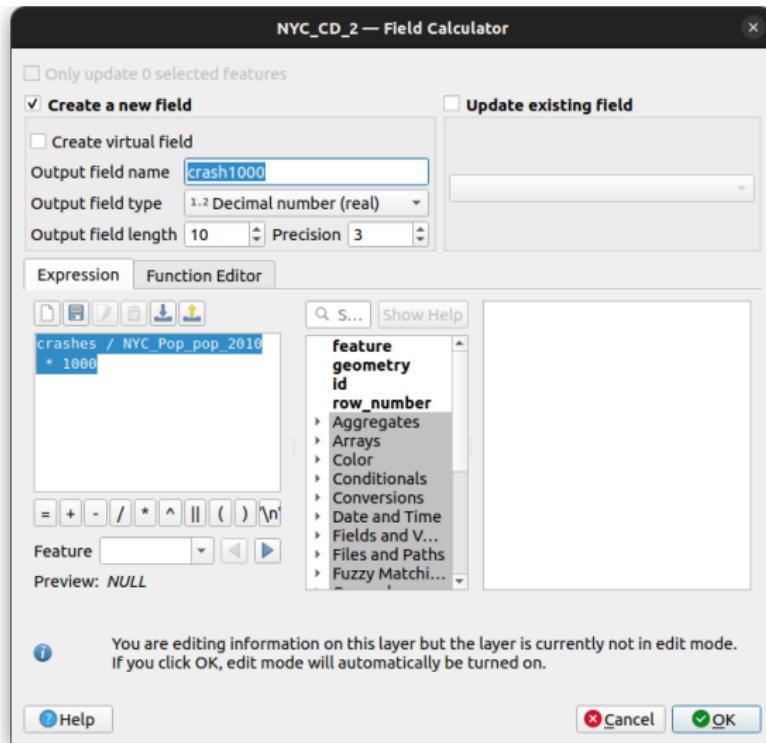
NYC_CD_2 — Features Total: 71, Filtered: 71, Selected: 0

The screenshot shows the QGIS Attribute Table for the NYC_CD_2 layer. The table has 15 columns and 15 rows. The first column is labeled 'boro_cd' and contains values like 308, 414, 204, etc. The second column is 'shape_area' and the third is 'shape_leng'. The fourth column is 'Pop_boro_cd' and the fifth is 'C_Pop_Boro'. The sixth column is 'Open field calculator (Ctrl+I)' which is currently active. The next four columns are 'Pop_pop_15C', 'C_Pop_pop_15C', 'Pop_pop_15C', and 'C_Pop_pop_15C'. The last column is 'C_Pop_pop_15C'. The 'Open field calculator (Ctrl+I)' column contains the expression '\$area / \$length * 1000000'. The 'C_Pop_Boro' column contains the expression 'if(boro_cd = 1, "Brooklyn", if(boro_cd = 2, "Queens", if(boro_cd = 3, "Bronx", if(boro_cd = 4, "The Bronx", if(boro_cd = 5, "Coney Island", if(boro_cd = 6, "Bensonhurst", if(boro_cd = 7, "Bay Ridge", if(boro_cd = 8, "Crown Heights", if(boro_cd = 9, "University Heights", if(boro_cd = 10, "Ozone Park", if(boro_cd = 11, "Morrisania", if(boro_cd = 12, "East Tremont", if(boro_cd = 13, "South Bronx", if(boro_cd = 14, "North Bronx", if(boro_cd = 15, "Staten Island", null)))))))))))))))'. The 'C_Pop_Boro' column is highlighted with a green border.

	boro_cd	shape_area	shape_leng	Pop_boro_cd	C_Pop_Boro	Open field calculator (Ctrl+I)	Pop_pop_15C	C_Pop_pop_15C	Pop_pop_15C	C_Pop_pop_15C
1	308	45603787...	38232.886...	3	Brooklyn	\$area / \$length * 1000000	8	Crown Hei...	121821	88796
2	414	195576601...	210918.17...	4	Queens		14	The Rocka...	98228	100592
3	204	55522140...	31358.450...	2	Bronx		4	Highbridge...	144207	114312
4	228	92938101...	122141.65...	NULL	NULL		NULL	NULL	NULL	NULL
5	205	38316975...	29443.048...	2	Bronx		5	University ...	121807	107995
6	313	88180564.15	65789.792...	3	Brooklyn		13	Coney Islan...	97750	100030
7	311	103208266...	51534.144...	3	Brooklyn		11	Bensonhur...	170119	155072
8	484	120185760...	227190.54...	NULL	NULL		NULL	NULL	NULL	NULL
9	310	111345162...	44781.826...	3	Brooklyn		10	Bay Ridge, ...	129822	118187
10	410	172077374...	105822.37...	4	Queens		10	Ozone Park...	113857	105651
11	164	38312378...	32721.097...	NULL	NULL		NULL	NULL	NULL	NULL
12	203	44796444...	33500.037...	2	Bronx		3	Morrisania,...	150636	53635
13	206	42664319...	35875.709...	2	Bronx		6	East Tremo...	114137	65016
14	226	50565798...	32820.482...	NULL	NULL		NULL	NULL	NULL	NULL
15	480	32777561...	47338.739...	NULL	NULL		NULL	NULL	NULL	NULL
16										
17										
18										
19										
20										
21										
22										
23										
24										
25										
26										
27										
28										
29										
30										
31										
32										
33										
34										
35										
36										
37										
38										
39										
40										
41										
42										
43										
44										
45										
46										
47										
48										
49										
50										
51										
52										
53										
54										
55										
56										
57										
58										
59										
60										
61										
62										
63										
64										
65										
66										
67										
68										
69										
70										
71										

Show All Features

Create new field called crash1000 of type Decimal number (real).
For the Expression, write crashes / NYC_Pop_popP2010 * 1000. Click OK



You should now see the new variable in the attribute table.

NYC_CD_2 — Features Total: 71, Filtered: 71, Selected: 0

The screenshot shows the QGIS attribute table for the 'NYC_CD_2' layer. The table has 71 features and 15 columns. The columns are: boro_cd, shape_area, shape_len, C_Pop_boro_c/C_Pop_Borou /C_Pop_cd_cov/C_Pop_cd_nar, C_Pop_pop_1%, C_Pop_pop_1%, C_Pop_pop_1%, C_Pop_pop_2%, C_Pop_pop_2%, crashes, and crash1000. The 'crashes' column contains numerical values ranging from 2 to 12,744, while the 'crash1000' column contains values ranging from 3.479 to 13.227. The 'shape_area' and 'shape_len' columns show geographic dimensions for each borough. The 'borocd' column lists the borough codes (e.g., 308 for Brooklyn, 414 for Queens, etc.). The 'C_Pop...' columns represent population statistics for each borough. The 'crashes' and 'crash1000' columns represent the number of crashes and the crash rate per 1000 inhabitants respectively. The table includes standard QGIS interface elements like toolbar icons, a search bar, and buttons for 'Update All' and 'Update Selected'.

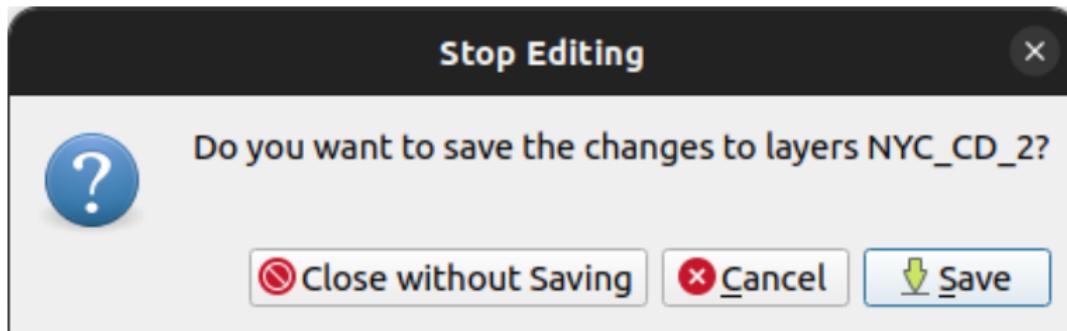
	boro_cd	shape_area	shape_len	C_Pop_boro_c/C_Pop_Borou /C_Pop_cd_cov/C_Pop_cd_nar	C_Pop_pop_1%	C_Pop_pop_1%	C_Pop_pop_1%	C_Pop_pop_2%	C_Pop_pop_2%	crashes	crash1000	
1	308	45603787...	38232.886...	3 Brooklyn	8 Crown Hel...	121821	88796	96400	96076	96317	1274	13.227
2	414	195576601...	210918.17...	4 Queens	14 The Rocka...	98228	100592	100596	106686	114978	400	3.479
3	204	55522140...	31358.450...	2 Bronx	4 Highbridge...	144207	114312	119962	139563	146441	820	5.600
4	228	92938101...	122141.65...	NULL NULL	NULL NULL	NULL	NULL	NULL	NULL	NULL	21	NULL
5	205	38316975...	29443.048...	2 Bronx	5 University ...	121807	107995	118435	128313	128200	724	5.647
6	313	88180564.15	65789.792...	3 Brooklyn	13 Coney Islan...	97750	100030	102596	106120	104278	527	5.054
7	311	103208266...	51534.144...	3 Brooklyn	11 Bensonhur...	170119	155072	149994	172129	181981	1117	6.138
8	484	120185760...	227190.54...	NULL NULL	NULL NULL	NULL	NULL	NULL	NULL	NULL	5	NULL
9	310	111345162...	44781.826...	3 Brooklyn	10 Bay Ridge, ...	129822	118187	110612	122542	124491	803	6.450
10	410	172077374...	105822.37...	4 Queens	10 Ozone Park...	113857	105651	107768	127274	122396	517	4.224
11	164	38312378...	32721.097...	NULL NULL	NULL NULL	NULL	NULL	NULL	NULL	NULL	404	NULL
12	203	44796444...	33500.037...	2 Bronx	3 Morrisania,...	150636	53635	57162	68574	79762	572	7.171
13	206	42664319...	35875.709...	2 Bronx	6 East Tremo...	114137	65016	68061	75688	83268	608	7.302
14	226	50565798...	32820.482...	NULL NULL	NULL NULL	NULL	NULL	NULL	NULL	NULL	12	NULL
15	480	32777561...	47338.739...	NULL NULL	NULL NULL	NULL	NULL	NULL	NULL	NULL	2	NULL

Click the Edit button to save the dataset

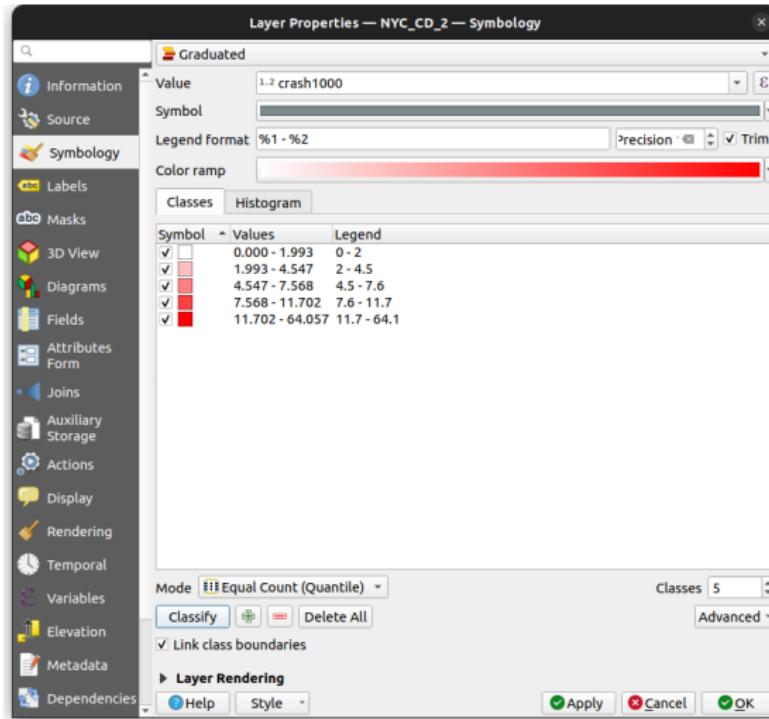
The screenshot shows the QGIS attribute table interface for the 'NYC_CD_2' dataset. The table has 71 features and 15 columns. The columns are: boro_cd, shape_area, shape_leng, Pop_boro_cd, Pop_Boro_cd, Pop_cd, Pop_cd_nar, Pop_pop_15C, Pop_pop_15C, Pop_pop_15C, Pop_pop_20C, Pop_pop_20C, crashes, and crash1000. The first few rows of data are as follows:

	boro_cd	shape_area	shape_leng	Pop_boro_cd	Pop_Boro_cd	Pop_cd	Pop_cd_nar	Pop_pop_15C	Pop_pop_15C	Pop_pop_15C	Pop_pop_20C	Pop_pop_20C	crashes	crash1000
1	308	45603787...	38232.886...	3	Brooklyn	8	Crown Hei...	121821	88796	96400	96076	96317	1274	13.227
2	414	195576601...	210918.17...	4	Queens	14	The Rocka...	98228	100592	100596	106686	114978	400	3.479
3	204	55522140...	31358.450...	2	Bronx	4	Highbridge...	144207	114312	119962	139563	146441	820	5.600
4	228	92938101...	122141.65...	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	21	NULL
5	205	38316975...	29443.048...	2	Bronx	5	University ...	121807	107995	118435	128313	128200	724	5.647
6	313	88180564.15	65789.792...	3	Brooklyn	13	Coney Islan...	97750	100030	102596	106120	104278	527	5.054
7	311	103208266...	51534.144...	3	Brooklyn	11	Bensonhur...	170119	155072	149994	172129	181981	1117	6.138
8	484	120185760...	227190.54...	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	5	NULL
9	310	111345162...	44781.826...	3	Brooklyn	10	Bay Ridge, ...	129822	118187	110612	122542	124491	803	6.450
10	410	172077374...	105822.37...	4	Queens	10	Ozone Park...	113857	105651	107768	127274	122396	517	4.224
11	164	38312378...	32721.097...	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	404	NULL
12	203	44796444...	33500.037...	2	Bronx	3	Morrisania,...	150636	53635	57162	68574	79762	572	7.171
13	206	42664319...	35875.709...	2	Bronx	6	East Tremo...	114137	65016	68061	75688	83268	608	7.302
14	226	50565798...	32820.482...	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	12	NULL
15	480	32777561...	47338.739...	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	2	NULL

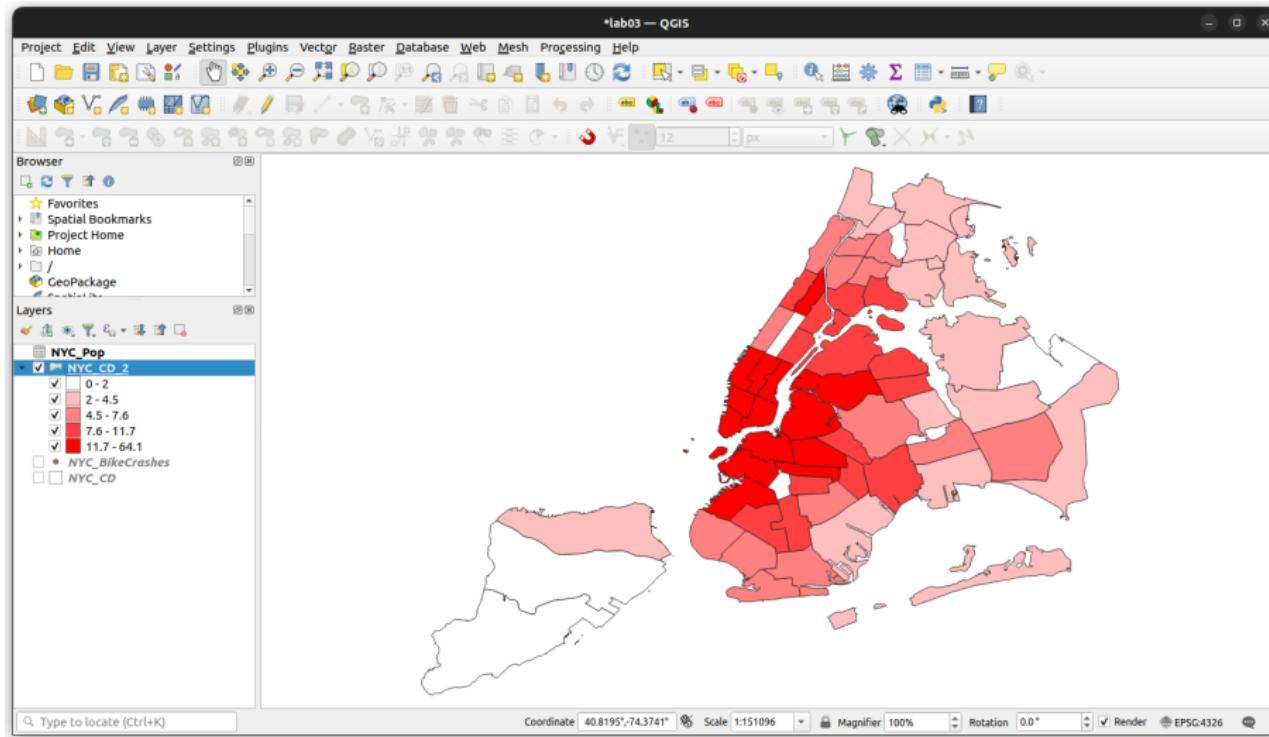
Click Save



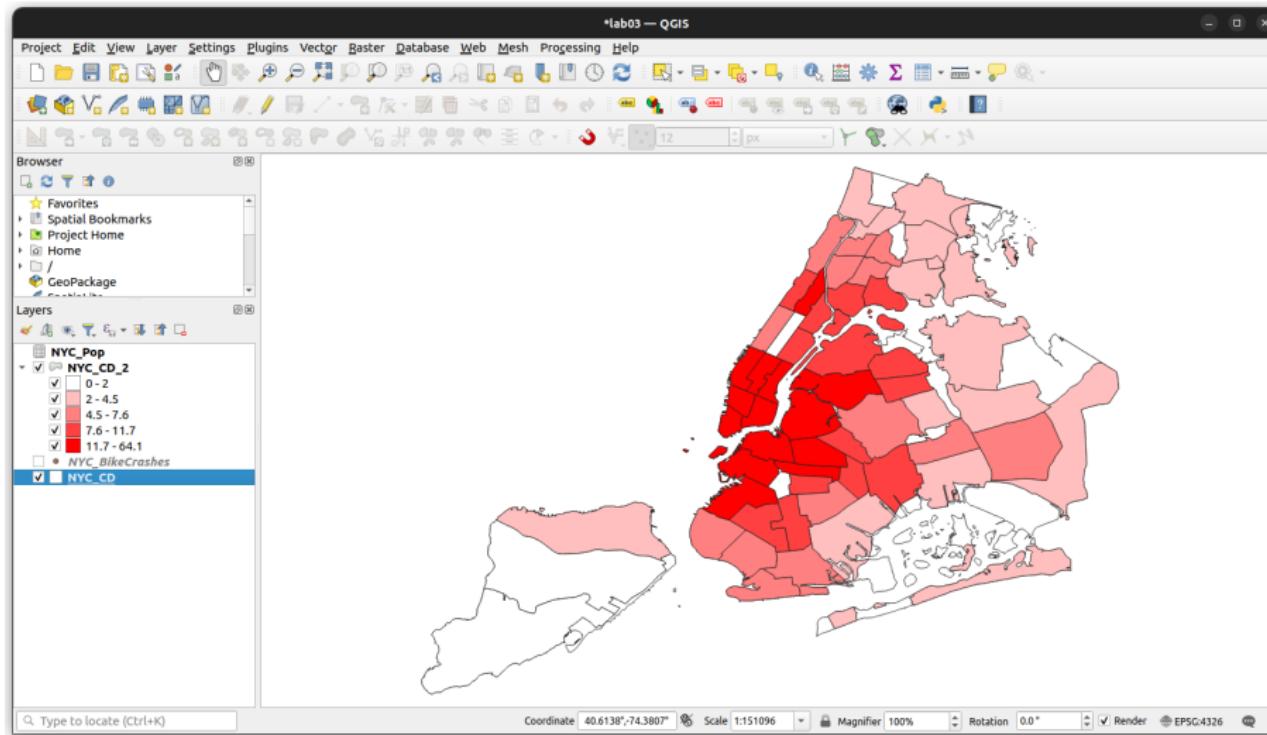
Change the symbology again, to color districts by crash1000. Remember to click Classify



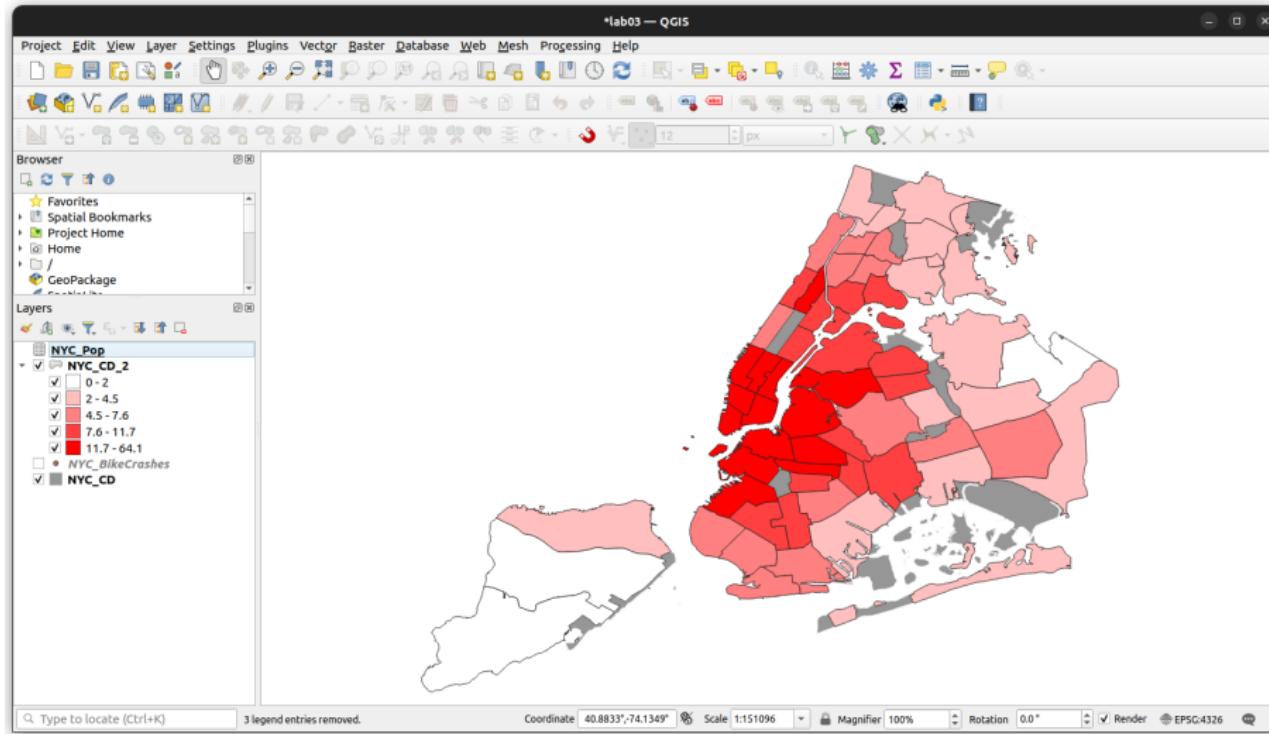
The distribution of crashes per 1000 residents (`crash1000`) looks similar to crashes. But there are a few missing districts (due to no population data)



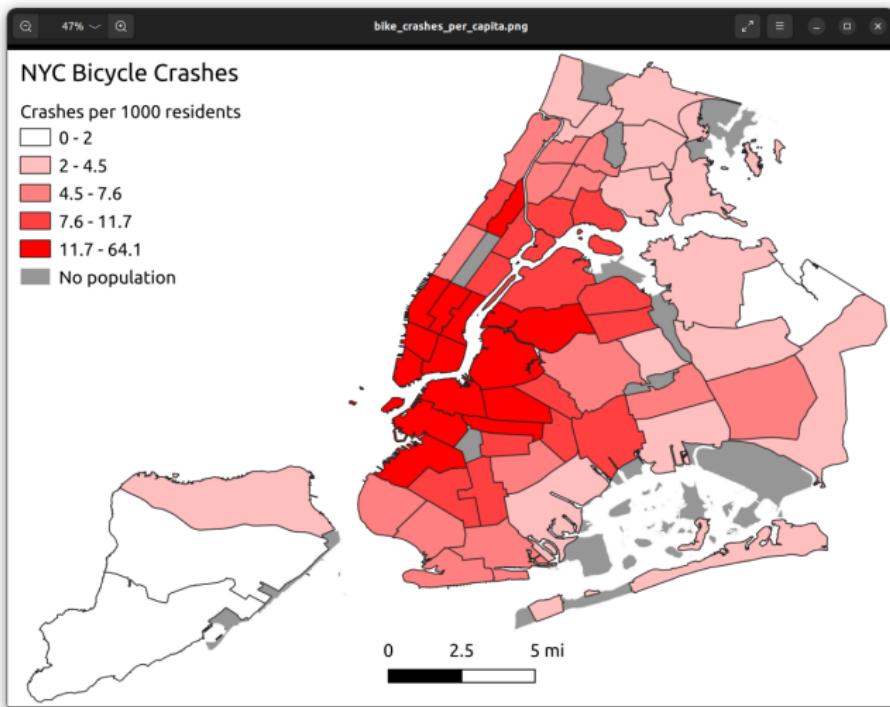
We can fix this (aesthetically) by un-hiding the NYC_CD layer.



You may want to change the color of NYC_CD to something neutral, like gray.



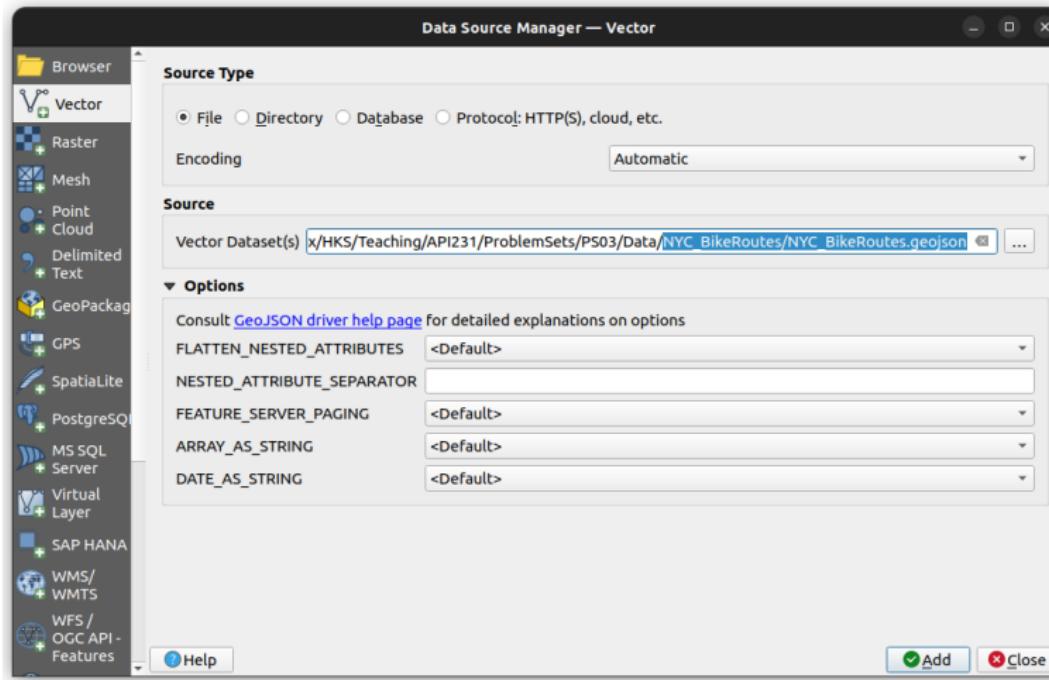
Export the map to image. Place map on a New Print Layout, add and properly format a legend, scale bar, etc. The end product should look vaguely like this.



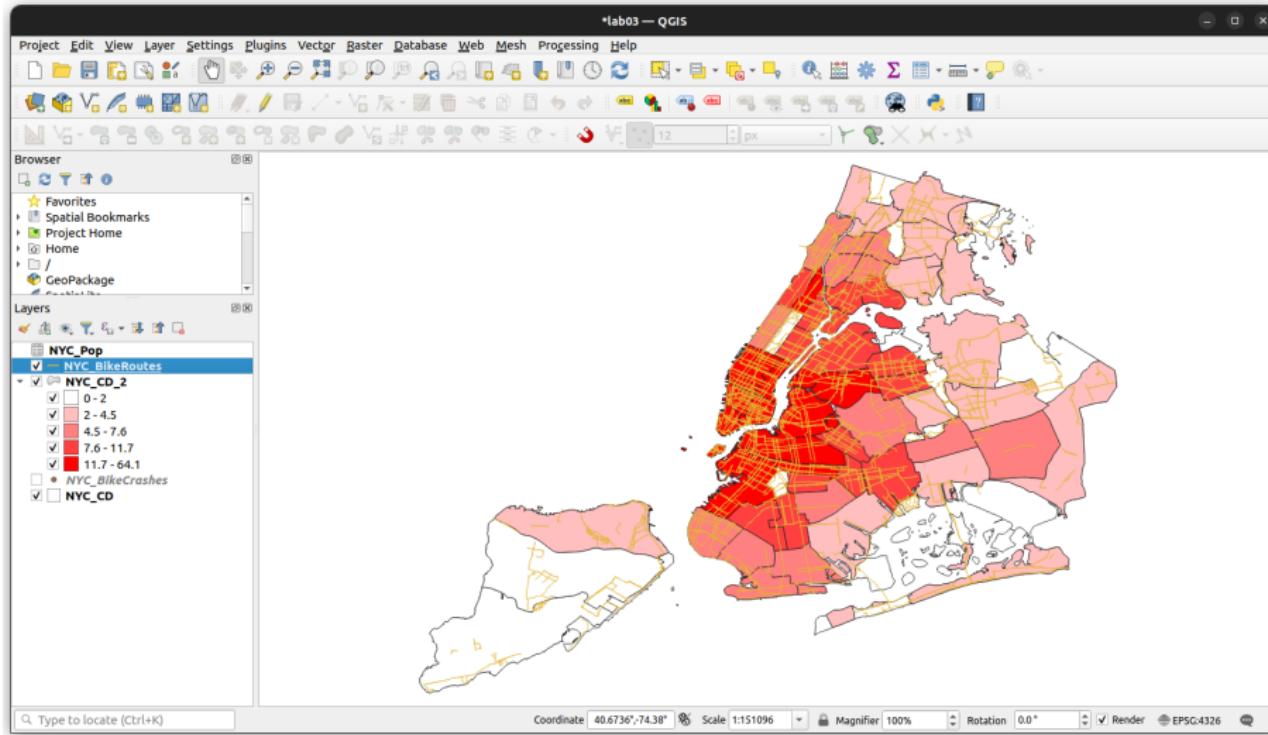
Map 2

Load bike routes: Layer → Add Layer → Add Vector Layer...

Navigate to NYC_BikeRoutes.geojson in the NYC_BikeRoutes folder. Click Add

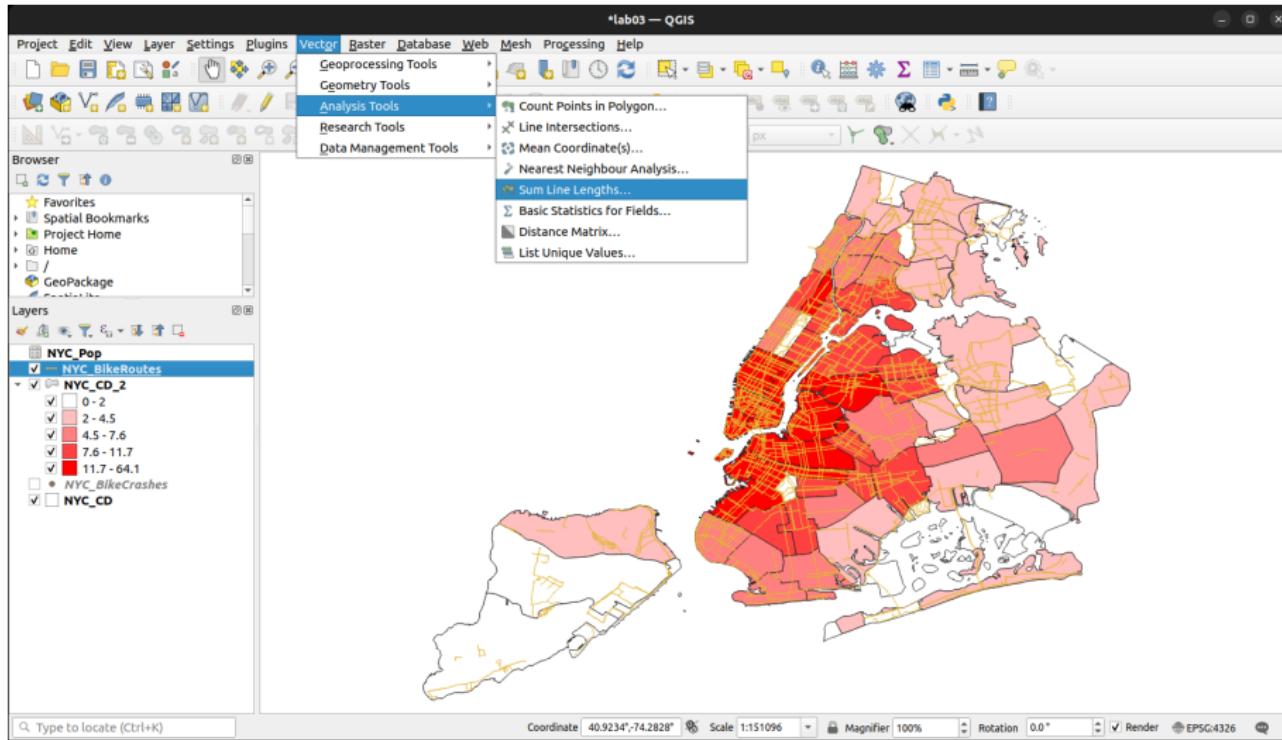


This is a polyline layer, representing NYC's cycling lanes and greenways.
Let's calculate how many miles of bike lanes are in each district.

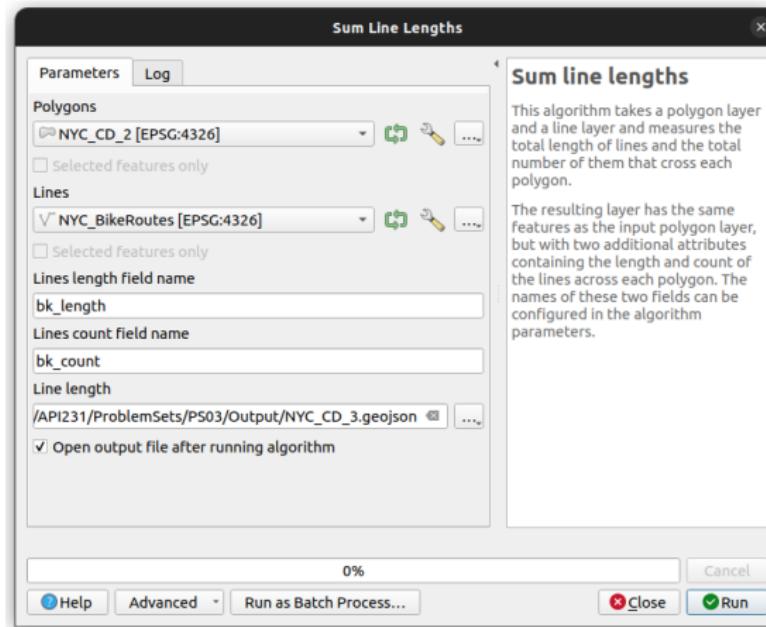


Navigate to the Sum Line Lengths tool.

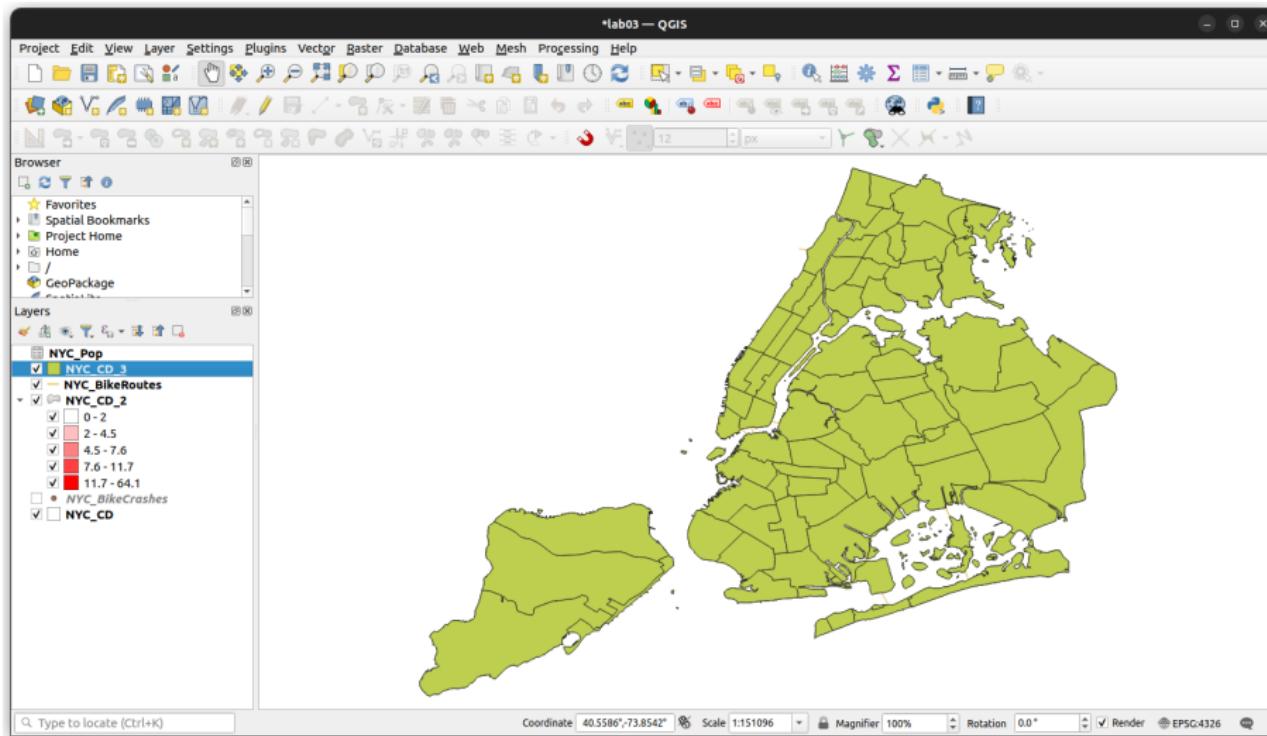
Vector menu → Analysis Tools → Sum Line Length



Select Lines = NYC_BikeRoutes, Polygons = NYC_CD_2.
Name the lengths and count fields bk_length and bk_count.
Save the output file as NYC_CD_3.geojson. Click Run



The new layer should appear as NYC_CD_3 on the map.
Let's take a look inside

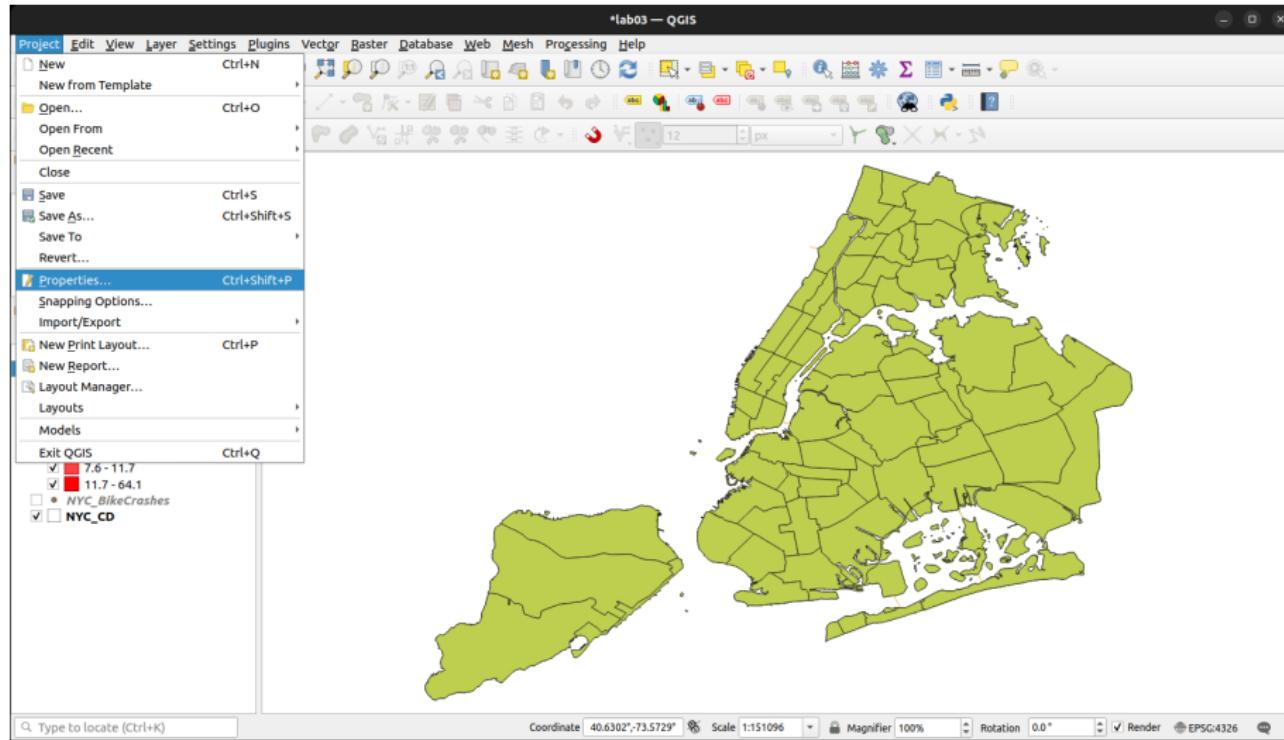


Open the Attribute Table for NYC_CD_3. The bk_length variable is there, but what are its units of measurement?

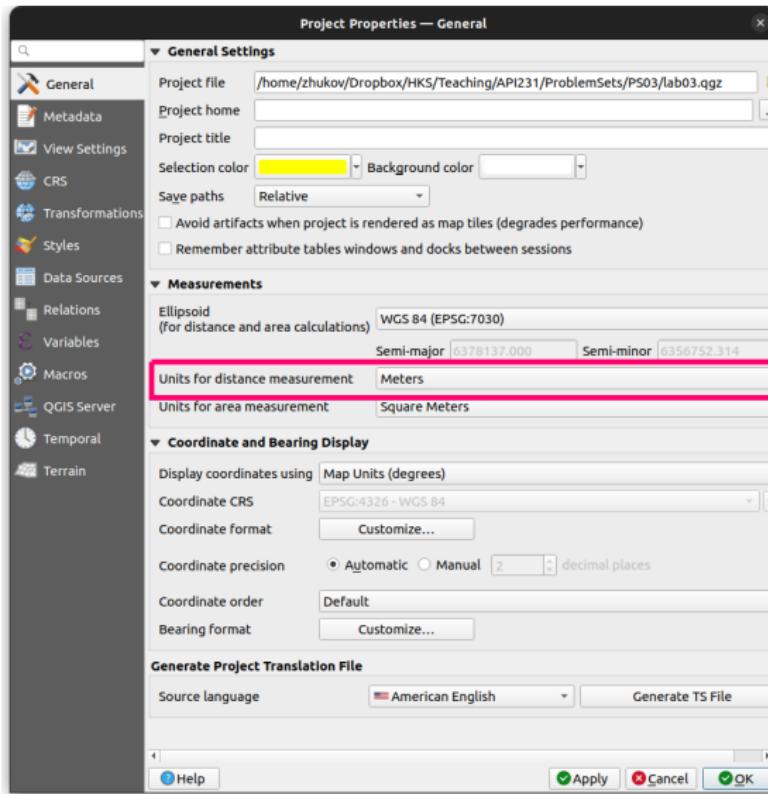
NYC_CD_3 — Features Total: 71, Filtered: 71, Selected: 0

	shape_leng	Pop_boro_c	Pop_Borou_c	Pop_cd_c	Pop_cd_c	Pop_cd_c	Pop_nar_c	Pop_pop_1c	Pop_pop_1c	Pop_pop_1c	Pop_pop_1c	Pop_pop_2c	Pop_pop_2c	crashes	crash1000	bk_length	bk_count
1	38232.886...	3	Brooklyn	8	Crown Hei...	121821	88796	96400	96076	96317	1274	13.227	25999.853...	280			
2	210918.17...	4	Queens	14	The Rocka...	98228	100592	100596	106686	114978	400	3.479	34645.022...	535			
3	31358.450...	2	Bronx	4	Highbridge...	144207	114312	119962	139563	146441	820	5.6	20138.531...	402			
4	122141.65...	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	21	NULL	14172.794...	107			
5	29443.048...	2	Bronx	5	University ...	121807	107995	118435	128313	128200	724	5.647	15096.169...	320			
6	65789.792...	3	Brooklyn	13	Coney Islan...	97750	100030	102596	106120	104278	527	5.054	18317.398...	313			
7	51534.144...	3	Brooklyn	11	Bensonhur...	170119	155072	149994	172129	181981	1117	6.138	11343.990...	134			
8	227190.54...	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	5	NULL	19093.918...	129			
9	44781.826...	3	Brooklyn	10	Bay Ridge, ...	129822	118187	110612	122542	124491	803	6.45	45523.898...	502			
10	105822.37...	4	Queens	10	Ozone Park...	113857	105651	107768	127274	122396	517	4.224	11655.278...	107			
11	32721.097...	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	404	NULL	16844.883...	280			
12	33500.037...	2	Bronx	3	Morrisania,...	150636	53635	57162	68574	79762	572	7.171	17305.319...	275			
13	35875.709...	2	Bronx	6	East Tremo...	114137	65016	68061	75688	83268	608	7.302	14486.147...	227			
14	32820.482...	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	12	NULL	6145.2135...	79			
15	47338.739...	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	2	NULL	51.460121...	3			

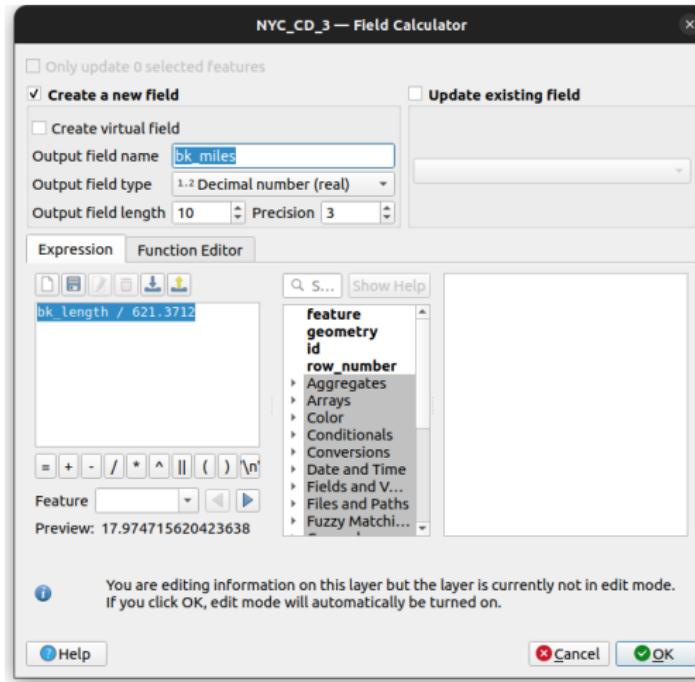
We can look up the project's units of measurement by going to Project menu → Properties



In Project Properties, we see Units for distance measurement = Meters



Let's convert bk_length from meters to miles. Go back to NYC_CD_3's Attribute Table → Field Calculator. Create new field called bk_miles of type Decimal number (real). For the Expression, write bk_length / 621.3712. Click OK



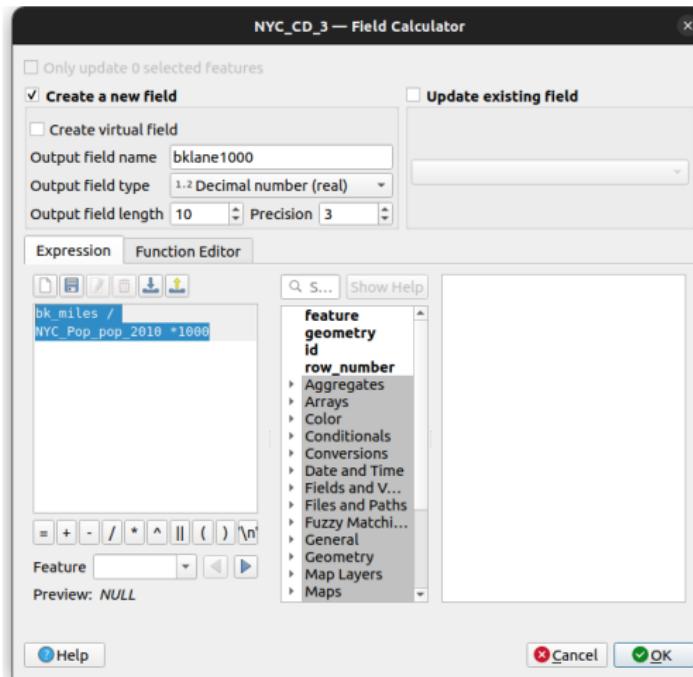
You should now see the new variable in the attribute table.

NYC_CD_3 — Features Total: 71, Filtered: 71, Selected: 0

The screenshot shows the QGIS attribute table for the 'NYC_CD_3' layer. The table has 71 features. A red box highlights the 'bk_miles' column header. The columns include: boro_cd, Pop_boro_cd, C_Pop_Boro, C_Pop_cd, coi_CD, C_Pop_cd_nar, C_Pop_pop_1, C_Pop_pop_15, C_Pop_pop_15C, C_Pop_pop_2, C_Pop_pop_20, crashes, crash1000, bk_length, bk_count, and bk_miles. The 'bk_miles' column contains values such as 41.843, 55.756, 32.410, etc.

	boro_cd	Pop_boro_cd	C_Pop_Boro	C_Pop_cd	coi_CD	C_Pop_cd_nar	C_Pop_pop_1	C_Pop_pop_15	C_Pop_pop_15C	C_Pop_pop_2	C_Pop_pop_20	crashes	crash1000	bk_length	bk_count	bk_miles
1	3 Brooklyn	8 Crown Hel...	121821	88796	96400	96076	96317	1274	13.227	25999.853...	280	41.843				
2	4 Queens	14 The Rocka...	98228	100592	100596	106686	114978	400	3.479	34645.022...	535	55.756				
3	2 Bronx	4 Highbridge...	144207	114312	119962	139563	146441	820	5.6	20138.531...	402	32.410				
4	NULL NULL	NULL NULL	NULL	NULL	NULL	NULL	NULL	NULL	21	NULL	14172.794...	107	22.809			
5	2 Bronx	5 University ...	121807	107995	118435	128313	128200	724	5.647	15096.169...	320	24.295				
6	3 Brooklyn	13 Coney Islan...	97750	100030	102596	106120	104278	527	5.054	18317.398...	313	29.479				
7	3 Brooklyn	11 Bensonhur...	170119	155072	149994	172129	181981	1117	6.138	11343.990...	134	18.256				
8	NULL NULL	NULL NULL	NULL	NULL	NULL	NULL	NULL	NULL	5	NULL	19093.918...	129	30.729			
9	3 Brooklyn	10 Bay Ridge, ...	129822	118187	110612	122542	124491	803	6.45	45523.898...	502	73.264				
10	4 Queens	10 Ozone Park...	113857	105651	107768	127274	122396	517	4.224	11655.278...	107	18.757				
11	NULL NULL	NULL NULL	NULL	NULL	NULL	NULL	NULL	NULL	404	NULL	16844.883...	280	27.109			
12	2 Bronx	3 Morrisania,...	150636	53635	57162	68574	79762	572	7.171	17305.319...	275	27.850				
13	2 Bronx	6 East Tremo...	114137	65016	68061	75688	83268	608	7.302	14486.147...	227	23.313				
14	NULL NULL	NULL NULL	NULL	NULL	NULL	NULL	NULL	NULL	12	NULL	6145.2135...	79	9.890			
4																

Now let's create a per-capita measure (miles of bike lane per 1000 residents). Go back to the Field Calculator. Create new field called bklane1000 of type Decimal number (real). For the Expression, write bk_miles / NYC_Pop_pop_2010 * 1000. Click OK



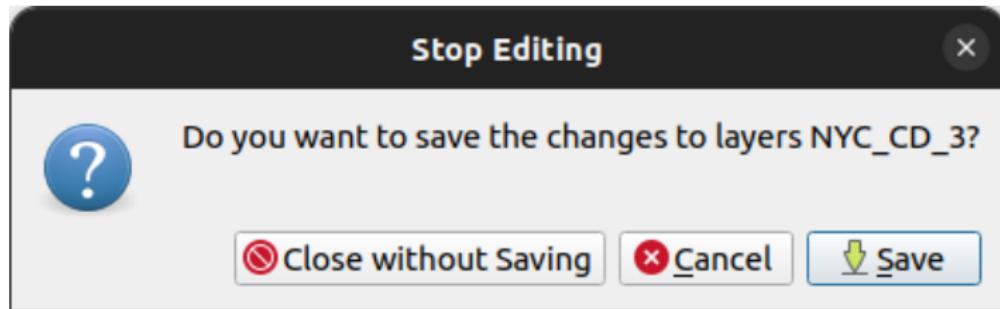
You should now see the new variable in the attribute table.

NYC_CD_3 — Features Total: 71, Filtered: 71, Selected: 0

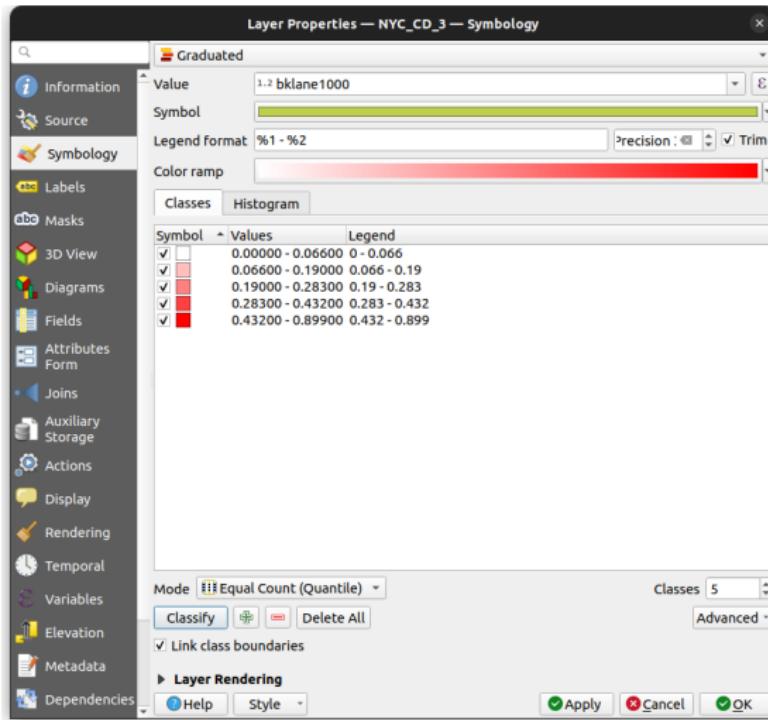
The screenshot shows the QGIS attribute table for the 'NYC_CD_3' layer. The table has 15 columns and 15 rows of data. The columns are: boro_cd, C_Pop_Borou..., C_Pop_cd_cir..., C_Pop_cd_nar..., C_Pop_pop_15..., C_Pop_pop_15..., C_Pop_pop_15..., C_Pop_pop_20..., crashes, crash1000, bk_length, bk_count, bk_miles, and bklane1000. The 'bklane1000' column is highlighted with a red box. The data includes various New York City boroughs and their corresponding statistics. The last row shows a total or summary value for all entries.

	boro_cd	C_Pop_Borou...	C_Pop_cd_cir...	C_Pop_cd_nar...	C_Pop_pop_15...	C_Pop_pop_15...	C_Pop_pop_15...	C_Pop_pop_20...	crashes	crash1000	bk_length	bk_count	bk_miles	bklane1000
1	Brooklyn	8 Crown He...	121821	88796	96400	96076	96317	1274	13.227	25999.853...	280	41.843	0.434	
2	Queens	14 The Rocka...	98228	100592	100596	106686	114978	400	3.479	34645.022...	535	55.756	0.485	
3	Bronx	4 Highbridge...	144207	114312	119962	139563	146441	820	5.6	20138.531...	402	32.410	0.221	
4	NULL	NULL NULL	NULL	NULL	NULL	NULL	NULL	21	NULL	14172.794...	107	22.809	NULL	
5	Bronx	5 University ...	121807	107995	118435	128313	128200	724	5.647	15096.169...	320	24.295	0.19	
6	Brooklyn	13 Coney Islan...	97750	100030	102596	106120	104278	527	5.054	18317.398...	313	29.479	0.283	
7	Brooklyn	11 Bensonhur...	170119	155072	149994	172129	181981	1117	6.138	11343.990...	134	18.256	0.1	
8	NULL	NULL NULL	NULL	NULL	NULL	NULL	NULL	5	NULL	19093.918...	129	30.729	NULL	
9	Brooklyn	10 Bay Ridge, ...	129822	118187	110612	122542	124491	803	6.45	45523.898...	502	73.264	0.589	
10	Queens	10 Ozone Park...	113857	105651	107768	127274	122396	517	4.224	11655.278...	107	18.757	0.153	
11	NULL	NULL NULL	NULL	NULL	NULL	NULL	NULL	404	NULL	16844.883...	280	27.109	NULL	
12	Bronx	3 Morrisania,...	150636	53635	57162	68574	79762	572	7.171	17305.319...	275	27.850	0.349	
13	Bronx	6 East Tremo...	114137	65016	68061	75688	83268	608	7.302	14486.147...	227	23.313	0.28	
14	NULL	NULL NULL	NULL	NULL	NULL	NULL	NULL	12	NULL	6145.2135...	79	9.890	NULL	
15														

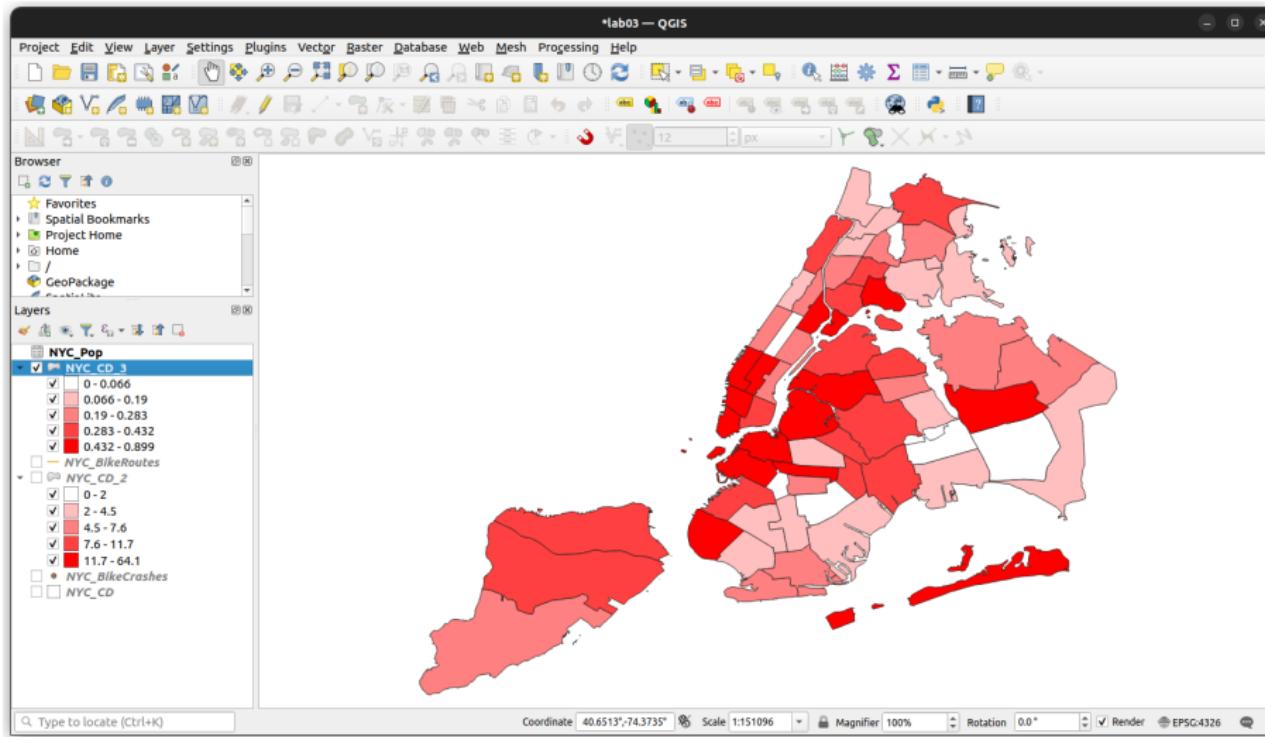
Remember to save your changes!



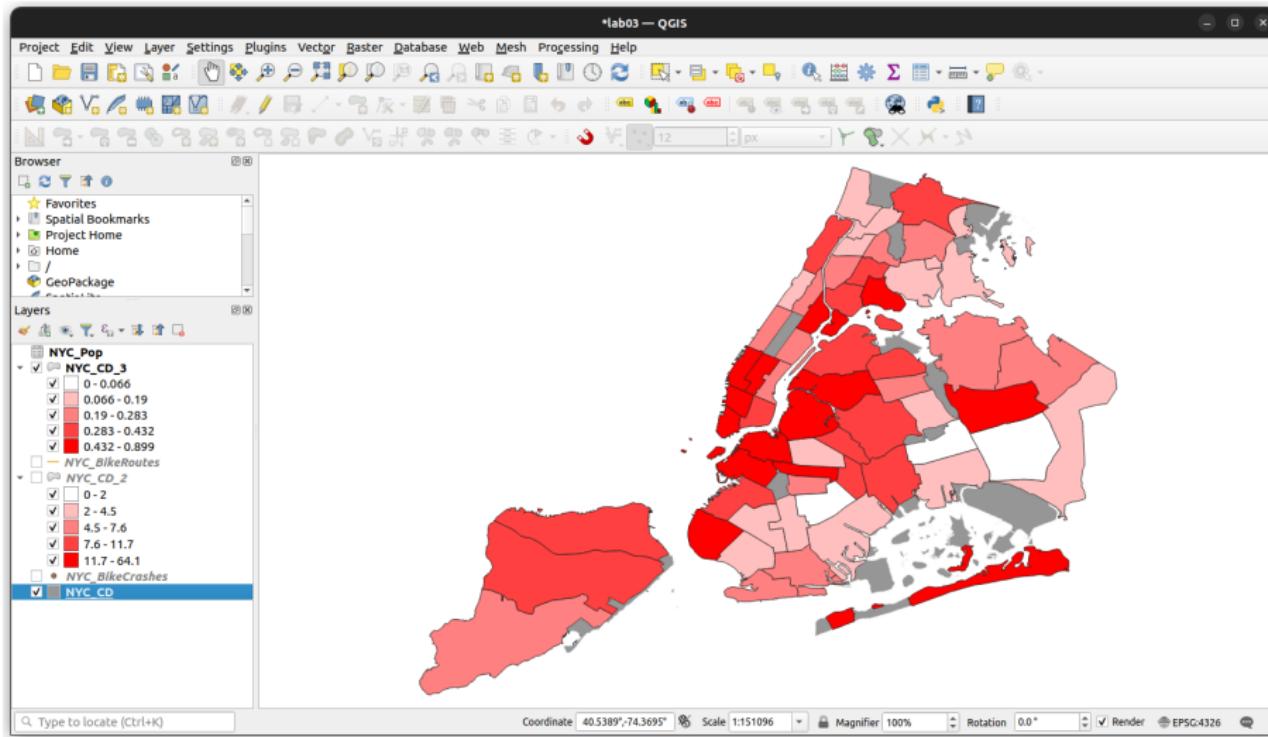
Change the symbology of NYC_CD_3, to color districts by bklane1000. Click Classify, then OK



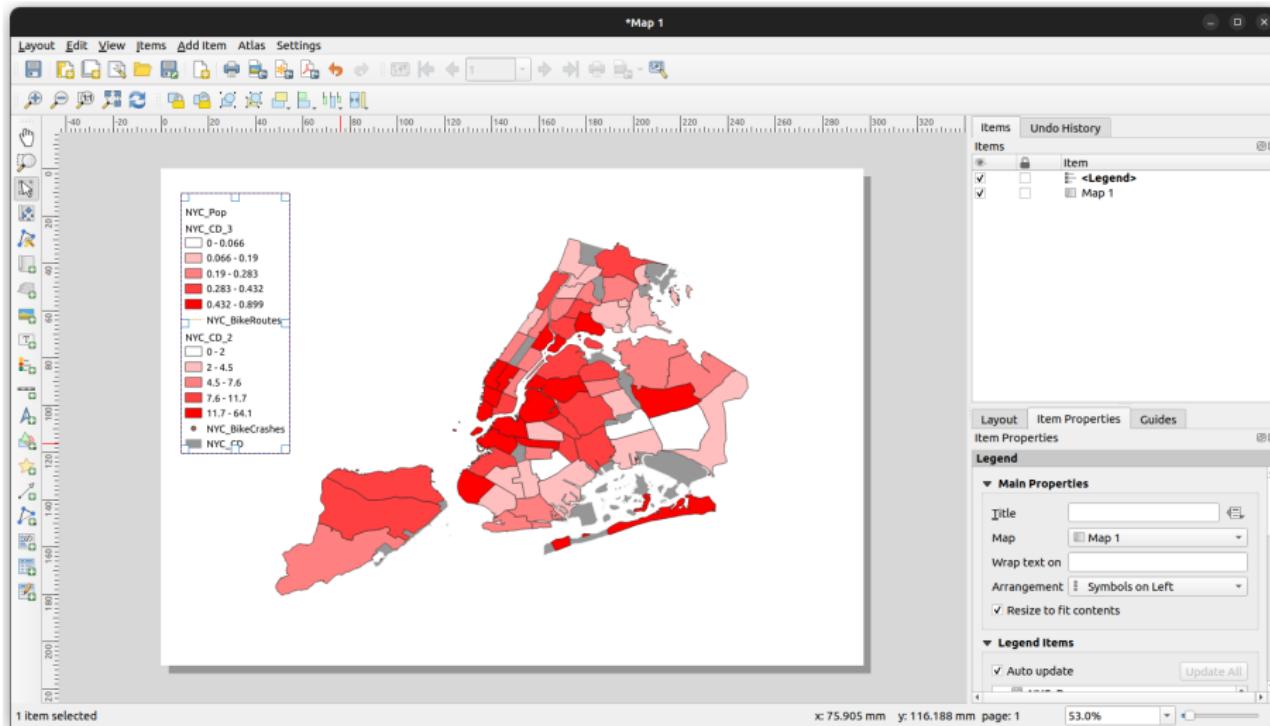
We now have a plot of “miles of bike lane per 1000 residents” (bklane1000). Again, there are a few missing districts (due to no population data)



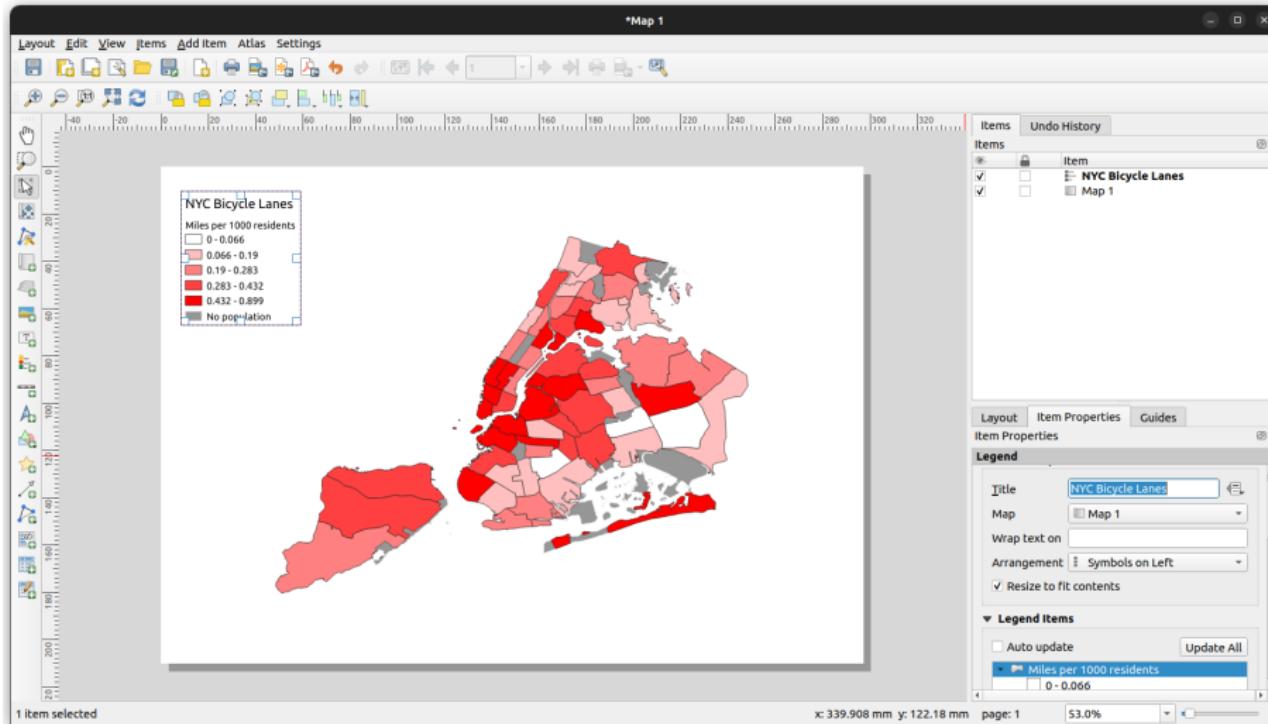
Let's un-hide the NYC_CD layer again to show the missing districts. Ideally, these should be colored gray or some other neutral color



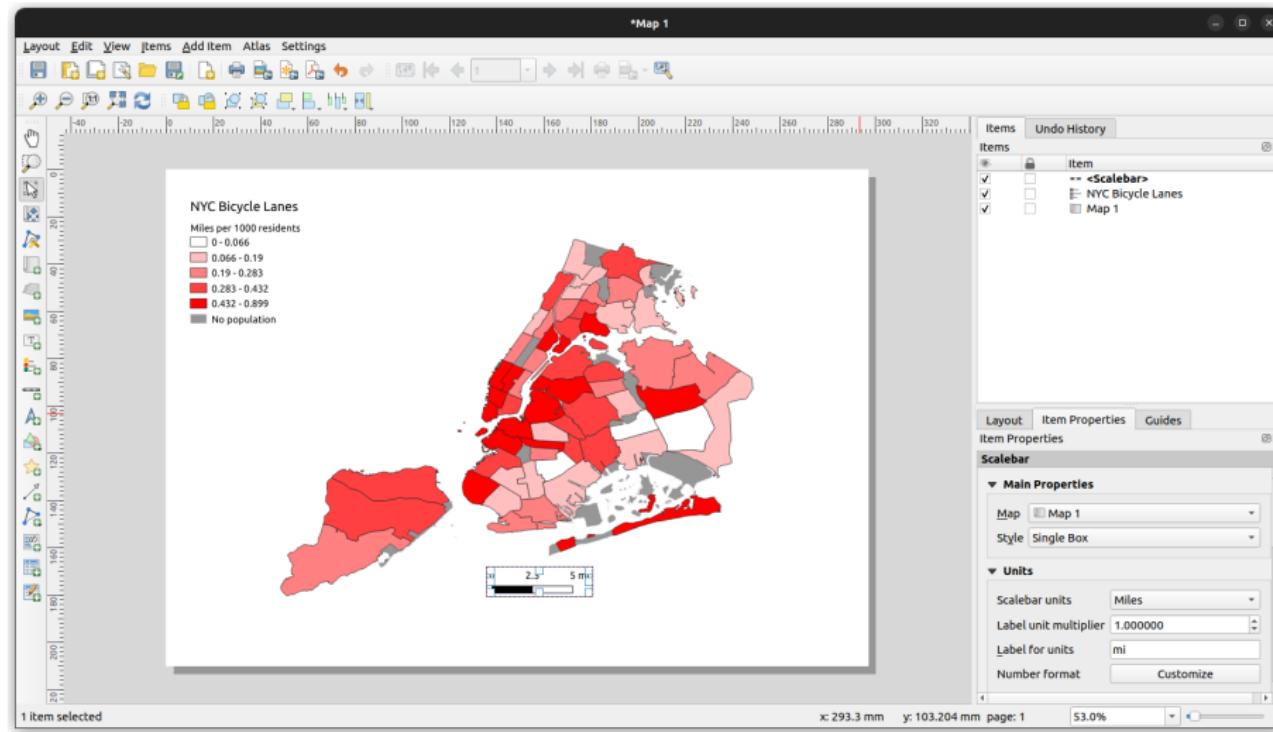
Export the map to image. Project → New Print Layout. Place map and legend to layout. The resulting legend includes several items we'll need to remove (everything but NYC_CD_3 and NYC_CD)



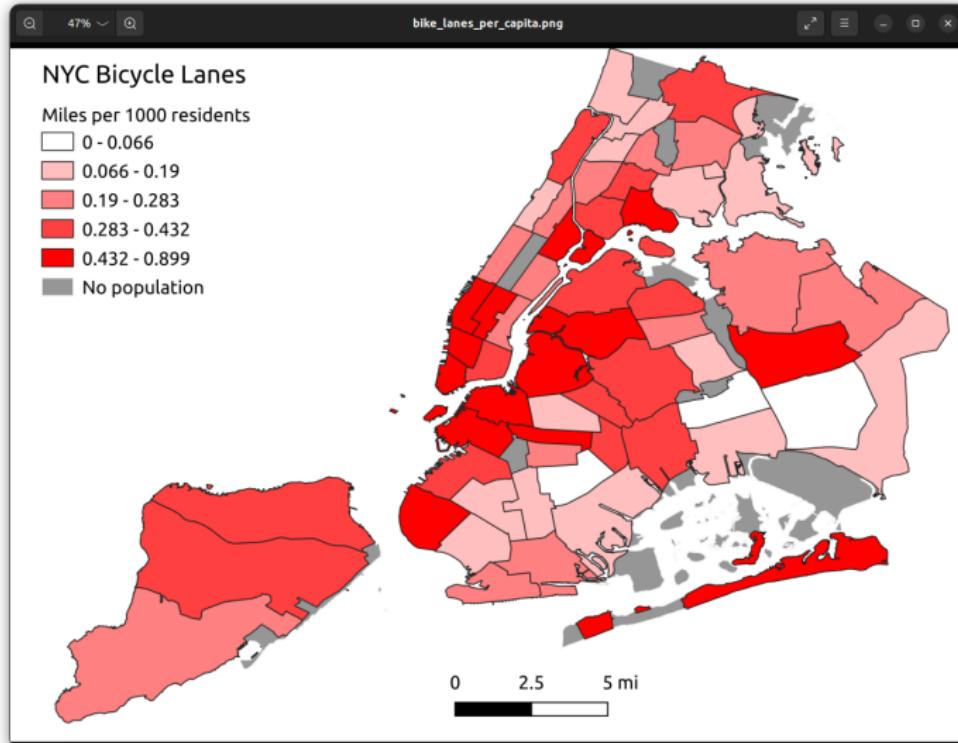
Remember how we did this in lab 2? (Item Properties → un-check Auto update, etc.) Remove everything but NYC_CD_3 and NYC_CD, change layers' names from NYC_CD_3 to “Miles per 1000 residents” and NYC_CD to “No population”.



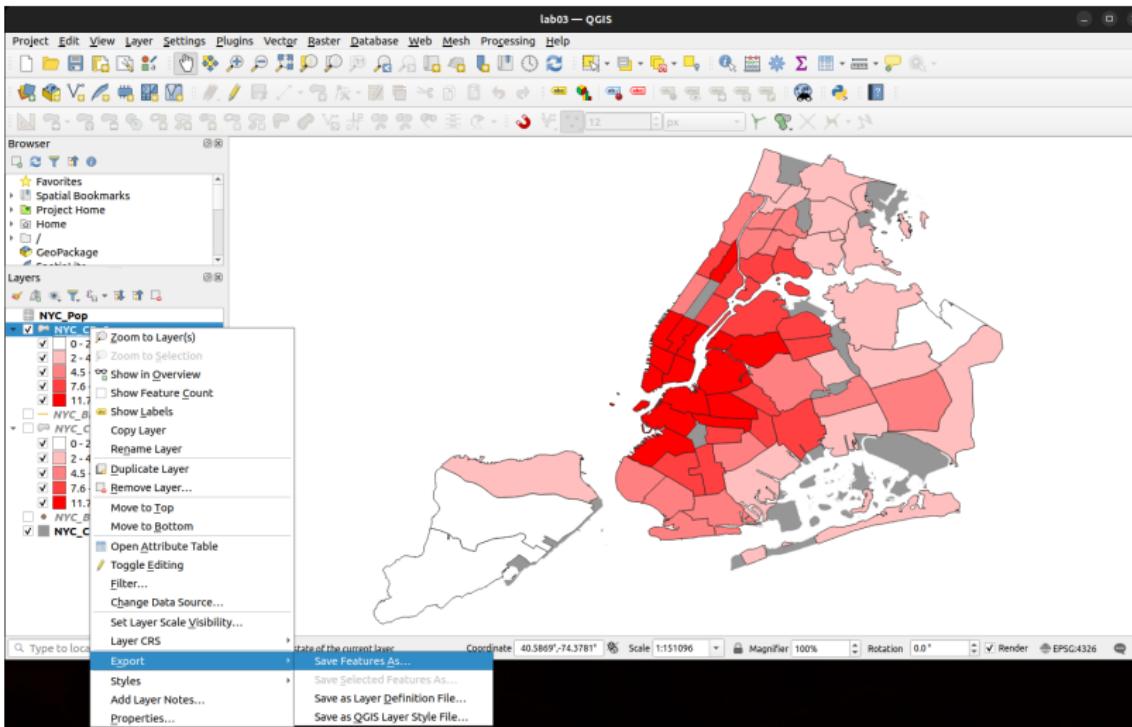
Change legend title to “NYC Bicycle Lanes”. Add scale bar. Change Scalebar units to Miles. Export as image. Name the file `bike_lanes_per_capita.png`



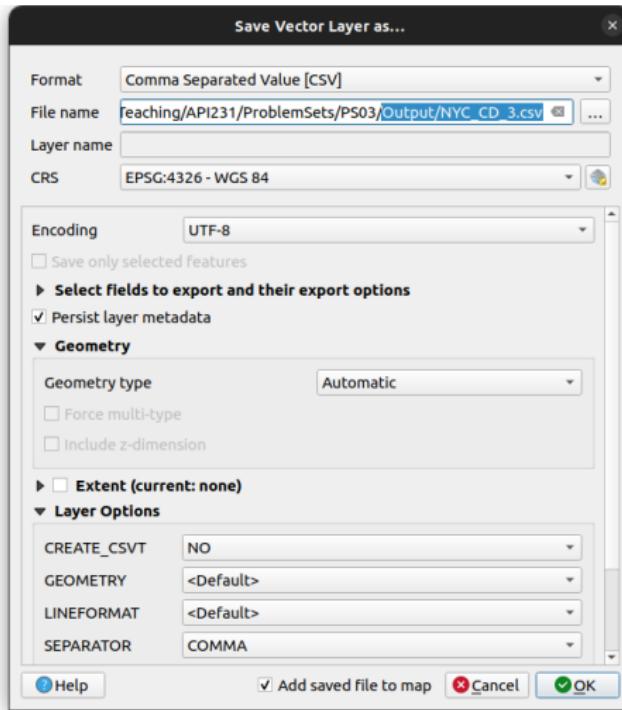
The output file should look roughly like this.



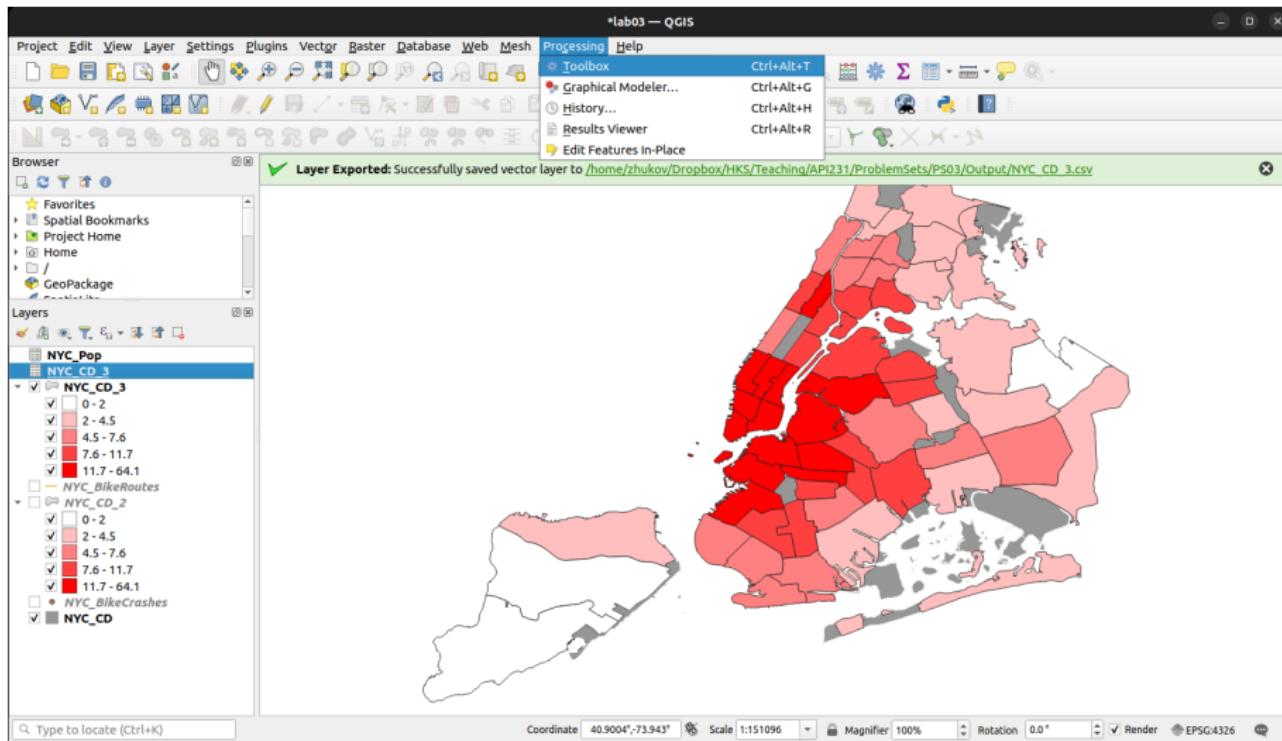
We can export the dataset we created, by right-clicking on NYC_CD_3 layer, and selecting Export → Save Features As...



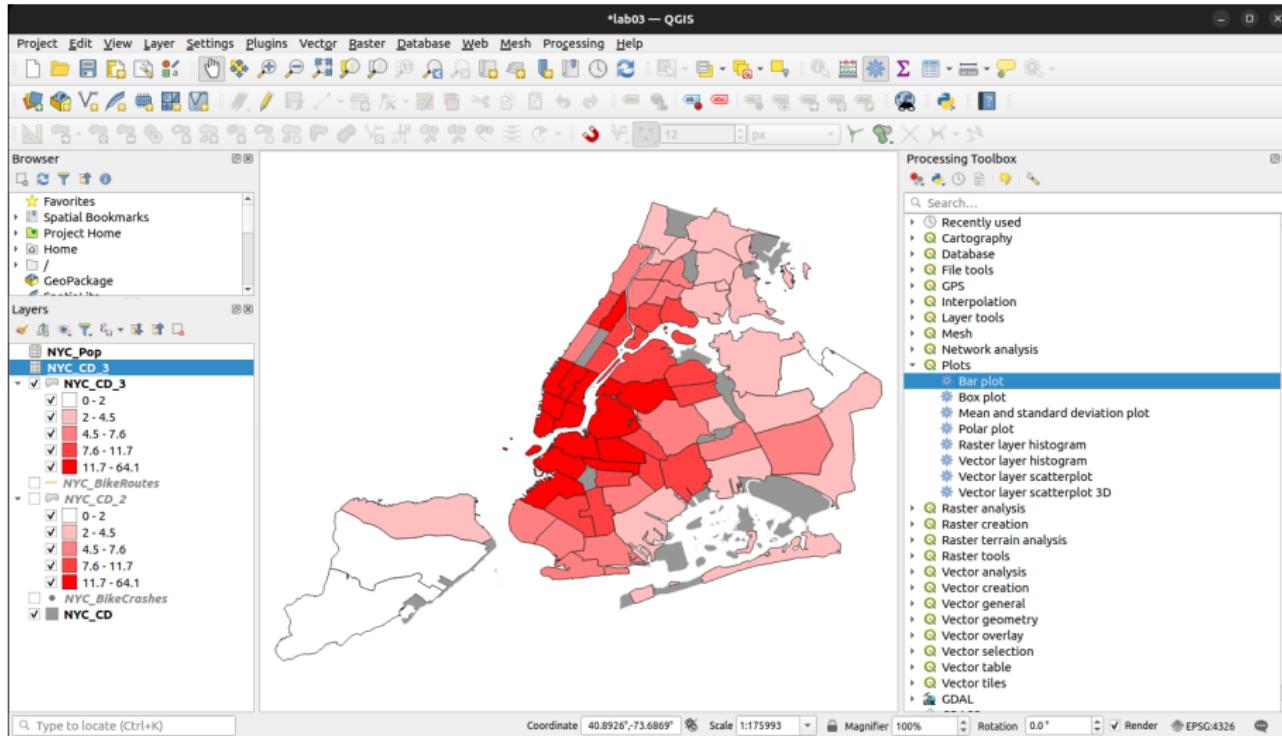
The file can be saved in a variety of formats, including .geojson, .shp, .csv and Excel



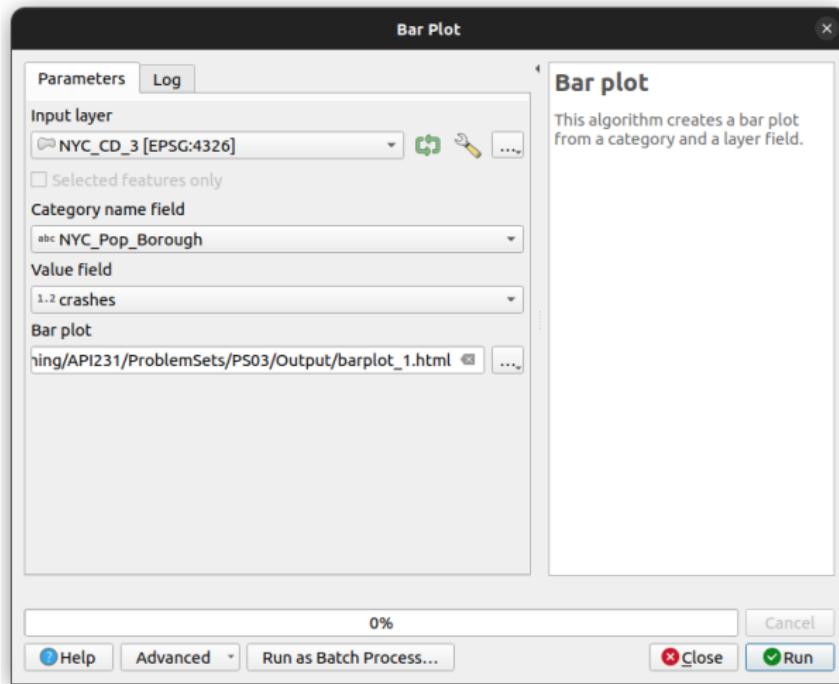
Let's create a bar plot. Go to Processing menu → Toolbox. This will open a Processing Toolbox menu on the right



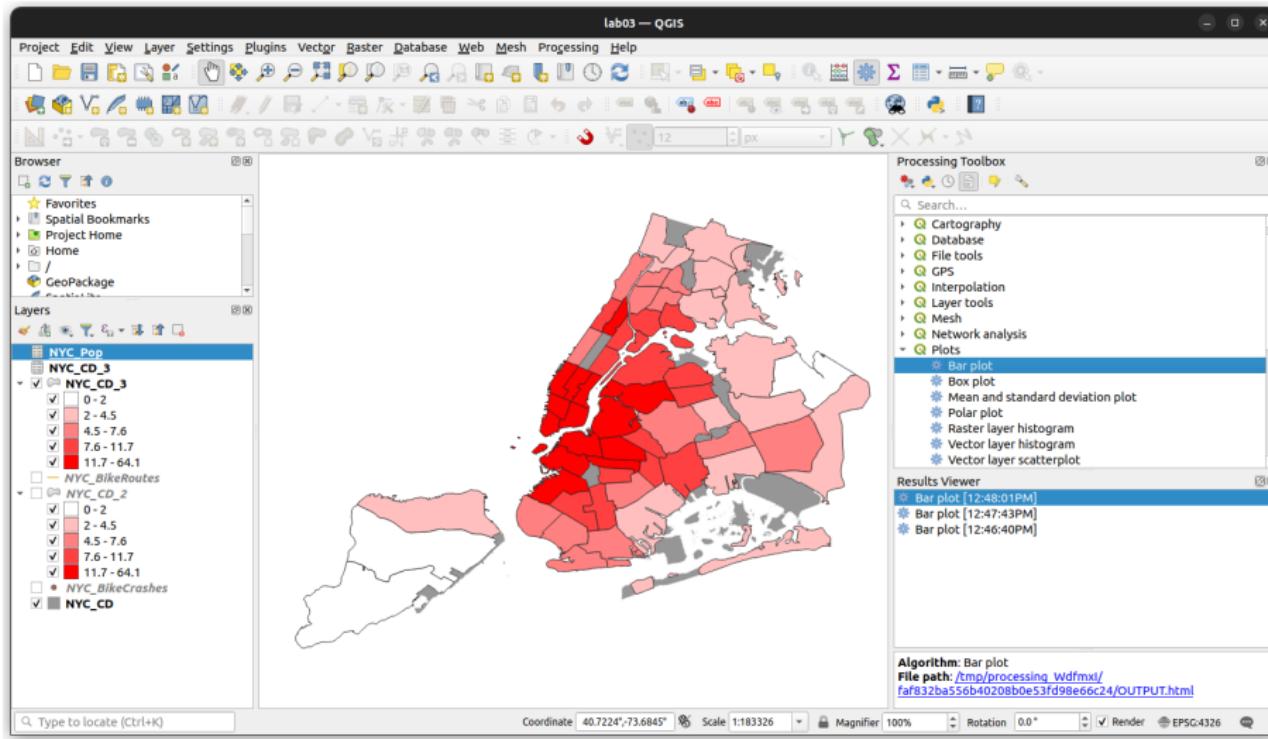
Double-click on Bar plot in the Plots submenu on the right



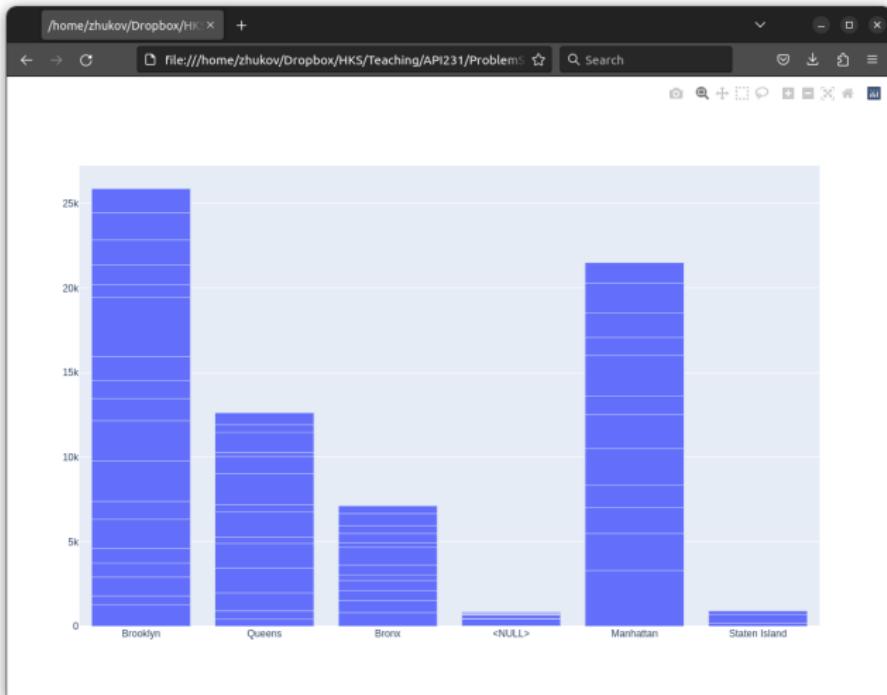
Set Input layer = NYC_CD_3, Category name field = NYC_Pop_Borough, and Value field = crashes. Set the output destination under Bar plot to barplot_1.html. Click Run



You can access the resulting plot by clicking on the Bar plot item in the Results Viewer on the right



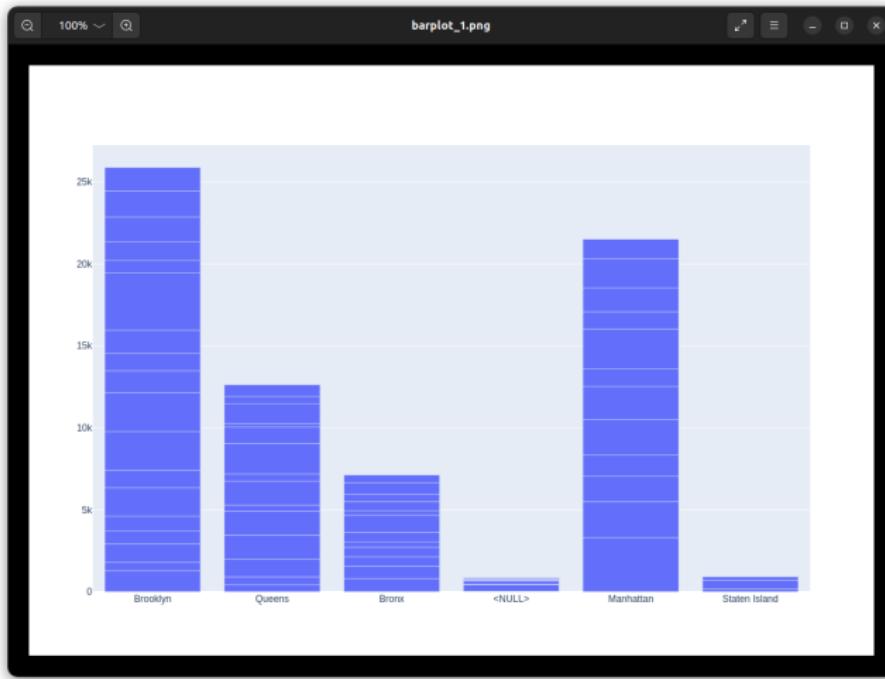
This will open a web browser window, with a bar plot of total bike crashes per district, grouped by borough. As expected, Brooklyn and Manhattan lead the pack.



Export the plot by clicking on the Download plot as png button in the top-right



The output file should look like this. For further analysis, you can also open the .csv file we just created in Excel or R, and start crunching numbers (see below)



Problem Set 3

Your assignment (if using QGIS):

- create a map of rat activity (per 1000 residents) in NYC!
- use this dataset:
 - NYC_Rats/NYC_Rats.csv
- follow the same steps as for “Map 1” above
 - (i.e. sum via point-in-polygon, then calculate per capita rate)
- name the file
 - rats_per_capita.png
- upload map to Canvas
 - (by next Wednesday)

NYC Rodent Inspections (2001-2023)

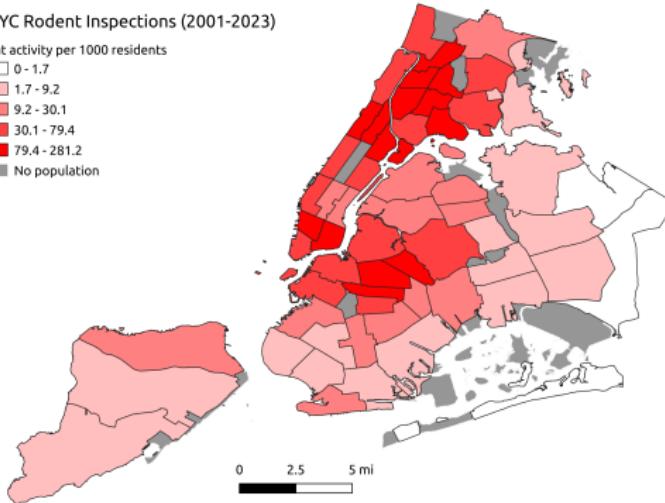
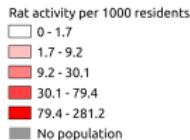


Figure 13: Can you make this map?

R

Loading R packages

To implement these steps in R, we will be using the `sf` package (again)

```
library(sf)
```

NOTE: The code to produce Map 1 and Map 2 in R is in `ps03_demo.R` on RStudio Cloud, and in `PS03.zip` (posted on Canvas).

Map 1

Loading spatial data

Let's load the *NYC community district boundaries* into R, using `sf::read_sf()`:

```
nyc_cd = sf::read_sf("Data/NYC_Communities/NYC_CD.geojson")
plot(nyc_cd["geometry"])
```



Loading non-spatial data

Load the tabular *population data* using `read.csv()`, and preview the first few rows:

```
pop = read.csv(file = "Data/NYC_Population/NYC_Pop.csv")
head(pop)
```

```
##   boro_code Borough cd_code          cd_name pop_1970
## 1          2 Bronx     1 Melrose, Mott Haven, Port Morris 138557
## 2          2 Bronx     2           Hunts Point, Longwood 99493
## 3          2 Bronx     3 Morrisania, Crotona Park East 150636
## 4          2 Bronx     4      Highbridge, Concourse Village 144207
## 5          2 Bronx     5 University Hts., Fordham, Mt. Hope 121807
## 6          2 Bronx     6           East Tremont, Belmont 114137
##   pop_1980 pop_1990 pop_2000 pop_2010
## 1    78441    77214    82159    91497
## 2    34399    39443    46824    52246
## 3    53635    57162    68574    79762
## 4   114312   119962   139563   146441
## 5   107995   118435   128313   128200
## 6    65016    68061    75688    83268
```

Joining spatial to non-spatial data

Similar to QGIS's Field Calculator, let's create a new `boro_cd` variable in `pop`:

```
pop$boro_cd = pop$boro_code*100 + pop$cd_code
```

We can join `nyc_cd` and `pop` by attribute `boro_cd` using the `merge()` command:

```
nyc_cd_2 = merge(x = nyc_cd, y = pop, by = "boro_cd")
print(nyc_cd_2, n=2) # Preview first 2 rows
```

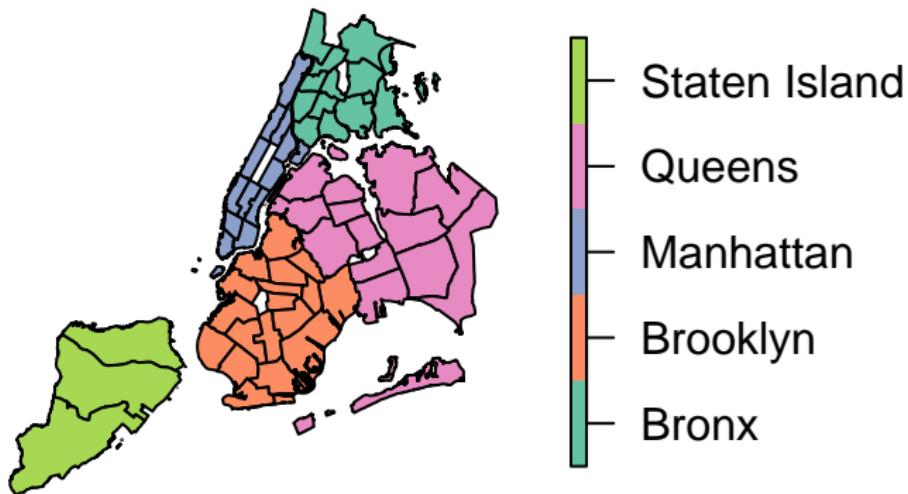
```
## Simple feature collection with 59 features and 12 fields
## Geometry type: MULTIPOLYGON
## Dimension: XY
## Bounding box: xmin: -74.25559 ymin: 40.49613 xmax: -73.70001 ymax: 40.91553
## Geodetic CRS: WGS 84
## First 2 features:
##   boro_cd     shape_area    shape_leng boro_code Borough cd_code
## 1      101 41692711.6252 69093.8528032          1 Manhattan      1
## 2      102 37604933.3646 32628.8675927          1 Manhattan      2
##           cd_name pop_1970 pop_1980 pop_1990 pop_2000 pop_2010
## 1 Battery Park City, Tribeca    7706    15918    25366    34420    60978
## 2 Greenwich Village, Soho    84337    87069    94105    93119    90016
##           geometry
## 1 MULTIPOLYGON (((-74.04388 4...
## 2 MULTIPOLYGON (((-73.99684 4...
```

Plotting merged features

Using the merged dataset `nyc_cd_2`, let's plot the districts by borough

```
plot(nyc_cd_2["Borough"])
```

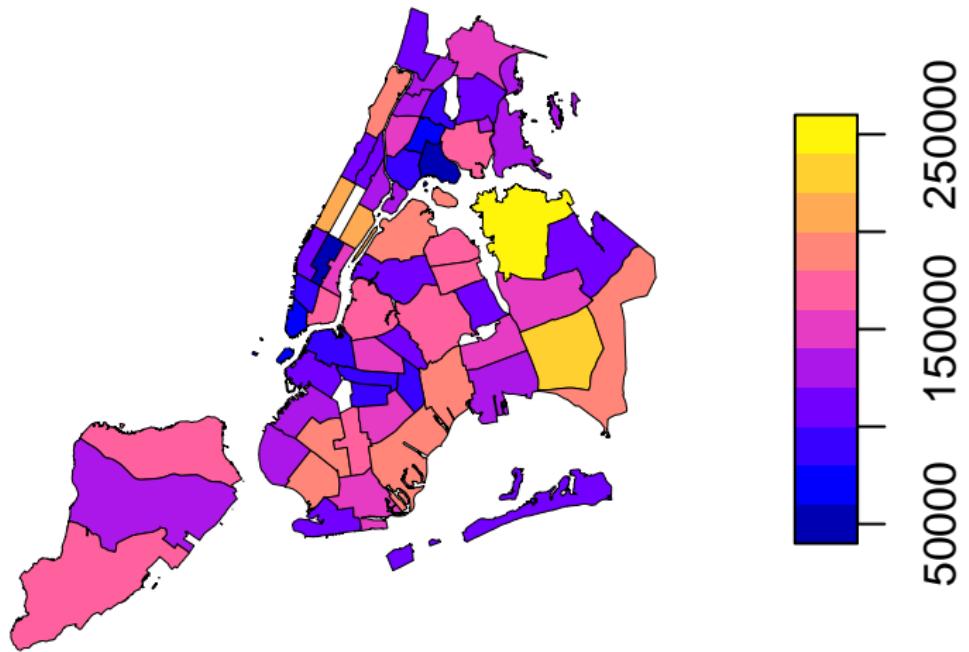
Borough



... or plot them by population size in 2010:

```
plot(nyc_cd_2["pop_2010"],main="Population size in 2010")
```

Population size in 2010



Loading spatial data from a CSV table

Our data on *bike collisions*, like population, is also a csv table:

```
crashes = read.csv(file = "Data/NYC_Collisions/NYC_BikeCrashes.csv")
```

But this file has geographic coordinates, which we can use to create a spatial object, borrowing the CRS from nyc_cd:

```
crashes = sf::st_as_sf(x = crashes, coords = c("LONGITUDE","LATITUDE"),  
                      crs = sf::st_crs(nyc_cd))  
plot(crashes["geometry"])
```



Point-in-polygon analysis

To count the number of collisions in each community district, we can overlay `nyc_cd_2` and `crashes` with `sf::st_intersects()`. The output is a classed list of feature IDs intersected:

```
o = sf::st_intersects(x = nyc_cd_2, y = crashes)
```

```
o
```

```
## Sparse geometry binary predicate list of length 59, where the predicate
## was `intersects'
## first 10 elements:
##  1: 15, 29, 188, 212, 236, 268, 284, 297, 625, 634, ...
##  2: 9, 10, 11, 18, 108, 135, 137, 169, 194, 274, ...
##  3: 79, 85, 165, 215, 242, 257, 260, 261, 330, 347, ...
##  4: 4, 12, 44, 101, 109, 116, 161, 210, 211, 233, ...
##  5: 8, 16, 47, 66, 75, 83, 105, 120, 149, 191, ...
##  6: 21, 58, 97, 227, 266, 280, 282, 307, 323, 331, ...
##  7: 76, 81, 110, 117, 164, 220, 229, 288, 301, 332, ...
##  8: 2, 186, 204, 285, 322, 354, 355, 361, 362, 390, ...
##  9: 89, 103, 129, 150, 286, 311, 349, 475, 529, 540, ...
## 10: 17, 54, 74, 82, 113, 156, 173, 228, 246, 249, ...
```

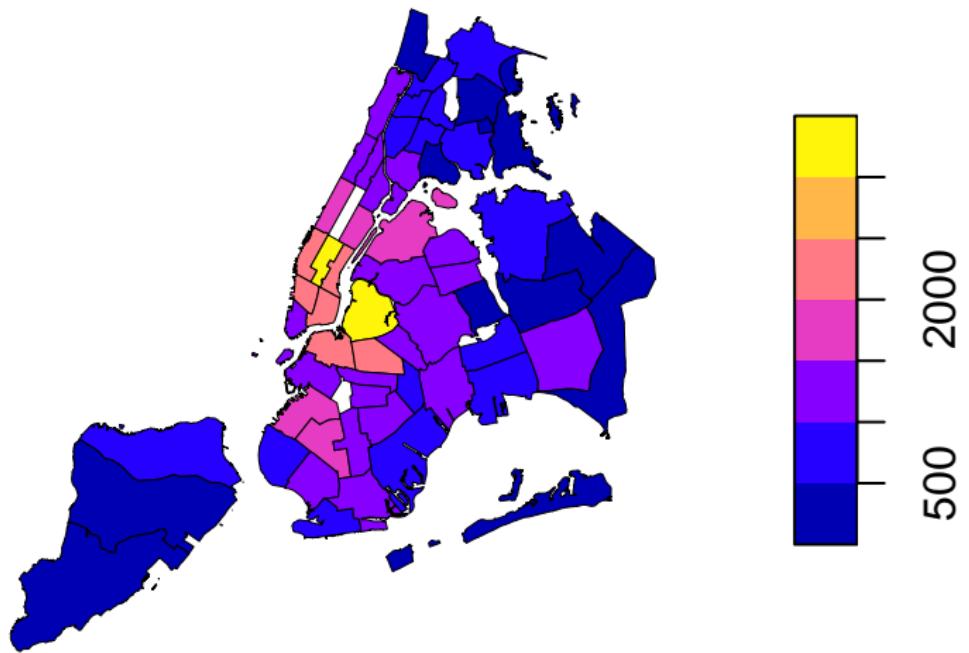
We can use the command `lengths(o)` to find the number of points in each polygon, and assign this as a new variable in `nyc_cd`:

```
nyc_cd_2$crashes = lengths(o)
```

Plot the new variable:

```
plot(nyc_cd_2["crashes"], main = "Bicycle Collisions (2013-2023)")
```

Bicycle Collisions (2013–2023)

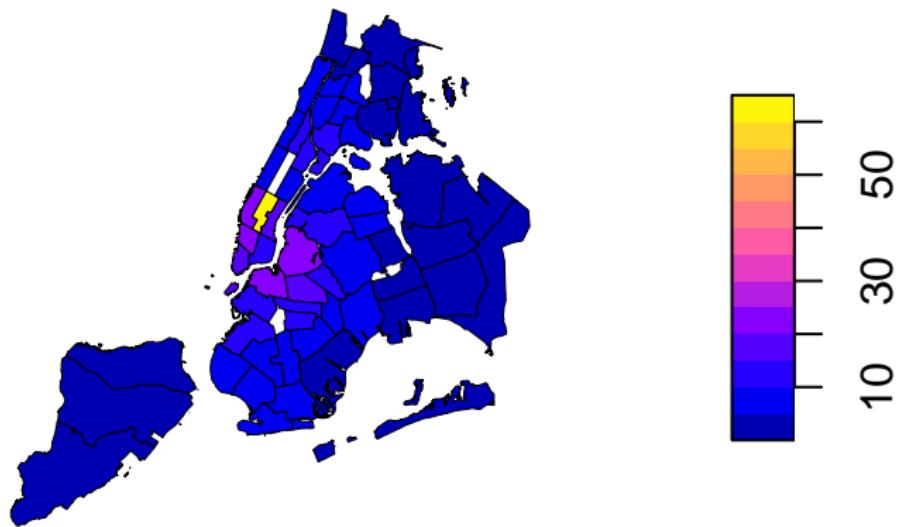


Creating per-capita measures

Let's calculate and plot a new field, crashes per 1000 residents:

```
nyc_cd_2$crash1000 = nyc_cd_2$crashes / nyc_cd_2$pop_2010 * 1000  
plot(nyc_cd_2["crash1000"], main = "Collisions per 1000 Residents")
```

Collisions per 1000 Residents

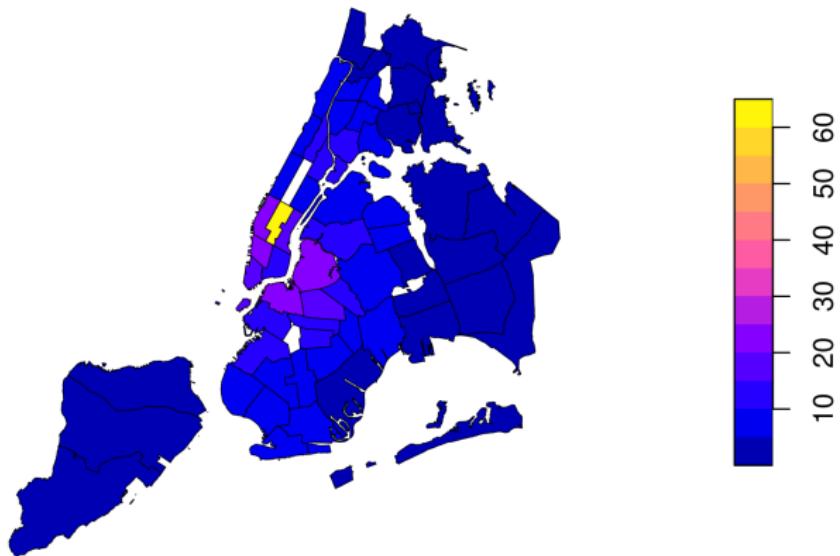


Exporting Map 1 to image file

To save the map, we will use the `png()` and `dev.off()` commands:

```
png(filename = "Output/bike_crashes_per_capita_r.png",
     width = 6, height = 4, units = "in", res = 300)
plot(nyc_cd_2["crash1000"],
      main = "Bicycle Collisions per 1000 Residents", lwd=.5)
dev.off()
```

Bicycle Crashes per 1000 Residents



The output file should look like this.

Map 2

Loading bike lane data

Import bike routes data

```
bikeroutes = sf::read_sf("Data/NYC_BikeRoutes/NYC_BikeRoutes.geojson",
  crs=sf::st_crs(nyc_cd_2))
plot(bikeroutes["geometry"], reset=FALSE, col="forestgreen", lwd=.5)
plot(nyc_cd_2["geometry"], add = TRUE)
```



Line-in-polygon analysis

We need to calculate the *miles of bike lane* in each community district. We can think of this as a two-step routine:

1. determine which line segments lie inside of which polygons
2. take a local sum of overlapping line lengths in each polygon

For step 1, we can use `sf::st_intersects()` to figure out which segments of `bikeroutes` lie in which districts (`nyc_cd_3`) (similar to point-in-polygon analysis):

```
o = sf::st_intersects(x = nyc_cd_2, y = bikeroutes)
```

```
o
```

```
## Sparse geometry binary predicate list of length 59, where the predicate
## was `intersects'
## first 10 elements:
##  1: 6, 35, 88, 100, 138, 139, 140, 171, 805, 806, ...
##  2: 11, 2484, 2485, 2486, 2487, 2488, 2489, 2490, 2491, 2492, ...
##  3: 81, 84, 110, 127, 177, 1505, 1506, 1507, 1508, 1533, ...
##  4: 41, 48, 78, 178, 2633, 2650, 2653, 2654, 2655, 2656, ...
##  5: 12, 47, 109, 2633, 2796, 2836, 2837, 2838, 2839, 2840, ...
##  6: 125, 126, 169, 2904, 2907, 2908, 2912, 2913, 2951, 2952, ...
##  7: 85, 3136, 3137, 3184, 3185, 3186, 3218, 3219, 3220, 3222, ...
##  8: 89, 90, 91, 168, 3546, 3599, 3600, 3601, 3602, 3603, ...
##  9: 3887, 3888, 3897, 3898, 3899, 3900, 3914, 3920, 3921, 3922, ...
## 10: 176, 180, 3867, 3868, 3915, 3916, 3917, 3918, 3919, 3920, ...
```

Line-in-polygon analysis (continued)

For step 2, we can write a loop that extracts the overlapping bike lane segments in each district, measures their (cumulative) length, and converts this sum to miles.

```
out <- c()                      # Create empty vector to store results
for(i in 1:length(o)){           # Open loop
  bike_subset = bikeroutes[o[i][[1]],]      # Take subset
  bk_meters = sum(sf::st_length(bike_subset)) # Sum lengths
  bk_km = bk_meters/1000                     # Convert to km
  bk_mi = bk_km * 0.6213712                  # Convert to miles
  out = c(out, bk_mi)                        # Concatenate
}
# Close loop
```

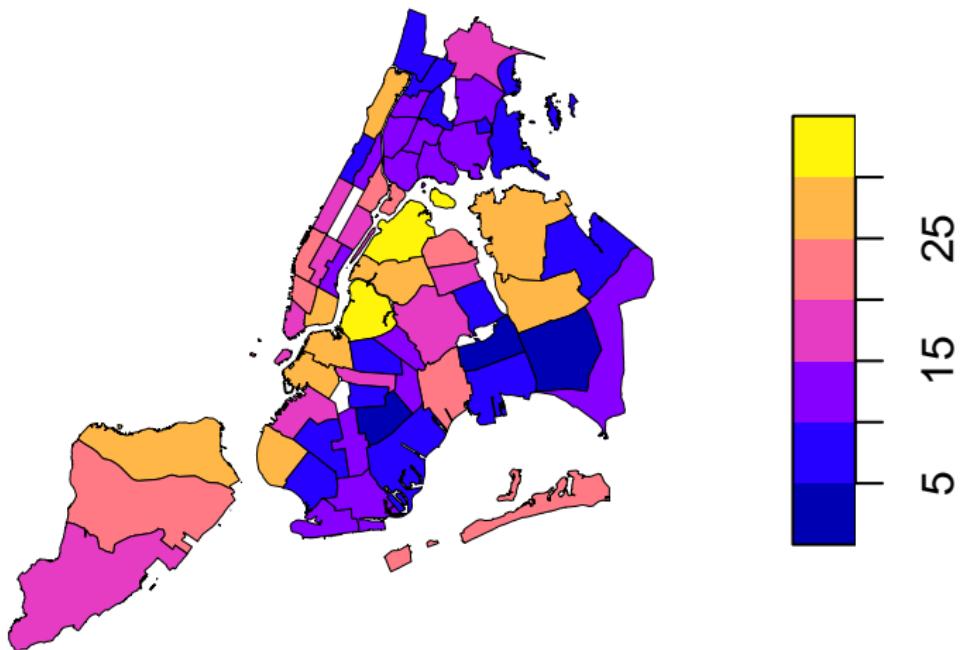
We then add this as a variable to nyc_cd_2

```
nyc_cd_2$bk_miles <- out
```

Let's quickly visualize the results:

```
plot(nyc_cd_2["bk_miles"],  
     main = "Miles of Bike Lane per District", lwd=.5)
```

Miles of Bike Lane per District

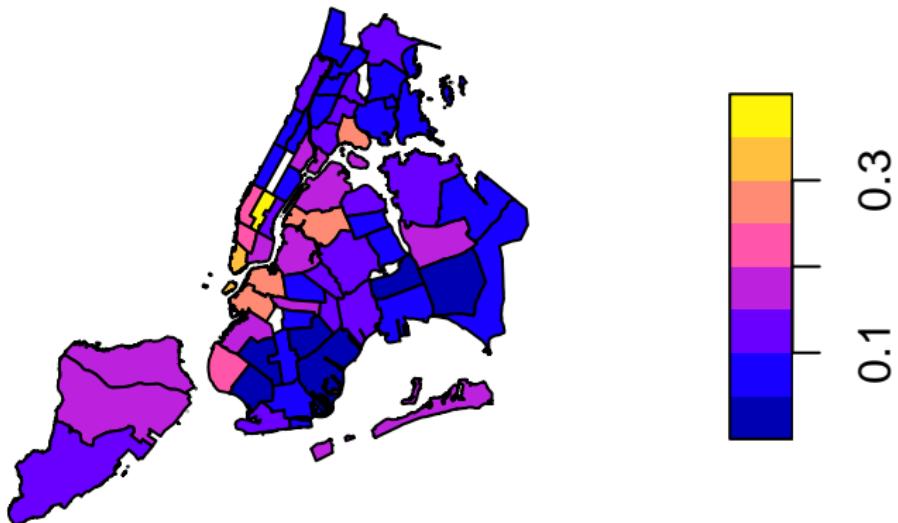


Creating per-capita measures

Let's calculate and plot a new field, miles of bicycle lane per 1000 residents:

```
nyc_cd_2$bklane1000 = nyc_cd_2$bk_miles / nyc_cd_2$pop_2010 * 1000  
plot(nyc_cd_2["bklane1000"], main = "Bike Lane per 1000 Residents (mi)")
```

Bike Lane per 1000 Residents (mi)

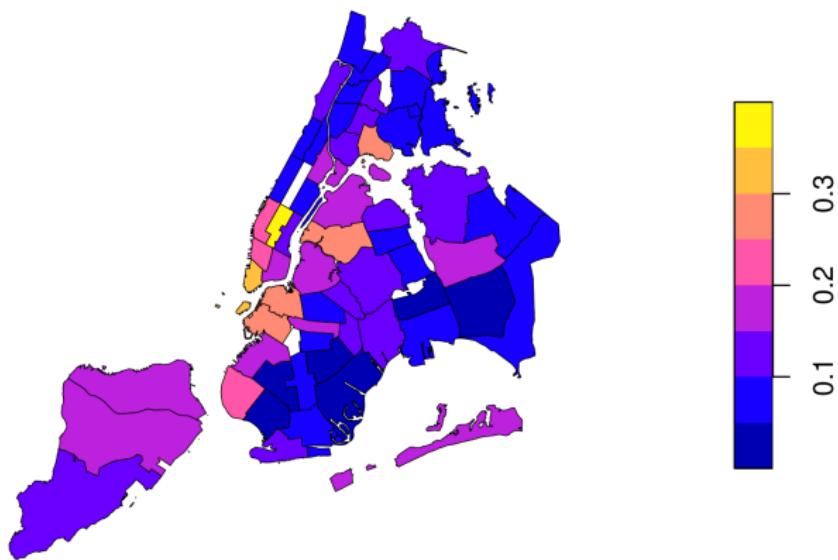


Exporting to image file

To save the map, we will use the `png()` and `dev.off()` commands:

```
png(filename = "Output/bike_lanes_per_capita_R.png",
     width = 6, height = 4, units = "in", res = 300)
plot(nyc_cd_2["bklane1000"],
      main = "Miles of Bicycle Lane per 1000 Residents", lwd=.5)
dev.off()
```

Miles of Bicycle Lane per 1000 Residents



The output file should look like this.

Exporting the dataset

To save the data attribute table as a csv file, we can use the `write.csv()` command:

```
write.csv(x = nyc_cd_2, file = "Output/NYC_CD_2_R.csv")
```

Making bar plots

For each of NYC's five boroughs, calculate averages of the variables we created:

```
borough_means = aggregate(  
  x = nyc_cd_2[,c("crashes","crash1000","bk_miles","bklane1000")],  
  by = list(Borough=nyc_cd_2$Borough),  
  FUN = mean  
)
```

This will create a new sf object with 5 rows:

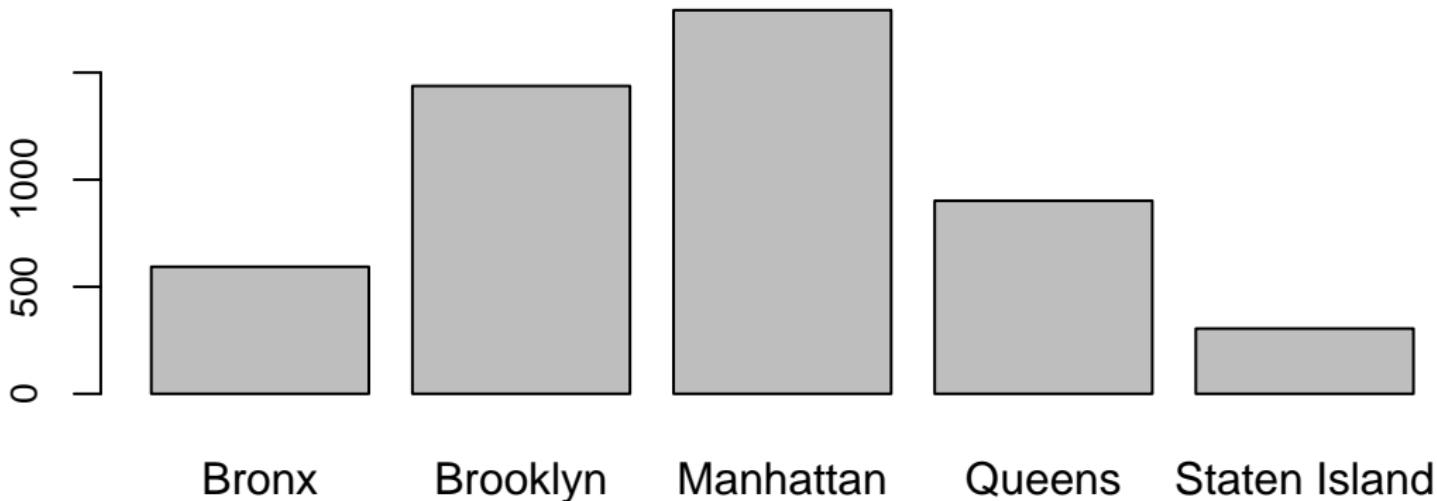
```
borough_means
```

```
borough_means  
  
## Simple feature collection with 5 features and 5 fields  
## Attribute-geometry relationships: aggregate (4), identity (1)  
## Geometry type: MULTIPOLYGON  
## Dimension: XY  
## Bounding box: xmin: -74.25559 ymin: 40.49613 xmax: -73.70001 ymax: 40.91553  
## Geodetic CRS: WGS 84  
  
##      Borough    crashes crash1000 bk_miles bklane1000  
## 1      Bronx  593.2500  5.581747 11.37986  0.1095067  
## 2    Brooklyn 1437.6667 10.709820 15.48712  0.1230569  
## 3   Manhattan 1791.5000 17.508324 18.69915  0.1717709  
## 4     Queens  901.7143  5.632056 16.58475  0.1081192  
## 5 Staten Island  304.6667  1.893391 22.25868  0.1441233
```

Make a bar plot of *bike crashes in an average community district*

```
barplot(height=borough_means$crashes, names.arg=borough_means$Borough,  
        main="Bicycle collisions per community district (mean)")
```

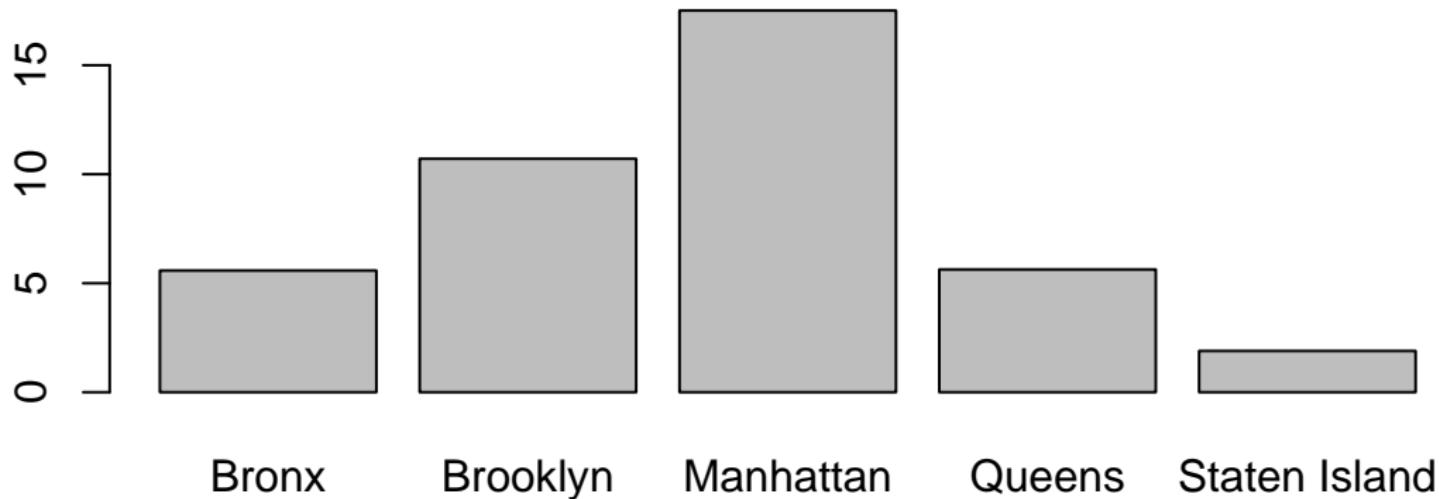
Bicycle collisions per community district (mean)



Make a bar plot of *bike crashes per 1000 residents*

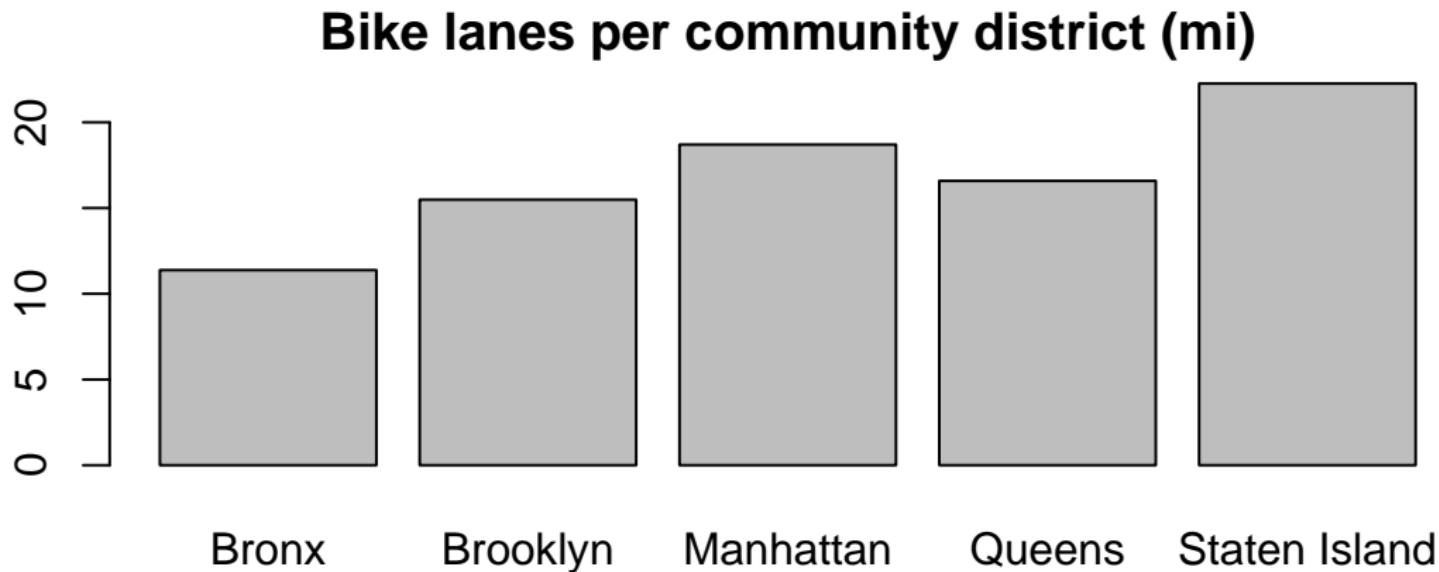
```
barplot(height = borough_means$crash1000,  
        names.arg = borough_means$Borough,  
        main = "Collisions per 1000 residents (mean)")
```

Collisions per 1000 residents (mean)



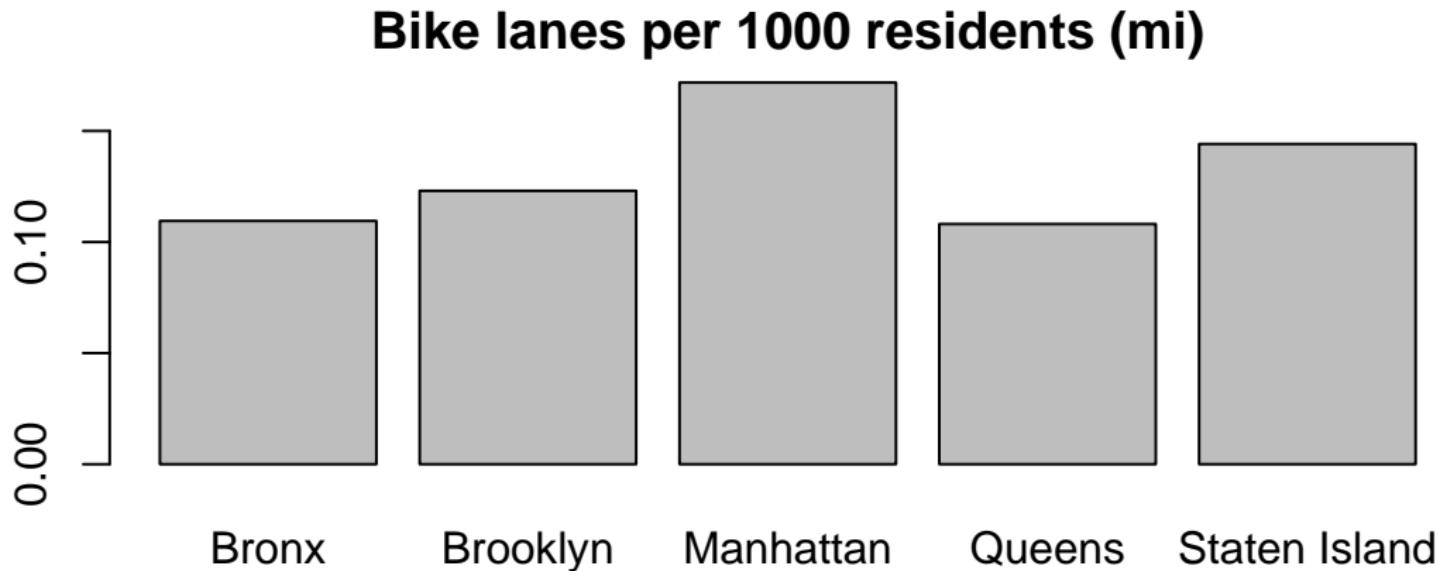
Make a bar plot of *miles of bike lane in an average community district*

```
barplot(height = borough_means$bk_miles,  
        names.arg = borough_means$Borough,  
        main = "Bike lanes per community district (mi)")
```



Make a bar plot of *miles of bike lane per 1000 residents*

```
barplot(height = borough_means$bklane1000,  
        names.arg = borough_means$Borough,  
        main = "Bike lanes per 1000 residents (mi)")
```



Basic regression analysis

You can use an `sf` object as you would a `data.frame` to run regressions in R.

For example, suppose we wanted to see if there are fewer bicycle collisions where there are more bike lanes. We could run a basic linear regression of bicycle collisions (per 1000) on bike lanes (per 1000), with fixed effects for each borough:

```
summary(lm(crash1000 ~ bklane1000 + as.factor(Borough), data=nyc_cd_2))

##
## Call:
## lm(formula = crash1000 ~ bklane1000 + as.factor(Borough), data = nyc_cd_2)
##
## Residuals:
##     Min      1Q  Median      3Q     Max 
## -11.094  -2.211   0.026   1.729  32.999 
##
## Coefficients:
##                               Estimate Std. Error t value Pr(>|t|)    
## (Intercept)              -1.5231    2.1642  -0.704  0.48468  
## bklane1000                 64.8800   11.1212   5.834 3.32e-07 ***
## as.factor(Borough)Brooklyn  4.2489    2.3146   1.836  0.07201 .  
## as.factor(Borough)Manhattan 7.8869    2.6232   3.007  0.00403 ** 
## as.factor(Borough)Queens    0.1403    2.4381   0.058  0.95432  
## as.factor(Borough)Staten Island -5.9343   4.0189  -1.477  0.14570 
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

This analysis suggests there are *more collisions where there are more bike lanes.*

Adjusting for differences across boroughs, each additional mile of bike lane (per 1000 residents) is associated with 65 additional collisions (per 1000).

Does this mean bike lanes are “causing” collisions?

How should we interpret this result?

What are some alternative explanations of the positive correlation?

What sort of data would we need to test these alternative arguments?

Problem Set 3

Your assignment (if using R):

- create a map of rat activity (per 1000 residents) in NYC!
- use this dataset:
 - NYC_Rats/NYC_Rats.csv
- follow the same steps as for “Map 1” (crashes per 1000)
(i.e. sum via point-in-polygon,
then calculate per capita rate)
- name the file
`rats_per_capita.png`
- upload map to Canvas
(by next Wednesday)

Rat Activity per 1000 Residents

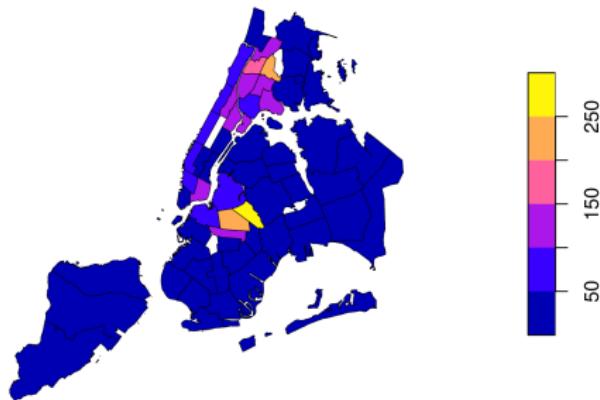


Figure 14: Can you make this map?