API-231 / GIS-PubPol
Meeting 15 (Lab Exercise + Problem Set 8)

Yuri M. Zhukov
Visiting Associate Professor of Public Policy
Harvard Kennedy School

March 21, 2024

Goal: analyze flood risk models in New Orleans

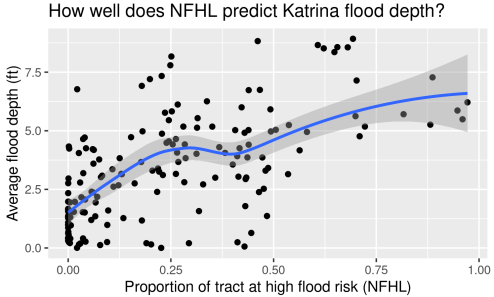

Figure 1: Predicting catastrophic events

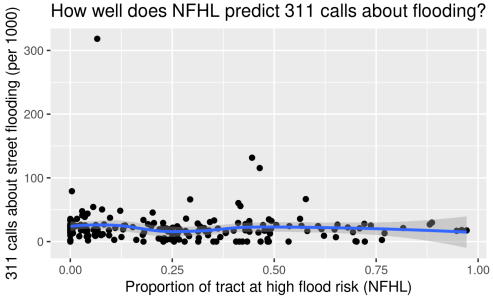

Figure 2: Predicting everyday events

Case study: Hurricane Katrina was a Category 5 storm that hit New Orleans in 2005

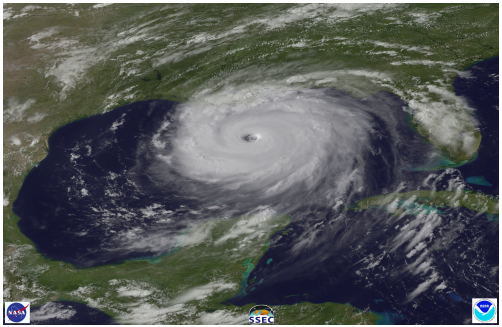

Figure 3: Direct hit

Figure 4: Levees break

It claimed over 1800 lives and over $100 billion in property damage

Figure 5: French Quarter

Figure 6: Lower 9th Ward

We will examine the geographic distribution of post-Katrina flooding in NOLA

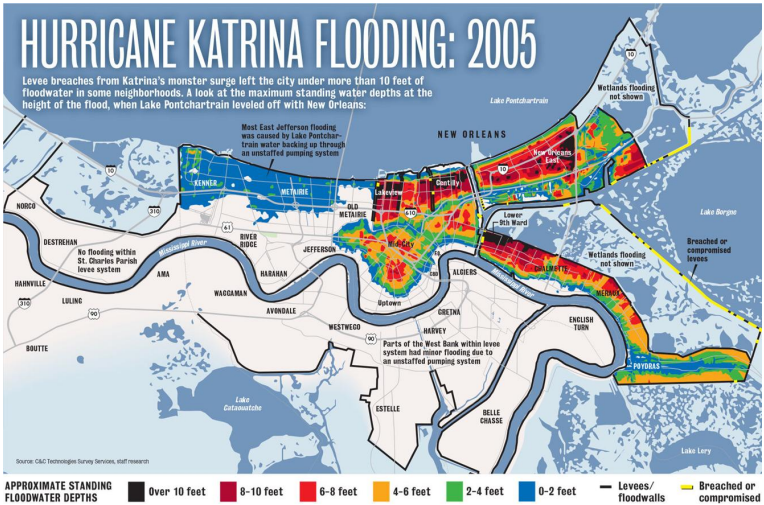

Figure 7: We will use these data, among others

We will look at whether communities of color were more vulnerable to flooding

Figure 8: Legacy of housing discrimination

We'll also evaluate the predictive accuracy of FEMA's National Flood Hazard Layer

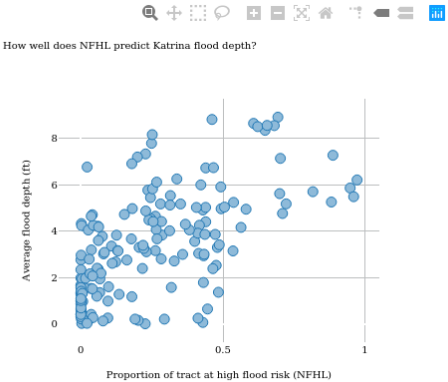

Figure 9: Predicting the past

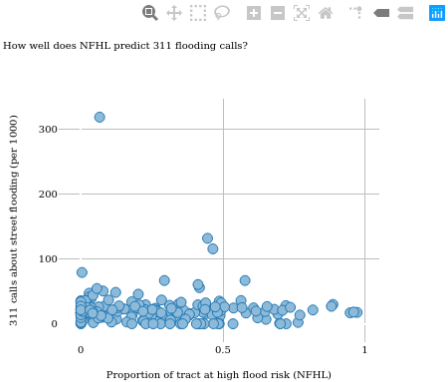

Figure 10: Predicting the everyday

Overview of lab exercise and problem set

1. Lab exercise
   a) Hands-on experience analyzing and editing raster data
   b) Calculate zonal statistics of flood depth for Census tracts
   c) Re-classify and subset flood risk raster data
   d) Integrate flood data with data on housing discrimination, race, 311 calls
2. Problem set
   a) Create statistical graphics (scatterplots) evaluating performance of flood hazard model in predicting:
      - Katrina flood levels
      - per capita 311 calls about street flooding

You can make these plots in QGIS or in R. Instructions for both are below

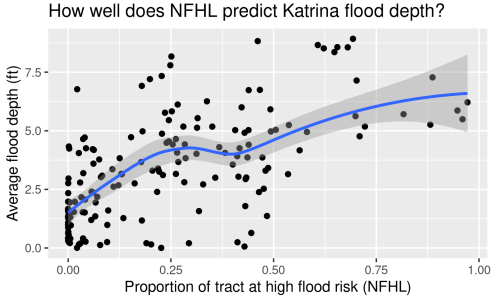

Figure 11: Scatterplot 1 in R

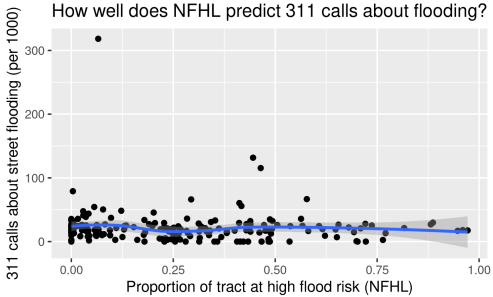

Figure 12: Scatterplot 2 in R

We have five sources of data:

Category	Type	Format	Data source
Hurricane Katrina flood depth	Raster	`.tif`	`NOAA`
2000 Census Tracts	Vector (polygon)	`.geojson`	`IPUMS NHGIS`
HOLC Redlining Maps	Vector (polygon)	`.geojson`	`Mapping Inequality`
National Flood Hazard Layer	Raster	`.tif`	`FEMA`
311 Calls	Table (geocoded)	`.csv`	`New Orleans Open Data`

These are all in the `PS08.zip` file posted on Canvas.

Let's open QGIS...

QGIS

Always save your progress!

Go to `Project` → `Save As...`

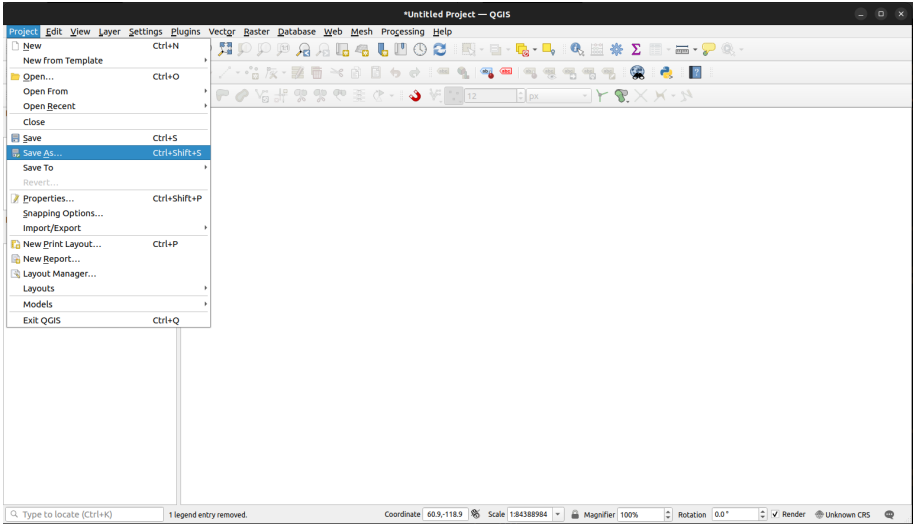

Zonal statistics

Let's begin by *calculating average flood depth in census tracts*.
Load the `CensusTracts2000.geojson` (vector) from the `Data/Census` directory

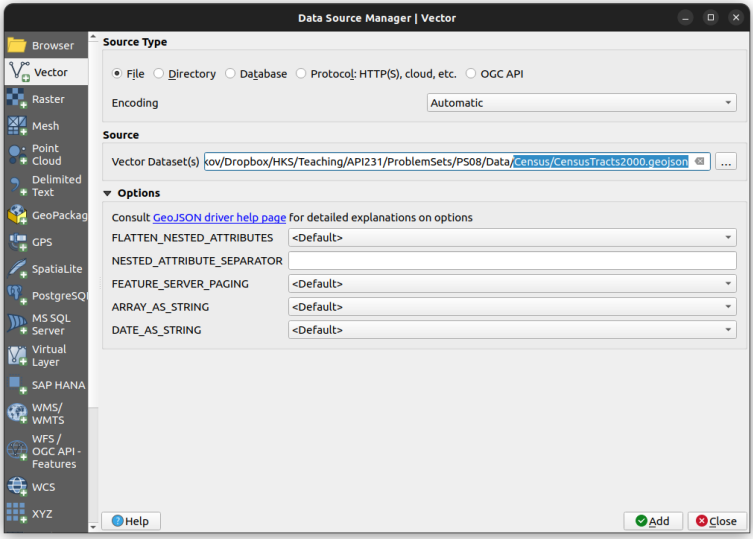

You may receive a "Ballpark Transform Occurred" warning. This sometimes happens when re-projecting from WGS84 (EPSG:4326, QGIS default) to UTM (EPSG:26915)

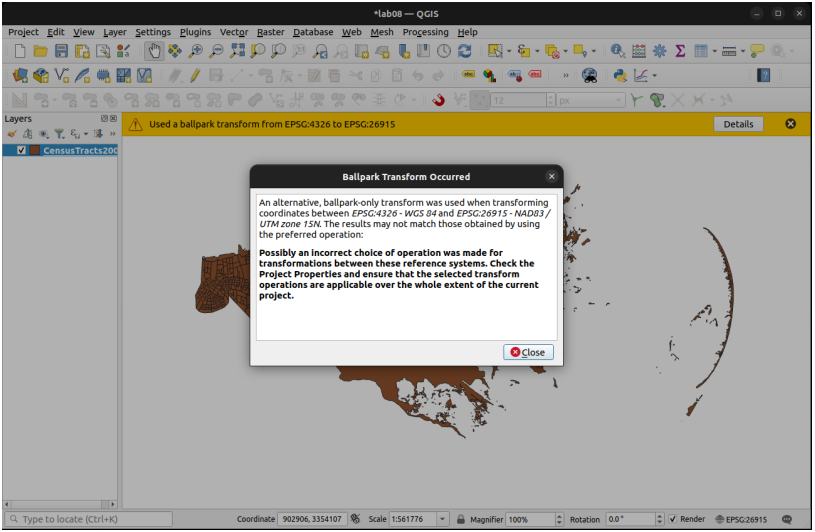

Load the raster `flood_depth.tif` from the `Data/Katrina` directory

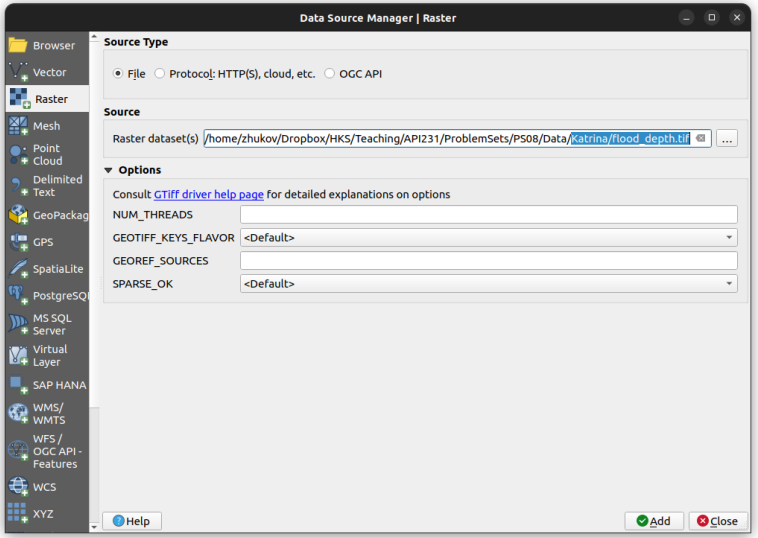

Note that these layers include both New Orleans and neighboring St. Bernard Parish

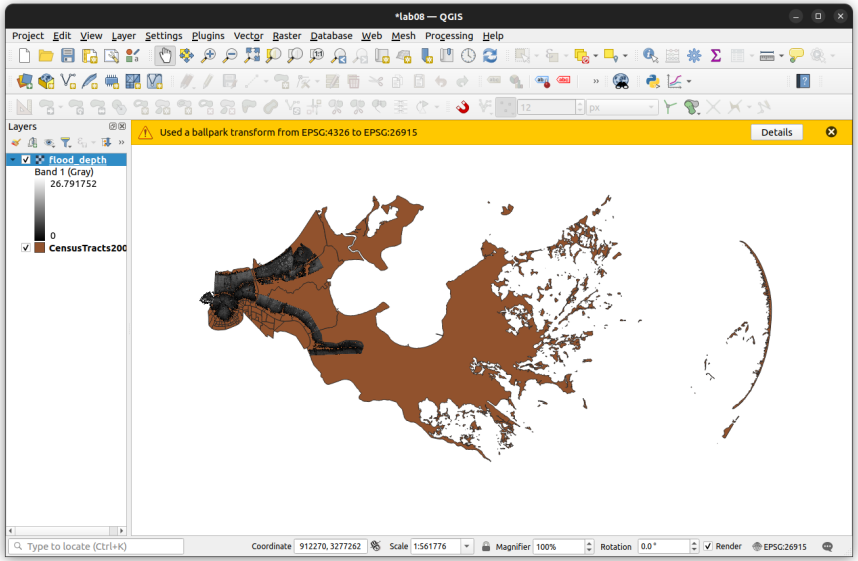

Let's calculate flood stats for each census tract. Open the `Processing Toolbox`

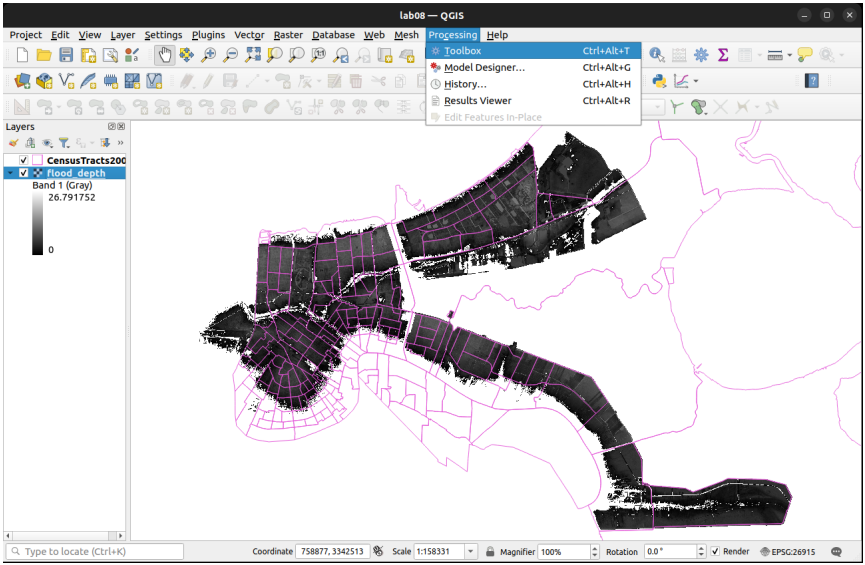

In the Toolbox, open `Raster analysis` submenu → `Zonal statistics` tool

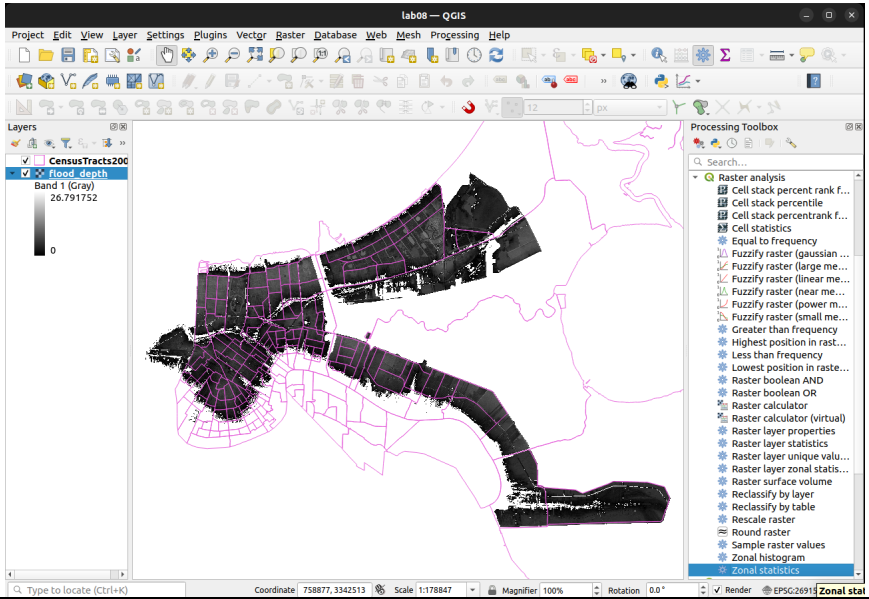

On the next screen, set

- `Input layer` = `CensusTracts2000`
- `Raster layer` = `flood_depth`
- `Output column prefix` = `flood_`

Click the `[...]` box next to `Statistics to calculate`

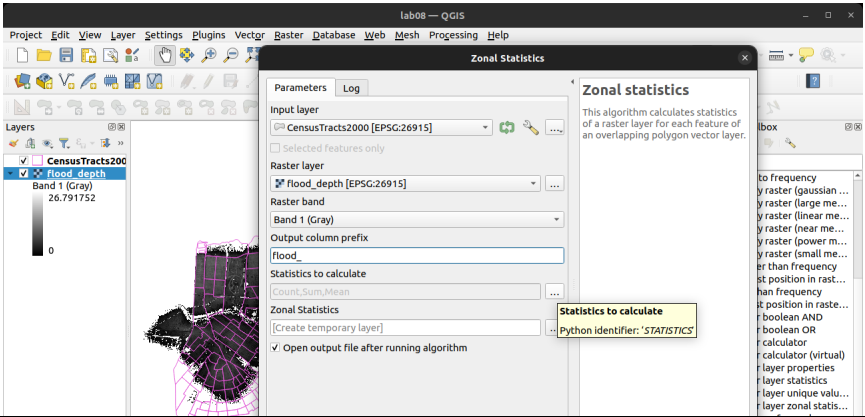

Select ✓ Mean, ✓ Maximum. Click OK

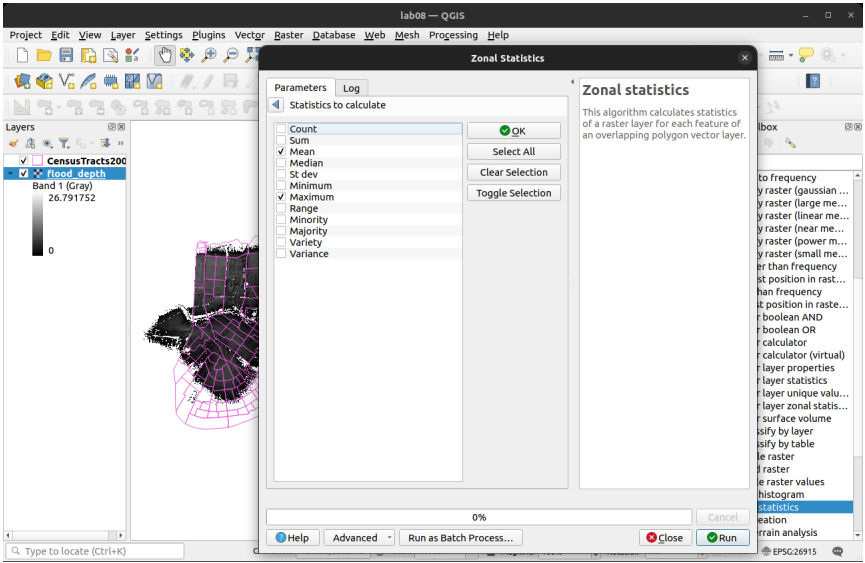

Save the output to file as `tracts_2.geojson`. Click `Run`

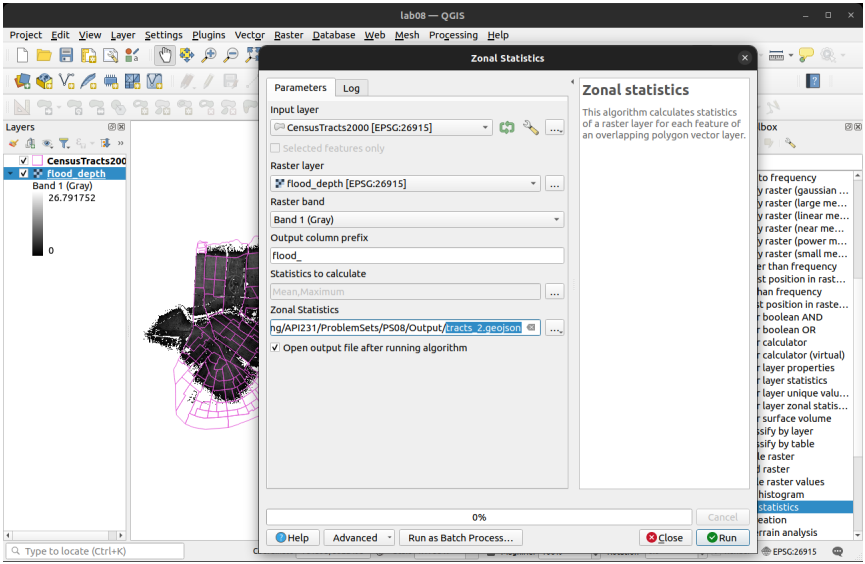

This will add a new layer, `tracts_2`, to your project window

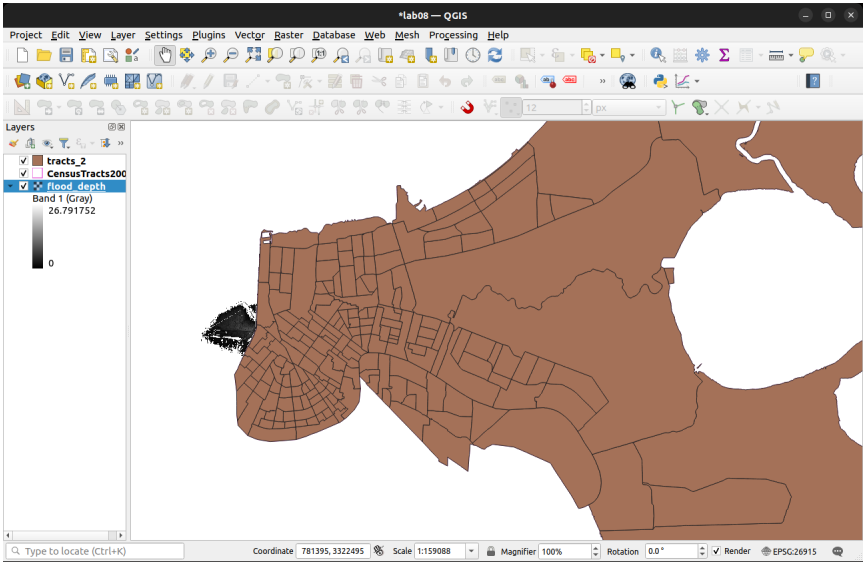

You can adjust the symbology to see how the `flood_mean` variable is distributed

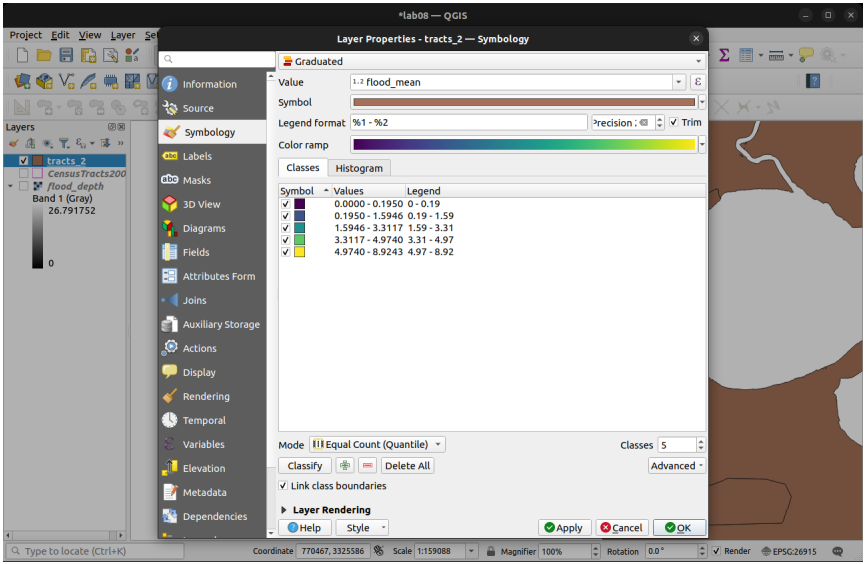

Some of the worst-affected areas seem to be in the city's north (Lakeview, Gentilly), southeast (lower 9th Ward) and wetlands to the east

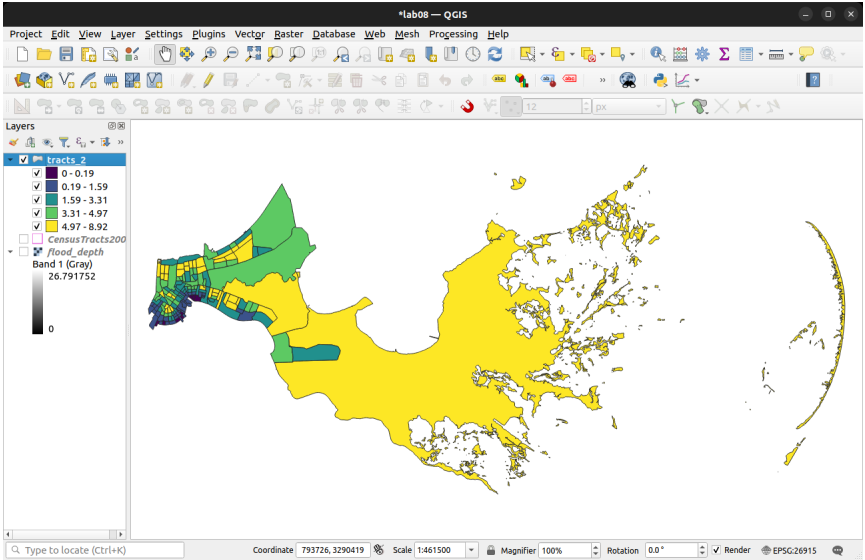

Let's now see if historically "redlined" areas were disproportionately affected. Load the `nola1939.geojson` file from `Data/Inequality/`

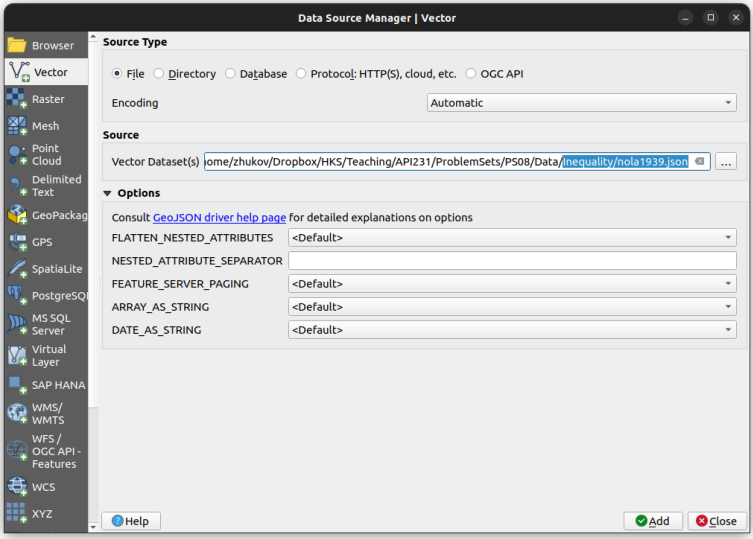

QGIS may prompt you to select a transformation method (to reproject this layer from WGS84 to UTM). Pick the top-listed one and click `OK`

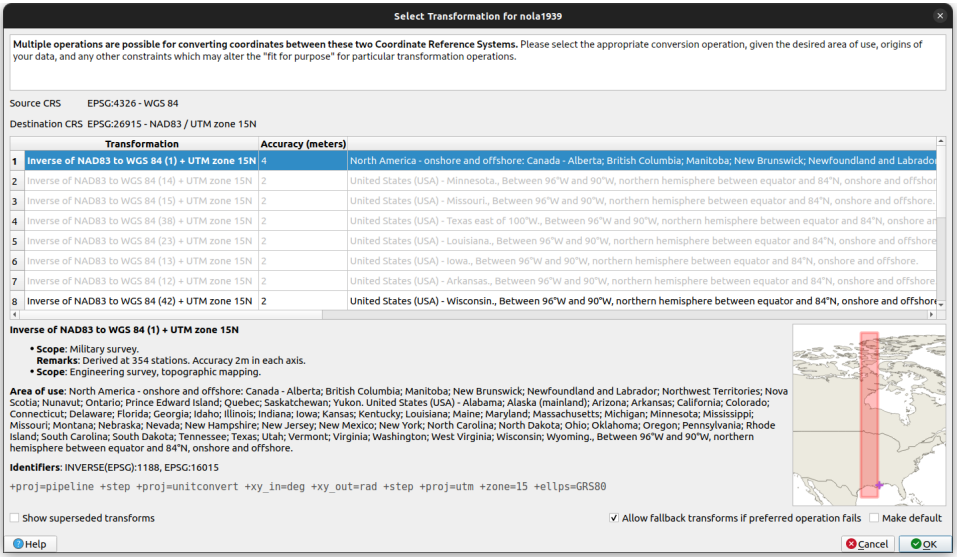

If all goes well, the `nola1939` layer should appear on the map

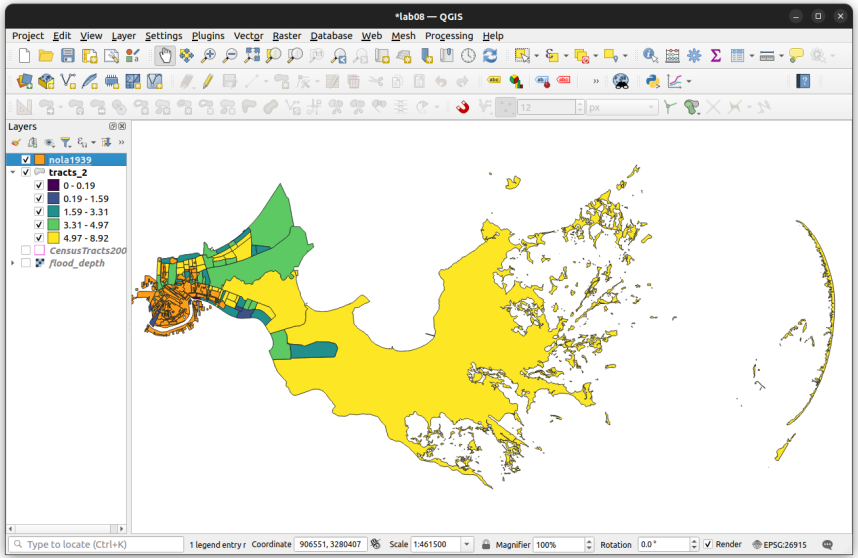

Let's color these polygons by `grade` (Categorized). Grade `D` represents the redlined areas that the HOLC considered too "hazardous" for home loans

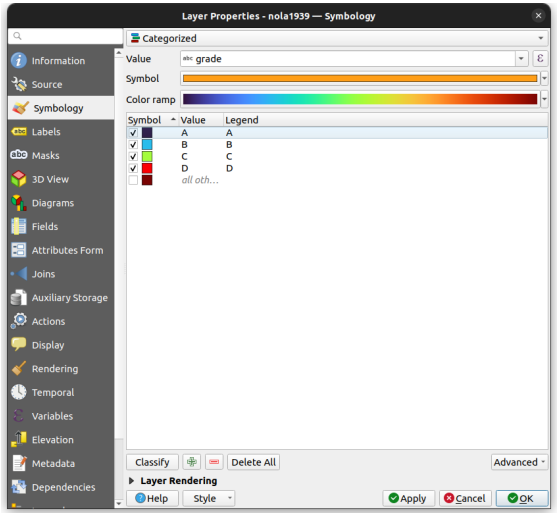

These grades made it difficult or impossible for people to access mortgage financing and become homeowners. The brunt of redlining fell on neighborhoods of color.

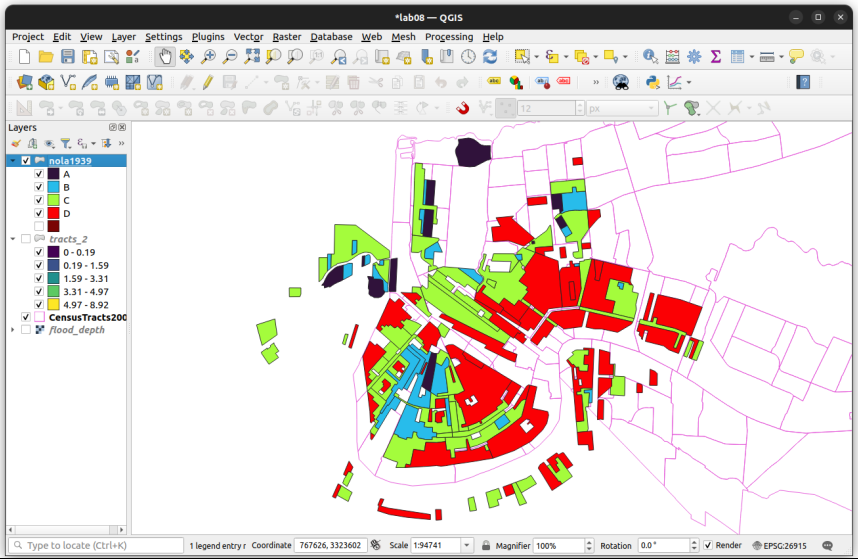

Let's separate out the areas with a grade of `D`. Go to `Edit` menu → `Select` → `Select Features by Expression...`

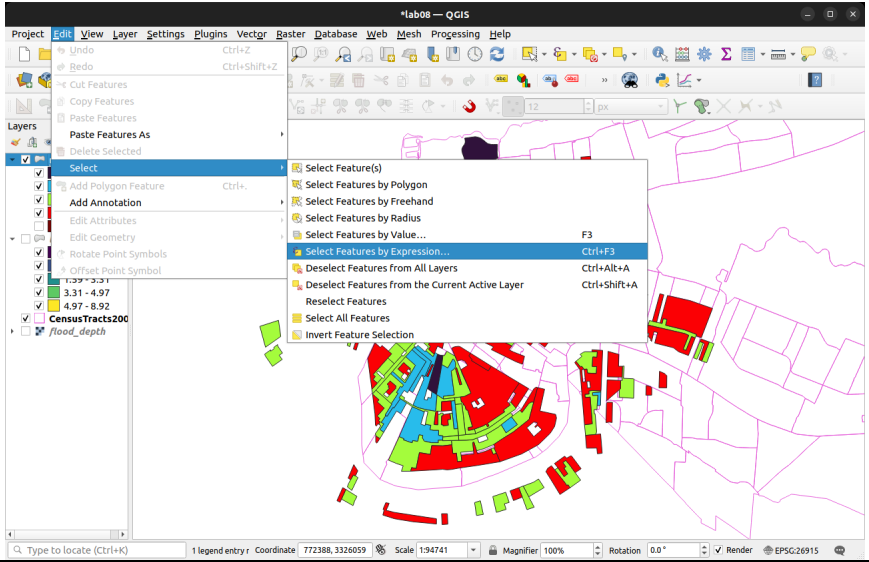

On the next screen, set `Expression` to `grade = 'D'`
Click `Select Features` and close this window

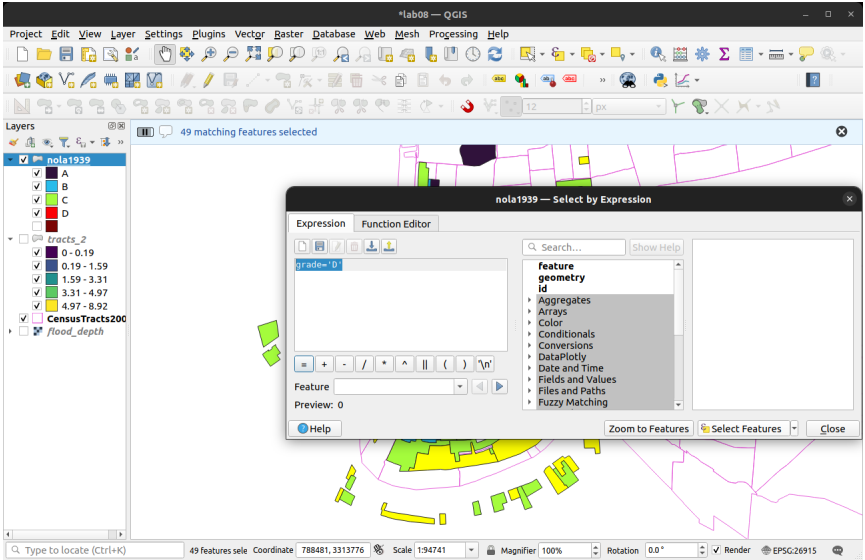

Let's export the selected features to a new file. Right-click on `nola1939` in the layer menu, then `Export` → `Save Selected Features As...`

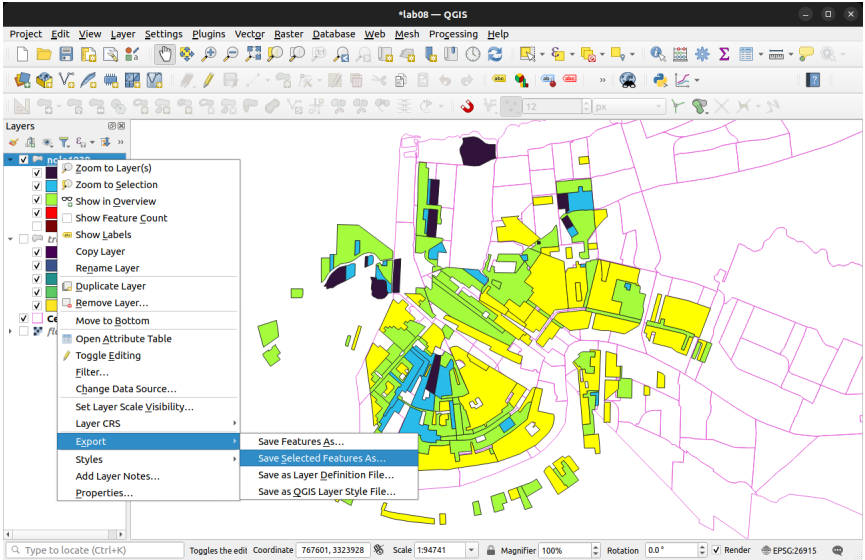

On the next screen set

- `Format` = `GeoJSON`
- `File name` = `redlined.geojson`
- ✓ `Save only selected features`
- `Geometry type` = `Polygon`

Click `OK`

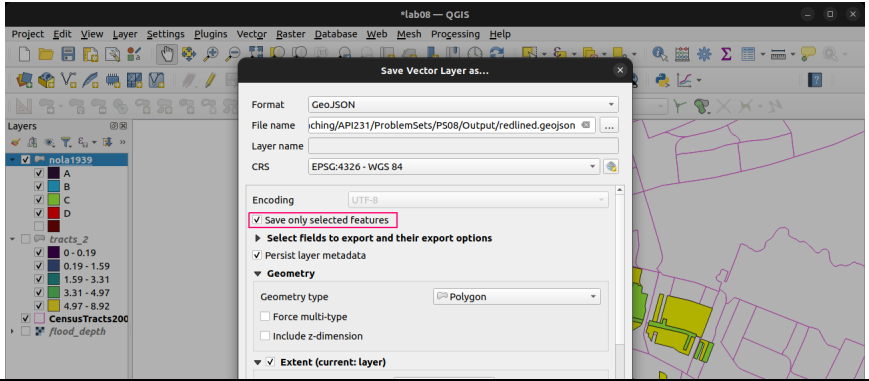

The `redlined` layer should appear in your project window.
Let's calculate the proportion of each census block that was redlined in 1930s

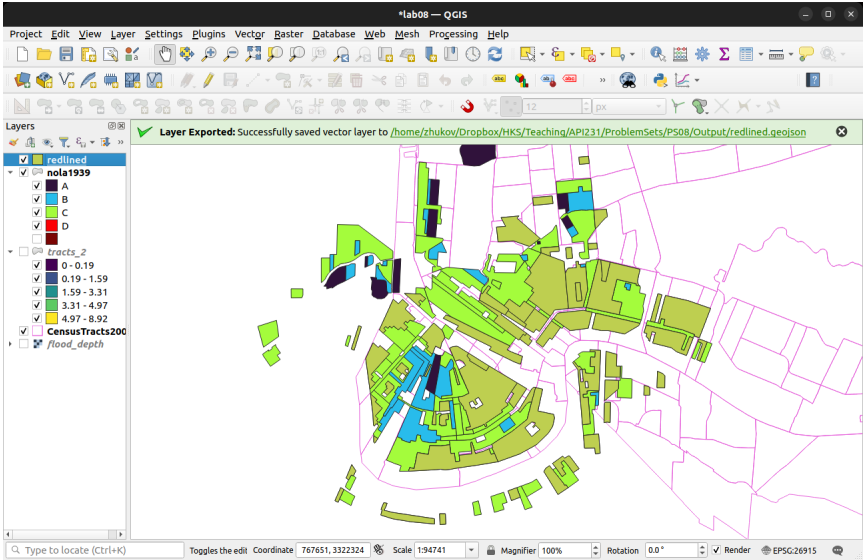

Go to `Processing Toolbox`, then `Vector analysis` → `Overlap analysis`

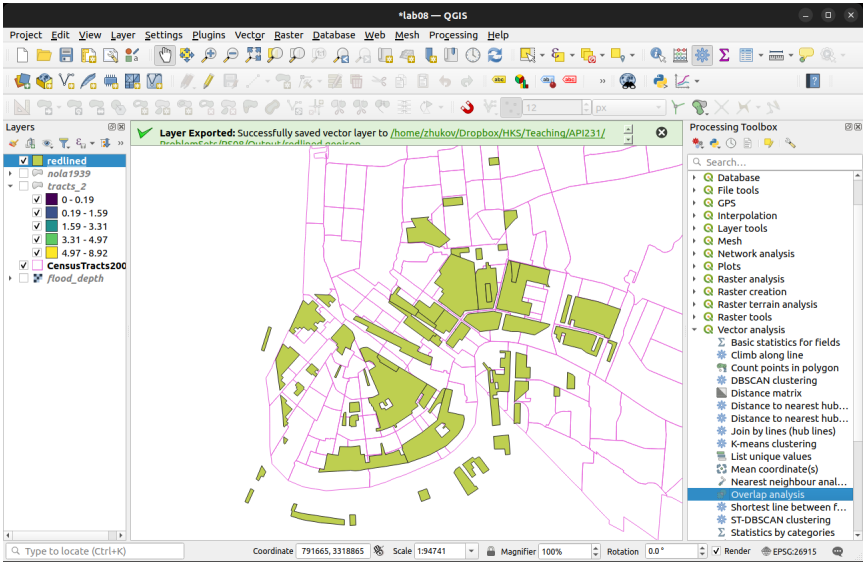

In the Overlap Analysis window, set

- `Input layer` = `tracts_2`
- `Overlap` (save to file) = `tracts_3.geojson`

Click on the `[...]` button next to `Overlay layers`

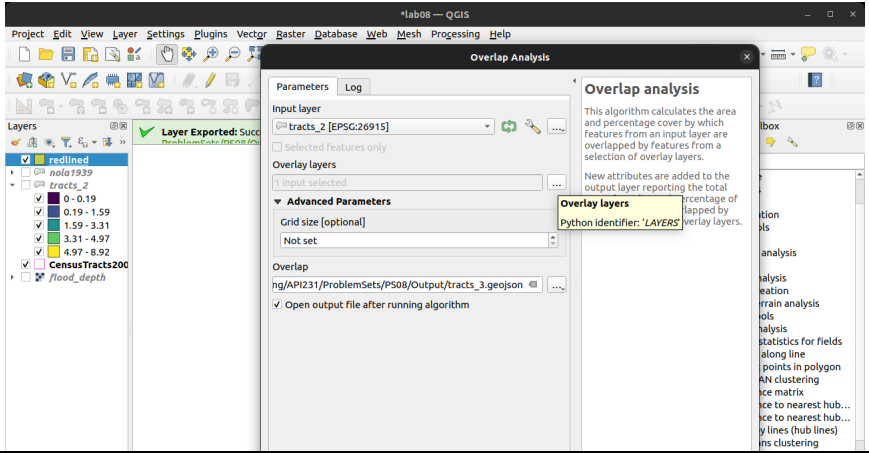

Select ✓ `redlined`. Click `OK`

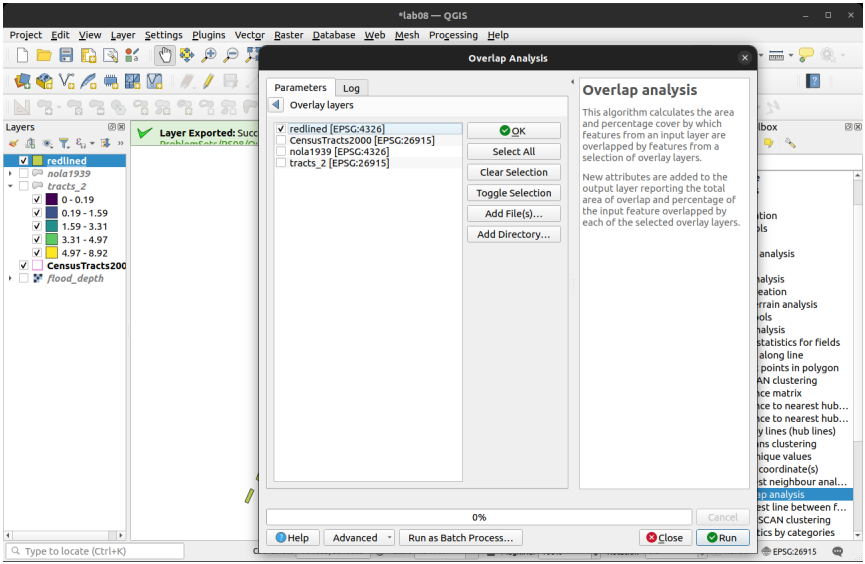

Click `Run`. This will add a new layer, `tracts_3`, to your project

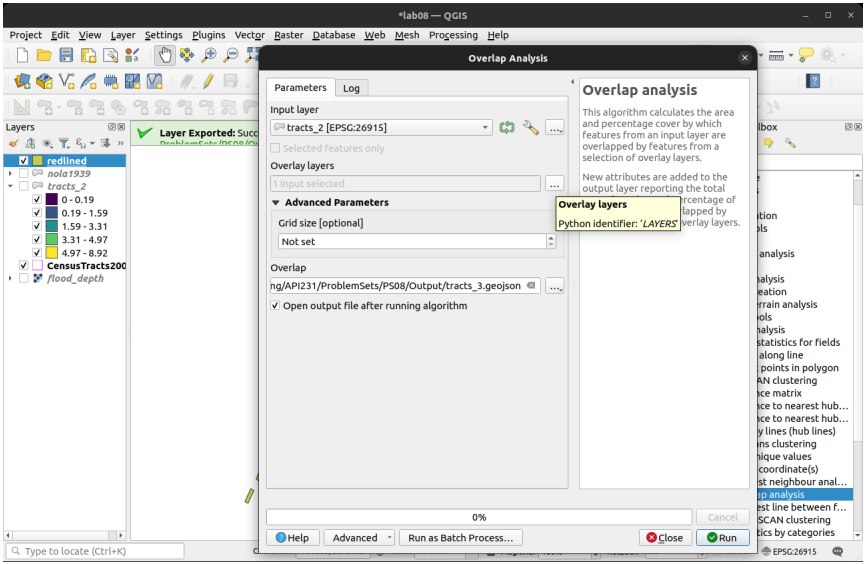

You can adjust the symbology to see how the `redlined_pc` variable (percent of tract with grade D) is distributed

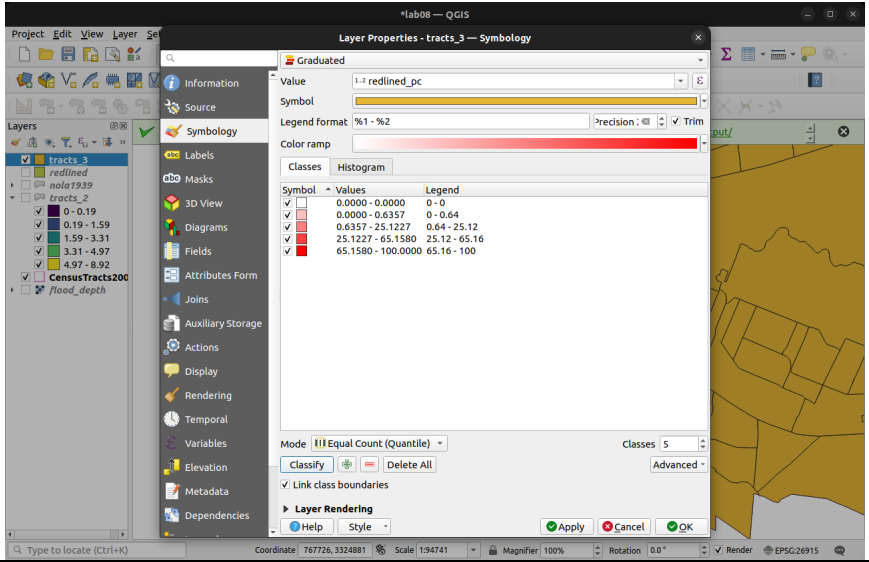

Let's make a quick scatterplot to see if redlined areas saw more flooding. Go to `View` menu → `Panels` → `DataPlotly` (if you need to install Plotly, see Lab 6)

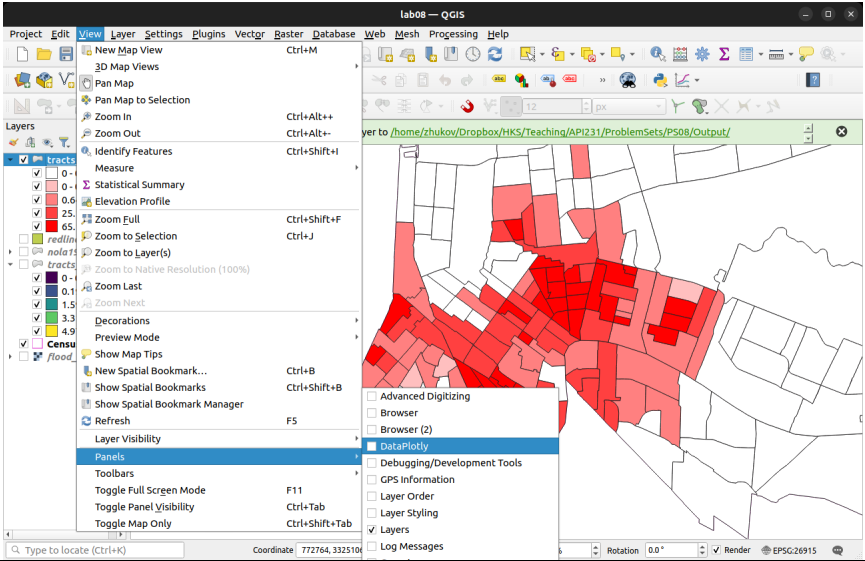

In the DataPlotly panel, set `Plot type` = `Scatterplot`; `Layer` = `tracts_3`; `X field` = `redlined_pc`; `Y field` = `flood_mean`; Click `Create Plot`

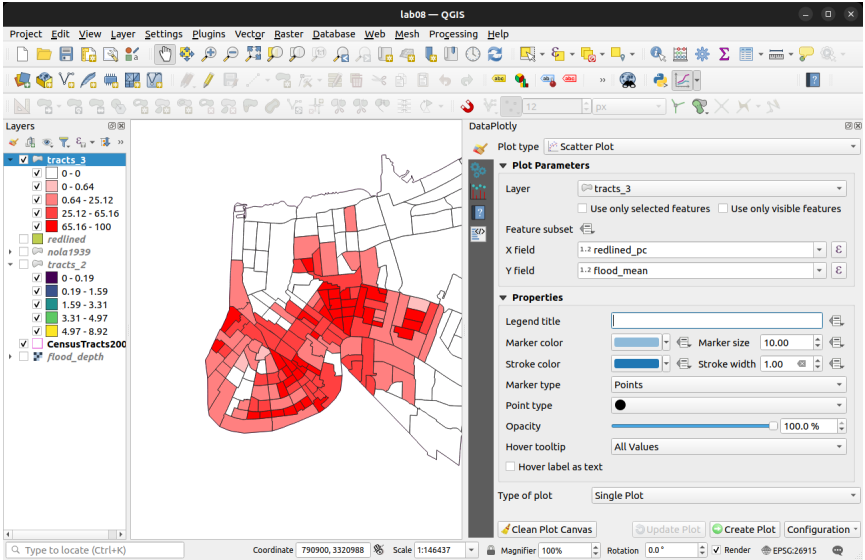

There doesn't seem to be much of a relationship here. Census tracts with more redlined areas did not experience more flooding than less-redlined areas

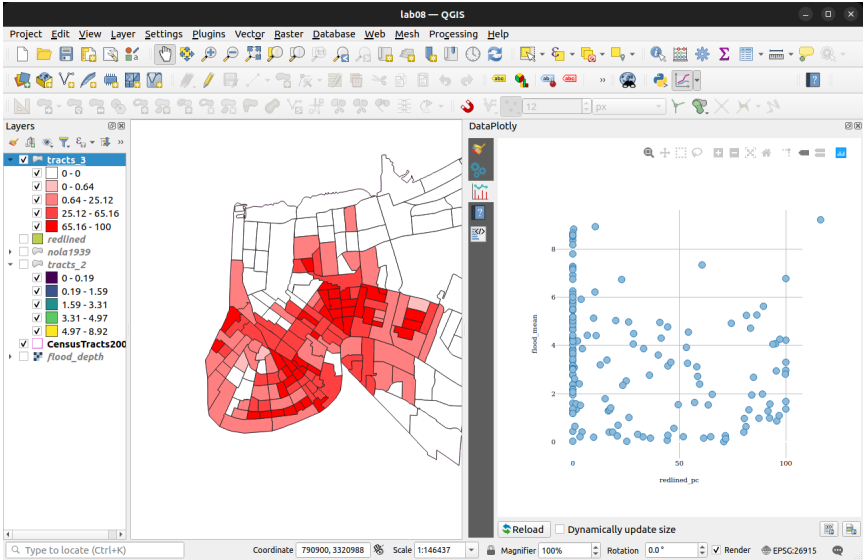

What if we compared tracts with *any* redlining to those with no redlining? Open the Field Calculator for `tracts_3` and set `name`: `redlined_any`, `type`: `Integer`, `Expression`: `redlined_pc > 0`, Click `OK`

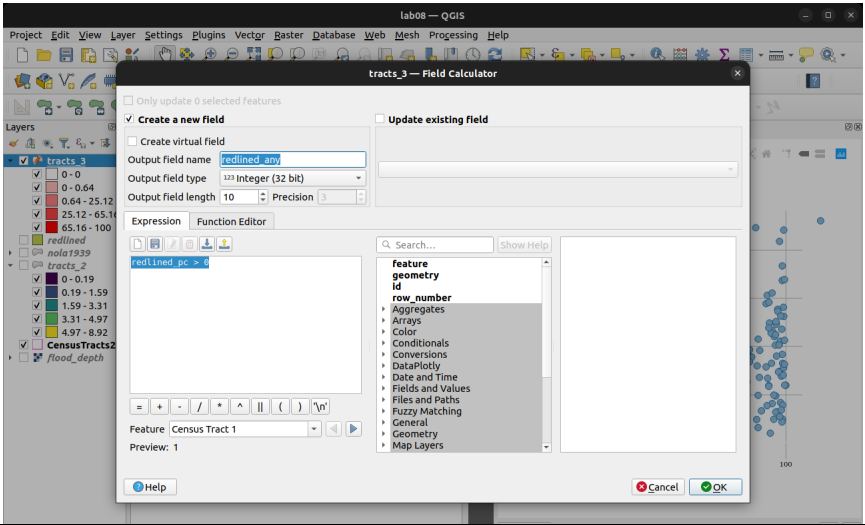

Back in the DataPlotly panel, set `Plot type` = `Box Plot`; `Layer` = `tracts_3`; `Grouping field` = `redlined_any`; `Y field` = `flood_mean`. Click `Clean Plot Canvas`, then `Create Plot`

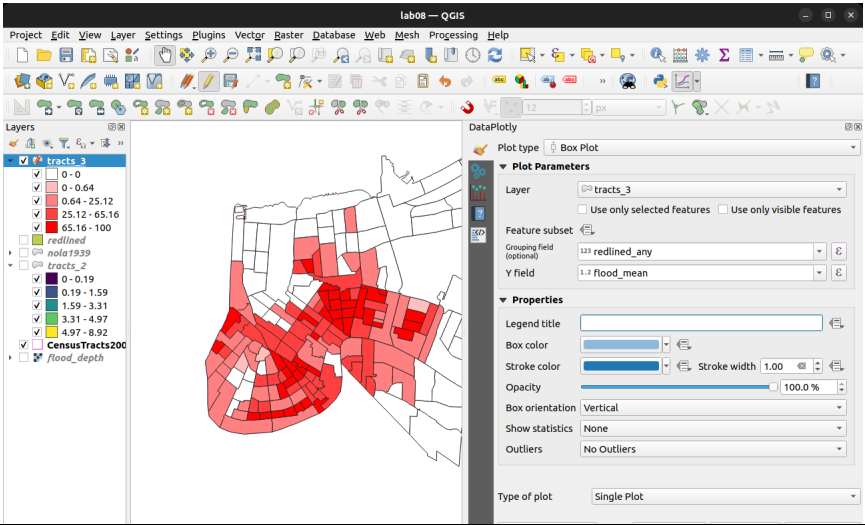

The boxplot suggests that, if anything, redlined areas saw less flooding, on average

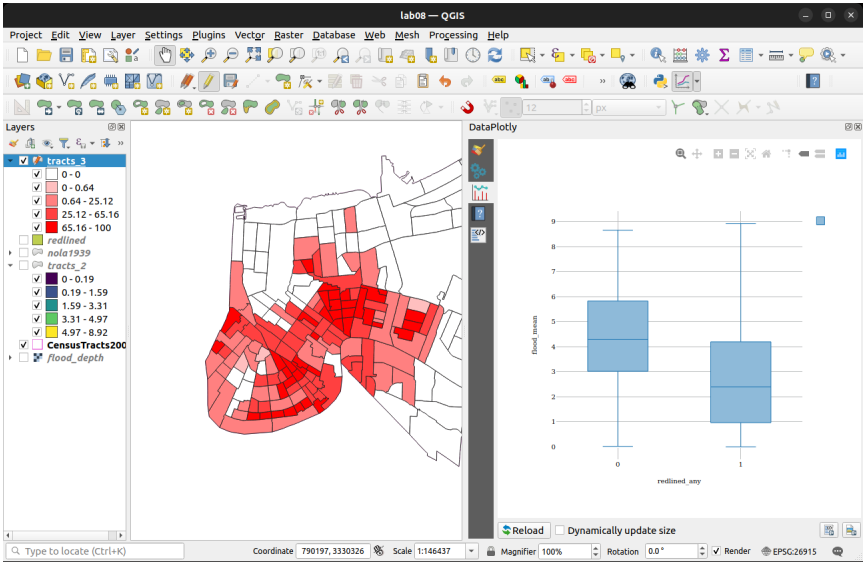

Let's look at redlining's relationship with neighborhood demographics. Go to Field Calculator for `tracts_3` and create a new variable, with `name`: `PCT_BLACK`, `type`: `Decimal number`, `Expression`: `RACE_BLACK / POP_TOTAL * 100`. Click `OK`

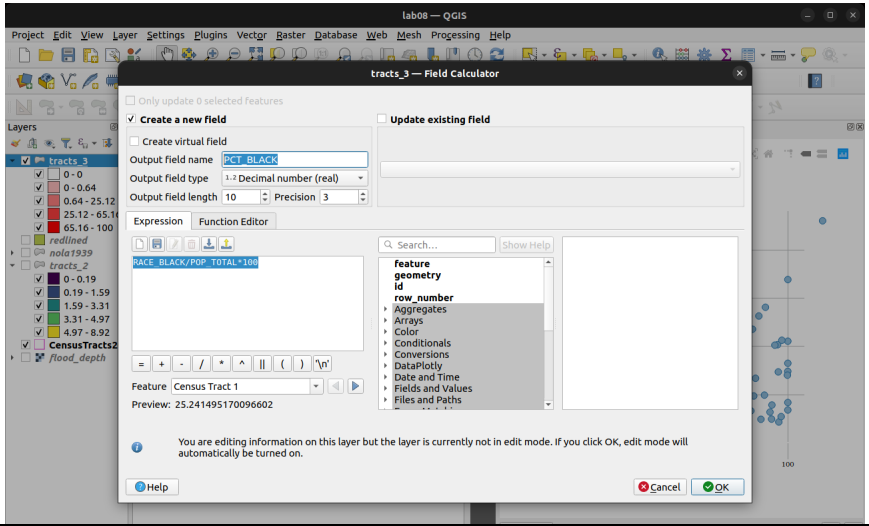

Let's make a scatterplot with `redlined_pc` on the X axis and `PCT_BLACK` on the Y axis. Click `Clean Plot Canvas`, then `Create Plot`

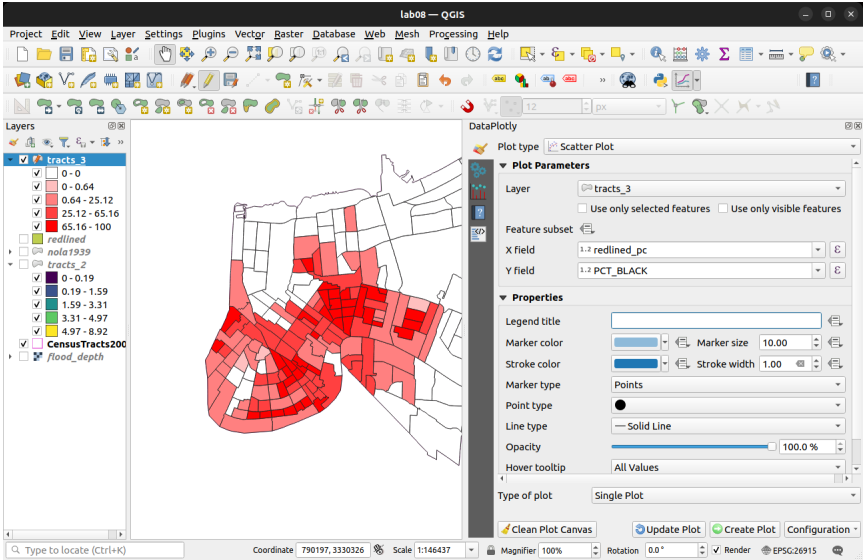

While there is a lot of variation among non-redlined neighborhoods, those closer to 100% redlining still have an overwhelmingly Black population

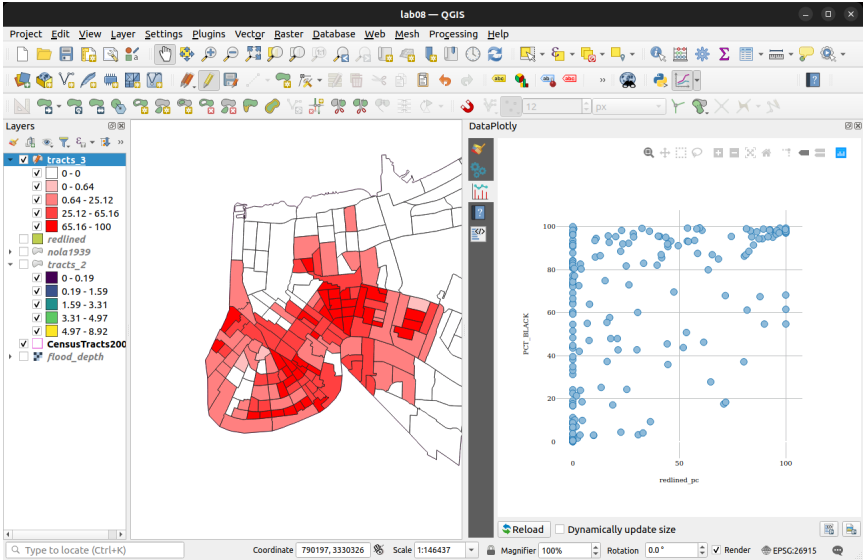

Raster reclassification

Sometimes, we need to do some pre-processing to prepare raster data for analysis. Suppose we want to know the proportion of a census tract classified by FEMA as "high flood risk". Load the raster `NFHL_NOLA.tif` from the folder `Data/FEMA/`

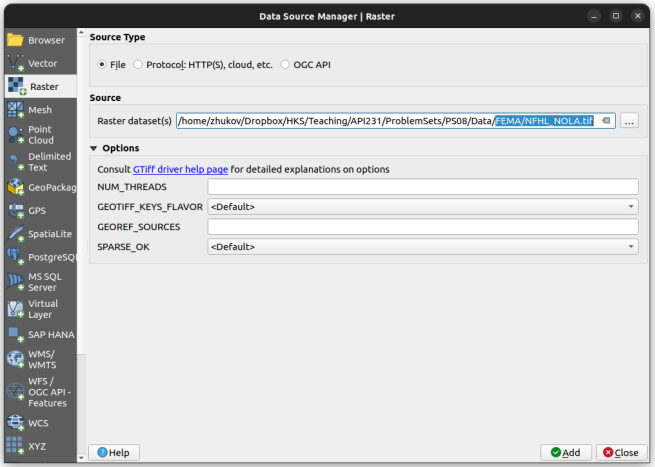

The raster appears to have four unique classes, titled `FLD_ZONE`. They are on a greyscale ramp by default

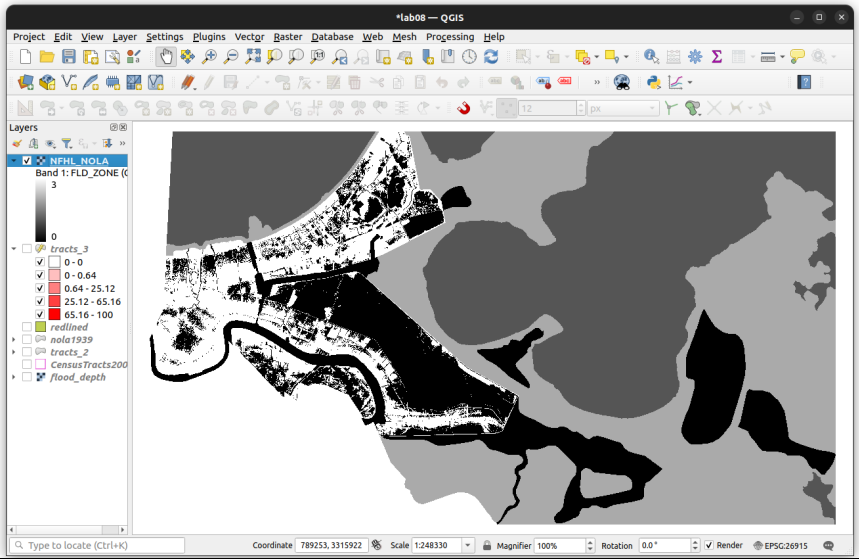

FEMA uses several designations of flood risk, which are used by flood insurance providers and home lenders

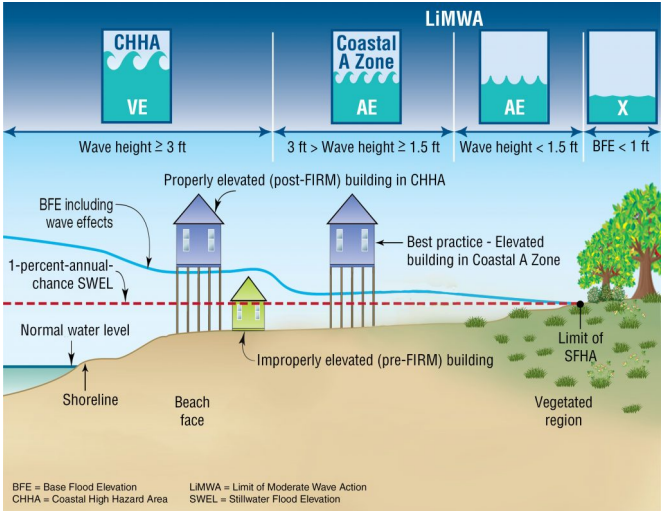

Figure 13: NFHL codes

QGIS reads these codes as integers, which correspond to...

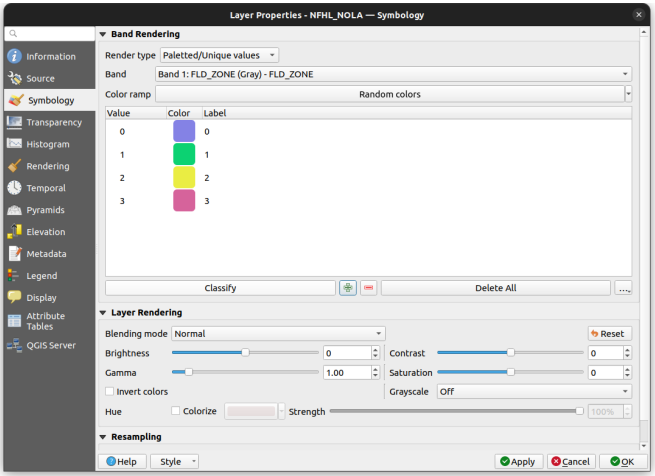

0 → “AE” (high flood risk), 1 → “OPEN WATER”,
2 → “VE” (coastal high hazard area), 3 → “X” (moderate-to-low risk)

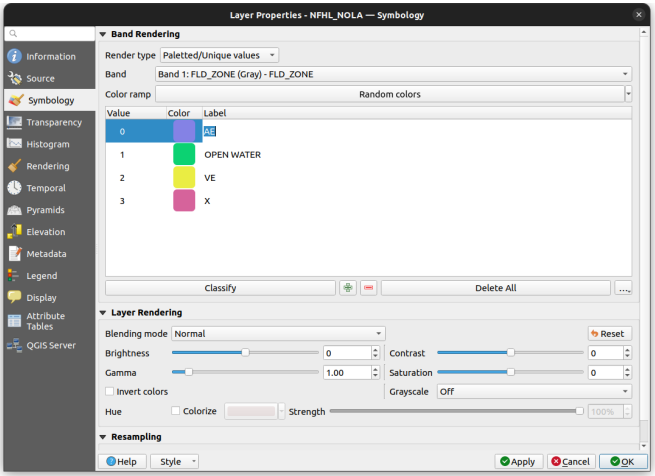

What we want to know is how much of a census tract is covered by "high flood risk" zones (AE, or 0). Let's extract just this part of the raster

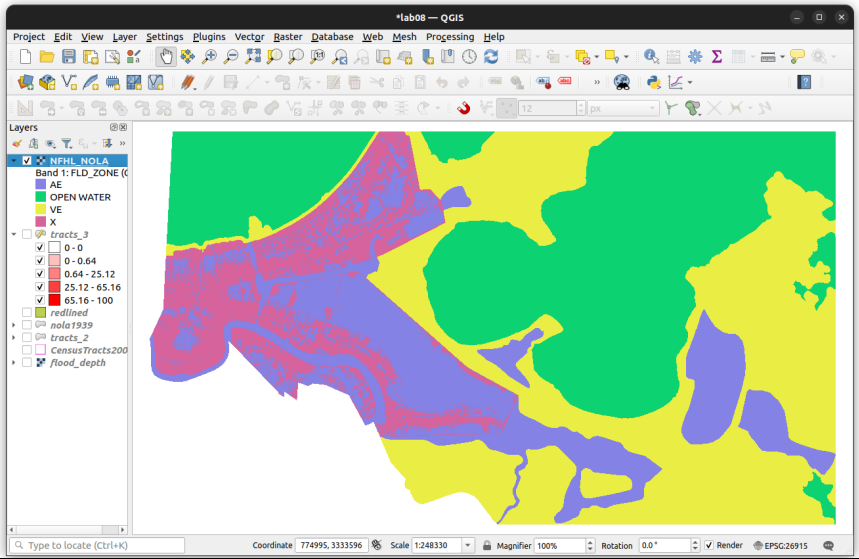

Go back to the `Processing Toolbox` → `Raster analysis` → `Raster calculator`

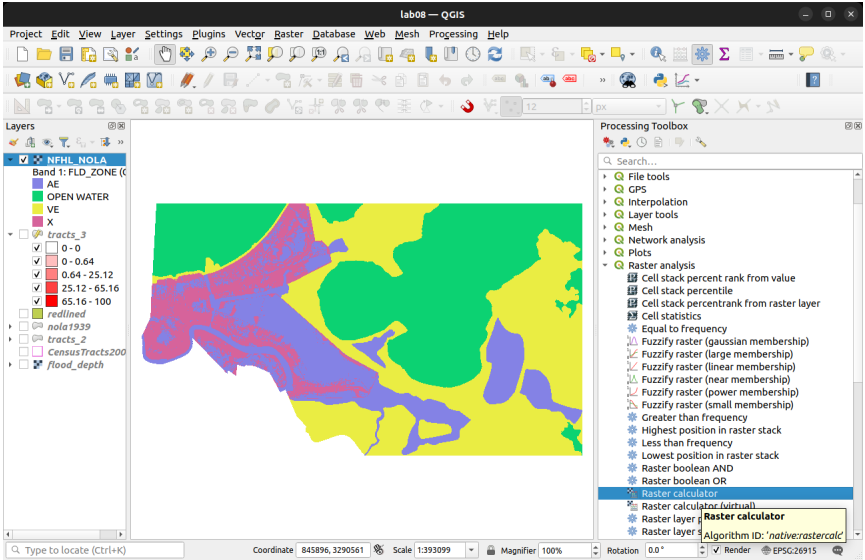

In the Raster Calculator window, click on the `[...]` box next to `Input layers`

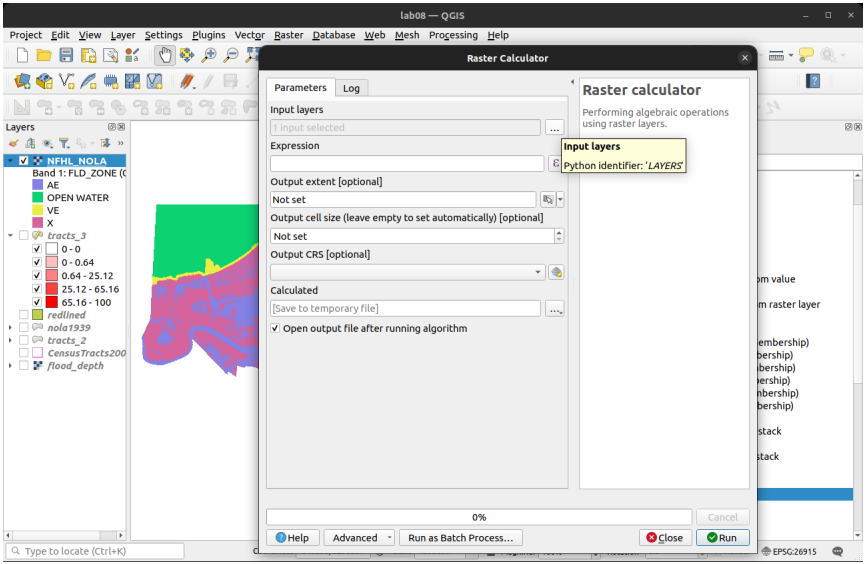

Select ✓ NFHL_NOLA. Click OK

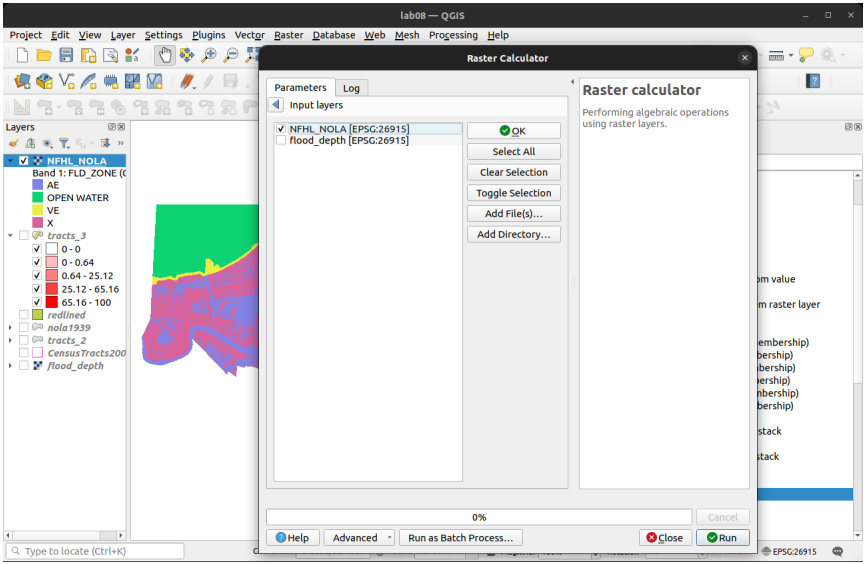

Click on the ϵ button next to `Expression`

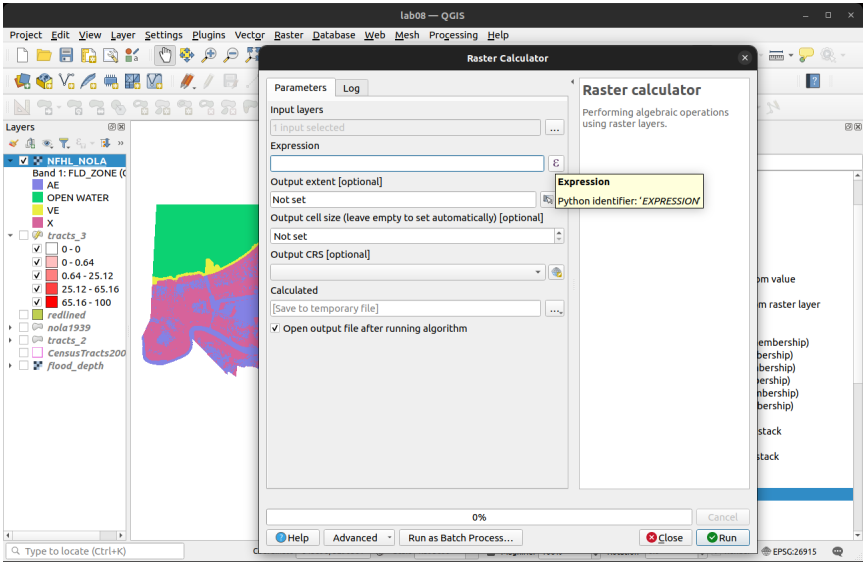

For the Expression, enter `"NFHL_NOLA@1" = 0` (make sure to include the quotation marks). Click `OK`

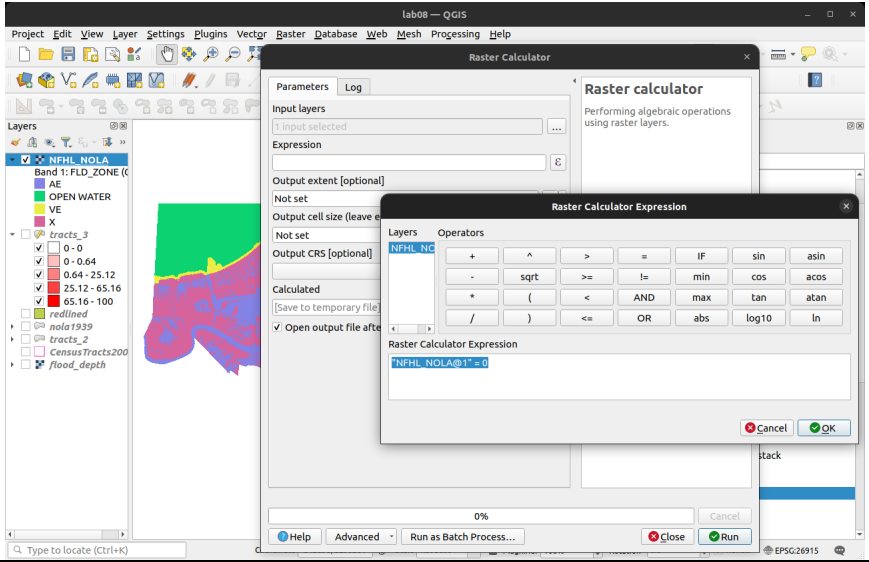

Leave the `Output extent`, `cell size` and `CRS` fields blank (accept defaults). For `Calculated`, save to file as `nfhl_highrisk.tif`. Click `Run`

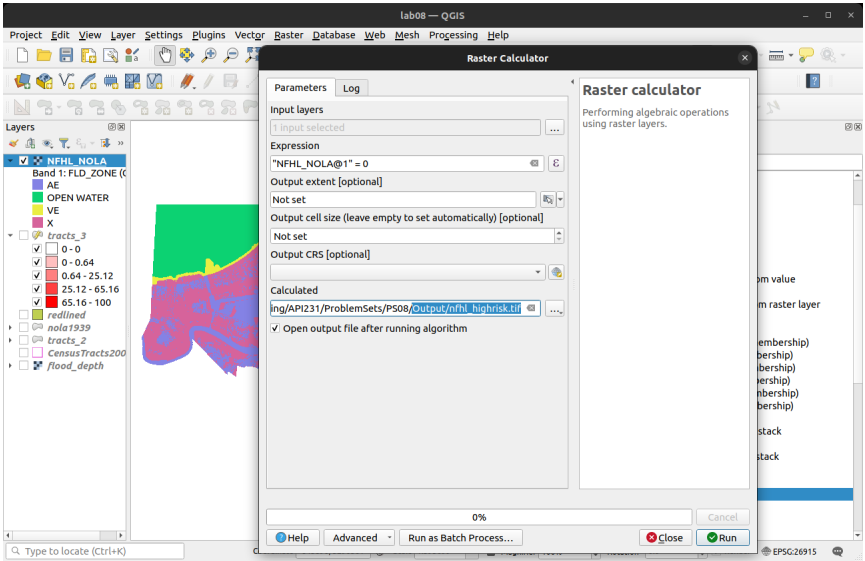

This creates a binary (1/0) raster indicating whether or not pixel is in high-risk zone

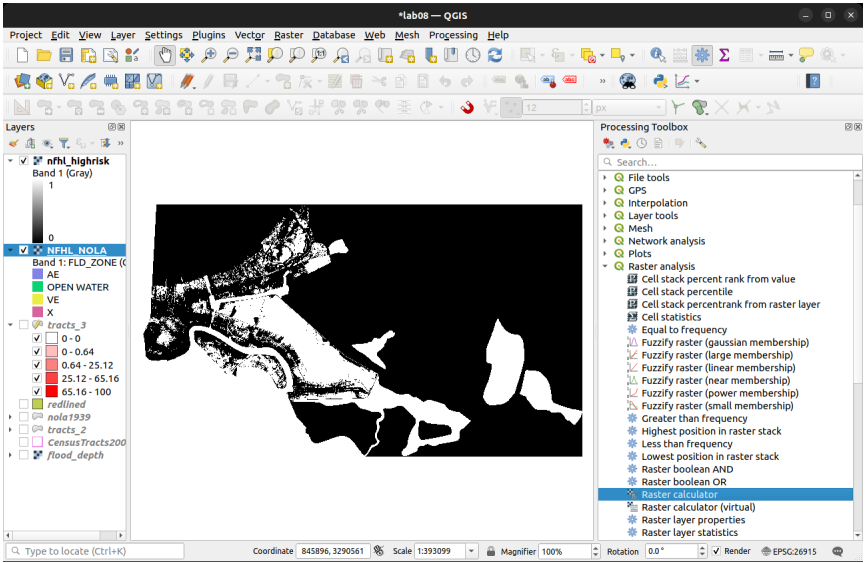

Go back to `Zonal statistics` in the Processing Toolbox

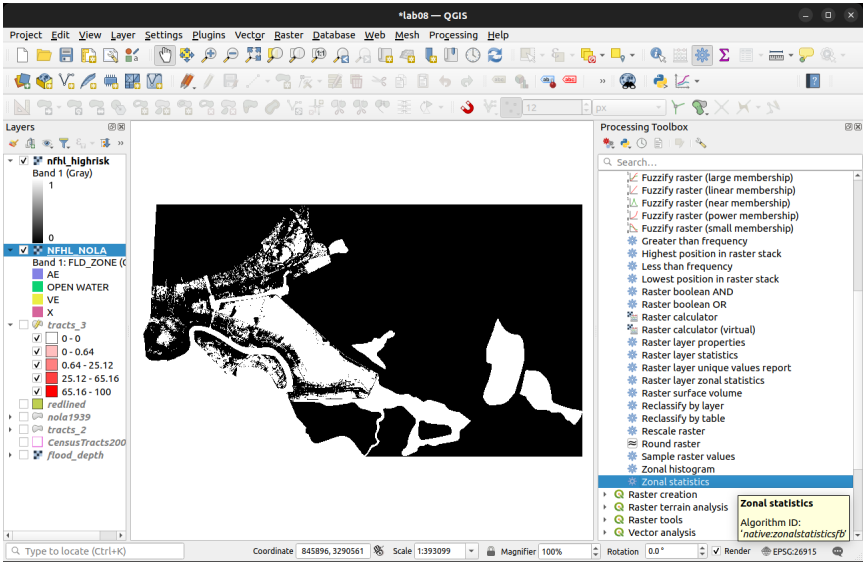

Select `tracts_3` as the `Input layer`, `nfhl_highrisk` as the `Raster layer`, set the column prefix to `highrisk_`. CLick on the `[...]` button next to `Statistics to calculate`

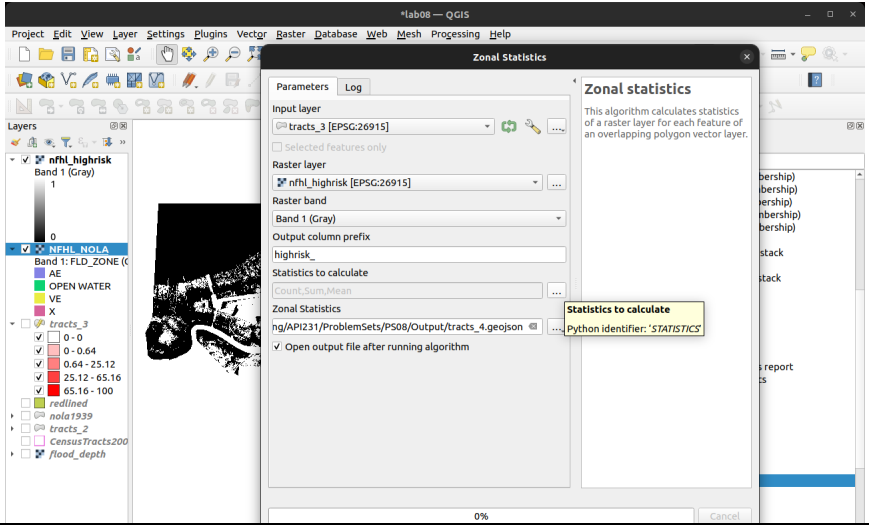

Select ✓ `Mean`. Click `OK`

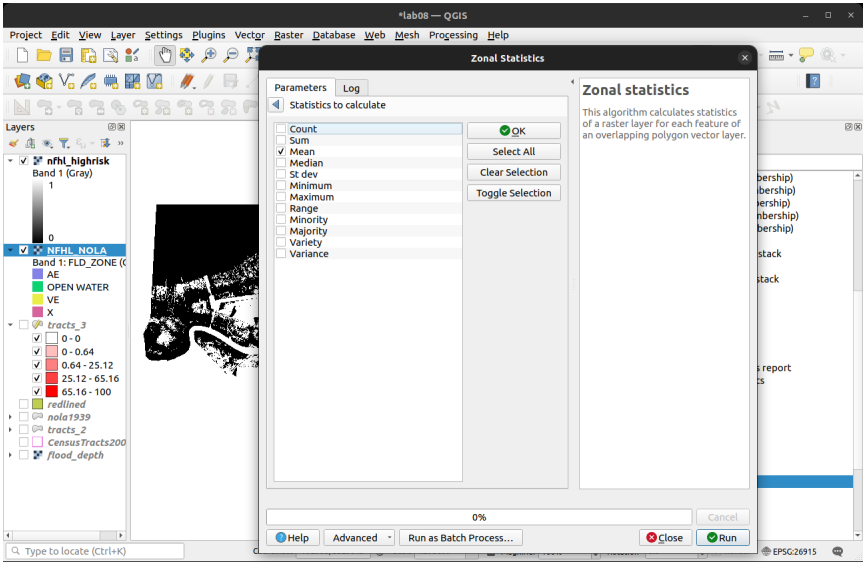

Save the output to a file named `tracts_4.geojson`. Click `Run`

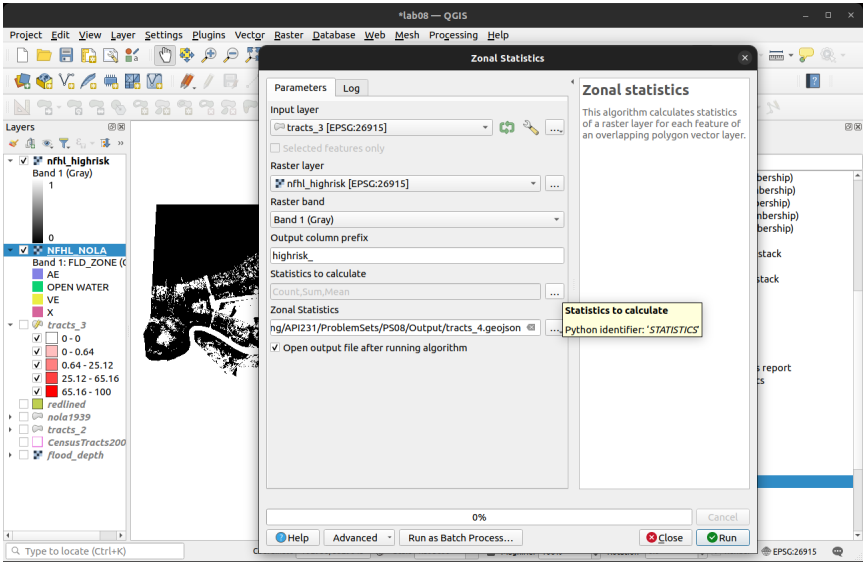

This may take a few minutes. Be strong. Don't give up

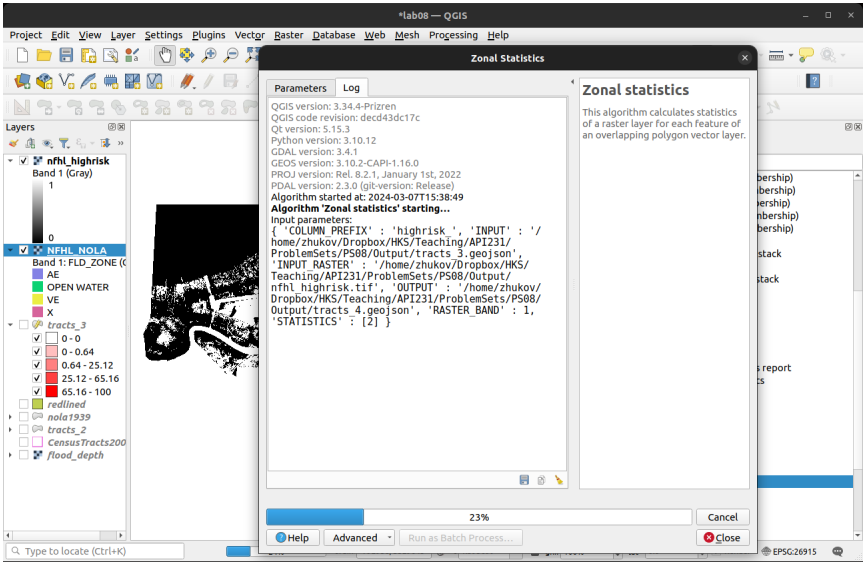

When the `tracts_4` layer finally loads, you can change the symbology and explore the distribution of the `highrisk_mean` variable

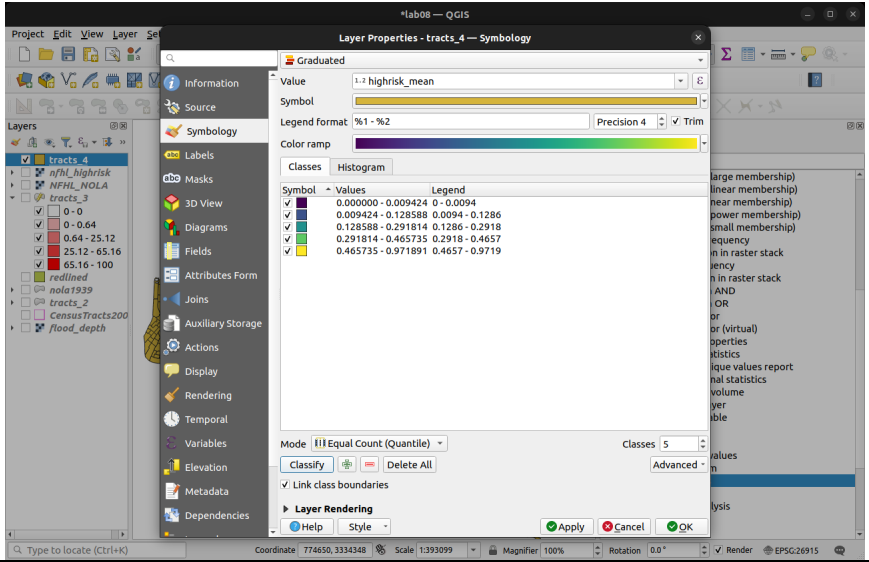

The final ingredient we will need is the 311 call data. Load the delimited text file `calls311_water_2012_2018.csv` from the `Data/Call311/` folder. Set the `Geometry Definition` to `Point coordinates` (as shown here), and check the box next to ✓ `Use spatial index` in `Layer Settings`. Click `Add`

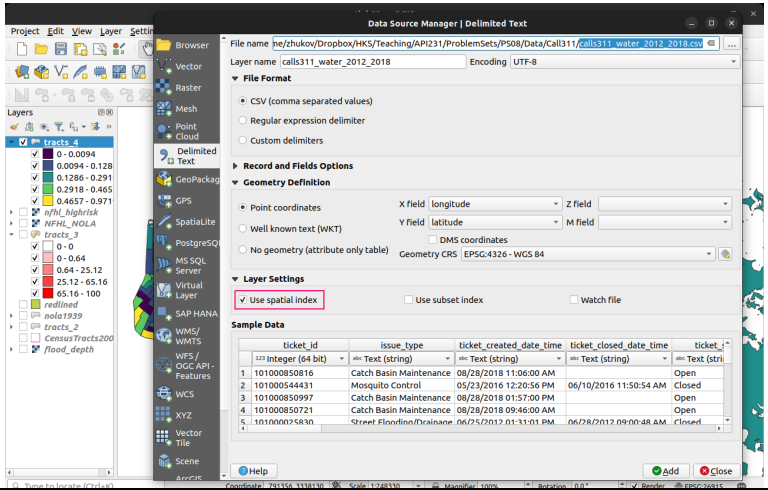

Note that the 311 calls are available only for New Orleans, not St. Bernard Parish

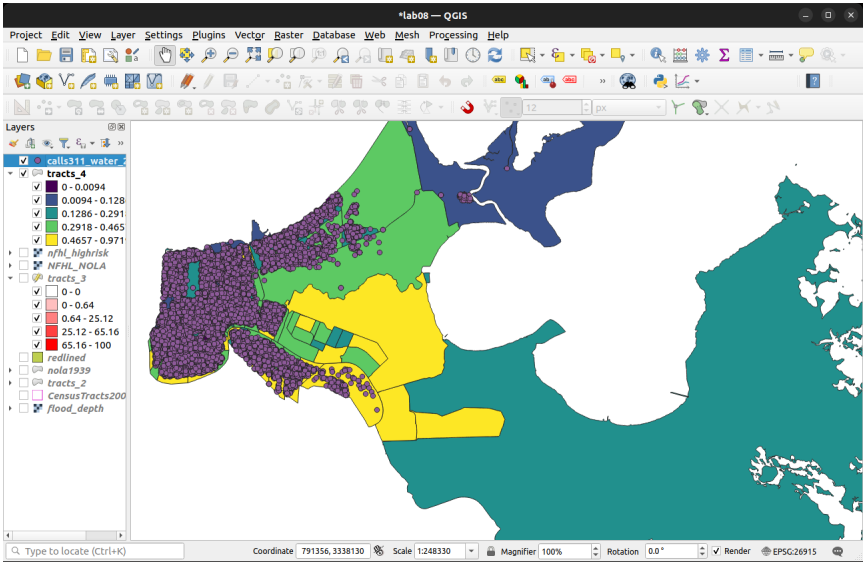

Open the Attribute Table for the `calls311_*` layer, and look at the `issue_type` field. There are three types here: "Catch Basin Maintenance", "Mosquito Control", and "Street Flooding/Drainage". Let's create point counts for each of these

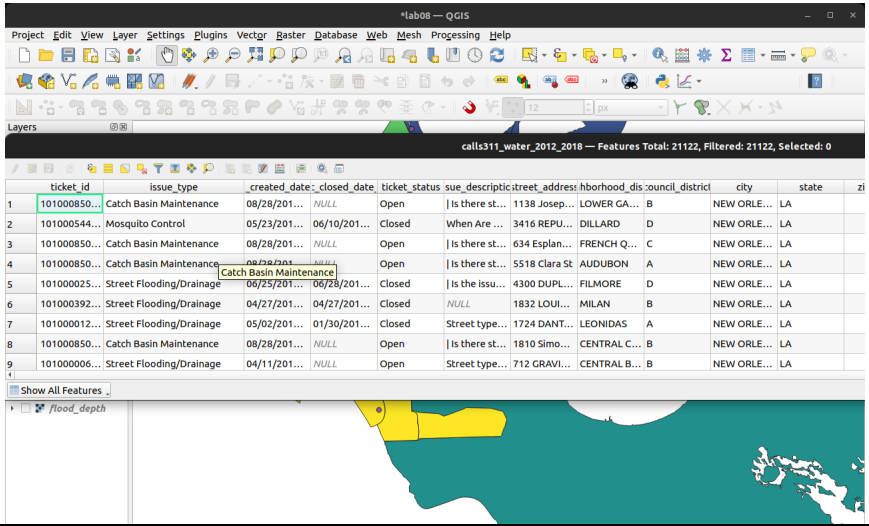

Go to `Edit` menu → `Select` → `Select Features by Expression...`

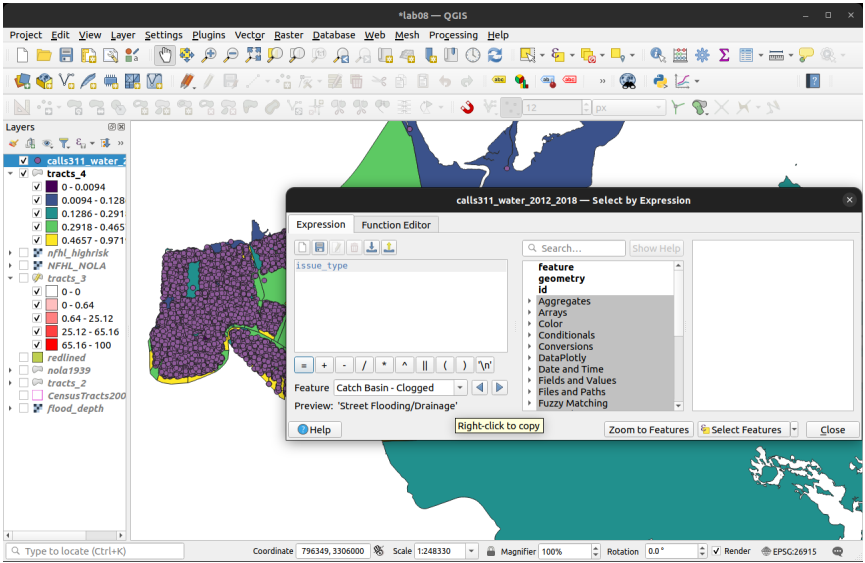

For `Expression`, enter `issue_type = 'Street Flooding/Drainage'` (with single quotation marks). Click `Select Features`

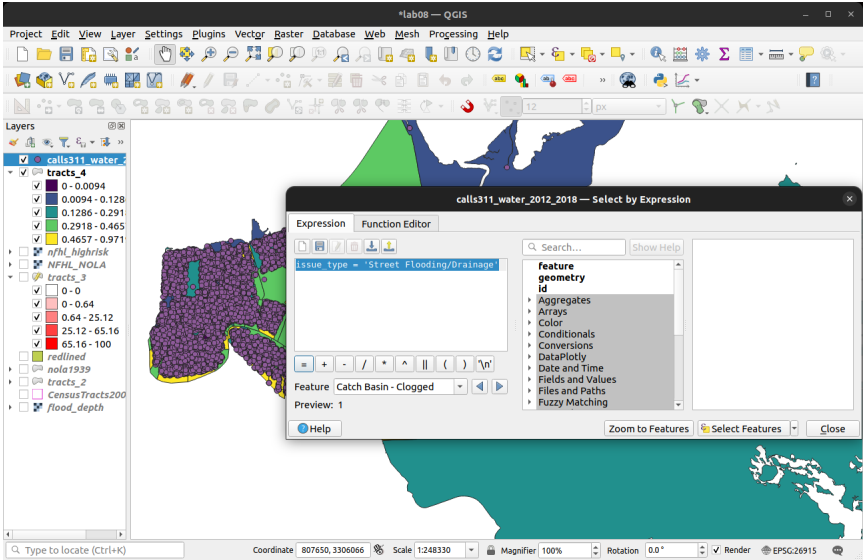

With the points selected, open the `Count Points in Polygon` tool in `Vector` → `Analysis Tools`

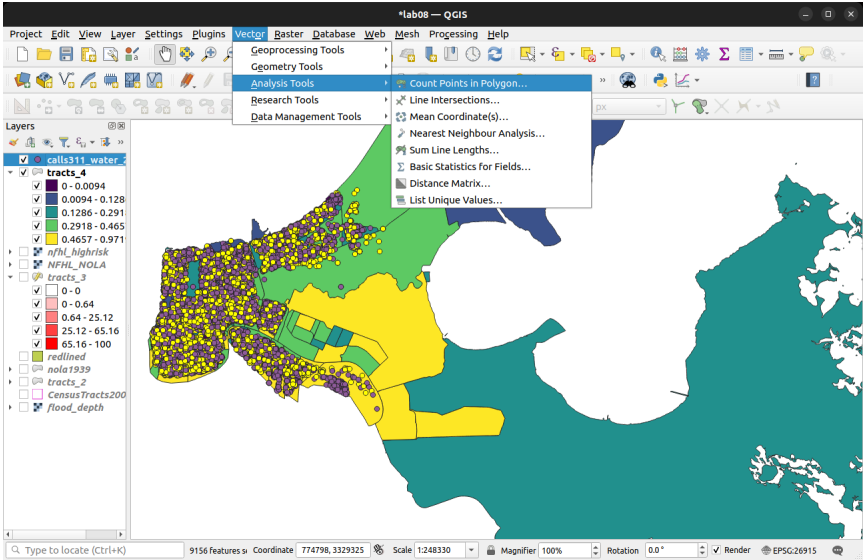

Select `tracts_4` as the Polygons layer and `calls311_*` as the Points layer. Make sure the box next to ✓ `Selected features only` is checked. Name the count field `calls311_streetflood` and save the output to a new file called `tracts_5.geojson`. Click `Run`

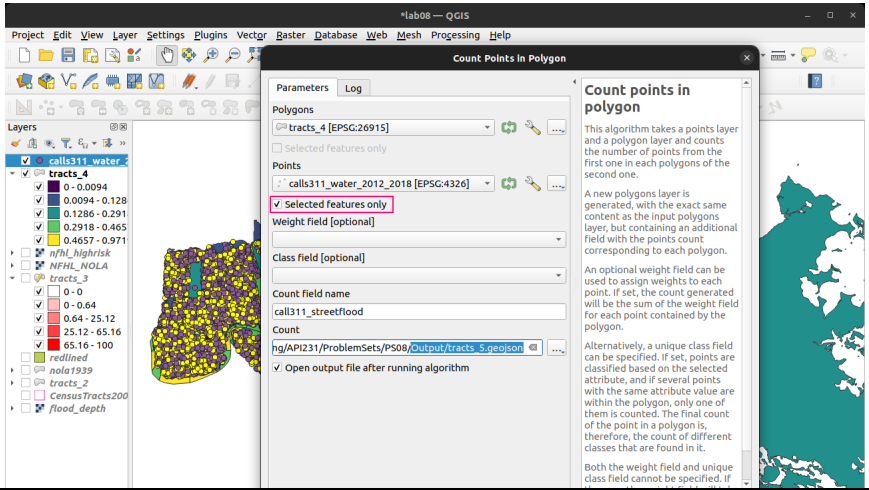

Repeat this process for the other two categories of calls, starting with "Catch Basin Maintenance"

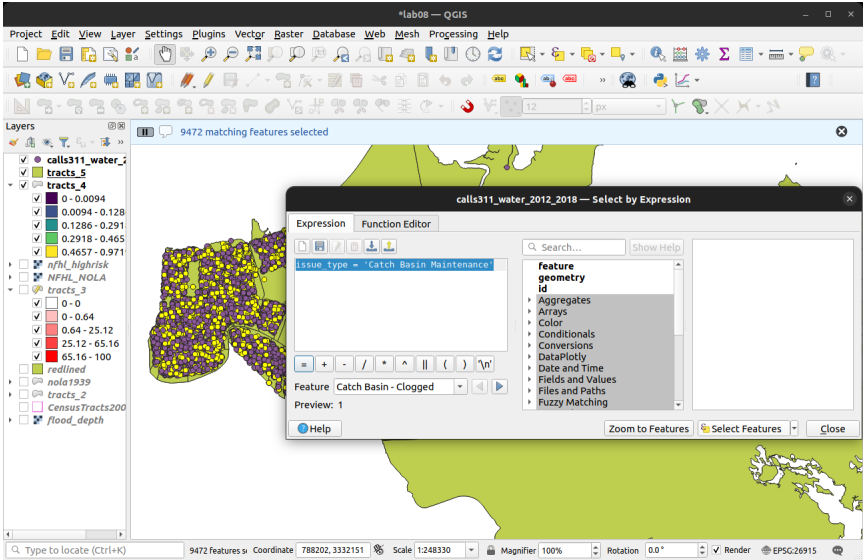

To save a bit of time, you can call up the most recently-used geoprocessing options by going to `Processing` menu → `History...`

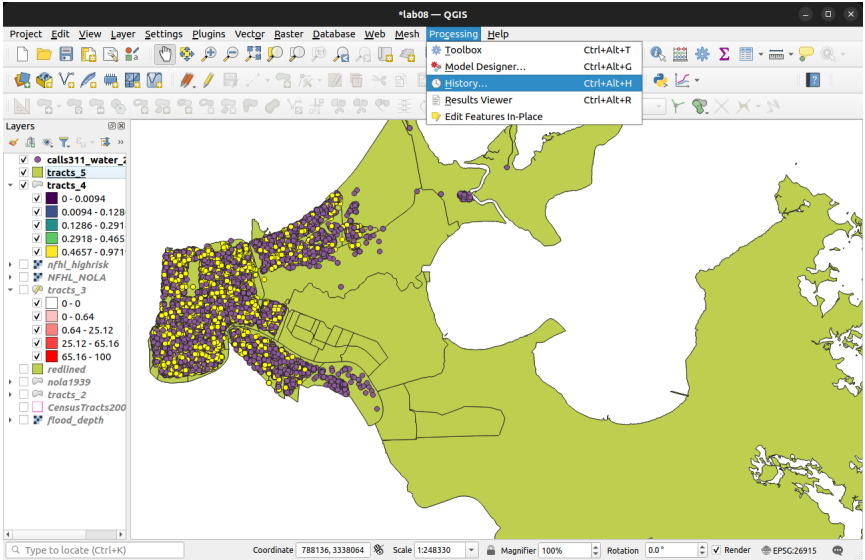

Double-click on the `Count points in polygon` operation

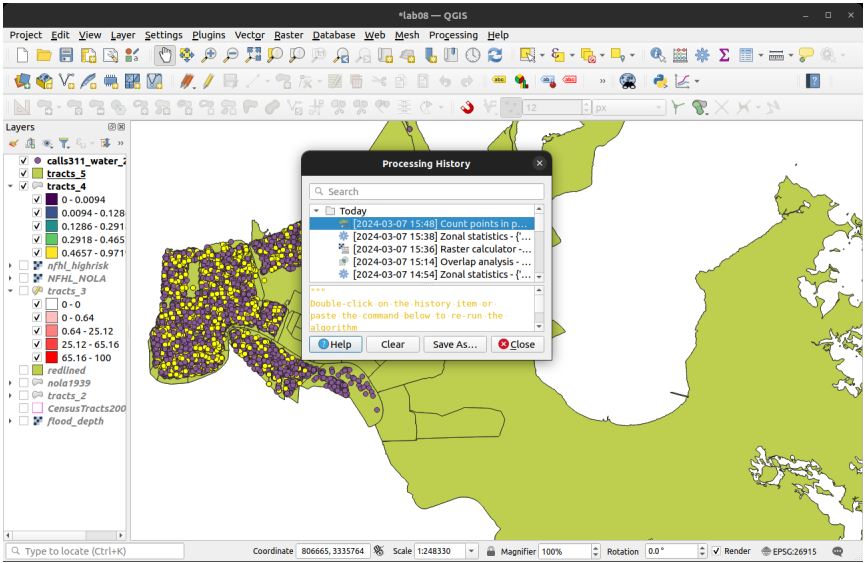

This will restore the settings you just used to create `tracts_4`. Now we just need to update the parameters

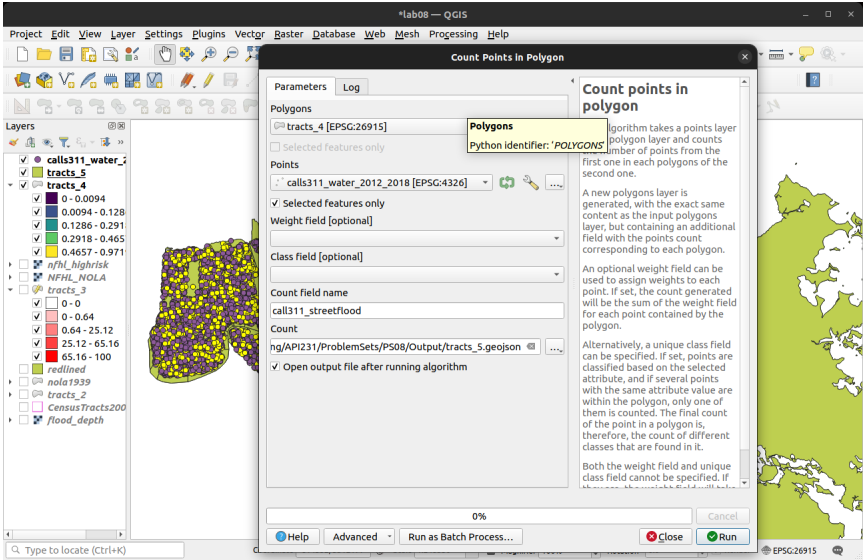

Change `Polygons` to `tracts_5`, change field name to `calls311_catchbasin`, and save the file as `tracts_6.geojson`. Click `Run`

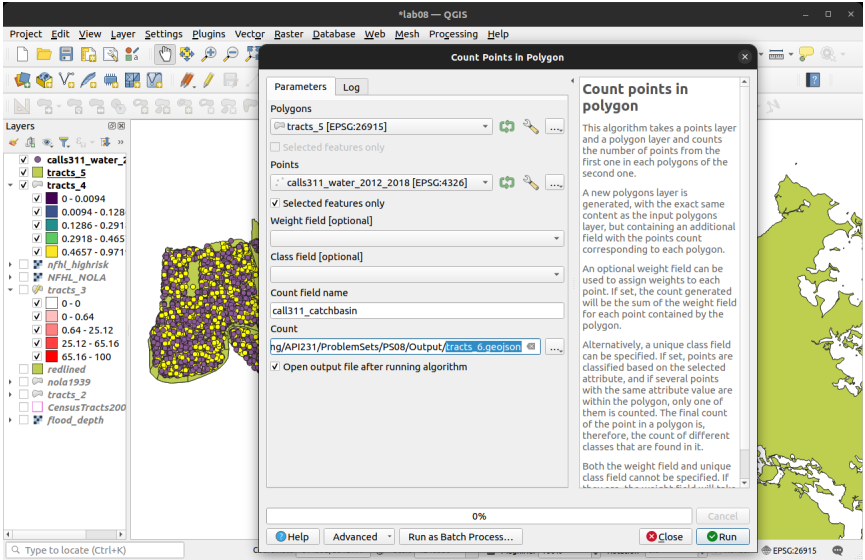

And one more time with `issue_type = 'Mosquito Control'`

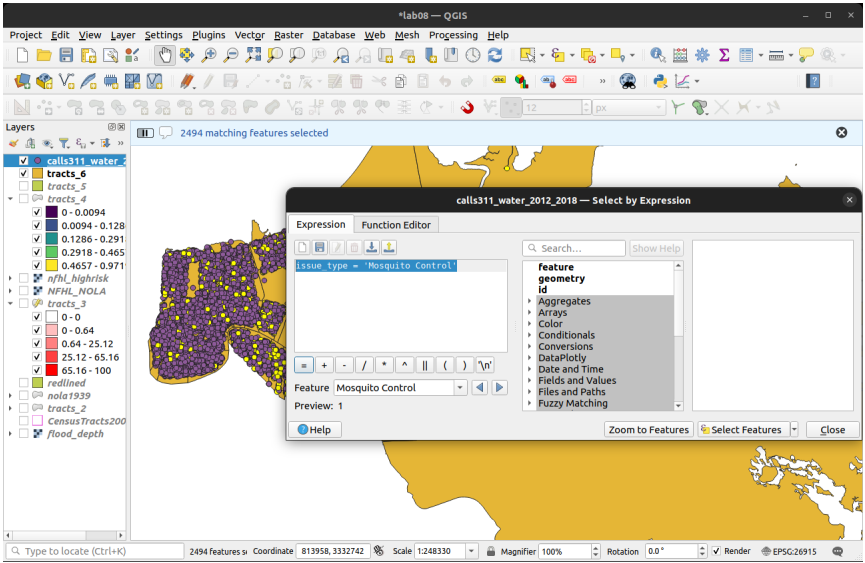

Set `Polygons` to `tracts_6`, field name to `calls311_mosquito`, save as `tracts_7`

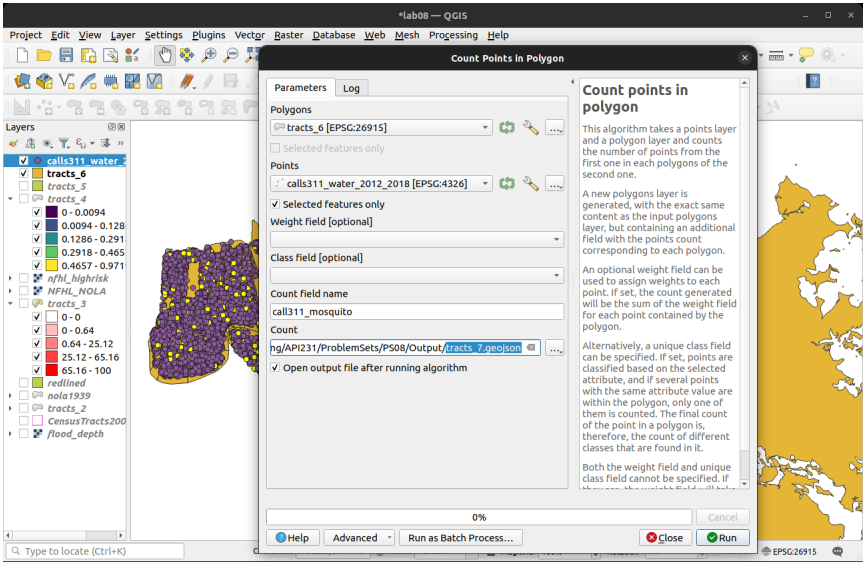

Let's make some per capita measures. Open the Field Calculator for `tracts_7`, create a new field with name `call311_flood_1000` of type `Decimal number`. Set Expression to `call311_streetflood / POP_TOTAL * 1000`

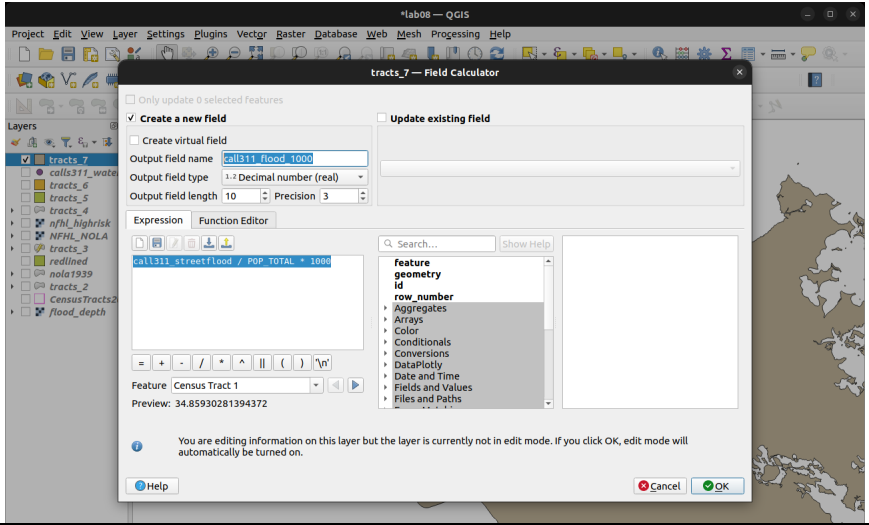

Also create a new field named `call311_basin_1000` of type `Decimal number`. Set Expression to `call311_catchbasin / POP_TOTAL * 1000`

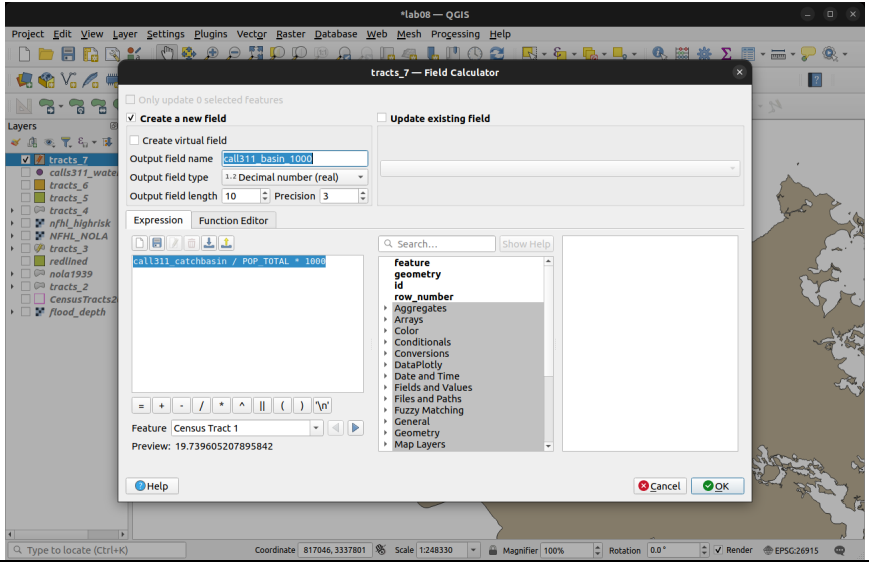

And finally, a new field with name `call311_mosquito_1000` of type `Decimal number`. Set Expression to `call311_mosquito / POP_TOTAL * 1000`

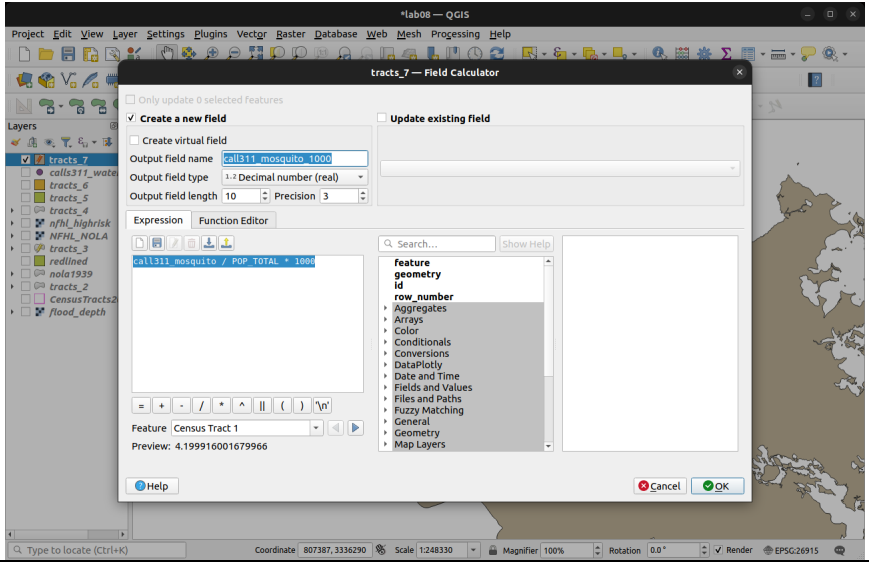

Now is a good time to save your progress, to `tracts_7` and to the project

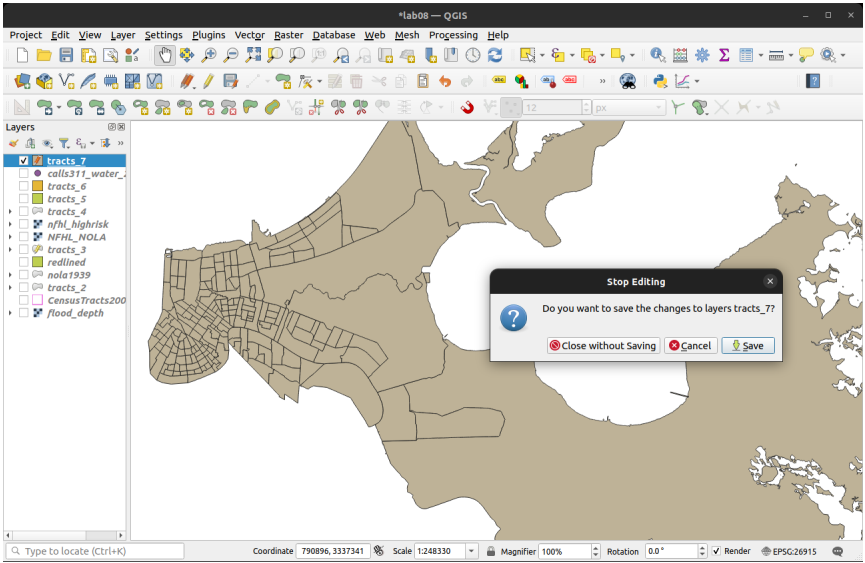

Problem Set 8

Your assignment (if using QGIS): create two scatterplots

1. Flood risk and Katrina flood depth
   - `highrisk_mean` on x-axis
   - `flood_mean` on y-axis
   - no legend
   - axes and plot title properly labeled as on next slide
   - name the file `nfhl_katrina.png`
2. Flood risk and per capita 311 calls about street flooding
   - `highrisk_mean` on x-axis
   - `call311_flood_1000` on y-axis
   - no legend
   - axes and plot title properly labeled as on next slide
   - name the file `nfhl_311flood.png`

- upload both plots to Canvas

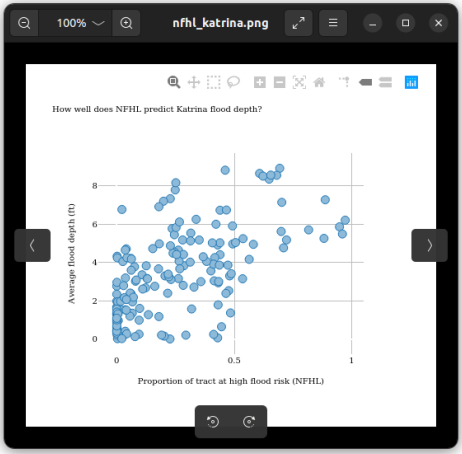

Figure 14: Can you make this?

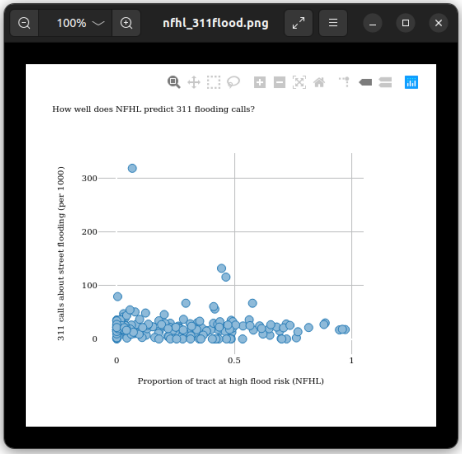

Figure 15: And this?

R

Loading R packages

To implement these steps in R, we will be using the `sf`, `terra` and `ggplot2` packages

```
library(sf)
library(terra)
library(ggplot2)
```

NOTE: The demo code for R is in `ps08_demo.R` on RStudio Cloud, and in `PS08.zip` (posted on Canvas).

Zonal statistics

Let's load the *census tracts data* into R, using `sf::read_sf()`:

```
tracts2000 <- sf::read_sf("Data/Census/CensusTracts2000.geojson")
```

Check the coordinate reference system of these data

```
sf::st_crs(tracts2000)
```

```
## $input
## [1] "NAD83 / UTM zone 15N"
```

Load the *Katrina flood depth data* into R, using `terra::rast()`:

```
flood_depth <- terra::rast("Data/Katrina/flood_depth.tif")
```

Is the coordinate reference system the same as for `tracts2000`?

```
sf::st_crs(flood_depth)==sf::st_crs(tracts2000)
```

```
## [1] TRUE
```

Yay! That means we can move forward with zonal statistics.

Calculate *zonal statistics*: mean and maximum flood depth in each tract

```
tracts2000$flood_mean <- terra::zonal(x=flood_depth,
  z=terra::vect(tracts2000), fun=mean, na.rm=TRUE)[,1]
tracts2000$flood_max <- terra::zonal(x=flood_depth,
  z=terra::vect(tracts2000), fun=max, na.rm=TRUE)[,1]
```

Plot the results. Note that, unlike QGIS, R assigns `NA` values (instead of 0) to polygons that do not overlap with the raster layer.

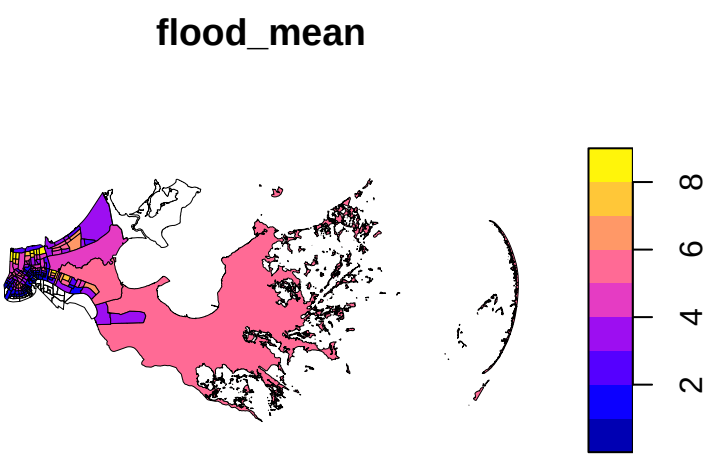

Let's load the *HOLC redlining data* into R, using `sf::read_sf()`:

```
nola1939 <- sf::read_sf("Data/Inequality/nola1939.json")
```

Is the coordinate reference system the same as for `tracts2000`?

```
sf::st_crs(nola1939)==sf::st_crs(tracts2000)
```

```
## [1] FALSE
```

Uh-oh! Looks like we need to reproject `nola1939`

```
nola1939 <- sf::st_transform(nola1939,crs=sf::st_crs(tracts2000))
sf::st_crs(nola1939)==sf::st_crs(tracts2000)
```

```
## [1] TRUE
```

All clear!

Subset the HOLC data to just “grade D”:

```
redlined <- nola1939[nola1939$grade=="D",]
```

```
plot(nola1939["grade"])
```

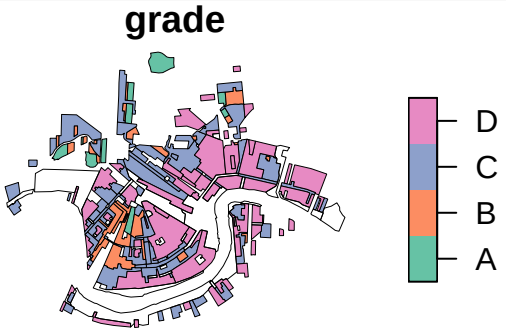


```
plot(redlined["grade"])
```

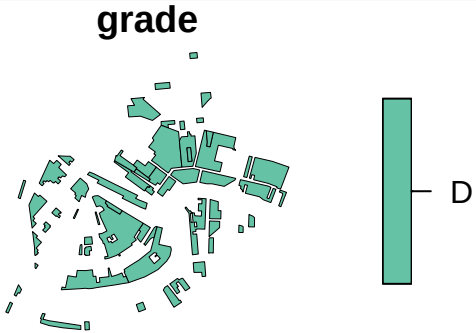

Overlap analysis in R takes a few more steps than in QGIS.
Calculate areas of `tracts2000`, and of `tracts2000`+`redlined` intersections:

```
tracts2000$area <- sf::st_area(tracts2000)
tract2red <- sf::st_intersection(tracts2000,sf::st_union(redlined))
tract2red$area_ix <- sf::st_area(tract2red)
```

Aggregate percentage area overlaps by census tract, add this field to `tracts2000`:

```
redlined_pc <-stats::aggregate(list(
  redlined_pc=as.numeric(tract2red$area_ix/tract2red$area)),
  by=list(GISJOIN=tract2red$GISJOIN),FUN=sum,na.rm=TRUE)
tracts2000 <- merge(tracts2000,redlined_pc,by="GISJOIN",all.x=TRUE)
```

Replace missing values with 0s:

```
tracts2000$redlined_pc[is.na(tracts2000$redlined_pc)] <- 0
```

Inspect the results.

```
plot(tracts2000["redlined_pc"])
```

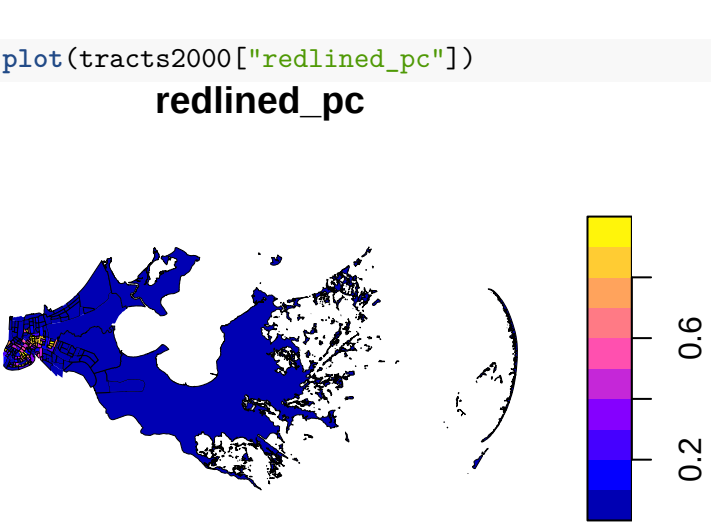

Make a *scatterplot* of flooding depth at different levels of redlining:

```
plot(x=tracts2000$redlined_pc,y=tracts2000$flood_mean)
```

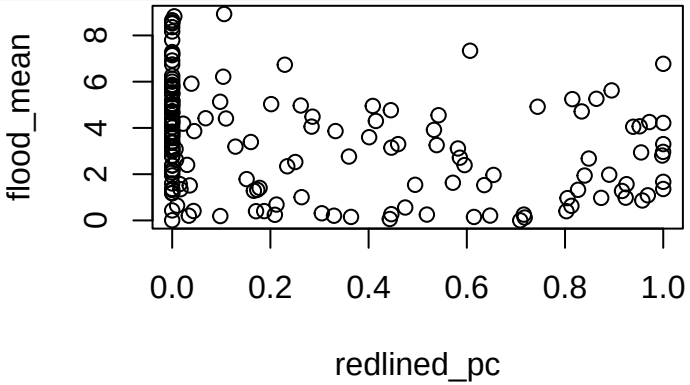

Make a *scatterplot* using ggplot2, with fitted LOESS curve:

```
ggplot2::ggplot(tracts2000, ggplot2::aes(x=redlined_pc,y=flood_mean))+
  ggplot2::geom_point()+
  ggplot2::geom_smooth()
```

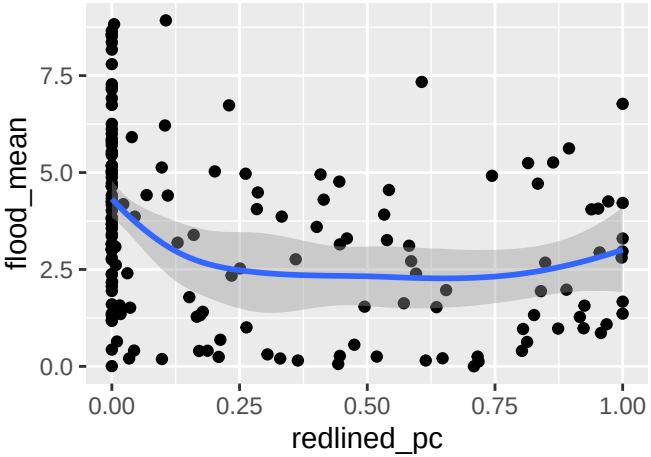

Make a *boxplot* of flooding depth with vs. without redlining:

```
tracts2000$redlined_any <- 1*(tracts2000$redlined_pc>0)
graphics::boxplot(flood_mean~redlined_any,data=tracts2000)
```

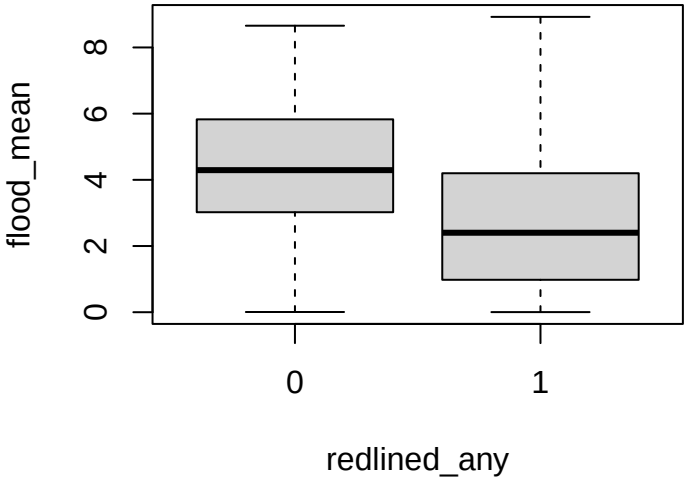

Make a *scatterplot* of race and redlining:

```
tracts2000$PCT_BLACK <- tracts2000$RACE_BLACK/tracts2000$POP_TOTAL*100
plot(x=tracts2000$PCT_BLACK,y=tracts2000$flood_mean)
```

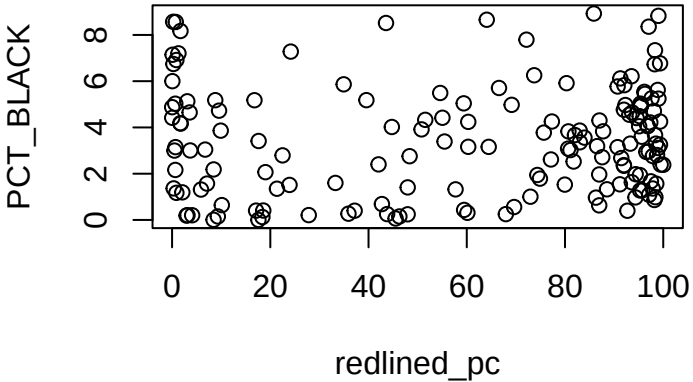

Raster reclassification

Load the *NFHL data* into R, and check CRS compatibility:

```
NFHL_NOLA <- terra::rast("Data/FEMA/NFHL_NOLA.tif")
sf::st_crs(NFHL_NOLA)==sf::st_crs(tracts2000)
```

Plot the raster:

```
terra::plot(NFHL_NOLA)
```

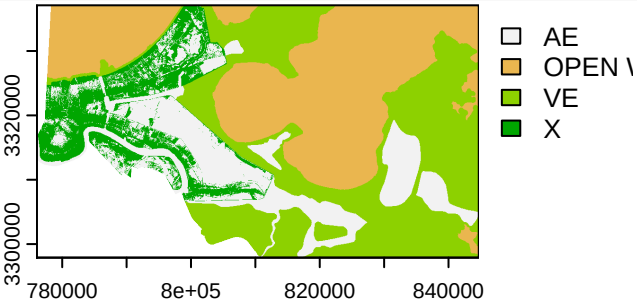

Reclassify the NFHL raster to include just zone "AE" (high risk):

```
nfhl_highrisk <- 1*(NFHL_NOLA=="AE")
terra::plot(nfhl_highrisk)
```

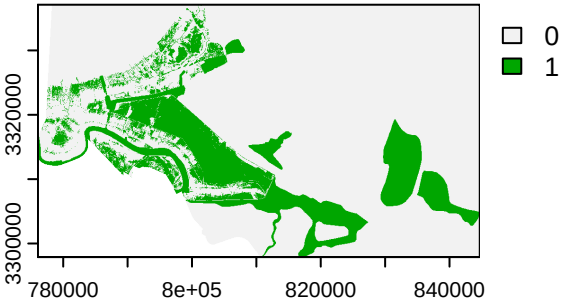

Calculate *zonal statistics*: proportion of each tract at high flood risk

```
tracts2000$highrisk_mean <- terra::zonal(x=nfhl_highrisk,
  z=terra::vect(tracts2000), fun=mean,na.rm=TRUE)[,1]
```

This may take a minute or two to run...

Inspect the results.

```
plot(tracts2000["highrisk_mean"])
```

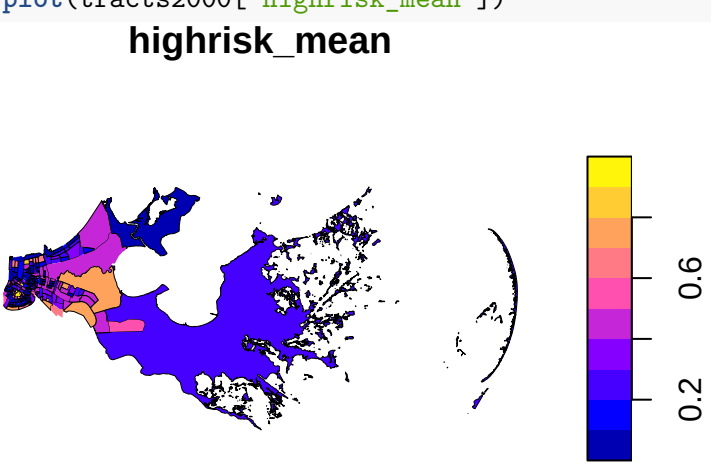

Load the *311 data* into R, and convert to spatial points:

```
calls311 <- read.csv("Data/Call311/calls311_water_2012_2018.csv")
calls311 <- sf::st_as_sf(calls311,coords=c("longitude","latitude"),
  crs=4326)
```

Reproject the 311 data to same CRS as `tracts2000`:

```
calls311 <- sf::st_transform(calls311,sf::st_crs(tracts2000))
```

Point-in-polygon analysis: 311 calls per 1000 residents, by type

```
tracts2000$call311_flood_1000 <- lengths(sf::st_intersects(
  tracts2000,calls311[calls311$issue_type=="Street Flooding/Drainage",]
  )/tracts2000$POP_TOTAL*1000
tracts2000$call311_basin_1000 <- lengths(sf::st_intersects(
  tracts2000,calls311[calls311$issue_type=="Catch Basin Maintenance",])
  )/tracts2000$POP_TOTAL*1000
tracts2000$call311_mosquito_1000 <- lengths(sf::st_intersects(
  tracts2000,calls311[calls311$issue_type=="Mosquito Control",])
  )/tracts2000$POP_TOTAL*1000
```

Inspect the results.

```
plot(tracts2000["call311_flood_1000"])
```

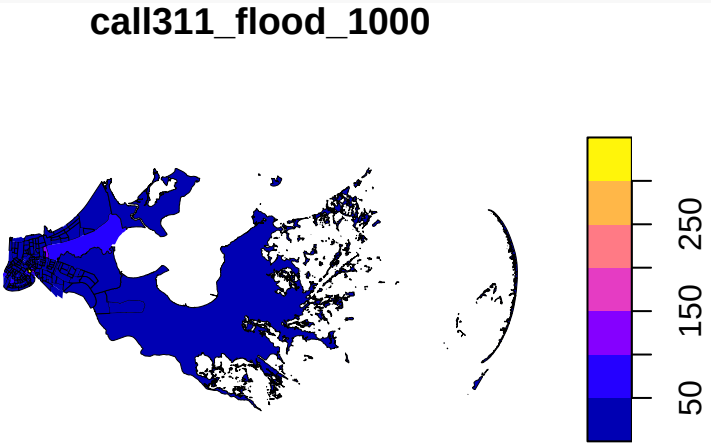

Problem Set 8

Your assignment (if using R): create two scatterplots

1. Flood risk and Katrina flood depth
   - `highrisk_mean` on x-axis
   - `flood_mean` on y-axis
   - no legend
   - axes and plot title properly labeled as on next slide
   - name the file `nfhl_katrina_R.png`
2. Flood risk and per capita 311 calls about street flooding
   - `highrisk_mean` on x-axis
   - `call311_flood_1000` on y-axis
   - no legend
   - axes and plot title properly labeled as on next slide
   - name the file `nfhl_311flood_R.png`

- use either R base plots or `ggplot2`
- upload both plots to Canvas

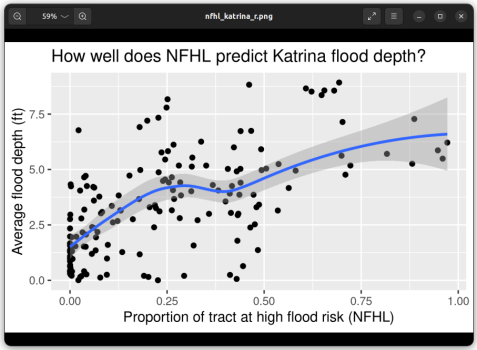

Figure 16: Can you make this?

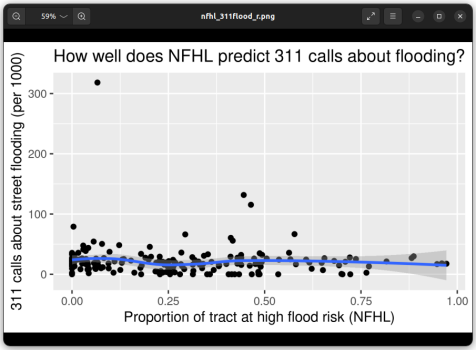

Figure 17: And this?