

# **API-231 / GIS-PubPol**

Meeting 20 (“Russian-Ukrainian War” Walk Through)

Yuri M. Zhukov  
Visiting Associate Professor of Public Policy  
Harvard Kennedy School

April 11, 2024



Figure 1: News = data about **who did what to whom, when and where**

Who	Did what	To whom	When	Where	Source
Russia	rocket strike	Ukraine	2/24/2022	Kyiv	CNN
Russia	rocket strike	Ukraine	2/24/2022	Kharkiv	CNN

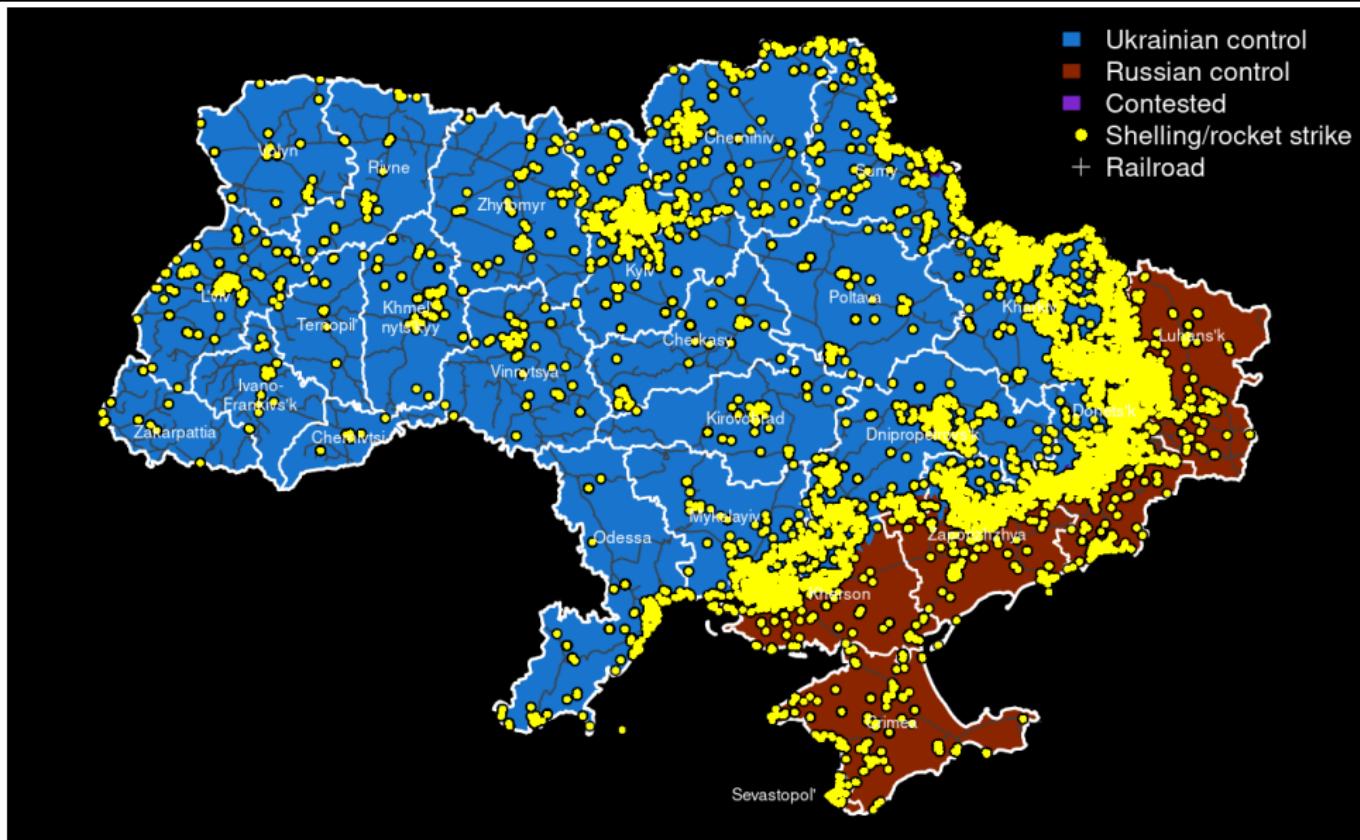


Figure 2: Near-real time event, territorial control data on Russian invasion of Ukraine

# Overview

**Motivation:** during Russia's 2022 invasion of Ukraine...

1. Russia has become a hermetically-sealed information environment
  - a) media required to stick to MoD press releases
  - b) cannot use word "war" when describing "special military operation"
  - c) up to 100K ruble fine for publicly "discrediting" Russian army
  - d) up to 15-year sentence for "knowingly false information" about war
  - e) last independent media shut down (e.g. TV Rain, Echo of Moscow)
  - f) Facebook, Twitter, Instagram, VPNs blocked
2. Ukrainian media more free, but vulnerable
  - a) TV news sometimes broadcasts from basements, bomb shelters
  - b) Russia has targeted TV towers, cut electricity, cell service
  - c) all national TV channels merged onto one platform under martial law
  - d) radio silence on Ukrainian casualties, ongoing operations

**Solution:** use machine learning, remote sensing to track events on the ground

## Tracking the War in Near-Real Time

## What are “event data”?

1. Incident-level data on “who did what to whom, when and where”
  - a) “who”: initiator of action (subject)
  - b) “did what”: description of action/tactic (verb)
  - c) “to whom”: target of action (object)
  - d) “when”: time/date of event
  - e) “where”: location of event
2. Types of events (examples we have used in this class)
  - a) political violence
  - b) bike crashes in NYC
  - c) crimes in DC
  - d) 311 calls about flooding in New Orleans
3. Sources of data
  - a) media/open sources (including social media)
  - b) government records/archives
  - c) remote sensing

## VIINA: Violent Incident Information from News Articles

1. Near-real time event data on Russian invasion of Ukraine (updated daily)
  - a) based on news reports from Ukrainian and Russian media, geocoded and classified with Bidirectional Encoder Representations from Transformers (BERT)
  - b) each event is accompanied by full source info, text and URLs
2. Data on territorial control at municipality level (updated daily)
  - a) based on vectorized georeferenced maps (e.g. DeepState, Wikipedia)
  - b) “boosted” by VIINA event data on changes in control

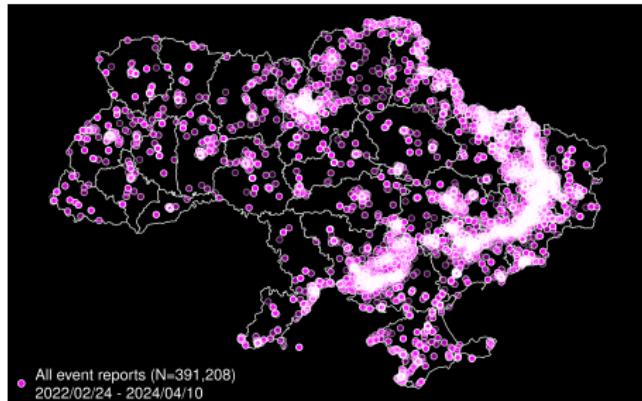


Figure 3: Events

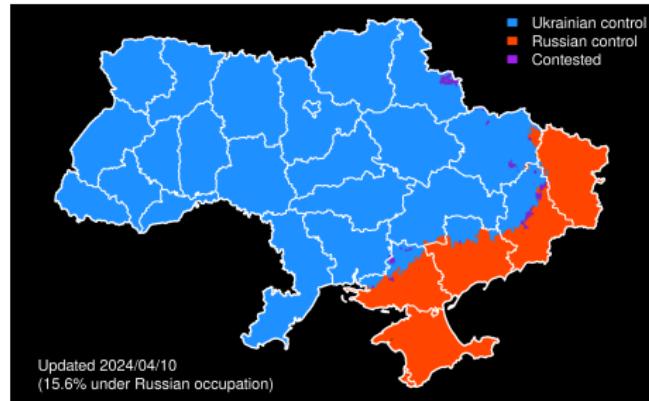


Figure 4: Control

**Problem:** 1 news report  $\neq$  1 unique event

1. We can characterize each event as a unique configuration of [subject]-[verb]-[object]-[time]-[location]  
(i.e. who did what to whom, when and where)
2. We can learn about these events from news reports
  - a) one report of one event ("A attacked B")
  - b) multiple reports of one event ("A attacked B", "B attacked by A")
  - c) multiple events in one report ("A attacked B, C attacked A")
  - ⋮
  - d)  $N$  reports of an unknown # of events

**Question:** do these news reports refer to the same event?

Date	Source	English translation
2022/02/25	Interfax.ua	Russian forces purposefully shelling residential buildings in Kharkiv – head of regional administration Synyehubov
2022/02/25	24tv.ua	Shell hits residential building in Kharkiv: casualties possible – frightening photos

Who	Did what	To whom	When	Where
✓	✓	✓	✓	

**Question:** do these news reports refer to the same event?

Date	Source	English translation
2022/08/10	liveuamap	In Kharkiv as a result of Russian shelling one person is wounded
2022/08/15	24tv.ua	Woman, wounded during shelling of Saltivka in Kharkiv, died in hospital

Who	Did what	To whom	When	Where
	✓			✓

## Coreference resolution (CR)

Process of resolving multiple references to same physical object or event

1. Why are duplicates a problem?
  - a) duplicates are a threat to causal inference
  - b) over-reporting of events may be correlated with unobservables (e.g. media presence, perceived "newsworthiness" or novelty)
  - c) duplicates make it harder to assess ground truth about violence
  - d) this problem affects both machine-coded and hand-coded data
2. Applications to event data:
  - a) remove exact textual duplicates ("bare minimum")
  - b) "1 per day" filter (if two reported events are of the same type, and were reported in same location on same day, then they are references to the same event)
  - c) MELTT spatio-temporal filter (match based on co-occurrence in space and time)
  - d) model-based methods (e.g. convolutional neural networks, transformers)

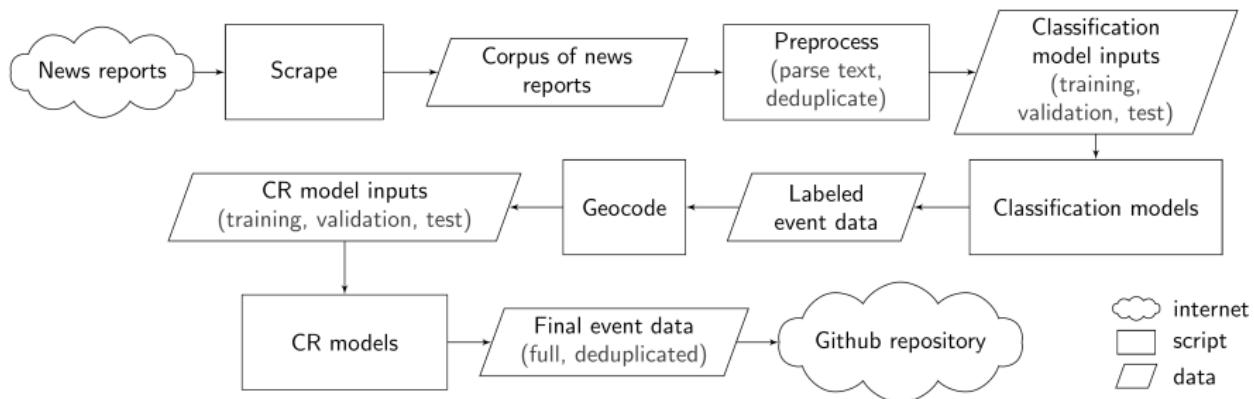


Figure 5: VIINA workflow

VIINA turns online news articles & social media into geocoded event data by:

1. Scraping online news & social media, preprocessing the raw text
2. Classifying the news reports by actor and tactics with large language models
3. Assigning geo coordinates based on locations mentioned in reports
4. Identifying (but not removing) likely duplicate events

## Other Near-Real Time Data Sources

## Historical weather & climate raster data (partial list)

Dataset	ACLED	GDELT	ICEWS	VIINA
Data on Violence?	✓	✓	✓	✓
Data on Control?				✓
Fully Automated?		✓	✓	✓
Text Descriptions?	✓			✓
Source URLs?		✓		✓
Events in 1st Year	40,448	778,350	27,858	113,446
Unique Event Locations	2,430	1,762	581	9,771
Update Frequency	1 week	Daily	1-2 months	Daily
Event Types	24	50	105	23
Sources	97	8,887	126	30
English-Only?	No	Yes	Yes	No
Ukrainian Sources (%)	74.2	10.1	4.2	92.5
Russian Sources (%)	14	11.8	12	7.5
Unknown Sources (%)	0	0	47.3	0

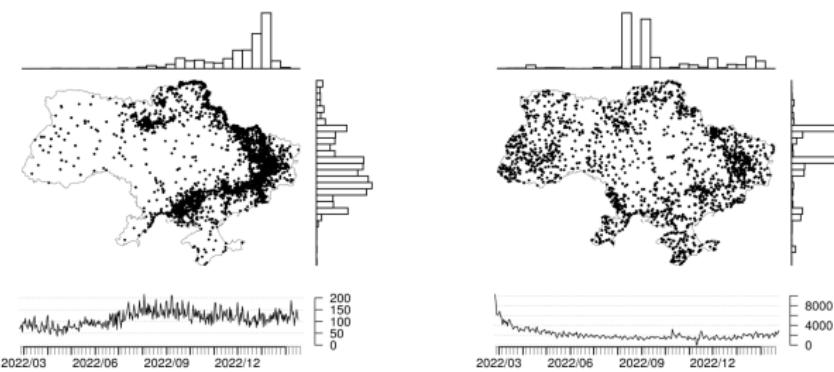


Figure 6: ACLED

Figure 7: GDELT

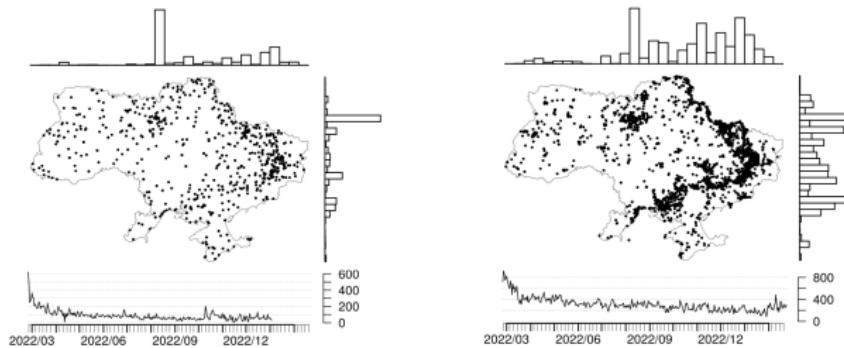


Figure 8: ICEWS

Figure 9: VIINA

## Near-real time remote sensing data (partial list)

Type	Source/link	Spatial resolution	Frequency	Free?
Fire anomalies	FIRMS	Points	Daily	✓
Night lights	VIIRS	Raster	Nightly	✓
Vegetation	NDVI	Raster	2 weeks	✓
Meteorological events	NASA Worldview	Raster	Daily	✓
Reflectance (photos)	NASA Worldview	Raster	Hourly/Daily	✓

What kinds of events can remote sensing capture that media cannot? (+ vice versa)

# Vignettes

## Overview of lab exercise

1. How much of Ukraine's territory does Russia occupy?
2. Compare media reports to remote sensing data on fire anomalies.

We will work with a (very large) dataset on territorial control in Ukraine

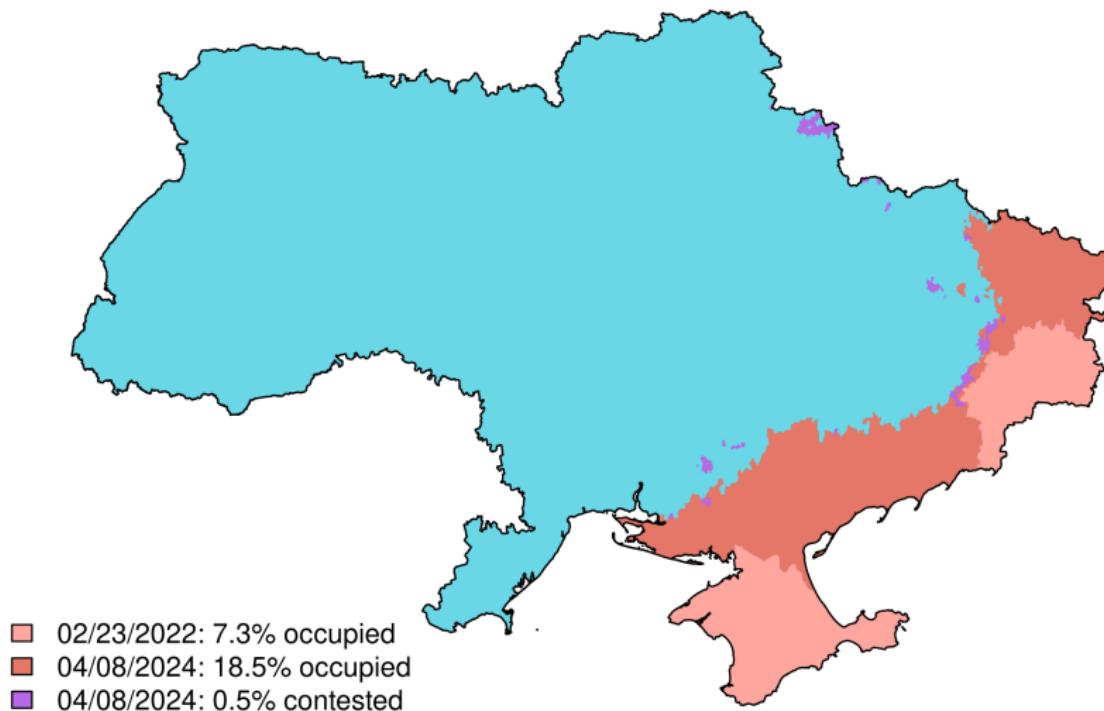


Figure 10: Vignette 1

We will then integrate NASA's data on active fires with VIINA event reports

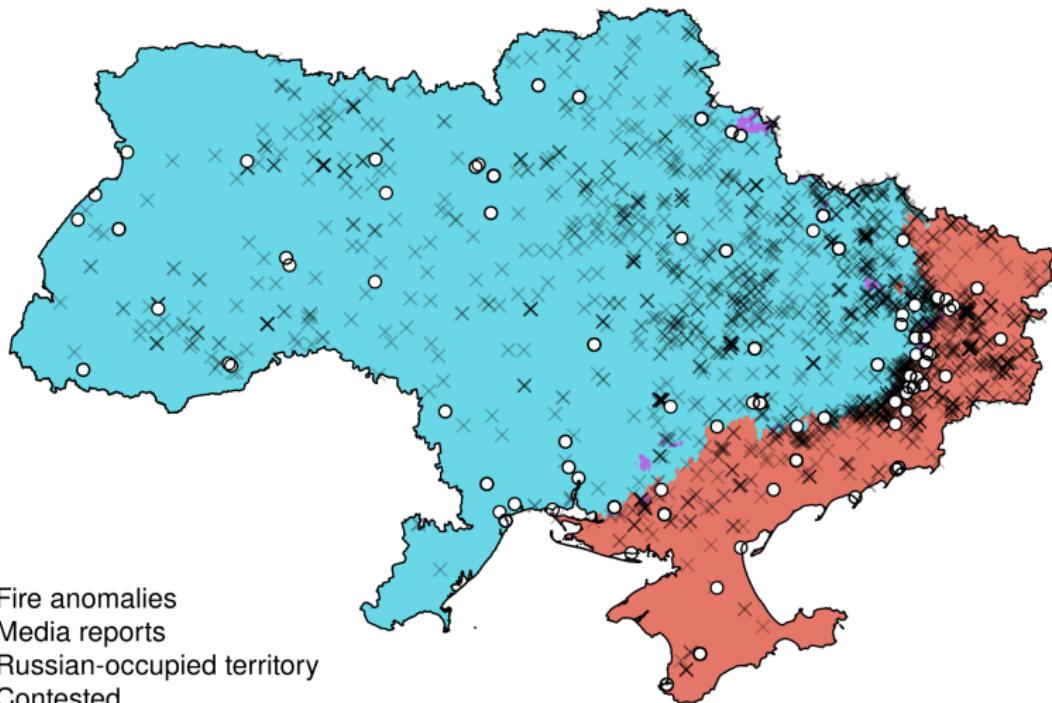


Figure 11: Vignette 2 / Step 1

... and identify locations that may be overlooked media vs. fires data

## Remote sensing vs. media reports on War in Ukraine

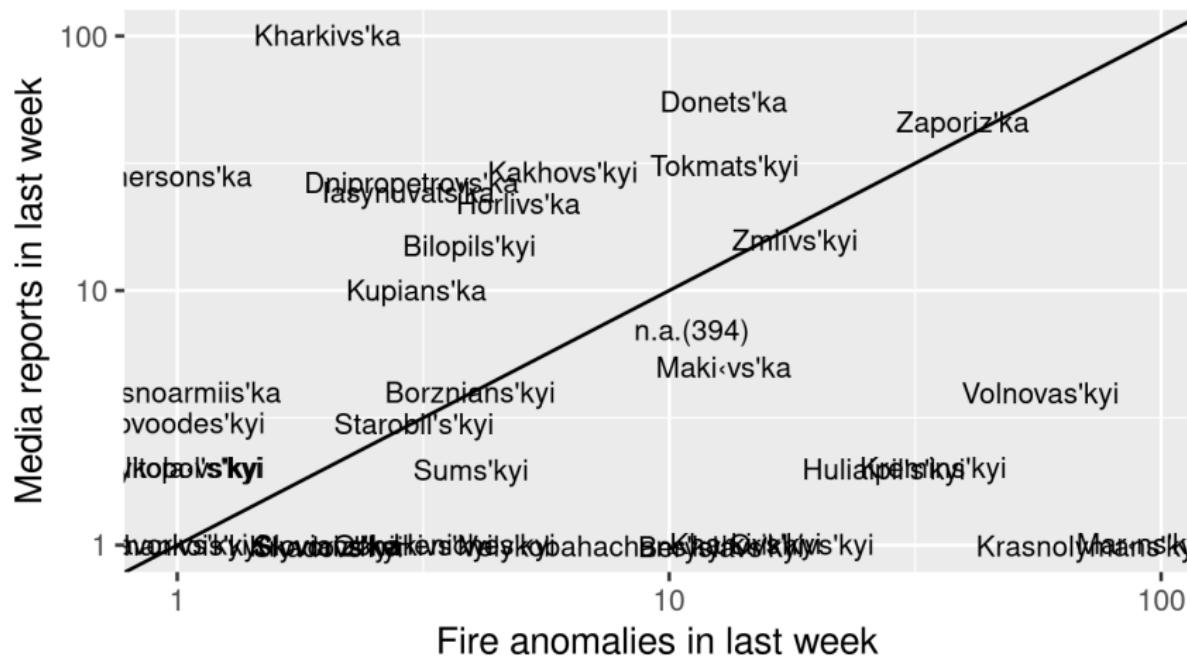
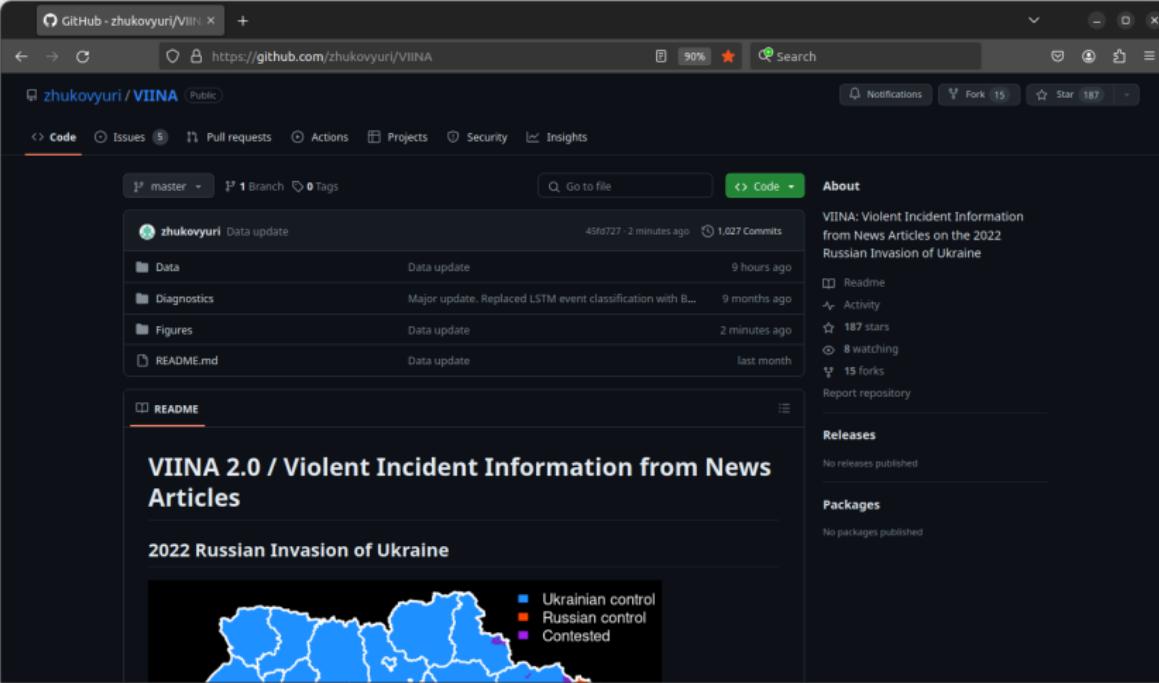


Figure 12: Vignette 2 / Step 2

We can obtain **data on territorial control and media reports** from  
[github.com/zhukovskyuri/VIINA](https://github.com/zhukovskyuri/VIINA)



The screenshot shows the GitHub repository page for 'zhukovskyuri/VIINA'. The repository is public and has 187 stars, 15 forks, and 8 watchers. It contains 1 branch and 1 commit. The 'Code' tab is selected, showing a list of recent updates:

Author	Commit Message	Time Ago
zhukovskyuri	Data update	45fd727 · 2 minutes ago
	Data	9 hours ago
	Diagnostics	Major update. Replaced LSTM event classification with B...
	Figures	Data update
	README.md	Data update

The 'About' section describes VIINA as 'Violent Incident Information from News Articles on the 2022 Russian Invasion of Ukraine'. It includes links to the Readme, Activity, and repository statistics (187 stars, 8 watching, 15 forks). The 'README' section highlights 'VIINA 2.0 / Violent Incident Information from News Articles' and '2022 Russian Invasion of Ukraine', featuring a map of Ukraine with territorial control status.

There are several datasets here. The ones we need are `control_latest` and `event_info_latest_2024`

VIINA/ВІИНА/ВОЙНА/WAR 2.0 is a near-real time multi-source event data system for the 2022 Russian Invasion of Ukraine. These data are based on news reports from Ukrainian and Russian media, which were geocoded and classified into standard conflict event categories through machine learning. In addition to raw events, VIINA also includes data on territorial control, at the level of individual populated places.

These data are GIS-ready, with temporal precision down to the minute. Each observation is accompanied by full source information, text and URLs.

VIINA is updated daily, and is freely available for use by students, journalists, policymakers, and everyday researchers.

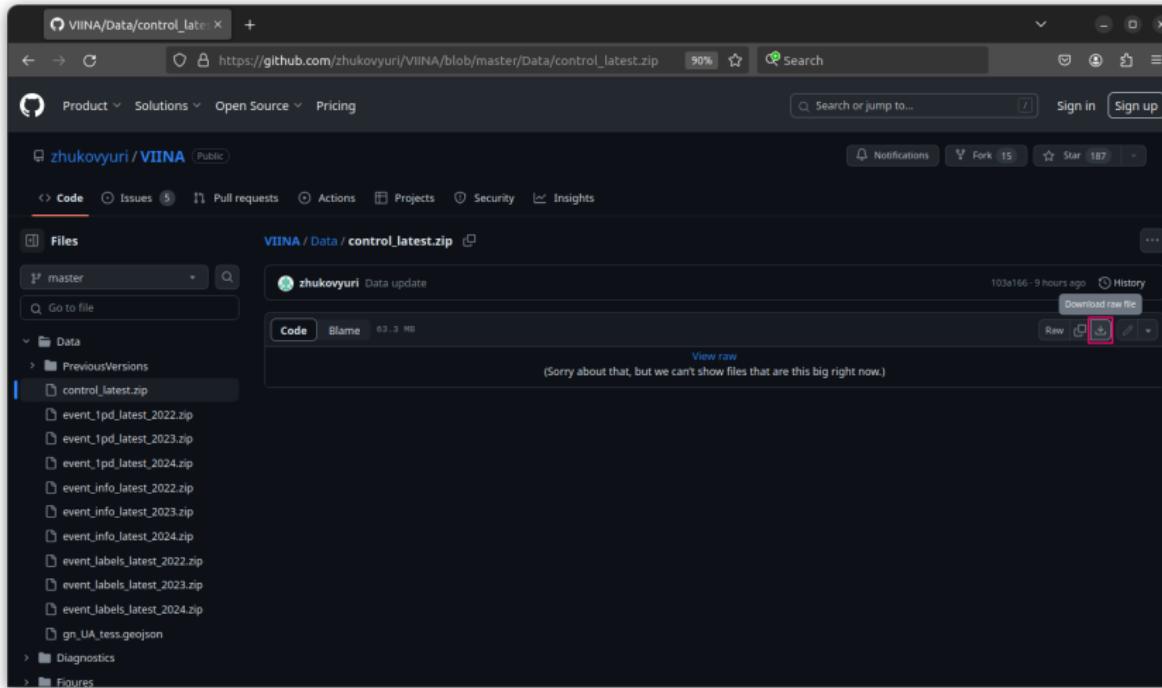
The most recent versions these data are available as a comma-delimited-text (csv) files within the following compressed ZIP archives:

- [Data/control\\_latest.zip](#) | Territorial control status
- [Data/event\\_info\\_latest\\_2022.zip](#) | Raw event reports for 2022 (locations, dates, urls, headlines)
- [Data/event\\_info\\_latest\\_2023.zip](#) | Raw event reports for 2023 (locations, dates, urls, headlines)
- [Data/event\\_info\\_latest\\_2024.zip](#) | **Raw event reports for 2024 (locations, dates, urls, headlines)**
- [Data/event\\_labels\\_latest\\_2022.zip](#) | Event reports for 2022, labeled by actor and tactic (from BERT model)
- [Data/event\\_labels\\_latest\\_2023.zip](#) | Event reports for 2023, labeled by actor and tactic (from BERT model)
- [Data/event\\_labels\\_latest\\_2024.zip](#) | Event reports for 2024, labeled by actor and tactic (from BERT model)
- [Data/event\\_1pd\\_latest\\_2022.zip](#) | De-duplicated event reports and labels for 2022 ("one-per-day" filter)
- [Data/event\\_1pd\\_latest\\_2023.zip](#) | De-duplicated event reports and labels for 2023 ("one-per-day" filter)
- [Data/event\\_1pd\\_latest\\_2024.zip](#) | De-duplicated event reports and labels for 2024 ("one-per-day" filter)

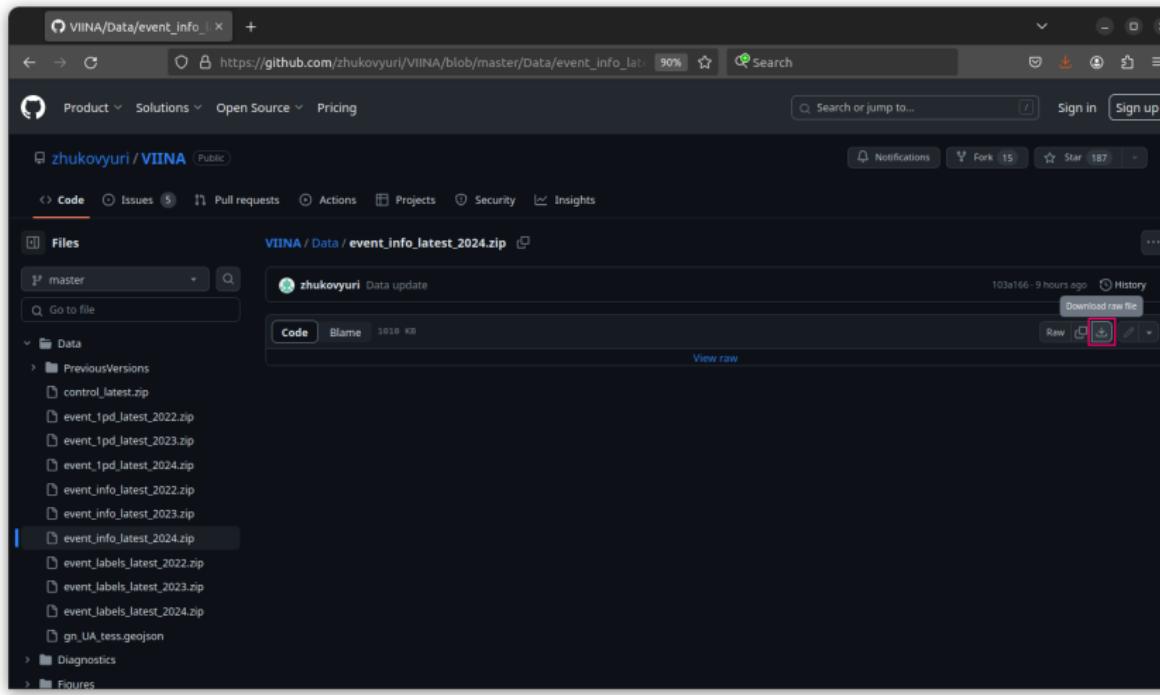
Note that each event data release includes both raw event reports (`event_info`, `event_labels`) and a simplified, de-duplicated data file (`event_1pd`). The latter uses a "one-per-day" filter to remove potential duplicate event reports, by treating multiple event reports of the same type (i.e. same combination of actor and tactic labels) in the same populated place on the same day as a single unique event.

[https://github.com/zhukovyuri/VIINA/tree/master/Data/control\\_latest.zip](https://github.com/zhukovyuri/VIINA/tree/master/Data/control_latest.zip)

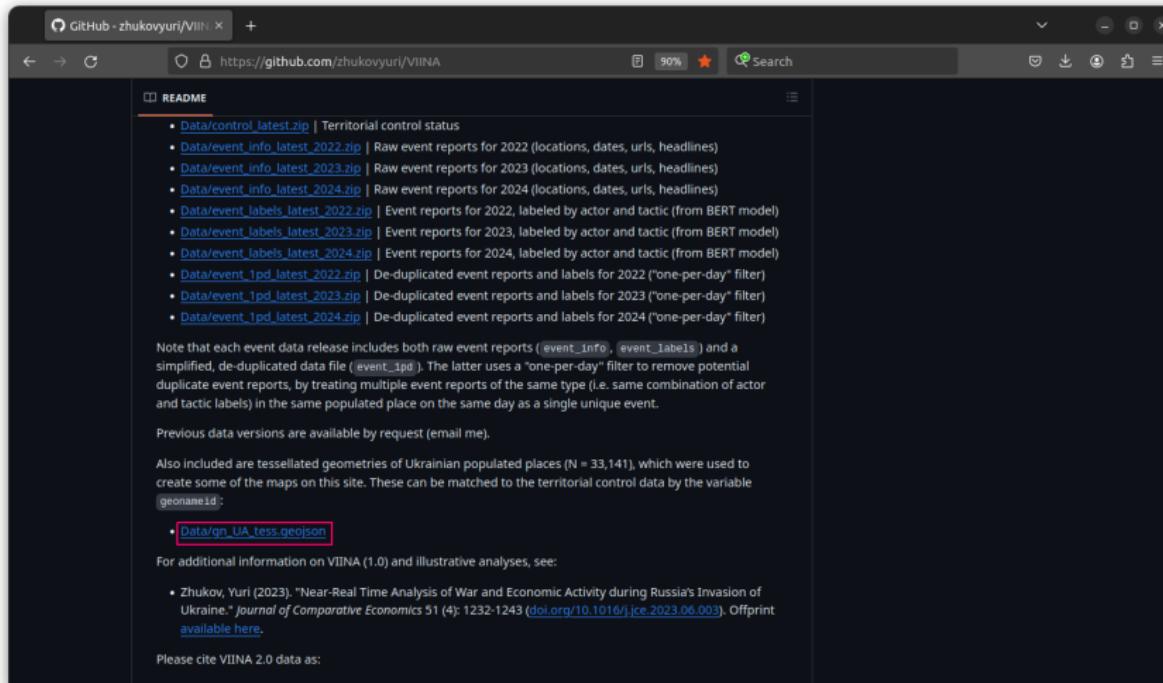
Go to `control_latest.zip` and download the file by clicking on the “Download raw file” button



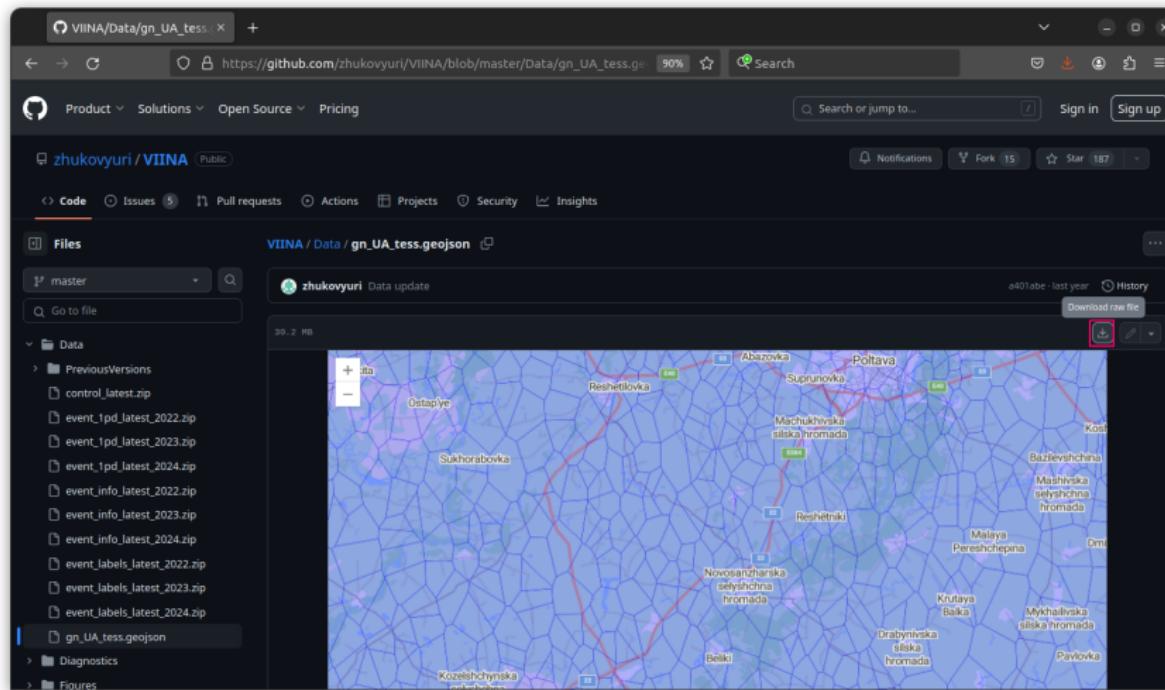
Do the same thing for event\_info\_latest\_2024.zip



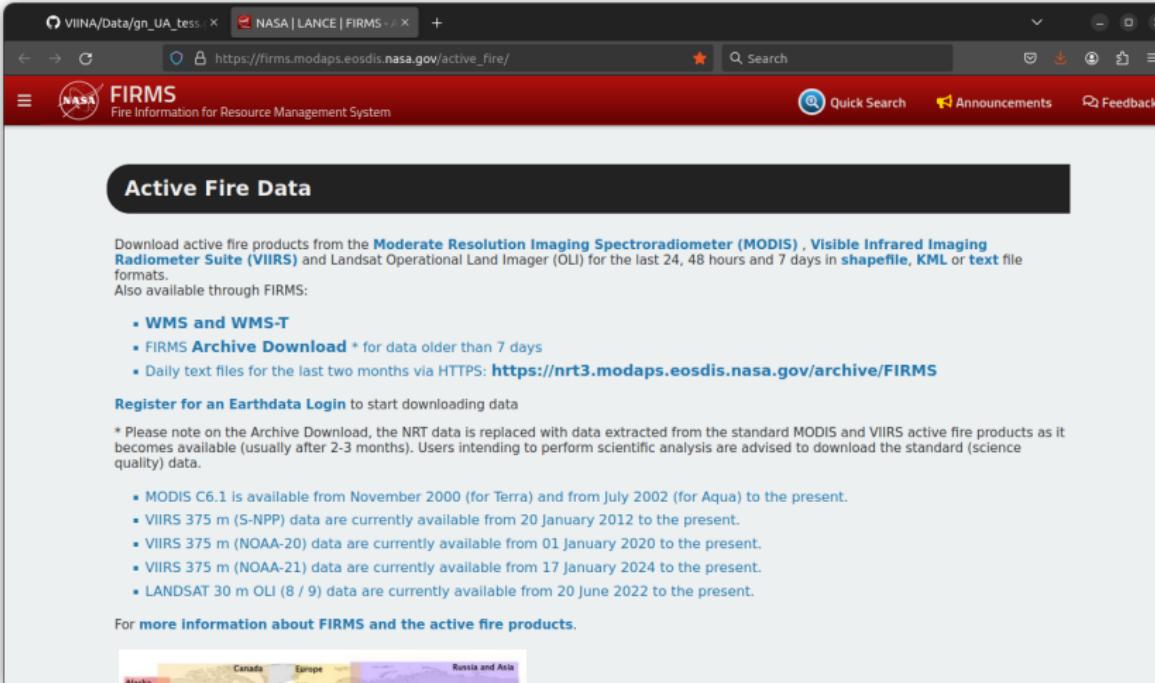
While we're here, let's also grab the GIS boundaries for Ukrainian populated places,  
`gn_UA_tess.geojson`



We can also get this file through the “Download raw file” link



Let's now get the **FIRMS Active Fires data** from  
[firms.modaps.eosdis.nasa.gov/active\\_fire/](https://firms.modaps.eosdis.nasa.gov/active_fire/)



The screenshot shows a web browser window with the URL [https://firms.modaps.eosdis.nasa.gov/active\\_fire/](https://firms.modaps.eosdis.nasa.gov/active_fire/). The page title is "FIRMS" and the subtitle is "Fire Information for Resource Management System". A red header bar contains links for "Quick Search", "Announcements", and "Feedback". The main content area has a dark header "Active Fire Data". Below it, text explains how to download active fire products from MODIS, VIIRS, and Landsat OLI, available in shapefile, KML, or text formats, and through WMS, FIRMS Archive Download, and daily text files. It also mentions the Earthdata Login for registration. A note about the archive download states that NRT data is replaced by standard MODIS and VIIRS data. A list of data availability includes MODIS C6.1 from November 2000, VIIRS 375 m from January 2012, VIIRS 375 m (NOAA-20) from January 2020, VIIRS 375 m (NOAA-21) from January 2024, and LANDSAT 30 m OLI from June 2022. At the bottom, there is a map showing the coverage areas of the data, including Alaska, Canada, Europe, and Russia and Asia.

Download active fire products from the **Moderate Resolution Imaging Spectroradiometer (MODIS)**, **Visible Infrared Imaging Radiometer Suite (VIIRS)** and Landsat Operational Land Imager (OLI) for the last 24, 48 hours and 7 days in **shapefile**, **KML** or **text** file formats.  
Also available through FIRMS:

- [WMS and WMS-T](#)
- [FIRMS Archive Download](#) \* for data older than 7 days
- Daily text files for the last two months via [HTTPS: https://nrt3.modaps.eosdis.nasa.gov/archive/FIRMS](https://nrt3.modaps.eosdis.nasa.gov/archive/FIRMS)

[Register for an Earthdata Login](#) to start downloading data

\* Please note on the Archive Download, the NRT data is replaced with data extracted from the standard MODIS and VIIRS active fire products as it becomes available (usually after 2-3 months). Users intending to perform scientific analysis are advised to download the standard (science quality) data.

- MODIS C6.1 is available from November 2000 (for Terra) and from July 2002 (for Aqua) to the present.
- VIIRS 375 m (S-NPP) data are currently available from 20 January 2012 to the present.
- VIIRS 375 m (NOAA-20) data are currently available from 01 January 2020 to the present.
- VIIRS 375 m (NOAA-21) data are currently available from 17 January 2024 to the present.
- LANDSAT 30 m OLI (8 / 9) data are currently available from 20 June 2022 to the present.

For more information about **FIRMS** and the active fire products.

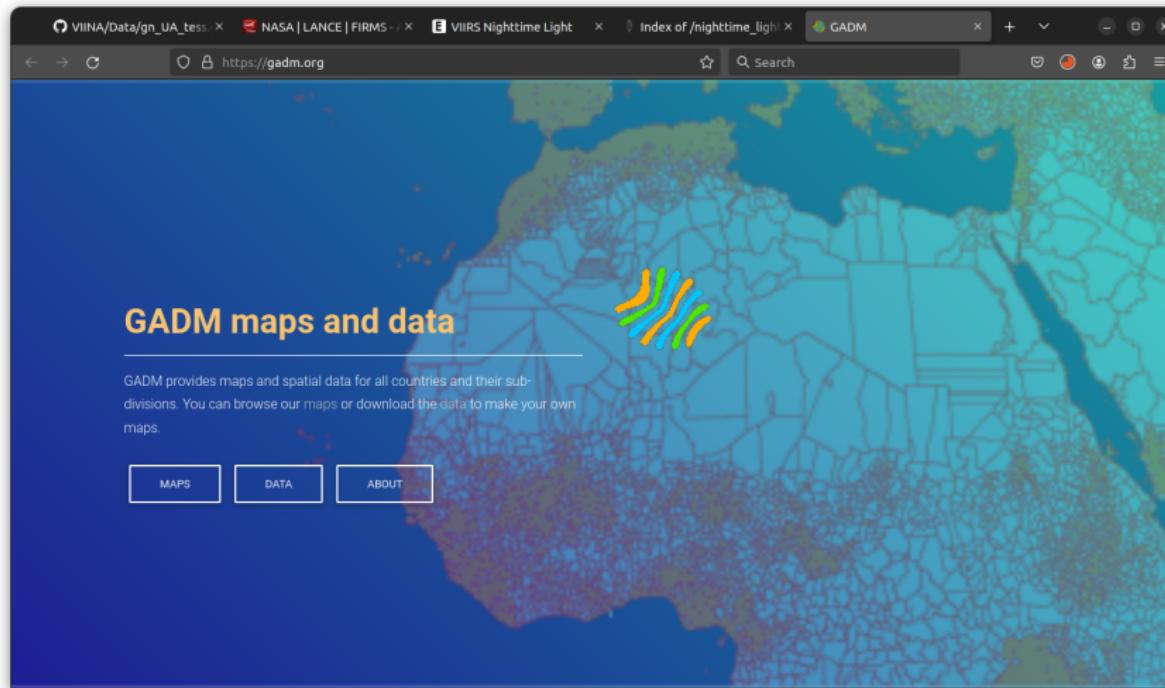
Alaska Canada Europe Russia and Asia

Scroll down to “Text Files (CSV)” and download the latest weekly (7d) data for “World” from “VIIRS 375m/NOAA-21”

The screenshot shows a web browser window with the URL [https://firms.modaps.eosdis.nasa.gov/active\\_fire/](https://firms.modaps.eosdis.nasa.gov/active_fire/). The page title is "Text Files (CSV)". A sidebar on the left lists "MODIS Collection 6.1" and "VIIRS 375m / S-NPP", "VIIRS 375m / NOAA-20", and "VIIRS 375m / NOAA-21". Below this, a note says "To keep file sizes to a minimum, the data are provided by region." A table compares file sizes for MODIS 1km, VIIRS 375m / S-NPP, VIIRS 375m / NOAA-20, VIIRS 375m / NOAA-21, and LANDSAT 30m across nine regions: World, Canada, Alaska, USA (Conterminous) and Hawaii, Central America, South America, Europe, North and Central Africa, and Southern Africa. The "VIIRS 375m / NOAA-21" column shows "24h 48h 7d" for all regions except the World, which is highlighted with a red border.

	MODIS 1km	VIIRS 375m / S-NPP	VIIRS 375m / NOAA-20	VIIRS 375m / NOAA-21	LANDSAT 30m
World	<a href="#">24h 48h 7d</a>	N/A			
Canada	<a href="#">24h 48h 7d</a>				
Alaska	<a href="#">24h 48h 7d</a>	N/A			
USA (Conterminous) and Hawaii	<a href="#">24h 48h 7d</a>				
Central America	<a href="#">24h 48h 7d</a>	N/A			
South America	<a href="#">24h 48h 7d</a>	N/A			
Europe	<a href="#">24h 48h 7d</a>	N/A			
North and Central Africa	<a href="#">24h 48h 7d</a>	N/A			
Southern Africa	<a href="#">24h 48h 7d</a>	N/A			

We will use **country-level and district-level boundaries** data from [gadm.org](https://gadm.org)



## Download the level-0 and level-2 files for Ukraine, in GeoJSON format

The screenshot shows a web browser window with the URL [https://gadm.org/download\\_country.html](https://gadm.org/download_country.html). The page title is "GADM". The main content is titled "Download GADM data (version 4.1)". A dropdown menu under "Country" is set to "Ukraine". Below it, under "Geopackage", "Shapefile", and "GeoJSON", the "level-0" and "level2" options are highlighted with red boxes. To the right is a map of Ukraine showing administrative boundaries at level 2, with blue regions and red outlines. Below the map, text states: "The coordinate reference system is longitude/latitude and the WGS84 datum." and "Description of file formats." At the bottom, a copyright notice reads "© 2018-2022 GADM - license".

Here is the full list of data sources and links:

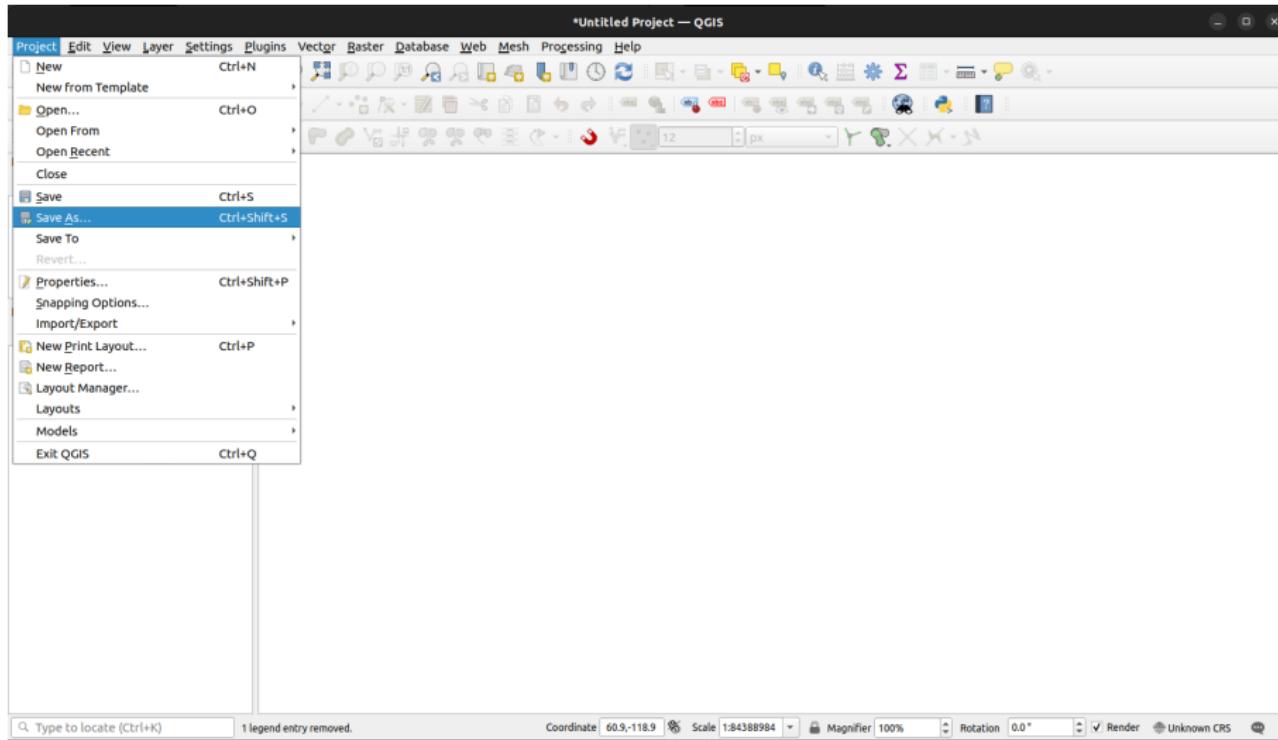
Category	Type	Format	Data source
Territorial control	Vector (polygons)	.csv, .geojson	VIINA
Media event reports	Vector (points)	.csv	VIINA
Active fires	Vector (points)	.csv	NASA FIRMS
Administrative borders	Vector (polygons)	.geojson	GADM

These are all in the WT04.zip file posted on Canvas.

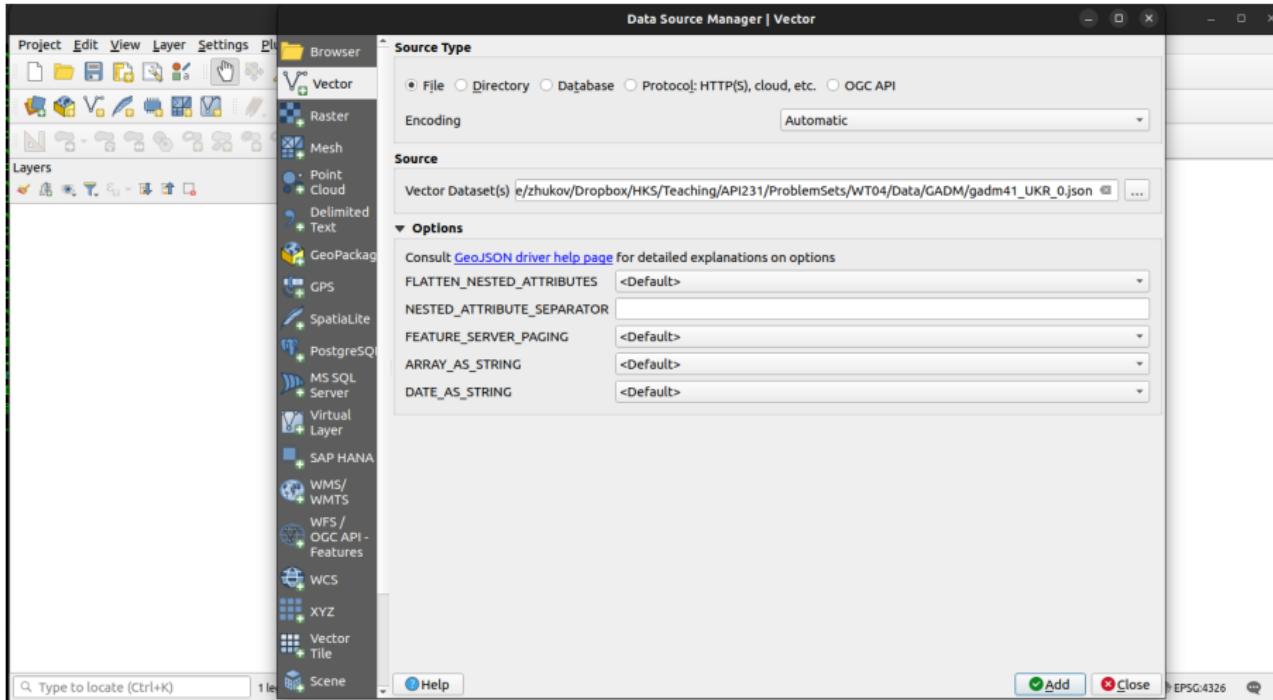
## How much of Ukraine's territory does Russia occupy?

# Always save your progress!

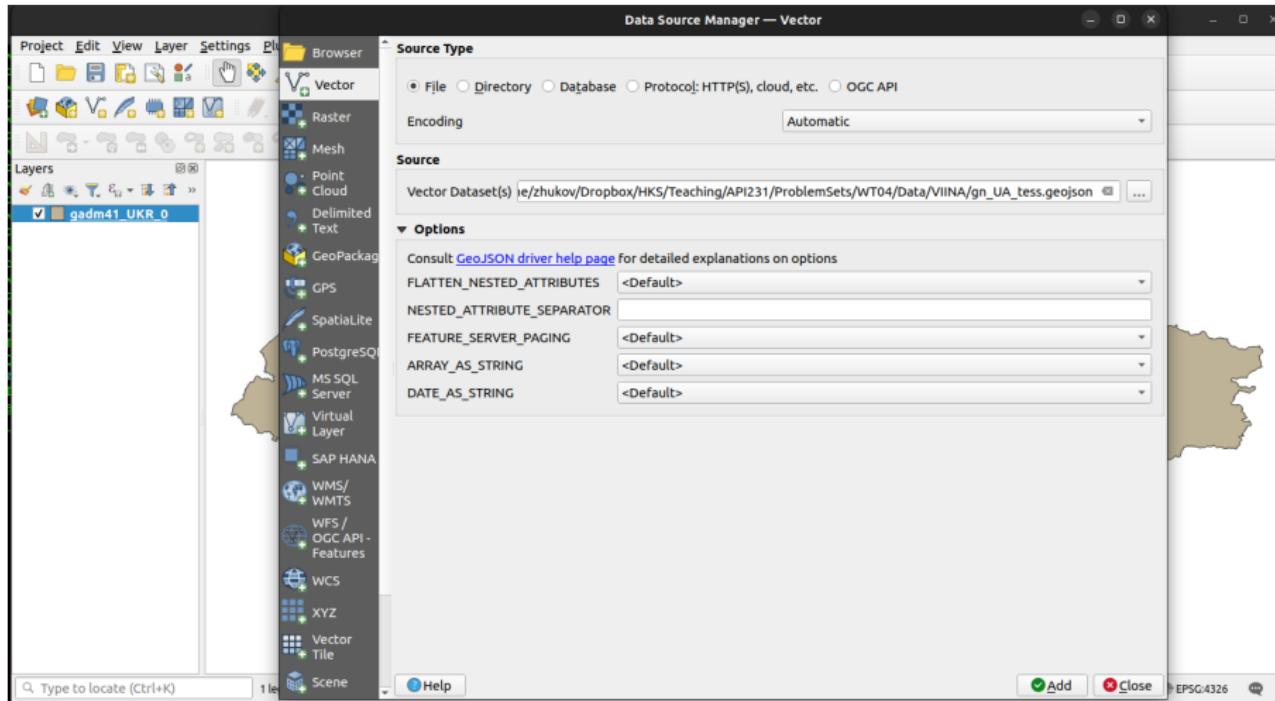
Go to Project → Save As...



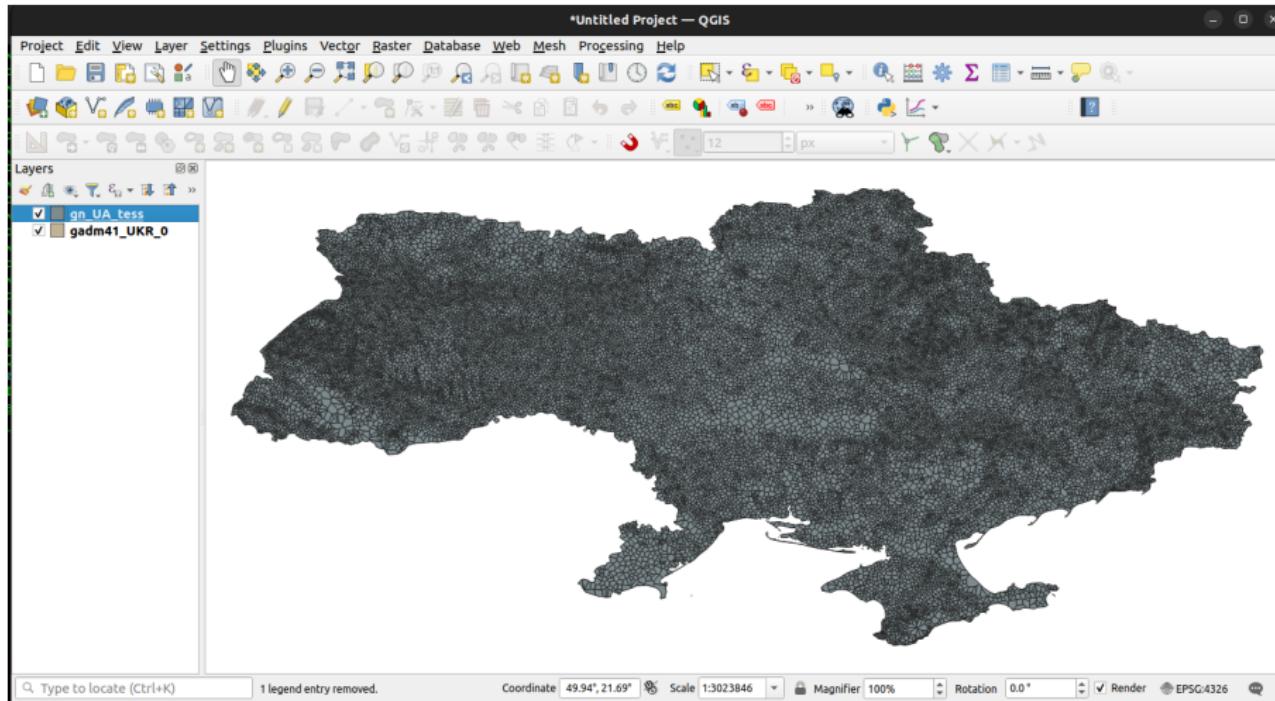
## Vignette 1. Load Ukraine's national borders (Layer → Add Layer → Add Vector Layer). gadm41\_UKR\_0.json file in Data/GADM



Also load the *populated place borders* (Layer → Add Layer → Add Vector Layer). gn\_UA\_tess.geojson file in Data/VIINA



There are 33,141 populated places in Ukraine. This is the level at which territorial control is measured

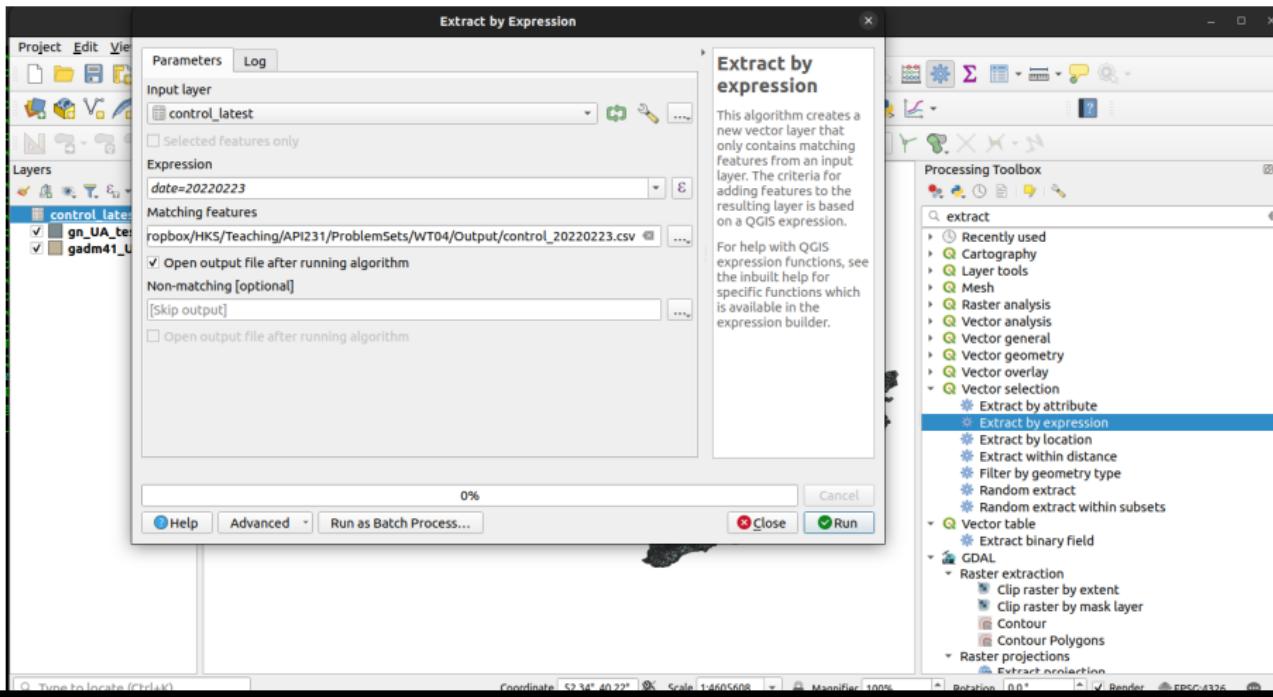


Load the *territorial control data* (`control_latest.csv`) as a delimited text file with no geometries. This is a **HUGE** table (> 25M rows), of which we'll be using only a small part. Take a note of how the date field is formatted (YYYYMMDD)

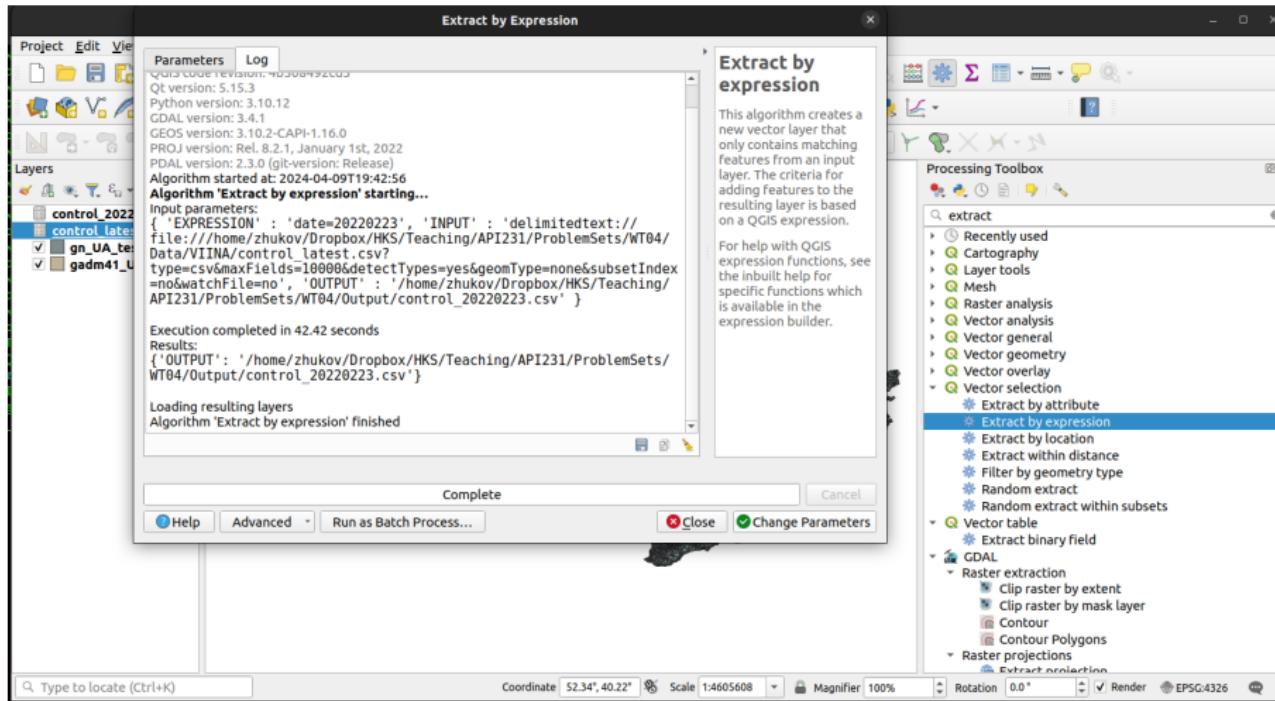
The screenshot shows the QGIS Data Source Manager interface. On the left, the project tree displays layers: 'gn\_UA\_tess' and 'gadm41\_UKR\_0'. The main panel shows the 'Data Source Manager | Delimited Text' configuration for a layer named 'control\_latest'. The 'File Format' section is set to 'CSV (comma separated values)'. Under 'Record and Fields Options', 'No geometry (attribute only table)' is selected. In the 'Layer Settings' section, 'Use spatial index' is checked, while 'Use subset index' and 'Watch file' are unchecked. The 'Sample Data' table shows the following data:

	geonameid	date	status_wiki	status_boost	status_dsm
1	461727	20220223	RU		
2	461727	20220224	RU		
3	461727	20220225	RU		
4	461727	20220226	RU		
5	461727	20220227	RU		RU

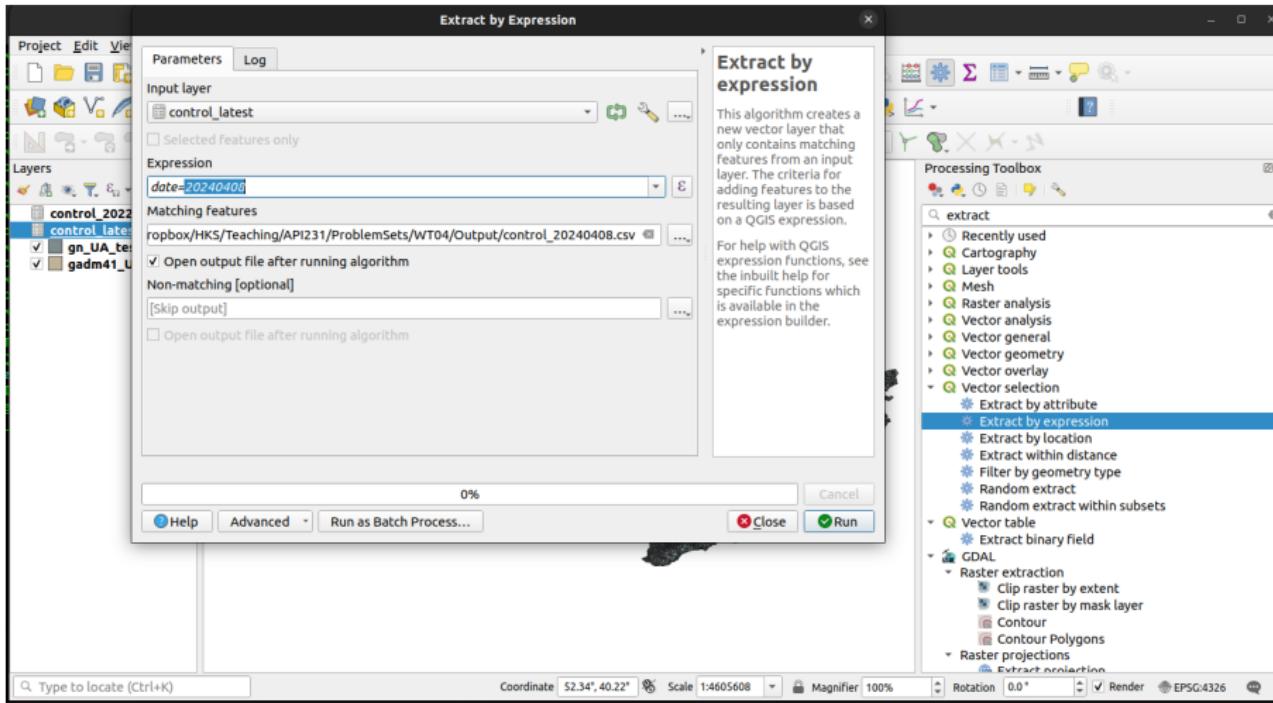
Let's take a subset of this file by date, starting with the day before the full scale invasion (23 Feb 2024). Go to the “Extract by Expression” tool in “Processing Toolbox” → “Vector selection”. Set Input layer: control\_latest and set Expression: date=20220223. Save as control\_20220223.csv



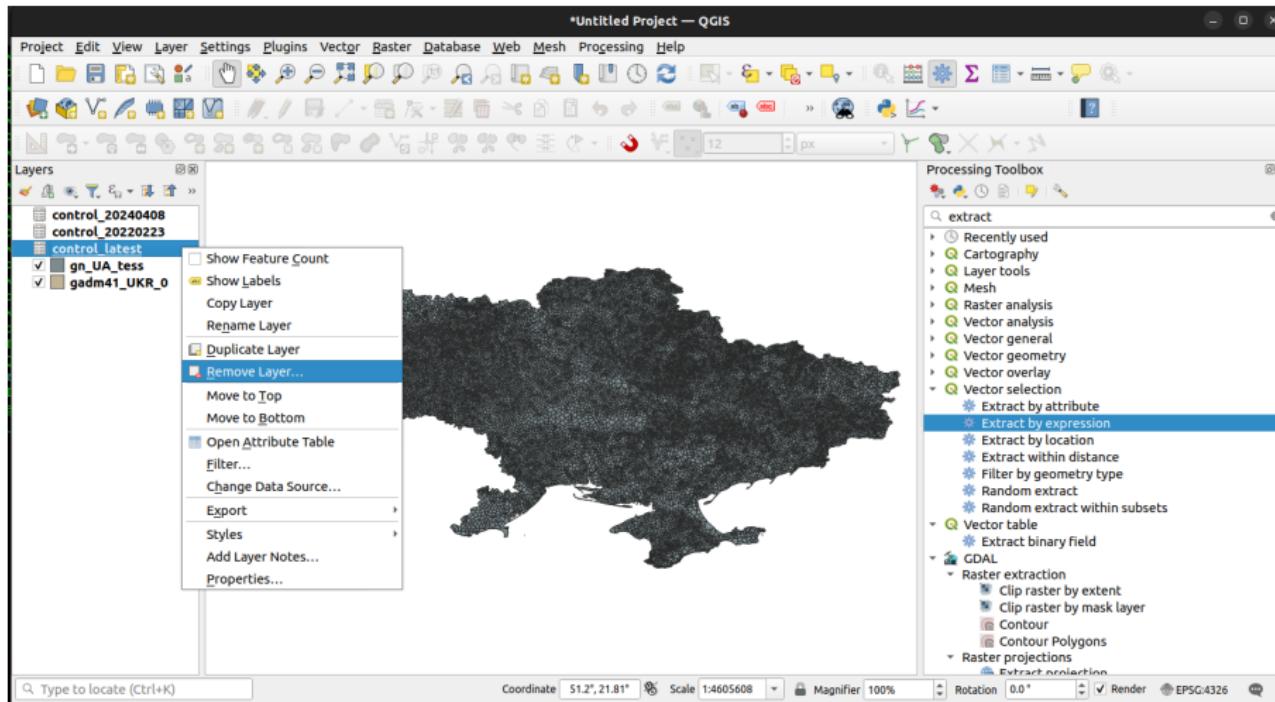
This will take a few minutes to run due to the file size.



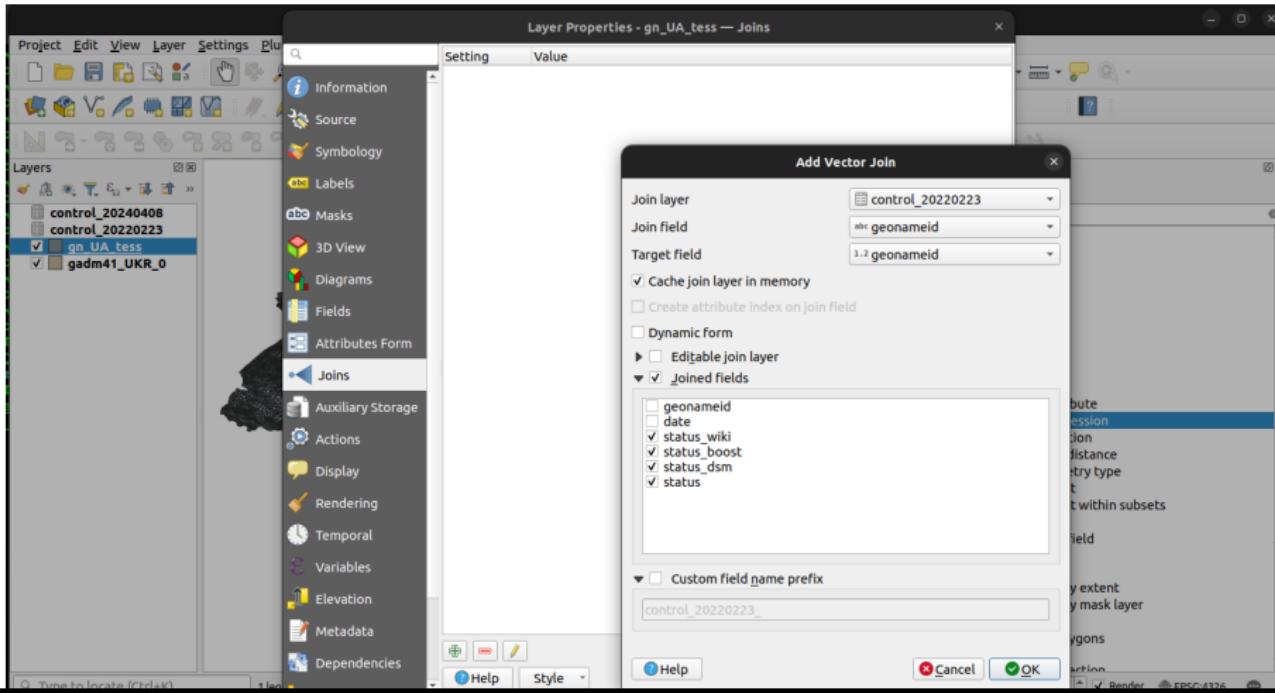
Repeat this process for the most recent date (8 Apr 2024). Set Input layer: control\_latest and Expression: date=20240408. Save as control\_20240408.csv



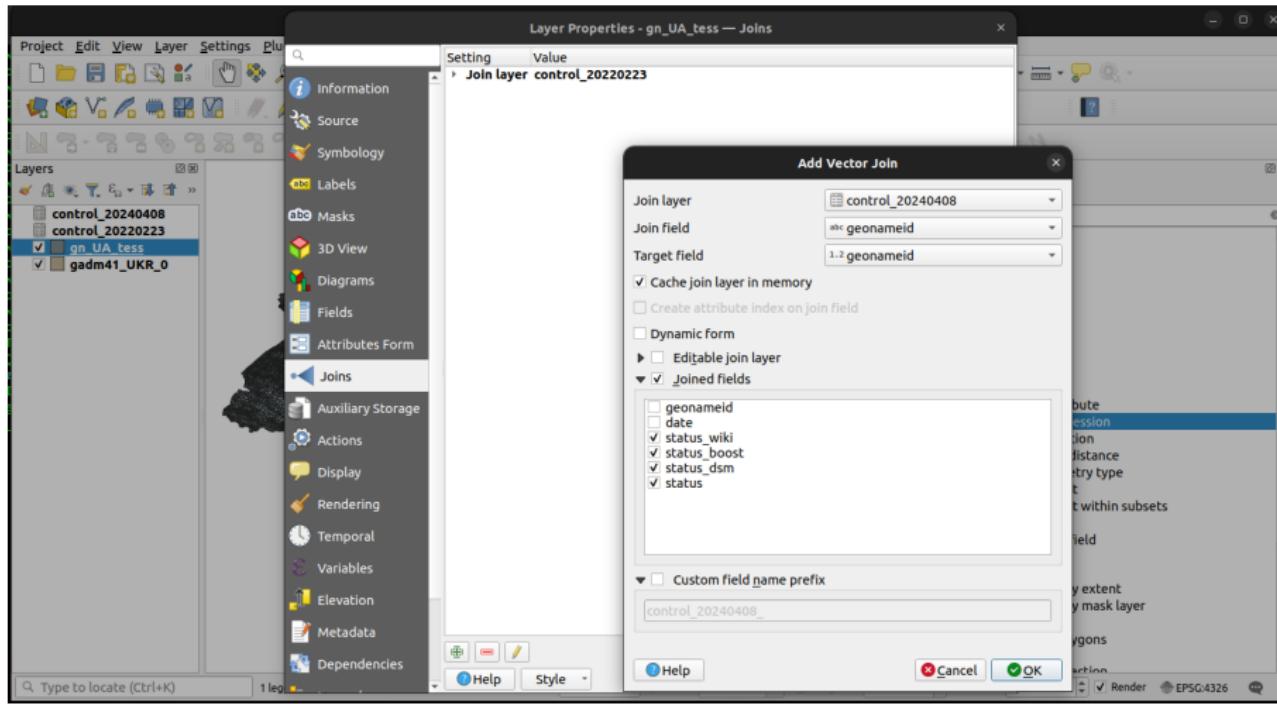
Once both subset tables are loaded, you can remove the original `control_latest` from memory



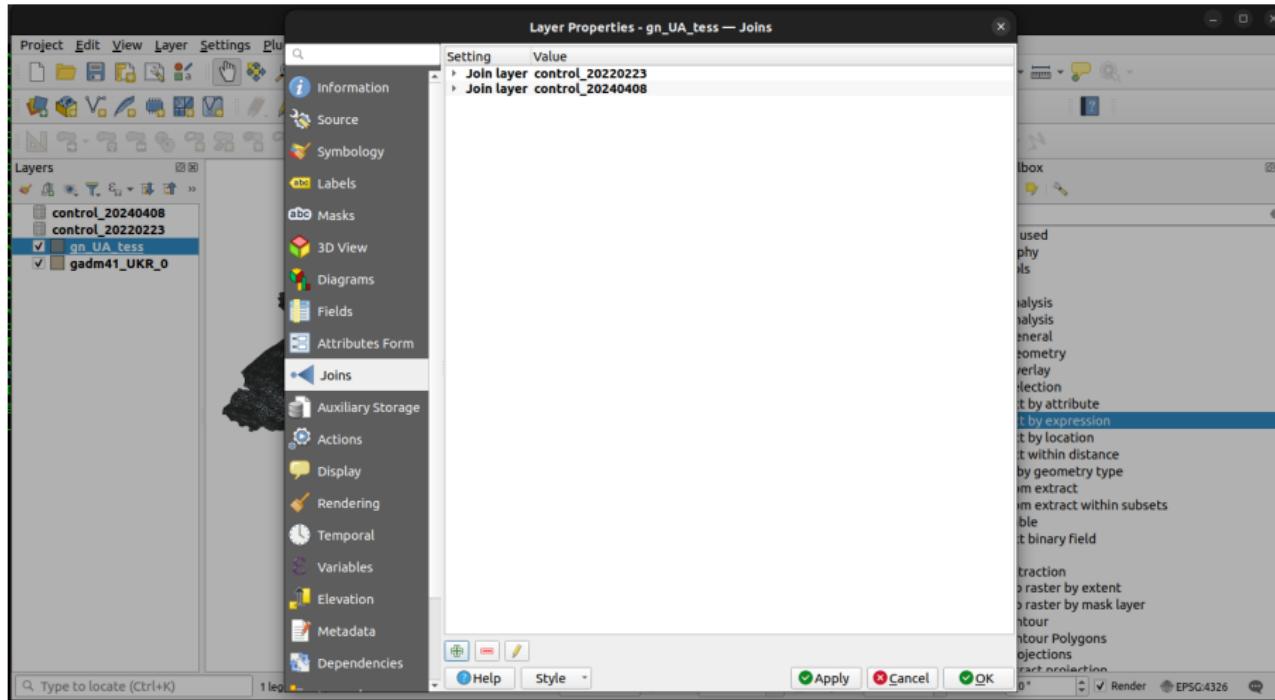
Let's now join the two subset tables to the populated place geometries. Go to the "Joins" tab in layer "Properties" for gn\_UA\_tess, and add a new join with Join layer = control\_20220223 and geonameid as the Join field and Target field. Select the four status\* fields as Joined fields



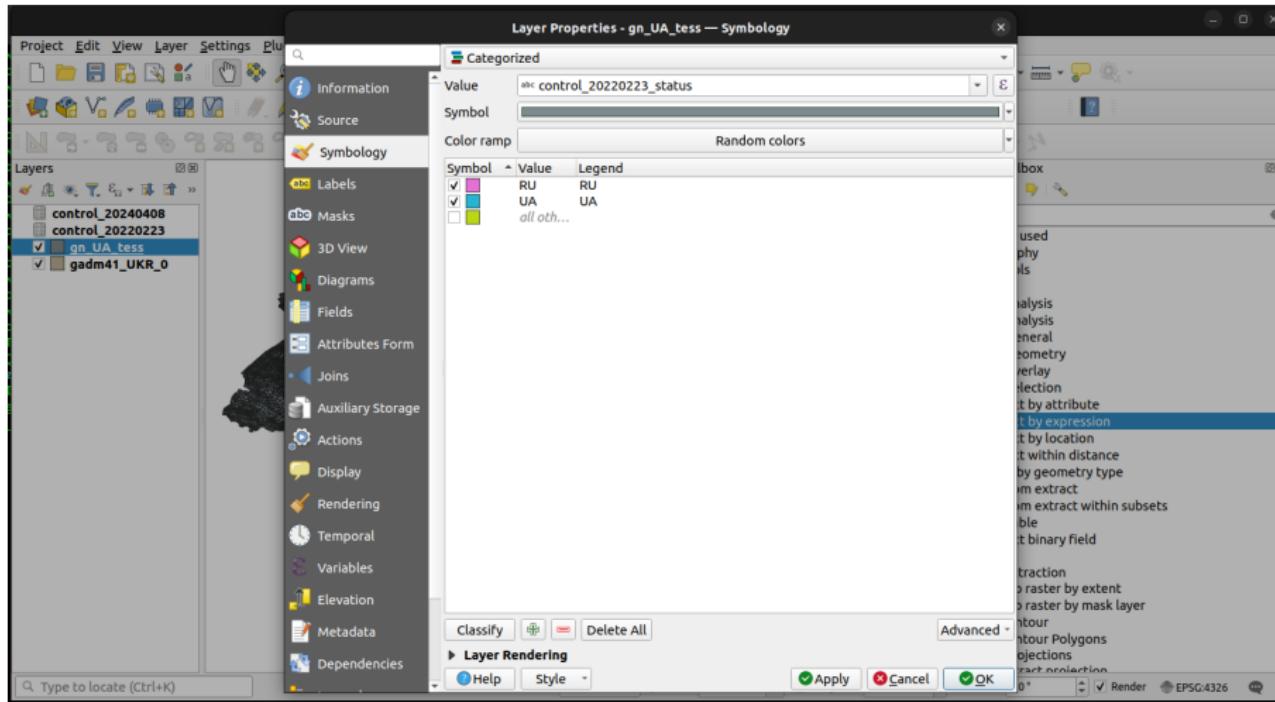
## Add a second join with control\_20240408 as Join layer



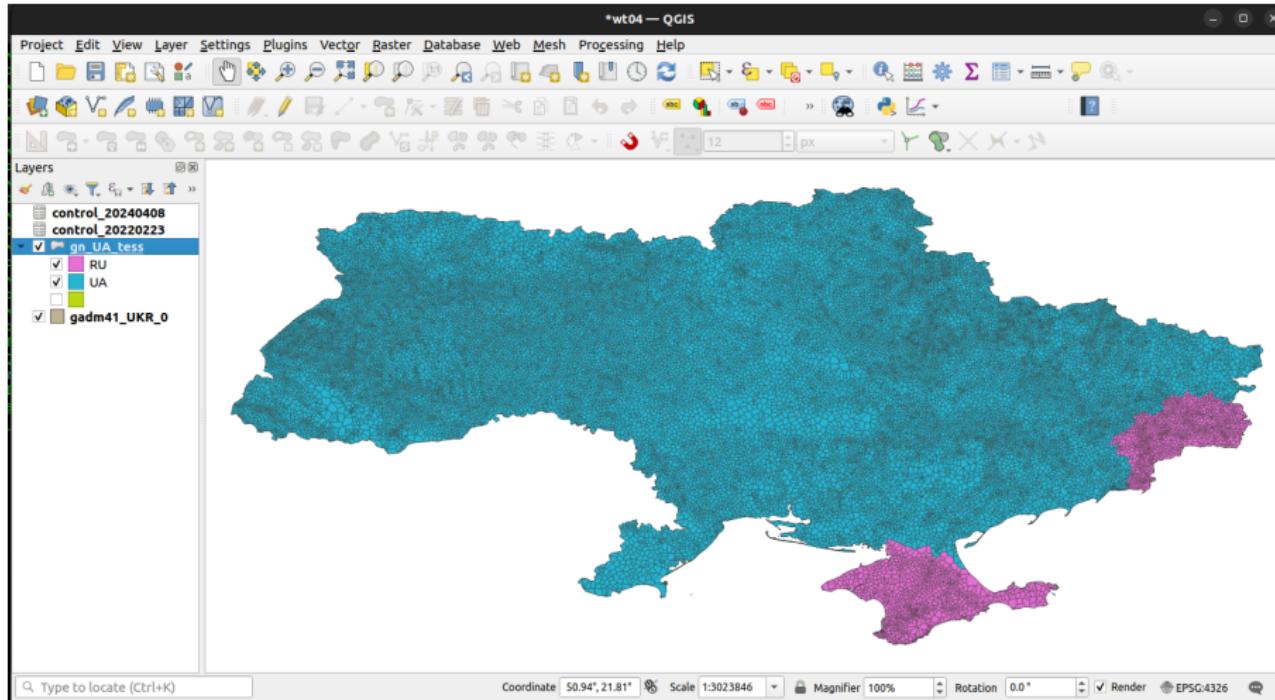
The two join layers should now appear in the “Joins” tab



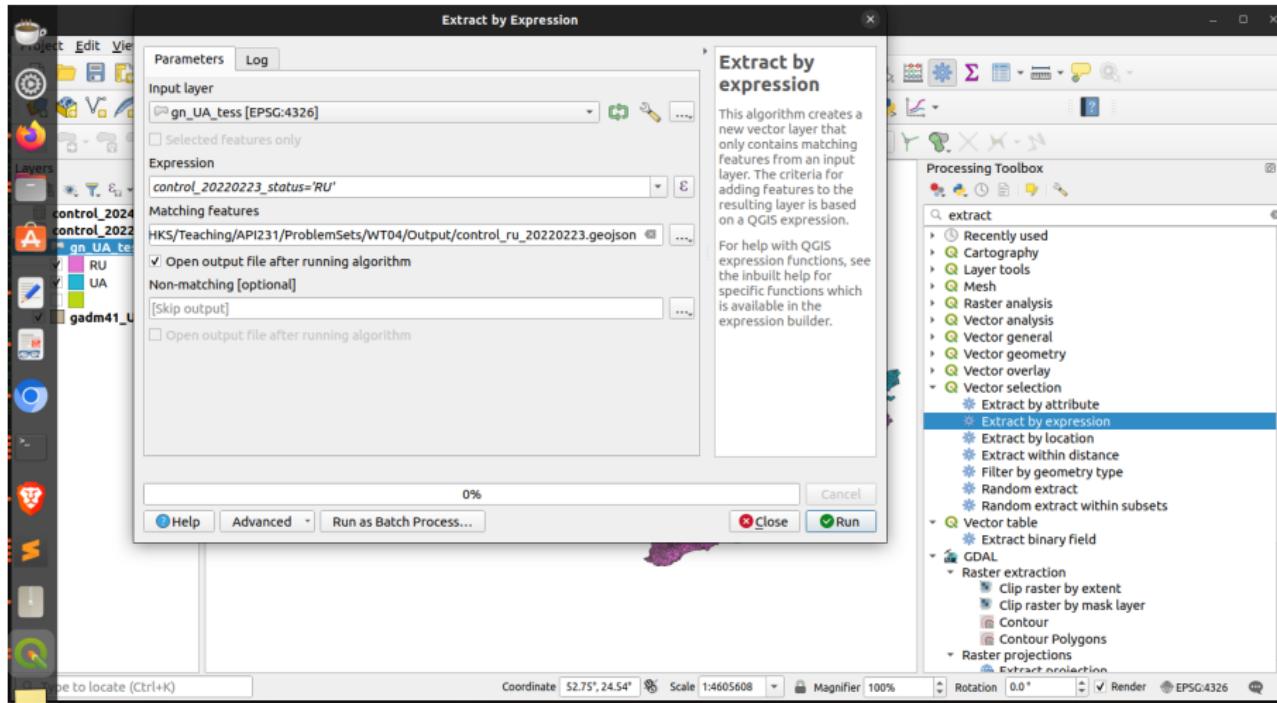
Let's visualize the control status, to make sure everything is right. Change to symbology to Categorized with Value = control\_20220223\_status and click Classify and OK



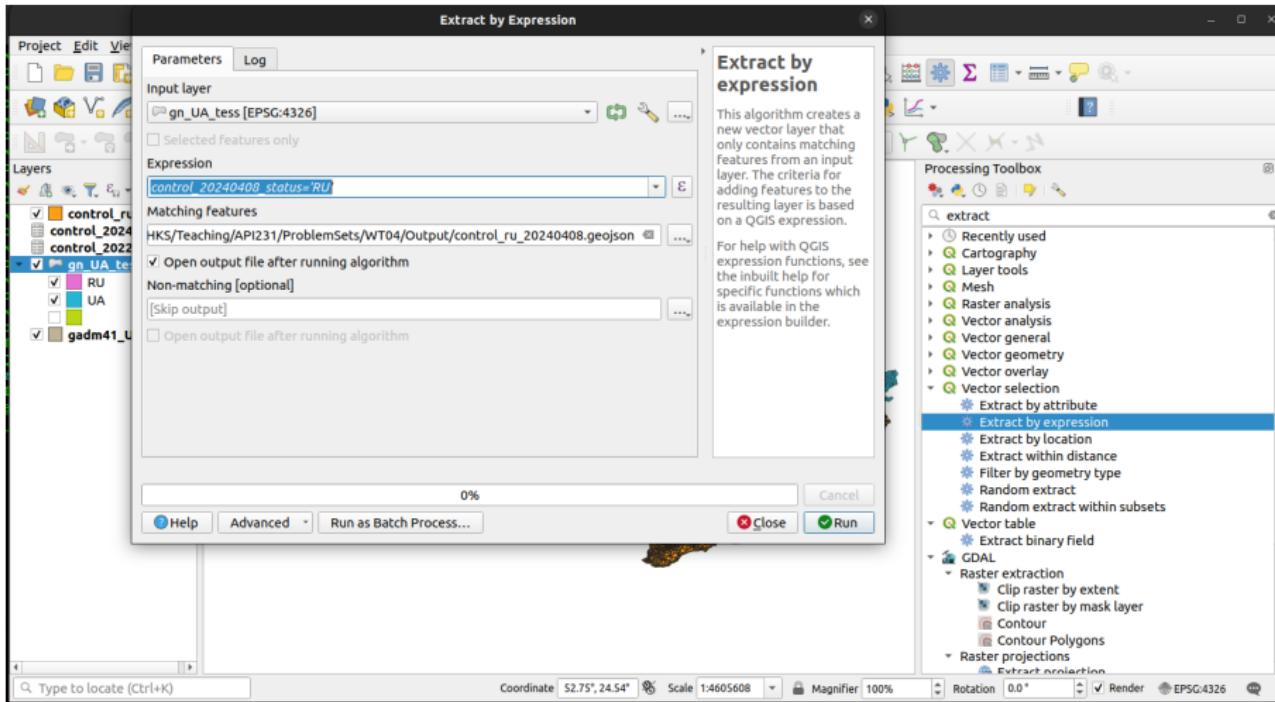
This looks about right. Now we just need to extract these occupied areas and calculate how much of Ukraine's territory they represent



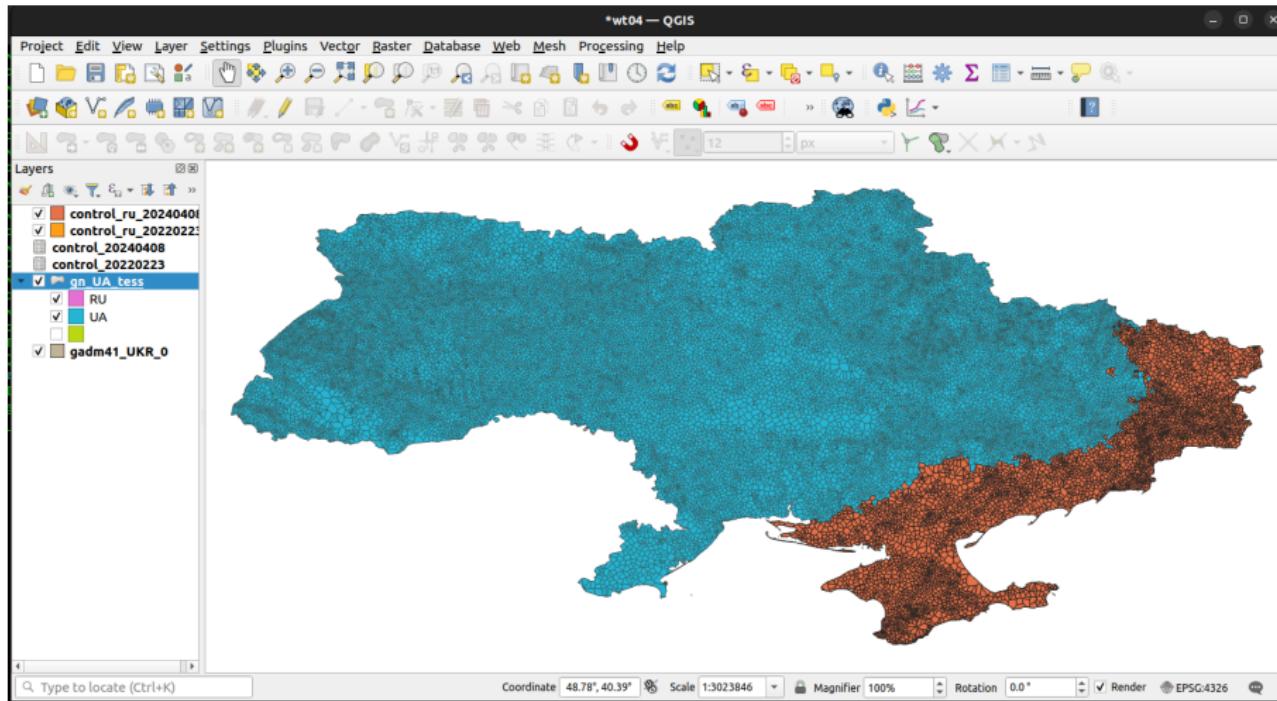
Go back to the “Extract by Expression” tool, with Input layer: gn\_UA\_tess and Expression: control\_20220223\_status='RU'. Save the output as control\_ru\_20220223.geojson



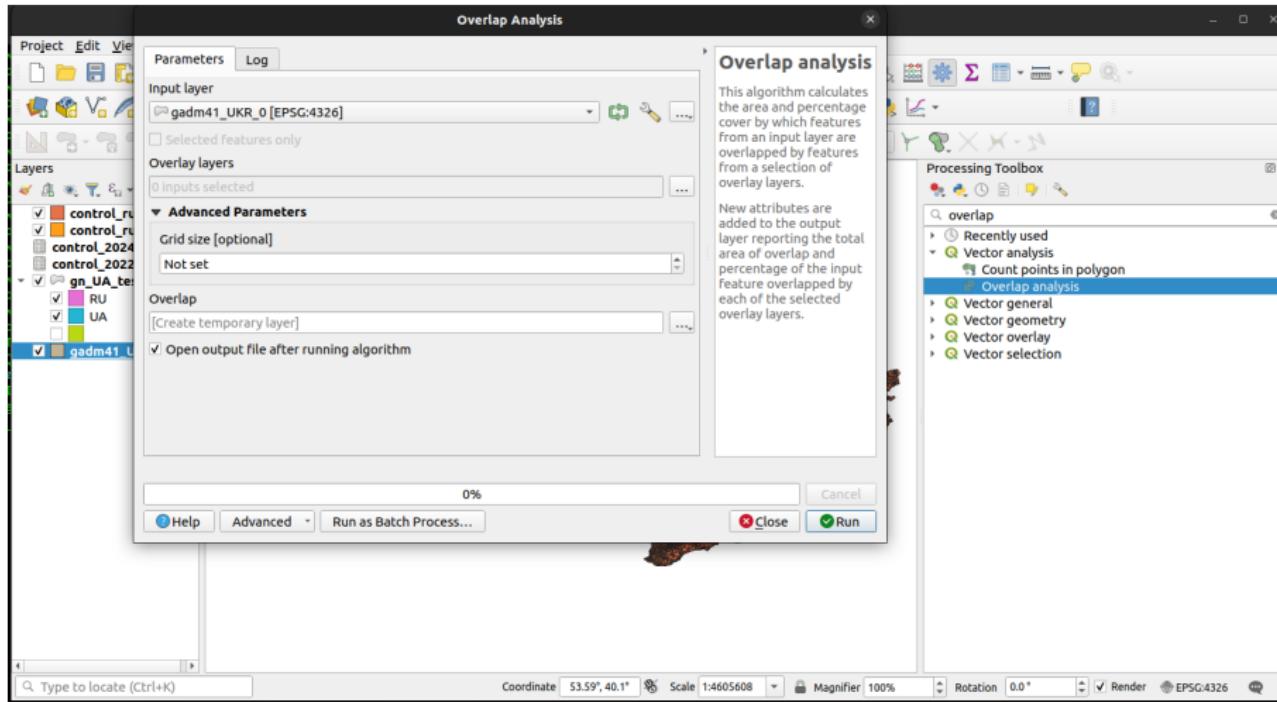
Repeat for the latest date, with Expression: `control_20240408_status='RU'`.  
Save the output as `control_ru_20240408.geojson`



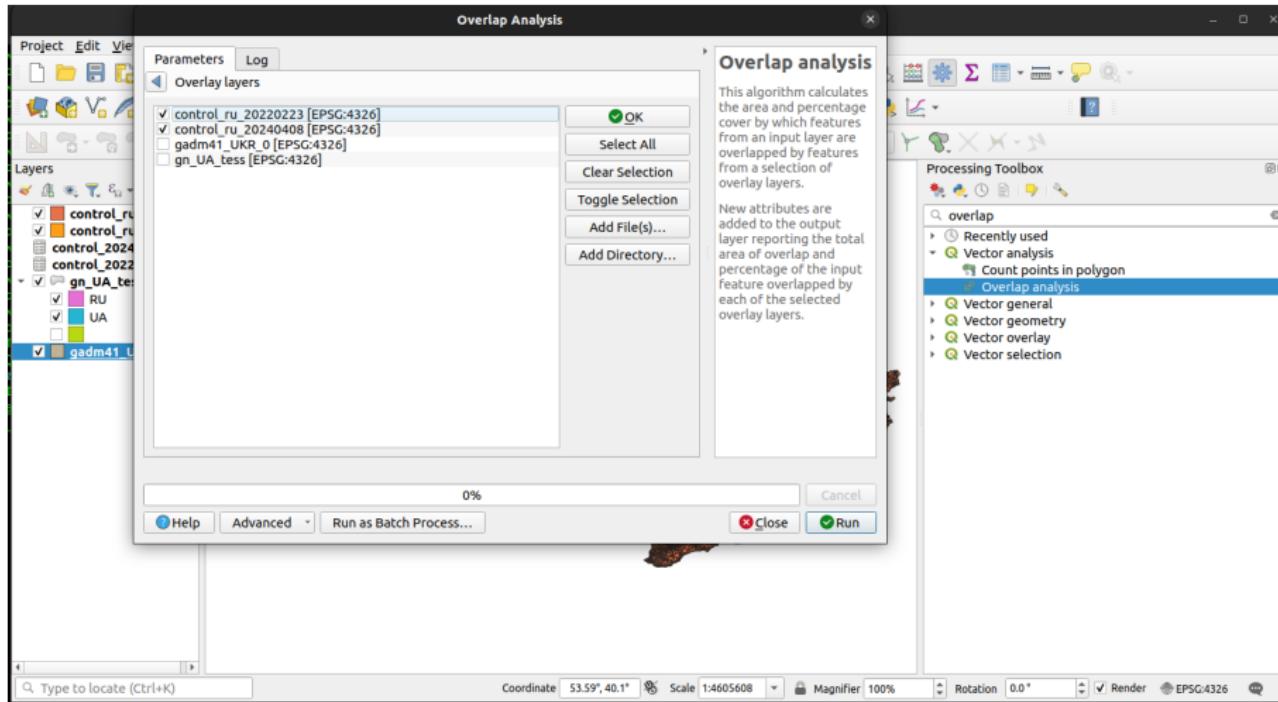
The extracted areas should appear in the project window.



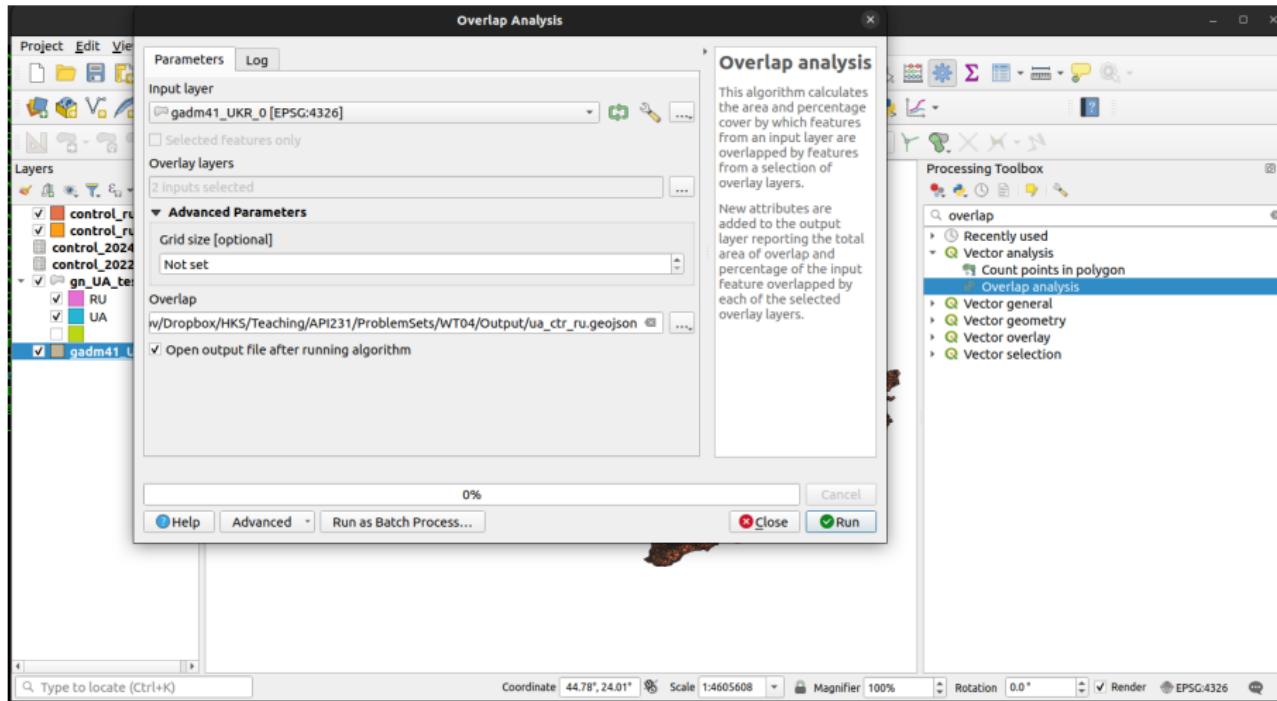
Go to the “Overlap Analysis” tool in “Processing Toolbox” → “Vector analysis”. Set Input layer: gadm41\_UKR\_0 (country-level borders). Click on the ... button next to Overlay layers



Select control\_ru\_20220223 and control\_ru\_20240408 as Overlay layers.  
Click OK



Save the output as ua\_ctr\_ru.geojson



Open the attribute table for the newly-created ua\_ctr\_ru layer. The \*\_pc variables indicate that Russia occupied 7.23% of Ukraine's territory on 23 Feb 2024 and 18.45% on 8 Apr 2024

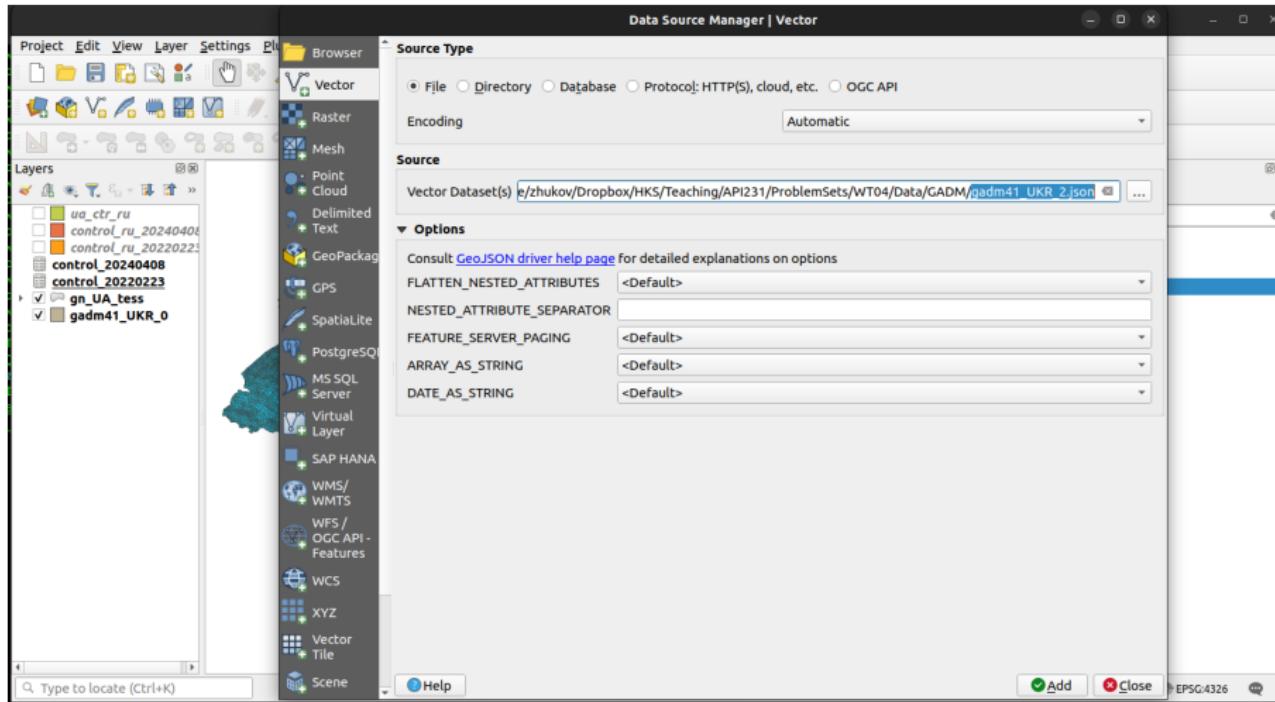
The screenshot shows the QGIS interface with the title bar "\*wt04 — QGIS". The menu bar includes Project, Edit, View, Layer, Settings, Plugins, Vector, Raster, Database, Web, Mesh, Processing, and Help. The toolbar contains various icons for file operations, selection, and editing. The Layers panel shows two layers: "ua\_ctr\_ru" (selected) and "control\_ru\_20240401". The Processing Toolbox is open with a search bar containing "overlap". The attribute table for "ua\_ctr\_ru" is displayed, showing one feature for Ukraine with the following data:

GID_0	COUNTRY	l_ru_2022022	control_ru_20220223_pc	l_ru_20240401	control_ru_20240408_pc	
1	UKR	Ukraine	435025531...	7.233974208118402	110937125...	18.447567946969357

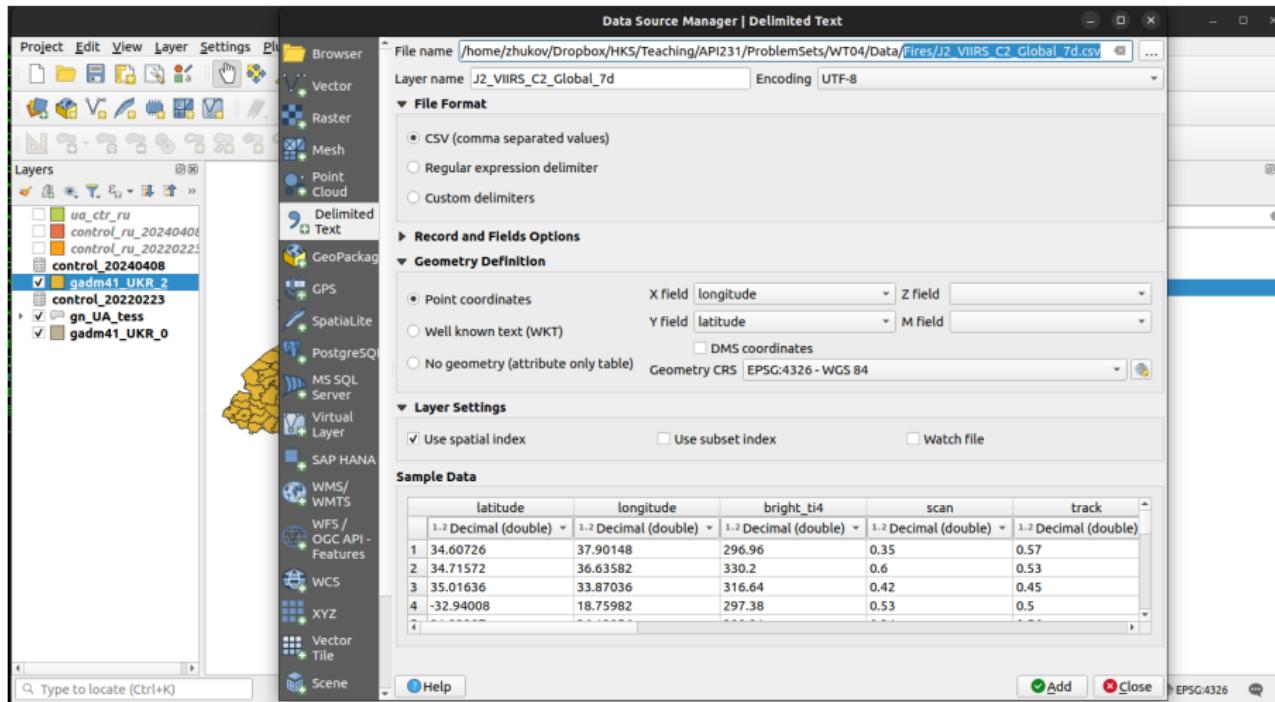
The bottom status bar shows coordinates 50.58°, 21.74°, scale 1:4605608, magnifier 100%, rotation 0.0°, render, and EPSG:4326.

## Comparing media event reports to remote sensing data on fires

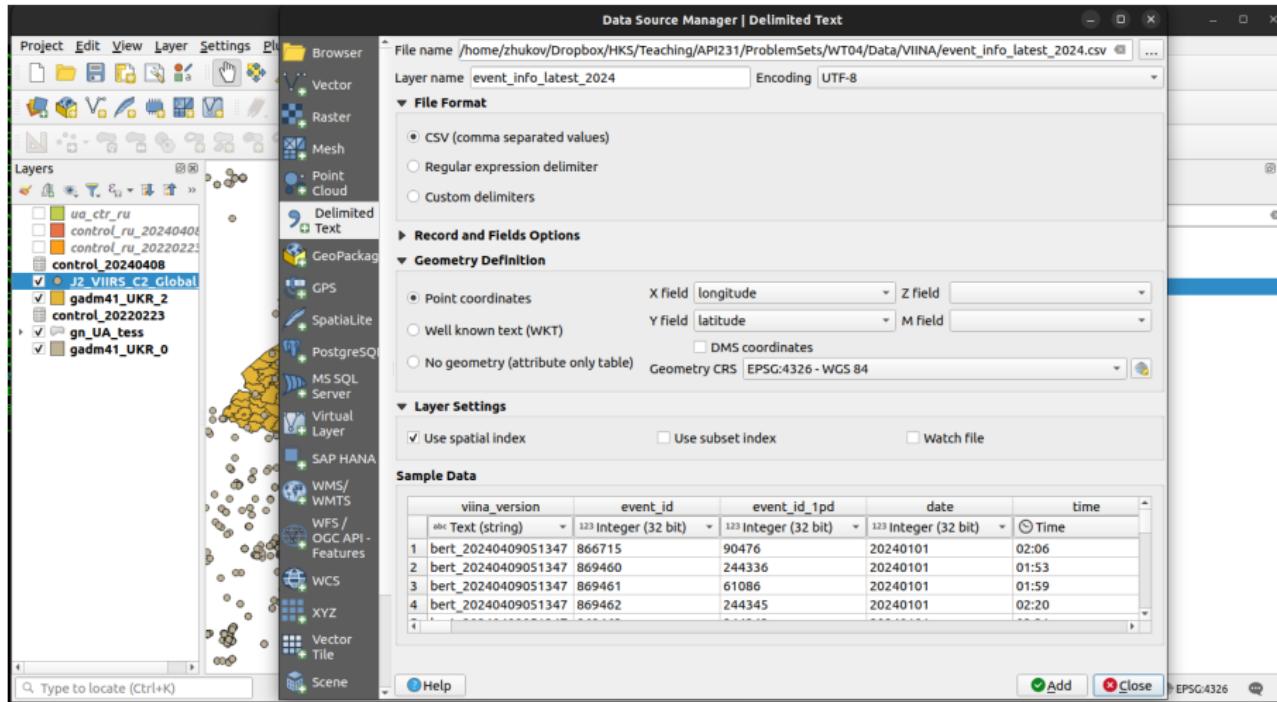
**Vignette 2!** Let's load *Ukrainian district borders: gadm41\_UKR\_2.json* from Data/GADM. These will be our spatial units of analysis



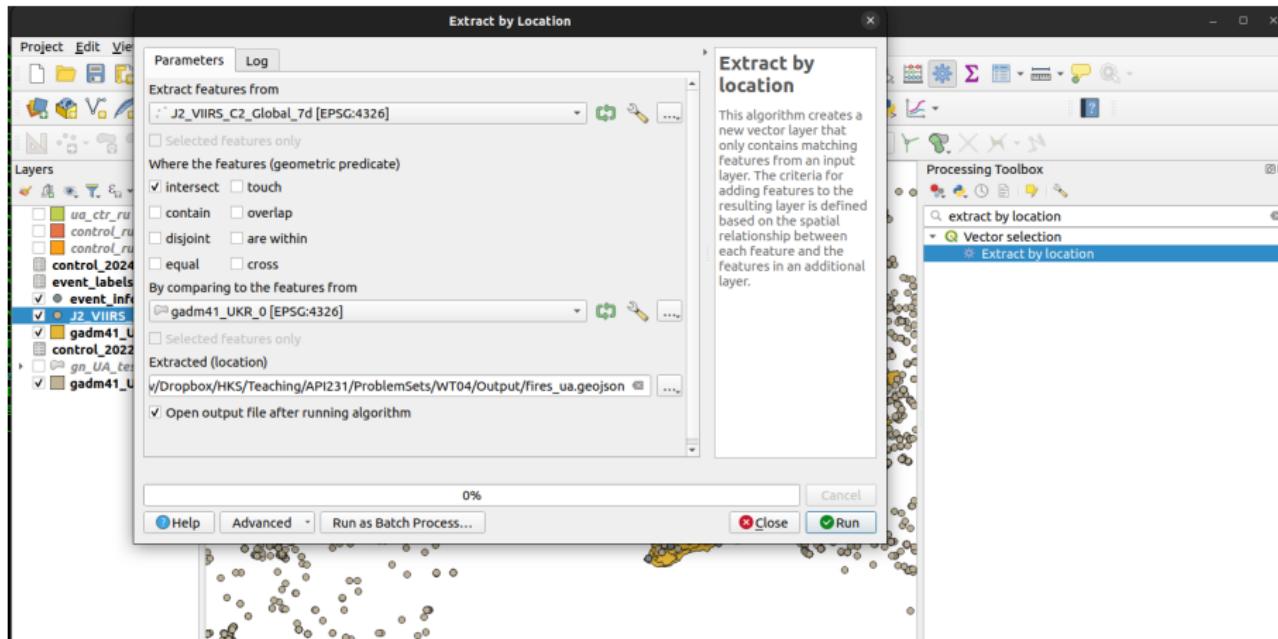
Load *active fires data* as delimited text: J2\_VIIRS\_C2\_Global\_7d.csv. Make sure the X and Y fields are properly specified, check box next to  Use spatial index



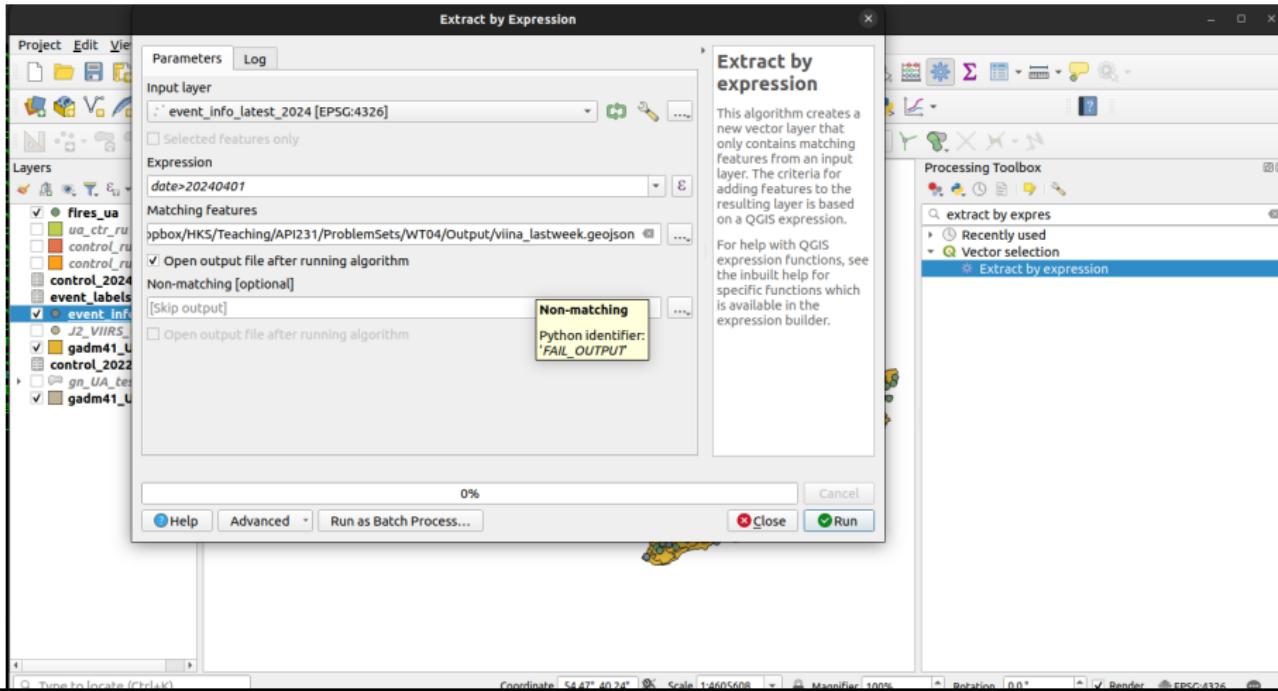
Load *media event reports* as delimited text: `event_info_latest_2024.csv`. Here, too, specify the X and Y fields and check the box next to  Use spatial index



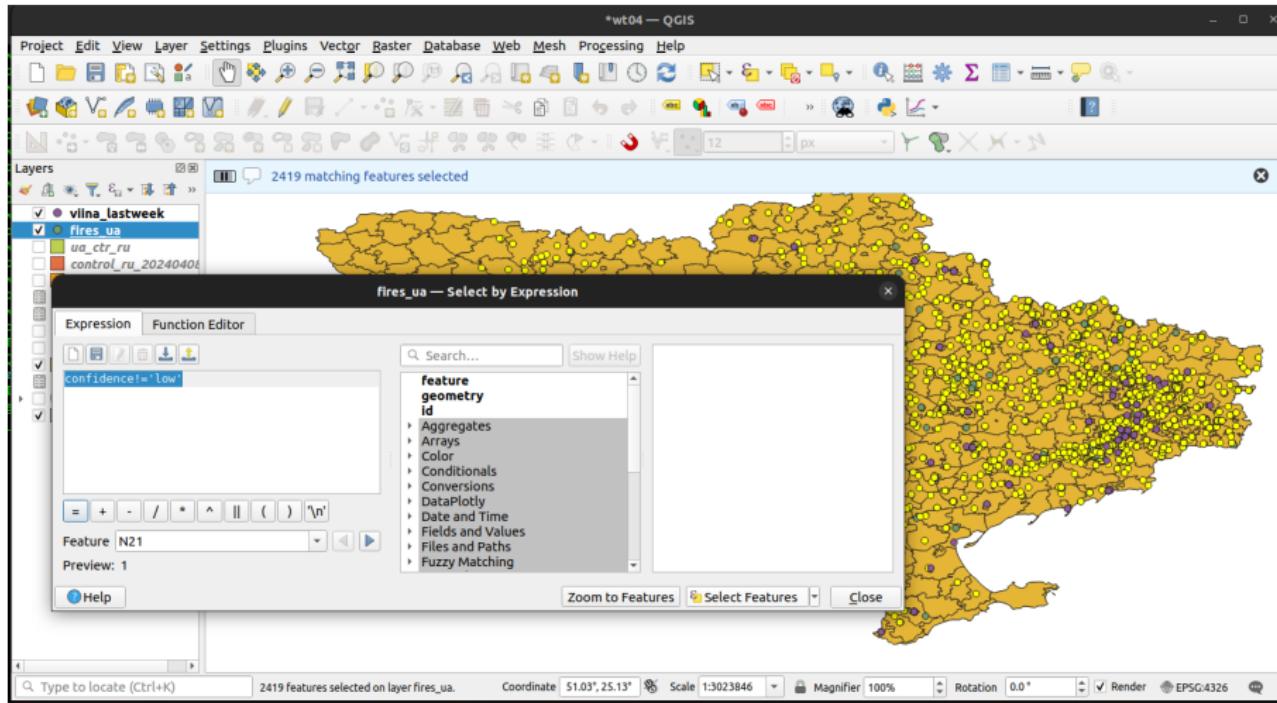
Let's extract just the fires inside of Ukraine. Go to the "Extract by Location" tool in "Processing Toolbox" → "Vector selection". Set Extract features from: J2\_VIIRS\_C2\_Global\_7d, check  Intersect, By comparing to the features from: gadm41\_UKR\_0. Save the extracted features as fires\_ua.geojson



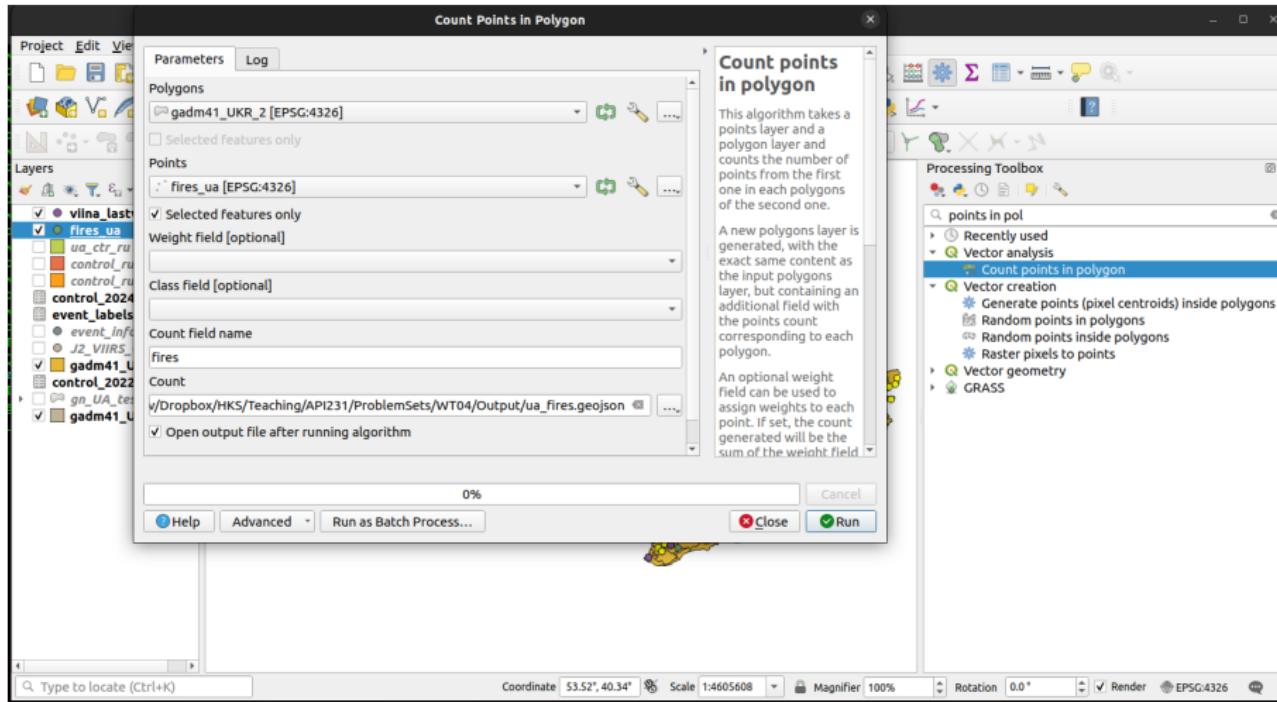
Let's also extract the media reports for the same time period as the fires (last week). Go to the "Extract by Expression" tool and set Input layer: event\_info\_latest\_2024, and Expression: date>20240401. Save the extracted features as viina\_lastweek.geojson



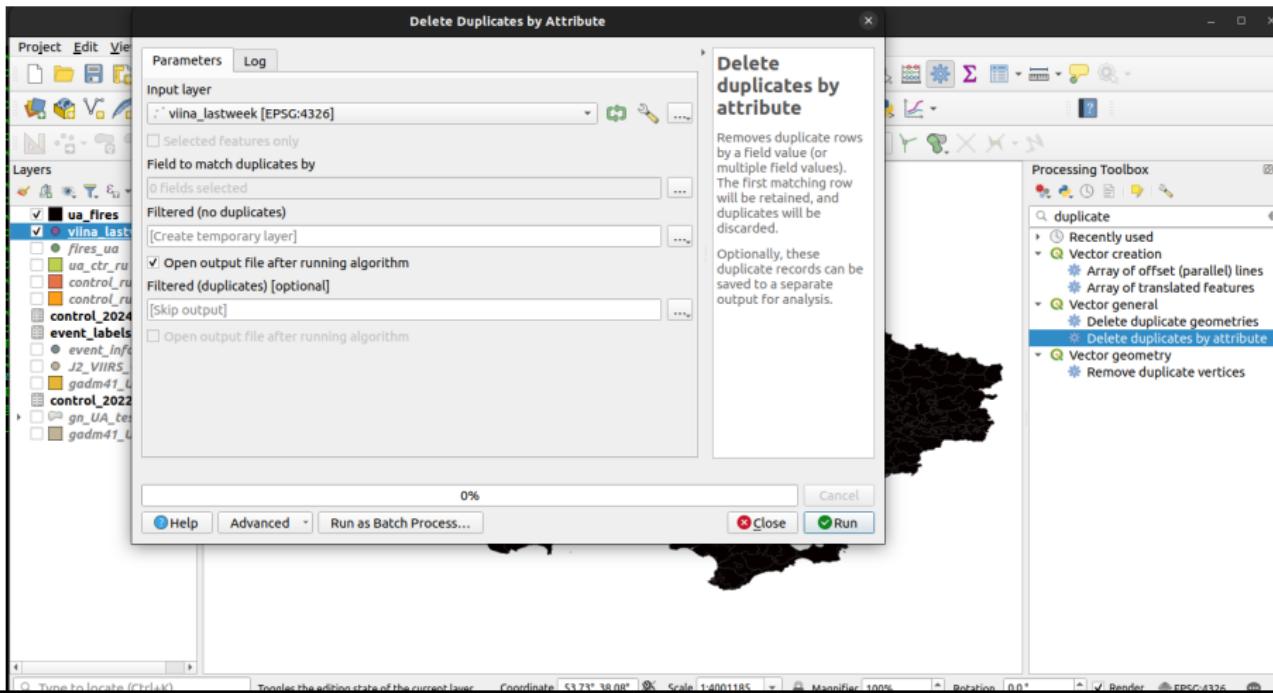
Before counting the fires, let's remove the “low confidence” fire anomalies. With `fires_ua` as the active layer, go to Select by Expression, with Expression: `confidence != 'low'` (`!=` means “ $\neq$ ”). This should select about 2419 features



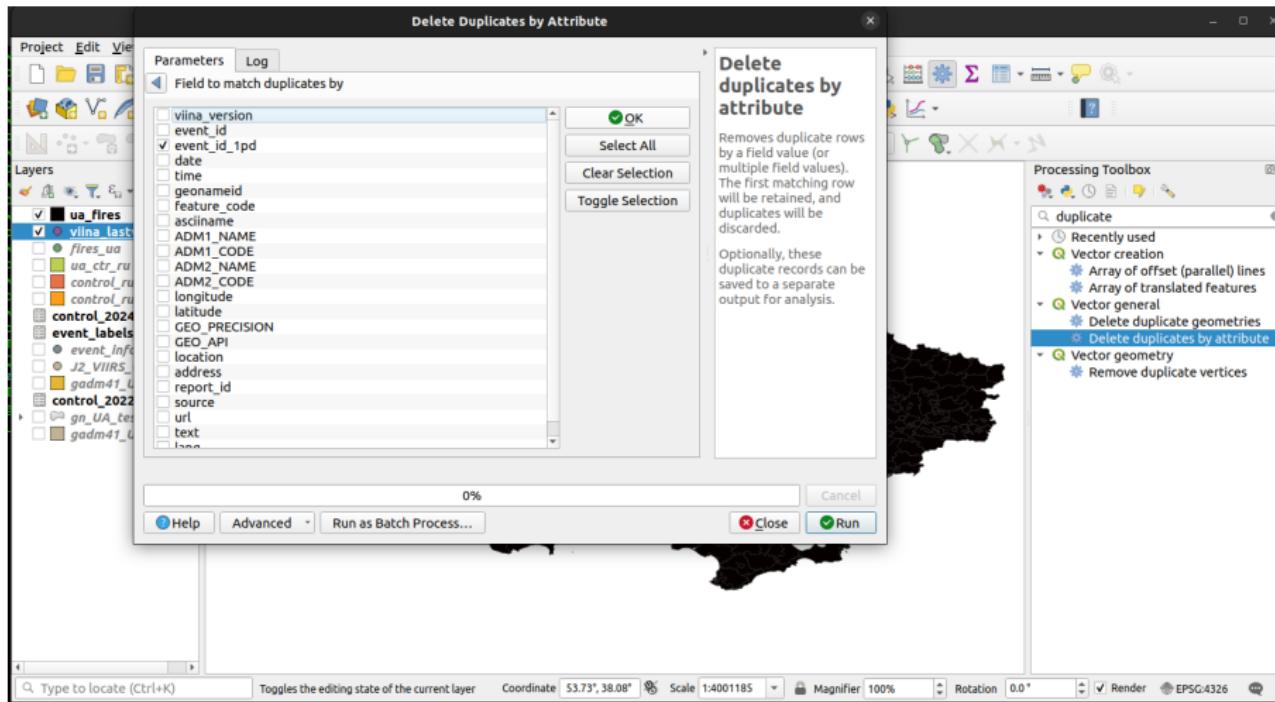
Open the “Count Points in Polygon” tool, set Polygons: gadm41\_UKR\_2, Points: fires\_ua, check the box  Selected features only, name the count field fires and save the output as ua\_fires.geojson



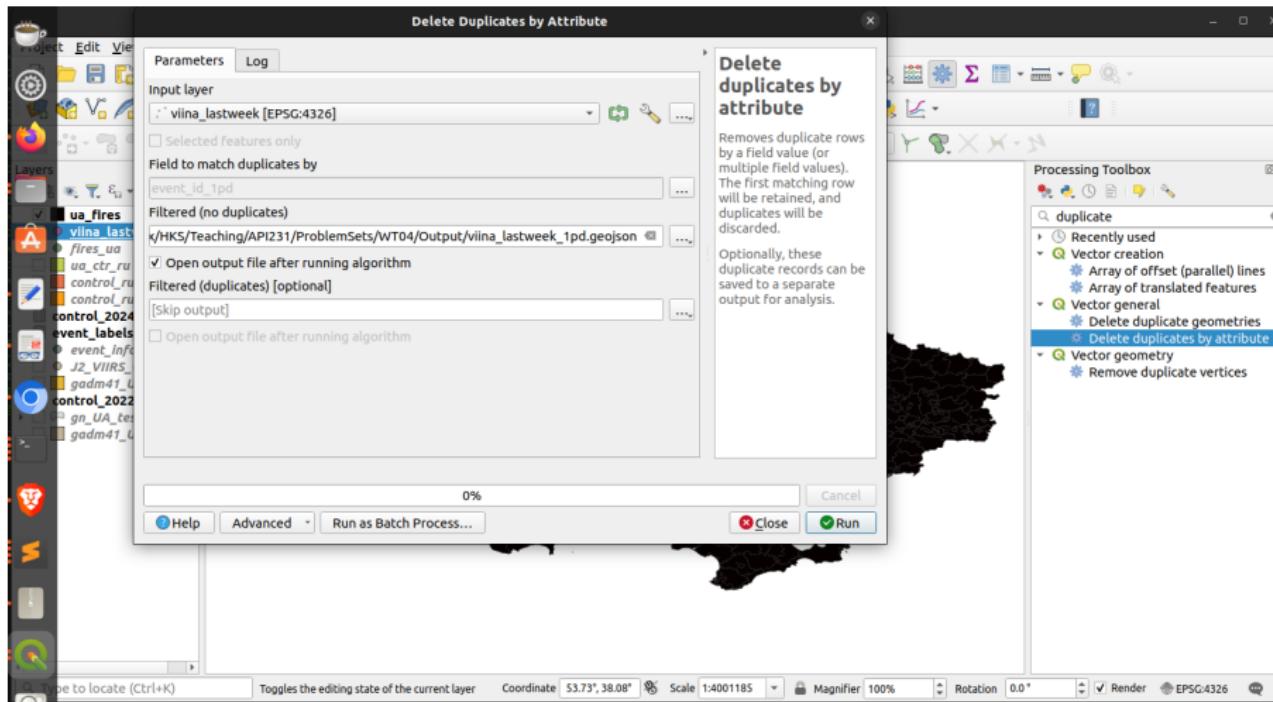
Before counting the media reports, let's remove the obvious duplicates. Open the “Delete Duplicates by Attribute” tool in “Processing Toolbox” → “Vector general”. Set Input layer: viina\_lastweek. Click the ... button next to Field to match duplicates by



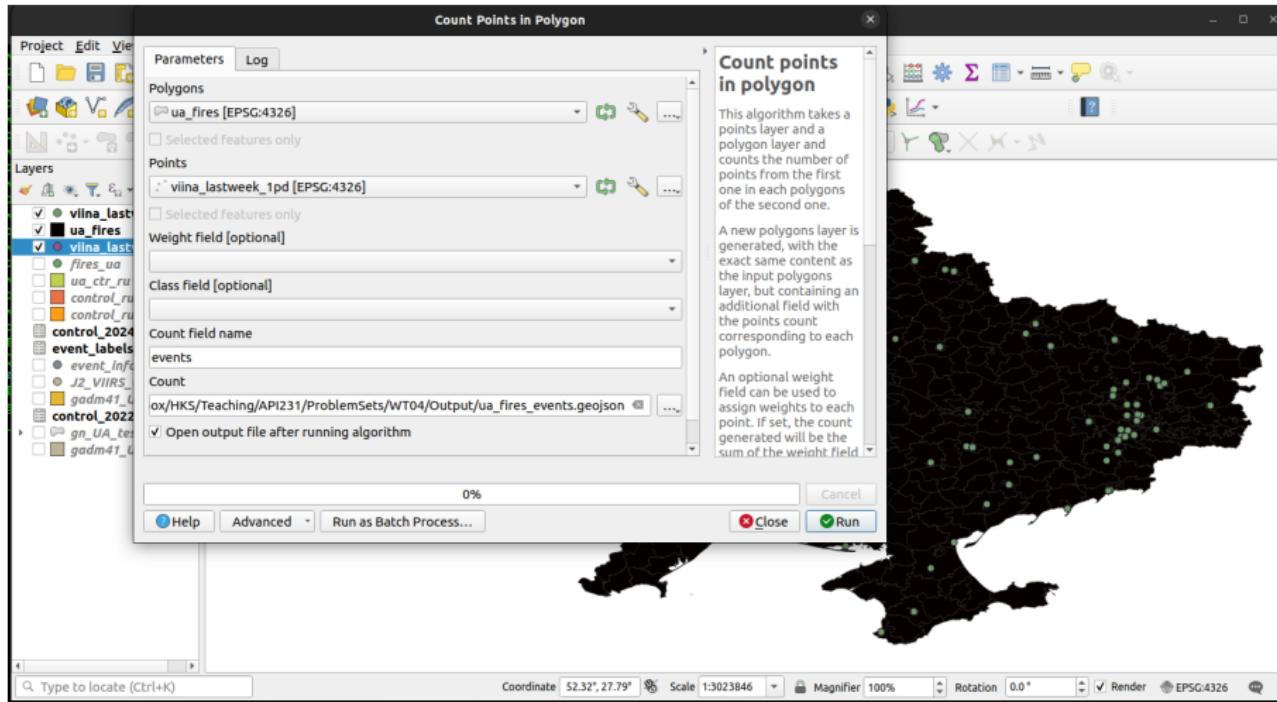
Select event\_id\_1pd as the field, click OK



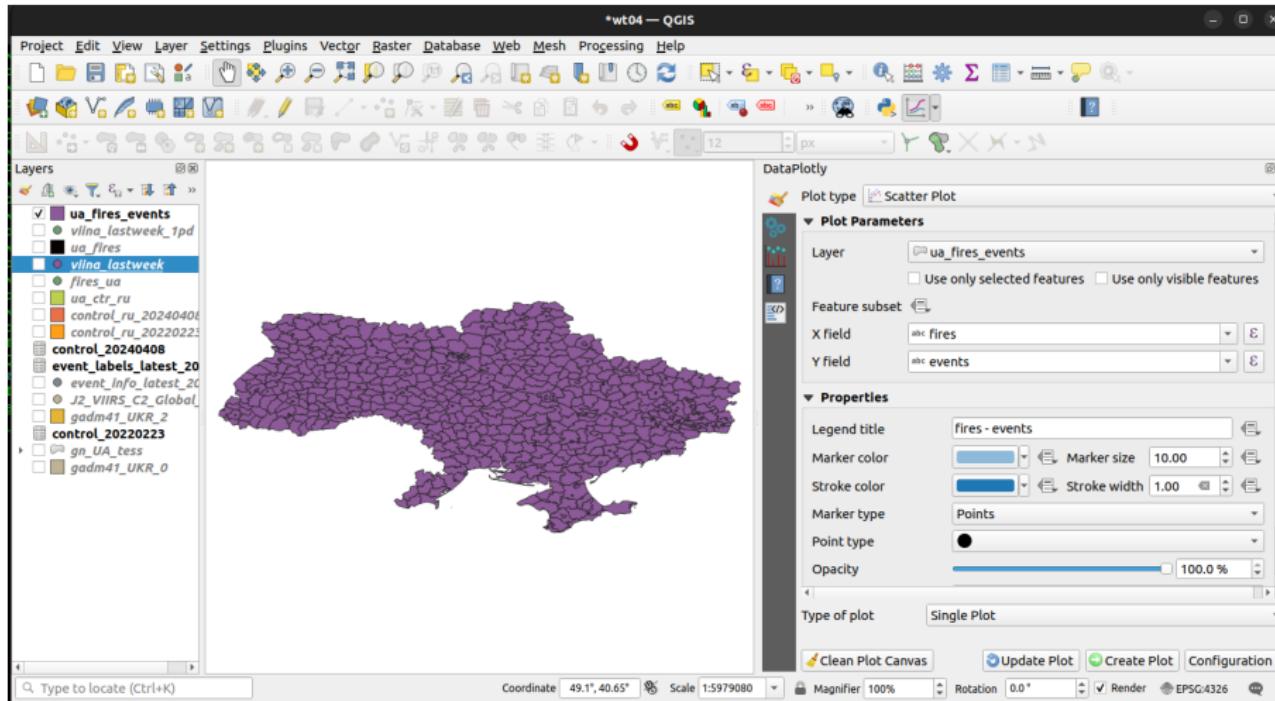
Save the deduplicated output as viina\_lastweek\_1pd.geojson



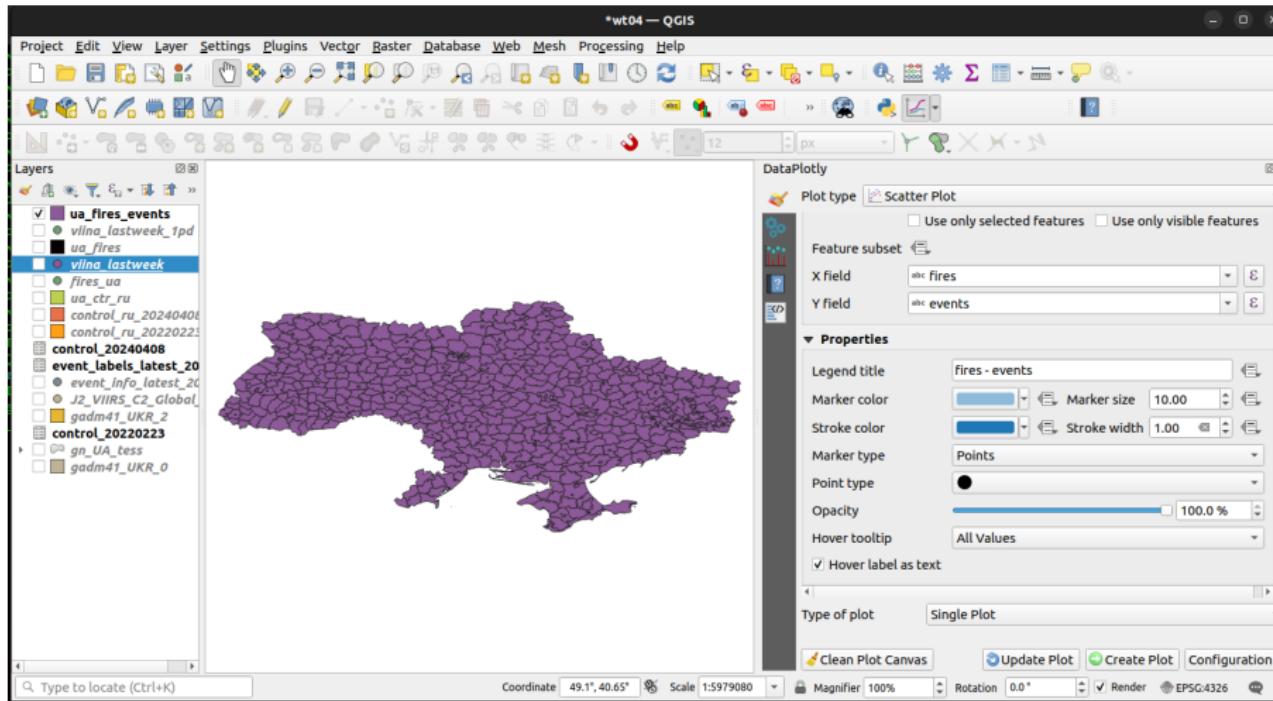
Go back to “Count points in polygon”. Set Polygons: ua\_fires, Points: viina\_lastweek\_1pd. Name the count field events and save as ua\_fires\_events.geojson



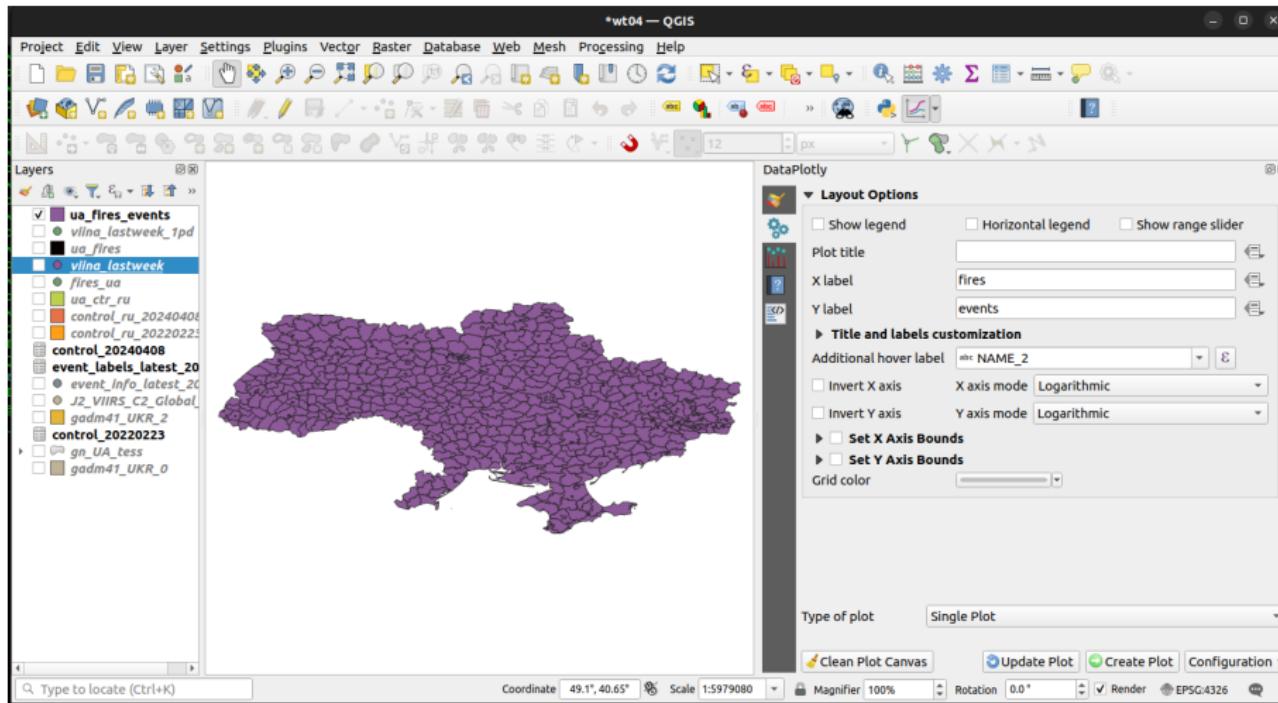
Open the Plotly tool. Set Plot type = Scatter Plot, Layer = ua\_fires\_events, X field = fires, Y field = events



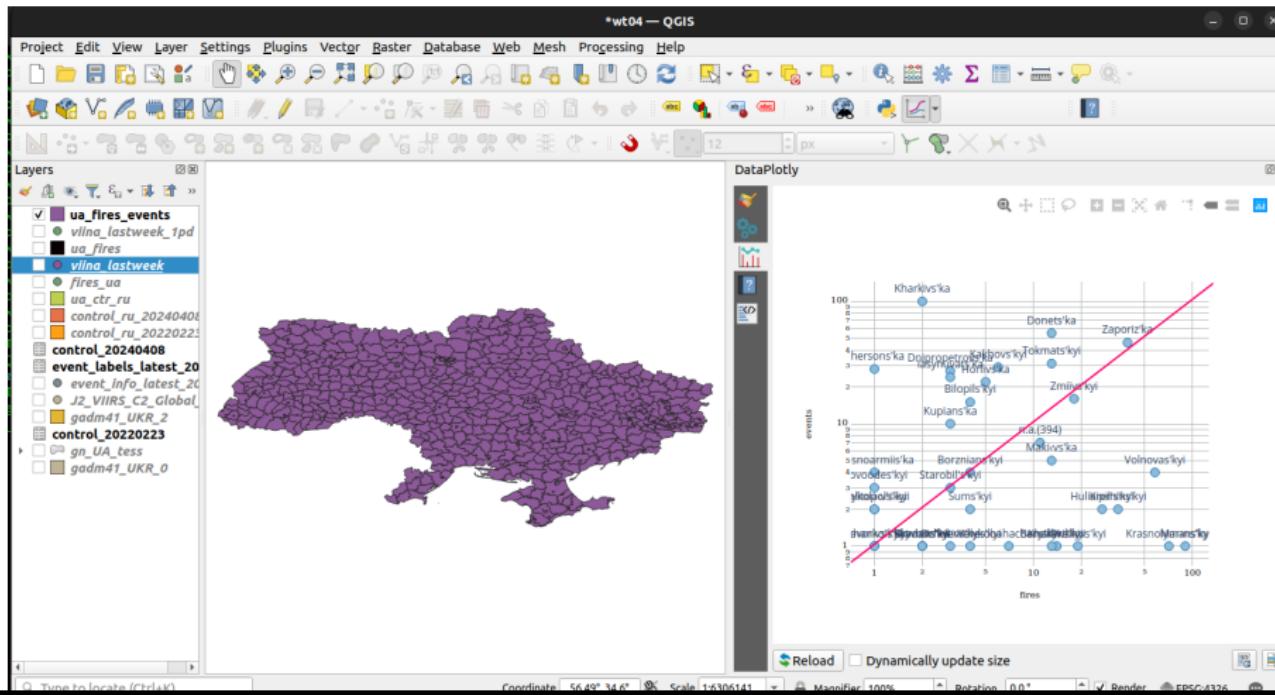
Check the box next to ✓ Hover label as text



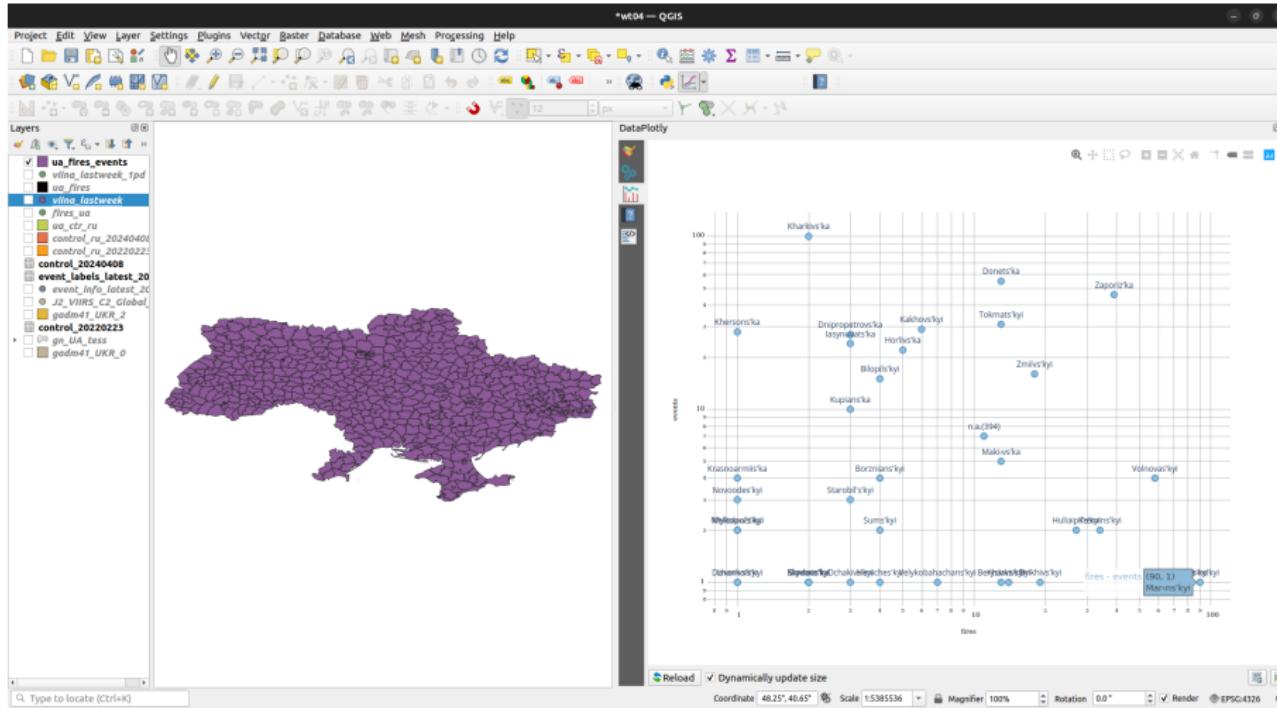
In “Layout Options”, uncheck the box next to  Show legend. Change the X and Y axis mode to Logarithmic



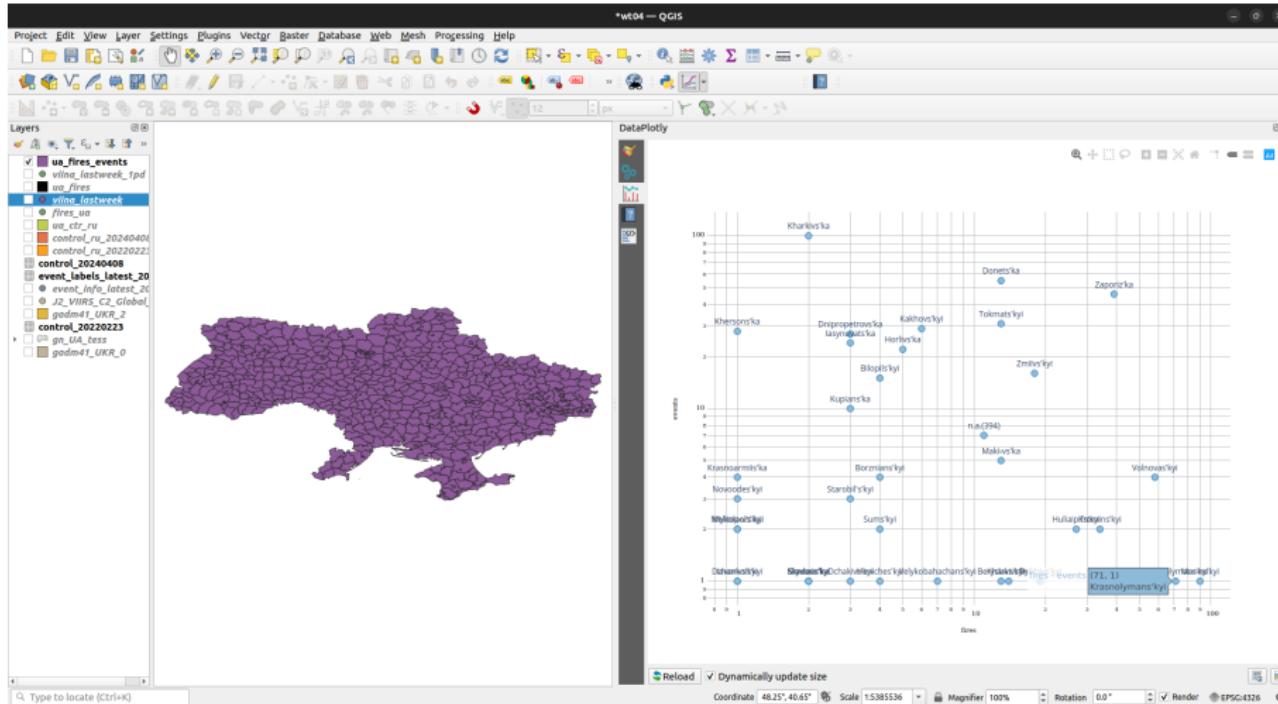
You can think of the points falling along the (imaginary) red line here as locations where both active fires data and media reports captured the same number of incidents. Points below (above) this line are ones where the fires data caught more (fewer) incidents than media reports



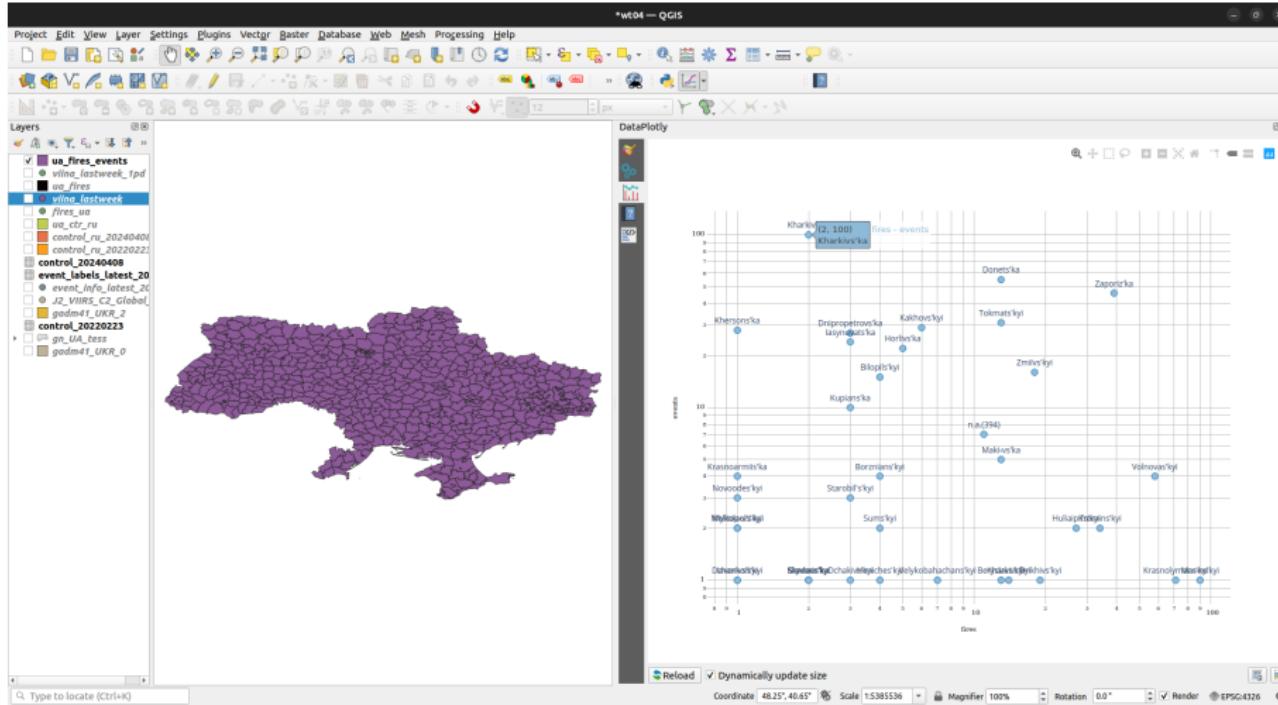
Locations with more active fires than media reports include places like Mar'inka...



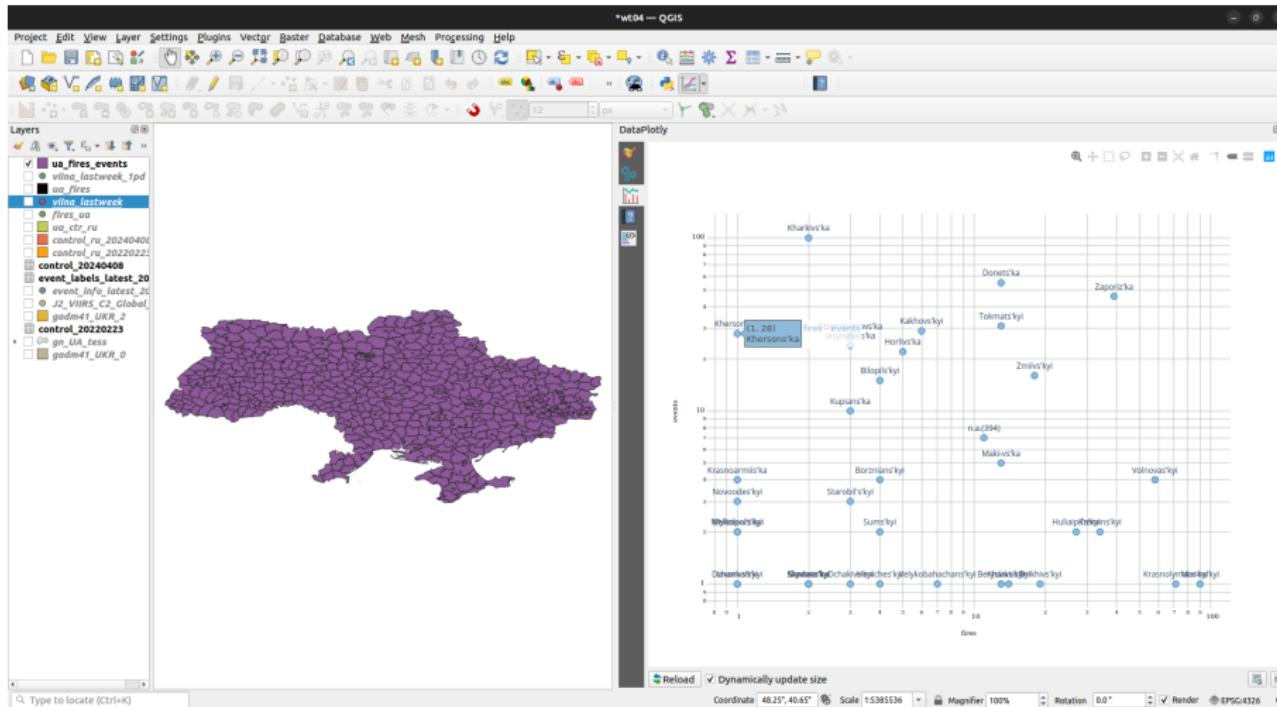
... and Lyman (in former Krasnolymans'kyy district). These are mostly-destroyed front line towns that are hard for journalists to access



Places where media reports capture more incidents include Kharkiv (Ukraine's second largest city, under Ukrainian control)...



...and Kherson (another province capital near the front line, but under Ukrainian control). Both of these places are far easier for journalists to access than the “no man’s land” towns in the lower-right corner



You can perform all these steps in R  
(see replication code `wt04_demo.R` in `WT04.zip`)

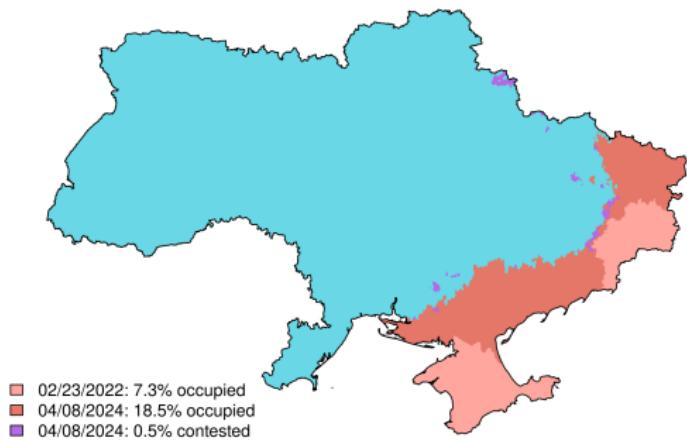


Figure 13: Vignette 1

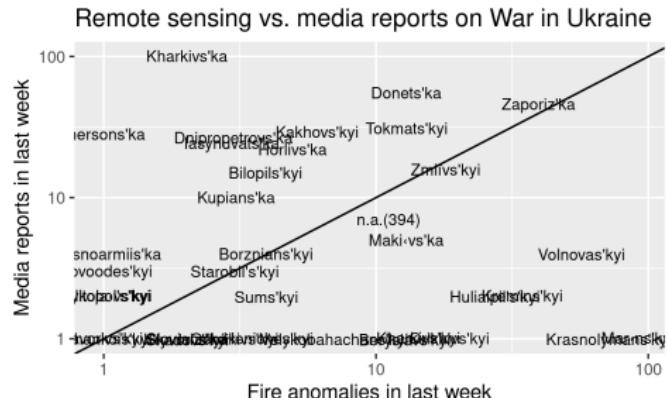


Figure 14: Vignette 2