

SEST-6577

Geographic Information Systems for Security Studies

Lab 03 (+ Problem Set 3)

Yuri M. Zhukov
Associate Professor
Georgetown University

September 30, 2025

Analysis and Geoprocessing with NYC OpenData



Figure 1: A lot of cool stuff here

Overview of lab exercise and problem set

1. Lab exercise
 - a) Map of **bicycle crashes** (per capita) in each NYC community district
 - b) Map of **bike lanes** per district (normalized for population)
2. Problem set
 - a) Map of **rat activity** (per capita) in each NYC community district

We will first make **this map**

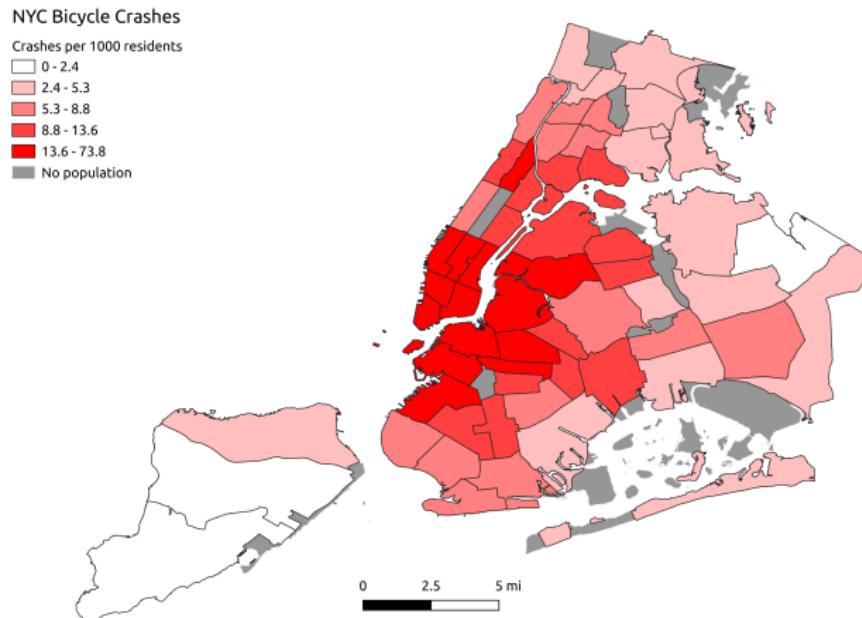


Figure 2: Bike crashes per capita in NYC

This will require joining community districts polygons ...

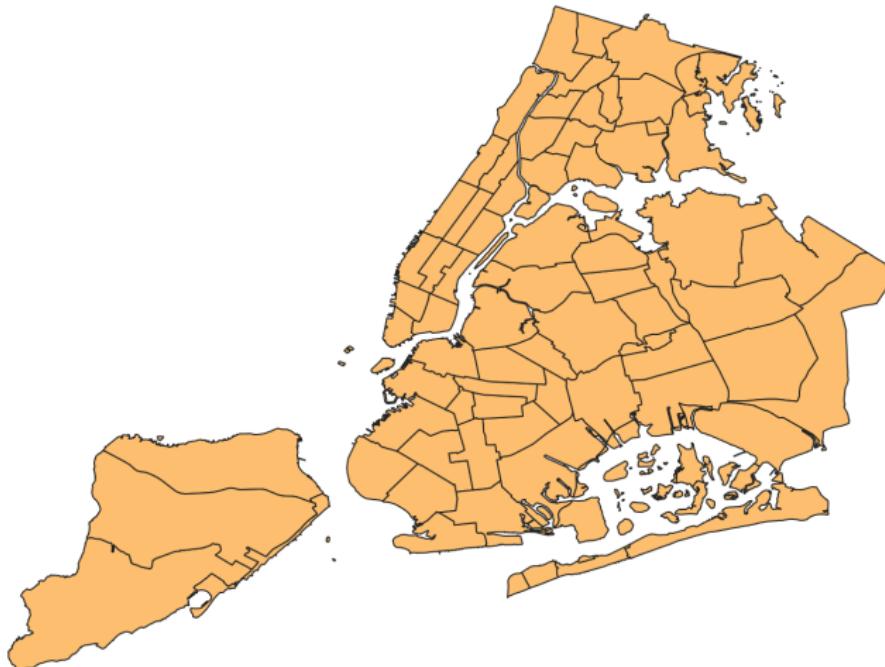


Figure 3: Community districts in NYC

... to a (non-geographic) spreadsheet of population statistics

The screenshot shows a LibreOffice Calc spreadsheet titled "NYC_Pop.csv". The data consists of 27 rows and 10 columns. Column A is labeled "boro_code" and contains integers from 1 to 3. Column B is labeled "Borough" and lists various neighborhoods. Column C is labeled "cd_code" and column D is labeled "cd_name". Columns E through I represent population counts for the years 1970, 1980, 1990, 2000, and 2010 respectively. The data includes many neighborhoods from all five boroughs of New York City.

A	Borough	cd_code	cd_name	pop_1970	pop_1980	pop_1990	pop_2000	pop_2010
1	2:Bronx	1:Melrose	Mott Haven, Port Morris	138557	78441	77214	82159	91497
2	2:Bronx	2:Hunts Point, Longwood		99493	34399	39443	46824	52246
3	2:Bronx	3:Morrisania, Crotona Park East		150636	53635	57162	68574	79762
4	2:Bronx	4:Hibridge, Concourse Village		144207	114312	119962	139563	146441
5	2:Bronx	5:University Hts., Fordham, Mt. Hope		121807	107995	118435	128313	128200
6	2:Bronx	6:East Tremont, Belmont		114137	65016	68061	75688	83268
7	2:Bronx	7:Bedford Park, Norwood, Fordham		113764	116827	128588	141411	139286
8	2:Bronx	8:Riverdale, Kingsbridge, Marble Hill		103543	98275	97030	101332	101731
9	2:Bronx	9:Soundview, Parkchester		166442	167627	155970	167859	172298
10	2:Bronx	10:Troy Hills Nc., Co-op City, Pelham Bay		84948	106516	108093	115948	120392
11	2:Bronx	11:Pelham Parkway, Morris Park, Laconia		105980	99085	97842	110706	113232
12	2:Bronx	12:Wakefield, Williamsbridge		135010	128226	129620	149077	152344
13	3:Brooklyn	1:Williamsburg, Greenpoint		178990	142942	155972	160238	173083
14	3:Brooklyn	2:Brooklyn Heights, Fort Greene		110221	92732	94534	98620	99617
15	3:Brooklyn	3:Bedford Stuyvesant		203380	133379	138896	143867	152985
16	3:Brooklyn	4:Bushwick		137902	92497	102572	104358	112634
17	3:Brooklyn	5:East New York, Stannet City		170791	154931	161350	173198	182896
18	3:Brooklyn	6:Park Slope, Carroll Gardens		138933	110228	102724	104054	104709
19	3:Brooklyn	7:Sunset Park, Windsor Terrace		111607	98567	102553	120063	126230
20	3:Brooklyn	8:Crown Heights North		121821	88796	96400	96076	96317
21	3:Brooklyn	9:Crown Heights South, Wrigate		101047	96669	110715	104014	98429
22	3:Brooklyn	10:Bay Ridge, Dyker Heights		129822	118187	110612	122542	124491
23	3:Brooklyn	11:Bensonhurst, Bath Beach		170119	155072	149994	172129	181981
24	3:Brooklyn	12:Borough Park, Ocean Parkway		166301	155899	160018	185046	191382
25	3:Brooklyn	13:Coney Island, Brighton Beach		97750	100030	102596	106120	104278
26	3:Brooklyn	14:Flatbush, Midwood		137041	143859	159825	168806	160664

Figure 4: NYC population data

... followed by a point-in-polygon analysis and some field calculations

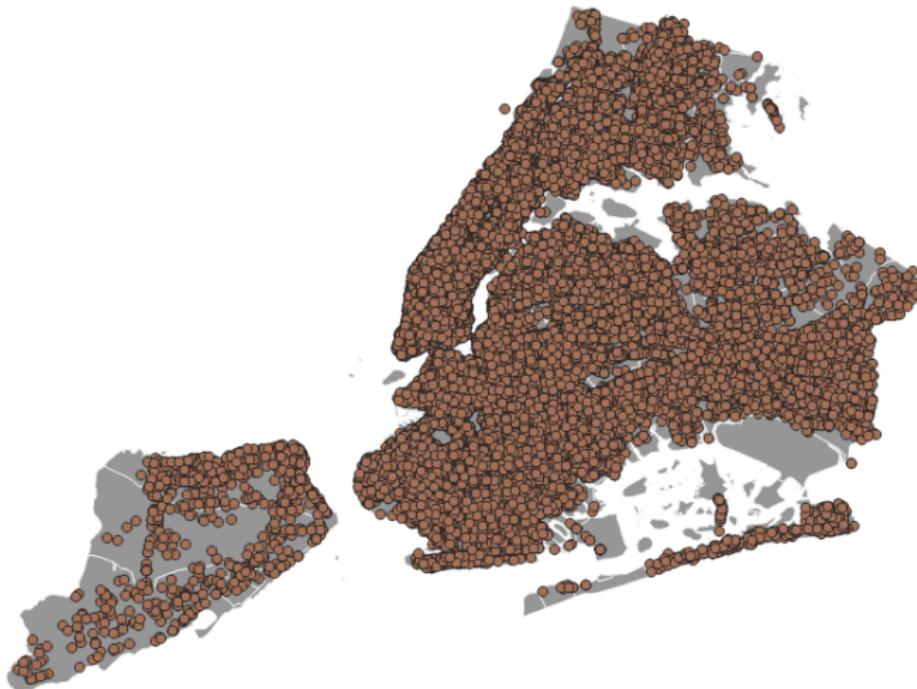


Figure 5: NYC bike crash data

We will then make **this map**

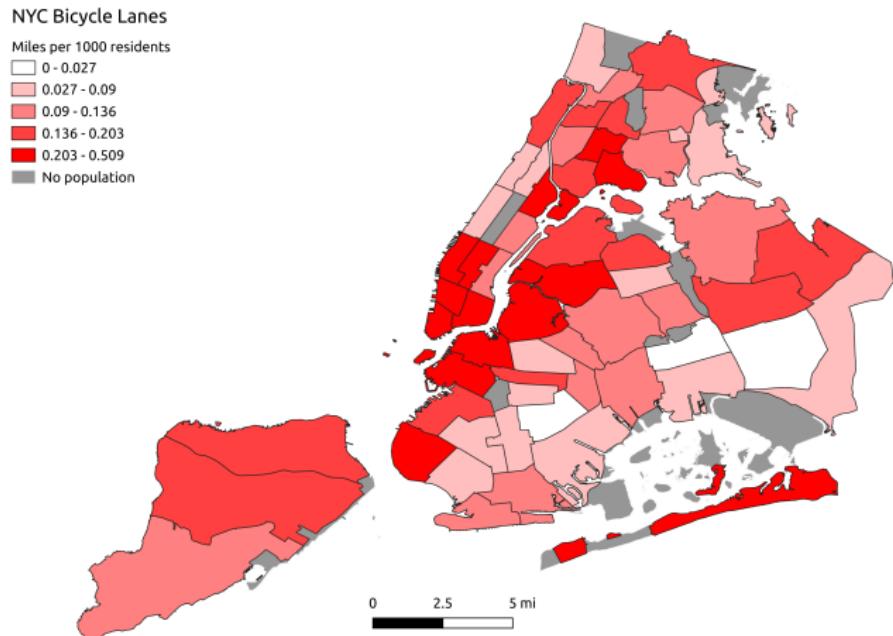


Figure 6: Bike lanes per capita in NYC

This will require some line-in-polygon analysis



Figure 7: Bike lanes in NYC

You will make this map for your **problem set**

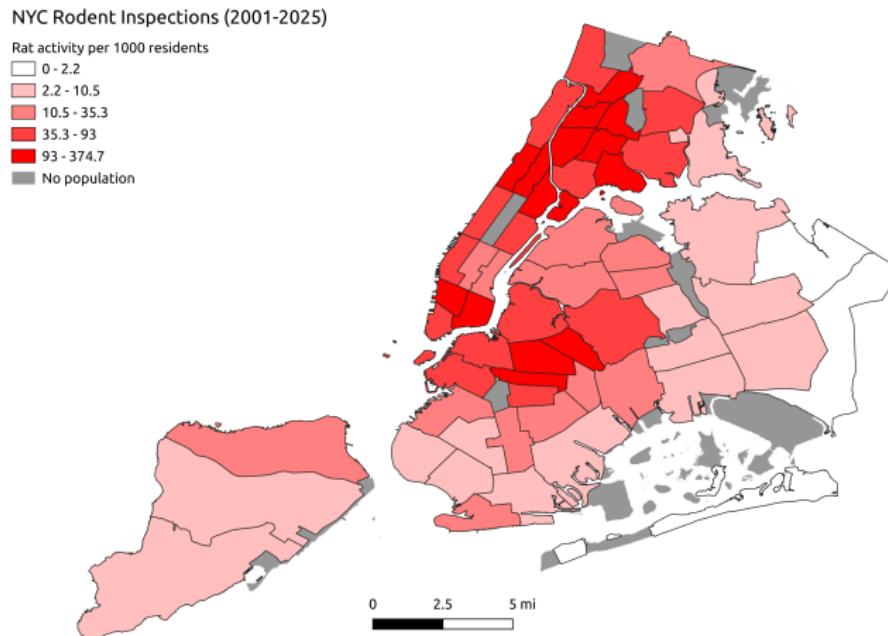


Figure 8: Rat activity per capita in NYC

This will involve (again) point-in-polygon analysis

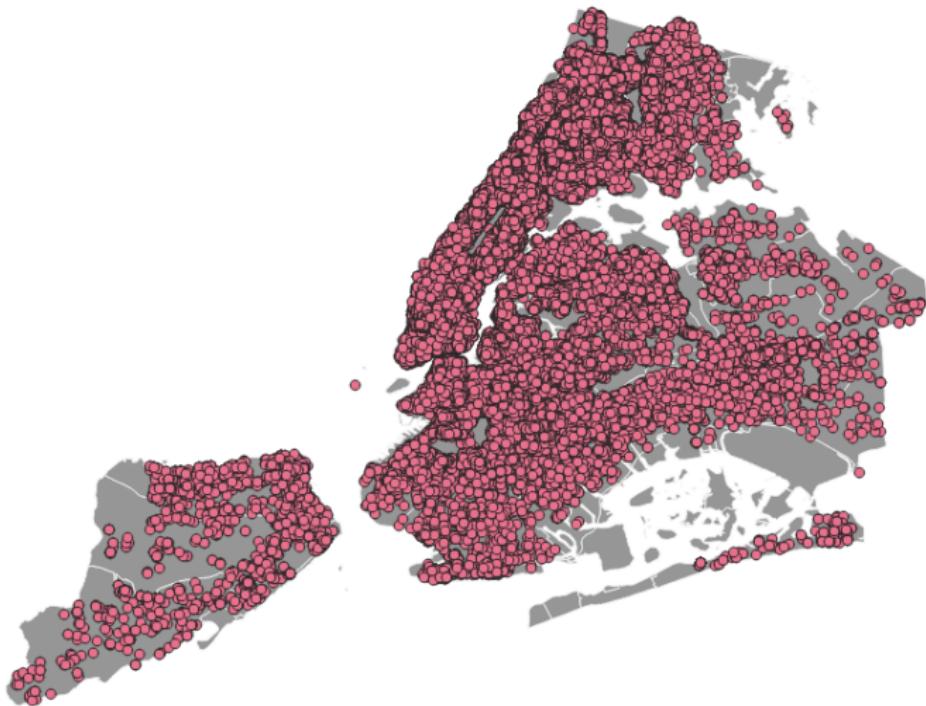


Figure 9: Rodent inspections in NYC

You can make these plots in QGIS, R, or Python. Instructions for QGIS and R are below.

Bicycle Crashes per 1000 Residents

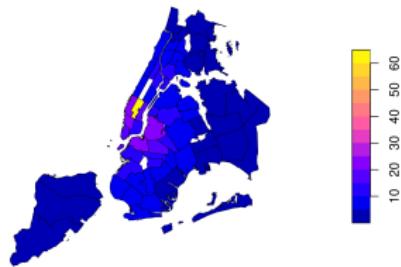


Figure 10: Map 1 in R

Miles of Bicycle Lane per 1000 Residents

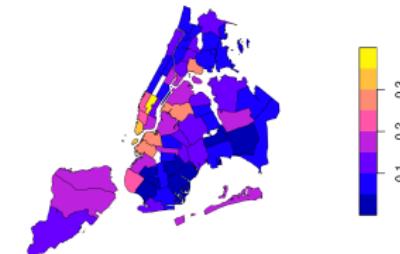


Figure 11: Map 2 in R

Rat Activity per 1000 Residents

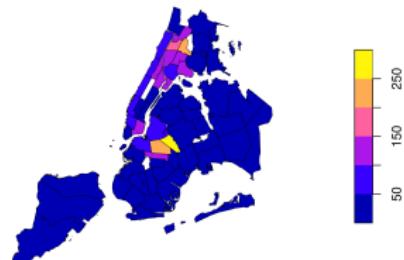


Figure 12: PS03 in R

Python replication code is also in the ZIP file.

Bicycle Collisions per 1000 Residents

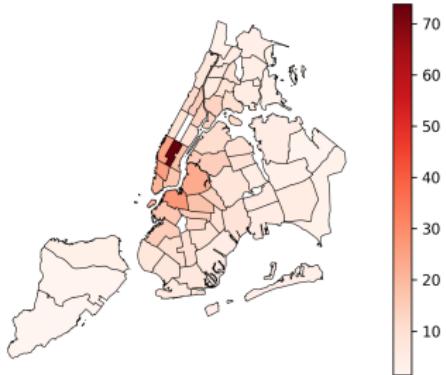


Figure 13: Map 1 in py

Bike Lane per 1000 Residents (mi)

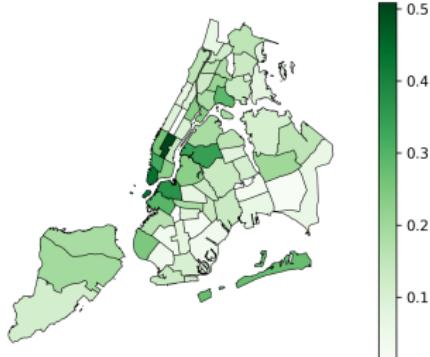


Figure 14: Map 2 in py

Rat Activity per 1000 Residents

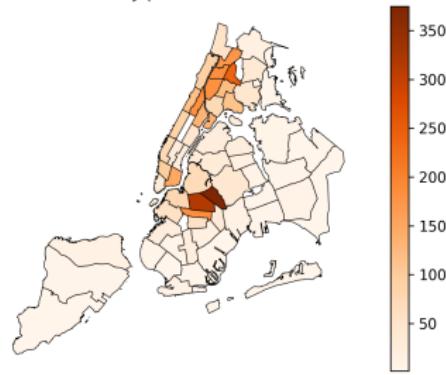


Figure 15: PS03 in py

We will put multiple types of data on same map:

Category	Type	Format	Data source
Community district borders	Vector (polygon)	.geojson	NYC OpenData
Community district population	Table	.csv	NYC OpenData
Vehicle collisions involving bicycles	Vector (point)	.csv	NYC OpenData
Bike routes	Vector (polyline)	.geojson	NYC OpenData
Rodent inspections finding "rat activity"	Vector (point)	.csv	NYC OpenData

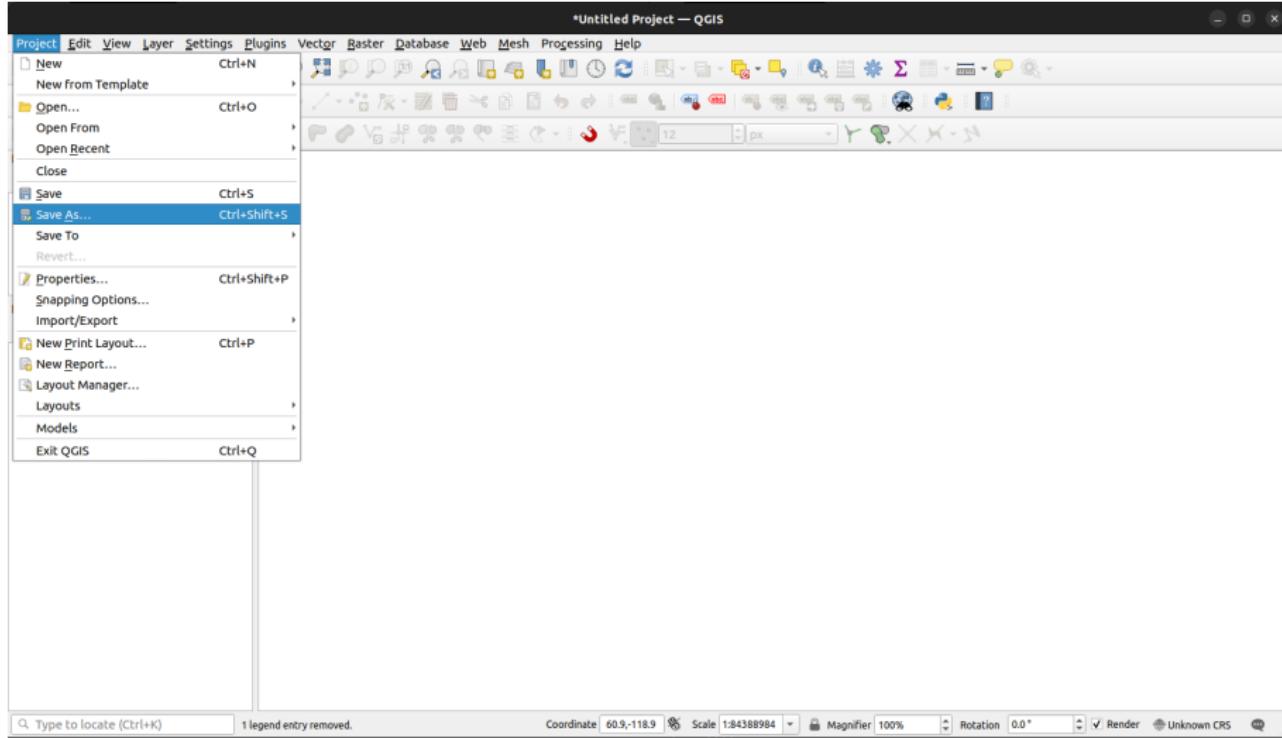
These are all in the Lab03PS03.zip file posted on Canvas.

Let's open QGIS...

QGIS

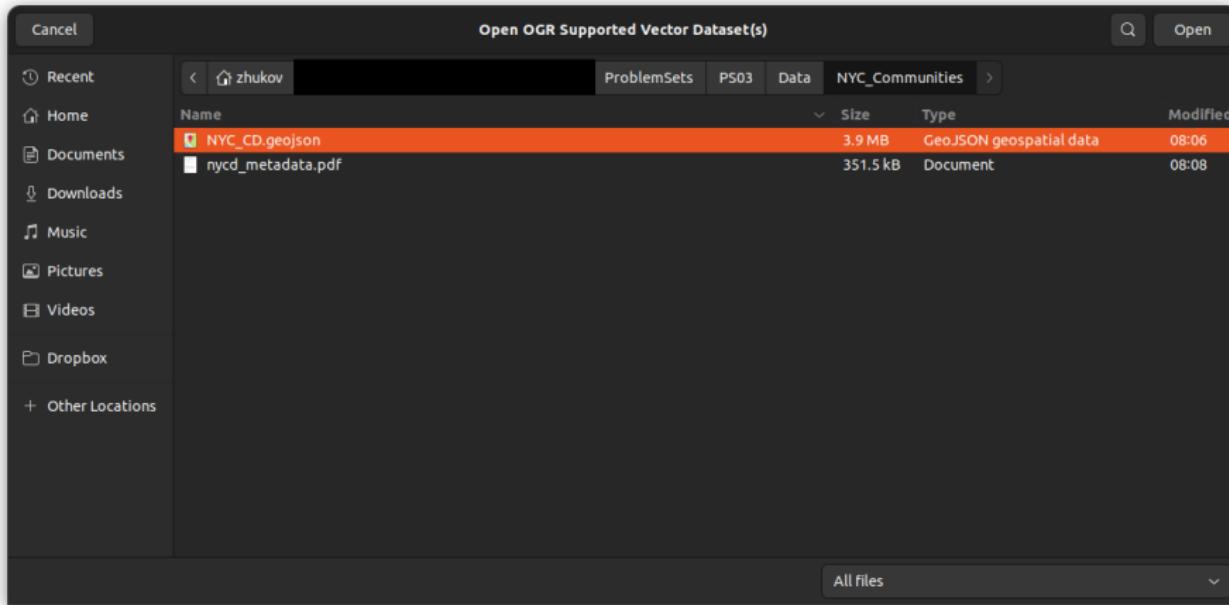
Always save your progress!

Go to Project → Save As...

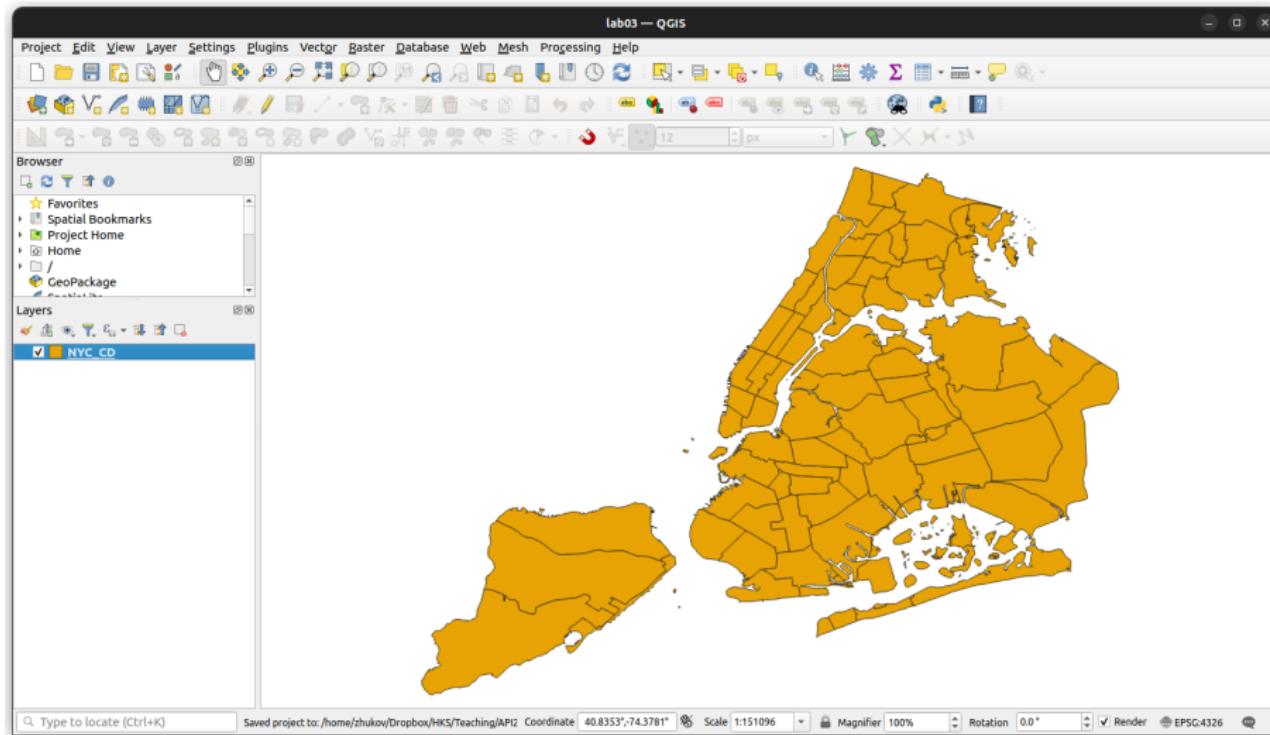


Map 1

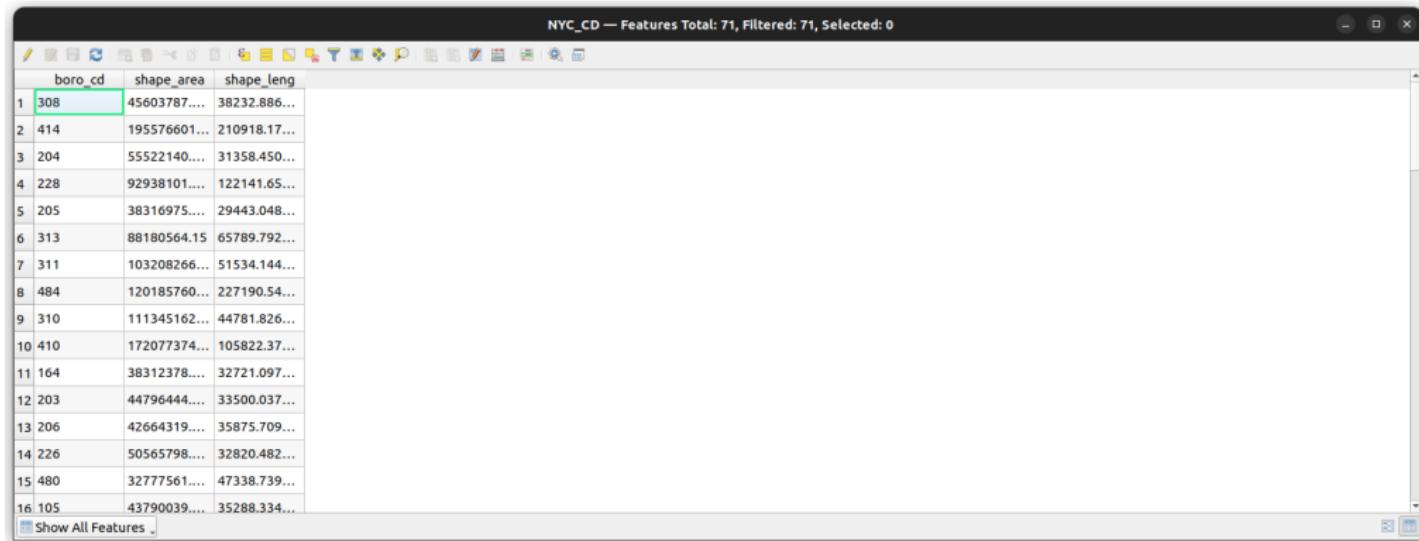
Load community districts: Layer → Add Layer → Add Vector Layer...
Open the file **NYC_CD.geojson** in the **NYC_Communities** folder. Click Add



This dataset includes 71 polygons, including New York's 59 community districts and 12 areas outside of community districts' jurisdiction (parks, waterfront)



Open the Attribute Table for this layer. We see a borough-district code (boro_cd) and geometric statistics (shape_area, shape_leng)



The screenshot shows the QGIS Attribute Table for the 'NYC_CD' layer. The table has three columns: 'boro_cd', 'shape_area', and 'shape_leng'. The first row, which contains the value '308' in the 'boro_cd' column, is highlighted with a green border. The table displays 16 rows of data, with the 17th row being a summary row labeled 'Show All Features'.

	boro_cd	shape_area	shape_leng
1	308	45603787...	38232.886...
2	414	195576601...	210918.17...
3	204	55522140...	31358.450...
4	228	92938101...	122141.65...
5	205	38316975...	29443.048...
6	313	88180564.15	65789.792...
7	311	103208266...	51534.144...
8	484	120185760...	227190.54...
9	310	111345162...	44781.826...
10	410	172077374...	105822.37...
11	164	38312378...	32721.097...
12	203	44796444...	33500.037...
13	206	42664319...	35875.709...
14	226	50565798...	32820.482...
15	480	32777561...	47338.739...
16	105	43790039...	35288.334...
Show All Features			

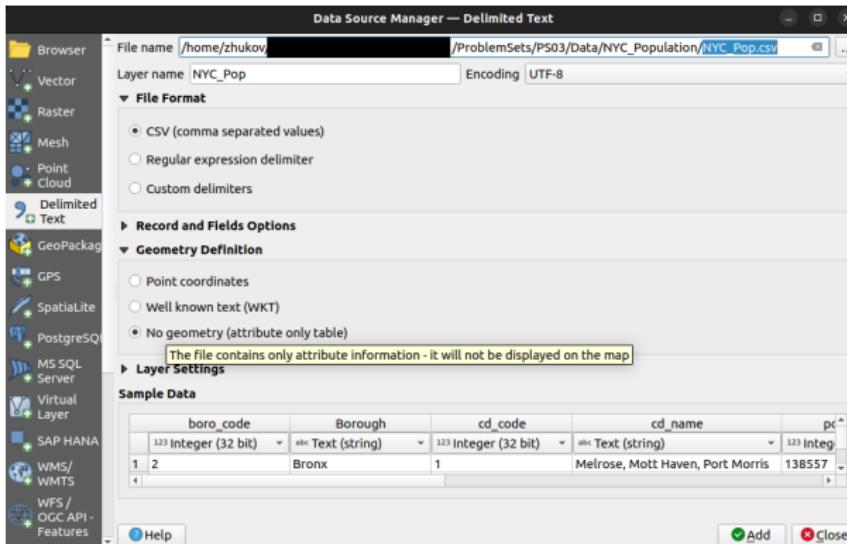
For a definition of boro_cd, we can look at the metadata for NYC_CD (nycd_metadata.pdf). The metadata seem out of date (BoroCD \neq boro_cd), but there's important info here about this code is constructed

```
FIELD BoroCD
▶
 * ALIAS BoroCD
 * DATA TYPE SmallInteger
 * WIDTH 2
 * PRECISION 0
 * SCALE 0
FIELD DESCRIPTION
Community district number preceded by BoroCode.
```

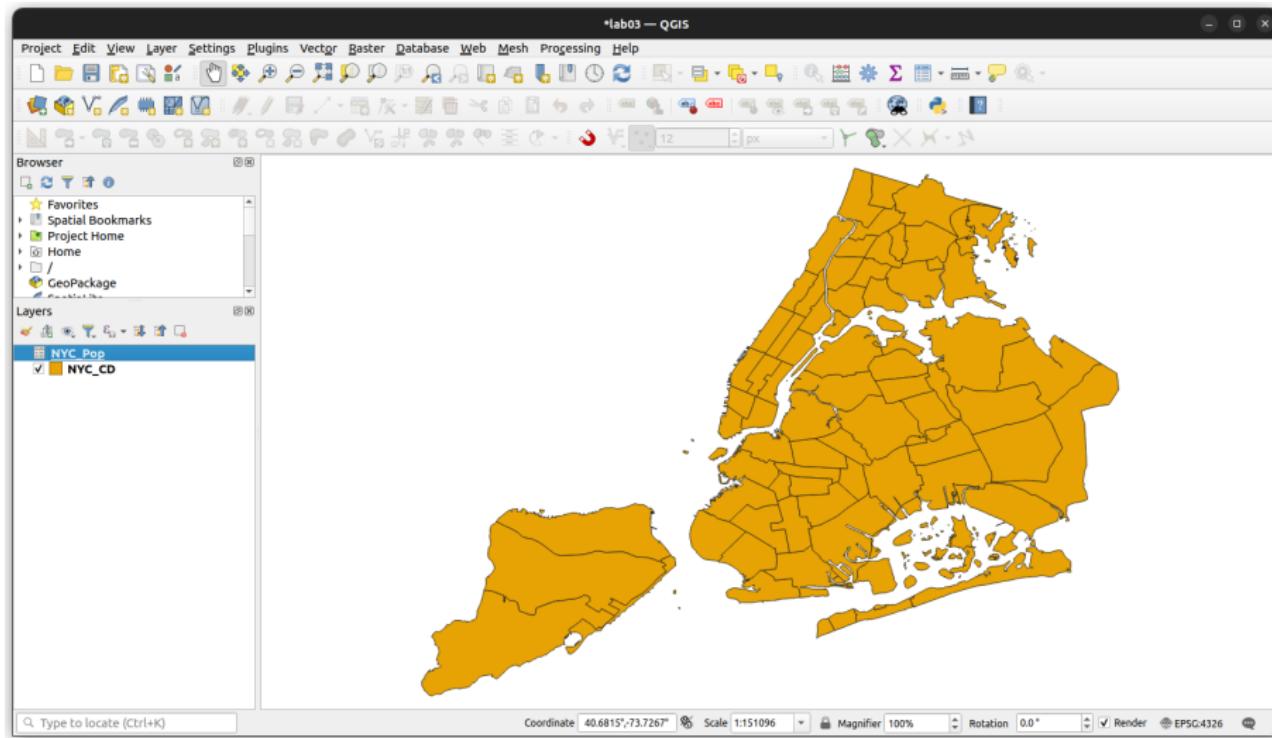
[Hide Field BoroCD ▲](#)

Load population data:

- Layer → Add Layer → Add Delimited Text Layer...
- Navigate to the file NY_Pop.csv in the NYC_Population folder
- Set Geometry Definition = No geometry (attribute table only)
- Click Add



The new layer is visible in the side menu. But it doesn't appear on the map, since it's not a spatial dataset. Right-click NYC_Pop → Open Attribute Table



This is a table of community district names and population statistics. Maybe we can link this table to NYC_CD... but there is no common variable

NYC_Pop — Features Total: 59, Filtered: 59, Selected: 0

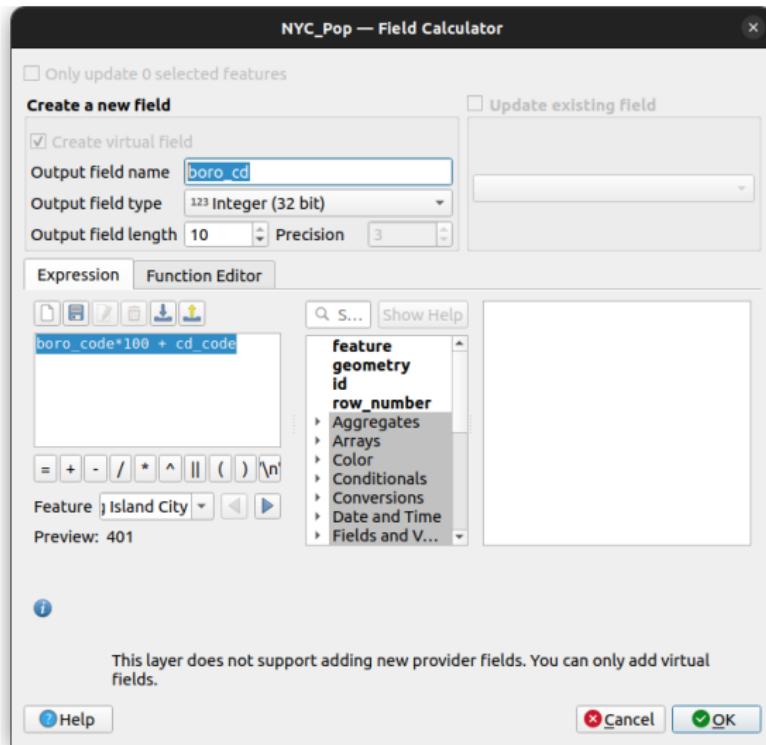
	boro_code	Borough	cd_code	cd_name	pop_1970	pop_1980	pop_1990	pop_2000	pop_2010
1	2	Bronx	1	Melrose, M...	138557	78441	77214	82159	91497
2	2	Bronx	2	Hunts Poin...	99493	34399	39443	46824	52246
3	2	Bronx	3	Morrisania,...	150636	53635	57162	68574	79762
4	2	Bronx	4	Highbridge...	144207	114312	119962	139563	146441
5	2	Bronx	5	University ...	121807	107995	118435	128313	128200
6	2	Bronx	6	East Tremo...	114137	65016	68061	75688	83268
7	2	Bronx	7	Bedford Pa...	113764	116827	128588	141411	139286
8	2	Bronx	8	Riverdale, ...	103543	98275	97030	101332	101731
9	2	Bronx	9	Soundview,...	166442	167627	155970	167859	172298
10	2	Bronx	10	Throgs Nk,...	84948	106516	108093	115948	120392
11	2	Bronx	11	Pelham Pk...	105980	99080	97842	110706	113232
12	2	Bronx	12	Wakefield, ...	135010	128226	129620	149077	152344
13	3	Brooklyn	1	Williamsbu...	179390	142942	155972	160338	173083
14	3	Brooklyn	2	Brooklyn H...	110221	92732	94534	98620	99617
15	3	Brooklyn	3	Bedford St...	203380	133379	138696	143867	152985
16	3	Brooklyn	4	Bushwick	137902	92497	102572	104358	112634

Wait. NYC_CD has a boro_cd code, which is “Community district number preceded by BoroCode”. Maybe we can create a boro_cd code in NYC_Pop, using boro_code and cd_code? Open Field Calculator

The screenshot shows the QGIS Field Calculator dialog for the 'NYC_Pop' layer. The title bar indicates 'NYC_Pop — Features Total: 59, Filtered: 59, Selected: 0'. The dialog contains a toolbar at the top with various icons. Below the toolbar is a table with columns: boro_code, Borough, cd_code, cd_name, Open field calculator (Ctrl+I) (highlighted with a green border), pop_1990, pop_2000, and pop_2010. The table lists 16 rows of data for Bronx and Brooklyn community districts. The 'Open field calculator' column is empty for all rows except the first one, where it contains the formula '\$boro_code + \$cd_code'. The data is as follows:

	boro_code	Borough	cd_code	cd_name	Open field calculator (Ctrl+I)	pop_1990	pop_2000	pop_2010
1	2	Bronx	1	Melrose, M...	138557	78441	77214	82159
2	2	Bronx	2	Hunts Poin...	99493	34399	39443	46824
3	2	Bronx	3	Morrisania,...	150636	53635	57162	68574
4	2	Bronx	4	Highbridge...	144207	114312	119962	139563
5	2	Bronx	5	University ...	121807	107995	118435	128313
6	2	Bronx	6	East Tremo...	114137	65016	68061	75688
7	2	Bronx	7	Bedford Pa...	113764	116827	128588	141411
8	2	Bronx	8	Riverdale, ...	103543	98275	97030	101332
9	2	Bronx	9	Soundview,...	166442	167627	155970	167859
10	2	Bronx	10	Throgs Nk,...	84948	106516	108093	115948
11	2	Bronx	11	Pelham Pk...	105980	99080	97842	110706
12	2	Bronx	12	Wakefield, ...	135010	128226	129620	149077
13	3	Brooklyn	1	Williamsbu...	179390	142942	155972	160338
14	3	Brooklyn	2	Brooklyn H...	110221	92732	94534	98620
15	3	Brooklyn	3	Bedford St...	203380	133379	138696	143867
16	3	Brooklyn	4	Bushwick	137902	92497	102572	104358

Create a new variable, called `boro_cd`, of type Integer. For the Expression, write `boro_code * 100 + cd_code`. Click OK

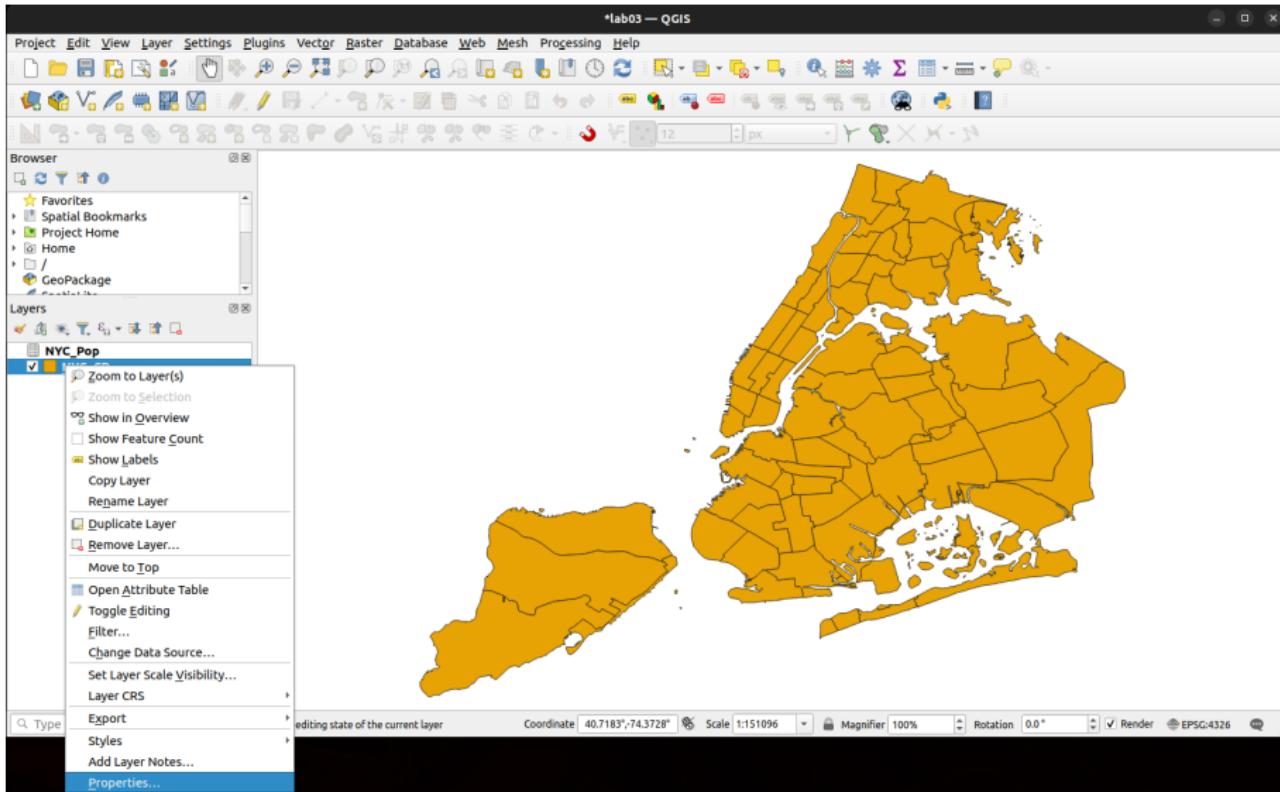


You should now see the variable `boro_cd` in the attribute table.

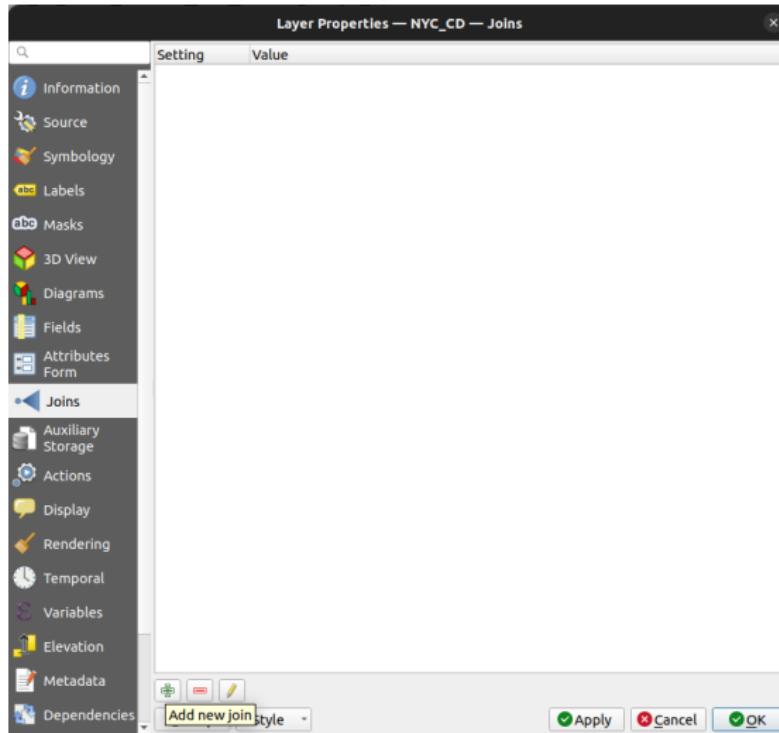
NYC_Pop — Features Total: 59, Filtered: 59, Selected: 0

	boro_code	Borough	cd_code	cd_name	pop_1970	pop_1980	pop_1990	pop_2000	pop_2010	boro_cd
1	2	Bronx	1	Melrose, M...	138557	78441	77214	82159	91497	201
2	2	Bronx	2	Hunts Poin...	99493	34399	39443	46824	52246	202
3	2	Bronx	3	Morrisania,...	150636	53635	57162	68574	79762	203
4	2	Bronx	4	Highbridge...	144207	114312	119962	139563	146441	204
5	2	Bronx	5	University ...	121807	107995	118435	128313	128200	205
6	2	Bronx	6	East Tremo...	114137	65016	68061	75688	83268	206
7	2	Bronx	7	Bedford Pa...	113764	116827	128588	141411	139286	207
8	2	Bronx	8	Riverdale, ...	103543	98275	97030	101332	101731	208
9	2	Bronx	9	Soundview,...	166442	167627	155970	167859	172298	209
10	2	Bronx	10	Throgs Nk,...	84948	106516	108093	115948	120392	210
11	2	Bronx	11	Pelham Pk...	105980	99080	97842	110706	113232	211
12	2	Bronx	12	Wakefield, ...	135010	128226	129620	149077	152344	212
13	3	Brooklyn	1	Williamsbu...	179390	142942	155972	160338	173083	301
14	3	Brooklyn	2	Brooklyn H...	110221	92732	94534	98620	99617	302
15	3	Brooklyn	3	Bedford St...	203380	133379	138696	143867	152985	303
16	3	Brooklyn	4	Bushwick	137902	92497	102572	104358	112634	304

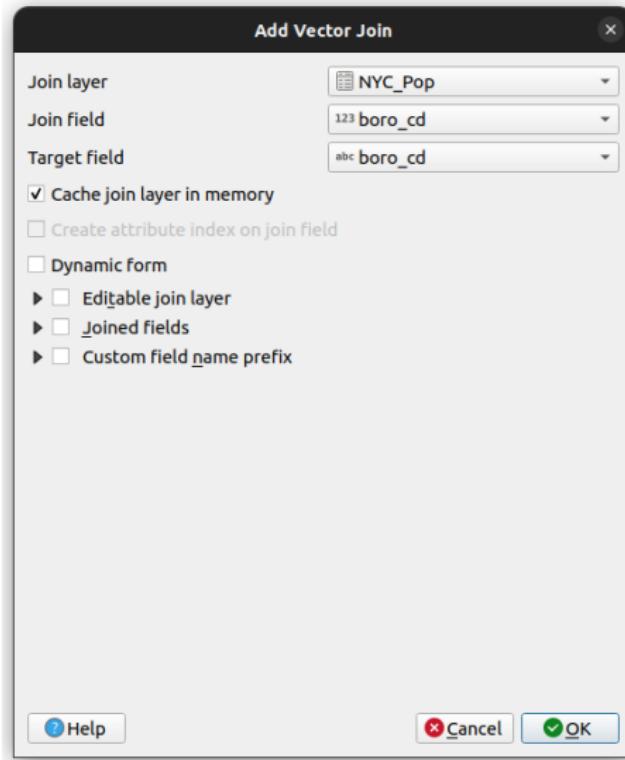
Let's join NYC_CD and NYC_Pop. Go to the Properties for NYC_CD



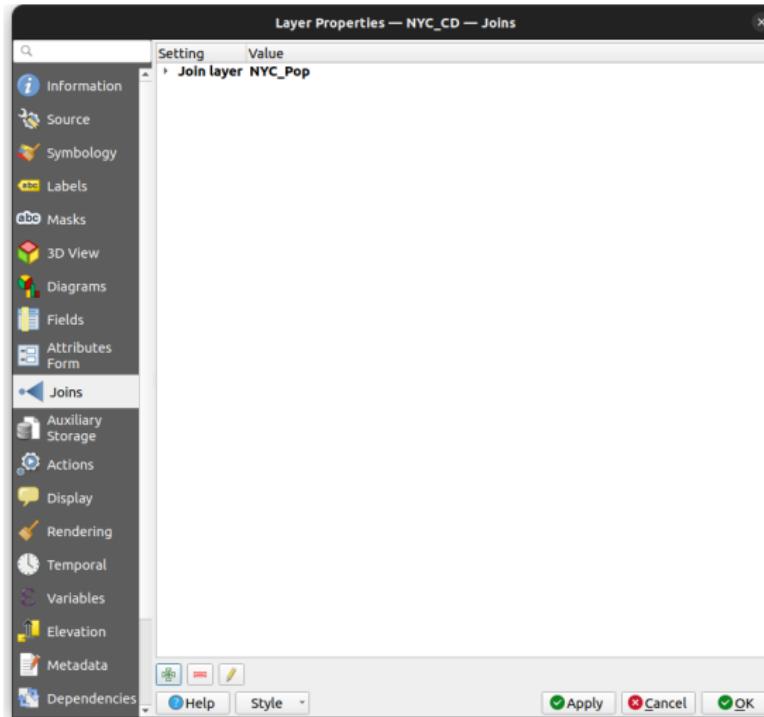
Open the Joins tab in Properties. Click on the “add” button (+)



Select NYC_Pop as the join layer, and boro_cd as the join and target field. Click OK



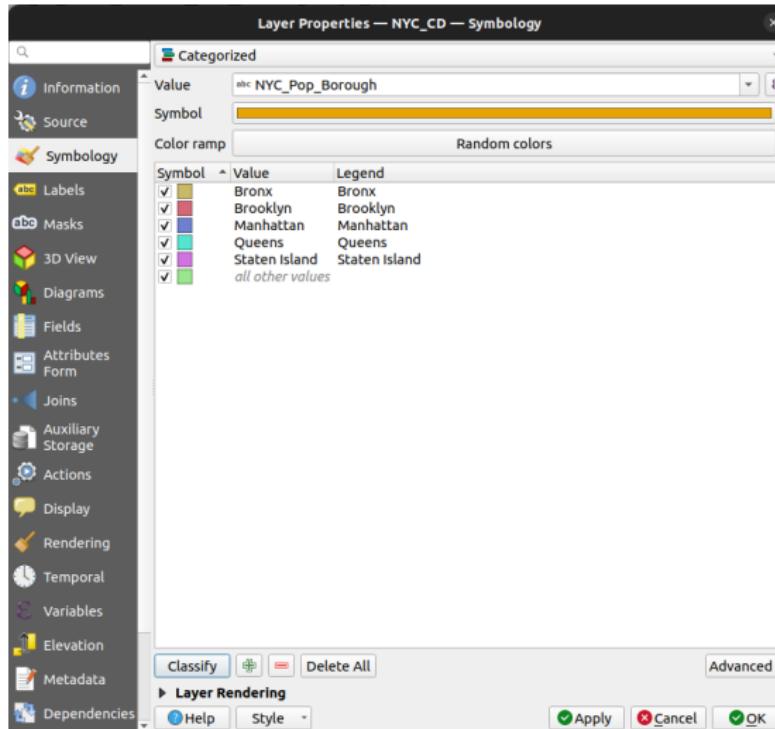
You should now see a new join in the Joins tab



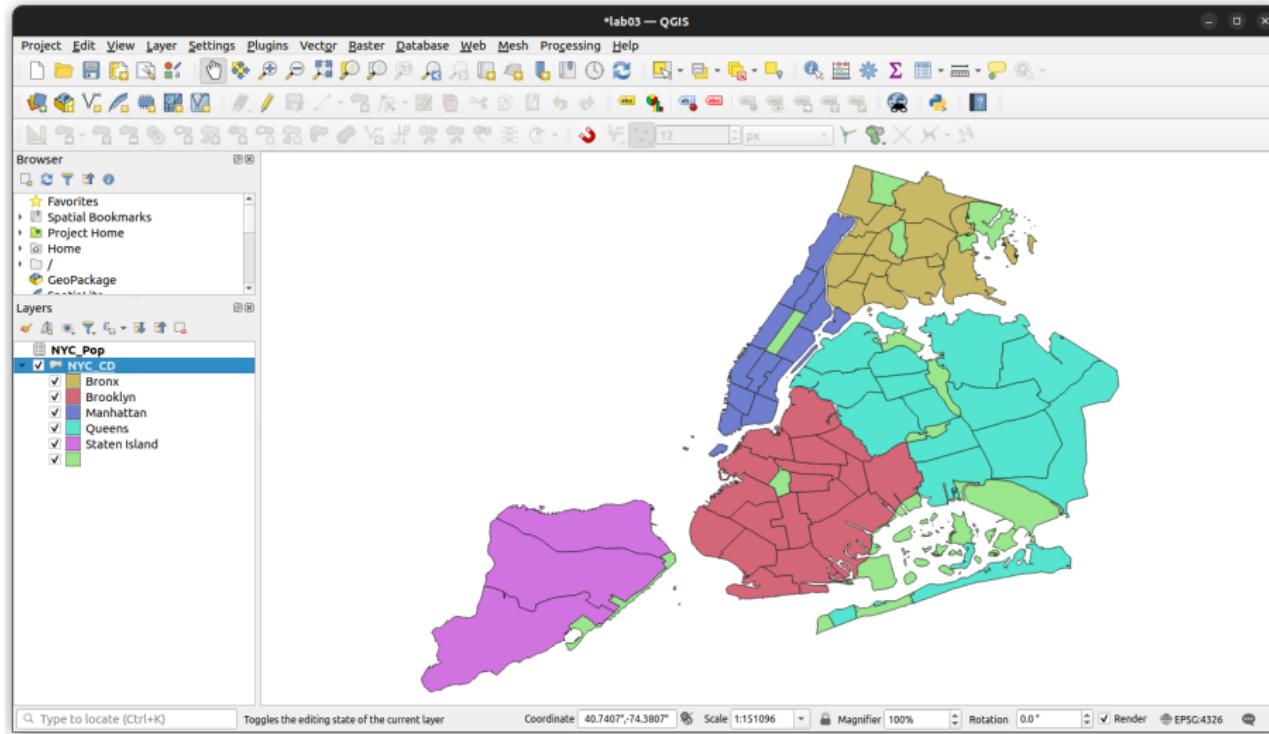
Now go to the Symbology tab in Properties.

Set Symbology type = Categorized, and Value = NYC_Pop_Borough.

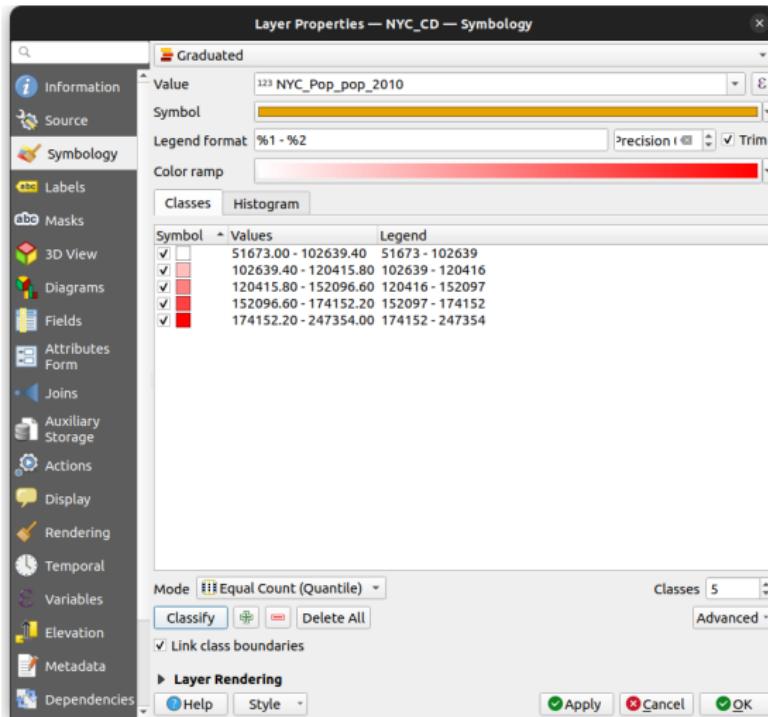
Click Classify, then OK



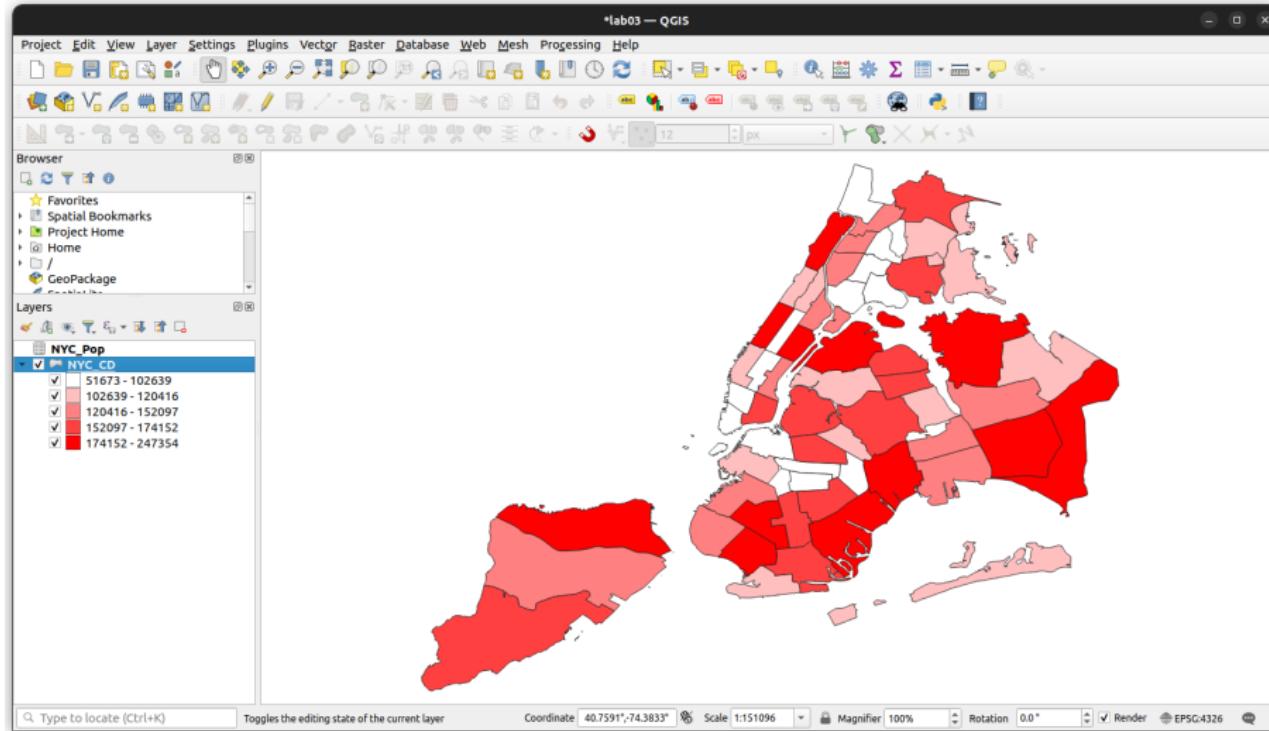
The districts should now be colored by borough



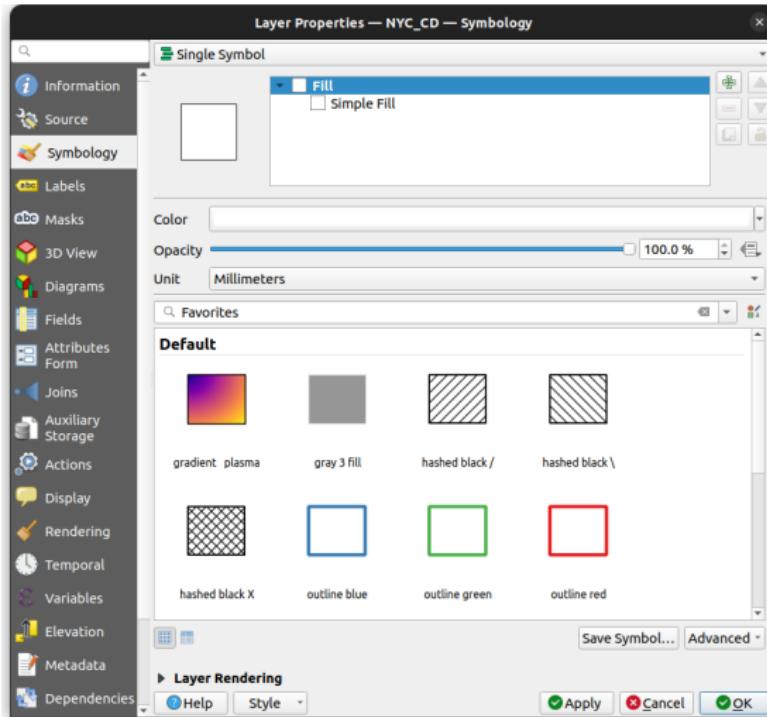
Let's now color them by population size. Go back to Properties → Symbology.
Set Symbology type = Graduated, and Value = NYC_Pop_pop_2010.
Click Classify, then OK



The districts should now be colored by their population in 2010

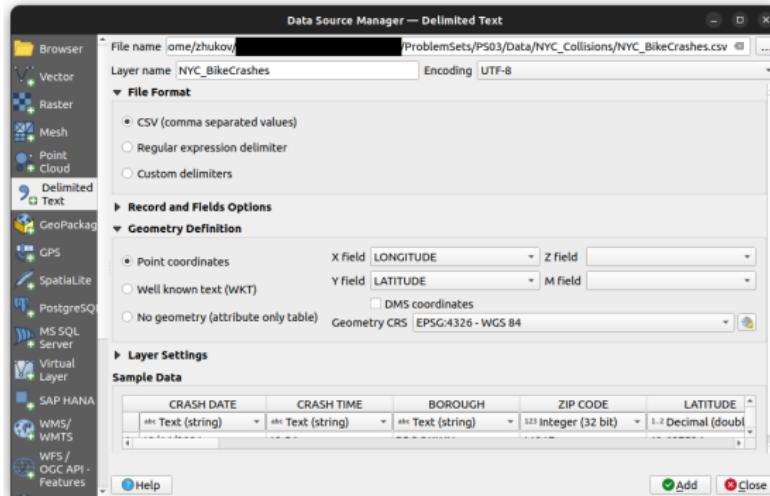


Set the color scheme back to Single symbol in Symbology

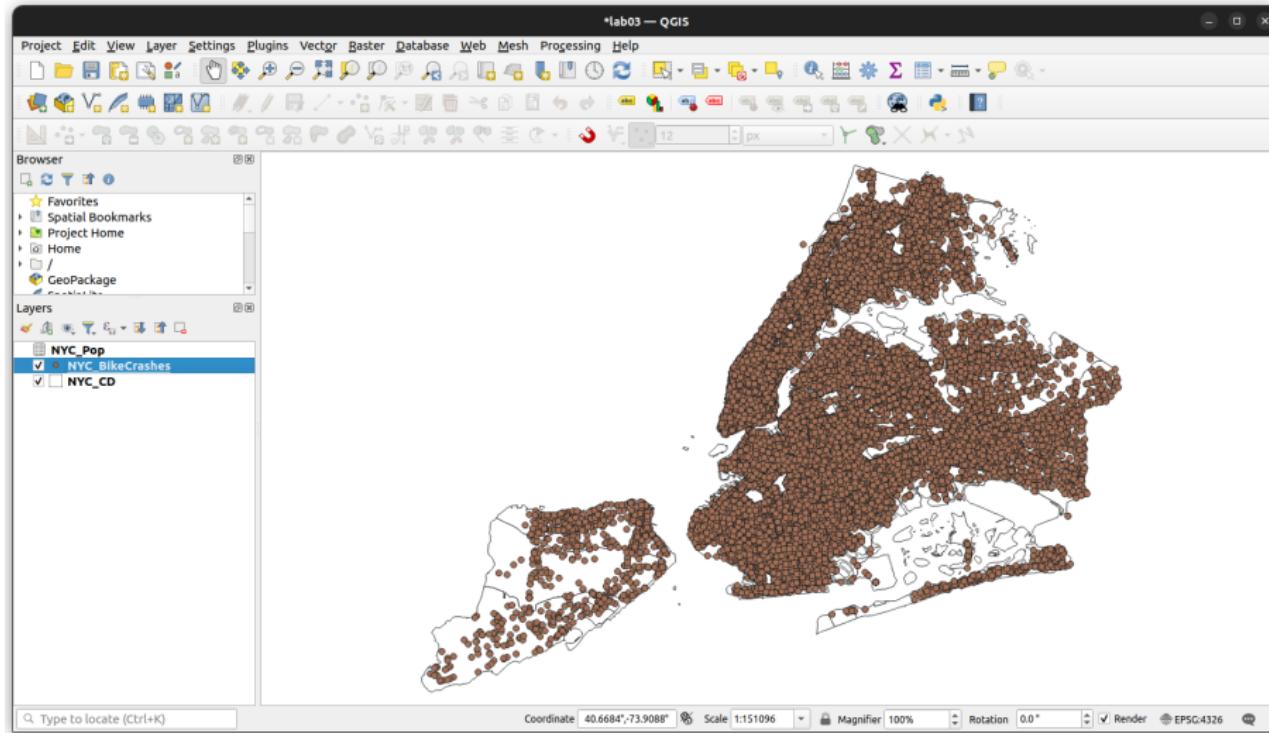


Load data on bicycle crashes:

- Go to Add Delimited Text Layer
- Navigate to NYC_BikeCrashes.csv in the NYC_Collisions folder
- Set Geometry Definition = Point coordinates
- Set X field = LONGITUDE, Y field = LATITUDE.
- Click Add

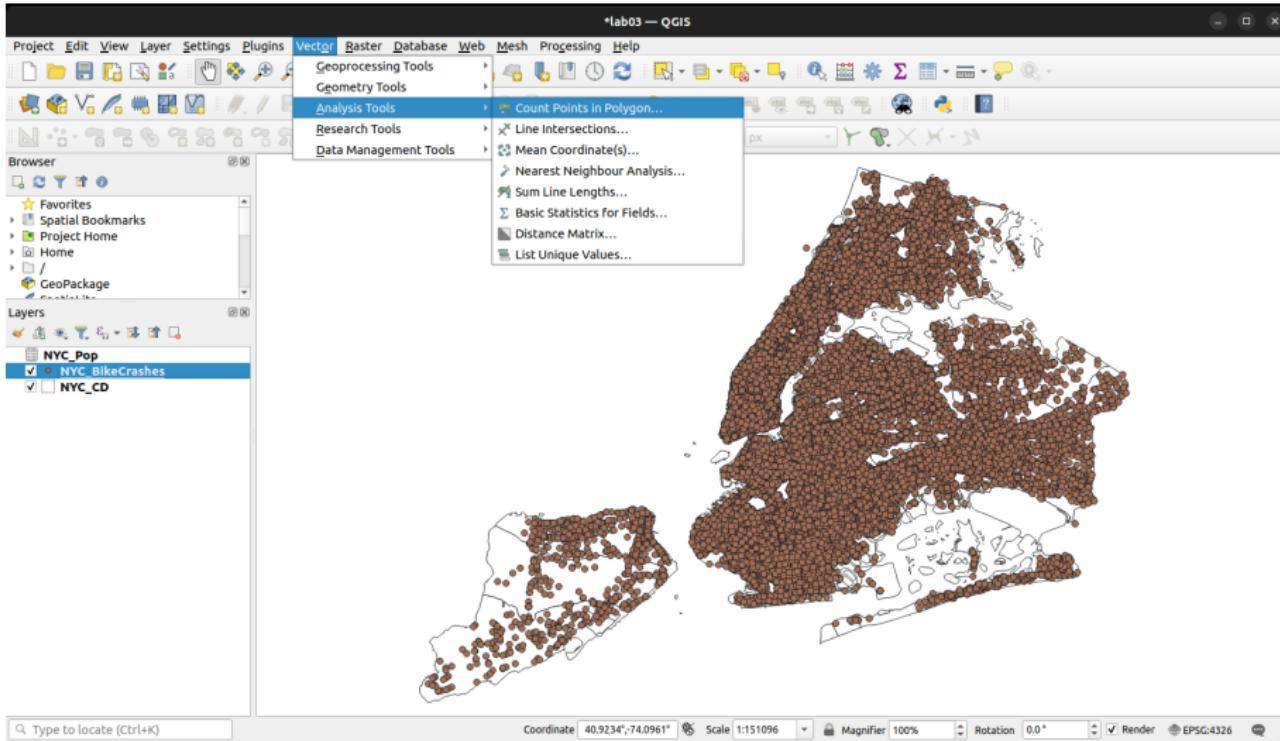


The new layer should be visible in the map, as points.

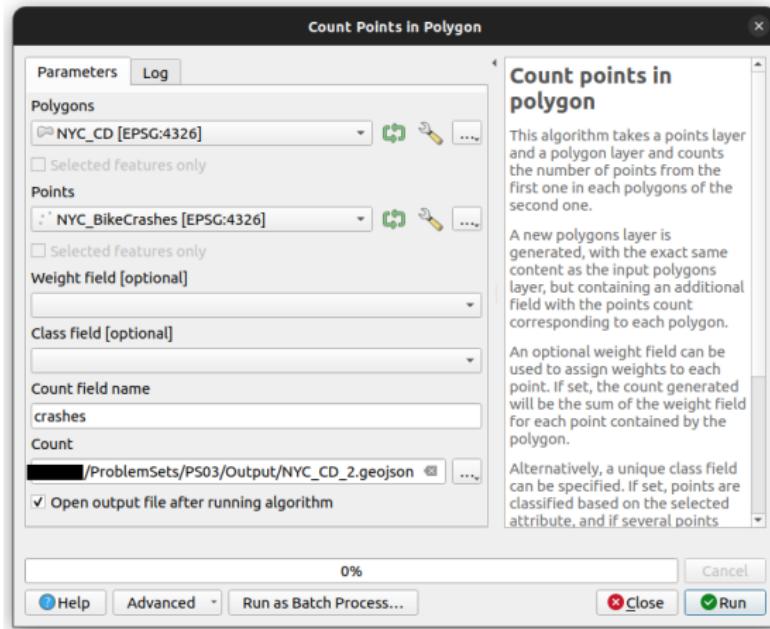


Navigate to the Count Points in Polygon tool.

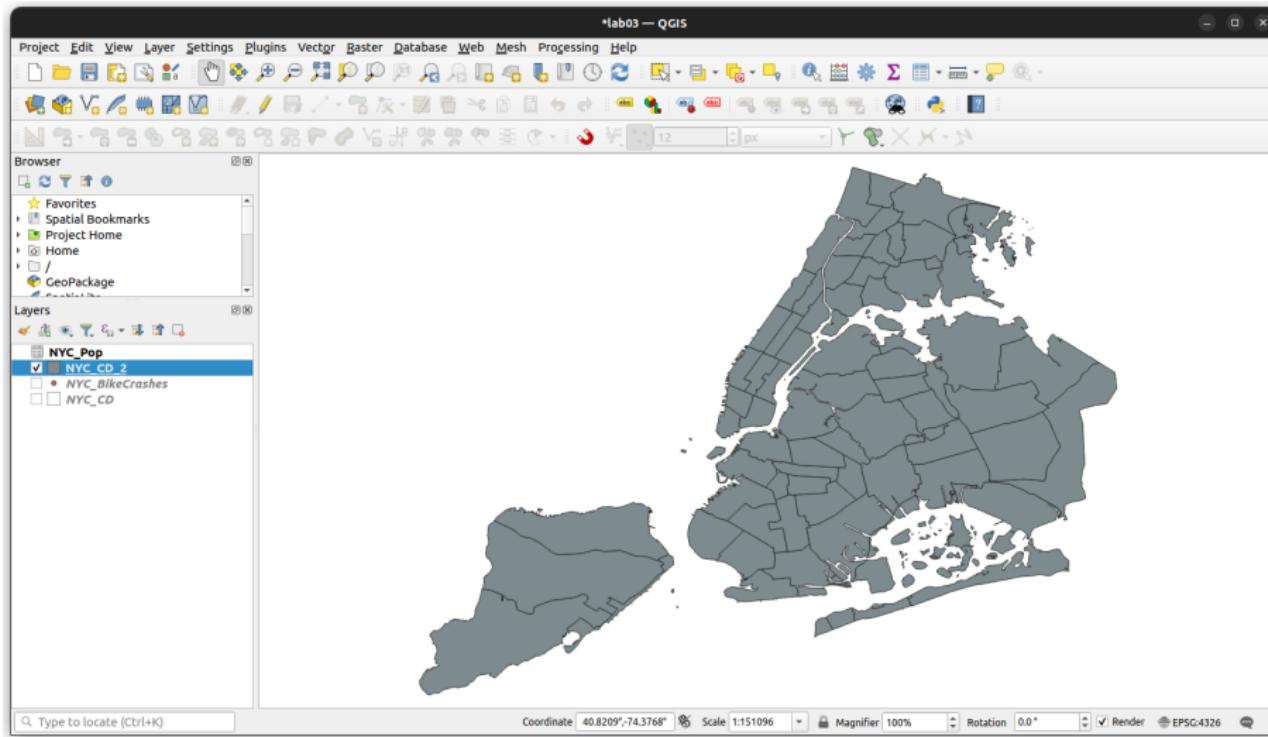
Vector menu → Analysis Tools → Count Points in Polygon



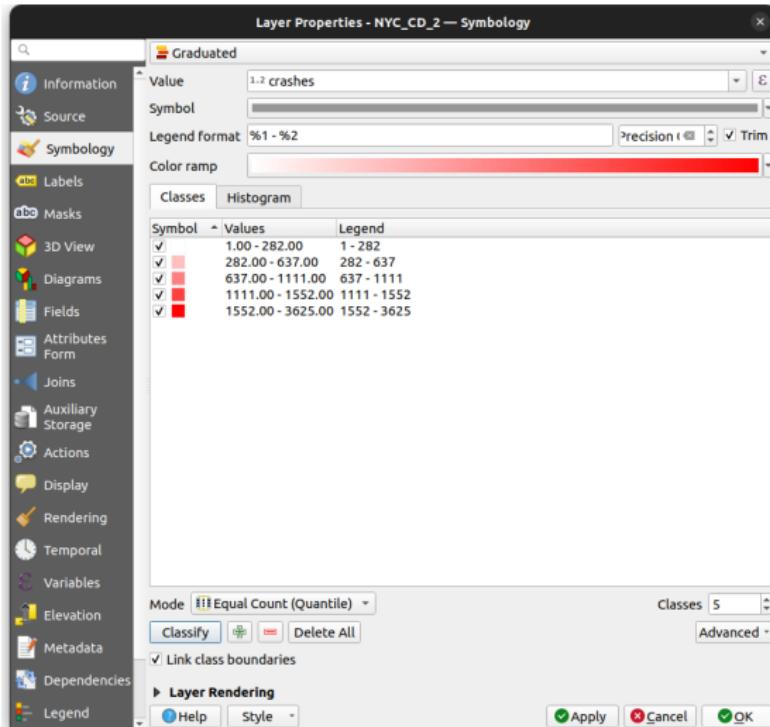
Select Polygons = NYC_CD, Points = NYC_BikeCrashes. Name the count field crashes, and save the output file as NYC_CD_2.geojson. Click Run



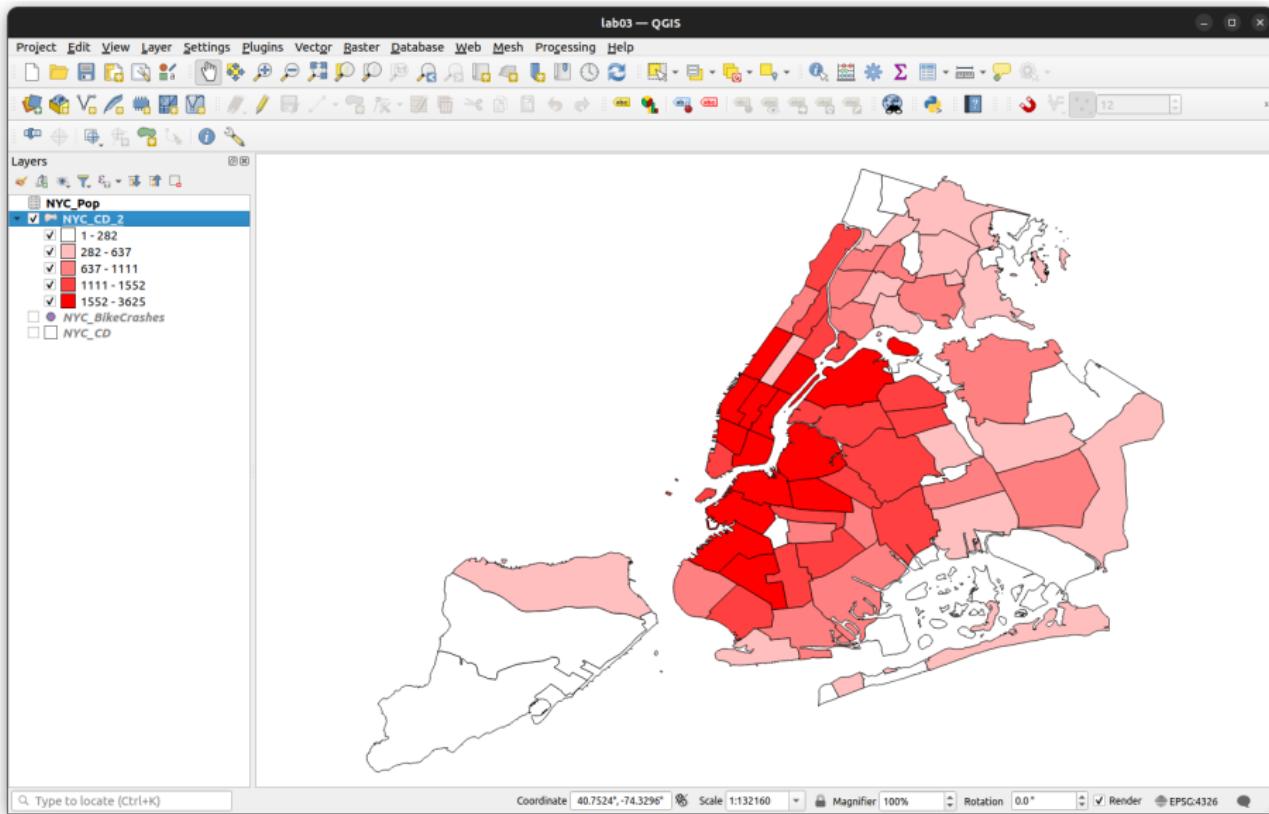
The new layer should appear on the map as `NYC_CD_2`.
Hide the other two layers by unchecking boxes next to them in the menu.



Let's color the districts by number of bike crashes. Go to the Symbology tab in the new layer's Properties. Set Symbology type = Graduated, and Value = crashes. Click Classify, then OK



In which parts of NYC are cars hitting the most cyclists?



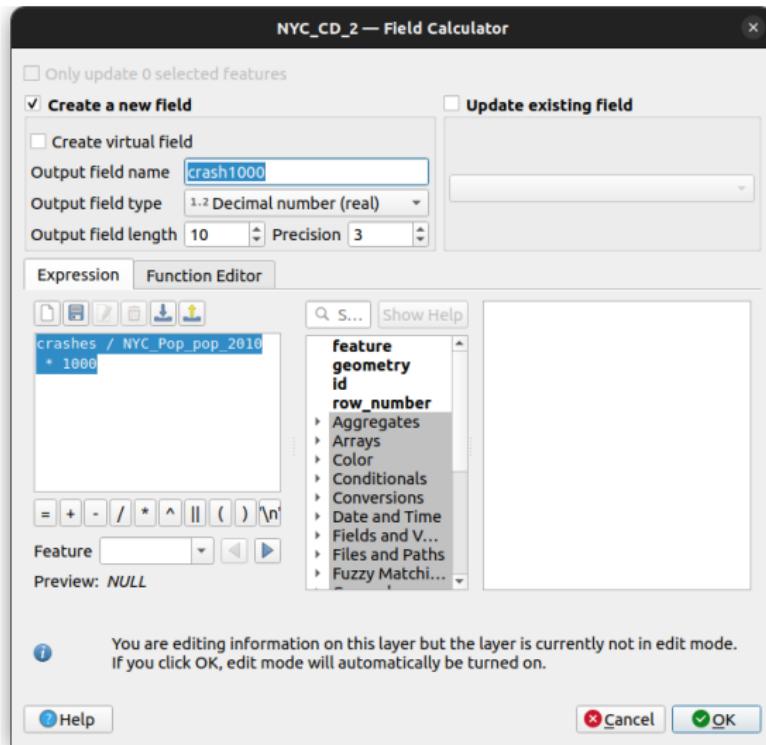
Let's combine this with population data to get a per-capita rate of bike crashes.
Open the Attribute Table for NYC_CD_2. Open the Field Calculator

NYC_CD_2 — Features Total: 71, Filtered: 71, Selected: 0

The screenshot shows the QGIS Attribute Table for the NYC_CD_2 layer. The table has 71 features. The columns include boro_cd, shape_area, shape_leng, and several population-related fields. A green box highlights the 'boro_cd' column header. The 'Field calculator' button is highlighted in yellow. The bottom of the table shows a toolbar with icons for 'Show All Features' and other GIS operations.

	boro_cd	shape_area	shape_leng	C_Pop_boro_cd	C_Pop_Boro_CD	Open field calculator (Ctrl+I)	Pop_pop_15C	C_Pop_pop_15C	C_Pop_pop_15C
1	308	45603787...	38232.886...	3 Brooklyn		8 Crown Hei...	121821	88796	96400
2	414	195576601...	210918.17...	4 Queens		14 The Rocka...	98228	100592	100596
3	204	55522140...	31358.450...	2 Bronx		4 Highbridge...	144207	114312	119962
4	228	92938101...	122141.65...	NULL NULL		NULL NULL	NULL	NULL	NULL
5	205	38316975...	29443.048...	2 Bronx		5 University ...	121807	107995	118435
6	313	88180564.15	65789.792...	3 Brooklyn		13 Coney Islan...	97750	100030	102596
7	311	103208266...	51534.144...	3 Brooklyn		11 Bensonhur...	170119	155072	149994
8	484	120185760...	227190.54...	NULL NULL		NULL NULL	NULL	NULL	NULL
9	310	111345162...	44781.826...	3 Brooklyn		10 Bay Ridge, ...	129822	118187	110612
10	410	172077374...	105822.37...	4 Queens		10 Ozone Park...	113857	105651	107768
11	164	38312378...	32721.097...	NULL NULL		NULL NULL	NULL	NULL	NULL
12	203	44796444...	33500.037...	2 Bronx		3 Morrisania,...	150636	53635	57162
13	206	42664319...	35875.709...	2 Bronx		6 East Tremo...	114137	65016	68061
14	226	50565798...	32820.482...	NULL NULL		NULL NULL	NULL	NULL	NULL
15	480	32777561...	47338.739...	NULL NULL		NULL NULL	NULL	NULL	NULL
16									
17									
18									
19									
20									
21									
22									
23									
24									
25									
26									
27									
28									
29									
30									
31									
32									
33									
34									
35									
36									
37									
38									
39									
40									
41									
42									
43									
44									
45									
46									
47									
48									
49									
50									
51									
52									
53									
54									
55									
56									
57									
58									
59									
60									
61									
62									
63									
64									
65									
66									
67									
68									
69									
70									
71									

Create new field called crash1000 of type Decimal number (real).
For the Expression, write crashes / NYC_Pop_pop_P2010 * 1000. Click OK



You should now see the new variable in the attribute table.

NYC_CD_2 — Features Total: 71, Filtered: 71, Selected: 0

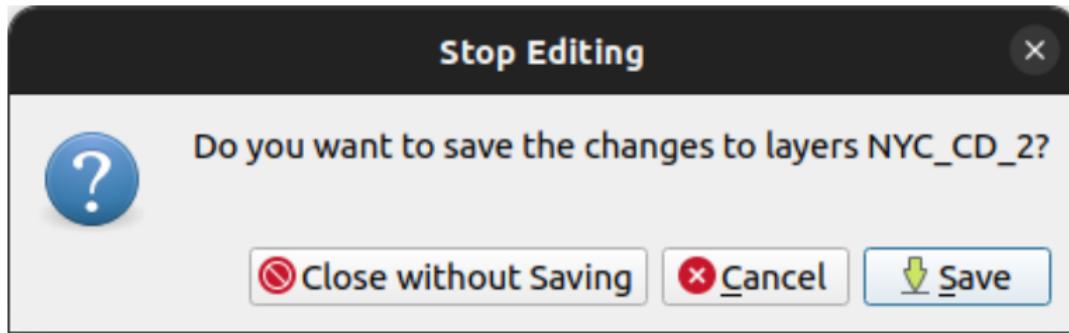
	boro_cd	shape_area	shape_len	C_Pop_Boro	C_Pop_cd	C_Pop_cd_nar	C_Pop_pop_1	C_Pop_pop_15	C_Pop_pop_15C	C_Pop_pop_2	crashes	crash1000
1	308	45603787...	38232.886...	3 Brooklyn	8 Crown Hel...	121821	88796	96400	96076	96317	1274	13.227
2	414	195576601...	210918.17...	4 Queens	14 The Rocka...	98228	100592	100596	106686	114978	400	3.479
3	204	55522140...	31358.450...	2 Bronx	4 Highbridge...	144207	114312	119962	139563	146441	820	5.600
4	228	92938101...	122141.65...	NULL NULL	NULL NULL	NULL	NULL	NULL	NULL	NULL	21	NULL
5	205	38316975...	29443.048...	2 Bronx	5 University ...	121807	107995	118435	128313	128200	724	5.647
6	313	88180564.15	65789.792...	3 Brooklyn	13 Coney Islan...	97750	100030	102596	106120	104278	527	5.054
7	311	103208266...	51534.144...	3 Brooklyn	11 Bensonhur...	170119	155072	149994	172129	181981	1117	6.138
8	484	120185760...	227190.54...	NULL NULL	NULL NULL	NULL	NULL	NULL	NULL	NULL	5	NULL
9	310	111345162...	44781.826...	3 Brooklyn	10 Bay Ridge, ...	129822	118187	110612	122542	124491	803	6.450
10	410	172077374...	105822.37...	4 Queens	10 Ozone Park...	113857	105651	107768	127274	122396	517	4.224
11	164	38312378...	32721.097...	NULL NULL	NULL NULL	NULL	NULL	NULL	NULL	NULL	404	NULL
12	203	44796444...	33500.037...	2 Bronx	3 Morrisania,...	150636	53635	57162	68574	79762	572	7.171
13	206	42664319...	35875.709...	2 Bronx	6 East Tremo...	114137	65016	68061	75688	83268	608	7.302
14	226	50565798...	32820.482...	NULL NULL	NULL NULL	NULL	NULL	NULL	NULL	NULL	12	NULL
15	480	32777561...	47338.739...	NULL NULL	NULL NULL	NULL	NULL	NULL	NULL	NULL	2	NULL

Click the Edit button to save the dataset

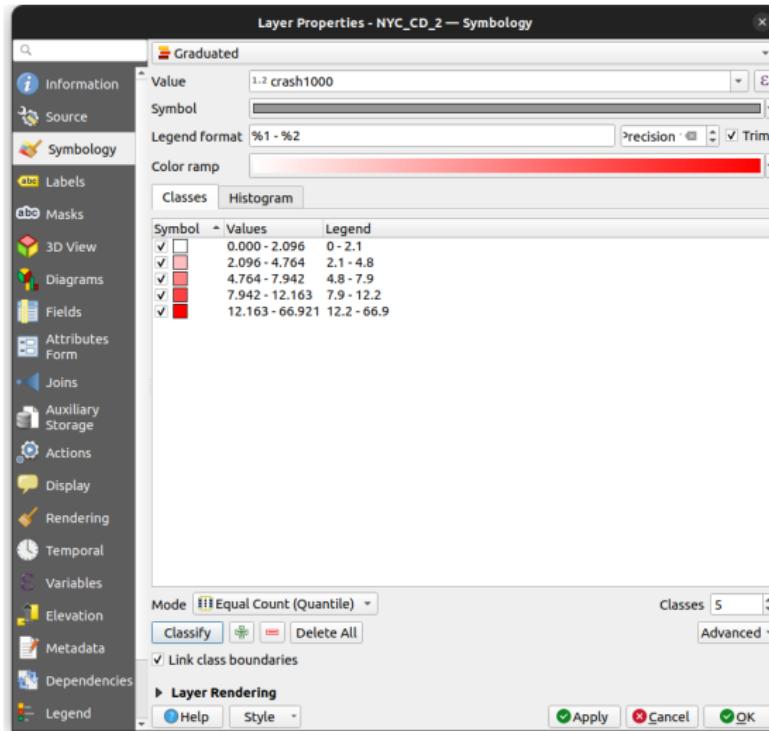
The screenshot shows the QGIS application window with a title bar 'NYC_CD_2 — Features Total: 71, Filtered: 71, Selected: 0'. Below the title bar is a toolbar with various icons. The main area contains a table with 15 rows and multiple columns. The columns include: boro_cd, shape_area, shape_leng, _Pop_boro_cd, C_Pop_Borou, C_Pop_cd, co, C_Pop_cd_nar, C_Pop_pop_15C, Pop_pop_15C, Pop_pop_15C, Pop_pop_20C, C_Pop_pop_20C, crashes, and crash1000. The first few rows show data for Brooklyn, Queens, Bronx, and other boroughs. The last row shows data for Bronx. At the bottom left of the table area is a button labeled 'Show All Features'.

	boro_cd	shape_area	shape_leng	_Pop_boro_cd	C_Pop_Borou	C_Pop_cd	co	C_Pop_cd_nar	C_Pop_pop_15C	Pop_pop_15C	Pop_pop_15C	Pop_pop_20C	C_Pop_pop_20C	crashes	crash1000
1	308	45603787...	38232.886...	3	Brooklyn	8	Crown Hei...	121821	88796	96400	96076	96317	1274	13.227	
2	414	195576601...	210918.17...	4	Queens	14	The Rocka...	98228	100592	100596	106686	114978	400	3.479	
3	204	55522140...	31358.450...	2	Bronx	4	Highbridge...	144207	114312	119962	139563	146441	820	5.600	
4	228	92938101...	122141.65...	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	21	NULL
5	205	38316975...	29443.048...	2	Bronx	5	University ...	121807	107995	118435	128313	128200	724	5.647	
6	313	88180564.15	65789.792...	3	Brooklyn	13	Coney Islan...	97750	100030	102596	106120	104278	527	5.054	
7	311	103208266...	51534.144...	3	Brooklyn	11	Bensonhur...	170119	155072	149994	172129	181981	1117	6.138	
8	484	120185760...	227190.54...	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	5	NULL
9	310	111345162...	44781.826...	3	Brooklyn	10	Bay Ridge, ...	129822	118187	110612	122542	124491	803	6.450	
10	410	172077374...	105822.37...	4	Queens	10	Ozone Park...	113857	105651	107768	127274	122396	517	4.224	
11	164	38312378...	32721.097...	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	404	NULL
12	203	44796444...	33500.037...	2	Bronx	3	Morrisania,...	150636	53635	57162	68574	79762	572	7.171	
13	206	42664319...	35875.709...	2	Bronx	6	East Tremo...	114137	65016	68061	75688	83268	608	7.302	
14	226	50565798...	32820.482...	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	12	NULL	
15	480	32777561...	47338.739...	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	2	NULL	

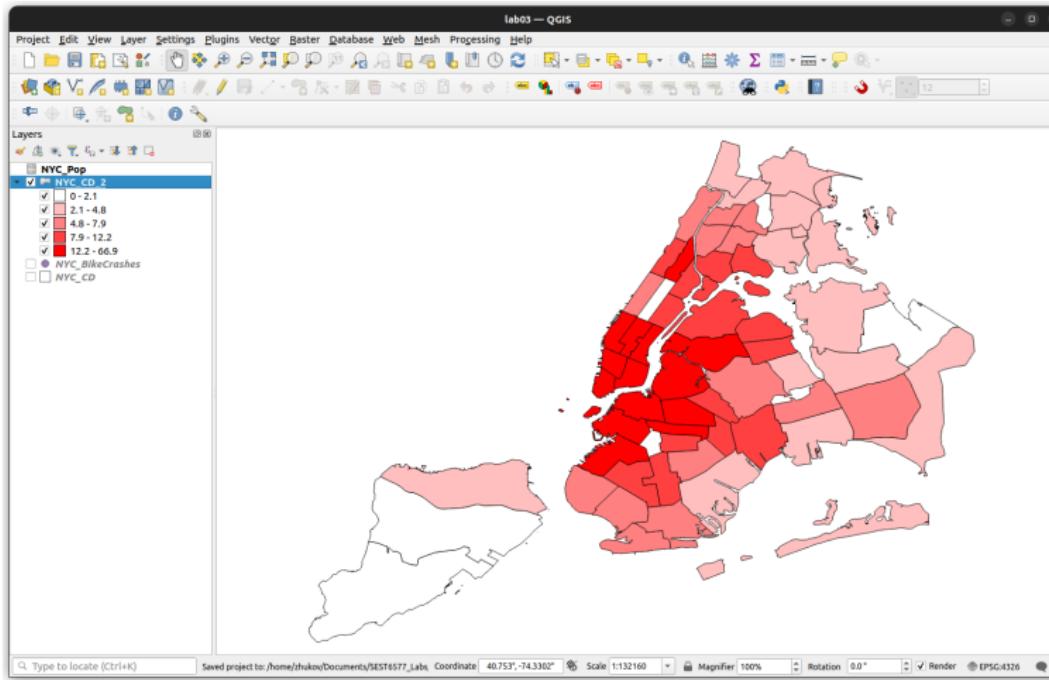
Click Save



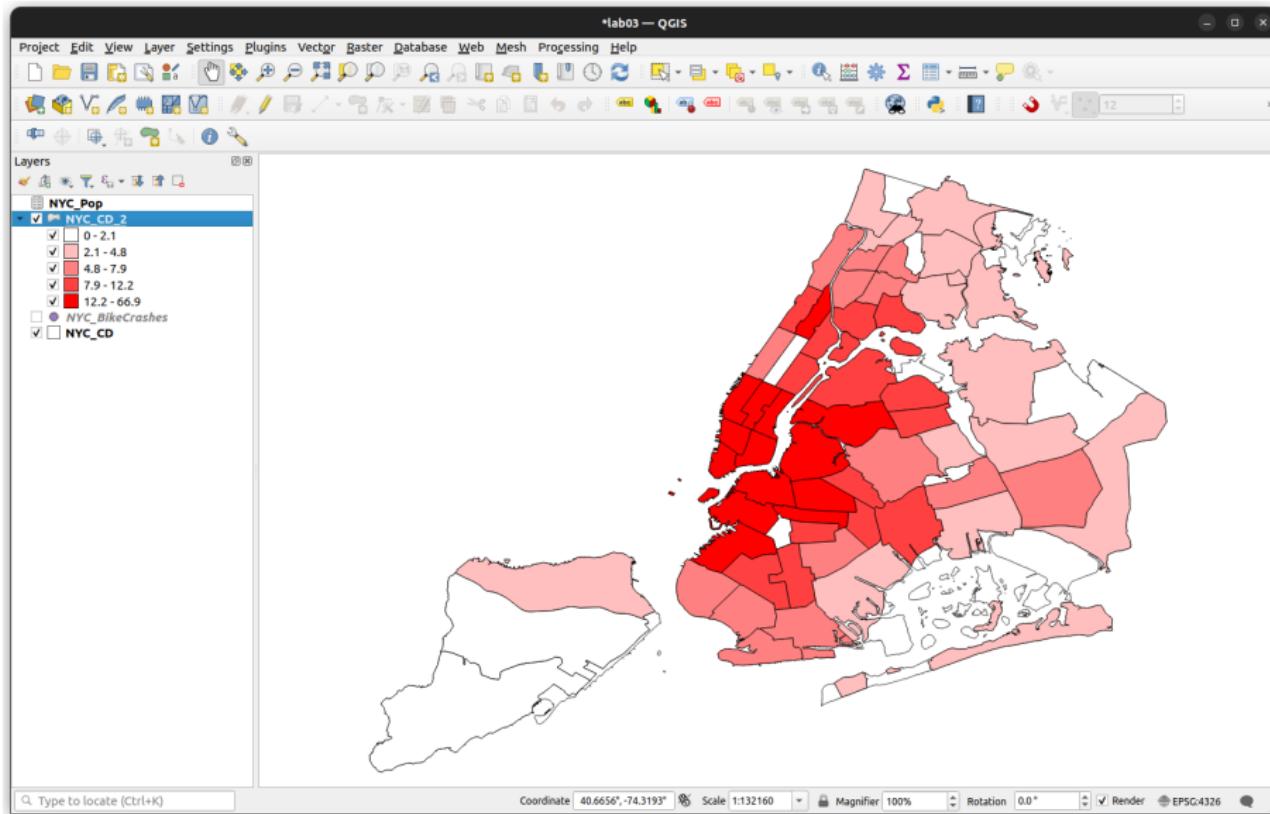
Change the symbology again, to color districts by crash1000. Remember to click Classify



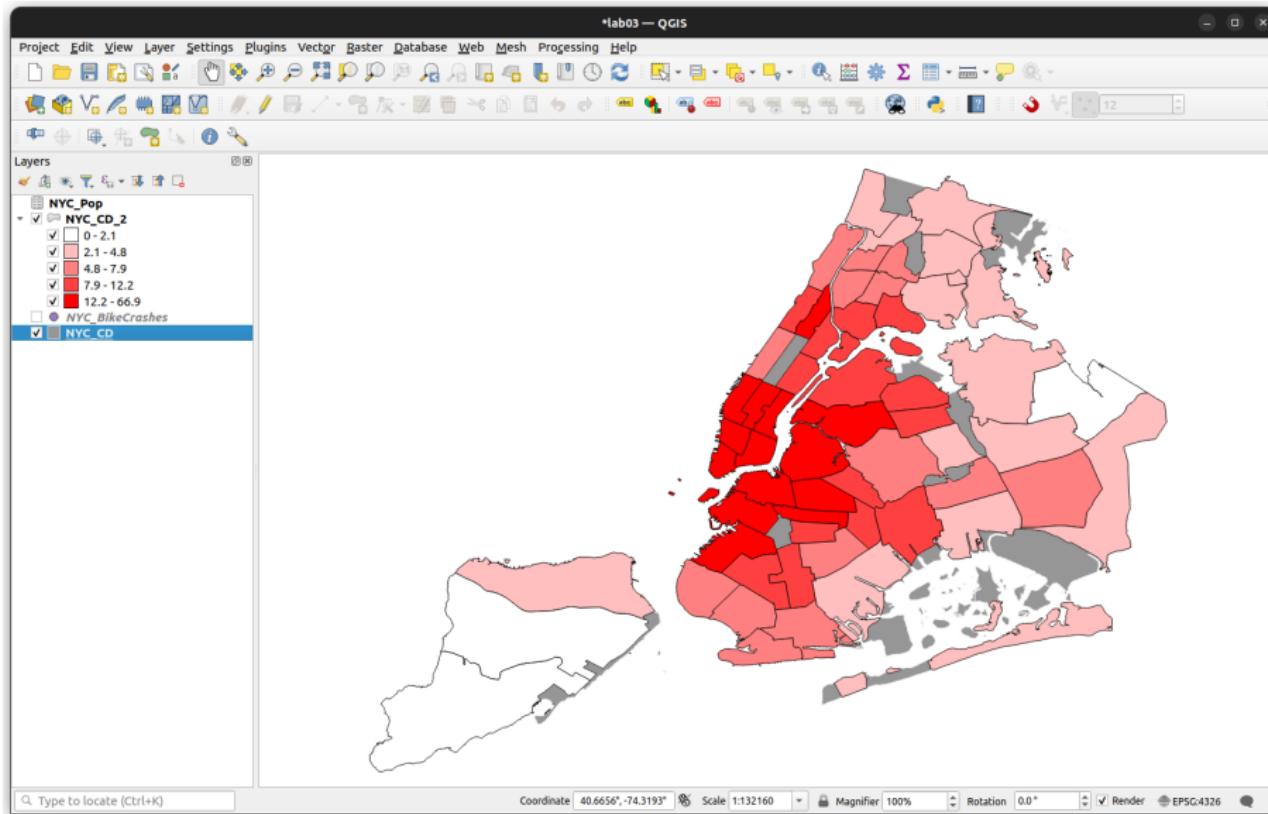
The distribution of crashes per 1000 residents (crash1000) looks similar to crashes. But there are a few missing districts (due to no population data)



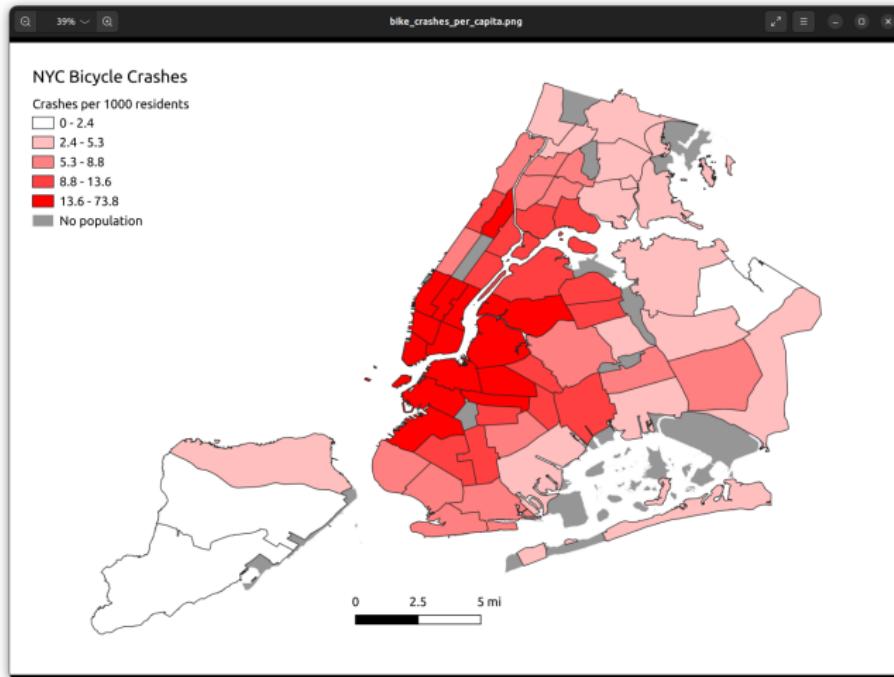
We can fix this (aesthetically) by un-hiding the NYC_CD layer.



You may want to change the color of NYC_CD to something neutral, like gray.



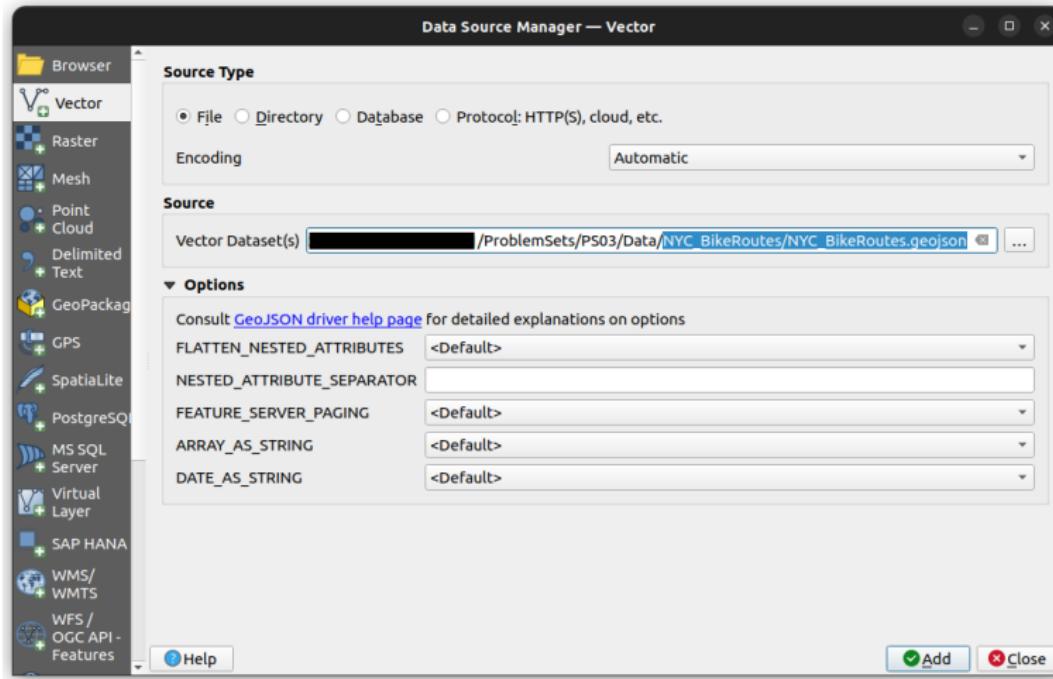
Export the map to image. Place map on a New Print Layout, add and properly format a legend, scale bar, etc. The end product should look vaguely like this.



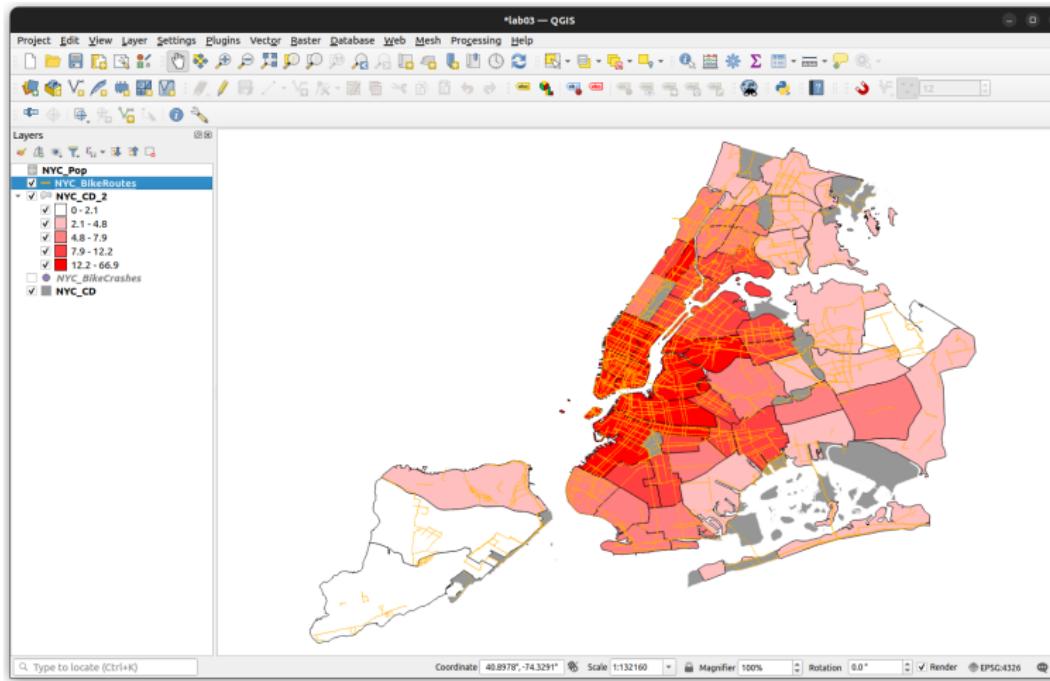
Map 2

Load bike routes: Layer → Add Layer → Add Vector Layer...

Navigate to NYC_BikeRoutes.geojson in the NYC_BikeRoutes folder. Click Add

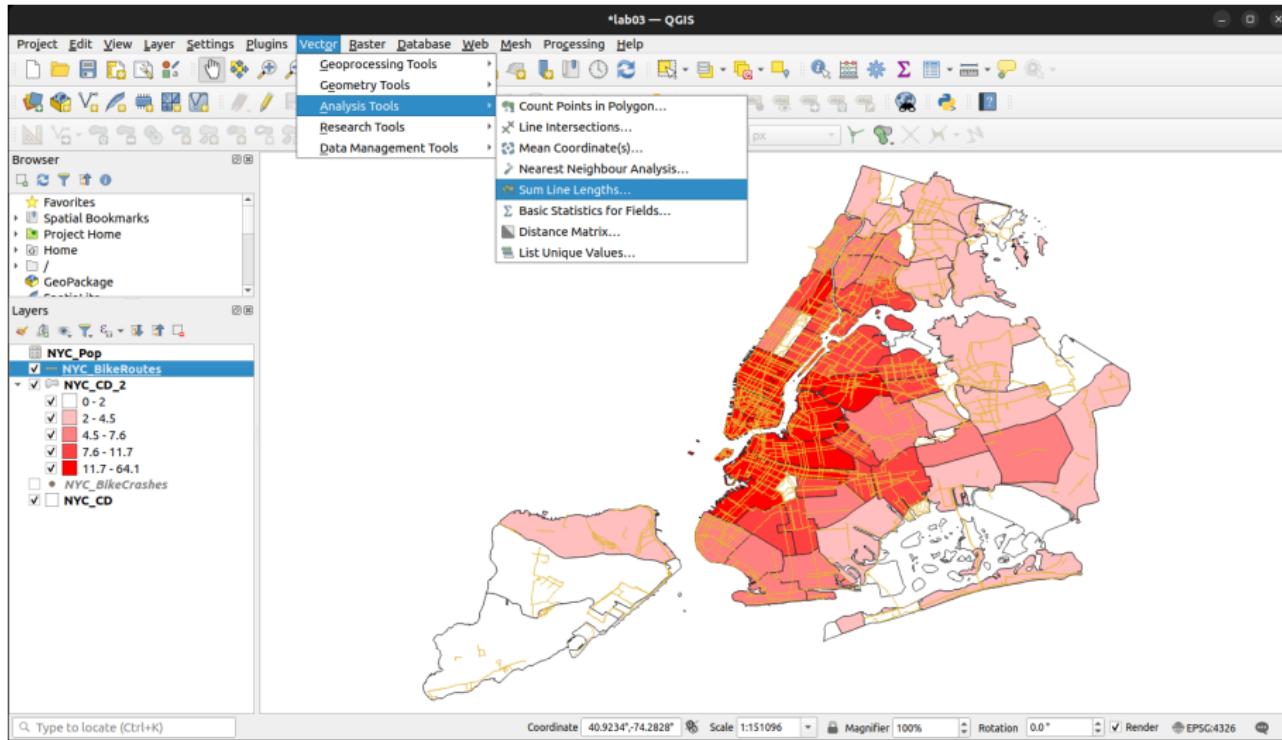


This is a polyline layer, representing NYC's cycling lanes and greenways.
Let's calculate how many miles of bike lanes are in each district.

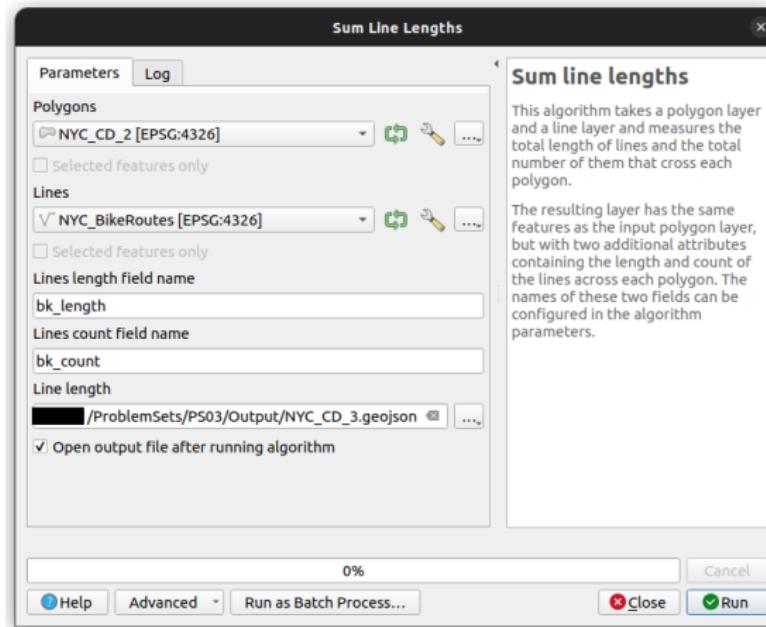


Navigate to the Sum Line Lengths tool.

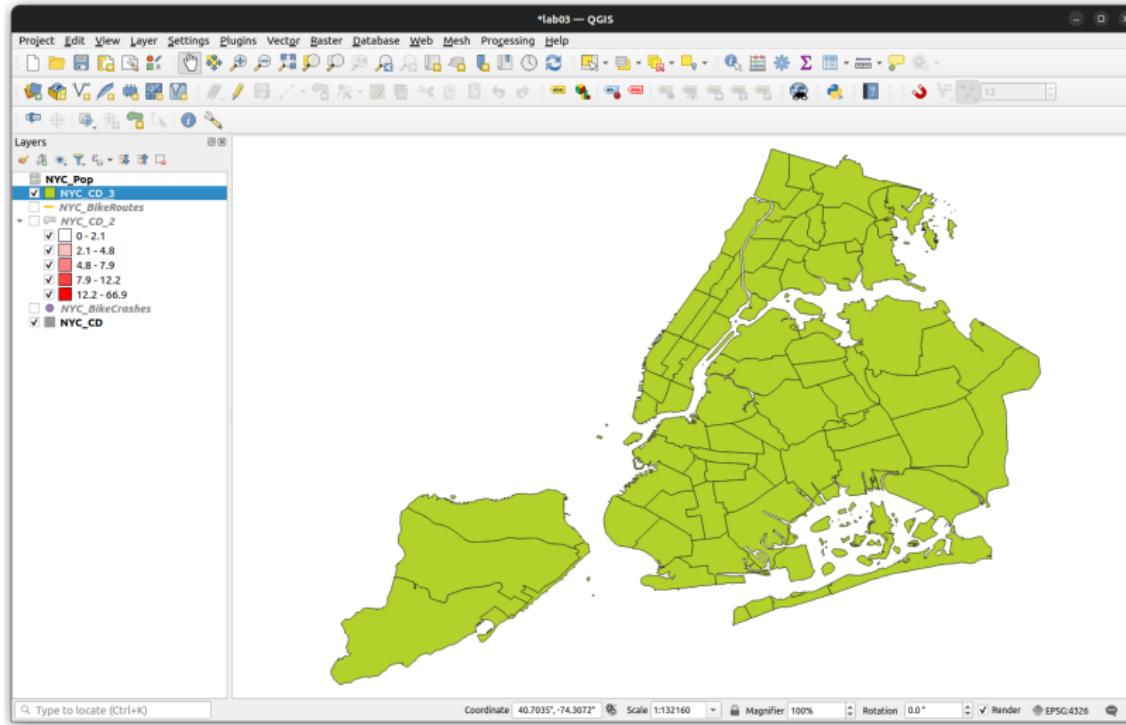
Vector menu → Analysis Tools → Sum Line Length



Select Lines = NYC_BikeRoutes, Polygons = NYC_CD_2.
Name the lengths and count fields bk_length and bk_count.
Save the output file as NYC_CD_3.geojson. Click Run



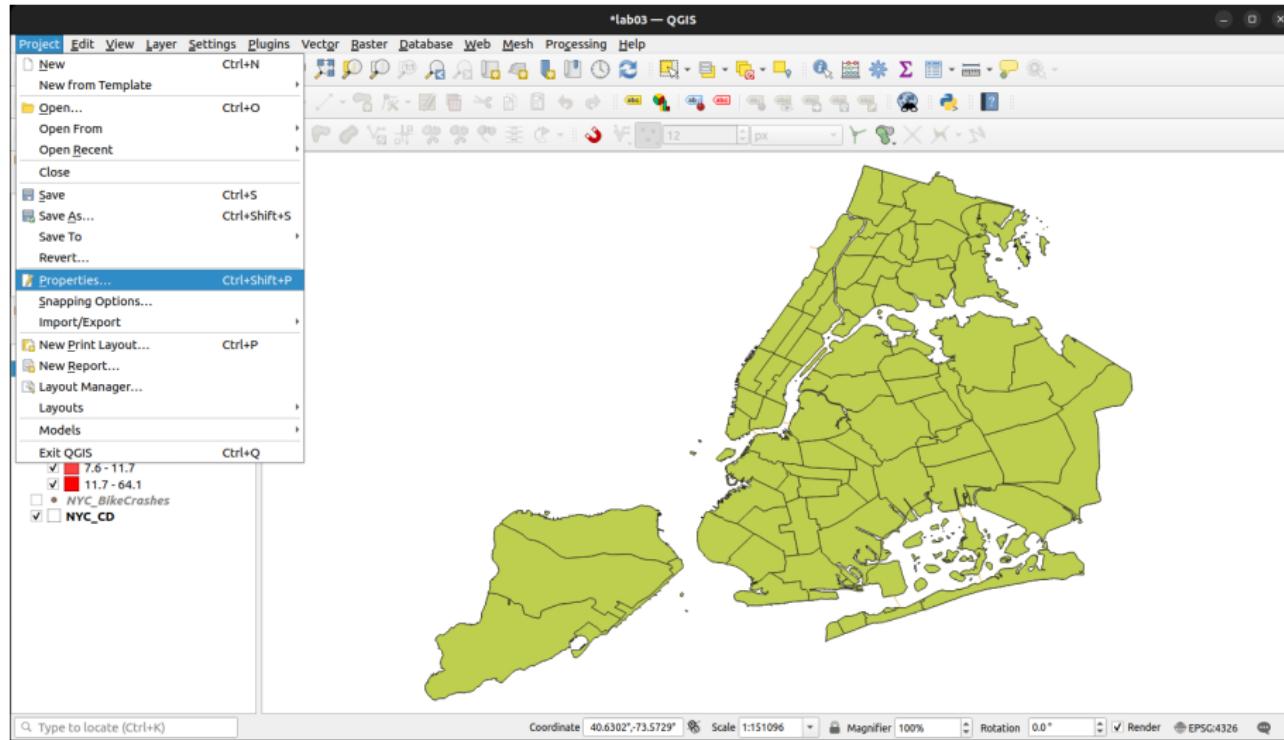
The new layer should appear as NYC_CD_3 on the map. Let's take a look



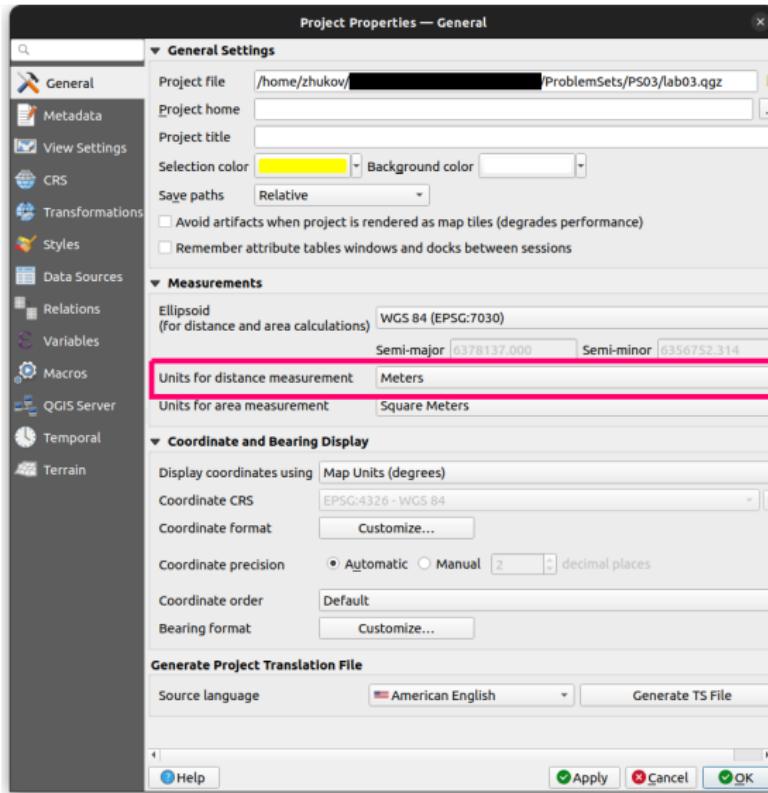
Open the Attribute Table for NYC_CD_3. The bk_length variable is there, but what are its units of measurement?

NYC_CD_3 — Features Total: 71, Filtered: 71, Selected: 0															
	shape_leng	Pop_boro_c/C	Pop_Borou/C	Pop_cd_c/C	Pop_cd_coi/C	Pop_cd_narC	Pop_pop_1%C	Pop_pop_15C	Pop_pop_15C	Pop_pop_2%C	Pop_pop_2C	crashes	crash1000	bk_length	bk_count
1	38232.886...	3 Brooklyn		8 Crown Hei...	121821	88796	96400	96076	96317	1274	13.227	25999.853...		280	
2	210918.17...	4 Queens		14 The Rocka...	98228	100592	100596	106686	114978	400	3.479	34645.022...		535	
3	31358.450...	2 Bronx		4 Highbridge...	144207	114312	119962	139563	146441	820	5.6	20138.531...		402	
4	122141.65...	NULL NULL		NULL NULL	NULL	NULL	NULL	NULL	NULL	21	NULL	14172.794...		107	
5	29443.048...	2 Bronx		5 University ...	121807	107995	118435	128313	128200	724	5.647	15096.169...		320	
6	65789.792...	3 Brooklyn		13 Coney Islan...	97750	100030	102596	106120	104278	527	5.054	18317.398...		313	
7	51534.144...	3 Brooklyn		11 Bensonhur...	170119	155072	149994	172129	181981	1117	6.138	11343.990...		134	
8	227190.54...	NULL NULL		NULL NULL	NULL	NULL	NULL	NULL	NULL	5	NULL	19093.918...		129	
9	44781.826...	3 Brooklyn		10 Bay Ridge, ...	129822	118187	110612	122542	124491	803	6.45	45523.898...		502	
10	105822.37...	4 Queens		10 Ozone Park...	113857	105651	107768	127274	122396	517	4.224	11655.278...		107	
11	32721.097...	NULL NULL		NULL NULL	NULL	NULL	NULL	NULL	NULL	404	NULL	16844.883...		280	
12	33500.037...	2 Bronx		3 Morrisania,...	150636	53635	57162	68574	79762	572	7.171	17305.319...		275	
13	35875.709...	2 Bronx		6 East Tremo...	114137	65016	68061	75688	83268	608	7.302	14486.147...		227	
14	32820.482...	NULL NULL		NULL NULL	NULL	NULL	NULL	NULL	NULL	12	NULL	6145.2135...		79	
15	47338.739...	NULL NULL		NULL NULL	NULL	NULL	NULL	NULL	NULL	2	NULL	51.460121...		3	

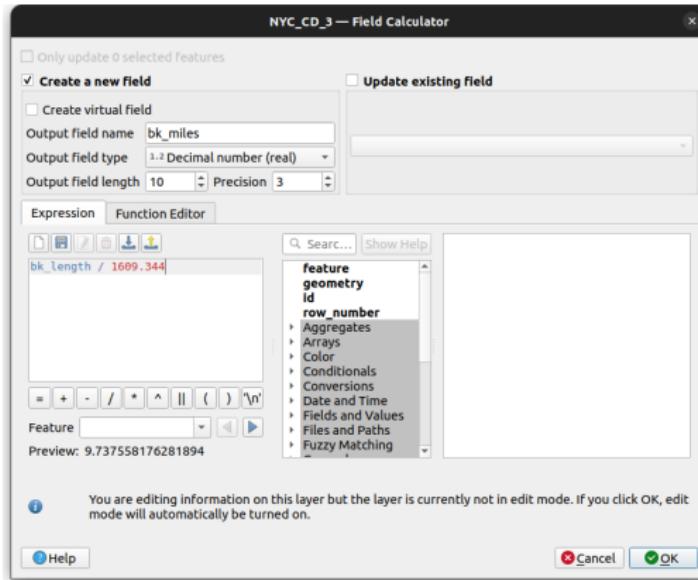
We can look up the project's units of measurement by going to Project menu → Properties



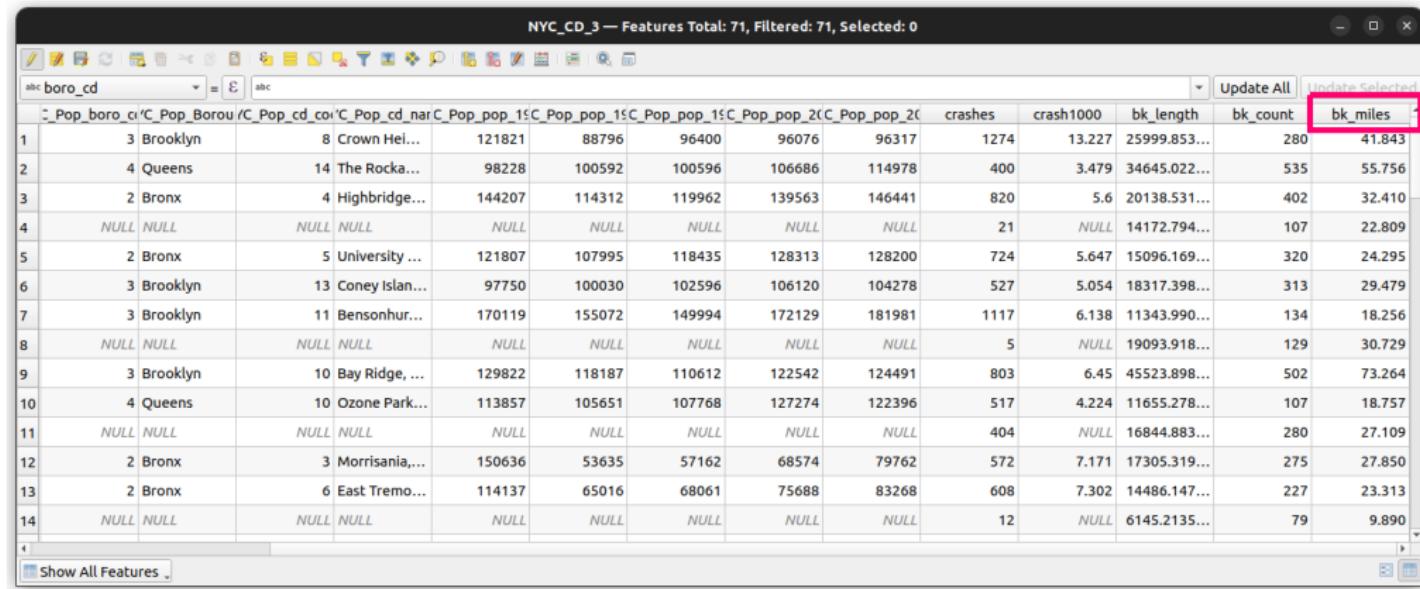
In Project Properties, we see Units for distance measurement = Meters



Let's convert bk_length from meters to miles. Go back to NYC_CD_3's Attribute Table → Field Calculator. Create new field called bk_miles of type Decimal number (real). For the Expression, write bk_length / 1609.344. Click OK



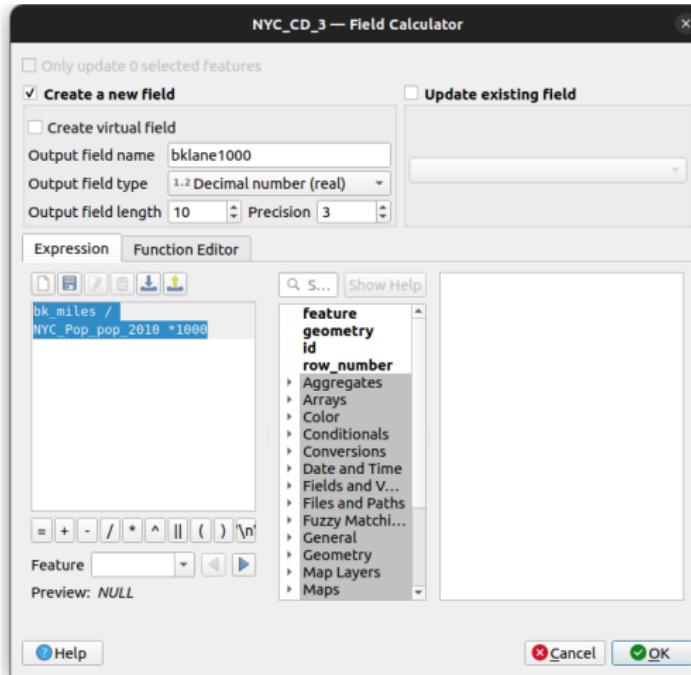
You should now see the new variable in the attribute table.



The screenshot shows the QGIS attribute table for the 'NYC_CD_3' layer. The table has 14 rows and 14 columns. The columns are labeled as follows: boro_cd, Pop_boro_cd, C_Pop_Boro, C_Pop_cd, coi_C_Pop_cd, nar_C_Pop_pop_1, C_Pop_pop_1, C_Pop_pop_1, C_Pop_pop_1, C_Pop_pop_2, C_Pop_pop_2, crashes, crash1000, bk_length, bk_count, and bk_miles. The 'bk_miles' column is highlighted with a red box. The data includes various borough names like Brooklyn, Queens, Bronx, and their corresponding statistics. For example, Brooklyn has 121821 population, 88796 households, and 1274 crashes, while Bronx has 170119 population, 155072 households, and 1117 crashes.

	boro_cd	Pop_boro_cd	C_Pop_Boro	C_Pop_cd	coi_C_Pop_cd	nar_C_Pop_pop_1	C_Pop_pop_1	C_Pop_pop_1	C_Pop_pop_1	C_Pop_pop_2	C_Pop_pop_2	crashes	crash1000	bk_length	bk_count	bk_miles
1	3	Brooklyn	8	Crown Hel...	121821	88796	96400	96076	96317	1274	13.227	25999.853...	280	41.843		
2	4	Queens	14	The Rocka...	98228	100592	100596	106686	114978	400	3.479	34645.022...	535	55.756		
3	2	Bronx	4	Highbridge...	144207	114312	119962	139563	146441	820	5.6	20138.531...	402	32.410		
4	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	21	NULL	14172.794...	107	22.809		
5	2	Bronx	5	University ...	121807	107995	118435	128313	128200	724	5.647	15096.169...	320	24.295		
6	3	Brooklyn	13	Coney Islan...	97750	100030	102596	106120	104278	527	5.054	18317.398...	313	29.479		
7	3	Brooklyn	11	Bensonhur...	170119	155072	149994	172129	181981	1117	6.138	11343.990...	134	18.256		
8	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	5	NULL	19093.918...	129	30.729		
9	3	Brooklyn	10	Bay Ridge, ...	129822	118187	110612	122542	124491	803	6.45	45523.898...	502	73.264		
10	4	Queens	10	Ozone Park...	113857	105651	107768	127274	122396	517	4.224	11655.278...	107	18.757		
11	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	404	NULL	16844.883...	280	27.109		
12	2	Bronx	3	Morrisania,...	150636	53635	57162	68574	79762	572	7.171	17305.319...	275	27.850		
13	2	Bronx	6	East Tremo...	114137	65016	68061	75688	83268	608	7.302	14486.147...	227	23.313		
14	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	12	NULL	6145.2135...	79	9.890		

Now let's create a per-capita measure (miles of bike lane per 1000 residents). Go back to the Field Calculator. Create new field called bklane1000 of type Decimal number (real). For the Expression, write bk_miles / NYC_Pop_pop_2010 * 1000. Click OK



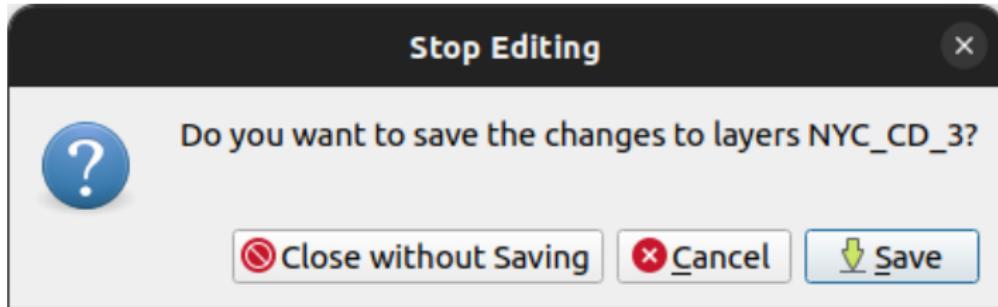
You should now see the new variable in the attribute table.

NYC_CD_3 — Features Total: 71, Filtered: 71, Selected: 0

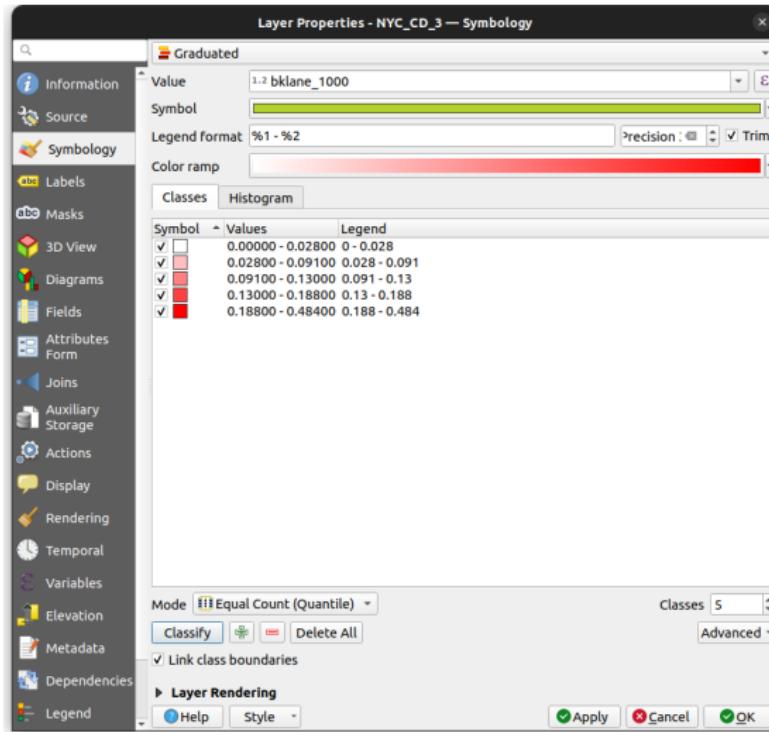
The screenshot shows the QGIS attribute table for the 'NYC_CD_3' layer. The table has 14 columns and 15 rows. The columns are: boro_cd, 'C_Pop_Borou / C_Pop_cd_cir', 'C_Pop_cd_nar', 'C_Pop_pop_1%', 'C_Pop_pop_1%', 'C_Pop_pop_1%', 'C_Pop_pop_2%', 'C_Pop_pop_2%', crashes, crash1000, bk_length, bk_count, bk_miles, and bklane1000. The 'bklane1000' column is highlighted with a red box. The data includes various boroughs like Brooklyn, Queens, Bronx, and null values, along with their respective population percentages and lane count statistics.

	boro_cd	'C_Pop_Borou / C_Pop_cd_cir'	'C_Pop_cd_nar'	'C_Pop_pop_1%'	'C_Pop_pop_1%'	'C_Pop_pop_1%'	'C_Pop_pop_2%'	'C_Pop_pop_2%'	crashes	crash1000	bk_length	bk_count	bk_miles	bklane1000
1	Brooklyn	8 Crown He...	121821	88796	96400	96076	96317	1274	13.227	25999.853...	280	41.843	0.434	
2	Queens	14 The Rocka...	98228	100592	100596	106686	114978	400	3.479	34645.022...	535	55.756	0.485	
3	Bronx	4 Highbridge...	144207	114312	119962	139563	146441	820	5.6	20138.531...	402	32.410	0.221	
4	NULL	NULL NULL	NULL	NULL	NULL	NULL	NULL	21	NULL	14172.794...	107	22.809	NULL	
5	Bronx	5 University ...	121807	107995	118435	128313	128200	724	5.647	15096.169...	320	24.295	0.19	
6	Brooklyn	13 Coney Islan...	97750	100030	102596	106120	104278	527	5.054	18317.398...	313	29.479	0.283	
7	Brooklyn	11 Bensonhur...	170119	155072	149994	172129	181981	1117	6.138	11343.990...	134	18.256	0.1	
8	NULL	NULL NULL	NULL	NULL	NULL	NULL	NULL	5	NULL	19093.918...	129	30.729	NULL	
9	Brooklyn	10 Bay Ridge, ...	129822	118187	110612	122542	124491	803	6.45	45523.898...	502	73.264	0.589	
10	Queens	10 Ozone Park...	113857	105651	107768	127274	122396	517	4.224	11655.278...	107	18.757	0.153	
11	NULL	NULL NULL	NULL	NULL	NULL	NULL	NULL	404	NULL	16844.883...	280	27.109	NULL	
12	Bronx	3 Morrisania,...	150636	53635	57162	68574	79762	572	7.171	17305.319...	275	27.850	0.349	
13	Bronx	6 East Tremo...	114137	65016	68061	75688	83268	608	7.302	14486.147...	227	23.313	0.28	
14	NULL	NULL NULL	NULL	NULL	NULL	NULL	NULL	12	NULL	6145.2135...	79	9.890	NULL	

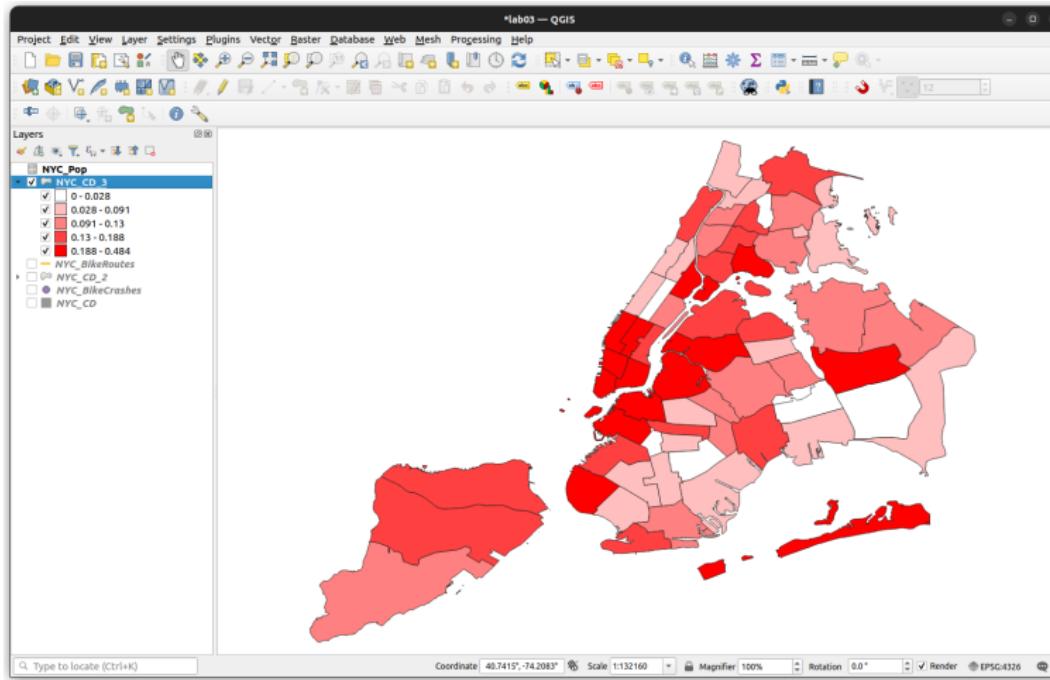
Remember to save your changes!



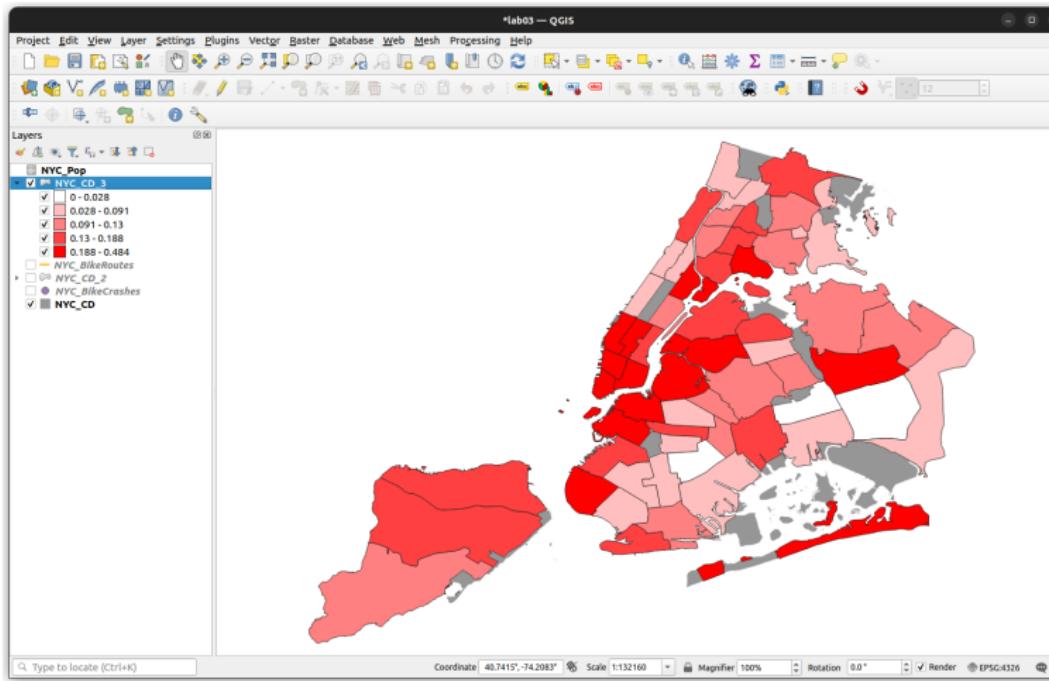
Change the symbology of NYC_CD_3, to color districts by bklane1000. Click Classify, then OK



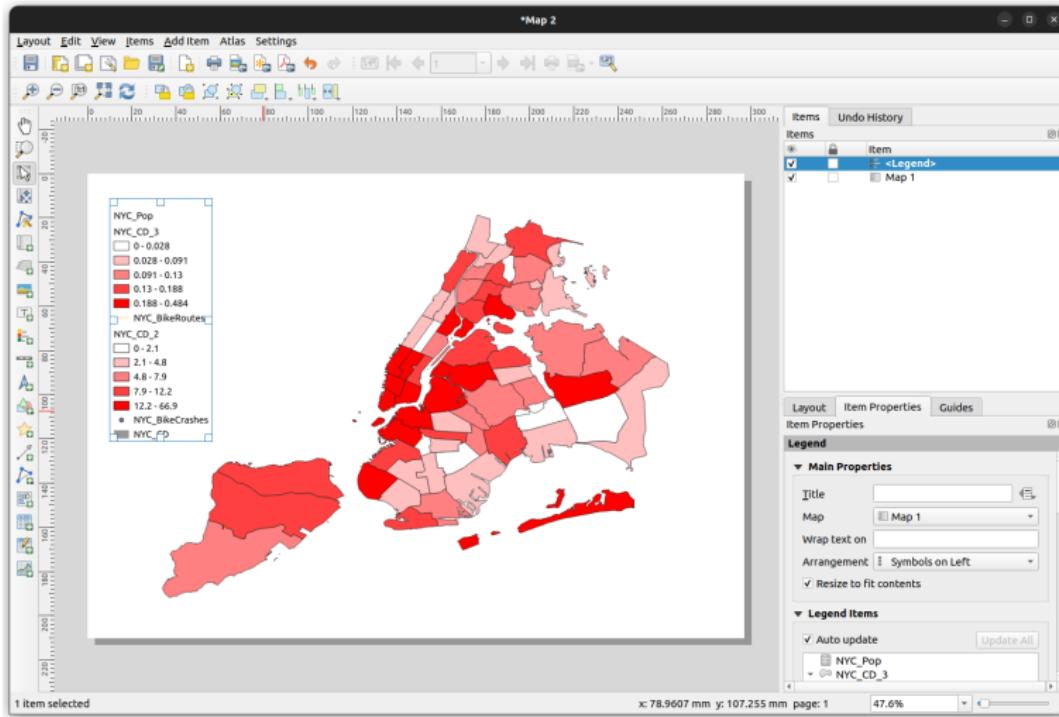
We now have a plot of “miles of bike lane per 1000 residents” (bklane1000). Again, there are a few missing districts (due to no population data)



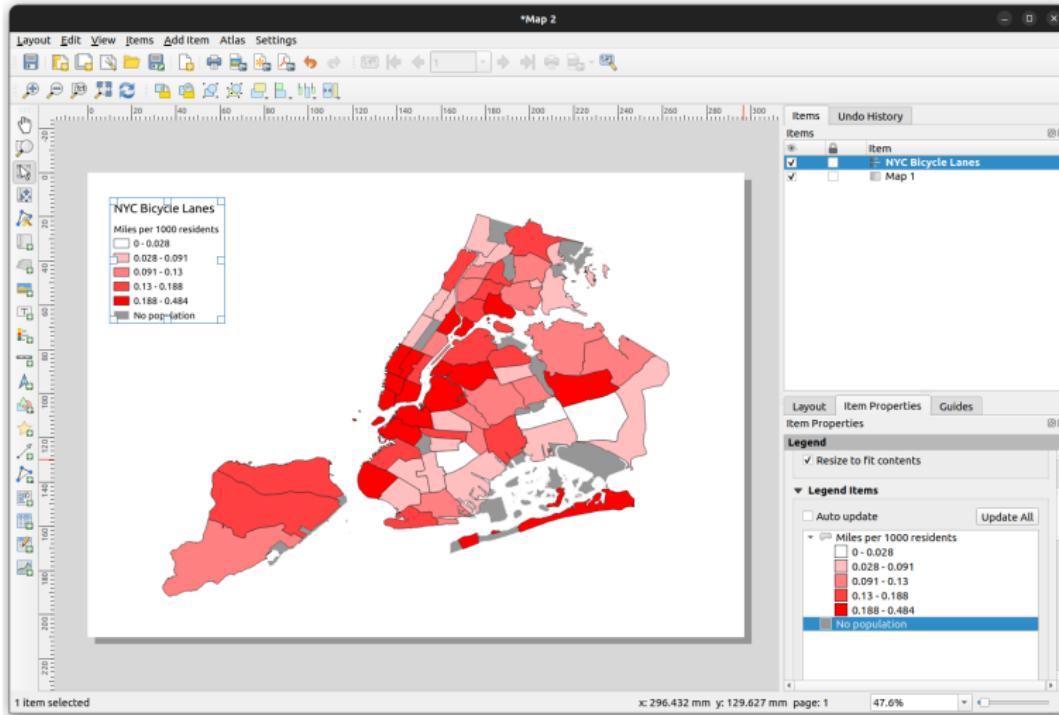
Let's un-hide the NYC_CD layer again to show the missing districts. Ideally, these should be colored gray or some other neutral color



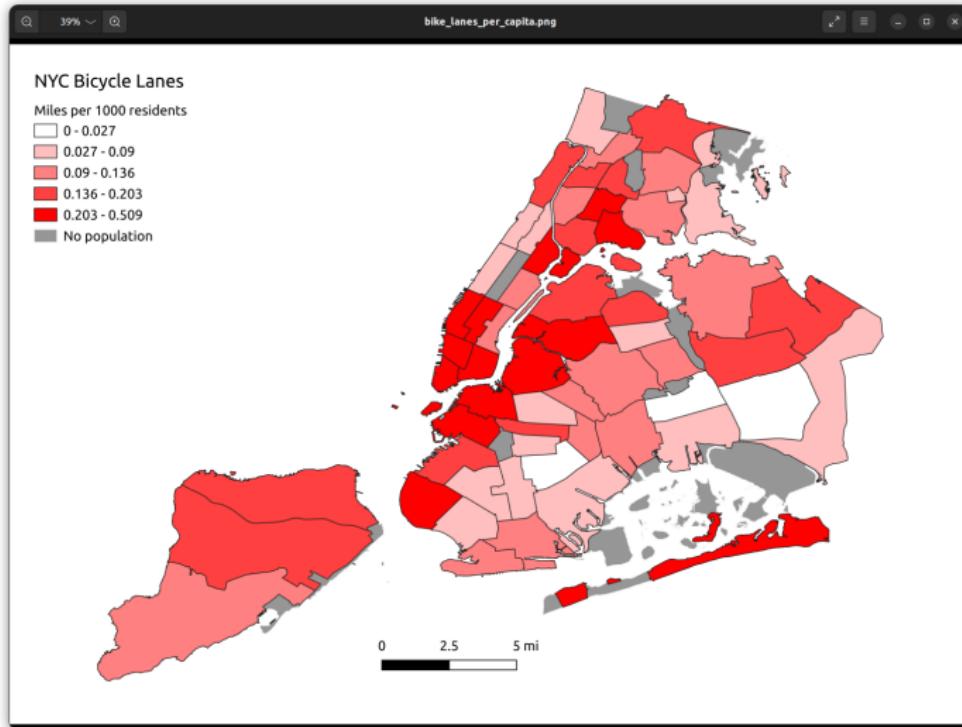
Export the map to image. Project → New Print Layout. Place map and legend to layout. The resulting legend includes several items we'll need to remove (everything but NYC_CD_3 and NYC_CD)



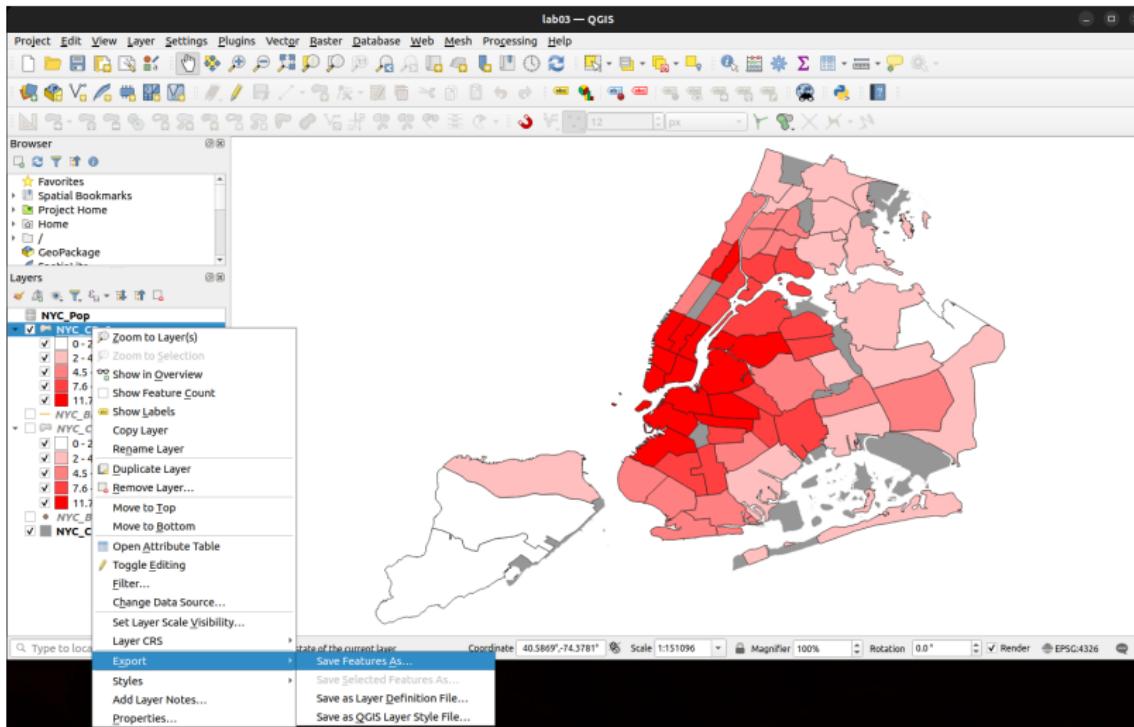
Remember how we did this in lab 2? (Item Properties → un-check Auto update, etc.) Remove everything but NYC_CD_3 and NYC_CD, change layers' names from NYC_CD_3 to “Miles per 1000 residents” and NYC_CD to “No population”.



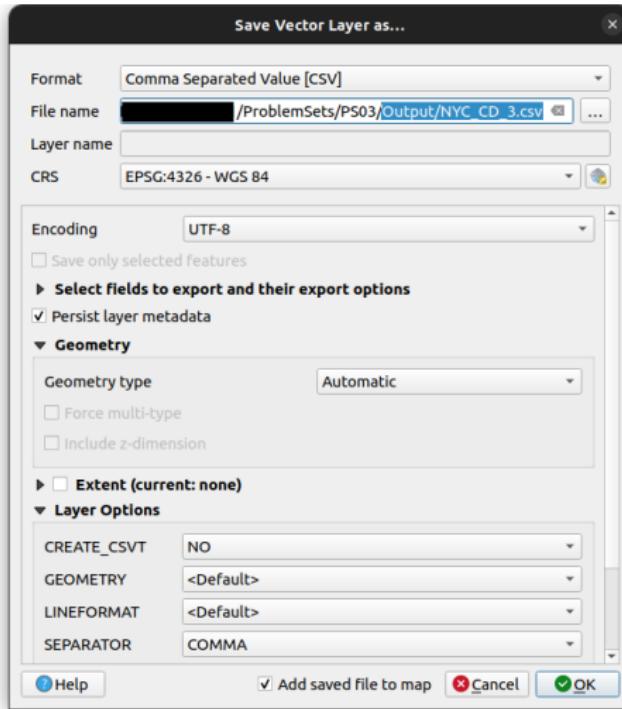
The output file should look roughly like this.



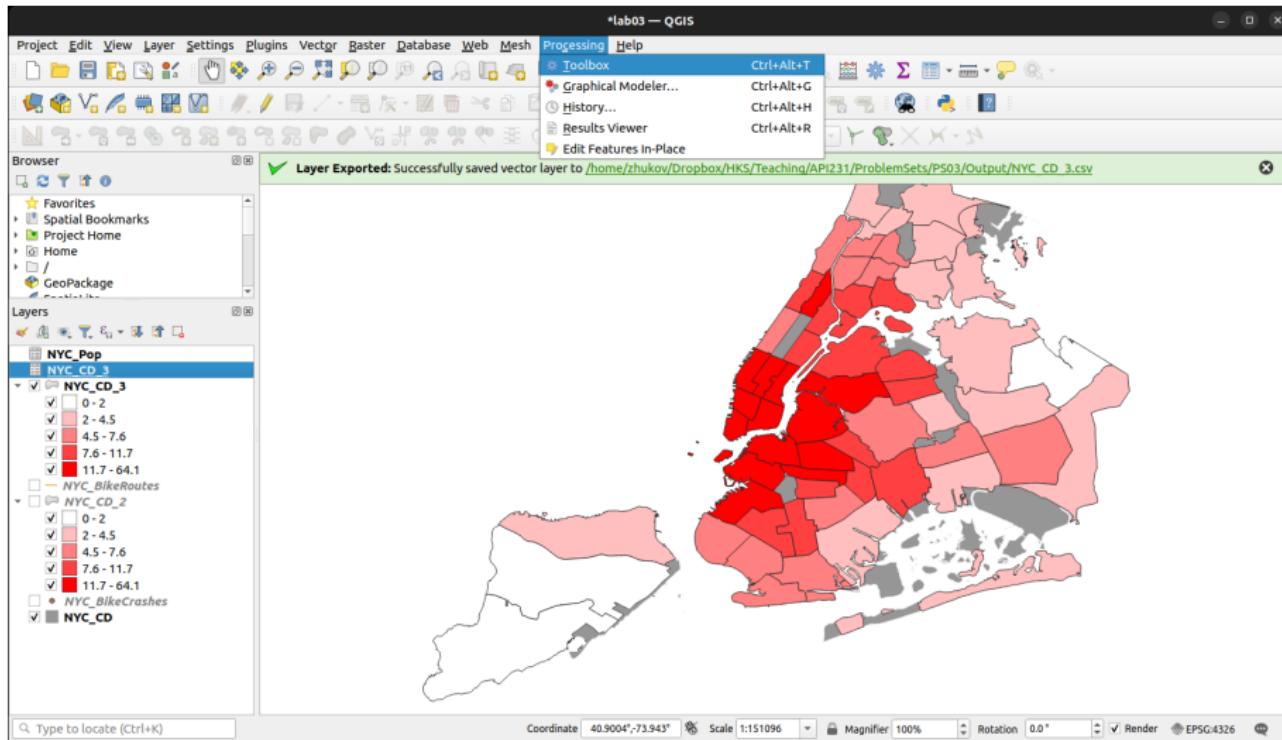
We can export the dataset we created, by right-clicking on NYC_CD_3 layer, and selecting Export → Save Features As...



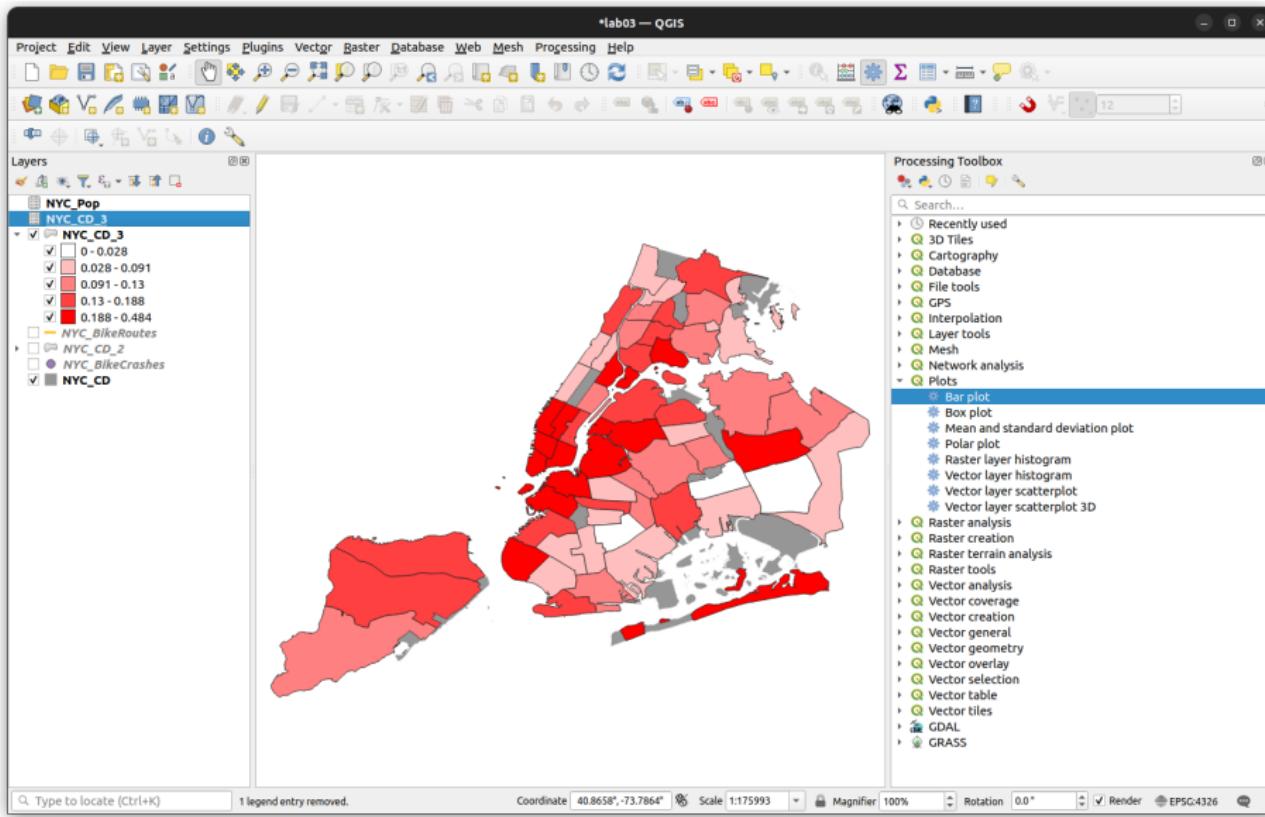
The file can be saved in a variety of formats, including .geojson, .shp, .csv and Excel



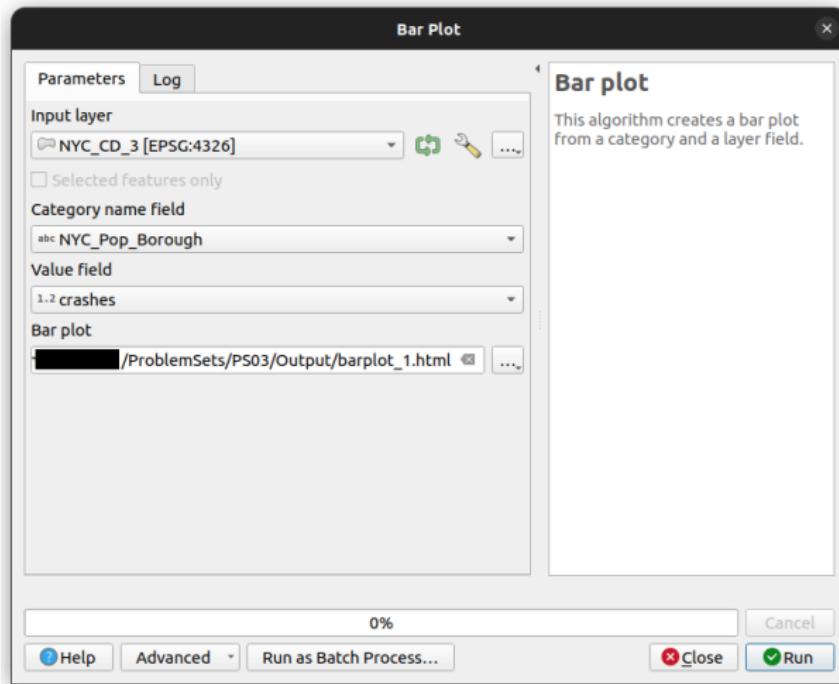
Let's create a bar plot. Go to Processing menu → Toolbox. This will open a Processing Toolbox menu on the right



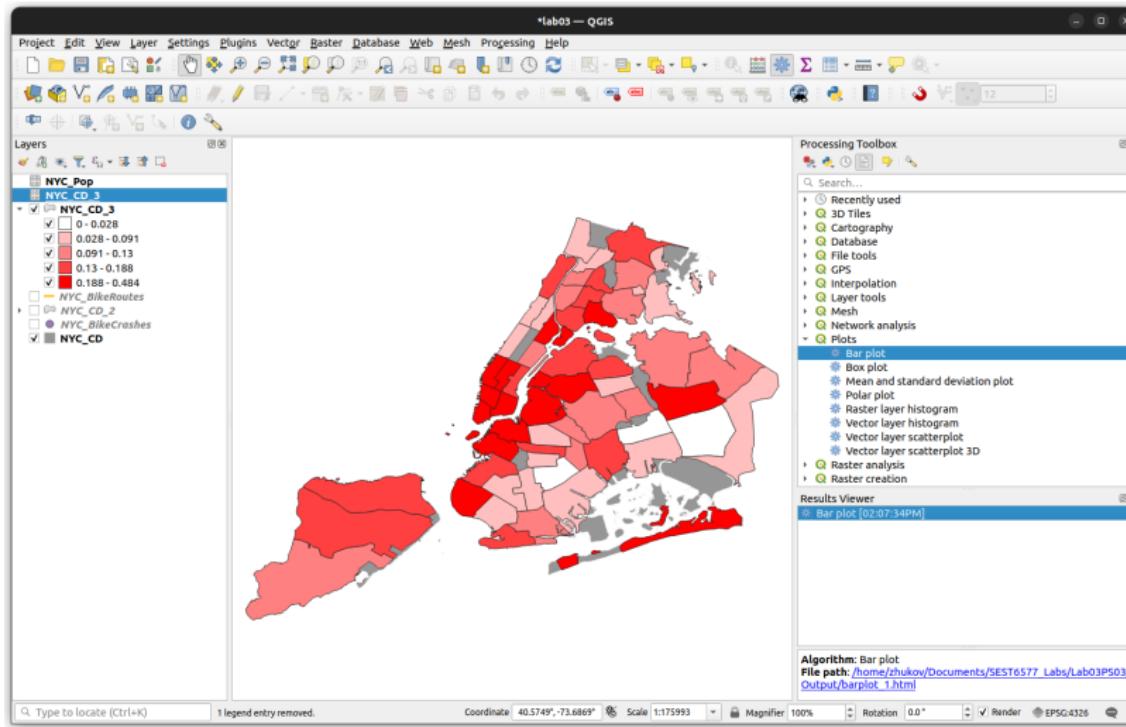
Double-click on Bar plot in the Plots submenu on the right



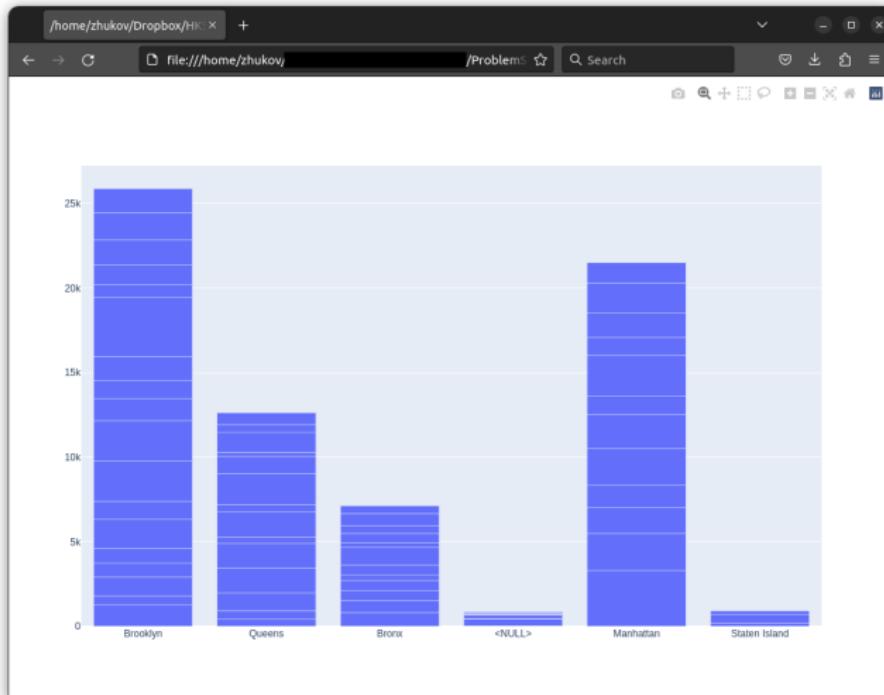
Set Input layer = NYC_CD_3, Category name field = NYC_Pop_Borough, and Value field = crashes. Set the output destination under Bar plot to barplot_1.html. Click Run



You can access the resulting plot by clicking on the Bar plot item in the Results Viewer on the right



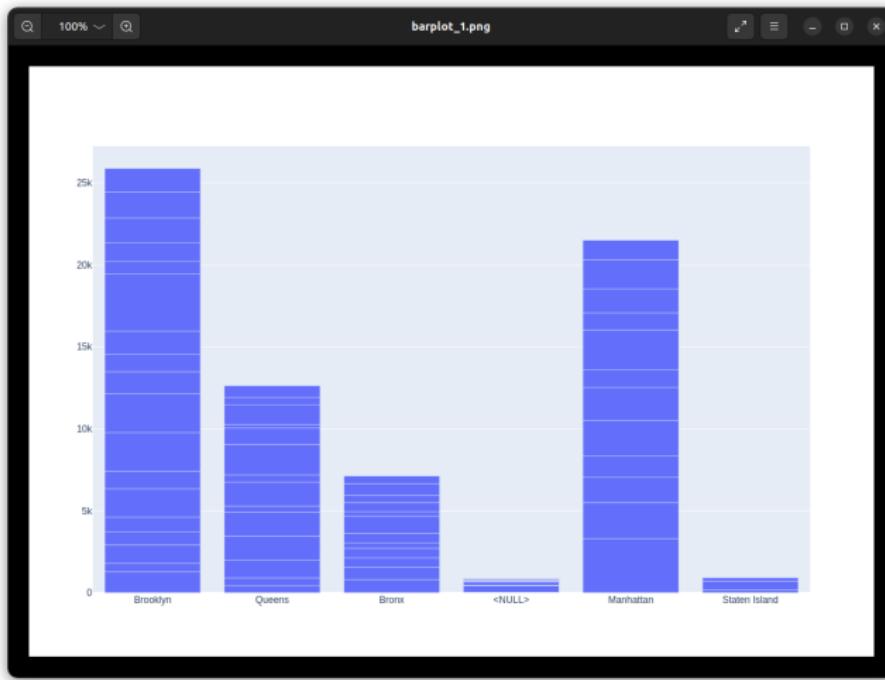
This will open a web browser window, with a bar plot of total bike crashes per district, grouped by borough. As expected, Brooklyn and Manhattan lead the pack.



Export the plot by clicking on the Download plot as png button in the top-right



The output file should look like this. For further analysis, you can also open the .csv file we just created in Excel or R, and start crunching numbers (see below)



Problem Set 3

Your assignment (if using QGIS):

- create a map of rat activity (per 1000 residents) in NYC!
- use this dataset:
 - NYC_Rats/NYC_Rats.csv
- follow the same steps as for “Map 1” above
(i.e. sum via point-in-polygon, then calculate per capita rate)
- name the file
`rats_per_capita.png`
- upload map to Canvas
(by next Monday)

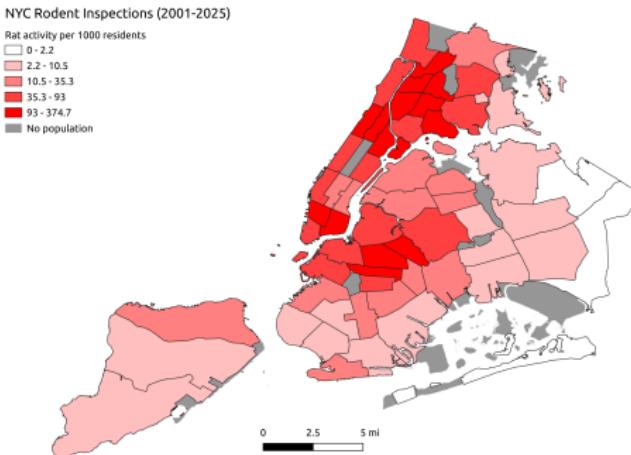


Figure 16: Can you make this map?

R

Loading R packages

To implement these steps in R, we will be using the `sf` package (again)

```
library(sf)
```

NOTE: The code to produce Map 1 and Map 2 in R is in `lab03_demo.R` on RStudio Cloud, and in `Lab03PS03.zip` (posted on Canvas).

Map 1

Loading spatial data

Let's load the *NYC community district boundaries* into R, using `sf::read_sf()`:

```
nyc_cd = sf::read_sf("Data/NYC_Communities/NYC_CD.geojson")
plot(nyc_cd["geometry"])
```



Loading non-spatial data

Load the tabular *population data* using `read.csv()`, and preview the first few rows:

```
pop = read.csv(file = "Data/NYC_Population/NYC_Pop.csv")
head(pop)

##   boro_code Borough cd_code          cd_name pop_1970
## 1          2 Bronx     1 Melrose, Mott Haven, Port Morris 138557
## 2          2 Bronx     2           Hunts Point, Longwood 99493
## 3          2 Bronx     3 Morrisania, Crotona Park East 150636
## 4          2 Bronx     4      Highbridge, Concourse Village 144207
## 5          2 Bronx     5 University Hts., Fordham, Mt. Hope 121807
## 6          2 Bronx     6           East Tremont, Belmont 114137
##   pop_1980 pop_1990 pop_2000 pop_2010
## 1    78441    77214    82159    91497
## 2    34399    39443    46824    52246
## 3    53635    57162    68574    79762
## 4   114312   119962   139563   146441
## 5   107995   118435   128313   128200
## 6    65016    68061    75688    83268
```

Joining spatial to non-spatial data

Similar to QGIS's Field Calculator, let's create a new `boro_cd` variable in `pop`:

```
pop$boro_cd = pop$boro_code*100 + pop$cd_code
```

We can join `nyc_cd` and `pop` by attribute `boro_cd` using the `merge()` command:

```
nyc_cd_2 = merge(x = nyc_cd, y = pop, by = "boro_cd")
print(nyc_cd_2, n=2) # Preview first 2 rows
```

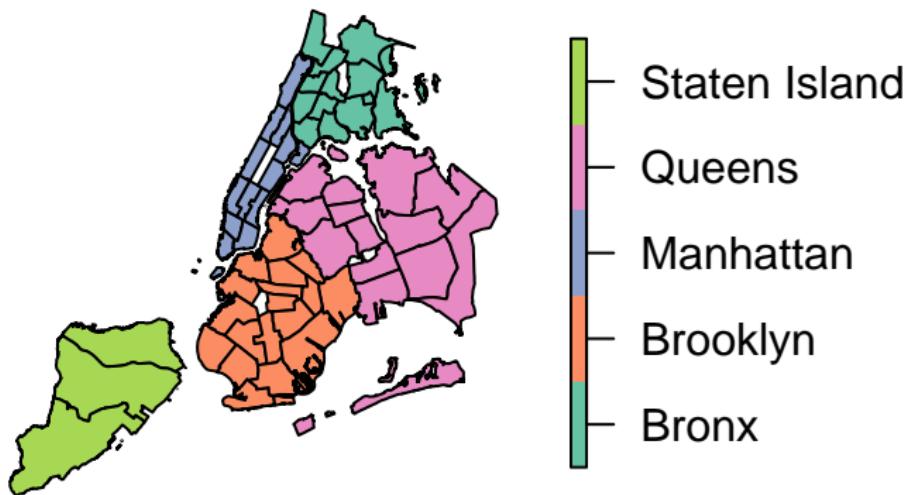
```
## Simple feature collection with 59 features and 12 fields
## Geometry type: MULTIPOLYGON
## Dimension: XY
## Bounding box: xmin: -74.25559 ymin: 40.49613 xmax: -73.70001 ymax: 40.91553
## Geodetic CRS: WGS 84
## First 2 features:
##   boro_cd     shape_area    shape_leng boro_code Borough cd_code
## 1      101 41692711.6252 69093.8528032          1 Manhattan      1
## 2      102 37604933.3646 32628.8675927          1 Manhattan      2
##           cd_name pop_1970 pop_1980 pop_1990 pop_2000 pop_2010
## 1 Battery Park City, Tribeca    7706    15918    25366    34420    60978
## 2 Greenwich Village, Soho    84337    87069    94105    93119    90016
##           geometry
## 1 MULTIPOLYGON (((-74.04388 4...
## 2 MULTIPOLYGON (((-73.99684 4...
```

Plotting merged features

Using the merged dataset `nyc_cd_2`, let's plot the districts by borough

```
plot(nyc_cd_2["Borough"])
```

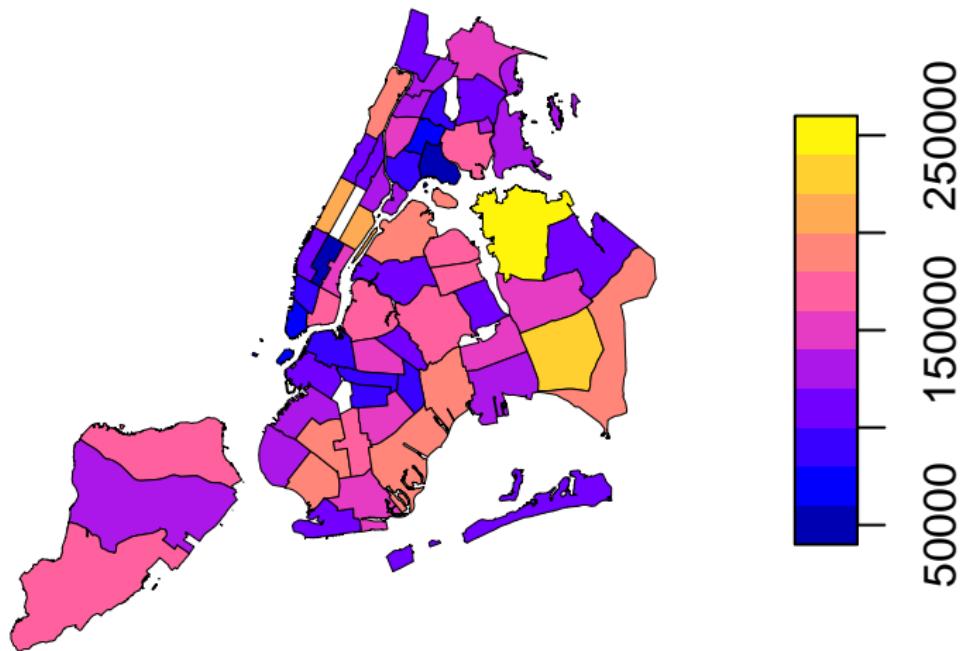
Borough



... or plot them by population size in 2010:

```
plot(nyc_cd_2["pop_2010"],main="Population size in 2010")
```

Population size in 2010



Loading spatial data from a CSV table

Our data on *bike collisions*, like population, is also a csv table:

```
crashes = read.csv(file = "Data/NYC_Collisions/NYC_BikeCrashes.csv")
```

But this file has geographic coordinates, which we can use to create a spatial object, borrowing the CRS from nyc_cd:

```
crashes = sf::st_as_sf(x = crashes, coords = c("LONGITUDE","LATITUDE"),  
                      crs = sf::st_crs(nyc_cd))  
plot(crashes["geometry"])
```



Point-in-polygon analysis

To count the number of collisions in each community district, we can overlay `nyc_cd_2` and `crashes` with `sf::st_intersects()`. The output is a classed list of feature IDs intersected:

```
o = sf::st_intersects(x = nyc_cd_2, y = crashes)
```

```
o
```

```
## Sparse geometry binary predicate list of length 59, where the predicate
## was `intersects'
## first 10 elements:
##  1: 15, 29, 195, 221, 245, 277, 293, 306, 643, 652, ...
##  2: 9, 10, 11, 18, 108, 137, 139, 173, 176, 201, ...
##  3: 79, 85, 168, 224, 251, 266, 269, 270, 338, 355, ...
##  4: 4, 12, 44, 101, 109, 116, 163, 219, 220, 242, ...
##  5: 8, 16, 47, 66, 75, 83, 105, 120, 151, 198, ...
##  6: 21, 58, 97, 204, 236, 275, 289, 291, 316, 332, ...
##  7: 76, 81, 110, 117, 166, 229, 238, 297, 310, 340, ...
##  8: 2, 193, 213, 294, 331, 362, 363, 369, 370, 399, ...
##  9: 89, 103, 131, 152, 167, 295, 320, 357, 492, 546, ...
## 10: 17, 54, 74, 82, 113, 158, 180, 237, 255, 258, ...
```

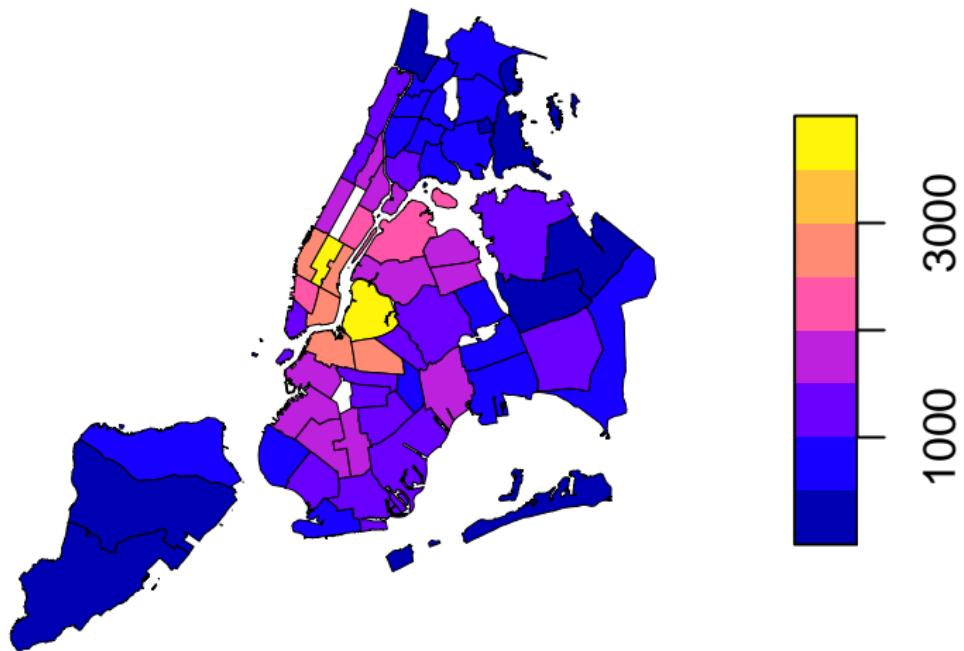
We can use the command `lengths(o)` to find the number of points in each polygon, and assign this as a new variable in `nyc_cd`:

```
nyc_cd_2$crashes = lengths(o)
```

Plot the new variable:

```
plot(nyc_cd_2["crashes"], main = "Bicycle Collisions (2013-2025)")
```

Bicycle Collisions (2013–2025)

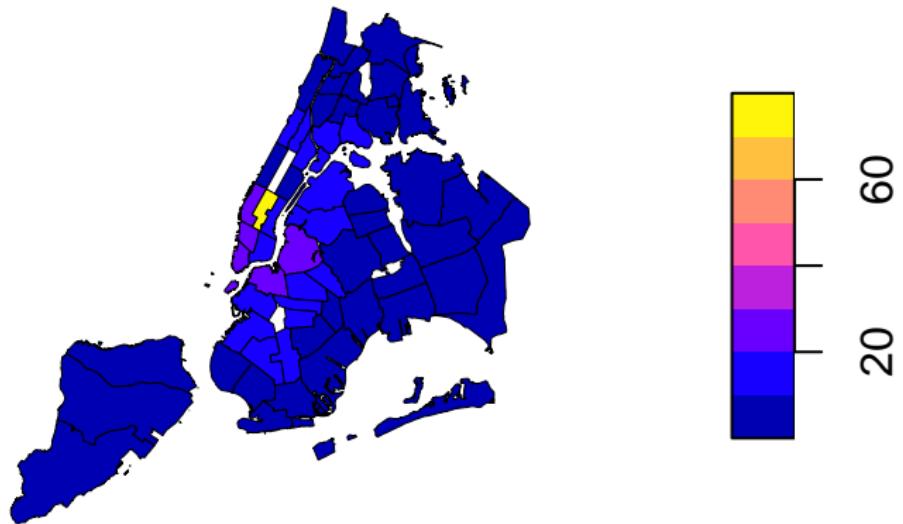


Creating per-capita measures

Let's calculate and plot a new field, crashes per 1000 residents:

```
nyc_cd_2$crash1000 = nyc_cd_2$crashes / nyc_cd_2$pop_2010 * 1000  
plot(nyc_cd_2["crash1000"], main = "Collisions per 1000 Residents")
```

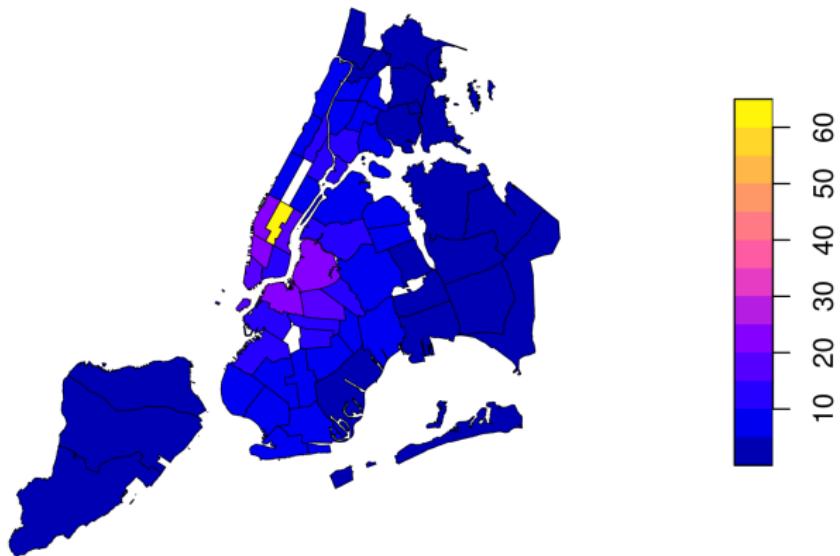
Collisions per 1000 Residents



Exporting Map 1 to image file

To save the map, we will use the `png()` and `dev.off()` commands:

```
png(filename = "Output/bike_crashes_per_capita_r.png",
     width = 6, height = 4, units = "in", res = 300)
plot(nyc_cd_2["crash1000"],
      main = "Bicycle Collisions per 1000 Residents", lwd=.5)
dev.off()
```

Bicycle Crashes per 1000 Residents

The output file should look like this.

Map 2

Loading bike lane data

Import bike routes data

```
bikeroutes = sf::read_sf("Data/NYC_BikeRoutes/NYC_BikeRoutes.geojson",
  crs=sf::st_crs(nyc_cd_2))
plot(bikeroutes["geometry"], reset=FALSE, col="forestgreen", lwd=.5)
plot(nyc_cd_2["geometry"], add = TRUE)
```



Line-in-polygon analysis

We need to calculate the *miles of bike lane* in each community district. We can think of this as a two-step routine:

1. determine which line segments lie inside of which polygons
2. take a local sum of overlapping line lengths in each polygon

For step 1, we can use `sf::st_intersects()` to figure out which segments of bikeroutes lie in which districts (`nyc_cd_3`) (similar to point-in-polygon analysis):

```
o = sf::st_intersects(x = nyc_cd_2, y = bikeroutes)
```

```
o
```

```
## Sparse geometry binary predicate list of length 59, where the predicate
## was `intersects'
## first 10 elements:
##  1: 68, 71, 86, 106, 107, 137, 160, 897, 898, 903, ...
##  2: 72, 73, 85, 1796, 2860, 2861, 2862, 2863, 2864, 2865, ...
##  3: 42, 53, 103, 154, 156, 158, 1696, 1697, 1698, 1699, ...
##  4: 44, 64, 167, 1961, 2735, 3033, 3112, 3113, 3116, 3117, ...
##  5: 167, 1961, 3325, 3326, 3370, 3371, 3372, 3373, 3374, 3375, ...
##  6: 3453, 3457, 3458, 3464, 3465, 3524, 3525, 3526, 3527, 3528, ...
##  7: 46, 63, 65, 3814, 3815, 3884, 3885, 3886, 3942, 3943, ...
##  8: 10, 15, 17, 18, 80, 4374, 4412, 4438, 4439, 4440, ...
##  9: 4820, 4821, 4830, 4831, 4832, 4853, 4854, 4855, 4856, 4857, ...
## 10: 21, 22, 99, 100, 4797, 4799, 4840, 4848, 4849, 4850, ...
```

Line-in-polygon analysis (continued)

For step 2, we can write a loop that extracts the overlapping bike lane segments in each district, measures their (cumulative) length, and converts this sum to miles.

```
out <- c()                      # Create empty vector to store results
for(i in 1:length(o)){           # Open loop
  bike_subset = bikeroutes[o[i][[1]],]      # Take subset
  bk_meters = sum(sf::st_length(bike_subset)) # Sum lengths
  bk_km = bk_meters/1000                     # Convert to km
  bk_mi = bk_km * 0.6213712                  # Convert to miles
  out = c(out, bk_mi)                        # Concatenate
}
# Close loop
```

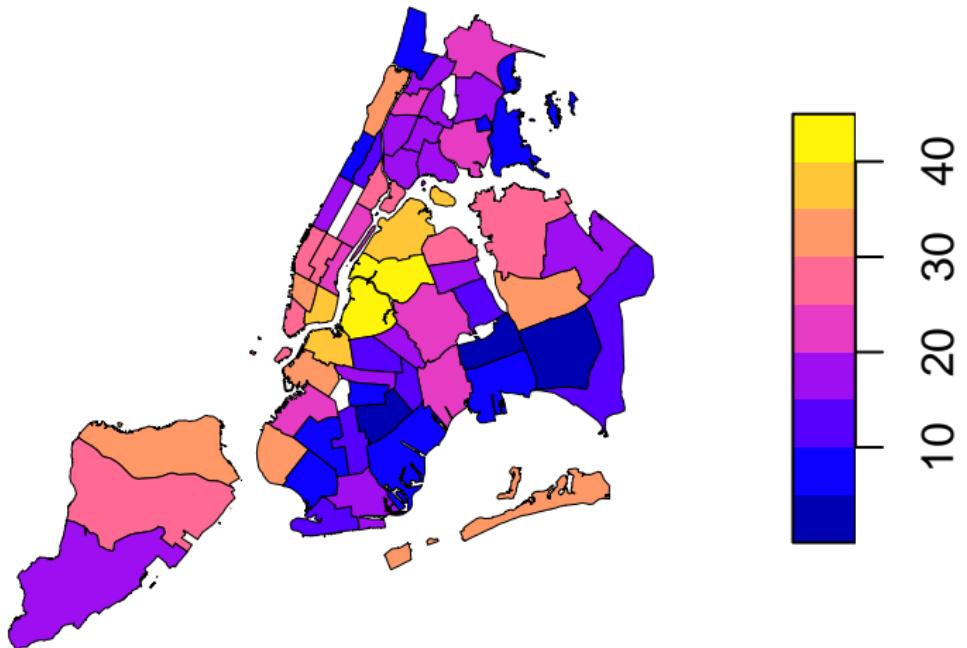
We then add this as a variable to nyc_cd_2

```
nyc_cd_2$bk_miles <- out
```

Let's quickly visualize the results:

```
plot(nyc_cd_2["bk_miles"],  
      main = "Miles of Bike Lane per District", lwd=.5)
```

Miles of Bike Lane per District

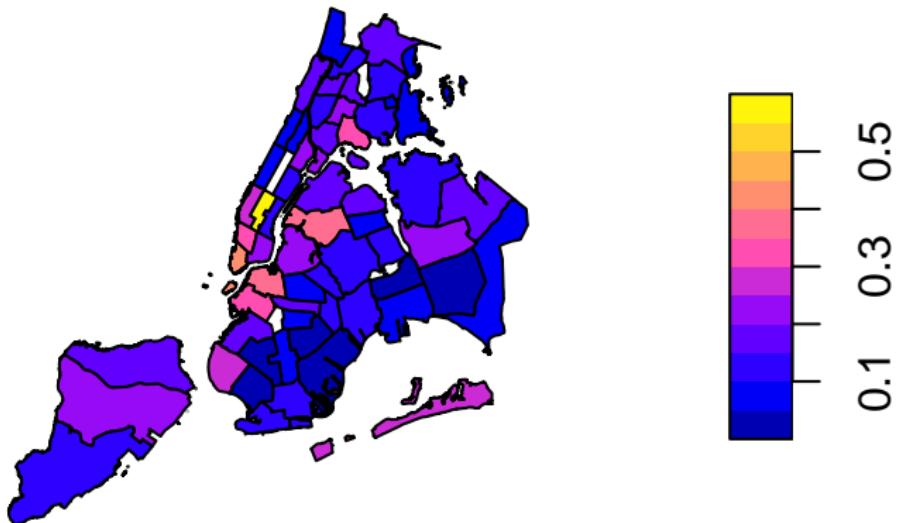


Creating per-capita measures

Let's calculate and plot a new field, miles of bicycle lane per 1000 residents:

```
nyc_cd_2$bklane1000 = nyc_cd_2$bk_miles / nyc_cd_2$pop_2010 * 1000  
plot(nyc_cd_2["bklane1000"], main = "Bike Lane per 1000 Residents (mi)")
```

Bike Lane per 1000 Residents (mi)

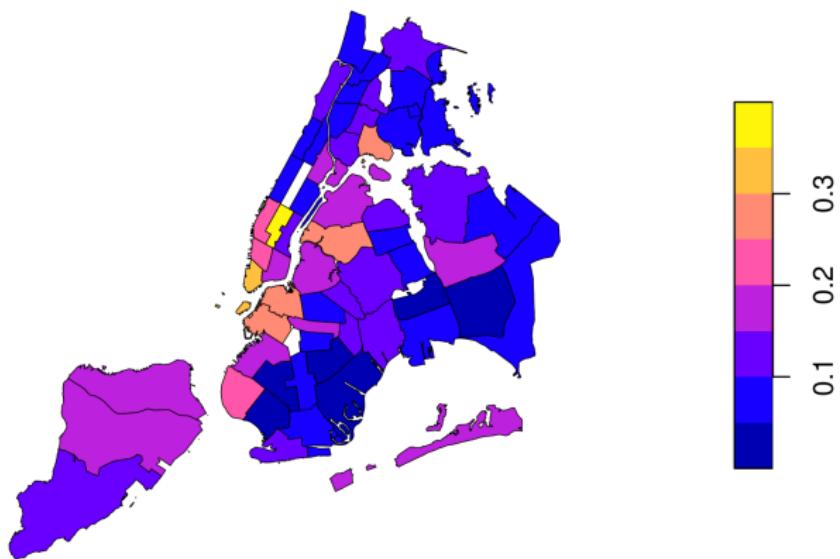


Exporting to image file

To save the map, we will use the `png()` and `dev.off()` commands:

```
png(filename = "Output/bike_lanes_per_capita_R.png",
     width = 6, height = 4, units = "in", res = 300)
plot(nyc_cd_2["bklane1000"],
      main = "Miles of Bicycle Lane per 1000 Residents", lwd=.5)
dev.off()
```

Miles of Bicycle Lane per 1000 Residents



The output file should look like this.

Exporting the dataset

To save the data attribute table as a csv file, we can use the `write.csv()` command:

```
write.csv(x = nyc_cd_2, file = "Output/NYC_CD_2_R.csv")
```

Making bar plots

For each of NYC's five boroughs, calculate averages of the variables we created:

```
borough_means = aggregate(  
  x = nyc_cd_2[,c("crashes","crash1000","bk_miles","bklane1000")],  
  by = list(Borough=nyc_cd_2$Borough),  
  FUN = mean )
```

This will create a new sf object with 5 rows:

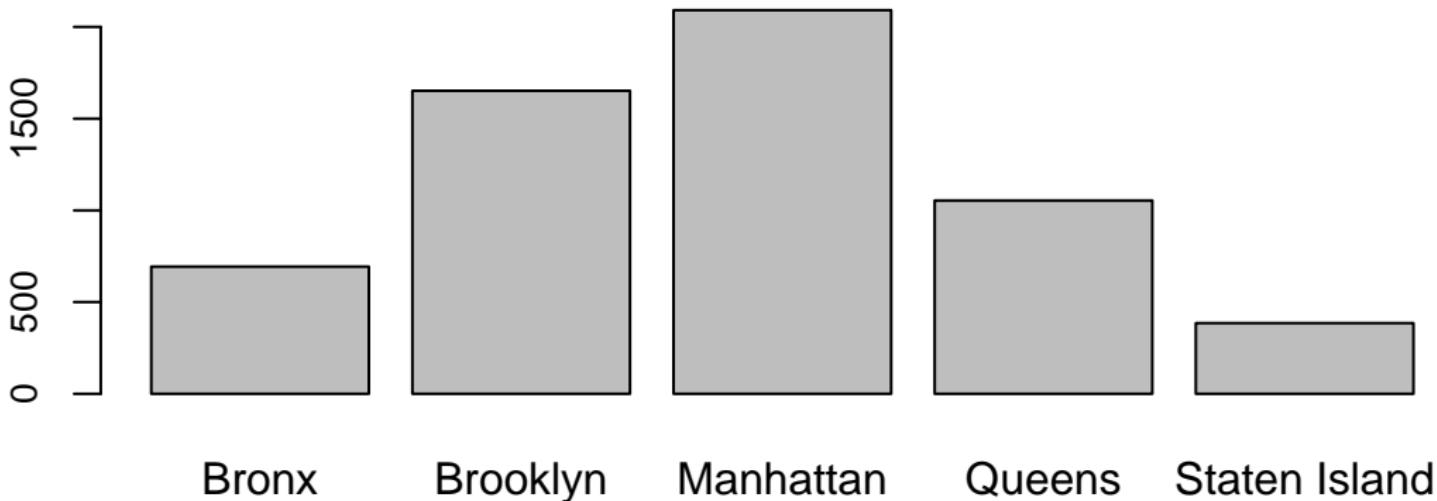
```
borough_means
```

```
borough_means  
  
## Simple feature collection with 5 features and 5 fields  
## Attribute-geometry relationships: aggregate (4), identity (1)  
## Geometry type: MULTIPOLYGON  
## Dimension: XY  
## Bounding box: xmin: -74.25559 ymin: 40.49613 xmax: -73.70001 ymax: 40.91553  
## Geodetic CRS: WGS 84  
##           Borough   crashes crash1000 bk_miles bklane1000  
## 1      Bronx    693.2500  6.556499 16.90893  0.1585769  
## 2     Brooklyn  1652.1667 12.305324 18.39499  0.1454636  
## 3   Manhattan  2091.2500 20.398060 24.66573  0.2314566  
## 4     Queens   1053.8571  6.576689 20.80431  0.1403516  
## 5 Staten Island   385.3333  2.402792 26.72134  0.1725632  
##           geometry  
## 1 MULTIPOLYGON (((-73.89857 4...
```

Make a bar plot of *bike crashes in an average community district*

```
barplot(height=borough_means$crashes, names.arg=borough_means$Borough,  
        main="Bicycle collisions per community district (mean)")
```

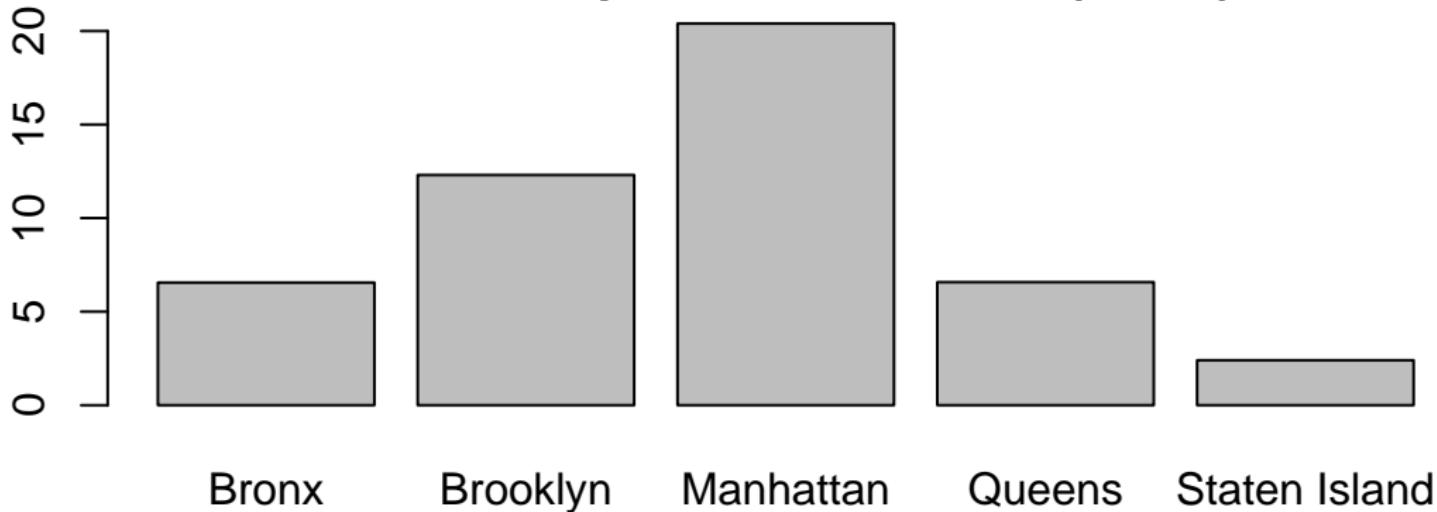
Bicycle collisions per community district (mean)



Make a bar plot of *bike crashes per 1000 residents*

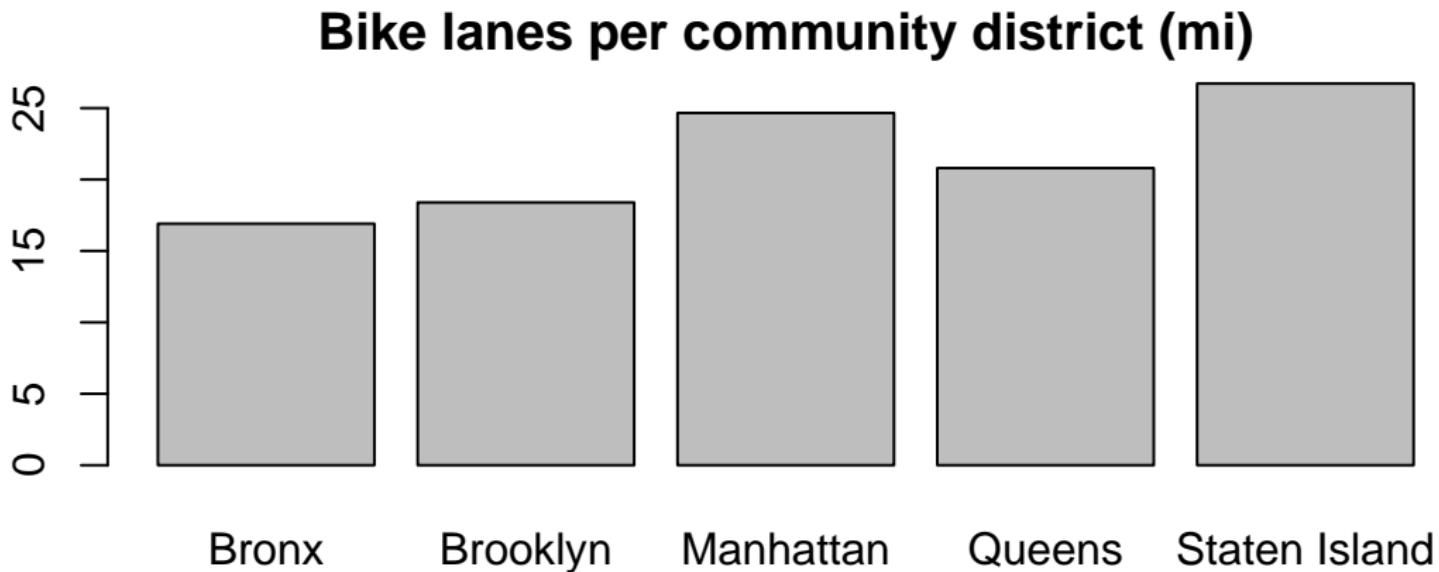
```
barplot(height = borough_means$crash1000,  
        names.arg = borough_means$Borough,  
        main = "Collisions per 1000 residents (mean)")
```

Collisions per 1000 residents (mean)



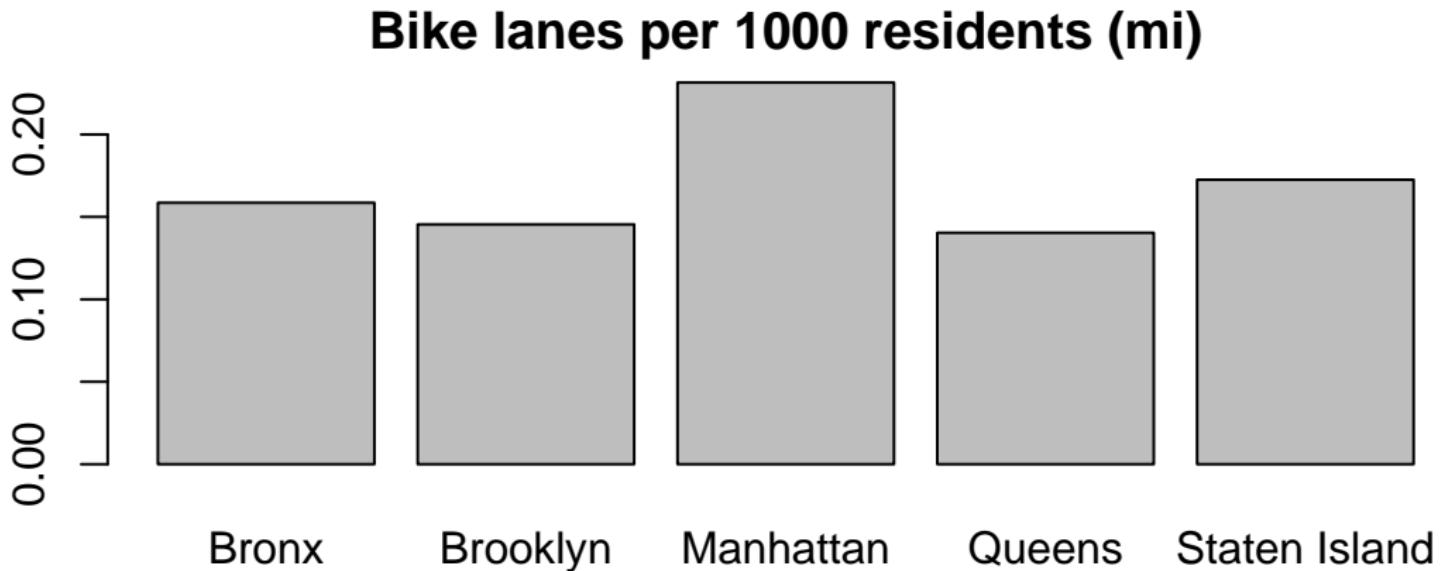
Make a bar plot of *miles of bike lane in an average community district*

```
barplot(height = borough_means$bk_miles,  
        names.arg = borough_means$Borough,  
        main = "Bike lanes per community district (mi)")
```



Make a bar plot of *miles of bike lane per 1000 residents*

```
barplot(height = borough_means$bklane1000,  
        names.arg = borough_means$Borough,  
        main = "Bike lanes per 1000 residents (mi)")
```



Basic regression analysis

You can use an `sf` object as you would a `data.frame` to run regressions in R.

For example, suppose we wanted to see if there are fewer bicycle collisions where there are more bike lanes. We could run a basic linear regression of bicycle collisions (per 1000) on bike lanes (per 1000), with fixed effects for each borough:

```
summary(lm(crash1000 ~ bklane1000 + as.factor(Borough), data=nyc_cd_2))

##
## Call:
## lm(formula = crash1000 ~ bklane1000 + as.factor(Borough), data = nyc_cd_2)
##
## Residuals:
##     Min      1Q  Median      3Q     Max 
## -12.690 -2.670   0.243   1.926  34.108 
##
## Coefficients:
##                               Estimate Std. Error t value Pr(>|t|)    
## (Intercept)              -2.887    2.345  -1.231  0.22372  
## bklane1000                 59.554    8.616   6.912 6.25e-09 *** 
## as.factor(Borough)Brooklyn  6.530    2.464   2.651  0.01057 *  
## as.factor(Borough)Manhattan 9.501    2.768   3.433  0.00117 ** 
## as.factor(Borough)Queens    1.106    2.603   0.425  0.67270  
## as.factor(Borough)Staten Island -4.987    4.264  -1.169  0.24747  
## ---                        
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

This analysis suggests there are *more collisions where there are more bike lanes.*

Adjusting for differences across boroughs, each additional mile of bike lane (per 1000 residents) is associated with 65 additional collisions (per 1000).

Does this mean bike lanes are “causing” collisions?

How should we interpret this result?

What are some alternative explanations of the positive correlation?

What sort of data would we need to test these alternative arguments?

Problem Set 3

Your assignment (if using R):

- create a map of rat activity (per 1000 residents) in NYC!
- use this dataset:
 - NYC_Rats/NYC_Rats.csv
- follow the same steps as for “Map 1” (crashes per 1000)
(i.e. sum via point-in-polygon,
then calculate per capita rate)
- name the file
`rats_per_capita.png`
- upload map to Canvas
(by next Monday)

Rat Activity per 1000 Residents

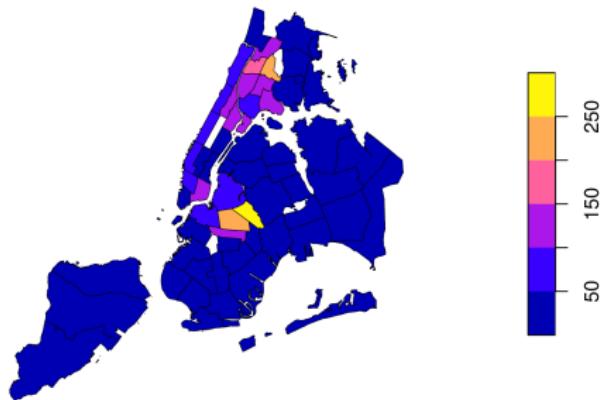


Figure 17: Can you make this map?

Problem Set 3

Your assignment (if using Python):

- create a map of rat activity (per 1000 residents) in NYC!
- use this dataset:
 - NYC_Rats/NYC_Rats.csv
- follow the same steps as for “Map 1” (crashes per 1000)
(i.e. sum via point-in-polygon, then calculate per capita rate)
- name the file `rats_per_capita.png`
- upload map to Canvas
(by next Monday)

Rat Activity per 1000 Residents

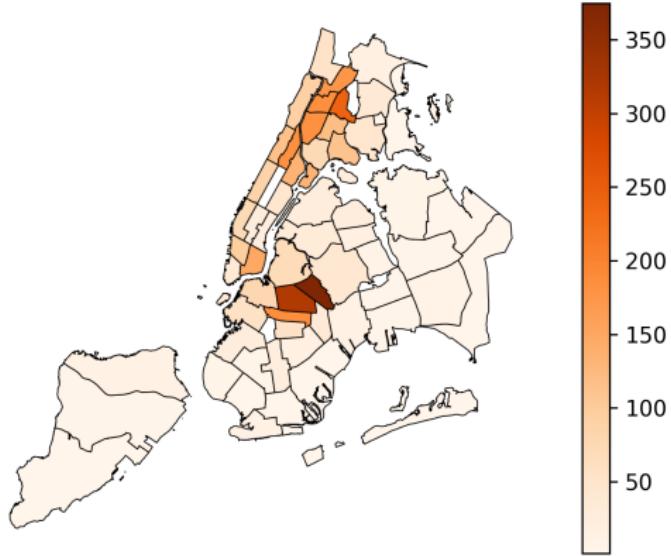


Figure 18: Can you make this map?