SEST-6577

Geographic Information Systems for Security Studies

Lab 09 (+ Walk Through 1)

Yuri M. Zhukov
Associate Professor
Georgetown University

October 31, 2024

Goal: explain geographic variation in Islamic State violence

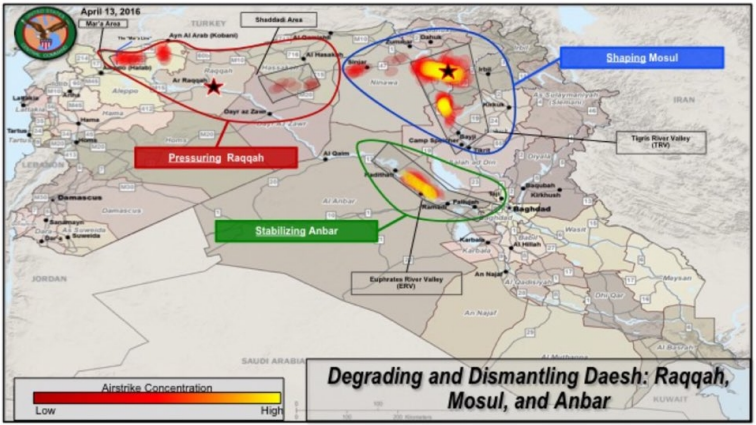

Figure 1: Operation Inherent Resolve

There is no formal Problem Set this week, but if you would like to replicate this in QGIS, follow along (for walk-through in R, see replication code on Canvas)

Overview of lab exercise

1. Data
   a) Collect raw data from open sources online
   b) Pre-processing to integrate data to common set of spatial units
2. Analysis
   a) Export pre-processed data as comma-separated file
   b) Run regression models in R

Data

Collection

Data on **political violence** are easily found online. If we search for `georeferenced violent events data`, a top result is for "UCDP Dataset Download Center"

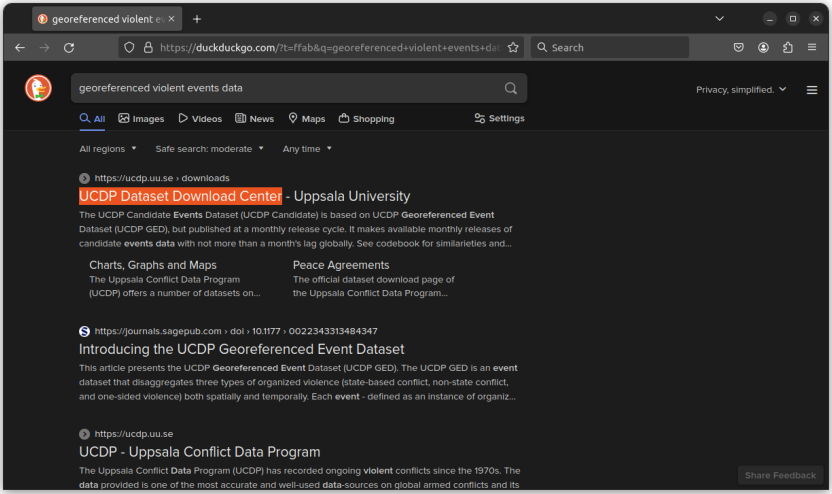

The Uppsala Conflict Data Program (UCDP) is one of the world's leading providers of open-source data on violence and armed conflict, based at Uppsala University

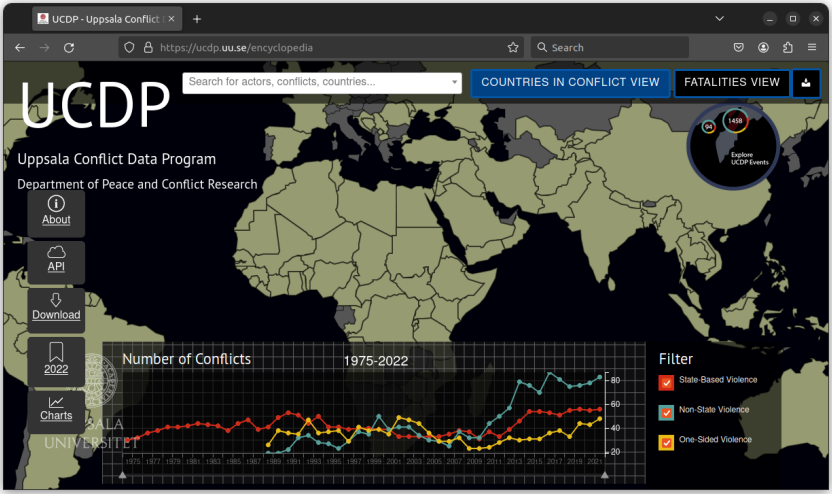

Download-able data are located at `ucdp.uu.se/downloads`.
Click on "Georeferenced Event Dataset (GED)"

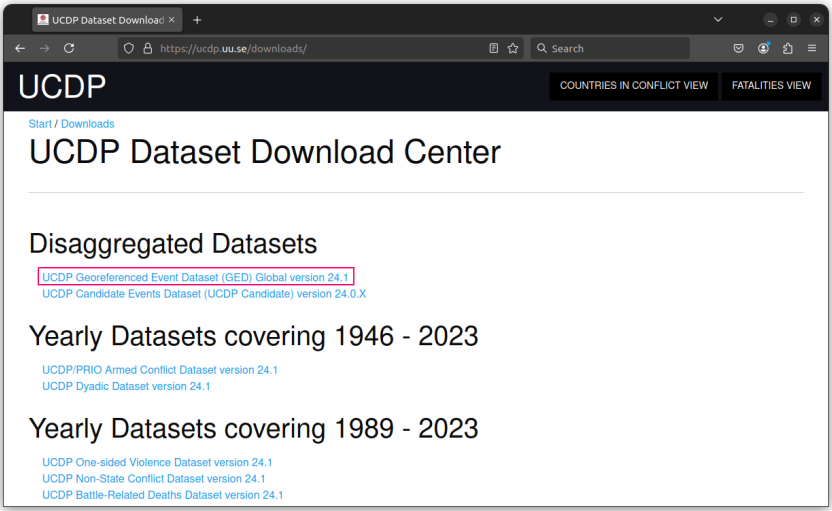

We have used GED before, for Problem Set 02 (Afghanistan and Yemen). Download GED as a `csv` file and extract it from the `zip` archive

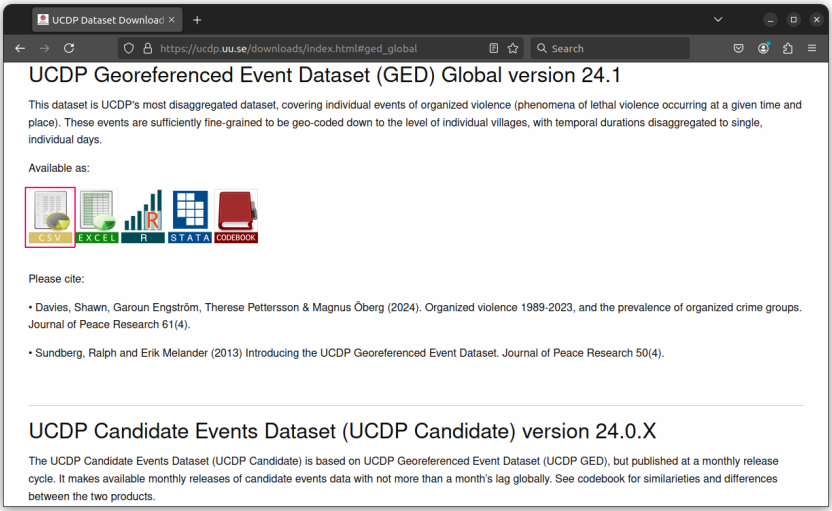

There are about 317K events here, each of which has point coordinates

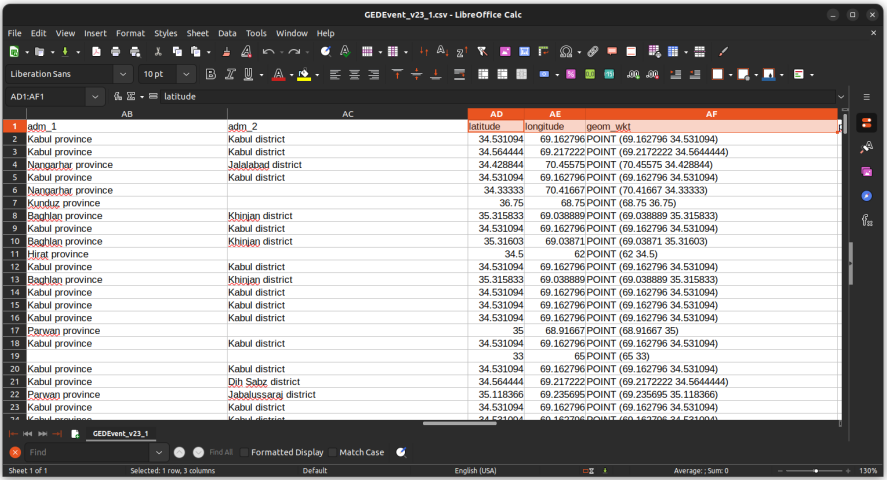

	AB	AC	AD	AE	AF
1	adm_1	adm_2	latitude	longitude	geom_wkt
2	Kabul province	Kabul district	34.531094	69.162796	POINT (69.162796 34.531094)
3	Kabul province	Kabul district	34.564444	69.217222	POINT (69.2172222 34.5644444)
4	Nangarhar province	Jalalabad district	34.428844	70.45575	POINT (70.45575 34.428844)
5	Kabul province	Kabul district	34.531094	69.162796	POINT (69.162796 34.531094)
6	Nangarhar province		34.33333	70.41667	POINT (70.41667 34.33333)
7	Kunduz province		36.75	68.75	POINT (68.75 36.75)
8	Baghlan province	Khinjan district	35.315833	69.038889	POINT (69.038889 35.315833)
9	Kabul province	Kabul district	34.531094	69.162796	POINT (69.162796 34.531094)
10	Baghlan province	Khinjan district	35.31603	69.03871	POINT (69.03871 35.31603)
11	Hirat province		34.5	62	POINT (62 34.5)
12	Kabul province	Kabul district	34.531094	69.162796	POINT (69.162796 34.531094)
13	Baghlan province	Khinjan district	35.315833	69.038889	POINT (69.038889 35.315833)
14	Kabul province	Kabul district	34.531094	69.162796	POINT (69.162796 34.531094)
15	Kabul province	Kabul district	34.531094	69.162796	POINT (69.162796 34.531094)
16	Kabul province	Kabul district	34.531094	69.162796	POINT (69.162796 34.531094)
17	Parwan province		35	68.91667	POINT (68.91667 35)
18	Kabul province	Kabul district	34.531094	69.162796	POINT (69.162796 34.531094)
19			33	65	POINT (65 33)
20	Kabul province	Kabul district	34.531094	69.162796	POINT (69.162796 34.531094)
21	Kabul province	Dih Sabz district	34.564444	69.217222	POINT (69.2172222 34.5644444)
22	Parwan province	Jabalussaraj district	35.118366	69.235695	POINT (69.235695 35.118366)
23	Kabul province	Kabul district	34.531094	69.162796	POINT (69.162796 34.531094)

For data on **administrative boundaries**, Rick Wilson's "Free GIS Data" site has links to multiple data sources. Of these, we will use `gadm.org`

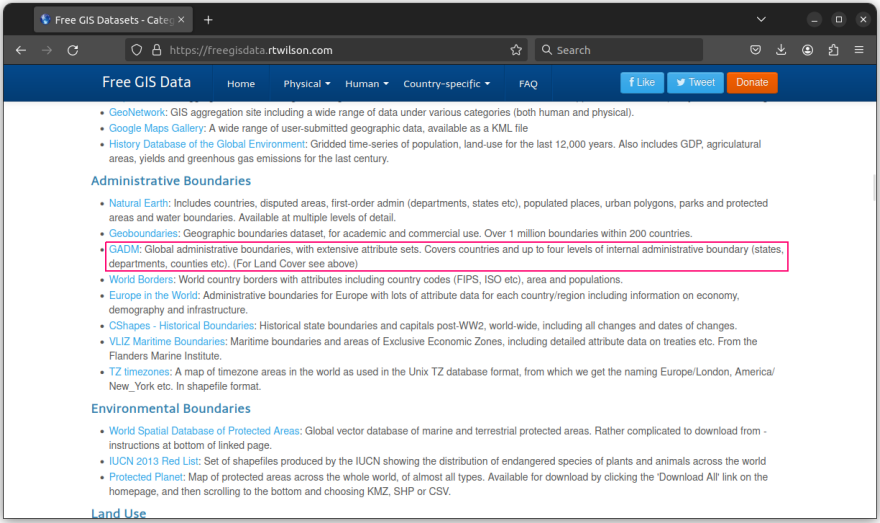

From GADM's website, we can get data on sub-national administrative boundaries. Download the `level2` `GeoJSON` file for Iraq...

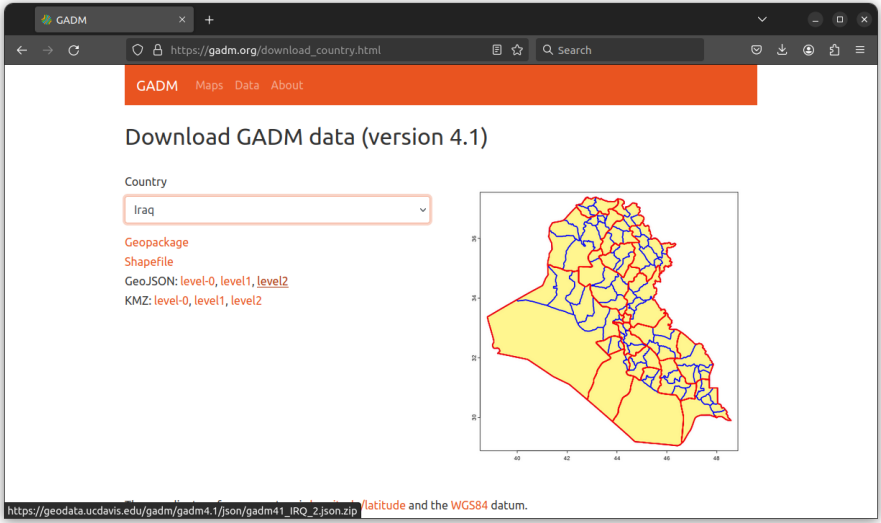

. . . and for Syria

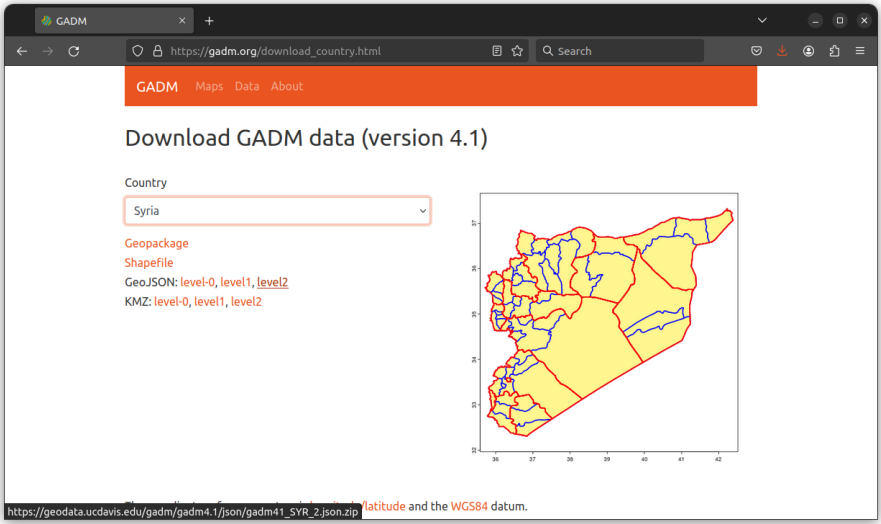

Multiple data sources exist for **roads**. Among the easiest to get are those from `diva-gis.org/gData`

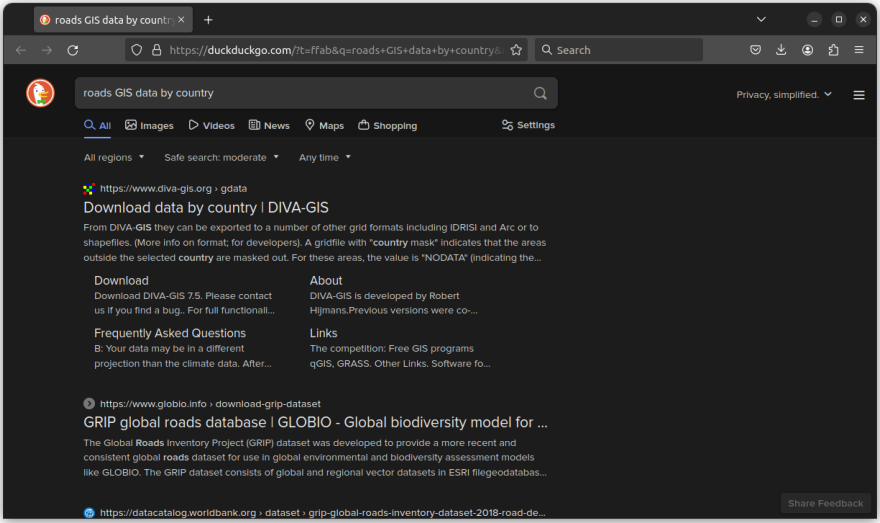

These are country-level road layers from “Digital Chart of the World”, download-able as zipped shapefiles

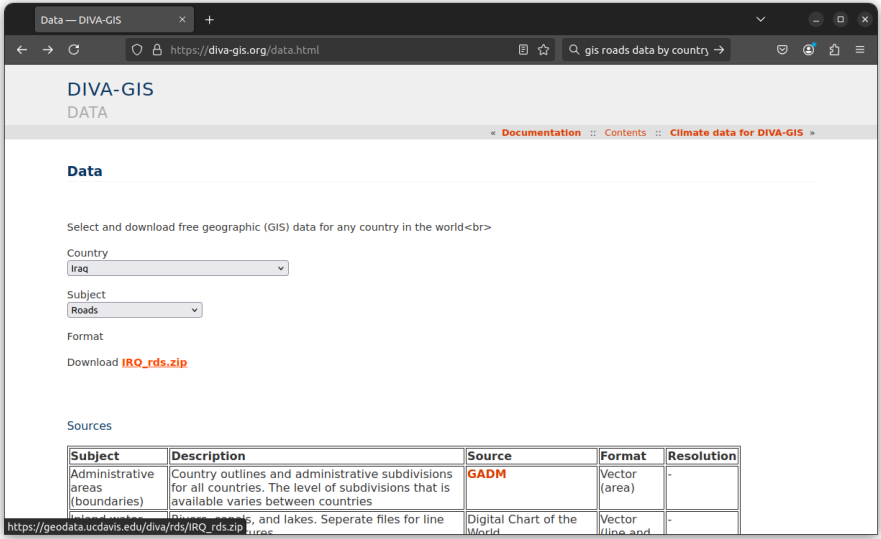

Download the roads data for both Iraq and Syria here

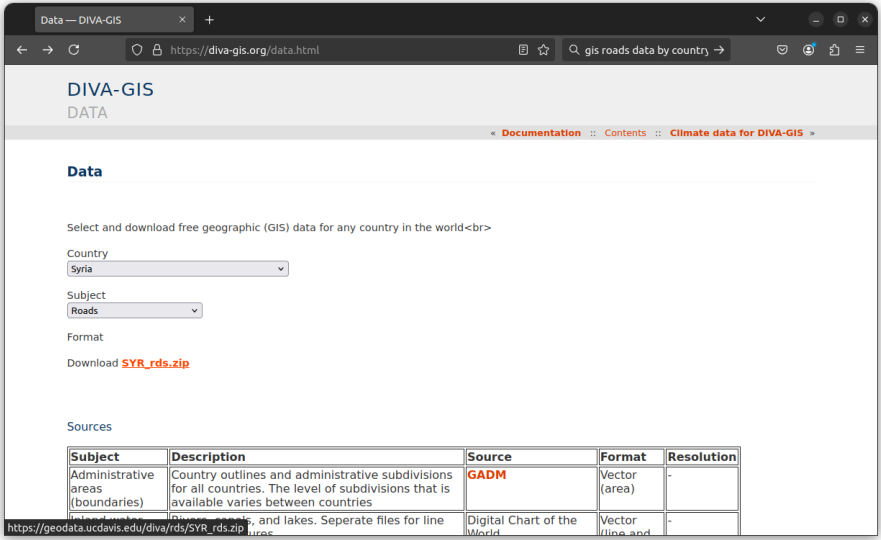

For **population data**, we can go back to Rick Wilson's page to find the link for "Gridded Population of the World (GPW)"

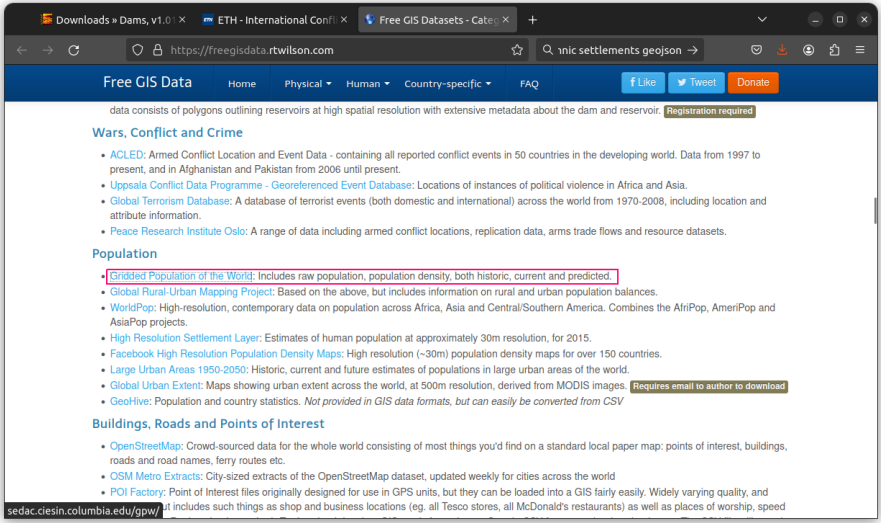

This will take us to the Center for International Earth Science Information Network (CIESIN) at Columbia, which hosts a large number of global, free GIS datasets

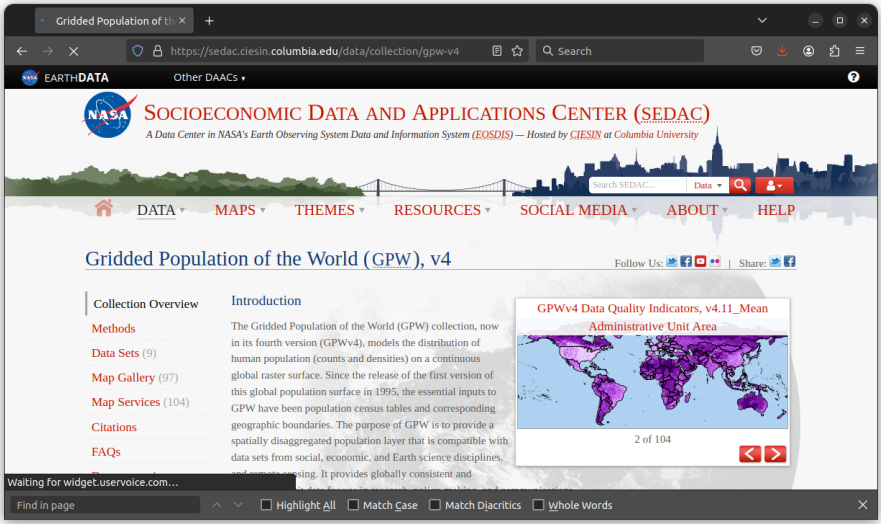

Navigate to GPW's "Population Count v4.11" dataset, and download the files for 2010 (you will need to create an account to login)

The data are available for multiple years, in multiple formats and resolutions. We will be using the 15-minute GeoTiff raster for 2010 (latest before ISIS emerged in 2014)

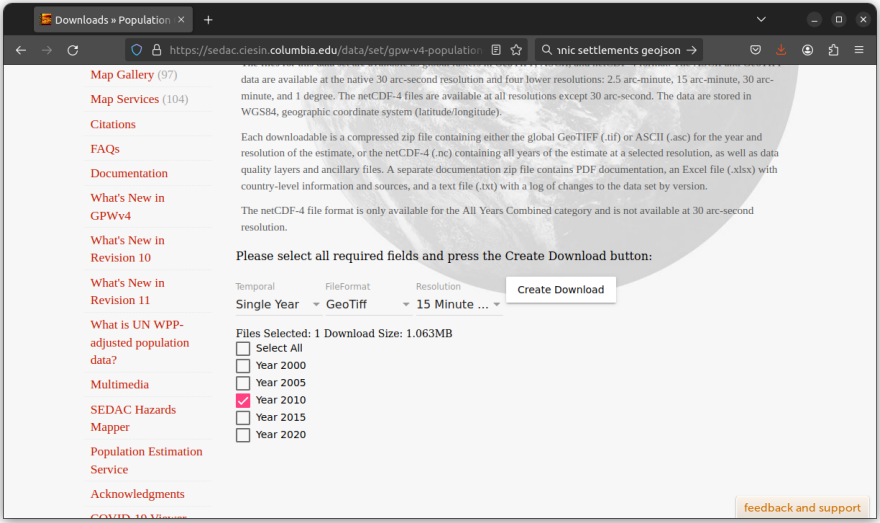

Also from CIESIN, we can download data on **cropland** (“Croplands, v1 (2000)”)

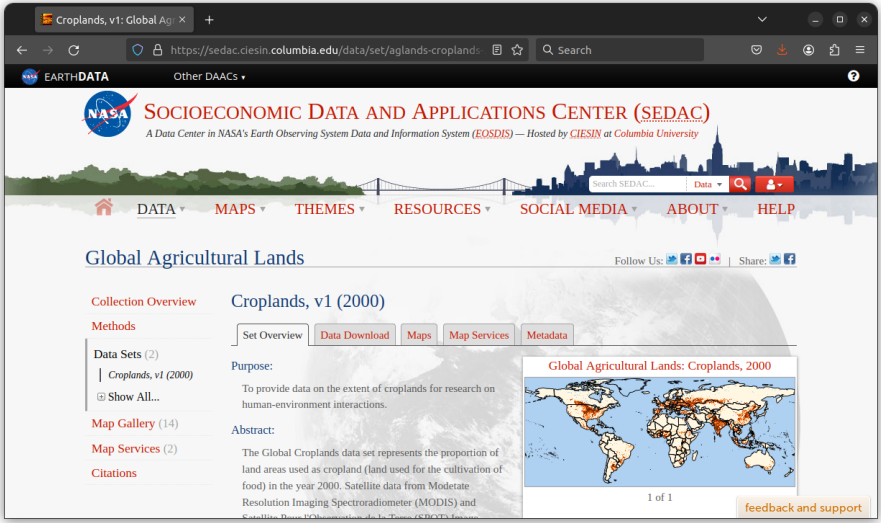

Download the global raster, in GeoTiff format

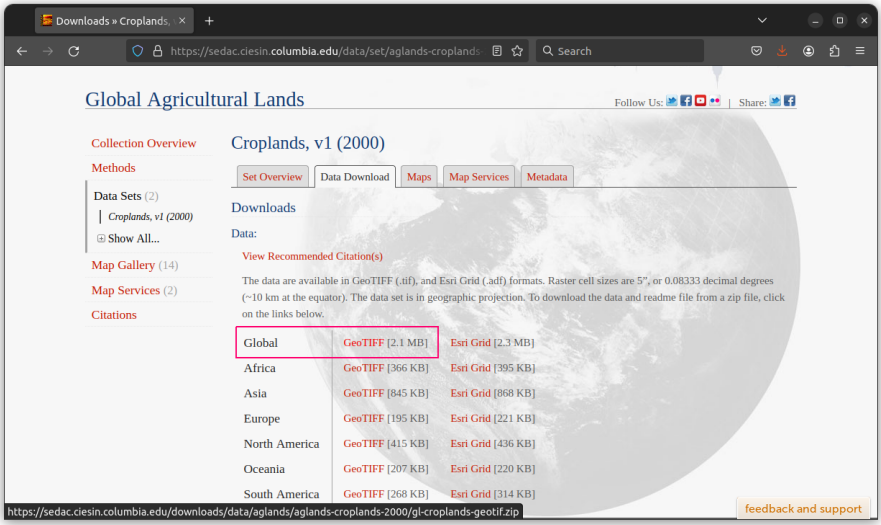

While we're here, let's also download data on **hydroelectric dams** ("Global Reservoir and Dam (GRanD), v1")

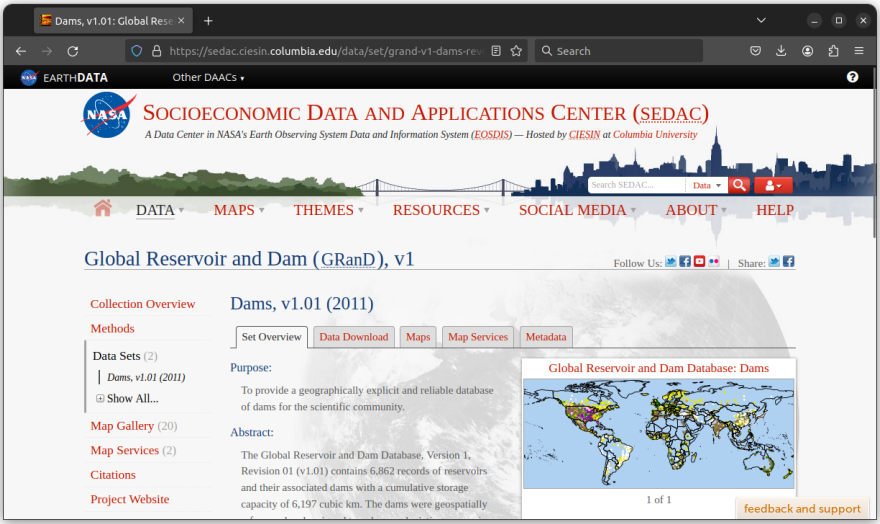

For data on **sectarian divisions**, a relatively reliable global source is the "Ethnic Power Relations" (geoEPR) dataset — we used this in Problem Set 02

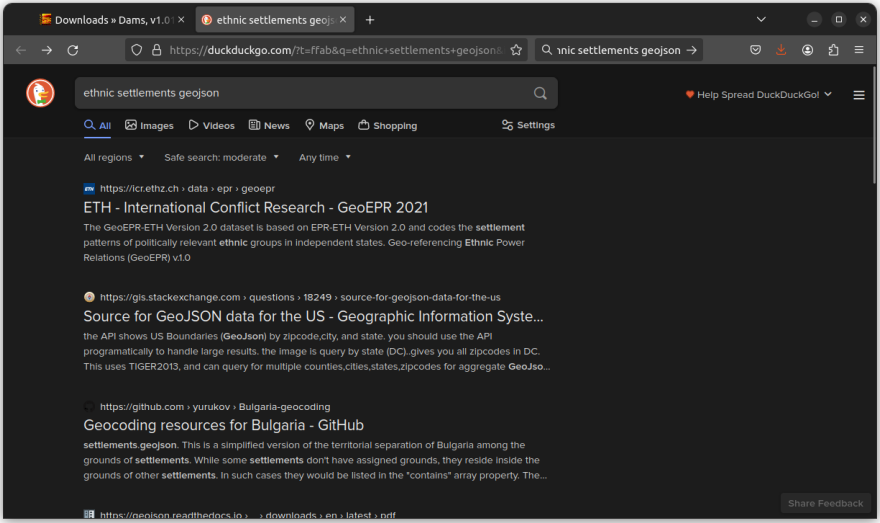

Specifically, we will download the `GeoEPR-2021.geojson` file, which contains polygons for each ethno-religious group (`icr.ethz.ch/data/epr/geoepr`)

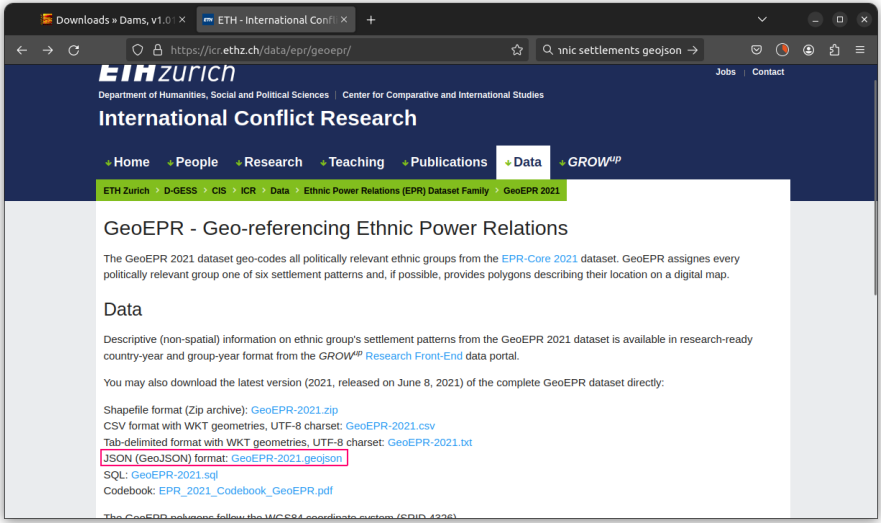

Here is the full list of data sources and links:

Category	Type	Format	Data source
ISIS violence	Table (non-geo)	`.csv`	UCDP GED
Administrative units	Vector (polygons)	`.geojson`	GADM
Roads	Vector (polylines)	`.shp`	DIVA-GIS
Population	Raster	`.tif`	CIESIN
Cropland	Raster	`.tif`	CIESIN
Dams	Vector (points)	`.shp`	CIESIN
Sectarian divisions	Vector (polygons)	`.geojson`	ETH-Zurich

These are all in the `WT01.zip` file posted on Canvas.

Pre-processing

Always save your progress!

Go to `Project` → `Save As...`

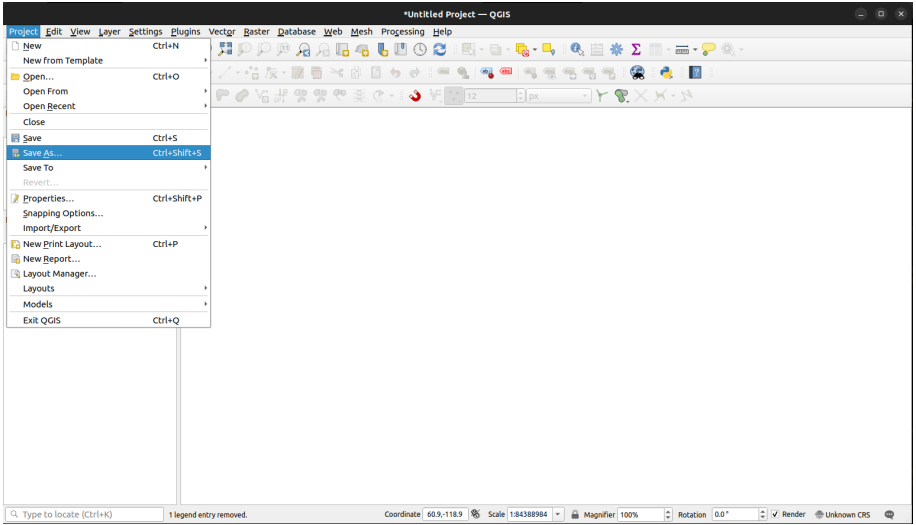

Load *administrative boundaries* (`Layer` → `Add Layer` → `Add Vector Layer`).
For **Iraq**: `gadm41_IRQ_2.json` file in `Data/GADM` folder.

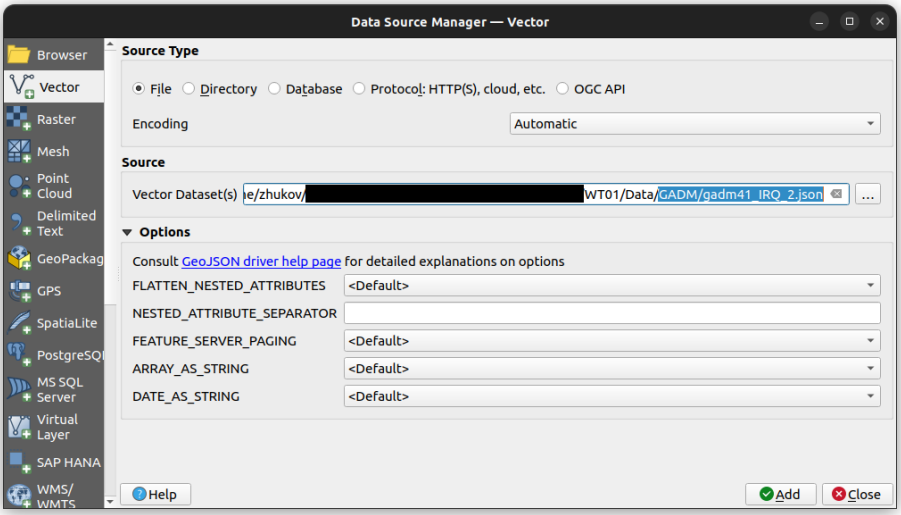

For **Syria**: `gadm41_SYR_2.json` file in `Data/GADM` folder.

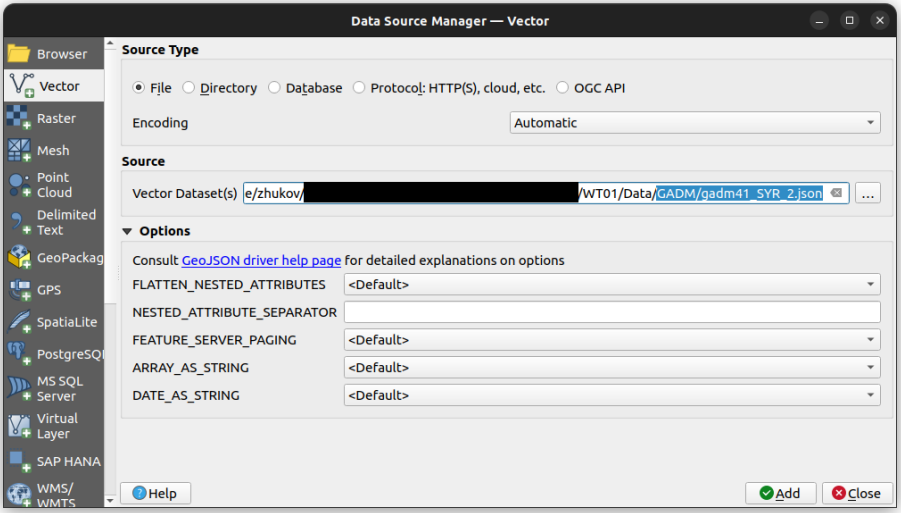

The two sets of district boundaries should be visible in the project window. Let's **merge them into a single layer**, with districts for both countries.

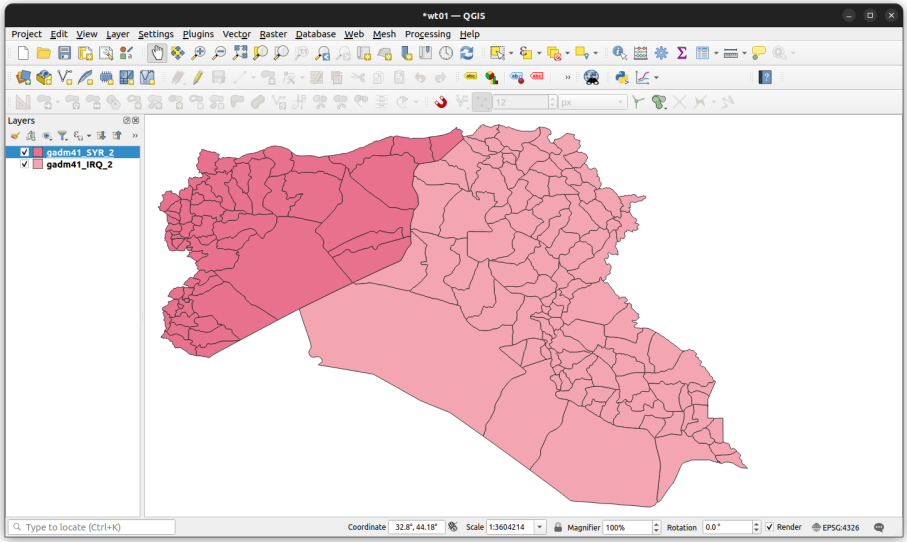

Open the `Merge Vector Layers` tool (`Vector` → `Data Management Tool` → `Merge Vector Layers`).

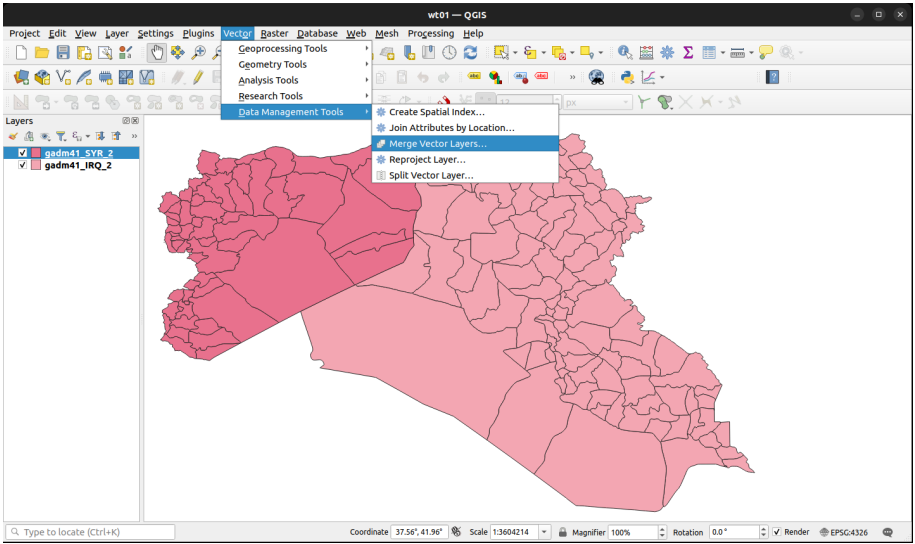

Click on the `[...]` box next to `Input layers`

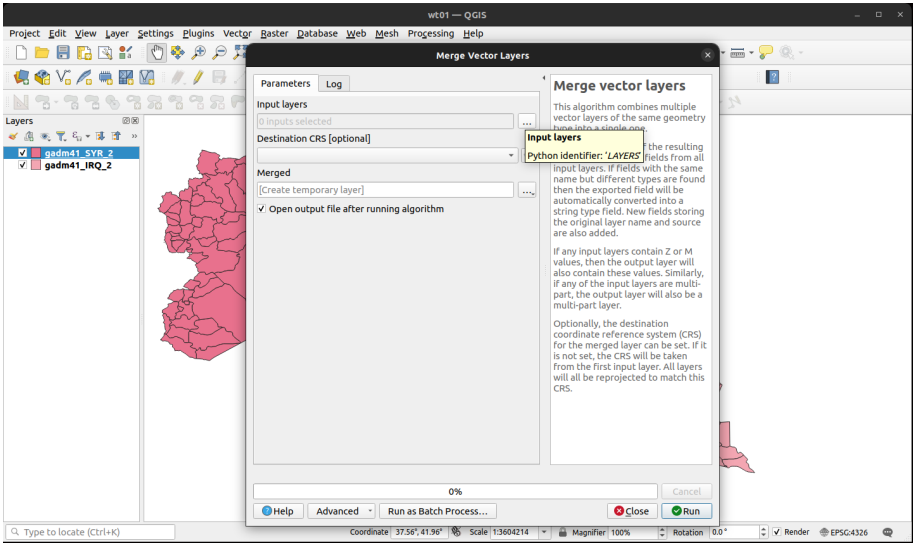

Check the boxes next to gadm41_IRQ_2 and gadm41_SYR_2. Click OK.

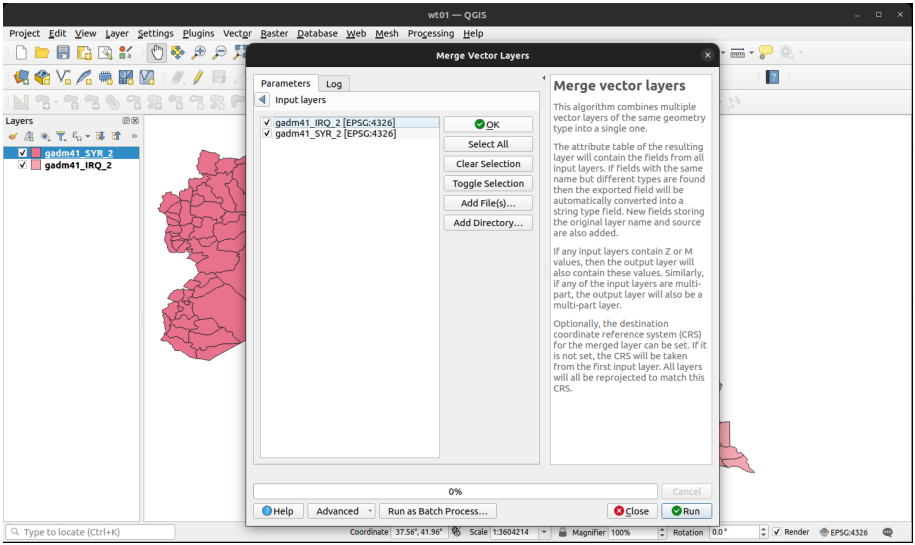

Click on the `[...]` box next to `Merged`. Find a location in which to save the output, and name the file `admin.geojson`. Click `Run`

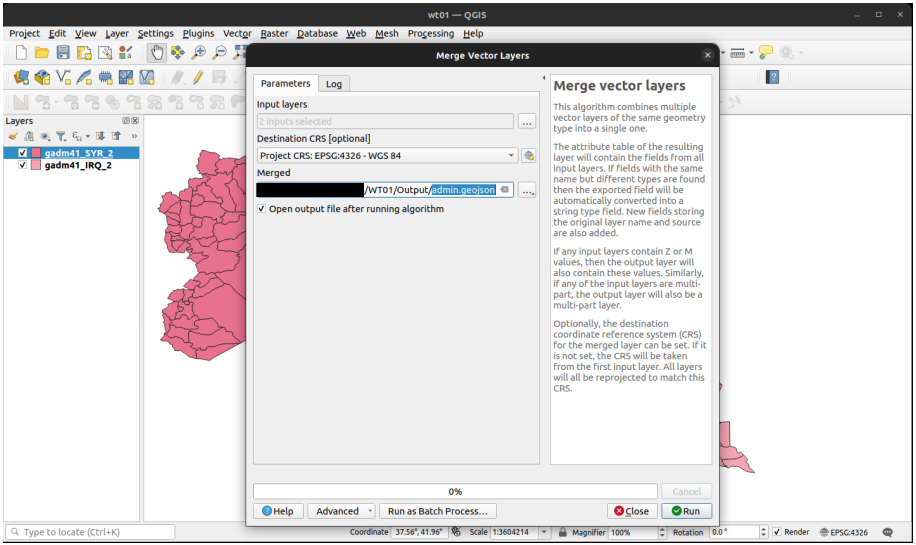

The merged layer should appear in the main project window.

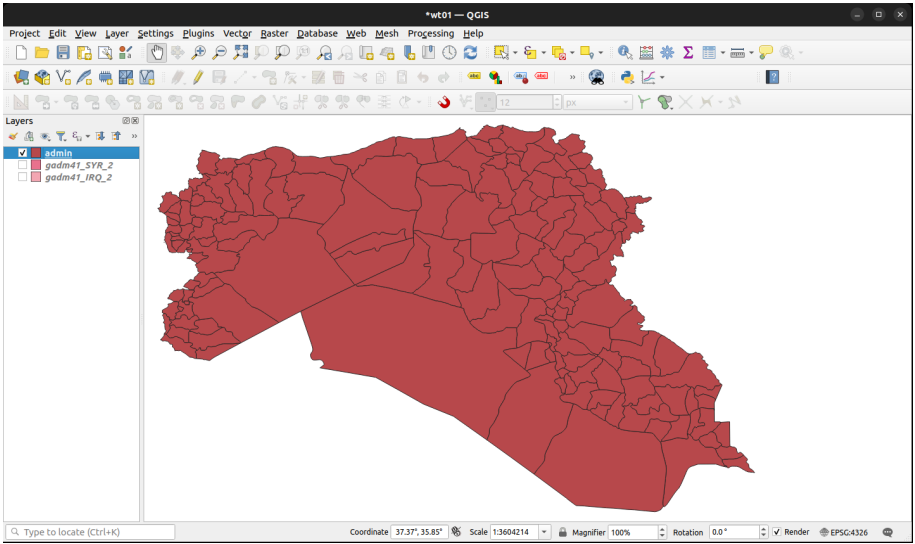

You can remove the `gadm41_IRQ_2` and `gadm41_SYR_2` layers from the project (`Right-click` on layer(s) in Layers menu → `Remove Layer...`)

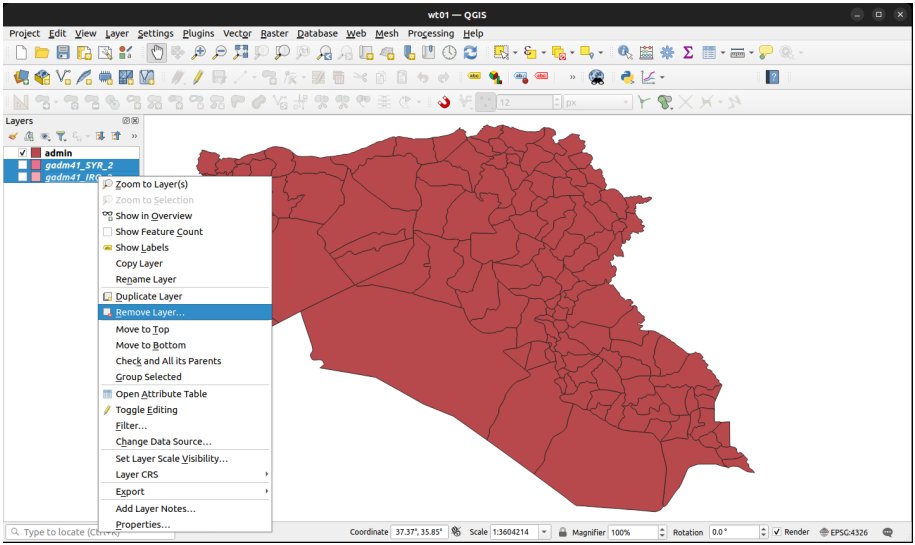

To conserve memory, it's good practice to remove data objects were are not actively using

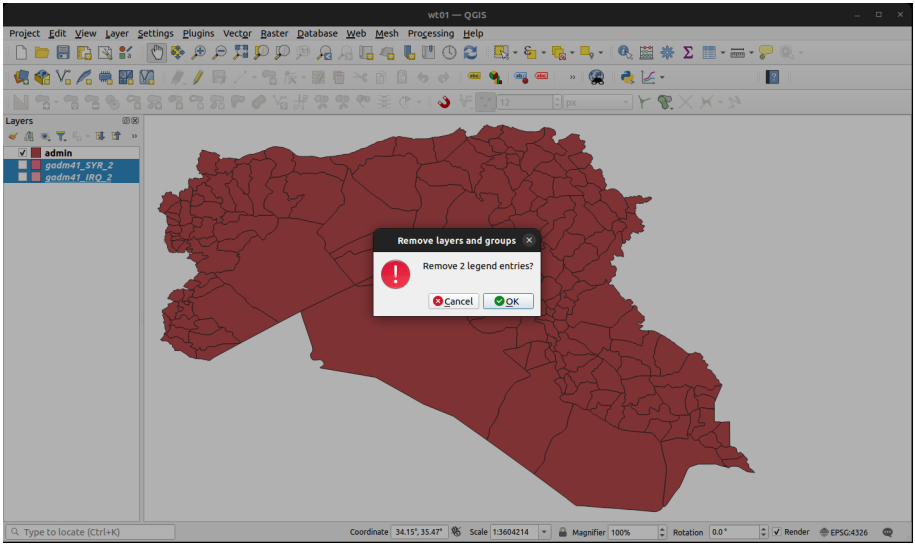

Load *roads data* (`Layer` → `Add Layer` → `Add Vector Layer`).
Try opening **two datasets at the same time**: `IRQ_roads.shp` and `SYR_roads.shp` files in `Data/Roads`

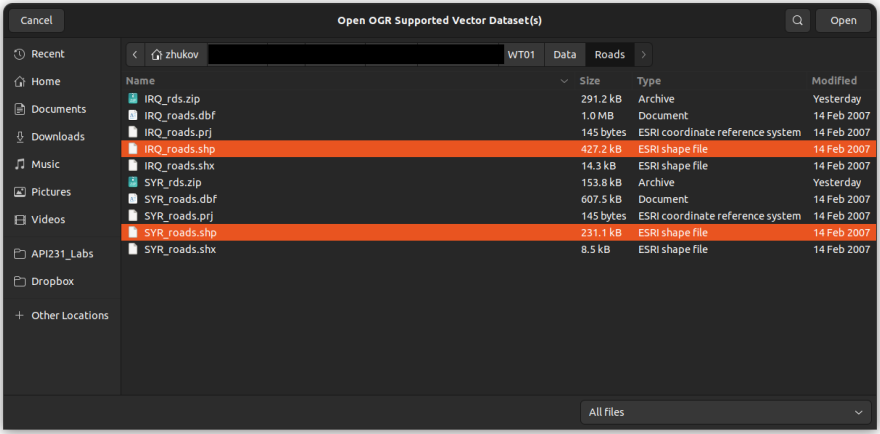

Paths to both files should appear in the `Source` box. Click `Add`

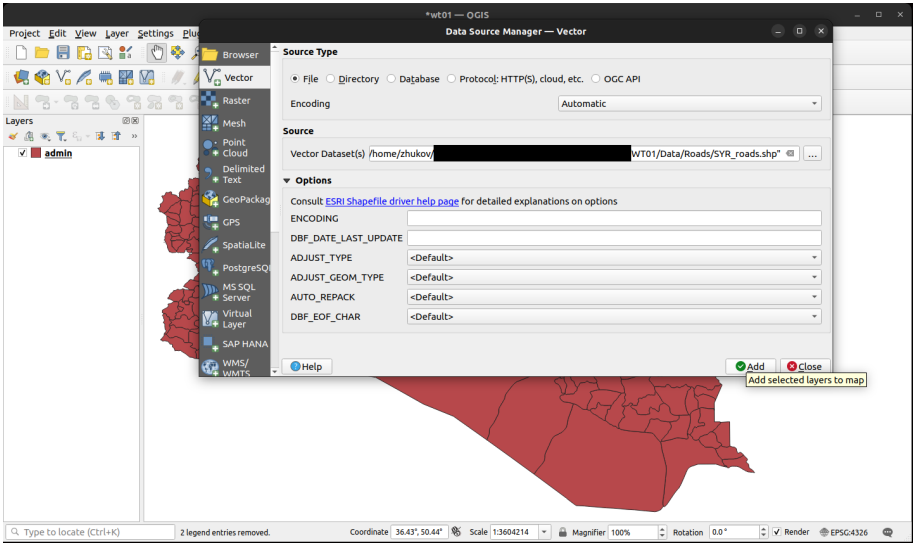

Let's repeat the `Merge Vector Layers` procedure for these two layers (`Vector` → `Data Management Tool` → `Merge Vector Layers`).

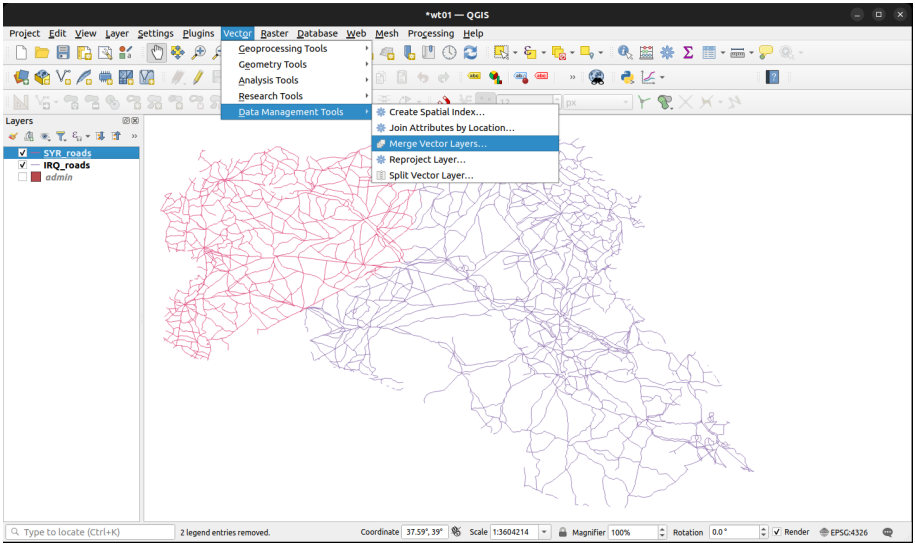

This time, perform the merge on the `IRQ_roads` and `SYR_roads` layers

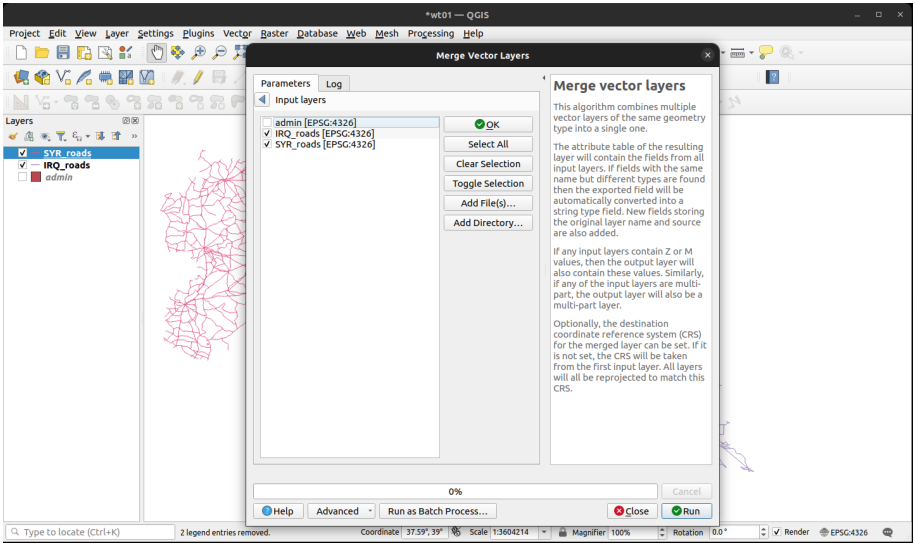

Save the output file as `roads.geojson`

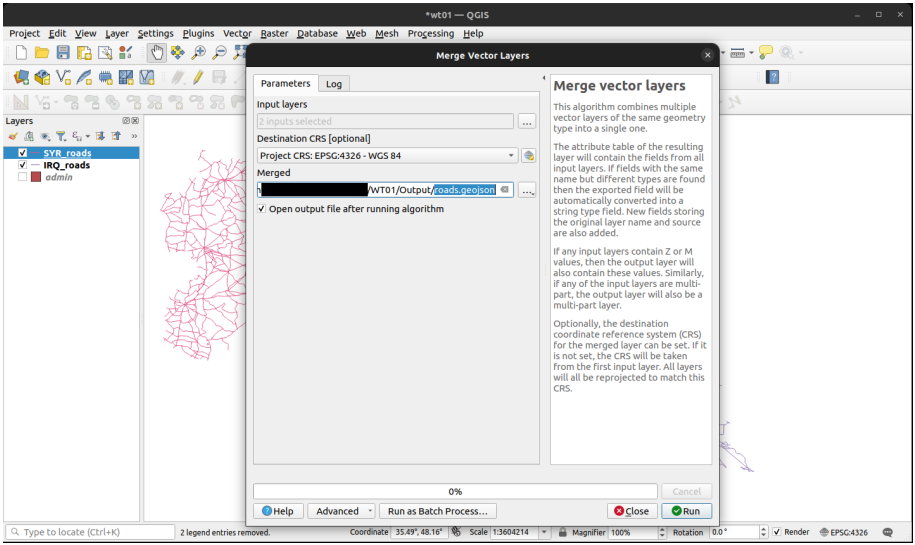

As before, you can remove the two country-specific roads layers (we no longer need them)

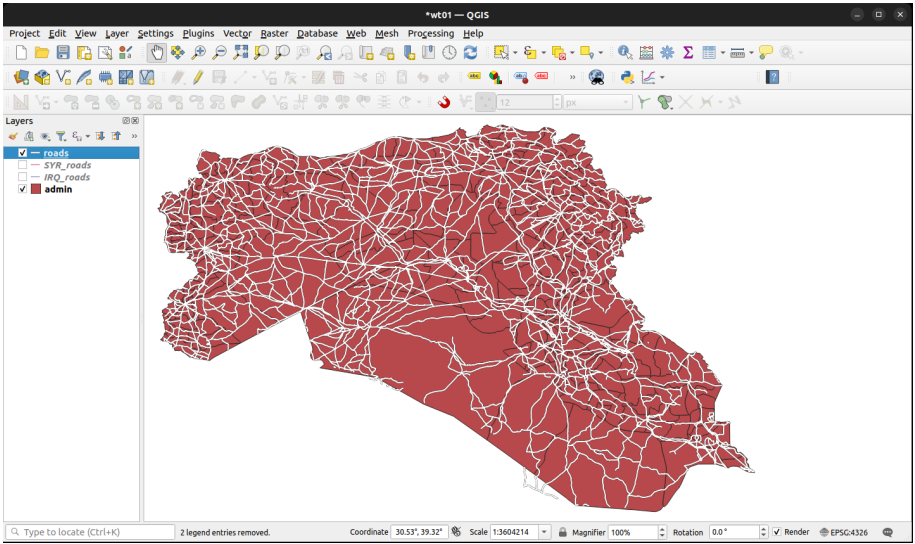

Let's **calculate each district's road density** with the `Sum Line Lengths` tool (`Vector` → `Analysis` → `Sum Line Lengths`)

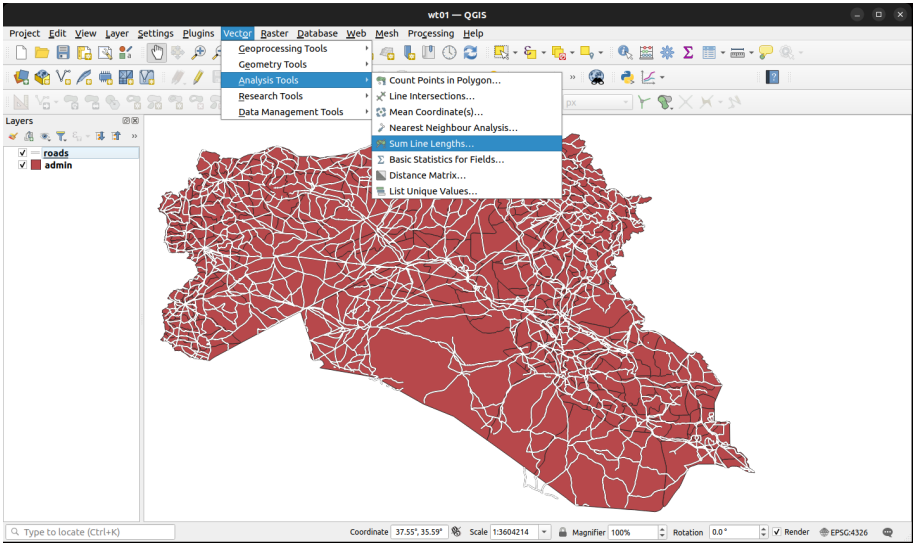

Set `Polygons` = `admin`, `Lines` = `roads`, `Line length field name` = `roads_m`, `Lines count field name` = `roads_ct`. Save as `admin_2.geojson`

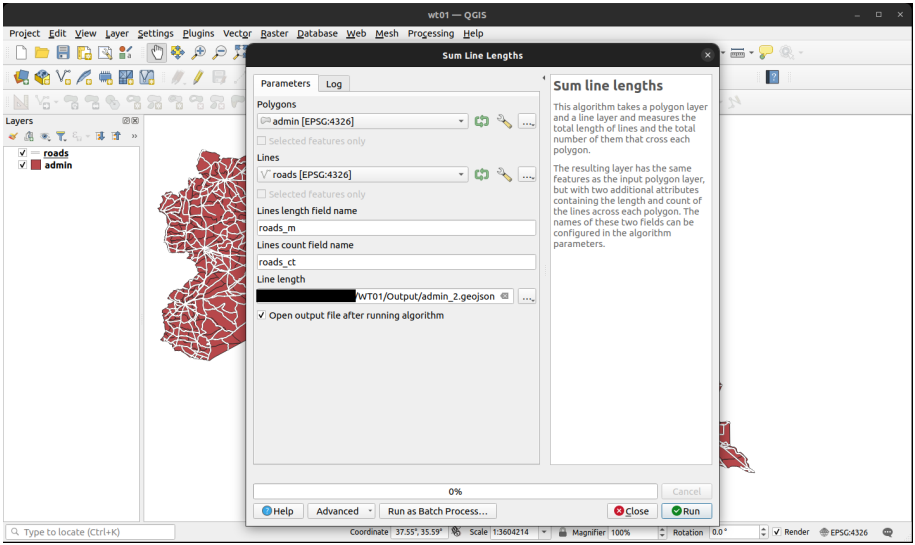

Open the attribute table for `admin_2`. The `roads_m` and `roads_ct` variables should be in the table. Let's convert from **meters to kilometers**

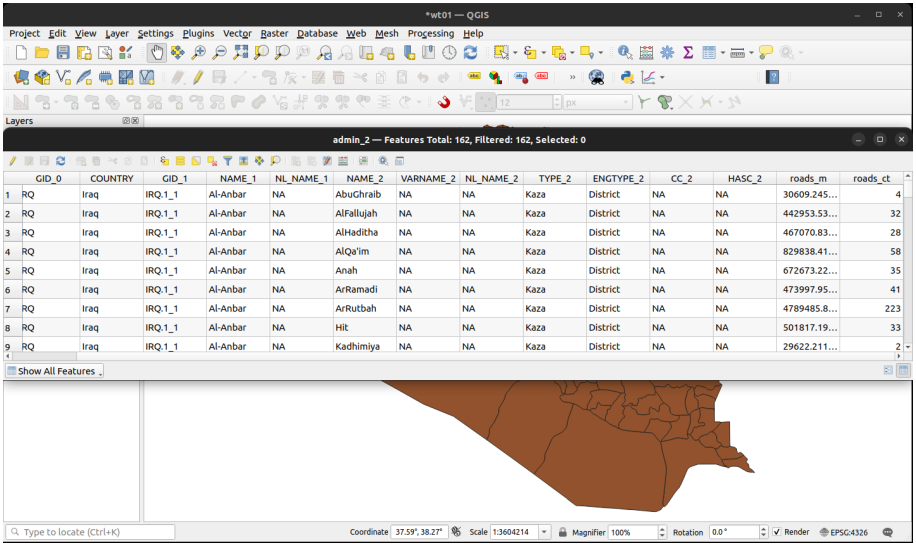

	GID_0	COUNTRY	GID_1	NAME_1	NL_NAME_1	NAME_2	VARNAME_2	NL_NAME_2	TYPE_2	ENGTYPE_2	CC_2	HASC_2	roads_m	roads_ct
1	RQ	Iraq	IRQ.1_1	Al-Anbar	NA	AbuGhraib	NA	NA	Kaza	District	NA	NA	30609.245…	4
2	RQ	Iraq	IRQ.1_1	Al-Anbar	NA	AlFallujah	NA	NA	Kaza	District	NA	NA	442953.53…	32
3	RQ	Iraq	IRQ.1_1	Al-Anbar	NA	AlHaditha	NA	NA	Kaza	District	NA	NA	467070.83…	28
4	RQ	Iraq	IRQ.1_1	Al-Anbar	NA	AlQa'im	NA	NA	Kaza	District	NA	NA	829838.41…	58
5	RQ	Iraq	IRQ.1_1	Al-Anbar	NA	Anah	NA	NA	Kaza	District	NA	NA	672673.22…	35
6	RQ	Iraq	IRQ.1_1	Al-Anbar	NA	ArRamadi	NA	NA	Kaza	District	NA	NA	473997.95…	41
7	RQ	Iraq	IRQ.1_1	Al-Anbar	NA	ArRutbah	NA	NA	Kaza	District	NA	NA	4789485.8…	223
8	RQ	Iraq	IRQ.1_1	Al-Anbar	NA	Hit	NA	NA	Kaza	District	NA	NA	501817.19…	33
9	RQ	Iraq	IRQ.1_1	Al-Anbar	NA	Kadhimiya	NA	NA	Kaza	District	NA	NA	29622.211…	2

Open `Field Calculator` for `admin_2`. Create new field, `roads_km`, of type `Decimal number (real)`. For the `Expression`, write `roads_m/1000`. The `Output preview` should show a number with decimals. Click `OK`

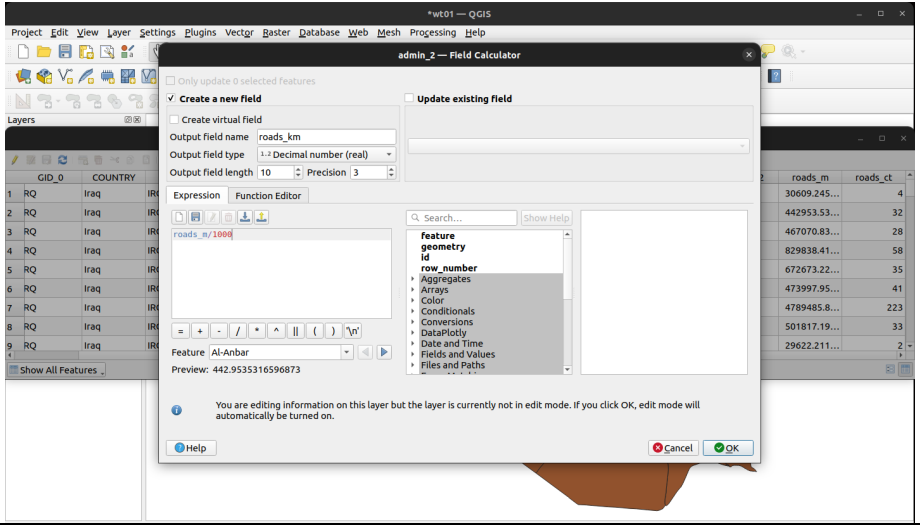

The resulting `roads_km` field should be visible in the attribute table

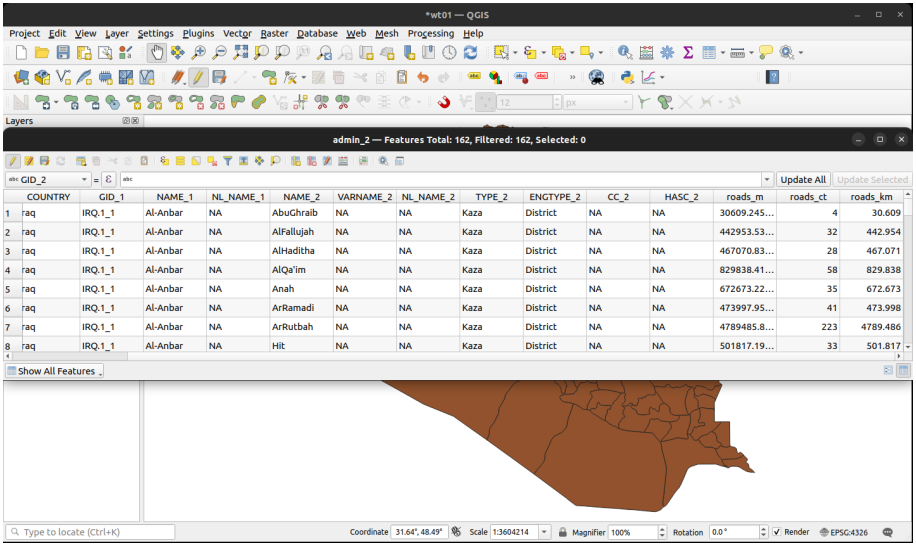

	COUNTRY	GID_1	NAME_1	NL_NAME_1	NAME_2	VARNAME_2	NL_NAME_2	TYPE_2	ENGTYPE_2	CC_2	HASC_2	roads_m	roads_ct	roads_km
1	raq	IRQ.1_1	Al-Anbar	NA	AbuGhraib	NA	NA	Kaza	District	NA	NA	30609.245…	4	30.609
2	raq	IRQ.1_1	Al-Anbar	NA	AlFallujah	NA	NA	Kaza	District	NA	NA	442953.53…	32	442.954
3	raq	IRQ.1_1	Al-Anbar	NA	AlHaditha	NA	NA	Kaza	District	NA	NA	467070.83…	28	467.071
4	raq	IRQ.1_1	Al-Anbar	NA	AlQa'im	NA	NA	Kaza	District	NA	NA	829838.41…	58	829.838
5	raq	IRQ.1_1	Al-Anbar	NA	Anah	NA	NA	Kaza	District	NA	NA	672673.22…	35	672.673
6	raq	IRQ.1_1	Al-Anbar	NA	ArRamadi	NA	NA	Kaza	District	NA	NA	473997.95…	41	473.998
7	raq	IRQ.1_1	Al-Anbar	NA	ArRutbah	NA	NA	Kaza	District	NA	NA	4789485.8…	223	4789.486
8	raq	IRQ.1_1	Al-Anbar	NA	Hit	NA	NA	Kaza	District	NA	NA	501817.19…	33	501.817

If we plot this variable, we can see its main limitations: smaller districts have less road length, so this becomes a proxy for district size. Let's create a **road density** measure that takes into account district size

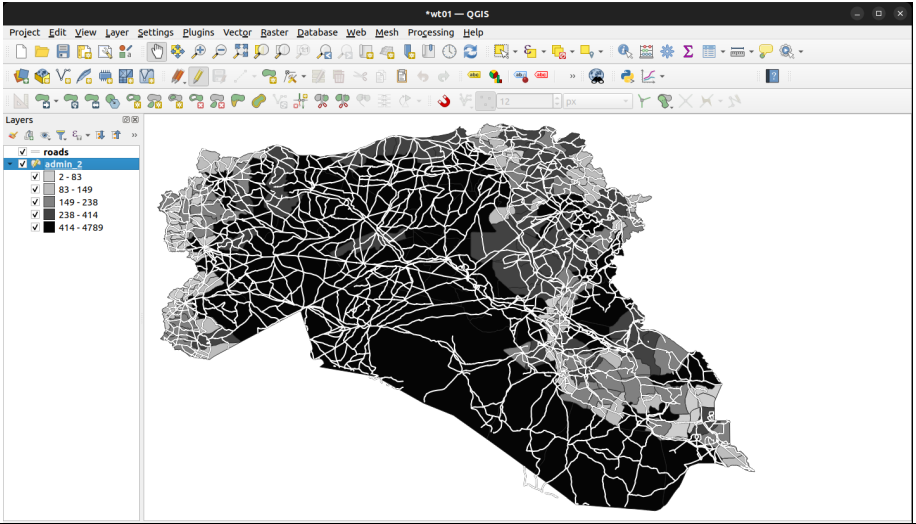

Back in the `Field Calculator`, create a new field called `area_km2` of type `Decimal number (real)`. For `Expression`, write `$area/1000000` (divide by 1M to convert from m^2 to km^2). Click `OK`

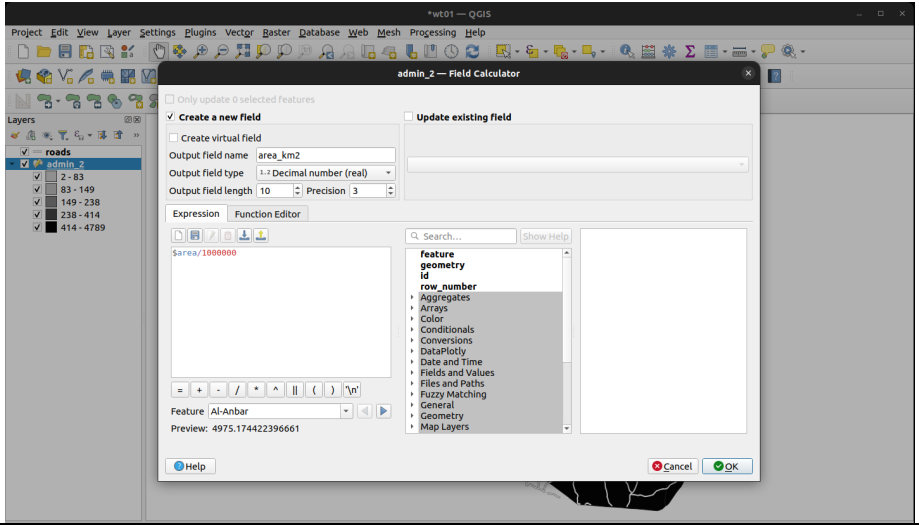

Open the `Field Calculator` again. Create a new field called `road_density` of type `Decimal number (real)`. For `Expression`, write `roads_km/area_km2`

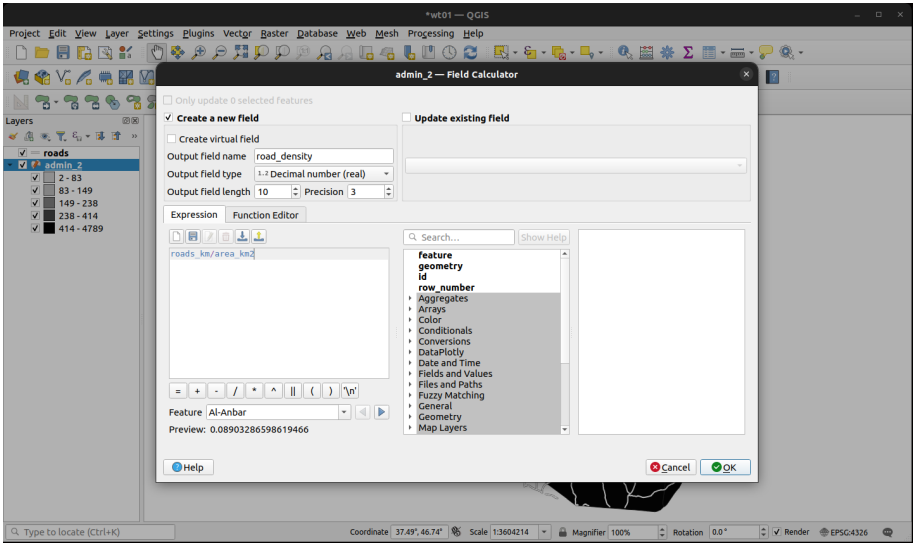

Check to make sure `road_density` is added to attribute table

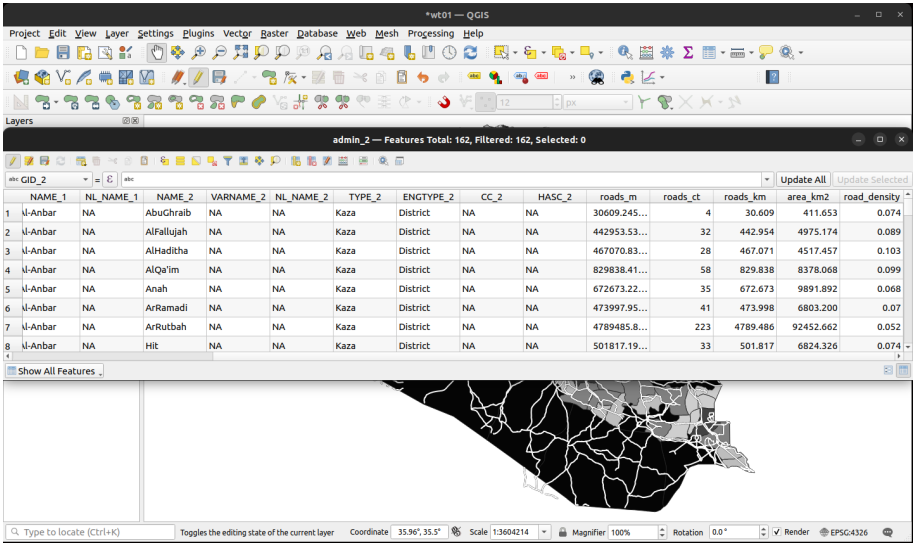

	NAME_1	NL_NAME_1	NAME_2	VARNAME_2	NL_NAME_2	TYPE_2	ENGTYPE_2	CC_2	HASC_2	roads_m	roads_ct	roads_km	area_km2	road_density
1	al-Anbar	NA	AbuGhraib	NA	NA	Kaza	District	NA	NA	30609.245…	4	30.609	411.653	0.074
2	al-Anbar	NA	AlFallujah	NA	NA	Kaza	District	NA	NA	442953.53…	32	442.954	4975.174	0.089
3	al-Anbar	NA	AlHaditha	NA	NA	Kaza	District	NA	NA	467070.83…	28	467.071	4517.457	0.103
4	al-Anbar	NA	AlQa'im	NA	NA	Kaza	District	NA	NA	829838.41…	58	829.838	8378.068	0.099
5	al-Anbar	NA	Anah	NA	NA	Kaza	District	NA	NA	672673.22…	35	672.673	9891.892	0.068
6	al-Anbar	NA	ArRamadi	NA	NA	Kaza	District	NA	NA	473997.95…	41	473.998	6803.200	0.07
7	al-Anbar	NA	ArRutbah	NA	NA	Kaza	District	NA	NA	4789485.8…	223	4789.486	92452.662	0.052
8	al-Anbar	NA	Hit	NA	NA	Kaza	District	NA	NA	501817.19…	33	501.817	6824.326	0.074

If we plot this variable, we see the opposite pattern from before (smaller, urban districts have higher road density)

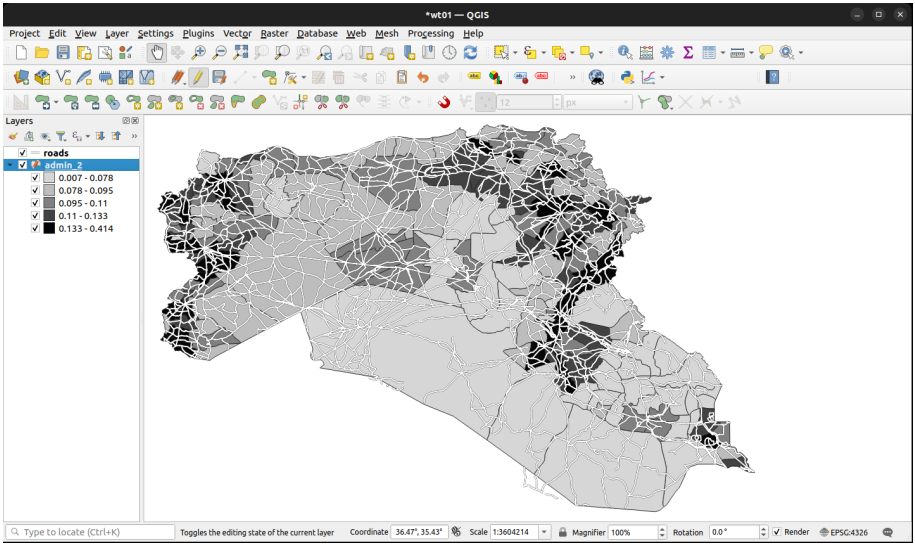

Save the edits you just made to the `admin_2` layer (right click → `Save Layer Edits`)

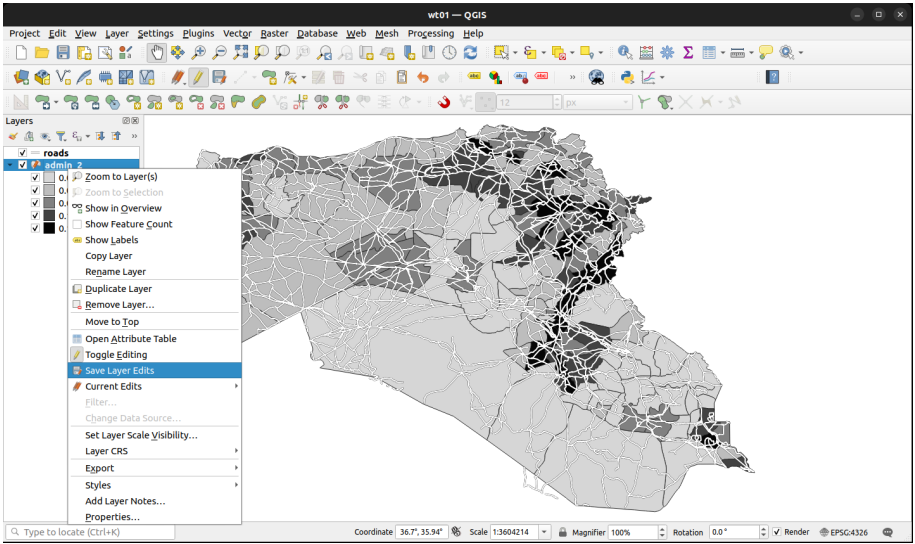

Load *population* raster (`Layer` → `Add Layer` → `Add Raster Layer`).\ Open the `gpw_v4_population_count_rev11_2010_15_min.tif` file in `Data/GPW` folder.

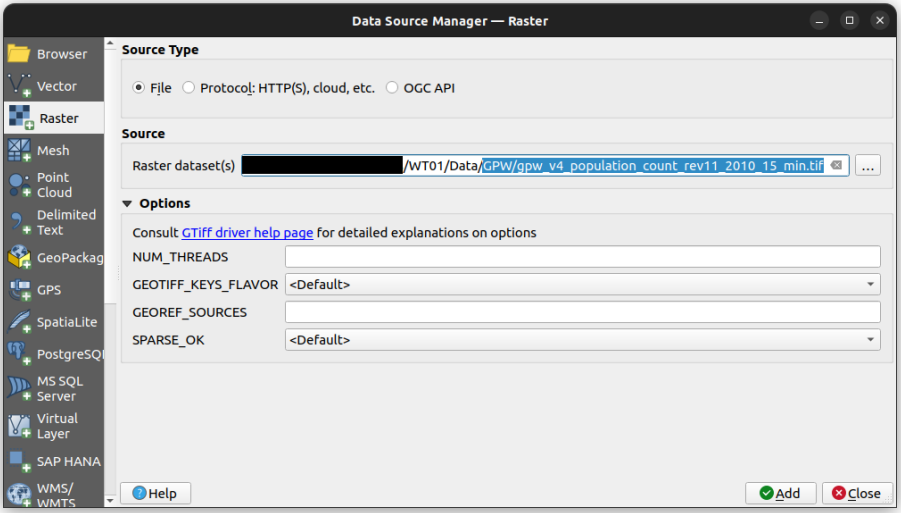

The (global) population layer should now be visible. Let's **extract a subset** of just the part of this layer that overlaps with Syria and Iraq.

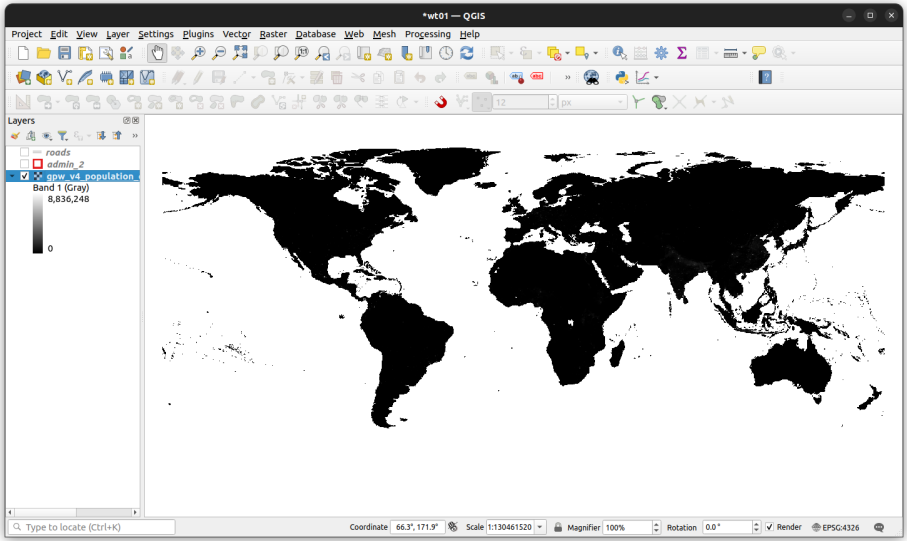

Open the `Clip Raster by Extent` tool (`Raster` → `Extension` → `Clip Raster by Extent`)

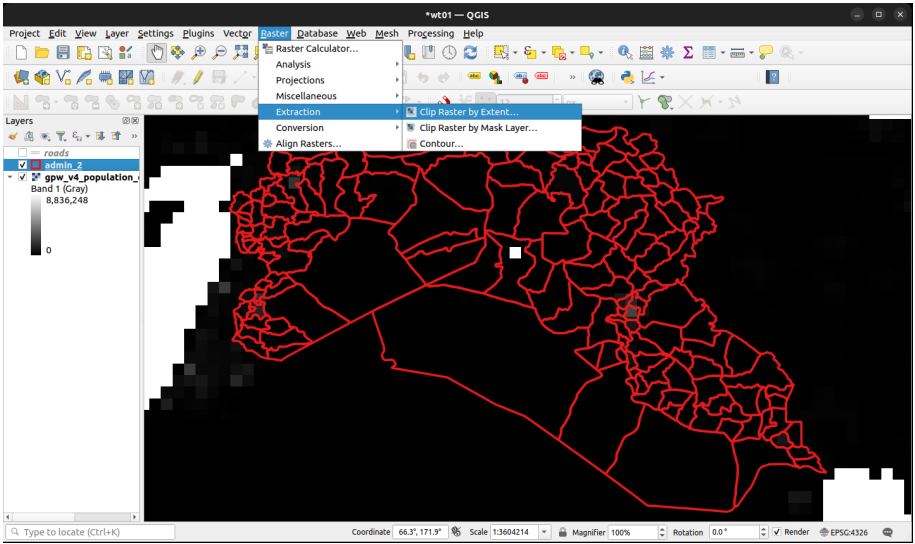

Set `Input layer` = `gpw_v4_population_count_rev11_2010_15_min`. Click the `[...]` button next to `Clipping extent` → `Calculate from Layer` → `admin_2`

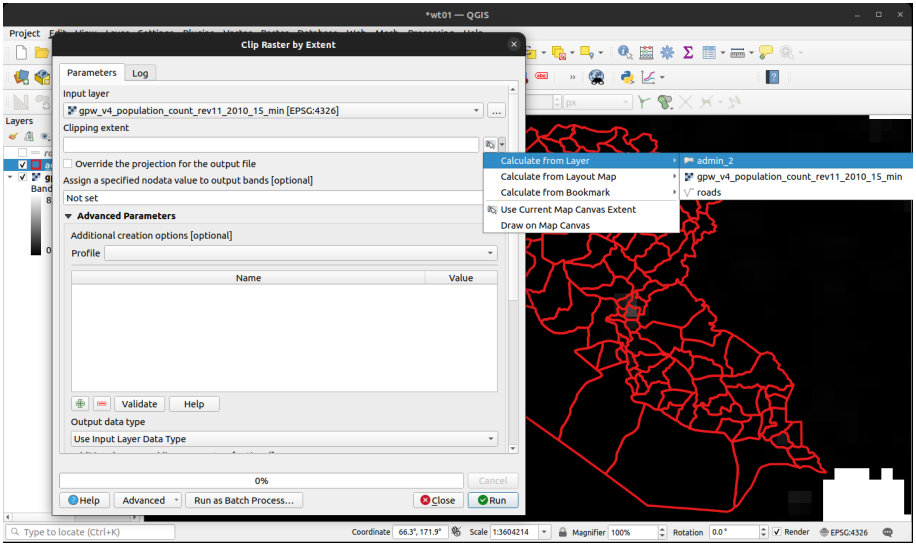

Save the file as `pop.tif`. Accept defaults for all other parameters, and click `Run`

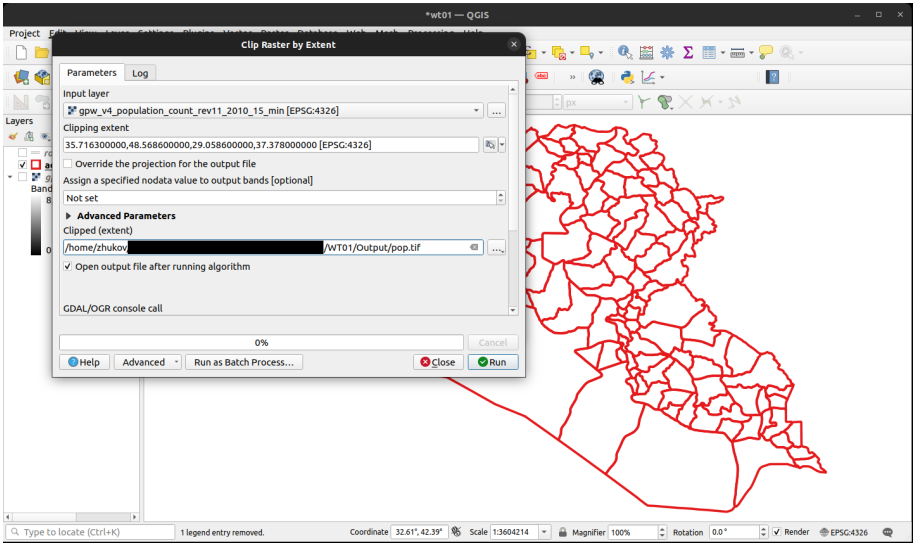

A new layer called pop should appear. You can remove the original, global gpw_v4... layer

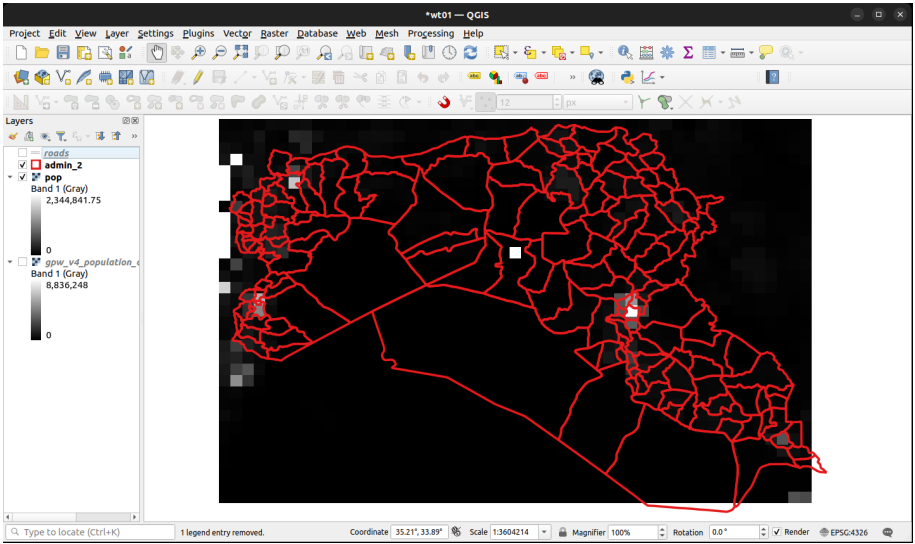

Let's **calculate the population size of each district**. Open the `Processing Toolbox` and open the `Zonal statistics` tool in the `Raster Analysis` menu

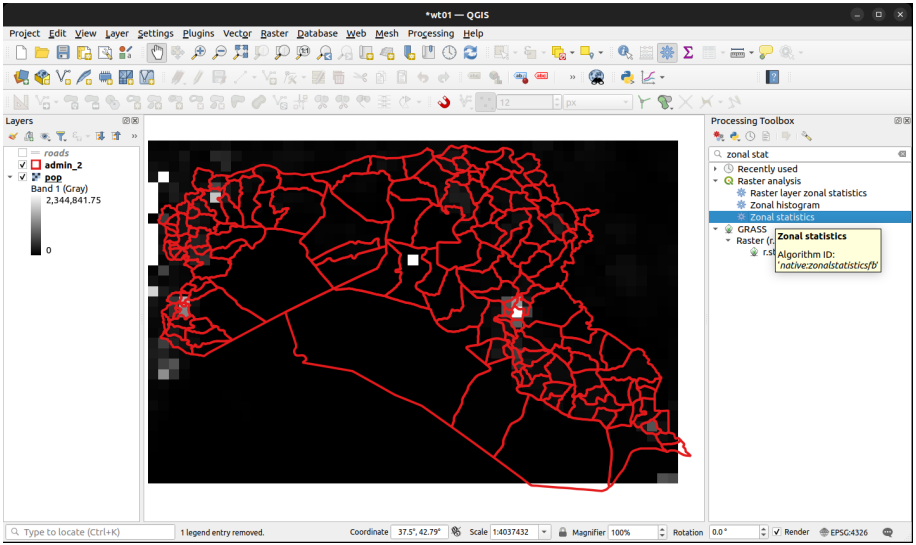

In the `Zonal Statistics` tool, set `Input layer=admin_2`, `Raster layer=pop`, and `Output column prefix=pop_`.
Click the `[...]` button next to `Statistics to calculate`

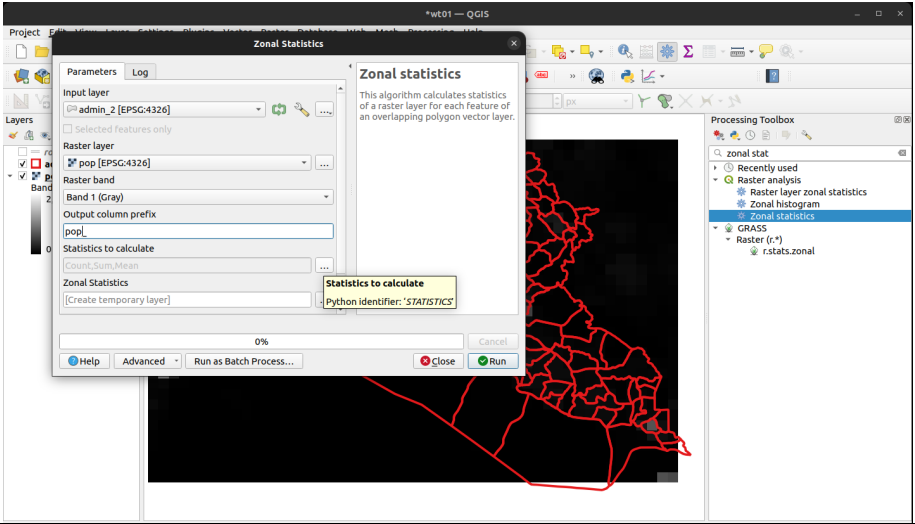

Check the box next to `Sum`. Leave all others un-checked. Click `OK`

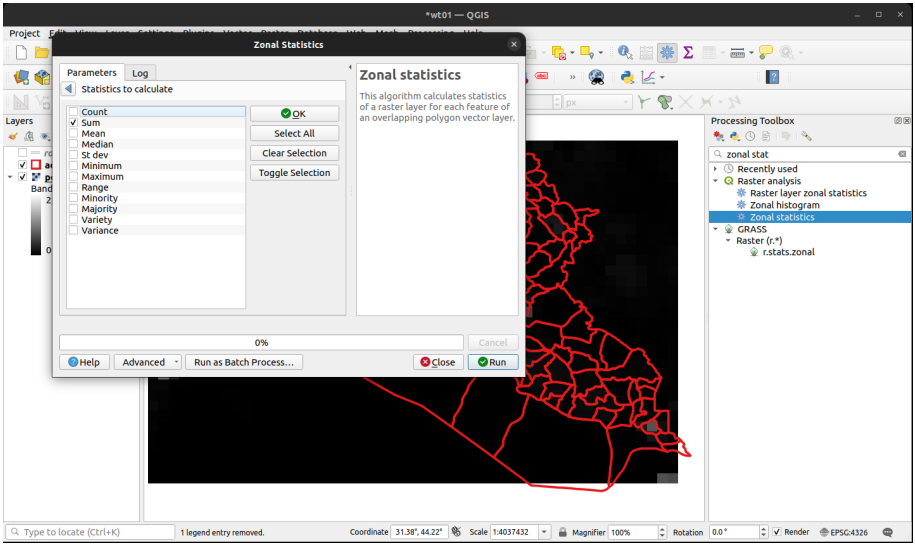

Save the zonal statistics output as `admin_3.geojson`. Click `Run`

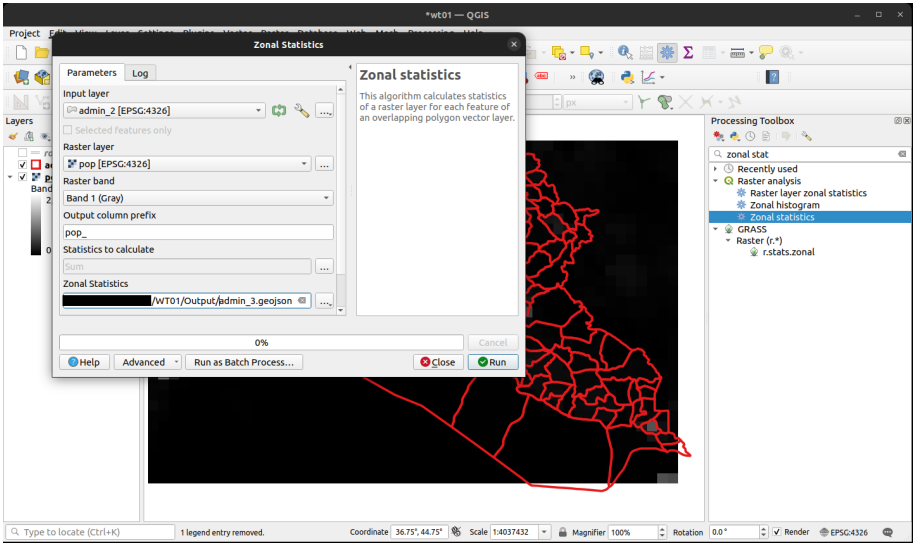

This should have generated a new layer, `admin_3`, with a new field, `pop_sum` (check the layer's attribute table)

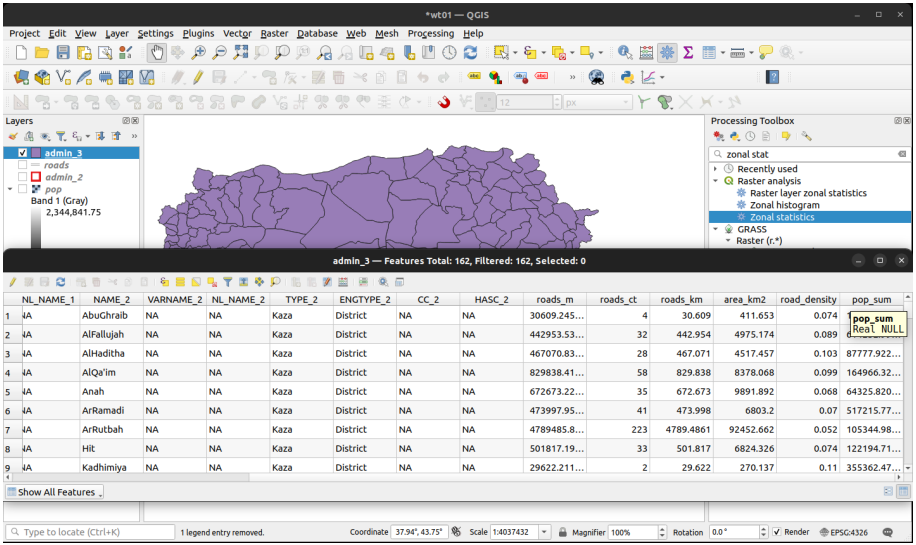

	NL_NAME_1	NAME_2	VARNAME_2	NL_NAME_2	TYPE_2	ENGTYPE_2	CC_2	HASC_2	roads_m	roads_ct	roads_km	area_km2	road_density	pop_sum
1	NA	AbuGhraib	NA	NA	Kaza	District	NA	NA	30609.245...	4	30.609	411.653	0.074	
2	NA	AlFallujah	NA	NA	Kaza	District	NA	NA	442953.53...	32	442.954	4975.174	0.089	
3	NA	AlHaditha	NA	NA	Kaza	District	NA	NA	467070.83...	28	467.071	4517.457	0.103	87777.922...
4	NA	AlQa'im	NA	NA	Kaza	District	NA	NA	829838.41...	58	829.838	8378.068	0.099	164966.32...
5	NA	Anah	NA	NA	Kaza	District	NA	NA	672673.22...	35	672.673	9891.892	0.068	64325.820...
6	NA	ArRamadi	NA	NA	Kaza	District	NA	NA	473997.95...	41	473.998	6803.2	0.07	517215.77...
7	NA	ArRutbah	NA	NA	Kaza	District	NA	NA	4789485.8...	223	4789.4861	92452.662	0.052	105344.98...
8	NA	Hit	NA	NA	Kaza	District	NA	NA	501817.19...	33	501.817	6824.326	0.074	122194.71...
9	NA	Kadhimiya	NA	NA	Kaza	District	NA	NA	29622.211...	2	29.622	270.137	0.11	355362.47...

Load *cropland* raster (`Layer` → `Add Layer` → `Add Raster Layer`).
Open the `cropland.tif` file in `Data/Cropland` folder.

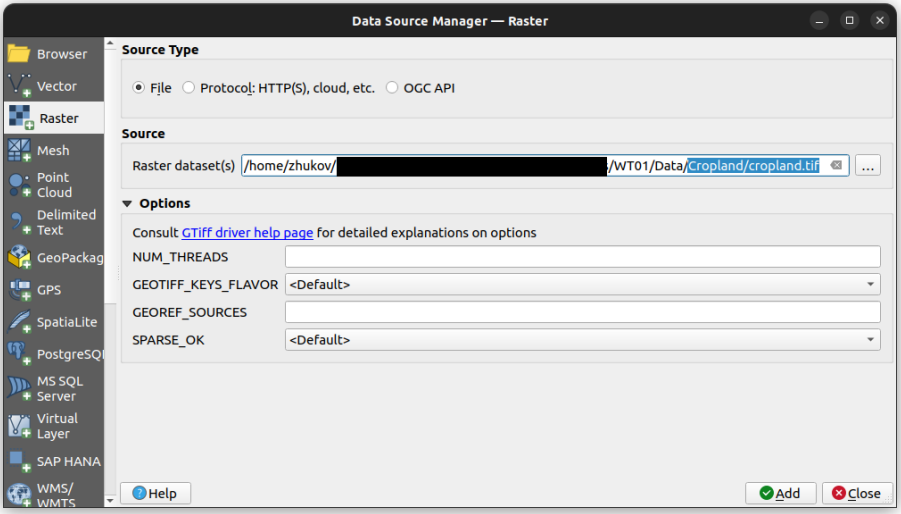

The (global) crops layer should now be visible. Let's **extract a subset** of it, as we just did for population

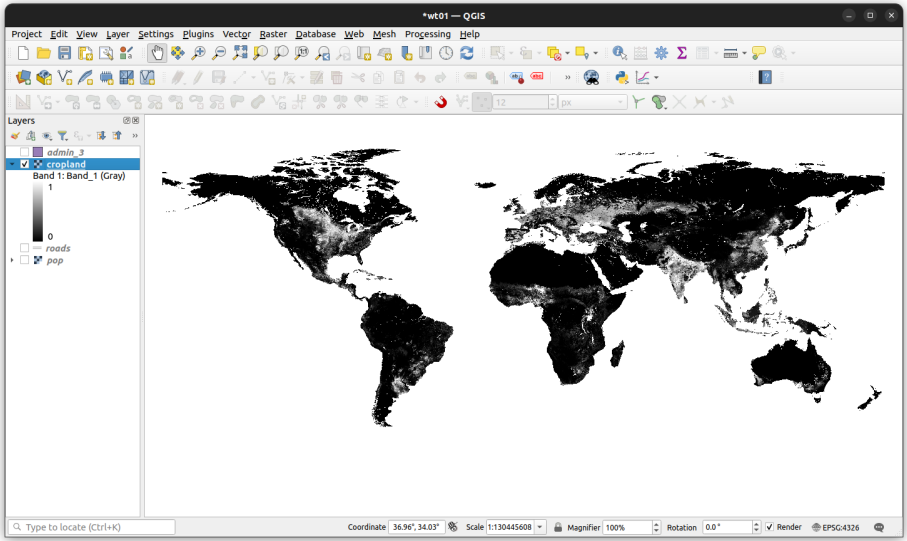

Open the `Clip Raster by Extent` tool (`Raster` → `Extension` → `Clip Raster by Extent`). Set `Input layer` = `cropland`. Click the `[...]` button next to `Clipping extent` → `Calculate from Layer` → `admin_3`

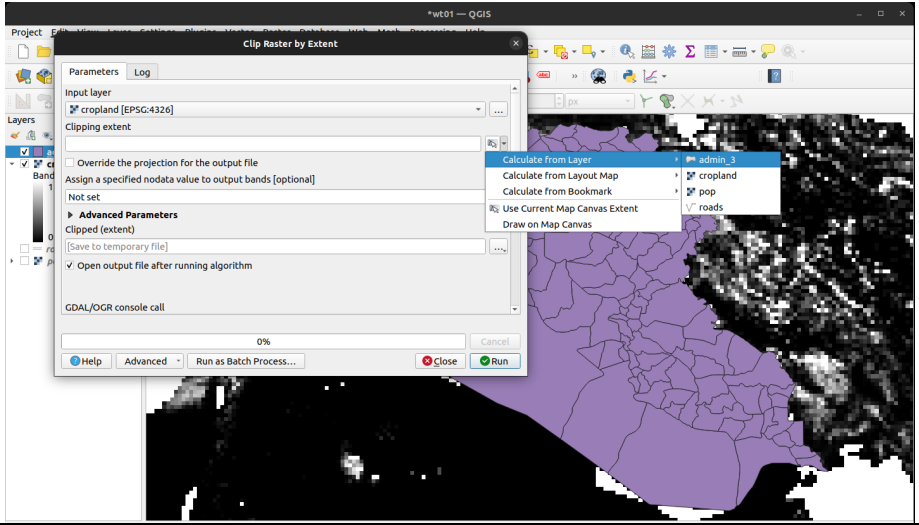

Save the output to a file called `crops.tif`

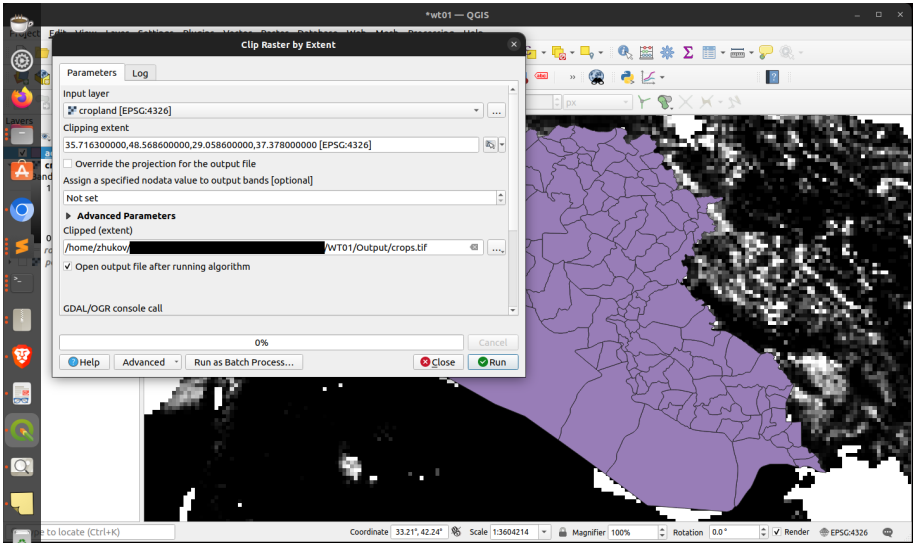

Make sure everything looks right. Remove the original, global `cropland` layer

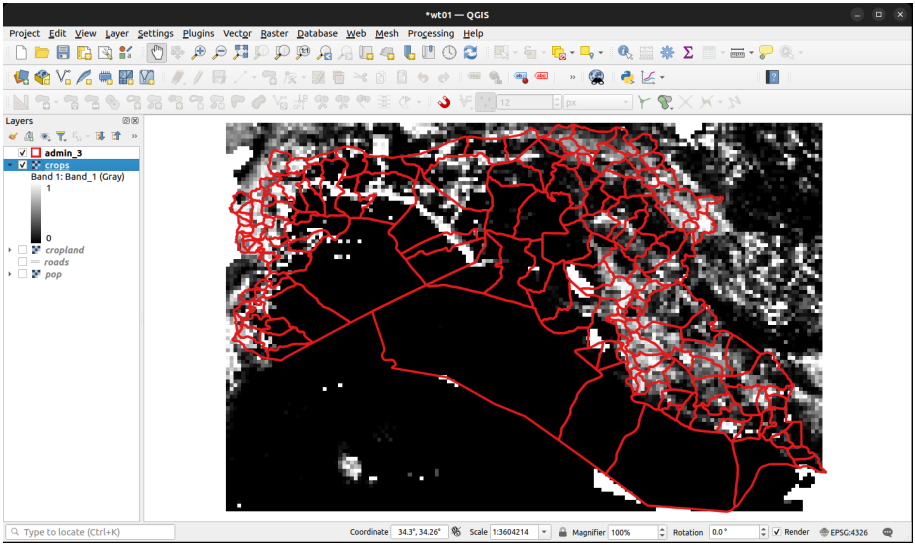

Let's **calculate crop cultivation in each district**. Open the `Zonal statistics` tool in `Processing Toolbox`. Set `Input layer`=`admin_3`, `Raster layer`=`crops`, and `Output column prefix`=`crops_`.
Click the `[...]` button next to `Statistics to calculate`

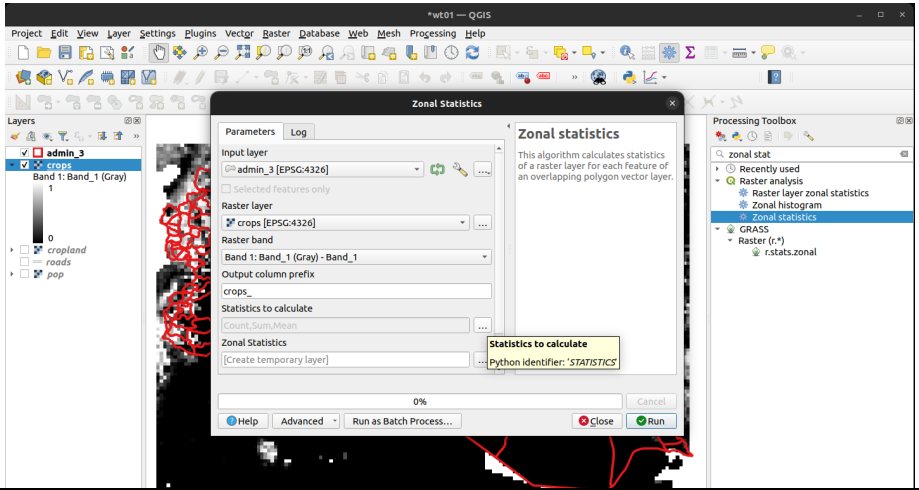

Check the box next to `Mean`. Leave all others un-checked. Click `OK`, then `Run` on the next screen

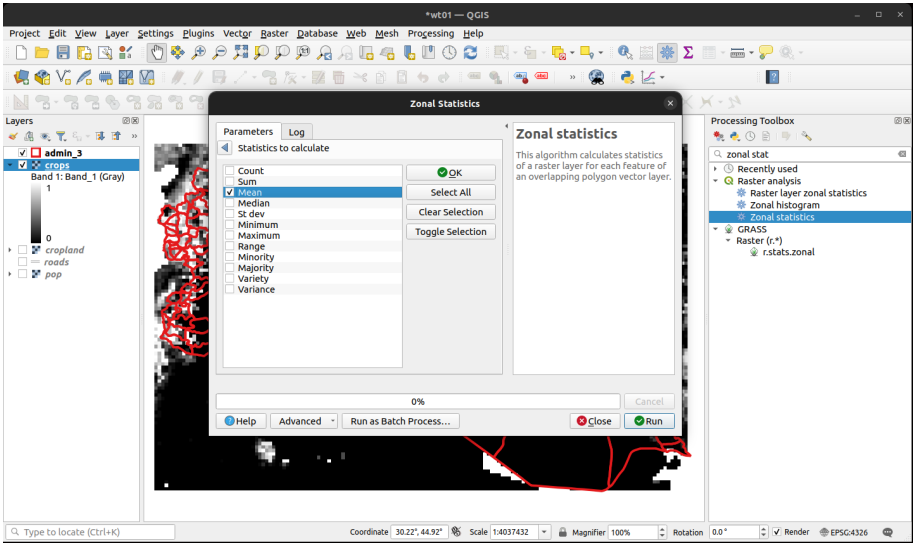

Save the output to a file called `admin_4.geojson`

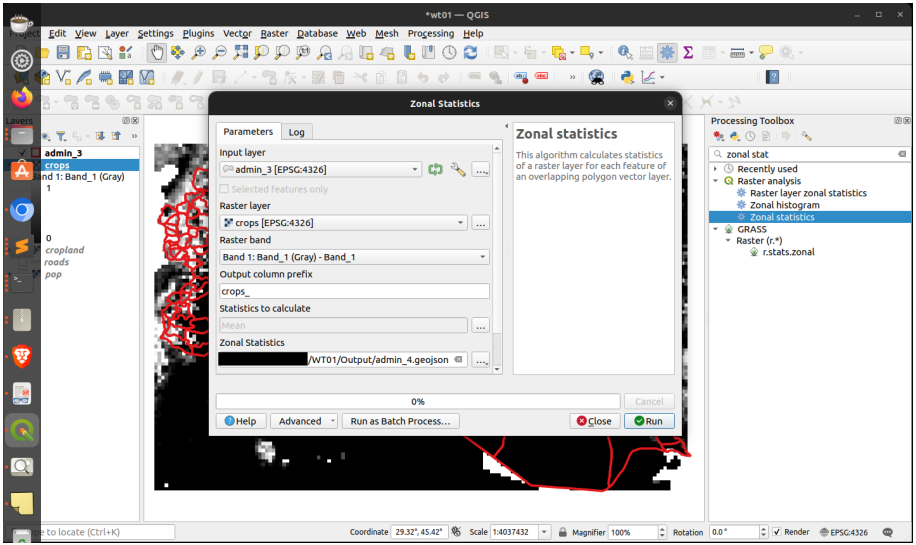

This should have generated a new field, `crops_mean` in the new `admin_4` layer (check the layer's attribute table)

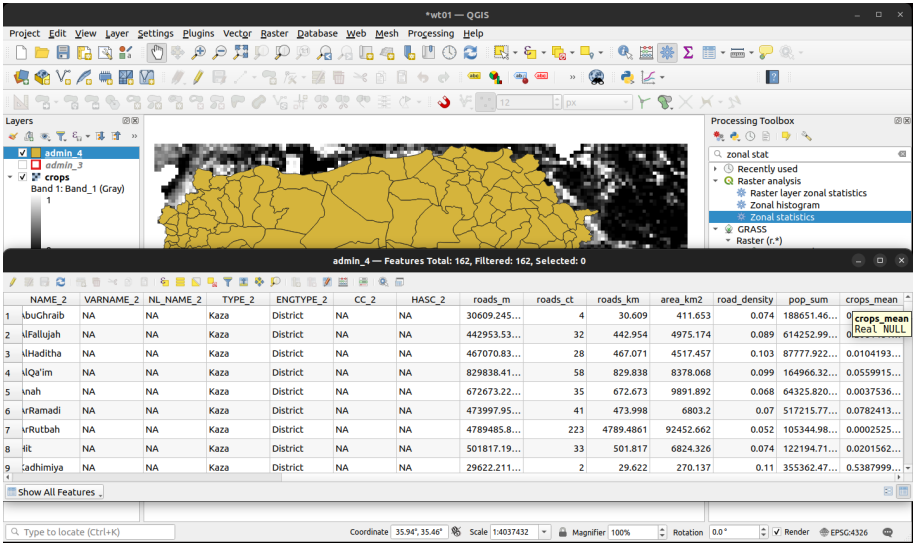

	NAME_2	VARNAME_2	NL_NAME_2	TYPE_2	ENGTYPE_2	CC_2	HASC_2	roads_m	roads_ct	roads_km	area_km2	road_density	pop_sum	crops_mean
1	AbuGhraib	NA	NA	Kaza	District	NA	NA	30609.245...	4	30.609	411.653	0.074	188651.46...	0
2	alFallujah	NA	NA	Kaza	District	NA	NA	442953.53...	32	442.954	4975.174	0.089	614252.99...	0
3	alHaditha	NA	NA	Kaza	District	NA	NA	467070.83...	28	467.071	4517.457	0.103	87777.922...	0.0104193...
4	alQa'im	NA	NA	Kaza	District	NA	NA	829838.41...	58	829.838	8378.068	0.099	164966.32...	0.0559915...
5	Anah	NA	NA	Kaza	District	NA	NA	672673.22...	35	672.673	9891.892	0.068	64325.820...	0.0037536...
6	arRamadi	NA	NA	Kaza	District	NA	NA	473997.95...	41	473.998	6803.2	0.07	517215.77...	0.0782413...
7	arRutbah	NA	NA	Kaza	District	NA	NA	4789485.8...	223	4789.4861	92452.662	0.052	105344.98...	0.0002525...
8	Hit	NA	NA	Kaza	District	NA	NA	501817.19...	33	501.817	6824.326	0.074	122194.71...	0.0201562...
9	Kadhimiya	NA	NA	Kaza	District	NA	NA	29622.211...	2	29.622	270.137	0.11	355362.47...	0.5387999...

Load *hydroelectric dam* locations (`Layer` → `Add Layer` → `Add Vector Layer`). Open the `GRanD_dams_v1_1.shp` file in `Data/Dams` folder.

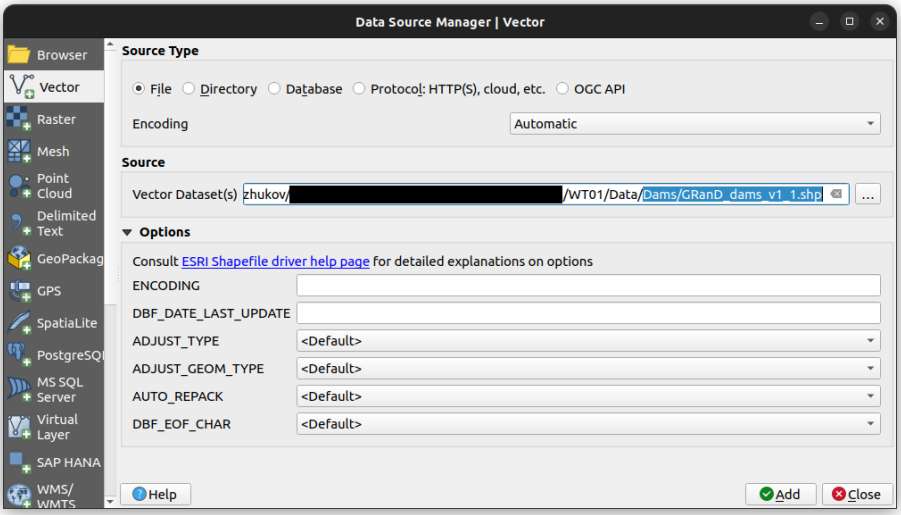

The dam locations should appear on the project window. Rename it from `GRanD_dams_v1_1` to `dams`

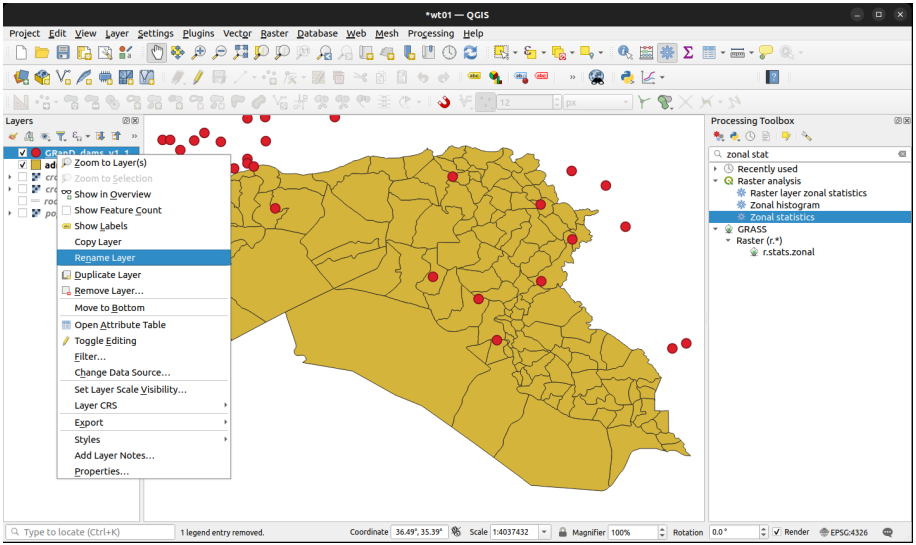

Let's calculate the **number of dams per district**. Open the `Count Points in Polygon` tool (`Vector` → `Analysis Tools` → `Count Points in Polygon`)

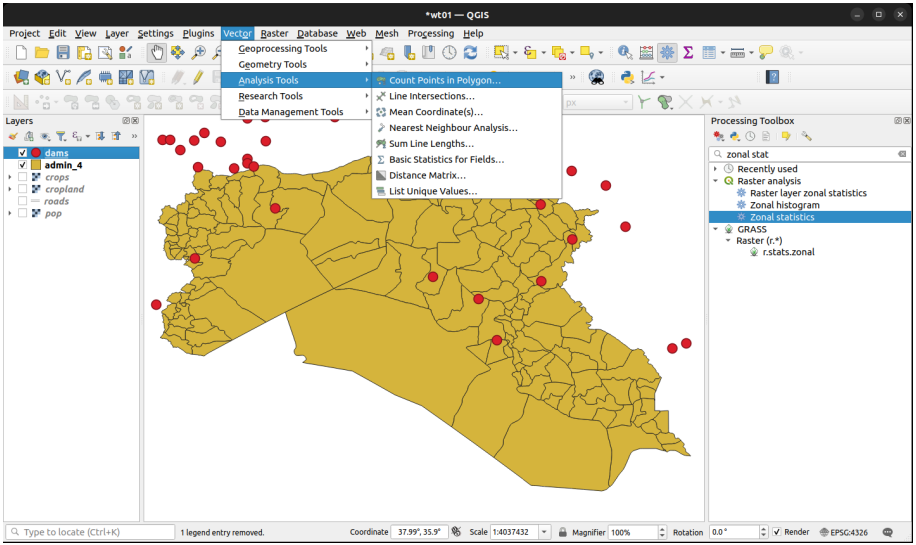

Select `Polygons` = `admin_4`, `Points` = `dams`. Name the count field `dams`, and save the output file as `admin_5.geojson`. Click `Run`

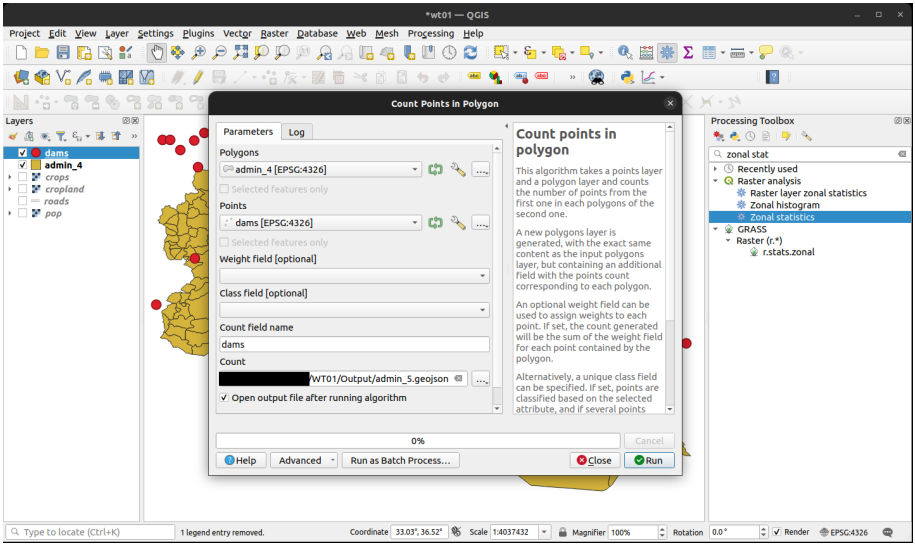

Add the *sectarian divisions* data to the project, using `Add Vector Layer...`\
Load the `GeoEPR-2021.geojson` file in `Data/GeoEPR` folder

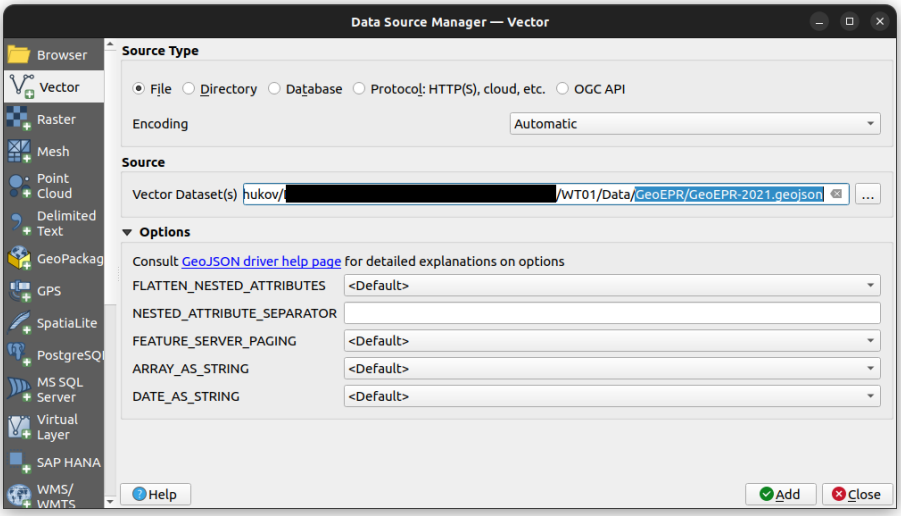

The (global) ethnicity layer should appear in the project window. Let's **calculate the proportion of each district populated by Sunni Arabs**. Let's first extract the subset that overlaps with our study area.

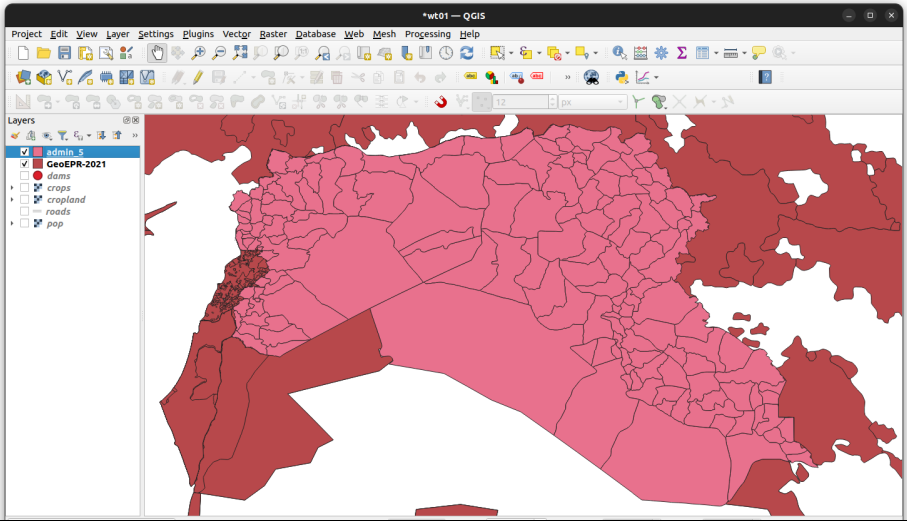

Open the `Clip` (vector) tool (`Vector` → `Geoprocessing Tools` → `Clip`)

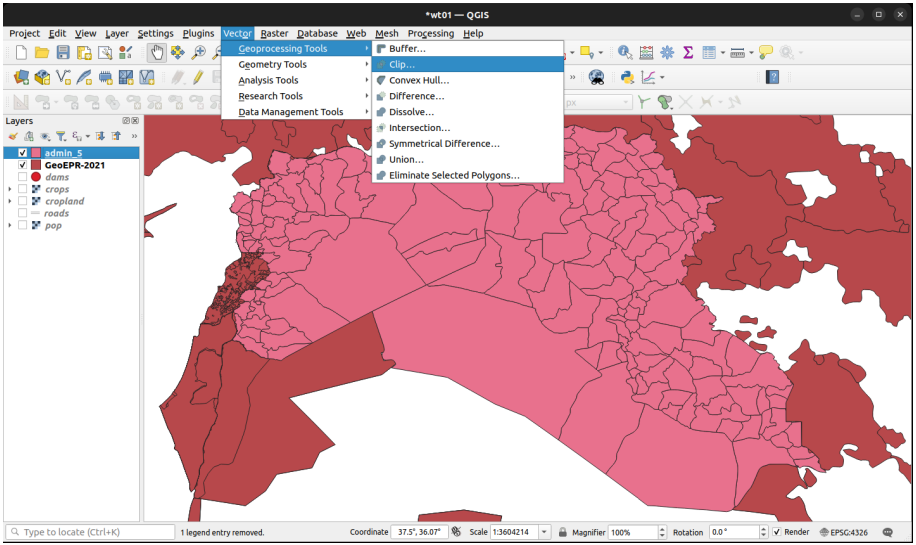

Set `Input layer` = `GeoEPR`, `Overlay layer` = `admin_5`. Save the file as `ethnicity.geojson`. Click `Run`

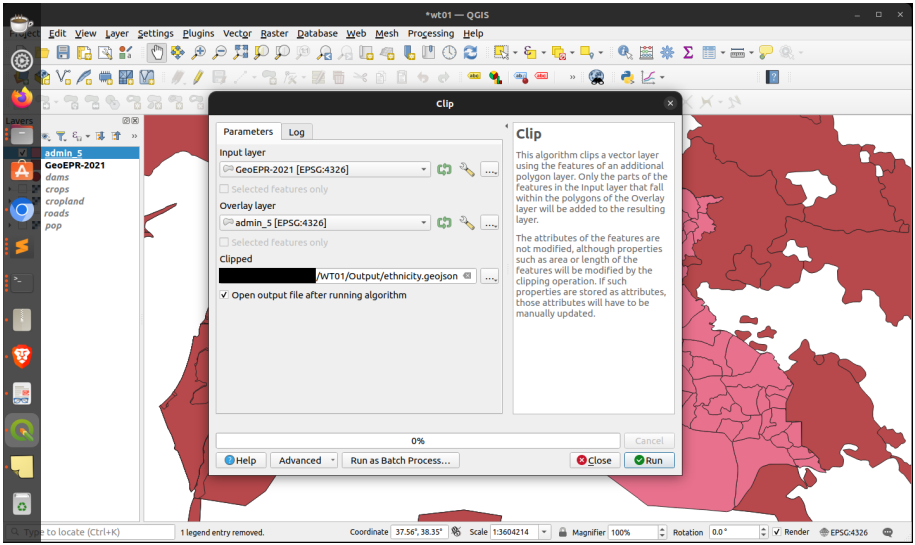

Uh oh! **Execution failed!** The error message says that `GeoEPR-2021` contains "invalid geometries". How do we fix this?

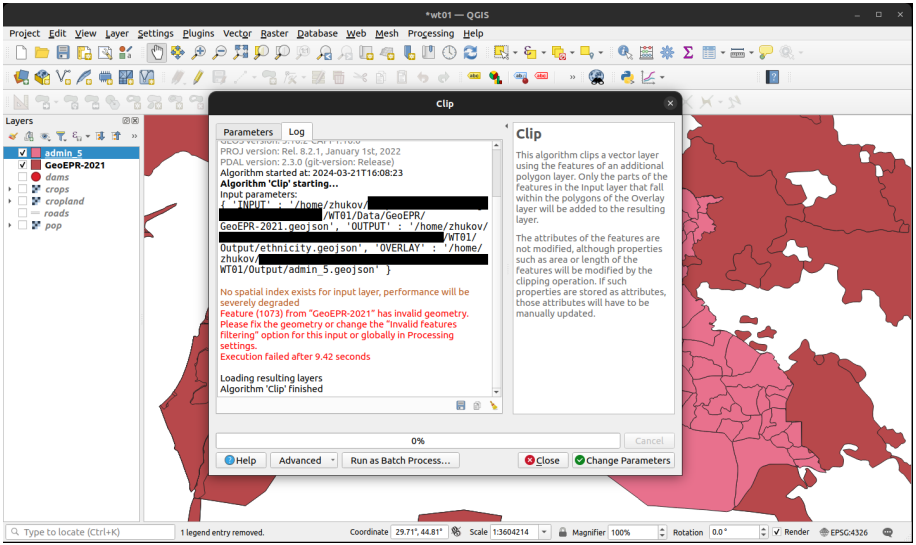

Many "invalid geometry" problems can be solved with the "zero buffer trick". Go to `Vector` → `Geoprocessing Tools` → `Buffer...`

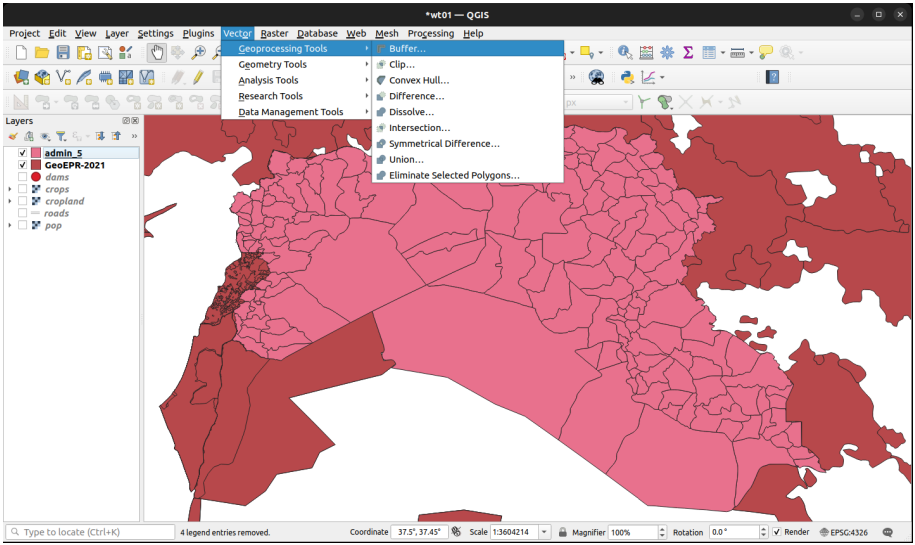

Set `Input layer` = `GeoEPR-2021` and `Distance` = `0`. Keep the defaults for the other parameters and save the file as `geoepr_0buffer.geojson`. Click `Run`

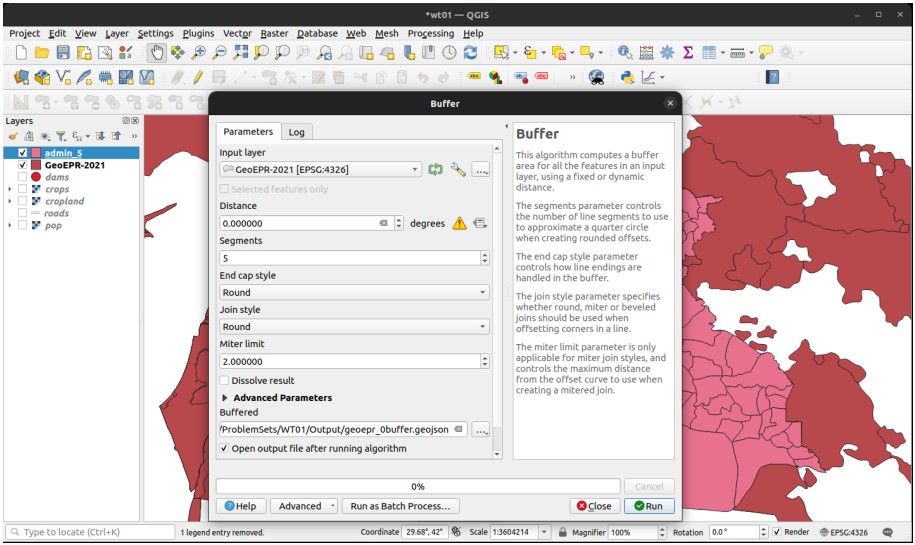

Now let's try the clipping operation again. To save time, go to `Processing` → `History...`

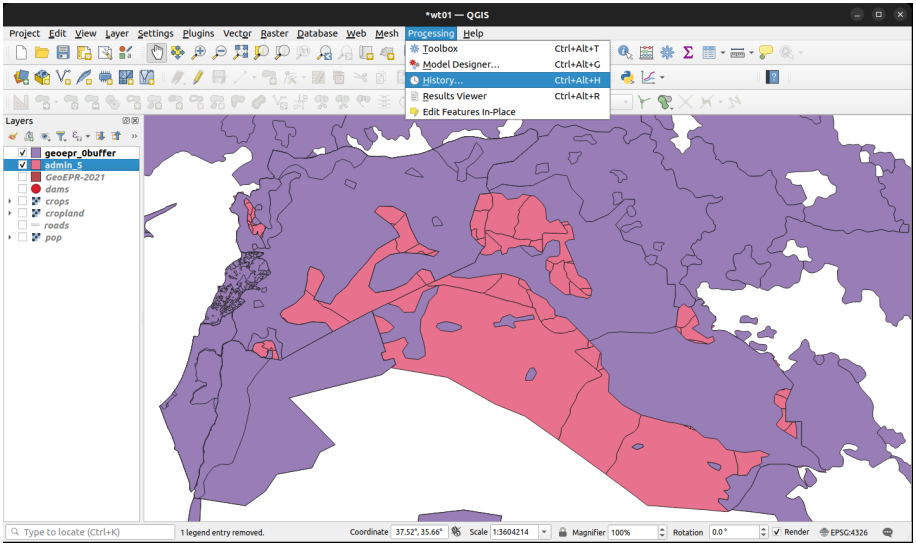

In the `Processing History` window, double-click on `Clip` (it should be the second-most recent item)

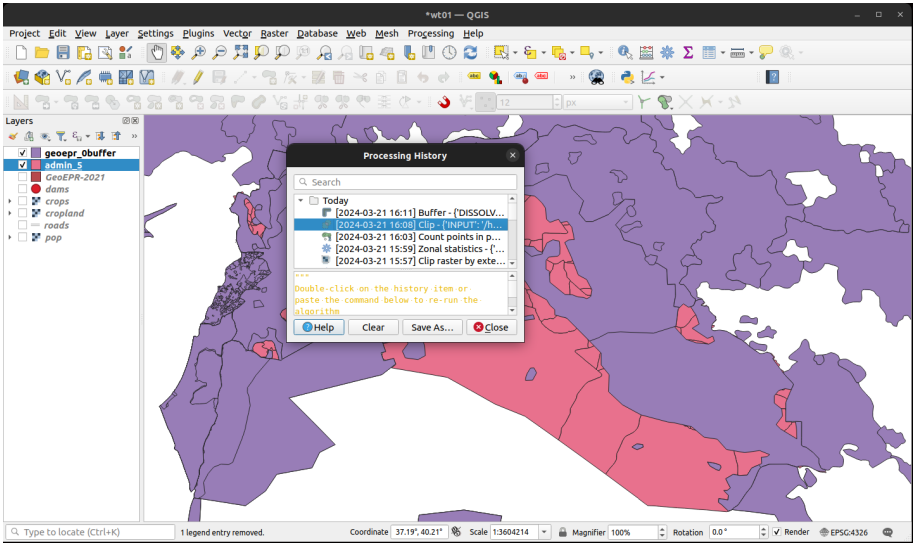

Set the `Input layer` = `geoepr_0buffer`. Click `Run`

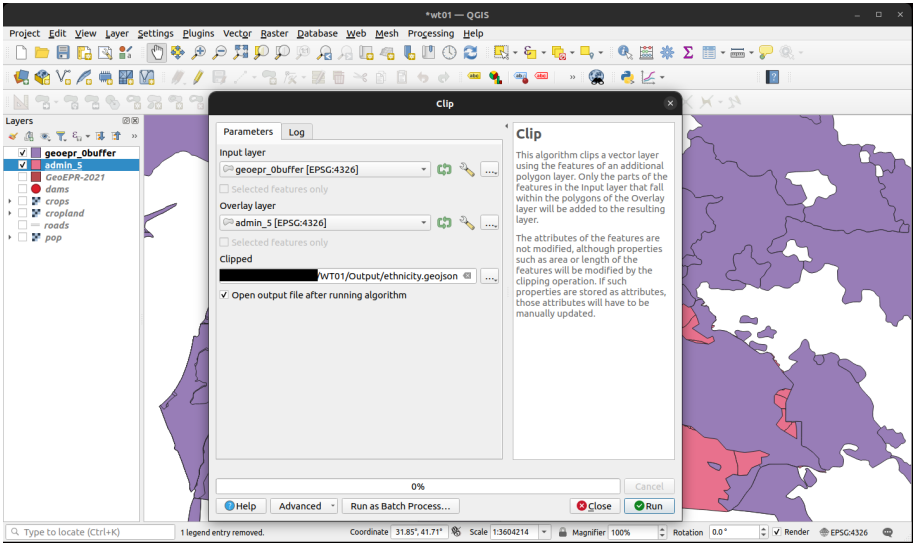

This may take a few minutes, but should finish without errors. . .

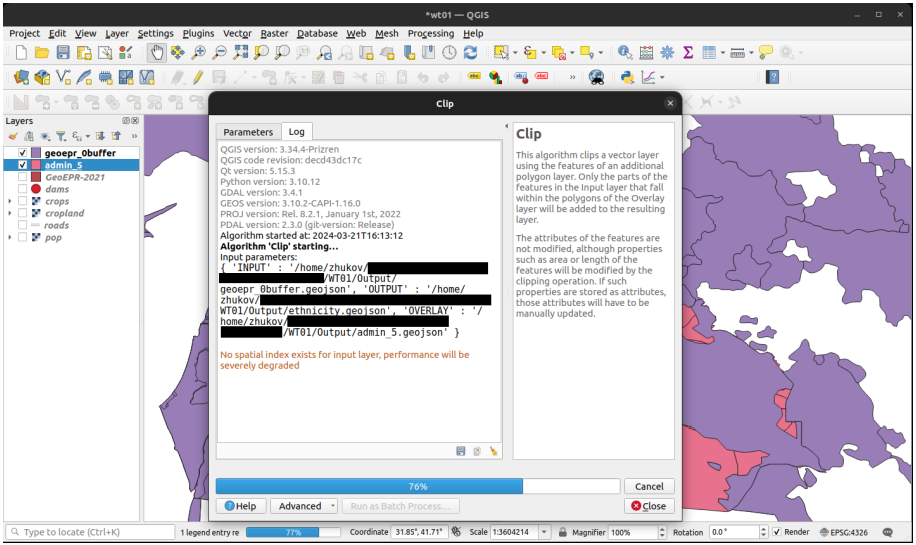

The clipped layer should appear in the project window. You can remove the full `GeoEPR` layer to save memory

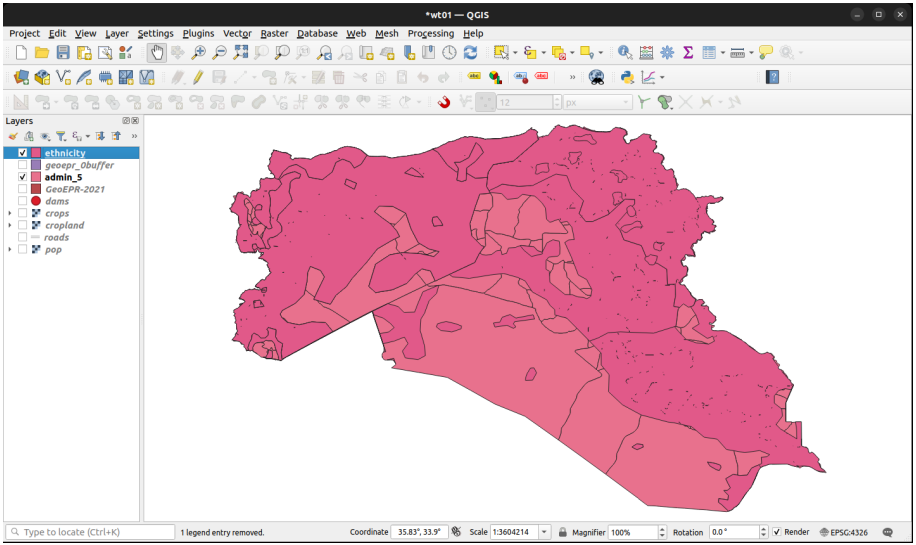

Look at the `group` field in the attribute table for `ethnicity`. There are multiple features for several groups (Kurds, Alawites, etc.). Let's **extract Sunni Arab polygons** only

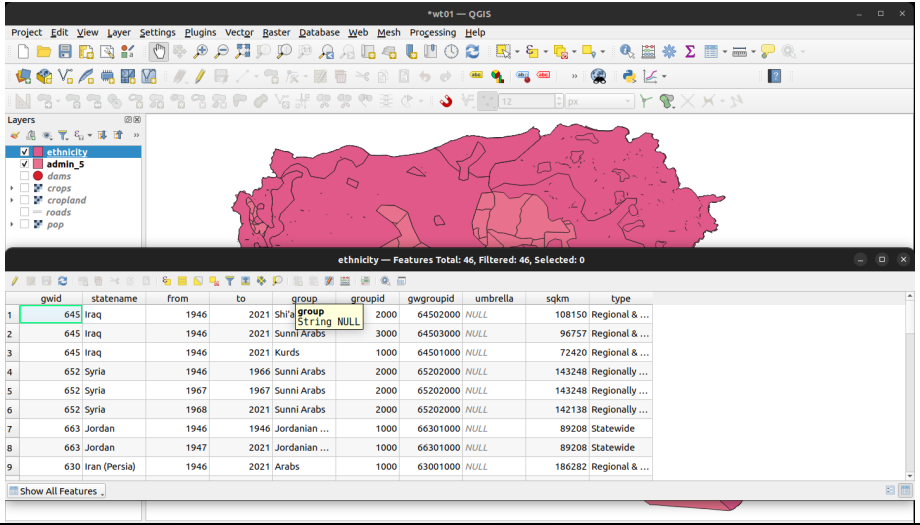

	gwid	statename	from	to	group	groupid	gwgroupid	umbrella	sqkm	type
1	645	Iraq	1946	2021	Shi'a	2000	64502000	NULL	108150	Regional & ...
2	645	Iraq	1946	2021	Sunni Arabs	3000	64503000	NULL	96757	Regional & ...
3	645	Iraq	1946	2021	Kurds	1000	64501000	NULL	72420	Regional & ...
4	652	Syria	1946	1966	Sunni Arabs	2000	65202000	NULL	143248	Regionally ...
5	652	Syria	1967	1967	Sunni Arabs	2000	65202000	NULL	143248	Regionally ...
6	652	Syria	1968	2021	Sunni Arabs	2000	65202000	NULL	142138	Regionally ...
7	663	Jordan	1946	1946	Jordanian ...	1000	66301000	NULL	89208	Statewide
8	663	Jordan	1947	2021	Jordanian ...	1000	66301000	NULL	89208	Statewide
9	630	Iran (Persia)	1946	2021	Arabs	1000	63001000	NULL	186282	Regional & ...

Open the `Select by Expression` tool (`Edit` menu → `Select` → `Select Features by Expression`)

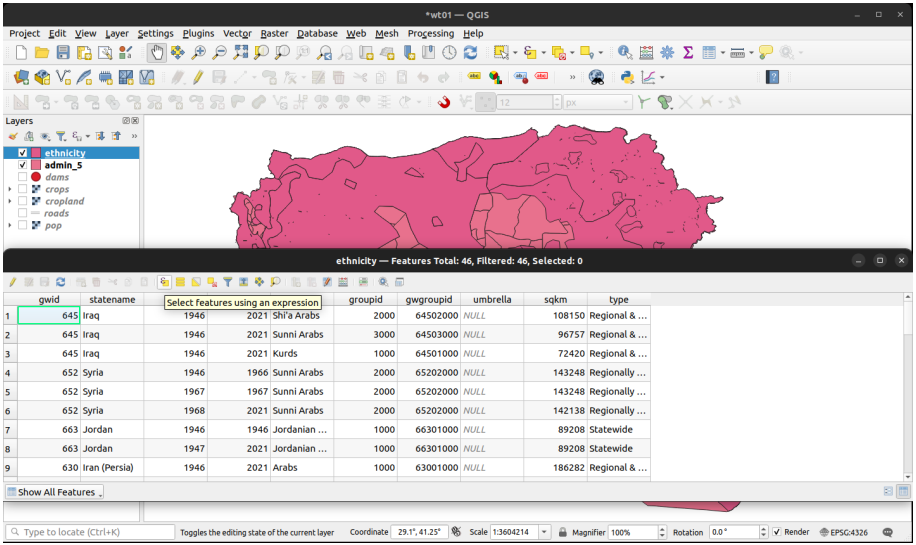

Let's use **regular expressions** to extract the polygons we need. Set `Expression = regexp_match(group,'Sunni')`. Click `Select Features`

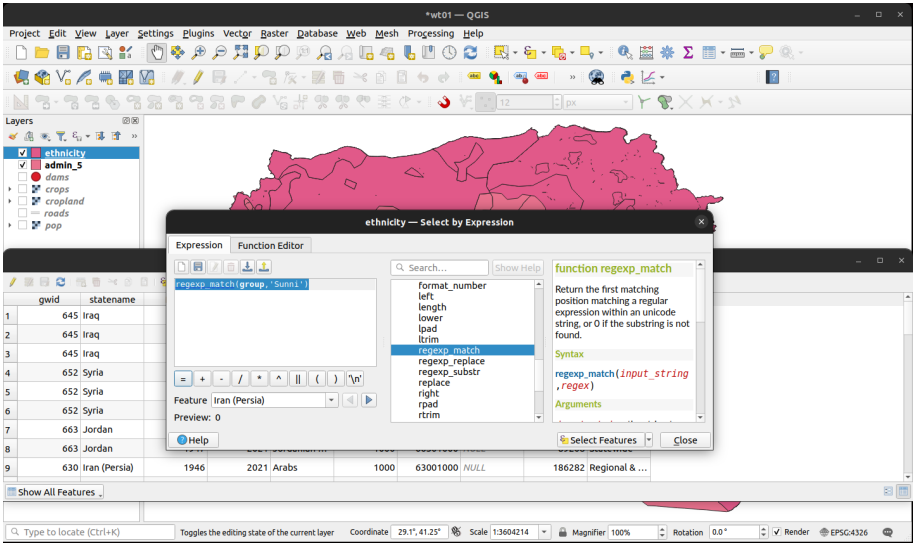

The selected polygons should appear yellow on the map. Let's **extract the selection into a new layer**. Right-click `ethnicity` in the Layer menu, go to `Export→Save Selected Features As`

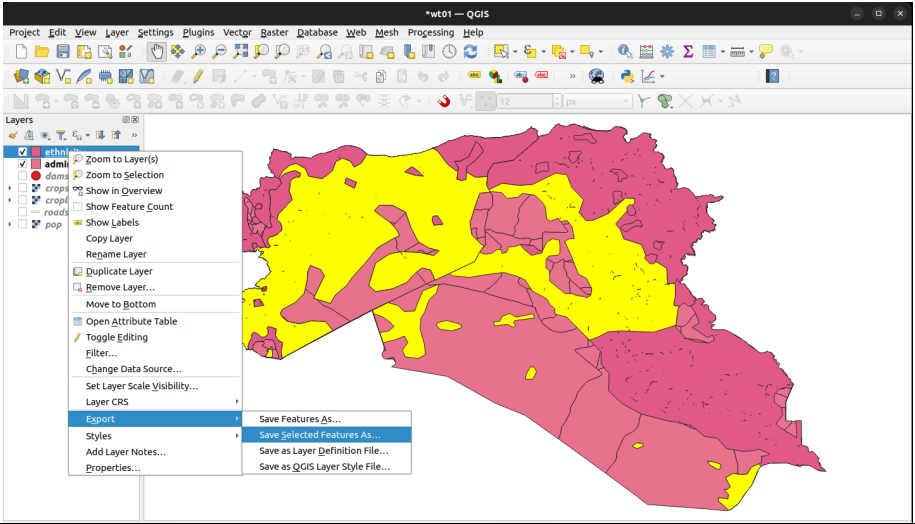

Save the layer as `sunni.geojson`. Make sure the box is checked next to `Save only selected features`. Click `OK`

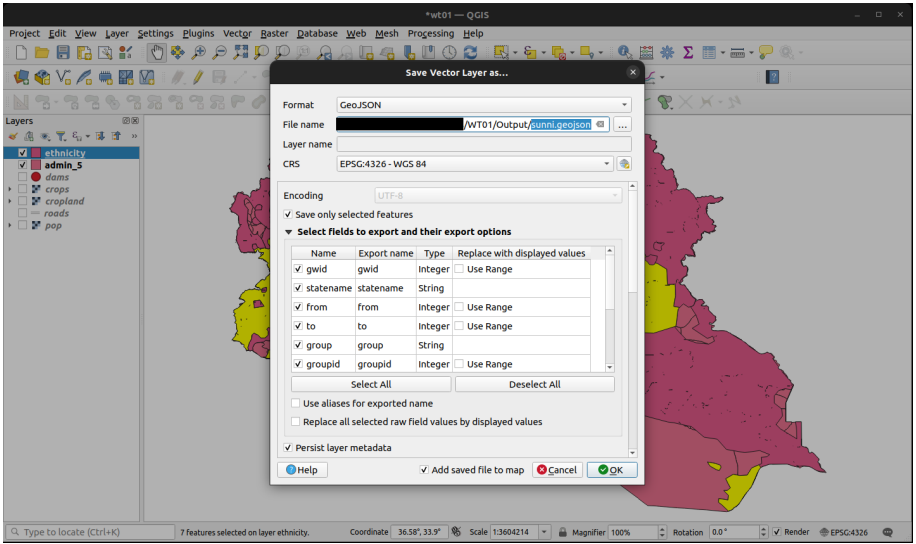

The extracted selection should appear. Hide the previous layer `ethnicity`

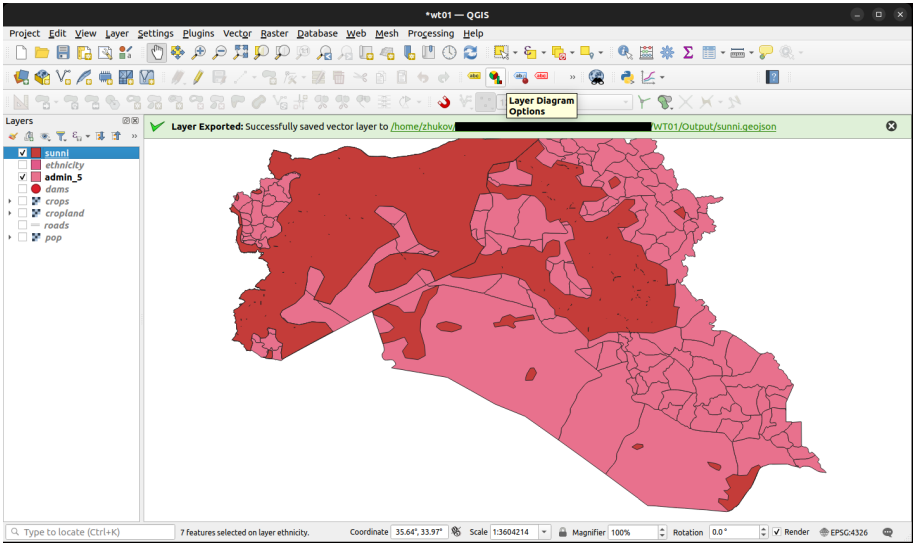

The attribute table for `sunni` should now contain only features where the `group` field contains the word `Sunni`. Let's **consolidate these features into one polygon**

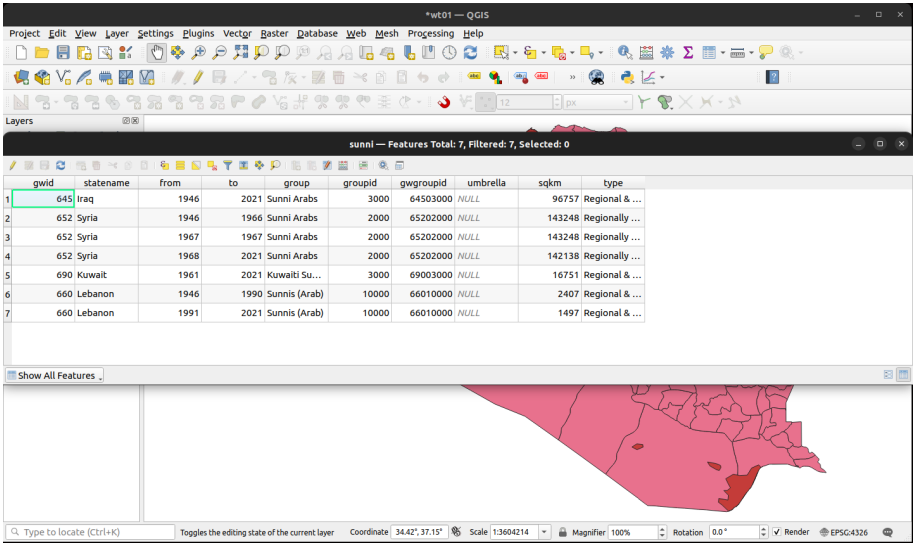

	gwid	statename	from	to	group	groupid	gwgroupid	umbrella	sqkm	type
1	645	Iraq	1946	2021	Sunni Arabs	3000	64503000	NULL	96757	Regionally ...
2	652	Syria	1946	1966	Sunni Arabs	2000	65202000	NULL	143248	Regionally ...
3	652	Syria	1967	1967	Sunni Arabs	2000	65202000	NULL	143248	Regionally ...
4	652	Syria	1968	2021	Sunni Arabs	2000	65202000	NULL	142138	Regionally ...
5	690	Kuwait	1961	2021	Kuwaiti Su...	3000	69003000	NULL	16751	Regional & ...
6	660	Lebanon	1946	1990	Sunnis (Arab)	10000	66010000	NULL	2407	Regional & ...
7	660	Lebanon	1991	2021	Sunnis (Arab)	10000	66010000	NULL	1497	Regional & ...

Open the `Dissolve` tool (`Vector` → `Geoprocessing Tools` → `Dissolve`).

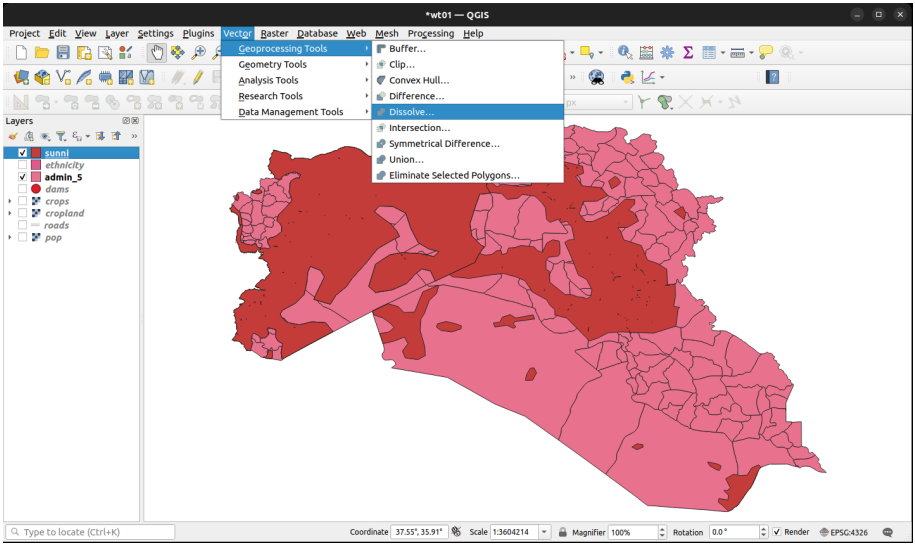

Set `Input layer=sunni`. Save the output as `sunni_2.geojson`

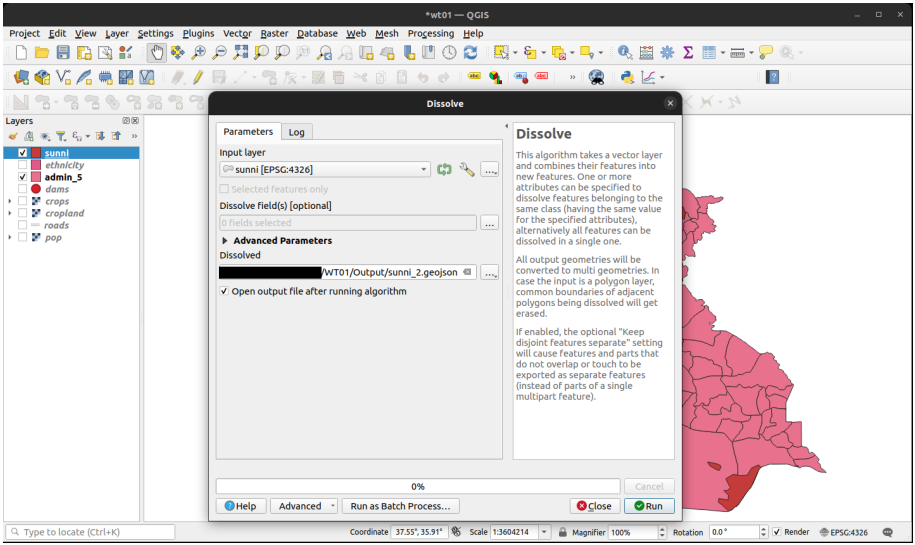

The dissolved polygon should appear in the project window. Now let's **calculate the proportion of each district populated by Sunnis**

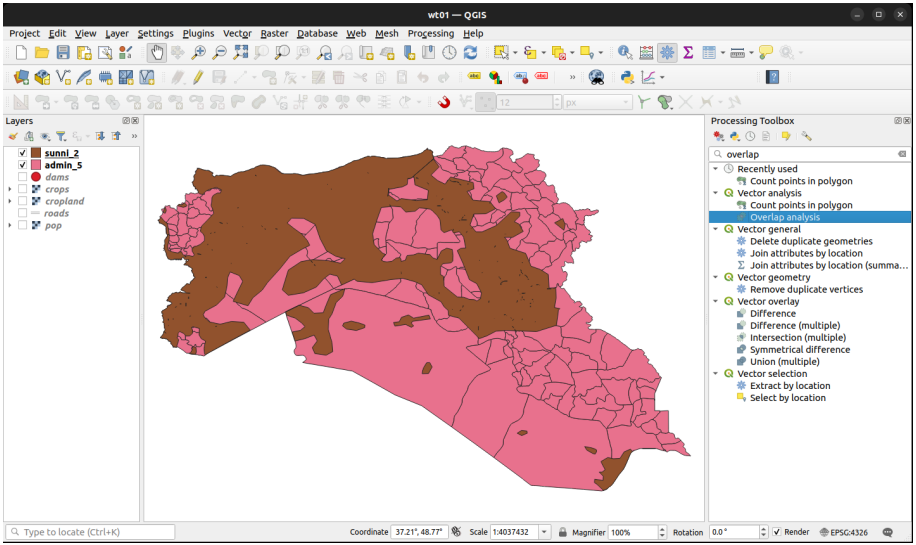

Open the `Overlap Analysis` tool (in `Processing Toolbox` → `Vector Analysis`). Set `Input layer` = `admin_5` and click the `[...]` button next to `Overlay layers`

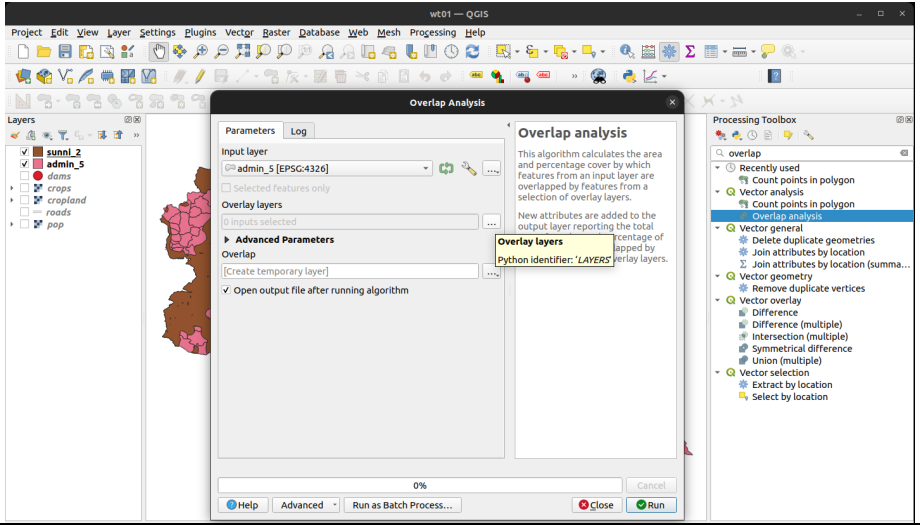

Check the box next to ✓`sunni_2`. Click `OK`

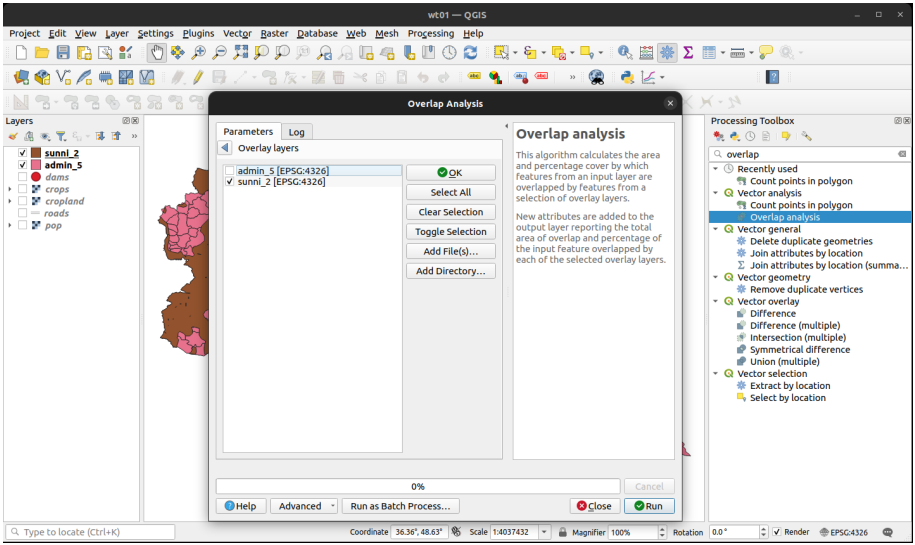

Save the output as `admin_6.geojson` and click `Run`

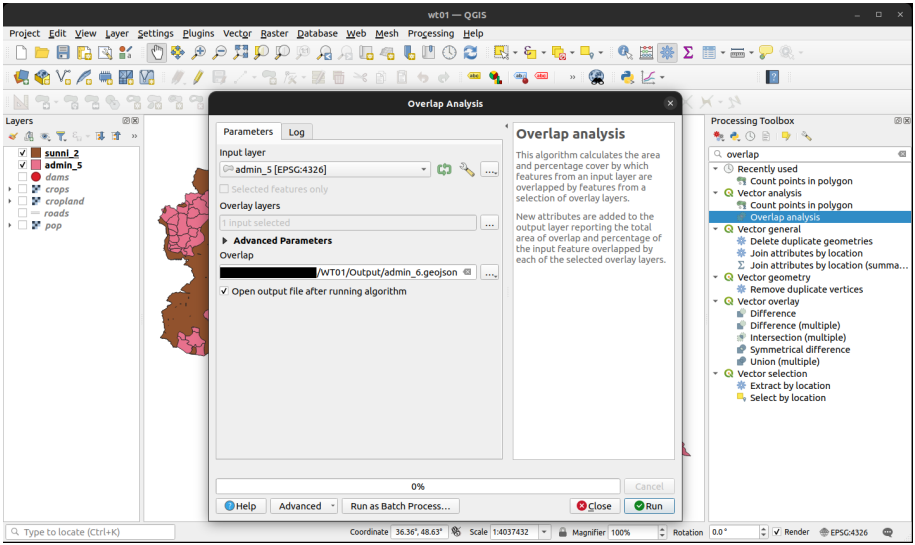

By default, the overlap fields in `admin_6` will be named `sunni_2_area` and `sunni_2_pc`. Let's **change the name of these fields**

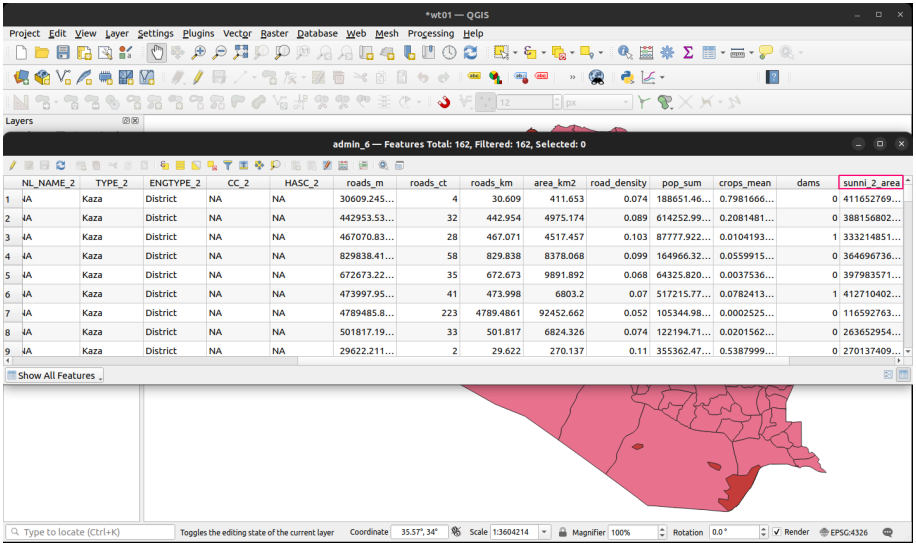

	NL_NAME_2	TYPE_2	ENGTYPE_2	CC_2	HASC_2	roads_m	roads_ct	roads_km	area_km2	road_density	pop_sum	crops_mean	dams	sunni_2_area
1	NA	Kaza	District	NA	NA	30609.245...	4	30.609	411.653	0.074	188651.46...	0.7981666...	0	411652769...
2	NA	Kaza	District	NA	NA	442953.53...	32	442.954	4975.174	0.089	614252.99...	0.2081481...	0	388156802...
3	NA	Kaza	District	NA	NA	467070.83...	28	467.071	4517.457	0.103	87777.922...	0.0104193...	1	333214851...
4	NA	Kaza	District	NA	NA	829838.41...	58	829.838	8378.068	0.099	164966.32...	0.0559915...	0	364696736...
5	NA	Kaza	District	NA	NA	672673.22...	35	672.673	9891.892	0.068	64325.820...	0.0037536...	0	397983571...
6	NA	Kaza	District	NA	NA	473997.95...	41	473.998	6803.2	0.07	517215.77...	0.0782413...	1	412710402...
7	NA	Kaza	District	NA	NA	4789485.8...	223	4789.4861	92452.662	0.052	105344.98...	0.0002525...	0	116592763...
8	NA	Kaza	District	NA	NA	501817.19...	33	501.817	6824.326	0.074	122194.71...	0.0201562...	0	263652954...
9	NA	Kaza	District	NA	NA	29622.211...	2	29.622	270.137	0.11	355362.47...	0.5387999...	0	270137409...

Open the layer `Properties` for `admin_6`, go to the `Fields` tab and click on the "pencil" button (`Toggle editing mode`)

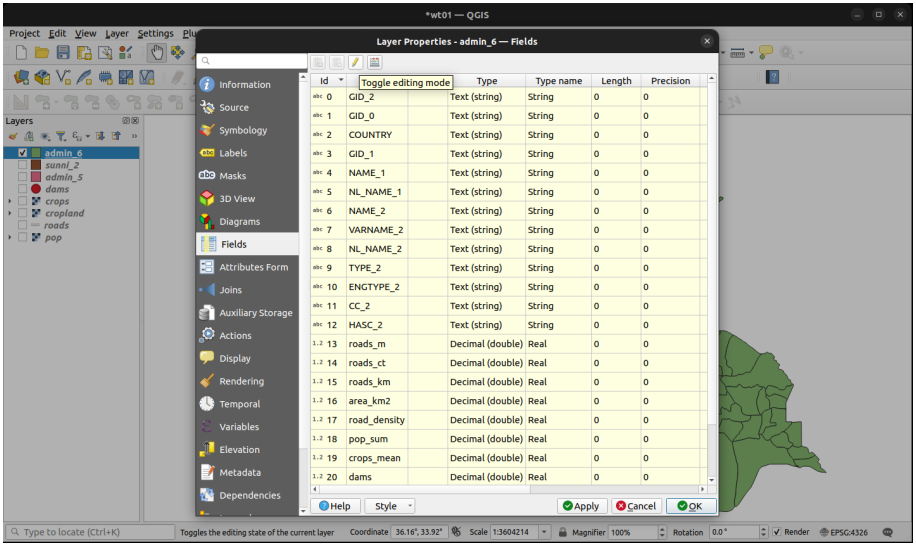

Scroll down to `sunni_2_pc` and double-click on its name

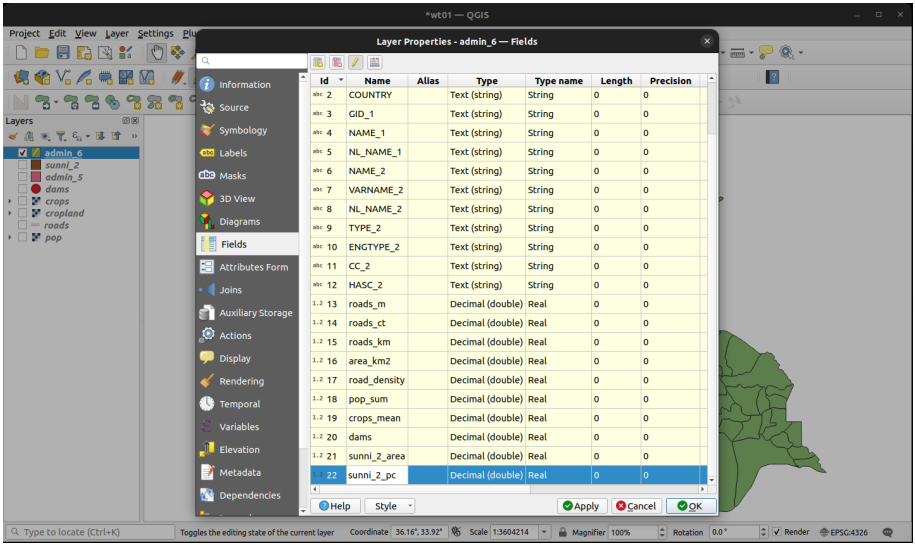

Rename the field `prop_sunni`

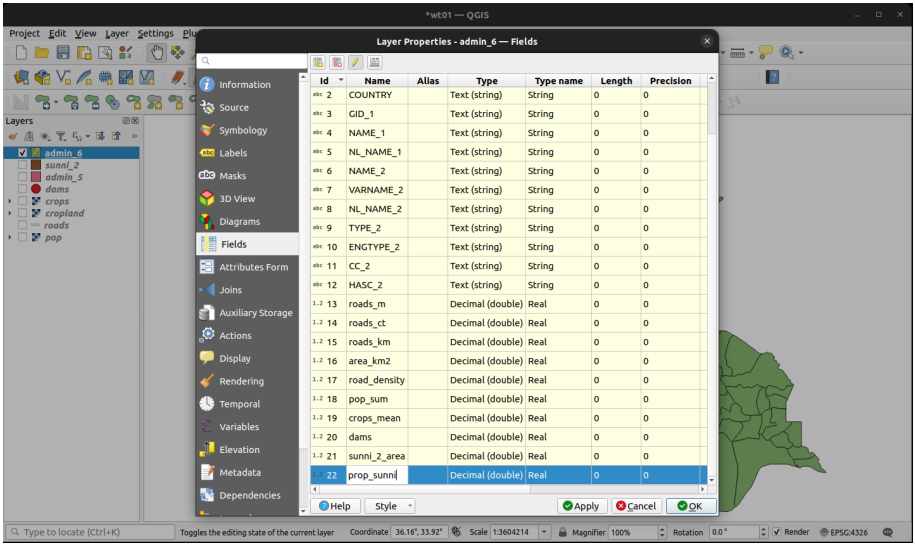

Hit Enter key (or Return key) to commit the name change. Click on the "pencil" again to leave editing mode

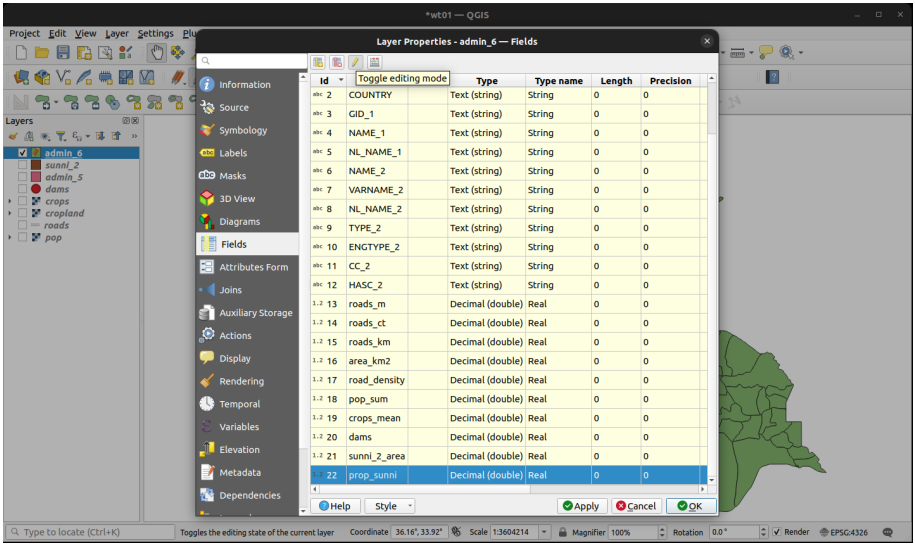

When prompted, save the changes you just made to `admin_6`

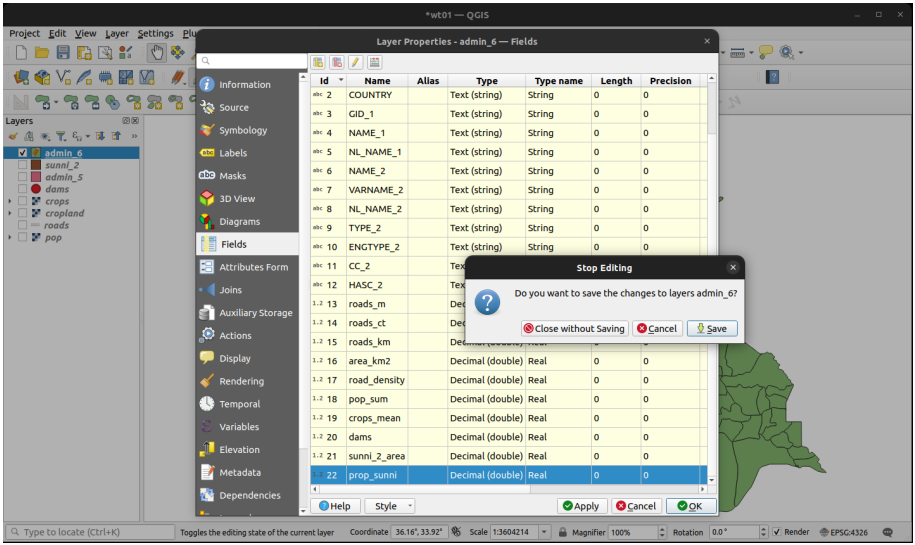

While we're at it, let's create a couple additional fields that could come in handy in the analysis. Create a **population density** field, named `pop_density`, of type `Decimal number (real)`, with `Expression` set to `pop_sum/area_km2`

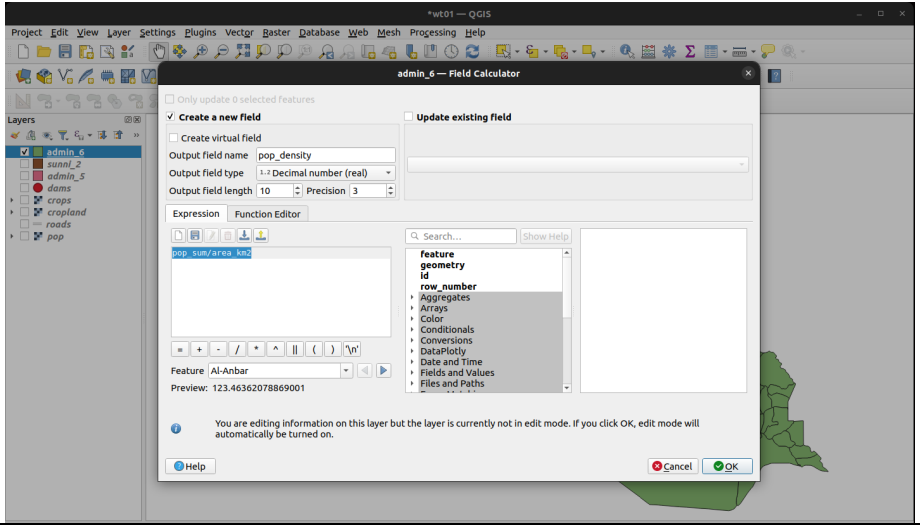

Let's also create a **rescaled population** field (1000's of residents). Name it `pop_1000`, with type `Decimal number (real)`. Set `Expression` to `pop_sum/1000`

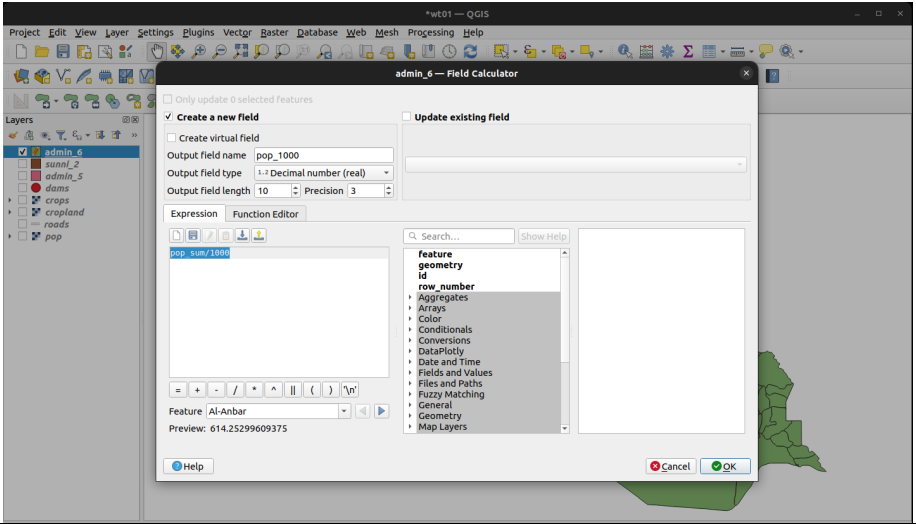

Save the layer edits to `admin_6`!

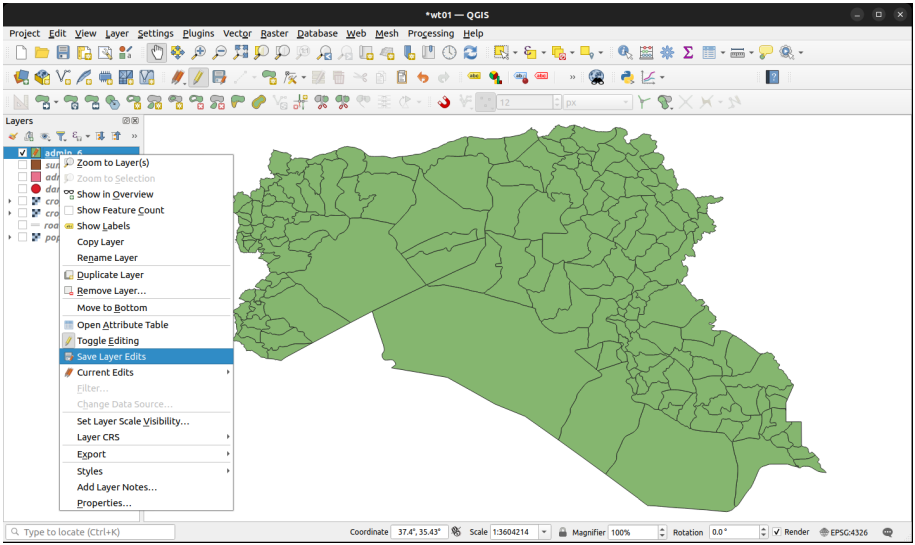

Add the *Islamic State violence data* to the project, using `Add Delimited Text Layer...`. Load the `GEDEvent_v24.csv` file in `Data/GED` folder. Set `X field` = `longitude` and `Y field` = `latitude`. Check the box next to ✓`Use spatial index`

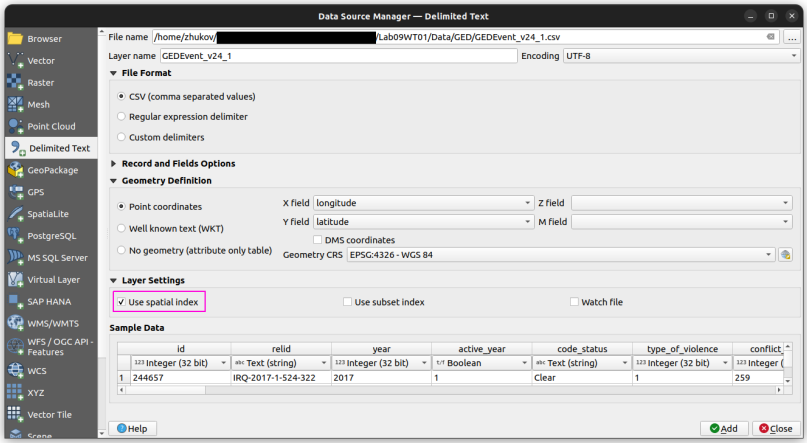

The (global) GED violent events layer should appear. There are several hundred thousand points here. We need to **extract events in Syria and Iraq involving the Islamic State**

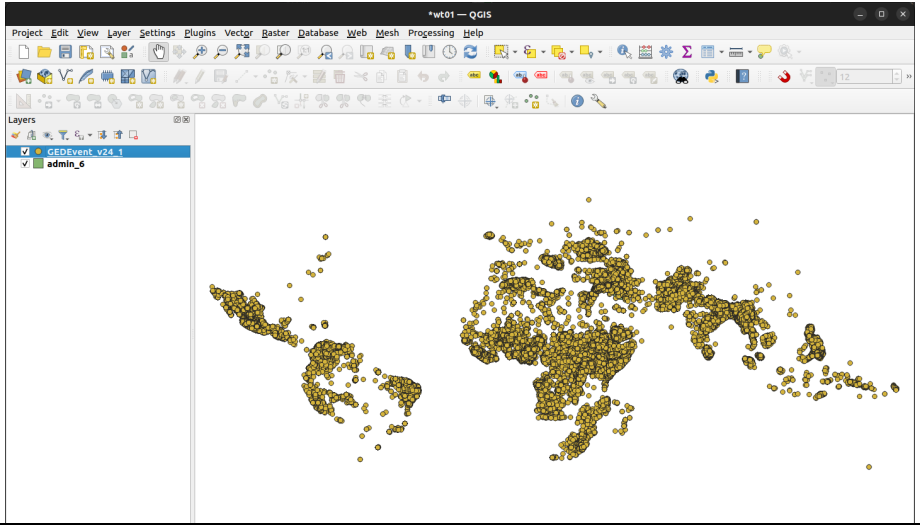

To figure out how to extract this data subset (by actor and location), let's explore the attribute table. We see fields for `side_a` and `side_b`, which list actors

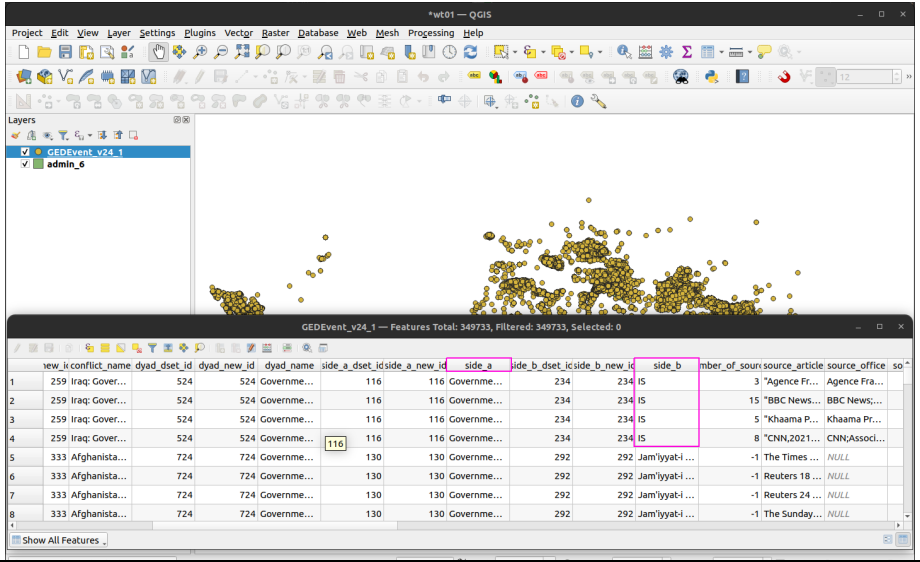

There is also a `country` field. So, we need to select points where `country` is Iraq or Syria and the Islamic State is `side_a` (or `side_b`)

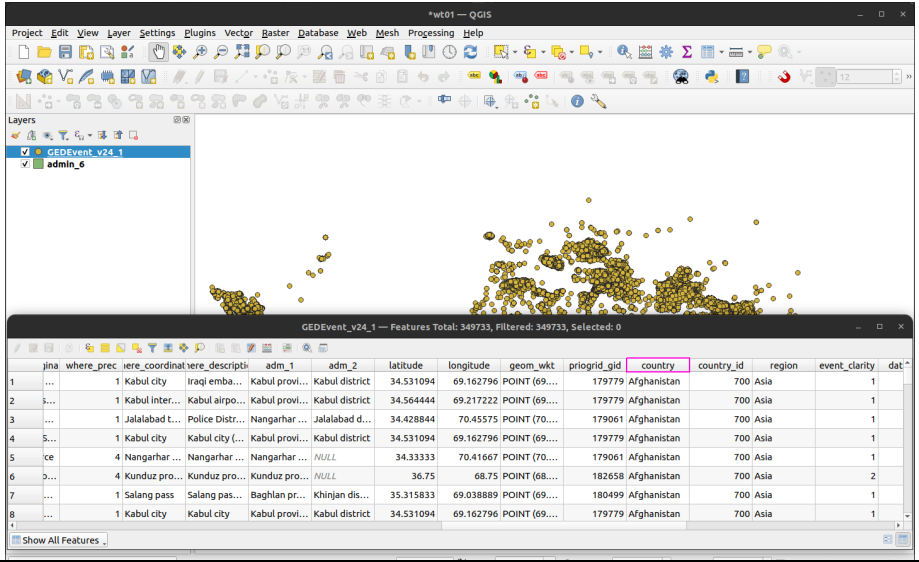

	jina	where_prec	iere_coordinat	iere_descripti	adm_1	adm_2	latitude	longitude	geom_wkt	priogrid_gid	country	country_id	region	event_clarity	dat
1	...	1	Kabul city	Iraqi emba...	Kabul provi...	Kabul district	34.531094	69.162796	POINT (69....	179779	Afghanistan	700	Asia	1	
2	s...	1	Kabul inter...	Kabul airpo...	Kabul provi...	Kabul district	34.564444	69.217222	POINT (69....	179779	Afghanistan	700	Asia	1	
3	...	1	Jalalabad t...	Police Distr...	Nangarhar ...	Jalalabad d...	34.428844	70.45575	POINT (70....	179061	Afghanistan	700	Asia	1	
4	S...	1	Kabul city	Kabul city (...	Kabul provi...	Kabul district	34.531094	69.162796	POINT (69....	179779	Afghanistan	700	Asia	1	
5	ce	4	Nangarhar ...	Nangarhar ...	Nangarhar ...	NULL	34.33333	70.41667	POINT (70....	179061	Afghanistan	700	Asia	1	
6	D...	4	Kunduz pro...	Kunduz pro...	Kunduz pro...	NULL	36.75	68.75	POINT (68....	182658	Afghanistan	700	Asia	2	
7	...	1	Salang pass	Salang pas...	Baghlan pr...	Khinjan dis...	35.315833	69.038889	POINT (69....	180499	Afghanistan	700	Asia	1	
8	...	1	Kabul city	Kabul city	Kabul provi...	Kabul district	34.531094	69.162796	POINT (69....	179779	Afghanistan	700	Asia	1	

Highlight the GED layer and go to `Edit` → `Select` → `Select by Expression...`

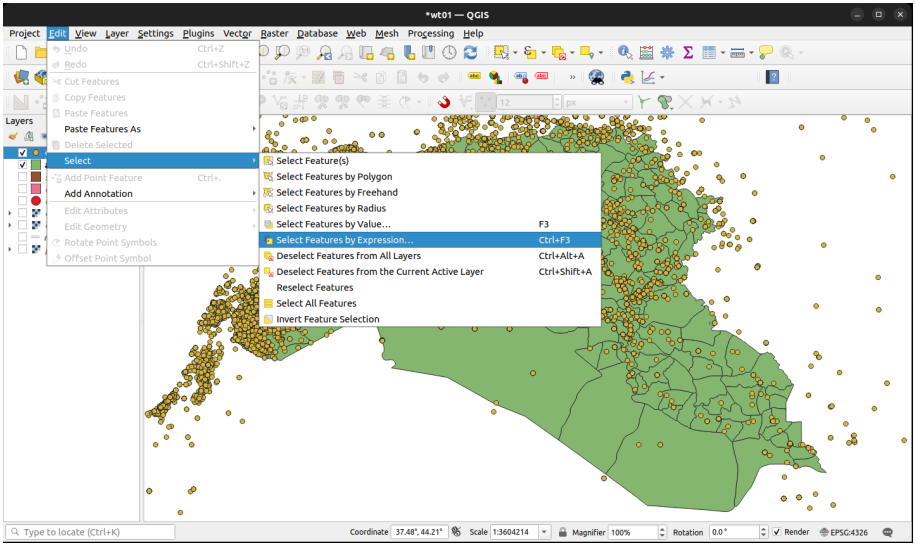

Let's combine regular expressions with logical operators. Set `Expression` to

```
regexp_match(country,'Syria|Iraq') AND (
regexp_match(side_a,'IS|Islamic State') OR
regexp_match(side_b,'IS|Islamic State') )
```

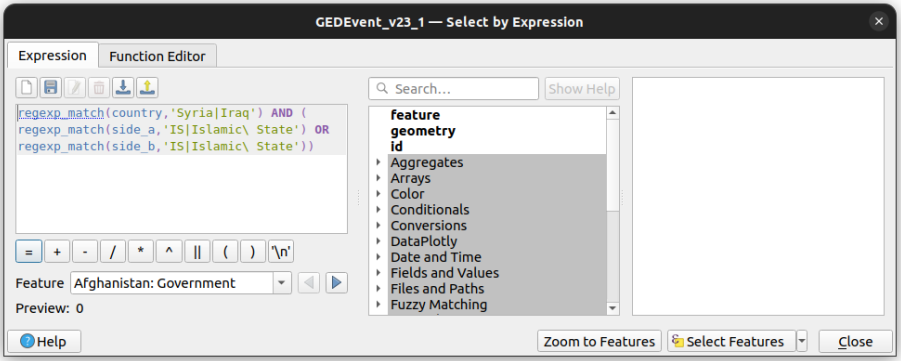

The vertical slash | is a regular expression for "OR". So, in English this expression means 'features where the field `country` contains "Iraq or Syria" and either `side_a` or `side_b` contains "IS or Islamic State" '

This should select a little over 23,000 events. Now let's calculate the **number of ISIS attacks per district**

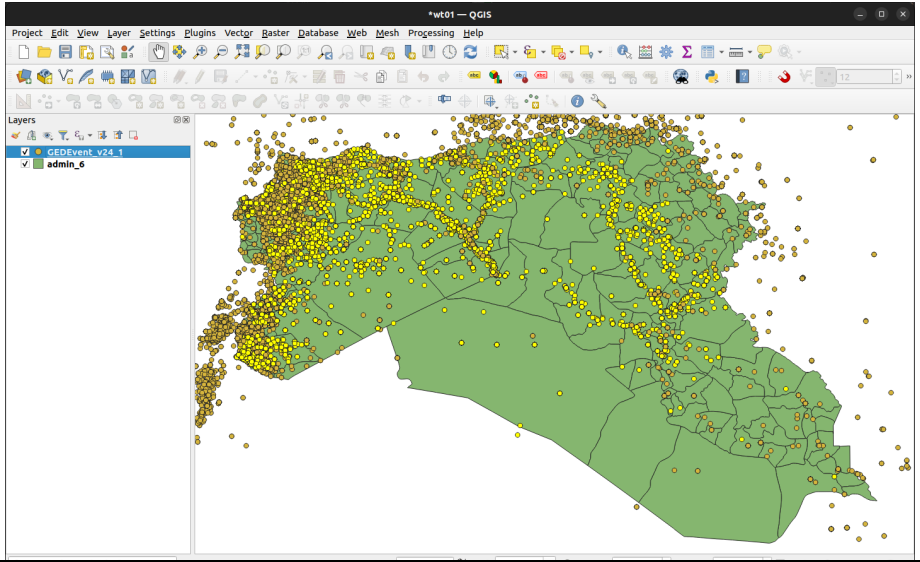

Open the `Count Points in Polygon` tool

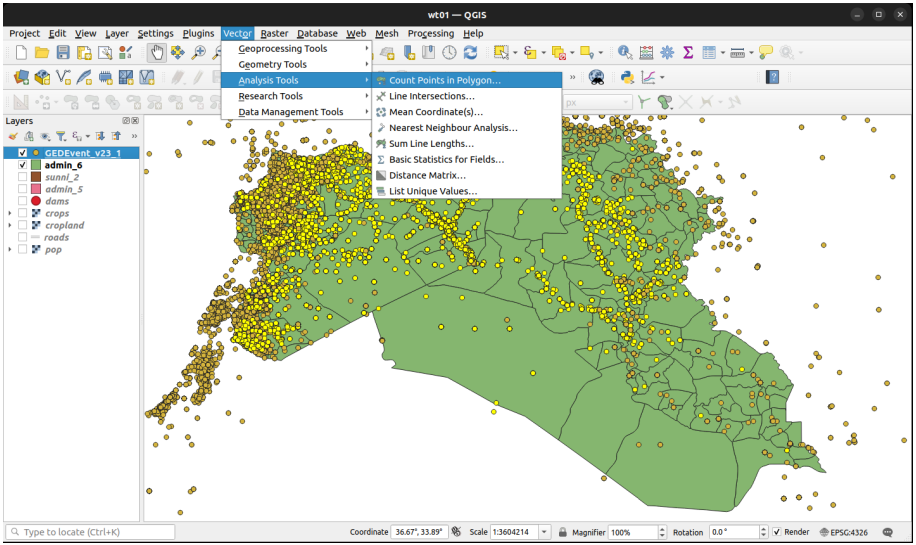

Select `Polygons` = `admin_6`, `Points` = `GEDEvent_v24_1`. Make sure the box is checked next to `Selected Features Only` for the points. Name the count field `is_violence`, and save the output file as `admin_7.geojson`. Click `Run`

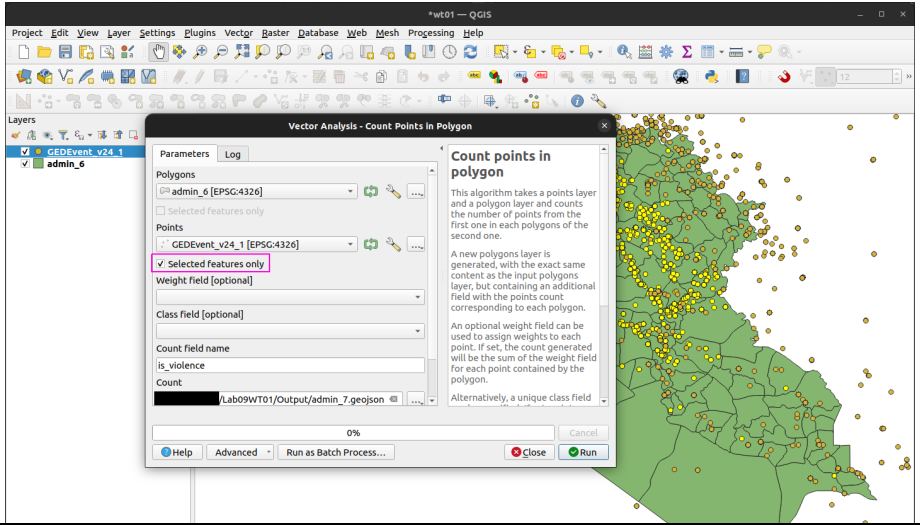

This may take a minute or two to run

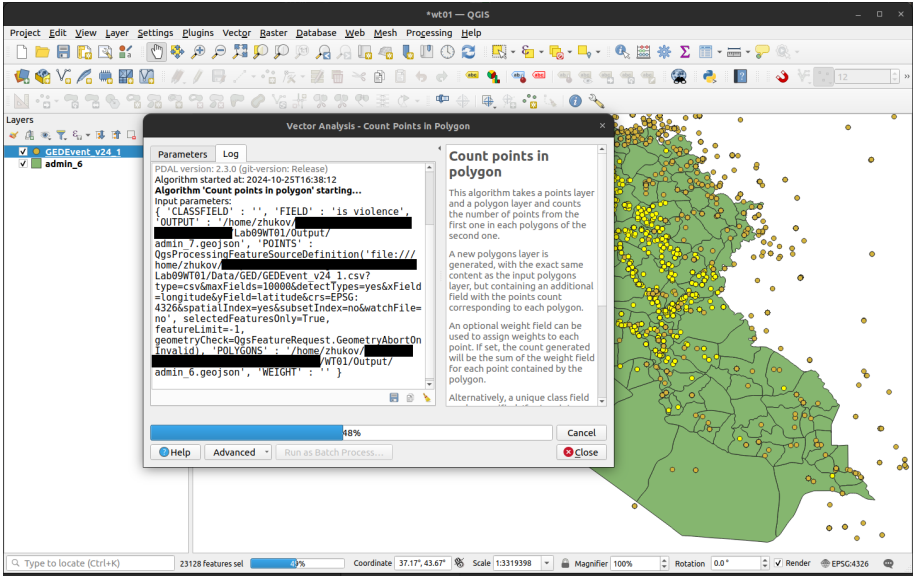

The attribute table for this new layer should include all the new fields we have generated. Let's now **export the table as CSV** for further analysis

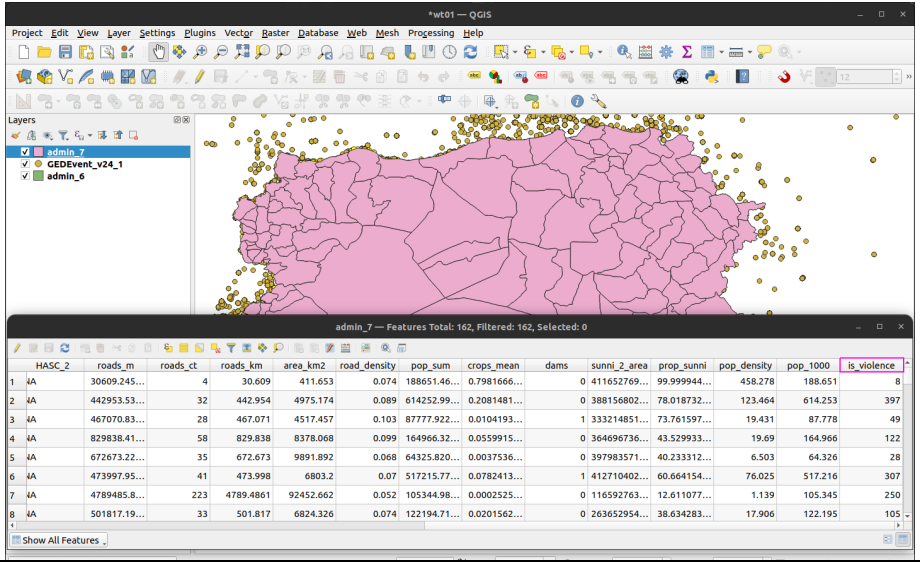

Right-click admin_7 in the Layer menu, go to Export→Save Features As.

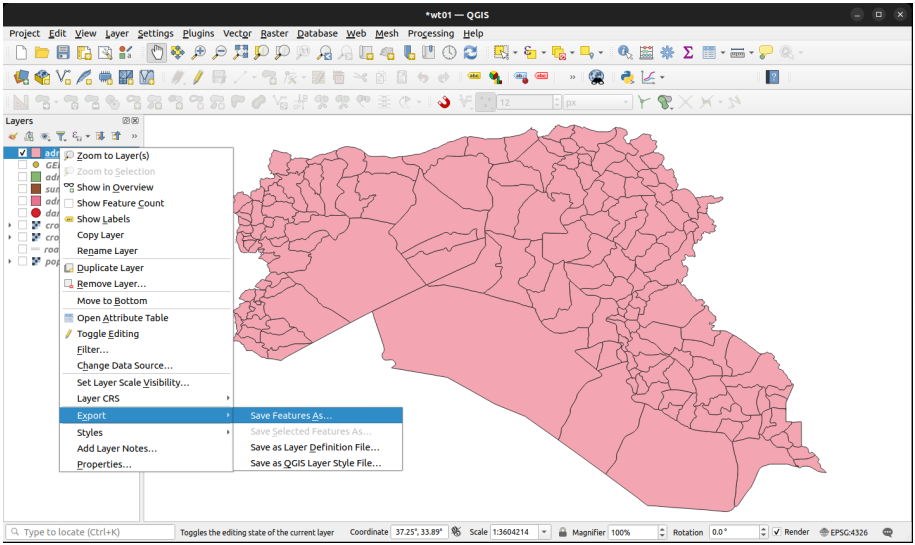

Save the layer as an `Comma Separated Values` file

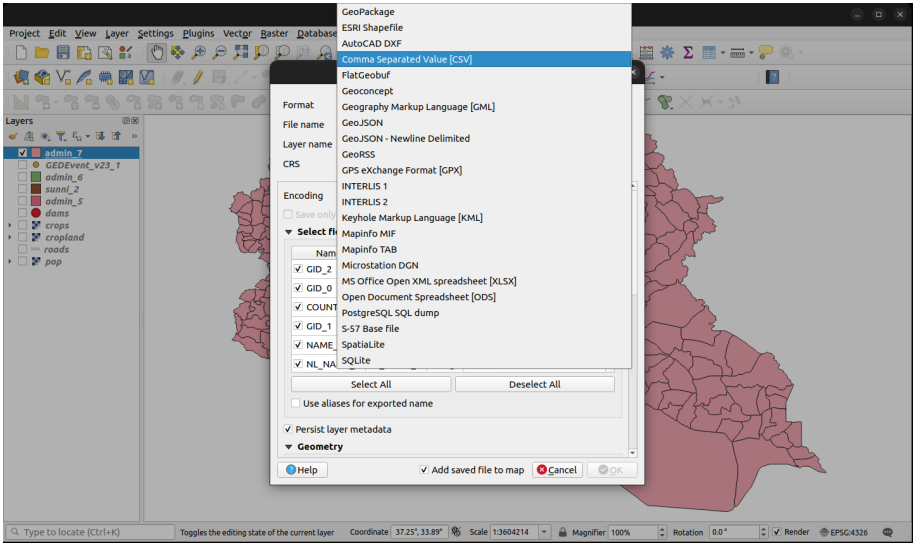

Name the file `final_data_qgis.csv`. Click `OK`

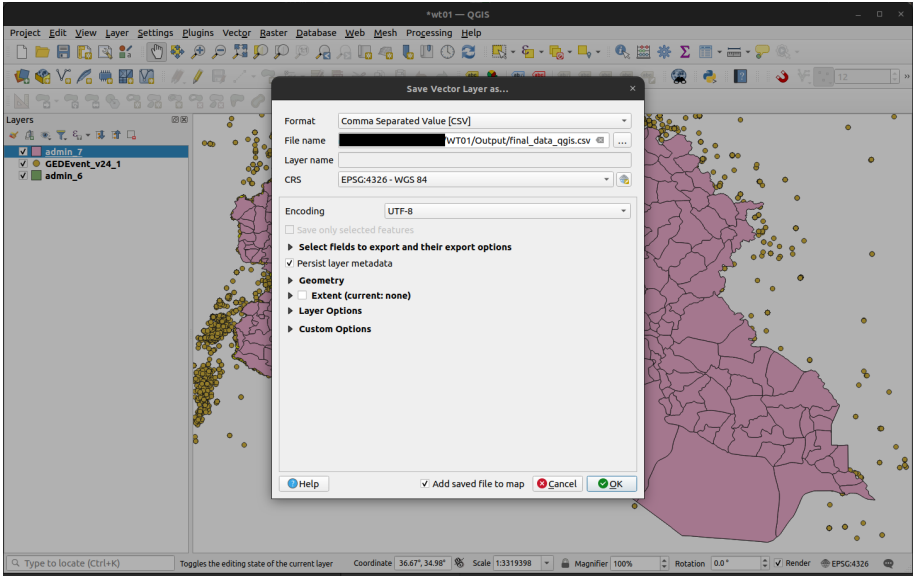

Analysis

Loading R packages

To conduct a regression analysis of these data in R, we will be using the `stargazer` packages

```
library(stargazer)
```

NOTE: The replication code for all of the preceding steps R is in `wt08_demo.R` on RStudio Cloud, and in `WT01.zip` (posted on Canvas).

Regression models

Now we're finally able to proceed to the *analysis stage*. For this we will need to **open the CSV file** we just created in R.

This code chunk imports the `final_data_qgis.csv` file into an object called `X`, and then lists the variable names:

```
X = read.csv("Output/final_data_qgis.csv")
names(X)
```

```
##  [1] "GID_2"        "GID_0"        "COUNTRY"      "GID_1"        "NAME_1"
##  [6] "NL_NAME_1"    "NAME_2"       "VARNAME_2"    "NL_NAME_2"    "TYPE_2"
## [11] "ENGTYPE_2"    "CC_2"         "HASC_2"       "roads_m"      "roads_ct"
## [16] "roads_km"     "area_km2"     "road_density" "pop_sum"      "crops_mean"
## [21] "dams"         "sunni_2_area" "prop_sunni"   "pop_density"  "pop_1000"
## [26] "is_violence"
```

All the variables we created seem to be here.
Let's now *run some regression models!*

Quick refresher: our regression analysis will test 5 hypotheses at once

$$\begin{aligned}\text{violence}_i = &\beta_1 \text{road density}_i + \beta_2 \text{population}_i + \beta_3 \text{cropland}_i \\ &+ \beta_4 \text{dams}_i + \beta_5 \text{Sunni presence}_i + \epsilon_i\end{aligned}$$

where

- violence_i is the observed number of ISIS attacks in district i (`is_violence`)
- road density_i, ..., Sunni presence_i are explanatory variables (`road_density`, `pop_1000`, `crops_mean`, `dams`, `prop_sunni`)
- β are coefficient estimates corresponding to each Hypothesis

Hypothesis	Expectation	Observation
1. Power projection	$\beta_1 < 0$	?
2. Demographics	$\beta_2 > 0$	?
3. Political economy	$\beta_3 < 0$	?
4. Key infrastructure	$\beta_4 > 0$	?
5. Sectarian divisions	$\beta_5 > 0$	?

The first model (`mod1`) is an **Ordinary Least Squares** model that regresses a logarithmically-transformed dependent variable `log(is_violence + 1)` on all of the explanatory variables that correspond to our hypotheses. The log-transform is useful here because the `is_violence` variable is highly skewed

```
mod1 = lm(log(is_violence+1) ~ road_density + pop_1000 + crops_mean +
    dams + prop_sunni, data=X)
summary(mod1)
```

```
##
## Call:
## lm(formula = log(is_violence + 1) ~ road_density + pop_1000 +
##     crops_mean + dams + prop_sunni, data = X)
##
## Residuals:
##     Min      1Q  Median      3Q     Max
## -4.8209 -1.1849 -0.1103  0.9159  3.6005
##
## Coefficients:
##                Estimate Std. Error t value Pr(>|t|)
## (Intercept)   1.2767764  0.3438320   3.713 0.000284 ***
## road_density  1.0261092  3.0441561   0.337 0.736513
## pop_1000      0.0012820  0.0003496   3.667 0.000336 ***
## crops_mean   -1.7270999  0.4814247  -3.587 0.000446 ***
## dams          1.0476508  0.5331035   1.965 0.051168 .
## prop_sunni    0.0348957  0.0030465  11.454  < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

The second (mod2) is a **Generalized Linear Model with a quasi-Poisson link**. This parameterization is designed to accommodate dependent variables that are (over-dispersed) event counts

```
mod2 = glm(is_violence ~ road_density + pop_1000 + crops_mean +
    dams + prop_sunni, data=X, family="quasipoisson")
summary(mod2)
```

```
##
## Call:
## glm(formula = is_violence ~ road_density + pop_1000 + crops_mean +
##     dams + prop_sunni, family = "quasipoisson", data = X)
##
## Coefficients:
##                Estimate Std. Error t value Pr(>|t|)
## (Intercept)   4.1398937  0.4813895   8.600 7.95e-15 ***
## road_density -7.5459241  3.9548121  -1.908   0.0582 .
## pop_1000      0.0010212  0.0002261   4.518 1.23e-05 ***
## crops_mean   -2.7261978  0.6513259  -4.186 4.74e-05 ***
## dams          0.7926585  0.3703707   2.140   0.0339 *
## prop_sunni    0.0287344  0.0047586   6.038 1.10e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for quasipoisson family taken to be 345.4133)
##
##     Null deviance: 69006  on 161  degrees of freedom
## Residual deviance: 33381  on 156  degrees of freedom
## AIC: NA
##
```

The third and fourth models (`mod3`, `mod4`) add province-level **fixed effects** (`NAME_1`), which allow each province (e.g. Anbar) to have a different baseline level of violence

```
mod3 = lm(log(is_violence+1) ~ road_density + pop_1000 + crops_mean +
    dams + prop_sunni + NAME_1, data=X)
mod4 = glm(is_violence ~ road_density + pop_1000 + crops_mean +
    dams + prop_sunni + NAME_1, data=X, family="quasipoisson")
```

`mod3` is OLS, `mod4` is quasi-Poisson

We can use the `stargazer()` command to export the models' coefficient estimates into a formatted table, which you could add to a paper or report:

```
stargazer::stargazer(mod1,mod2,mod3,mod4,type = "text",
          keep = c("road_density","pop_1000","crops_mean",
            "dams","prop_sunni"),
          add.lines = list(c("Province FE","N","N","Y","Y")))
```

```
##
## ================================================================================================
##                                              Dependent variable:
##                          ------------------------------------------------------------------------
##                          log(is_violence + 1)      is_violence      log(is_violence + 1)      is_violence
##                                   OLS           glm: quasipoisson           OLS           glm: quasipoisson
##                                                    link = log                                  link = log
##                                   (1)                (2)                (3)                (4)
## ------------------------------------------------------------------------------------------------
## road_density                      1.026             -7.546*             2.898             -11.922***
##                                  (3.044)            (3.955)            (3.679)             (4.533)
##
## pop_1000                        0.001***           0.001***           0.001***           0.001***
##                                 (0.0003)           (0.0002)           (0.0003)           (0.0002)
##
## crops_mean                      -1.727***          -2.726***          -2.684***          -2.537***
##                                  (0.481)            (0.651)            (0.528)            (0.606)
##
## dams                             1.048*             0.793**            0.805*             0.880***
##                                  (0.533)            (0.370)            (0.459)            (0.298)
##
## prop_sunni                      0.035***           0.029***           0.016***           0.012***
##                                  (0.003)            (0.005)            (0.005)            (0.004)
##
## ------------------------------------------------------------------------------------------------
## Province FE                         N                  N                  Y                  Y
## Observations                      162                162                162                162
## R2                                0.570                                 0.790
## Adjusted R2                       0.556                                 0.730
## Residual Std. Error    1.543 (df = 156)                      1.204 (df = 125)
## F Statistic         41.351*** (df = 5; 156)               13.079*** (df = 36; 125)
## ================================================================================================
## Note:                                                                *p<0.1; **p<0.05; ***p<0.01
```

What does this tell us about whether the data support our Hypotheses?

Hypothesis	Expectation	Confirm? (OLS)	Confirm? (QP)
1. Power projection	$\beta_{\texttt{road\_density}} < 0$	✗	✓
2. Demographics	$\beta_{\texttt{pop\_1000}} > 0$	✓	✓
3. Political economy	$\beta_{\texttt{crops\_mean}} < 0$	✓	✓
4. Key infrastructure	$\beta_{\texttt{dams}} > 0$	✓	✓
5. Sectarian divisions	$\beta_{\texttt{prop\_sunni}} > 0$	✓	✓